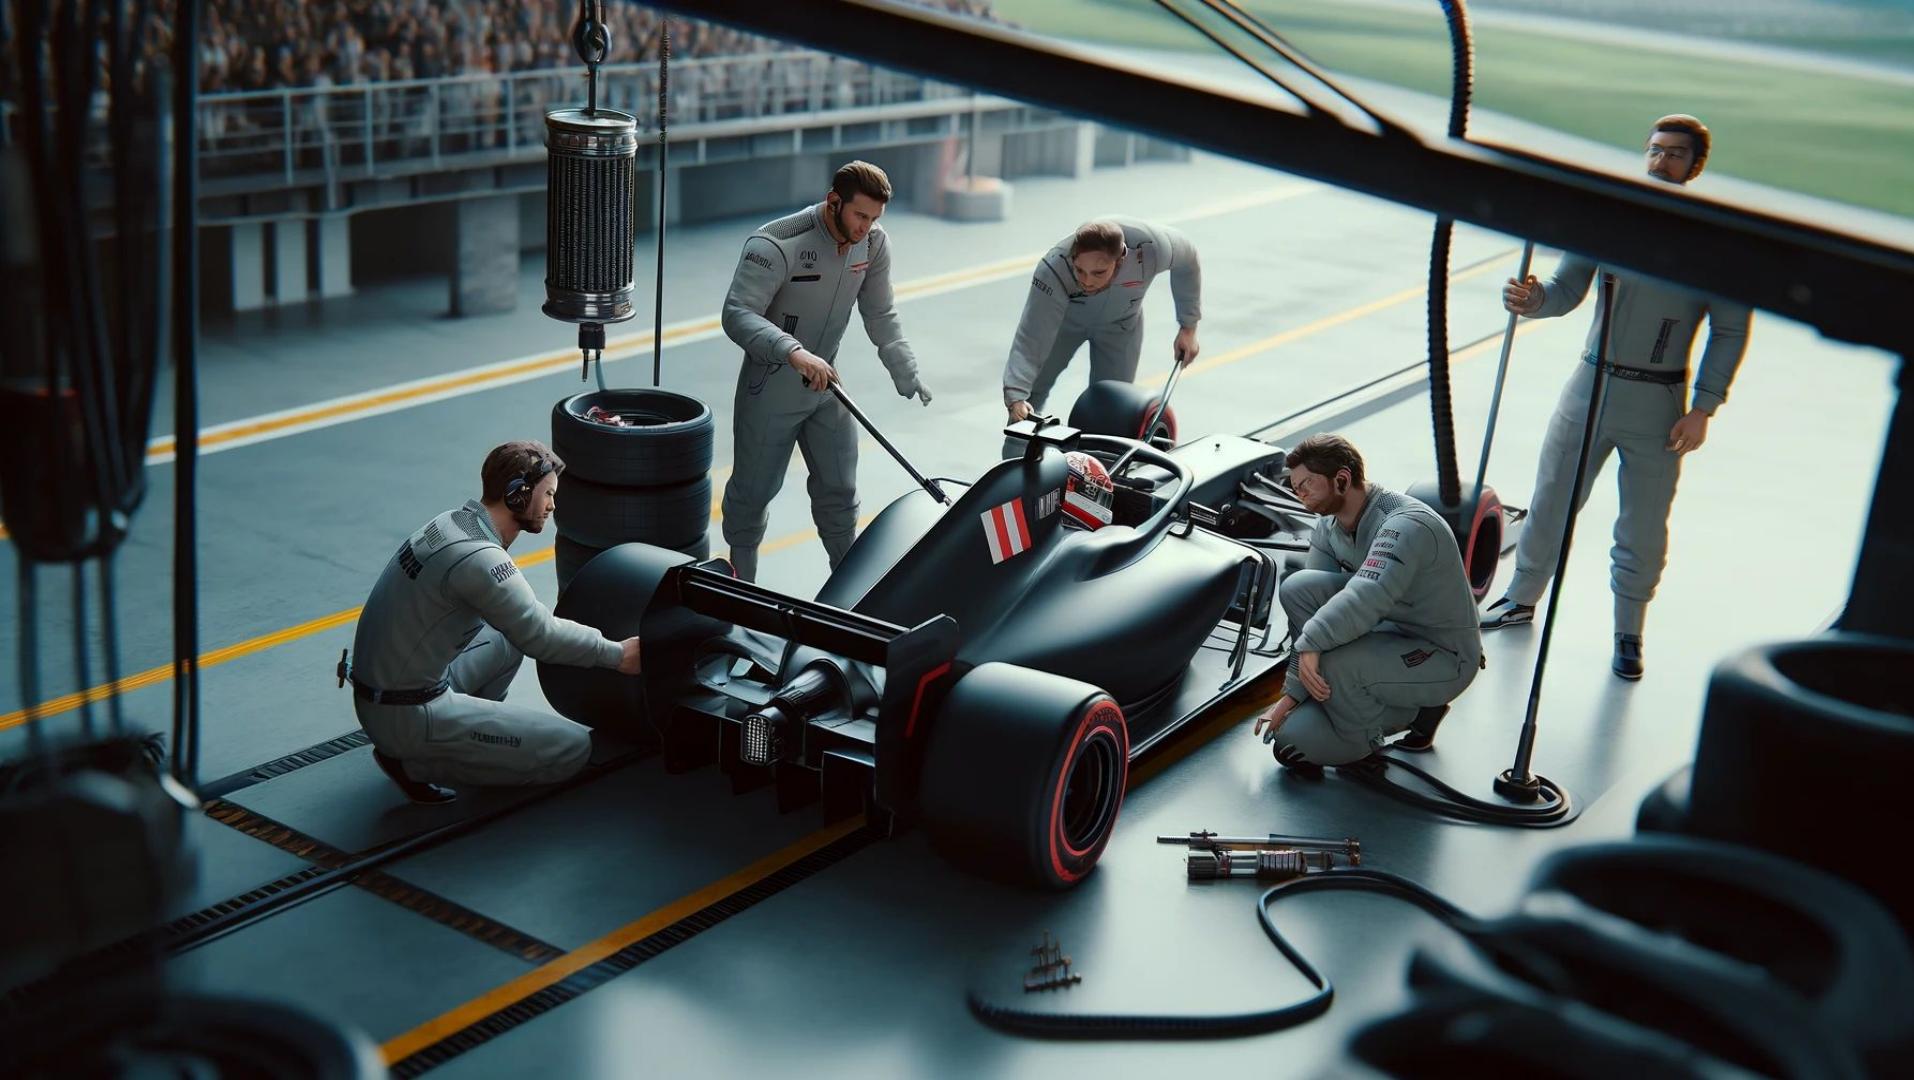


World's Best Racing Driver





**Would You Let Josh
Plan and Build the Engine ?**



Bridging the DevOps to Data Divide with a Common Cloud Native Stack

Elad Hirsch & Mey Beisaron



Hello!

I am **Elad Hirsch**

- Software Architect for over ten years
- Tech Lead , CTO Office @ **TERASKY**
- Public Cloud & Cloud-Native
- Tech Advocate



Hello!

I am **Mey Beisaron**

- Backend Engineer for over five years
- Senior DevOps Engineer at **FORTER**
- Clojure, Nodejs, Groovy, Python
- International public speaker & mentor
- Sworn Star Wars fan





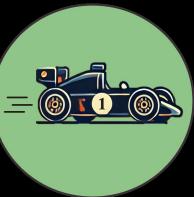
Provision, Set, Go!

Running DBs on K8s with CloudNativePG



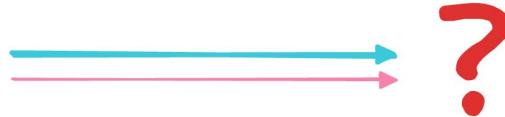
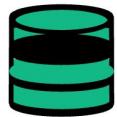
Shift left, Stay Ahead

Manage DB Schema as an Infrastructure



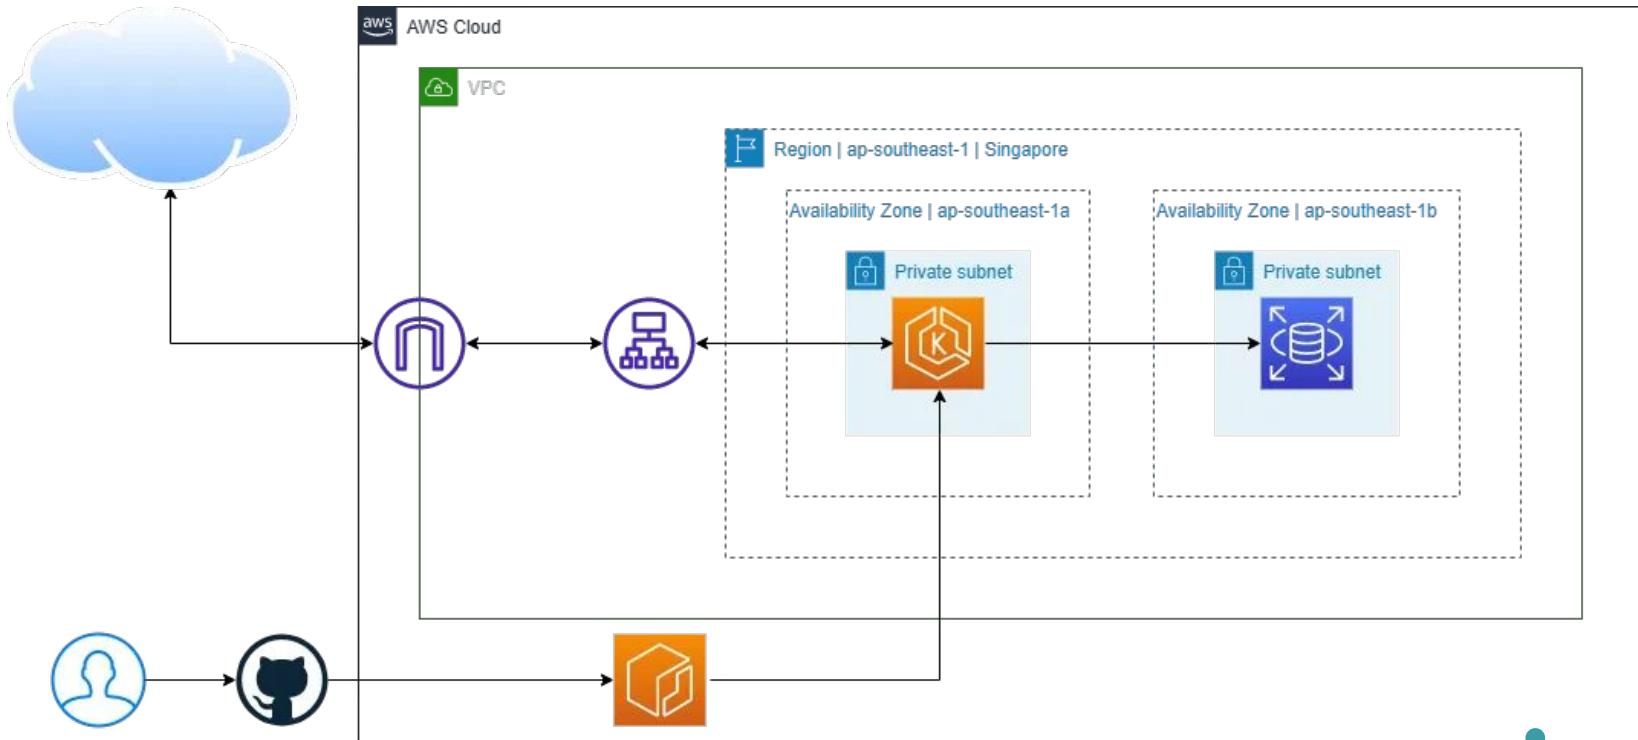
Fast Lane Monitoring

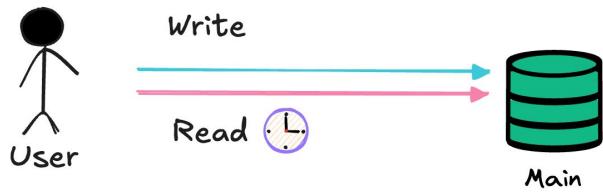
Adding visibility for DBs on K8s



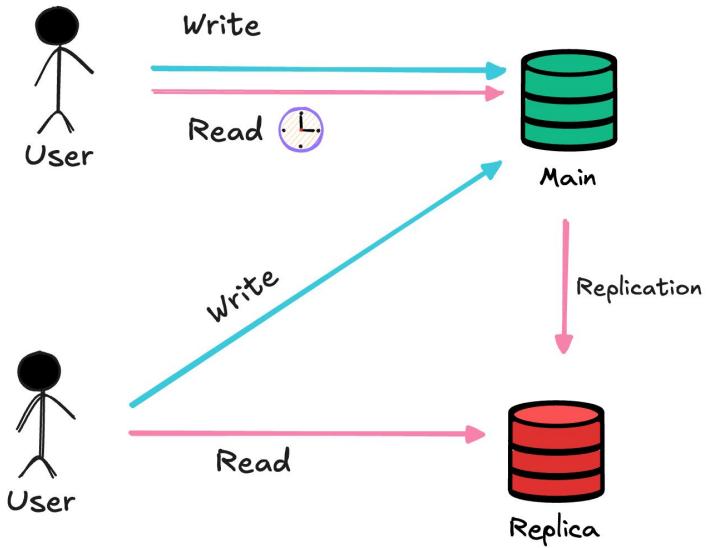
So You Need Good
Database Reliability ...

You went with a managed DB, right?

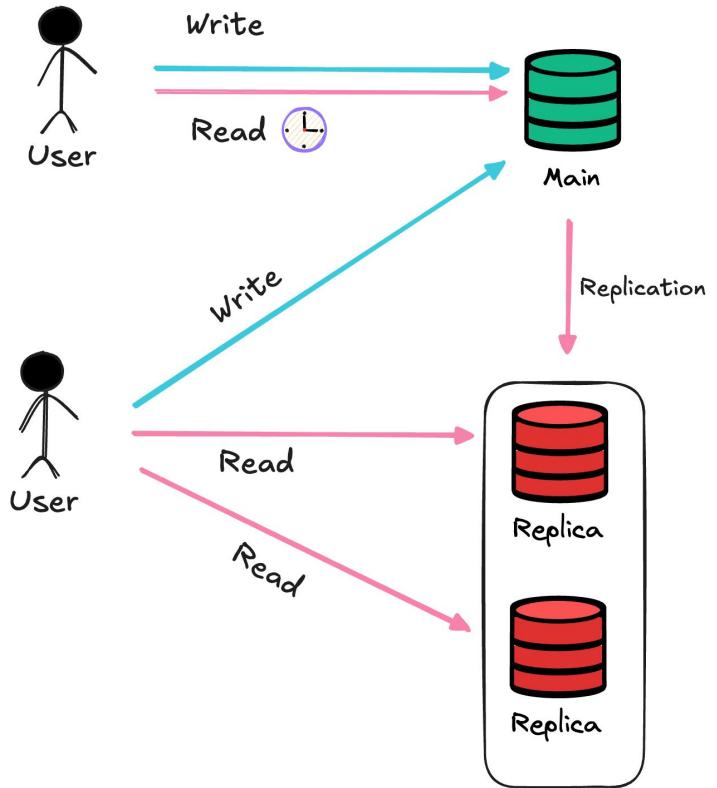




We'll need more than
one database
instance, right?

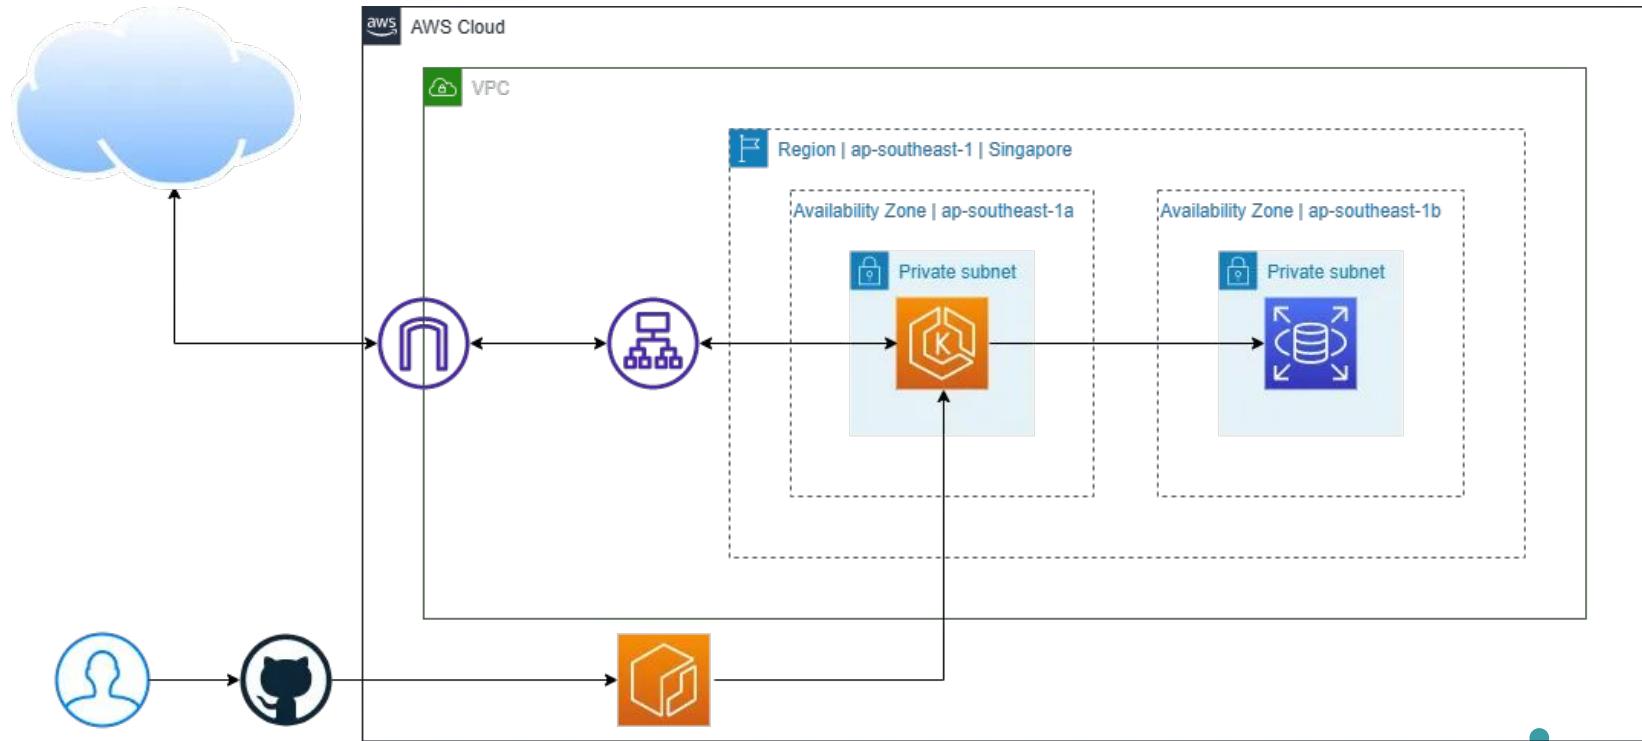


We'll need more than
one database
instance, right?

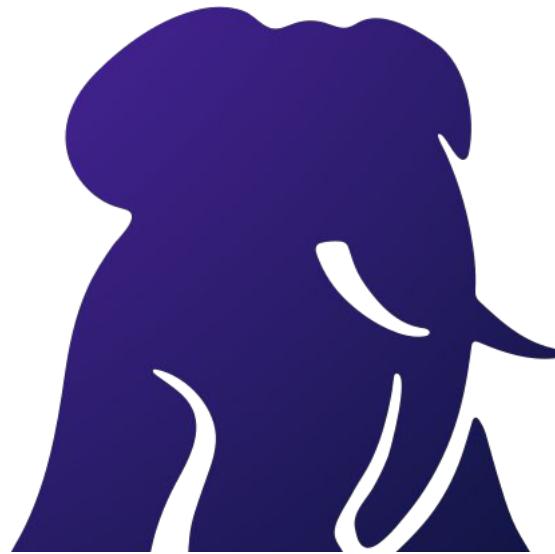


Do we need
multiple AZ's for
high availability?

So it has to be managed DB, right?



CloudNativePG Operator



—
Open Source and Vendor-Neutral

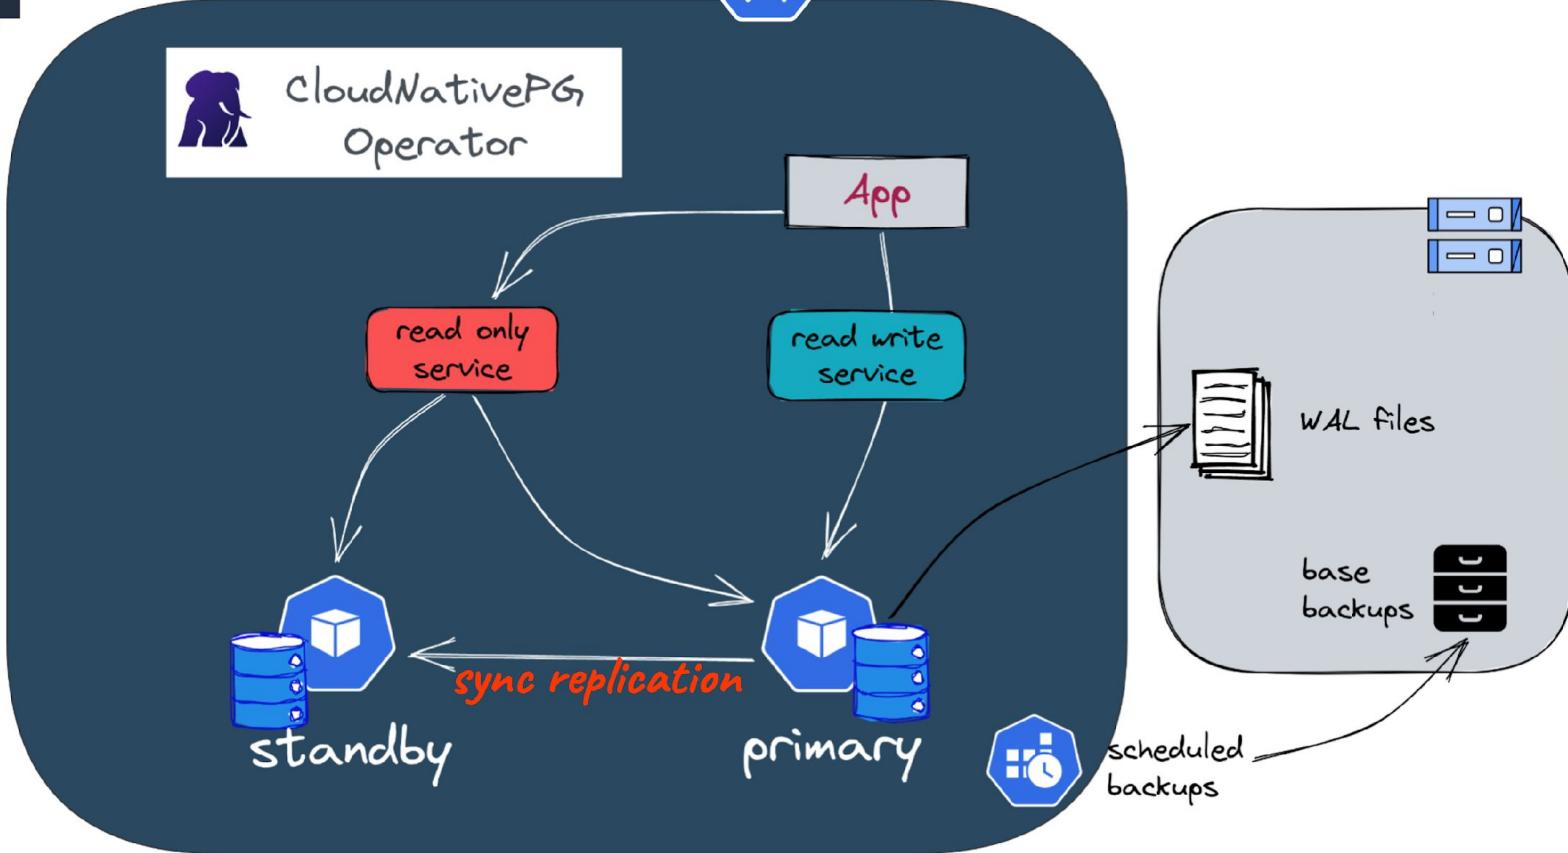
—
Declarative, Immutable Design

—
No Patroni and No StatefulSets

—
HA with automated Failover



My K8s Cluster



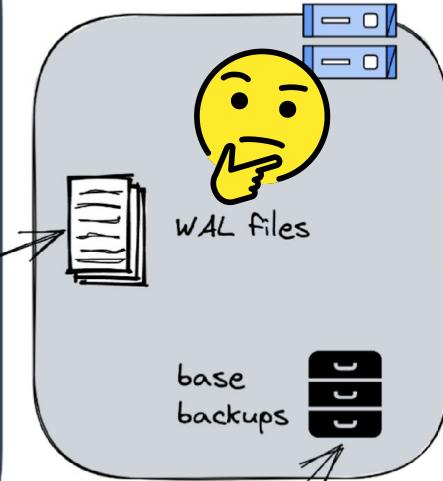
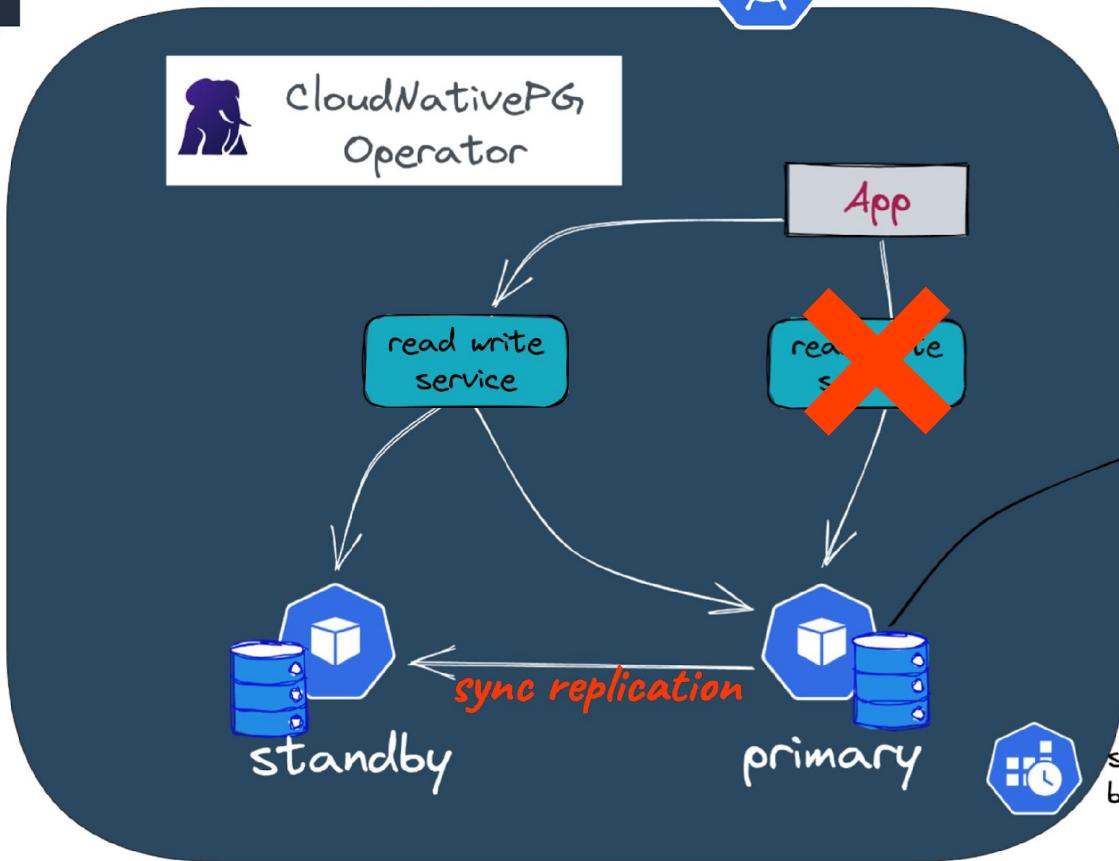


```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: cluster-example
spec:
  instances: 3
  minSyncReplicas: 1
  maxSyncReplicas: 1
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: topology.kubernetes.io/zone
                operator: In
                values:
                  - us-east-1a
                  - us-east-1b
                  - us-east-1c
```



```
tolerations:
  - key: node-role.kubernetes.io/postgres
    operator: Exists
    effect: NoSchedule
topologySpreadConstraints:
  - maxSkew: 1
    topologyKey: topology.kubernetes.io/zone
    whenUnsatisfiable: DoNotSchedule
    labelSelector:
      matchLabels:
        cnpq.io/cluster: cluster-example
storage:
  size: 40Gi
walStorage:
  size: 10Gi
```

My K8s Cluster





```
postgres@5d71369c69ff:~/data$ ls -l
total 132
drwx----- 2 postgres postgres 4096 Sep 30 17:24 pg_tblspc
drwx----- 2 postgres postgres 4096 Sep 30 17:24 pg_twophase
-rw----- 1 postgres postgres      3 Sep 30 17:24 PG_VERSION
drwx----- 4 postgres postgres 4096 Oct 14 17:00 pg_wal
drwx----- 2 postgres postgres 4096 Sep 30 17:24 pg_xact
-rw----- 1 postgres postgres     88 Sep 30 17:24 postgresql.auto.conf
-rw----- 1 postgres postgres 30827 Oct 16 12:27 postgresql.conf
-rw----- 1 postgres postgres      36 Oct 17 15:39 postmaster.opts
-rw----- 1 postgres postgres      94 Oct 17 15:39 postmaster.pid
postgres@5d71369c69ff:~/data$
```

Boosting Data Reliability and Transaction Efficiency



```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: cluster-example
spec:
  storage:
    size: 100Gi
  walStorage:
    storageClass: wal-fast-storage
    size: 10Gi
```



```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: wal-fast-storage
provisioner: kubernetes.io/aws-ebs
parameters:
  type: io2          # High-performance SSD for WAL
  iopsPerGB: "100"   # Increase IOPS;
                     # Suitable for heavy-write workloads
  fsType: ext4
  encrypted: "true"
reclaimPolicy: Delete
allowVolumeExpansion: true
volumeBindingMode: WaitForFirstConsumer
```



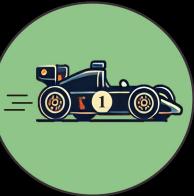
Provision, Set, Go!

Running DBs on K8s with CloudNativePG



Shift left, Stay Ahead

Manage DB Schema as an Infrastructure

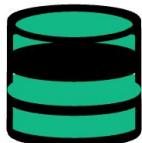


Fast Lane Monitoring

Adding visibility for DBs on K8s



alter table



Main



So You Need a Reliable
Way to Manage
Schema Changes?

-
-
- Schema update during startup
-



Schema update in *InitContainers*

```
// schema.prisma

model User {
    id    Int      @id
    @default(autoincrement())
    name  String
    email String  @unique
    bio   String?
}
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
spec:
  template:
    spec:
      containers:
        # Main application container

      initContainers:
        - name: prisma-migrate
          image: node:18
          command: ["npx", "prisma",
                    "migrate", "dev",
                    "--name", "change_set"]

      env:
        - name: DATABASE_URL
          valueFrom:
            secretKeyRef:
              name: my-database-secret
              key: DATABASE_URL
      volumeMounts:
        - name: app-source
          mountPath: /app

      volumes:
        - name: app-source
          emptyDir: {}
```



```
apiVersion: batch/v1
kind: Job
metadata:
  name: liquibase-migration-job
spec:
  template:
    spec:
      containers:
        - name: liquibase
          image: liquibase/liquibase:latest
          env:
          args:
            - "--changeLogFile=/changelog/change300.xml"
            - "update"
          volumeMounts:
            - name: changelog-volume
              mountPath: /liquibase/changelog
      restartPolicy: OnFailure
      volumes:
        - name: changelog-volume
      configMap:
        name: liquibase-changelog-config
```



Schema Update Managed with Kubernetes Job



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: liquibase-changelog-config
data:
  changelog.xml: |
    <databaseChangeLog>
      <changeSet id="300" author="developer">
        <dropColumn tableName="book" columnName="name"/>
      </changeSet>
    </databaseChangeLog>
```



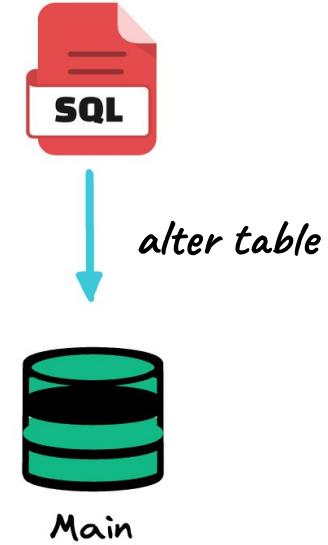
Back to Square One ?



Am I Clear on My Schema Changes and how it's going to R ?

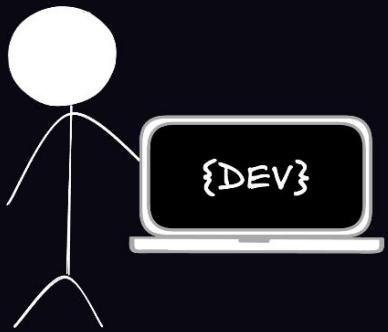
Can I Simply Declare My Evolving Schema's as a Target State ?

How Do I Ensure My Schema Stays Consistent with My Target State ?

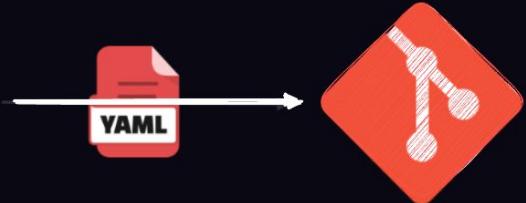




Cloud Native Operator

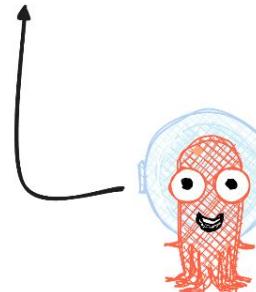


```
apiVersion: db.atlasgo.io/v1alpha1
kind: AtlasSchema
metadata:
  name: myapp-schema
spec:
  schema:
    sql: |
      create table users (
        id int not null auto_increment,
        name varchar(255) not null,
        email varchar(255) unique not null,
        primary key (id)
      );
```





```
kind: AtlasSchema  
metadata:  
  name: myapp-schema  
spec:  
  schema:  
    sql: |  
      create table users (  
        id int not null auto_increment,  
        name varchar(255) not null,  
        email varchar(255) unique not null,  
        primary key (id)  
      );
```



inspect
current
schema



Atlas Operator

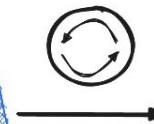
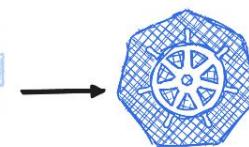
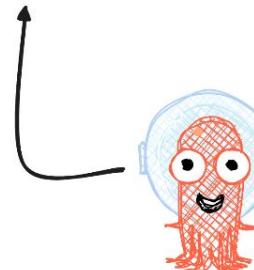


insert into ...



```
kind: AtlasSchema
metadata:
  name: myapp-schema
spec:
  schema:
    sql: |
      create table users (
        id int not null auto_increment,
        name varchar(255) not null,
        email varchar(255) unique not null,
        age int not null,
        primary key (id)
      );

```

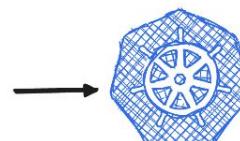
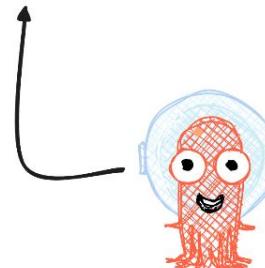


inspect
current
schema





```
kind: AtlasSchema  
metadata:  
  name: myapp-schema  
spec:  
  schema:  
    sql: |  
      create table users (  
        id int not null auto_increment,  
        name varchar(255) not null,  
        email varchar(255) unique not null,  
        age int not null,  
        primary key (id)  
      );
```



inspect
current
schema

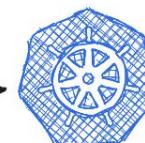
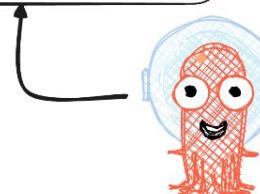
Atlas Operator

alter table...





```
kind: AtlasSchema
metadata:
  name: myapp-schema
spec:
  policy:
    lint:
      destructive:
        error: true
  schema:
    sql: |
      create table users (
        id int not null auto_increment,
        name varchar(255) not null,
        age int not null,
        primary key (id)
      );
```



inspect
current
schema





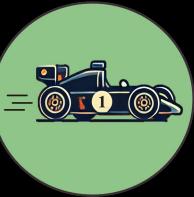
Provision, Set, Go!

Running DBs on K8s with CloudNativePG



Shift left, Stay Ahead

Manage DB Schema as an Infrastructure



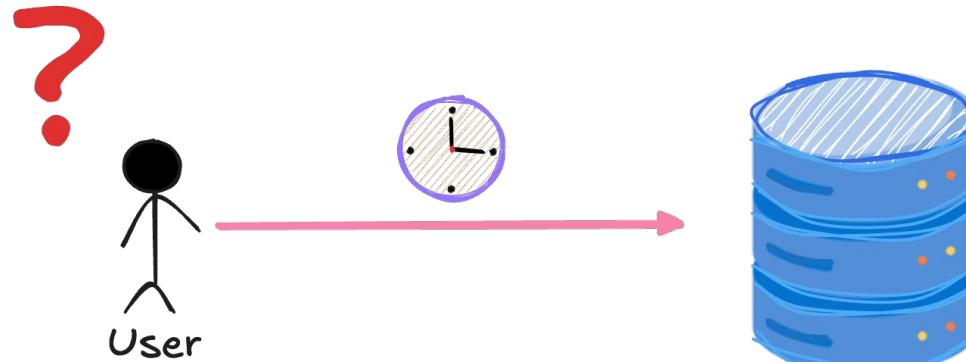
Fast Lane Monitoring

Adding visibility for DBs on K8s



So,

You've Added More Instances
but the Query is Still Slow—Why?





Golden Path



Enable Comprehensive Monitoring



```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: example-cluster
  namespace: monitoring
spec:
  instances: 1
  storage:
    size: 10Gi
  monitoring:
    enablePodMonitor: true
```



Enable Comprehensive Monitoring



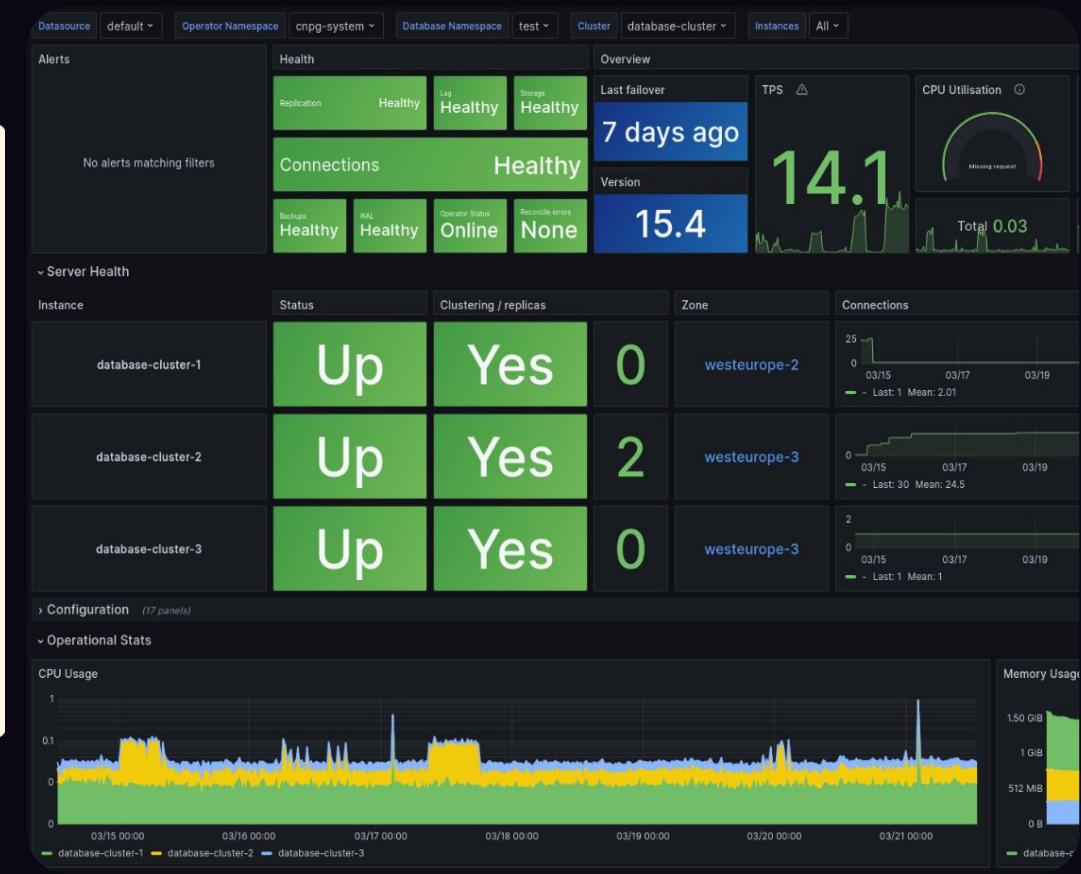
```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: example-cluster
  namespace: monitoring
spec:
  instances: 1
  storage:
    size: 10Gi
  monitoring:
    enablePodMonitor: true
```



Enable Comprehensive Monitoring



```
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata:
  name: example-cluster
  namespace: monitoring
spec:
  instances: 1
  storage:
    size: 10Gi
  monitoring:
    enablePodMonitor: true
```



Golden Path and Support for Super Users



User Defined Metrics for Advanced Use Cases



```
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata:
  name: example-cluster
  namespace: monitoring
spec:
  instances: 1
  storage:
    size: 10Gi
  monitoring:
    enablePodMonitor: true
    customQueriesConfigMap:
      - name: active-connections-metric
        key: custom-queries
```



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: active-connections-metric
  namespace: monitoring
data:
  custom-queries: |
    active_connections_query:
      query: |
        SELECT
          current_database() AS datname,
          state,
          COUNT(*) AS active_connections
        FROM pg_stat_activity
        WHERE state = 'active'
        GROUP BY state, datname;
  metrics:
    - datname:
        usage: "LABEL"
    - state:
        usage: "LABEL"
    - active_connections:
        usage: "GAUGE"
```

User Defined Metrics for Advanced Use Cases

```
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata:
  name: example-cluster
  namespace: monitoring
spec:
  instances: 1
  storage:
    size: 10Gi
  monitoring:
    enablePodMonitor: true
    customQueriesConfigMap:
      - name: active-connections-metric
        key: custom-queries
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: active-connections-metric
  namespace: monitoring
data:
  custom-queries: |
    active_connections_query:
      query: |
        SELECT
          current_database() AS datname,
          state,
          COUNT(*) AS active_connections
        FROM pg_stat_activity
        WHERE state = 'active'
        GROUP BY state, datname;
  metrics:
    - datname:
        usage: "LABEL"
    - state:
        usage: "LABEL"
    - active_connections:
        usage: "GAUGE"
```

User Defined Metrics for Advanced Use Cases



```
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata:
  name: example-cluster
  namespace: monitoring
spec:
  instances: 1
  storage:
    size: 10Gi
  monitoring:
    enablePodMonitor: true
    customQueriesConfigMap:
      - name: active-connections-metric
        key: custom-queries
```



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: active-connections-metric
  namespace: monitoring
data:
  custom-queries: |
    active_connections_query:
      query: |
        SELECT
          current_database() AS datname,
          state,
          COUNT(*) AS active_connections
        FROM pg_stat_activity
        WHERE state = 'active'
        GROUP BY state, datname;
metrics:
  - datname:
      usage: "LABEL"
  - state:
      usage: "LABEL"
  - active_connections:
      usage: "GAUGE"
```

Key Takeaways



Consider Cloud Native Databases

Shift Left and Treat Schema as Infra

No One-Size-Fits-All Solution

Bridging the Data and DevOps Gap

THANKS!

Any questions?

