



Observability Day
NORTH AMERICA

PROMETHEUS

QUERIES THAT MEASURE CPU AND MEMORY

**⚡ Lightning Talk: Vote on Your
Favorite--Standardize This!--10 PromQL Queries That
Measure CPU & Memory Usage**

AGENDA

01

INTRO

Me and why I
love hate
PromQL

02


WHAT WHY WHERE WHEN

The what, why,
when, where of
so many CPU
and Memory
Queries

03

VOTE NOW!

PICK your
favorite!


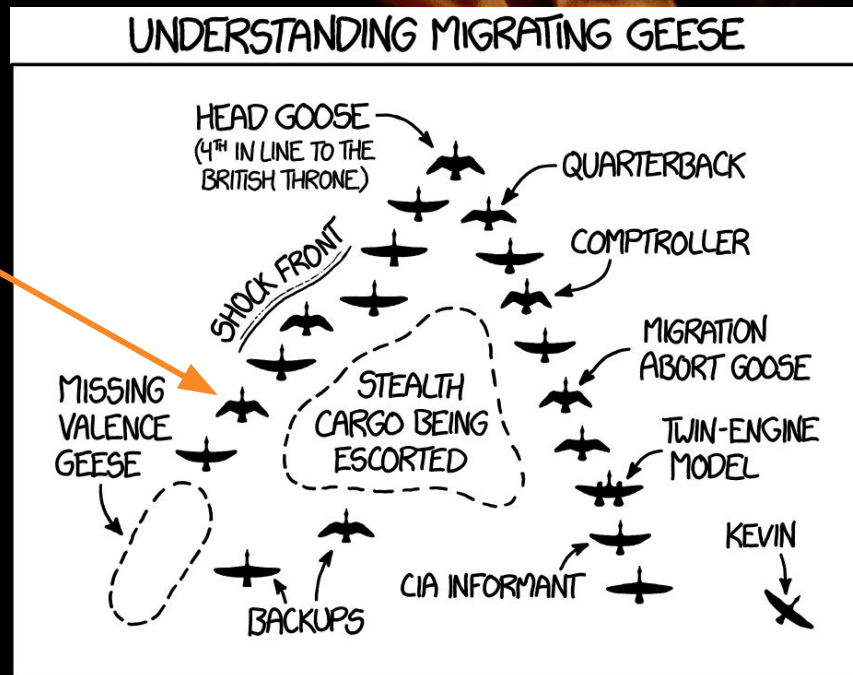


ME


Solutions Architect Migration Work to PROMQL

- SignalFX to Prom
- Wavefront to Prom
- DogstatsD to Prom
- NRQL to Prom

So many Queries




chronosphere



THE WHY

MIGRATIONS

Change is hard



THE WHERE

AI GENERATED EXPLANATION

"When a database and a coder love each other very much..."

THREE MAIN SOURCES OF QUERIES



CUSTOMER TRANSLATIONS

```
(sum by (kubernetes_cluster) (
  rate(
    container_cpu_usage_seconds_total{container="proxy",kubernetes_cluster=~"prod-.*|dev|staging"} [30s]
  ) *100))
/
(sum by (kubernetes_cluster) (
  kube_pod_container_resource_requests{container="proxy",kubernetes_cluster=~"prod-.*|dev|staging",resource="cpu"}
))
```



INTERNAL DASHBOARDS AND MONITORS

```
avg by (container, environment, az) (
  rate(container_cpu_usage_seconds_total{container=~".*",instance=~".*",pod=~".*"} [5m])
)
```



THE INTERNET

```
rate(container_cpu_usage_seconds_total{container!="",namespace="my-namespace",pod=~"my-pod.*"} [15m])
```



THE WHAT

QUERIES

CUSTOMER TRANSLATION QUERIES

SIGNALFX/WAVEFRONT

1. `(sum by (kubernetes_cluster)
(rate(container_cpu_usage_seconds_total{container="proxy",kubernetes_cluster=~"prod-.*|dev|staging"}[30s])
* 100))/ (sum by (kubernetes_cluster)
(kube_pod_container_resource_requests{kubernetes_cluster=~"prod-.*|dev|staging",container="proxy",resource
="cpu"})) % Percentage`
2. `max_over_time(kube_node_status_allocatable{resource="cpu",kubernetes_cluster=~".*prod.*"}[1h]) # of Cores`
3. `count(sum by (pod)
(avg_over_time(container_cpu_usage_seconds_total{kubernetes_cluster=~".*prod.*"}[$interval]))) Pod Count`
4. `sum(rate(container_cpu_usage_seconds_total{namespace=~"$namespace",instance=~"$node",kubernetes_cluster=~"
$kubernetes_cluster",container=~"$container"}[$interval])) Core Usage/per Second`
5. `avg(rate(container_cpu_usage_seconds_total{namespace=~"$namespace",instance=~"$node",name!="",kubernetes_c
luster=~"$kubernetes_cluster", pod=~"$pod"}[$interval]) / on(namespace, pod, id, kubernetes_cluster)
((container_spec_cpu_quota{namespace=~"$namespace",instance=~"$node",name!="",kubernetes_cluster=~"$kubern
etes_cluster",pod=~"$pod"} > 0) / 100000)) by (namespace, pod, kubernetes_cluster) % Percentage`

PRETTY PROM

```
(
  sum by (kubernetes_cluster) (
    rate(
      container_cpu_usage_seconds_total{container="proxy",kubernetes_cluster=~"prod-.*|dev|staging"}[30s]
    )
    *
    100
  )
)
/
(
  sum by (kubernetes_cluster) (
    kube_pod_container_resource_requests{container="proxy",kubernetes_cluster=~"prod-.*|dev|staging",resource="cpu"}
  )
)
)
```

```
sum(
  rate(
    container_cpu_usage_seconds_total{container=~"$container",instance=~"$node",kubernetes_cluster=~"$kubernetes_cluster",namespace=~"$namespace"}[15m]
  )
)
```

```
avg by (namespace, pod, kubernetes_cluster) (
  rate(
    container_cpu_usage_seconds_total{instance=~"$node",kubernetes_cluster=~"$kubernetes_cluster",name!="",namespace=~"$namespace",pod=~"$pod"}[15m]
  )
)
/ on (namespace, pod, id, kubernetes_cluster)
(
  (
    container_spec_cpu_quota{instance=~"$node",kubernetes_cluster=~"$kubernetes_cluster",name!="",namespace=~"$namespace",pod=~"$pod"}
    >
    0
  )
)
/
100000))
```

INTERNAL CHRONOSPHERE QUERIES

QUERIES PER PANEL WITH MULTIPLE PANELS PER DASHBOARD

6. CPU USAGE (cores)

- `avg by (container,environment,az) (rate(container_cpu_usage_seconds_total{container=~".*",instance=~".*",pod=~".*"}[5m]))` **Core usage/per second**
- `max by (container,environment,az) (rate(container_cpu_usage_seconds_total{container=~".*",instance=~".*",pod=~".*"}[5m]))` **Core usage/per second**
- `(max by (container,environment,kubernetes_cluster) ((container_spec_cpu_shares{container=~".*",instance=~".*",pod=~".*"} / 1024)) or (max by (container,environment,kubernetes_cluster) (label_replace(kube_pod_container_resource_requests{container=~".*",node=~".*",pod=~".*.*",resource="cpu"}, "instance", "$1", "node", "(.*)"))))` **# of Cores**

7. CPU Usage as a Percentage of Requests

- `max by (container) (max by (pod, container, instance,environment,kubernetes_cluster) (rate(container_cpu_usage_seconds_total{container=~".*",instance=~".*",lens_service="$service_name",pod=~".*.*"}[5m])) / (max by (pod, container, instance,environment,kubernetes_cluster) (container_spec_cpu_shares{container=~".*",instance=~".*",pod=~".*.*"} / 1024) or max by (pod, container, instance,environment,kubernetes_cluster) (label_replace(kube_pod_container_resource_requests{container=~".*",node=~".*",pod=~".*.*",resource="cpu"}, "instance", "$1", "node", "(.*)"))))` **% Percentage**

8. CPU Usage as a Percentage of Quota

- `avg by (container) (max by (pod, container, instance,environment,kubernetes_cluster) (rate(container_cpu_usage_seconds_total{container=~".*",instance=~".*",lens_service="$service_name",pod=~".*.*"}[5m])) / (max by (pod, container, instance,environment,kubernetes_cluster) (container_spec_cpu_quota{container=~".*",instance=~".*",pod=~".*.*"} / 1e+06 or max by (pod, container, instance,environment,kubernetes_cluster) (label_replace(kube_pod_container_resource_limits{container=~".*",node=~".*",pod=~".*.*",resource="cpu"}, "instance", "$1", "node", "(.*)"))))` **% Percentage**

THE INTERNET

CADVISOR, NODE EXPORTER DASHBOARDS

9. rate

```
(container_cpu_usage_seconds_total{pod=~"my-pod.*",namespace="my-namespace",  
container!=""}[$__rate_interval])
```

 Core usage/per second

10. $100 - (\text{avg}(\text{irate}(\text{node_cpu_seconds_total}\{\text{mode}=\text{"idle"}\}[30\text{m}])) * 100$ % Percentage

11. $(\text{sum}(\text{irate}(\text{node_cpu_seconds_total}\{\text{mode}\neq\text{"idle"}\}[1\text{m}])) \text{ without } (\text{cpu})) /$
 $\text{count}(\text{node_cpu_seconds_total}) \text{ without } (\text{cpu})$ % Percentage (per core)

12. $\text{sum}(\text{rate}(\text{container_cpu_usage_seconds_total}\{\text{instance}\sim\text{"\$host"},\text{name}\sim\text{"\$container"},$
 $\text{name}\sim\text{"\."}\}[5\text{m}])) \text{ by } (\text{name}) * 100$ % Percentage (per container)

13. $\text{sum}(\text{irate}(\text{node_cpu_seconds_total}\{\text{instance}=\text{"\$node"},\text{job}=\text{"\$job"},$
 $\text{mode}=\text{"system"}\}[\text{\$__rate_interval}])) /$
 $\text{scalar}(\text{count}(\text{count}(\text{node_cpu_seconds_total}\{\text{instance}=\text{"\$node"},\text{job}=\text{"\$job"}\}) \text{ by } (\text{cpu})))$
% Percentage (per core)



THE WHEN

USAGE, PERCENTAGE, BYTES

Options Options Options
Container, KubeStateMetrics,
Node exporter Oh my

THE TYPES OF QUERIES



POD COUNT

:shrug:



%PERCENTAGE

Per core, per
quota, per
container...



OF CORES

Millicores or
Cores



CORE USAGE PER SECOND

Usage

VOTE

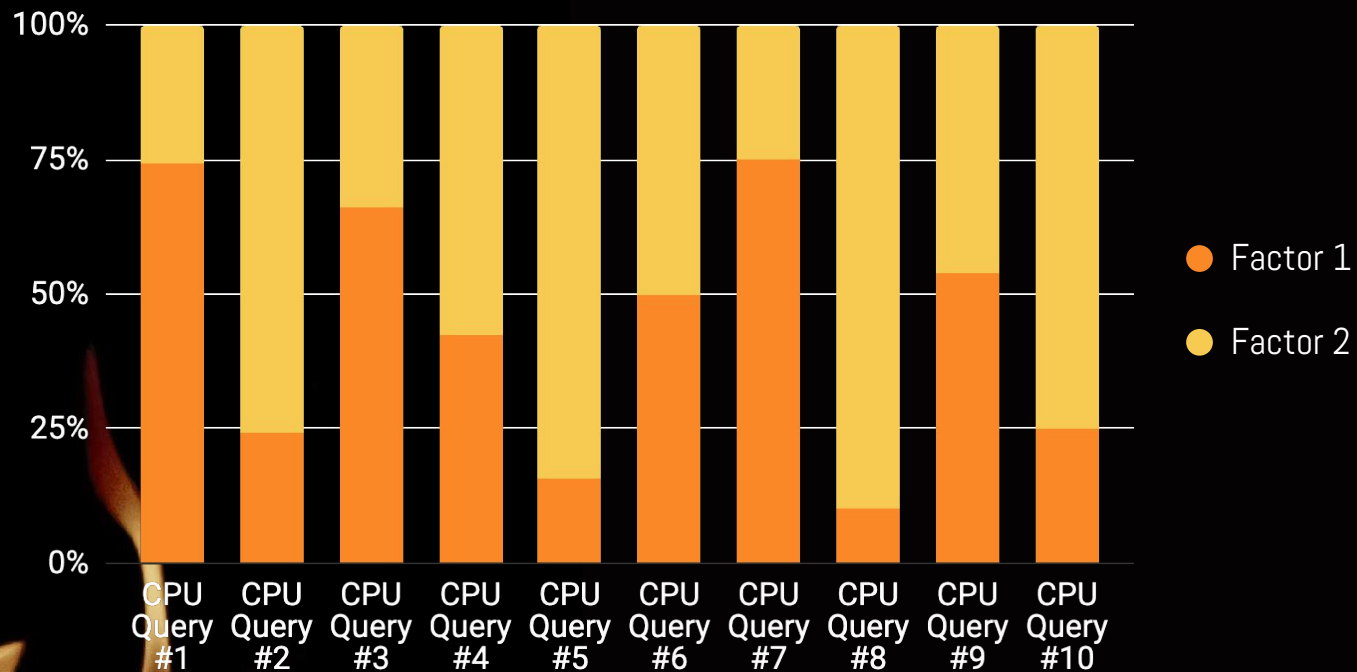
Go to

[https://app.sli.do/event/o3HQSXVyPCDmH
DwUU8EdA2](https://app.sli.do/event/o3HQSXVyPCDmHDwUU8EdA2)



Go to [slido.com #3245009](https://www.slido.com/join/#3245009)

DUMMY RESULTS



REPEAT FOR MEMORY

```
1. max by (container) (max by (pod, container, instance,environment,cluster,namespace)
   (container_memory_working_set_bytes{container=~".*",instance=~".*",pod=~".*.*"}) / (max by (pod, container,
   instance,environment,cluster,namespace) (container_spec_memory_limit_bytes{container=~".*",instance=~".*",pod=~".*.*"})) or
   max by (pod, container, instance,environment,k8s_cluster,namespace)
   (label_replace(kube_pod_container_resource_limits{container=~".*",node=~".*",pod=~".*.*",resource="memory"}, "instance", "$1",
   "node", "(.*)"))))
2. max by (pod, container, instance,environment,cluster,namespace)
   (container_memory_working_set_bytes{container=~".*",instance=~".*",pod=~".*.*"})
3. sum by (pod, container, instance,environment,cluster,namespace)
   (container_memory_rss{container=~".*",instance=~".*",pod=~".*.*"})
4. node_memory_MemTotal_bytes{instance="10.1.1.37:9100",job="node-exporter"}
5. node_memory_MemTotal_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_MemFree_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   (node_memory_Cached_bytes{instance="10.1.1.37:9100",job="node-exporter"} +
   node_memory_Buffers_bytes{instance="10.1.1.37:9100",job="node-exporter"} +
   node_memory_SReclaimable_bytes{instance="10.1.1.37:9100",job="node-exporter"})
6. node_memory_Cached_bytes{instance="10.1.1.37:9100",job="node-exporter"} +
   node_memory_Buffers_bytes{instance="10.1.1.37:9100",job="node-exporter"} +
   node_memory_SReclaimable_bytes{instance="10.1.1.37:9100",job="node-exporter"}
7. node_memory_MemFree_bytes{instance="10.1.1.37:9100",job="node-exporter"}
8. (node_memory_SwapTotal_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_SwapFree_bytes{instance="10.1.1.37:9100",job="node-exporter"})
9. node_memory_MemTotal_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_MemFree_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_Buffers_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_Cached_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_Slab_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_PageTables_bytes{instance="10.1.1.37:9100",job="node-exporter"} -
   node_memory_SwapCached_bytes{instance="10.1.1.37:9100",job="node-exporter"}
10. 100*(sum by(kubernetes_cluster, pod, container, cloud_region) (container_memory_working_set_bytes{container!="POD",
   namespace=~"$namespace", environment="prod", container=~"app"}) / sum by(kubernetes_cluster, pod, container, cloud_region)
   (kube_pod_container_resource_requests{container!="",container!="POD",namespace=~"$namespace", resource="memory", unit="byte",
   environment="prod", container=~"app"}))
```



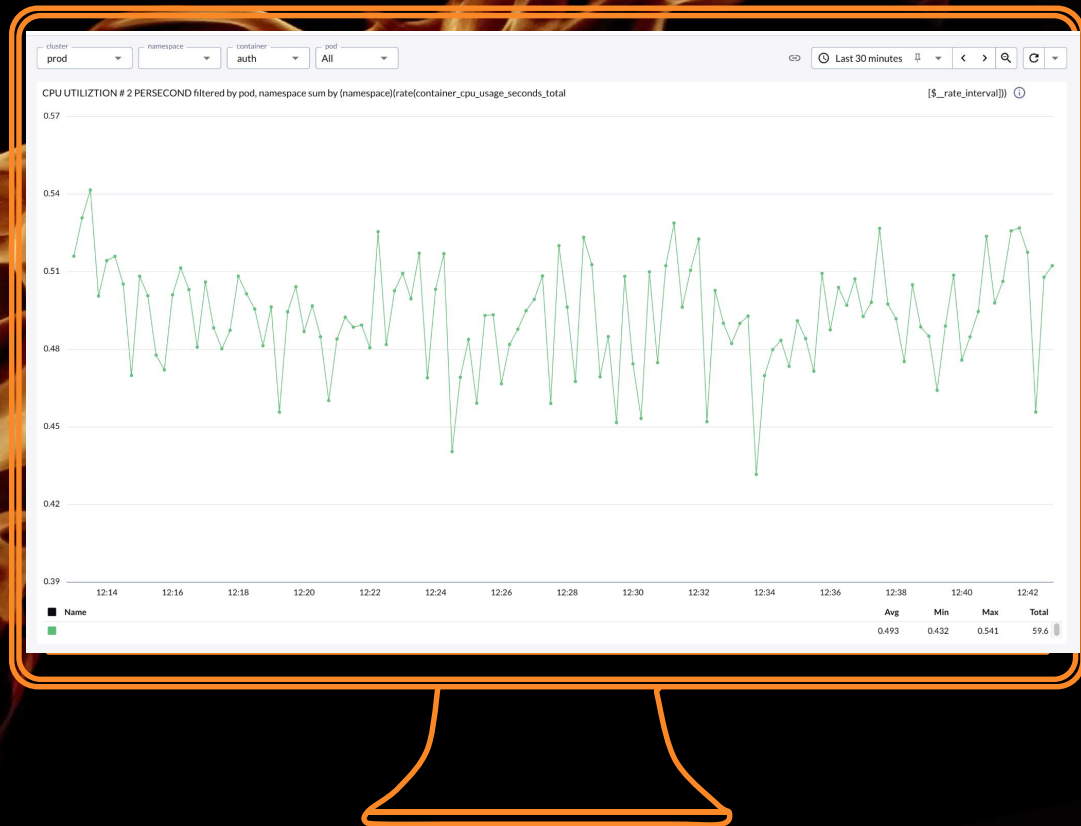
HIGH CPU USAGE

100

*

```
max by (deployment, k8s_cluster, cloud_region) (
  label_replace(
    max by (pod, k8s_cluster, cloud_region) (
      rate(
```

```
container_cpu_usage_seconds_total{container=~"$container"
,image!="",namespace="$namespace"} [5m]
)
),
"deployment",
"$1",
"pod",
"(.*)-.{8,10}-.{5}"
)
)
```



CPU USAGE PER SECOND PER CORE ACROSS NAMESPACE

```
sum by (namespace) (
```

```
rate(
```

```
container_cpu_usage_seconds_total{container=~"$container",
k8s_cluster=~"$cluster",namespace=~"$namespace"} [5m]
```

```
)
```

```
)
```




HIGH CPU USAGE %

```
(
  max by (cloud_region, chronosphere_k8s_cluster,
    namespace) (
    rate(
      container_cpu_usage_seconds_total{chronosphere_k8s_cluste
r=~"$cluster",container=~"$container",namespace="$namespa
ce"} [5m]
    )
  )
  /
  max by (cloud_region, chronosphere_k8s_cluster,
    namespace) (
    kube_pod_container_resource_requests{chronosphere_k8s_clu
ster=~"$cluster",container=~"$container",namespace="$name
space",resource="cpu",unit="core"}
  )
)
*
100
```

THANKS!

Do you have any questions?
sarah.hudspeth@chronosphere.io
Come see us at Booth #H11
Check out the blogs on
chronosphere.io

Observability Day NORTH AMERICA

CREDITS: This presentation template was created by Slidesgo, including icons by Flaticon, and infographics & images by Freepik.



chronosphere | kubecost | do it | Google Cloud

Cheers to Observability!

Thursday Nov 14, 2024 | 6:00-900 PM MT
The Rabbit Hole Lounge, Salt Lake City

[SAVE MY SEAT](#)

KubeCon | CloudNativeCon

A photograph of a bartender with a beard and tattoos, wearing a black shirt, standing behind a bar. The bar is well-stocked with various bottles of liquor and glasses.

Find Chronosphere @ Booth #H11

RESOURCES



- Grafana Open Source Dashboards
 - <https://grafana.com/grafana/dashboards/1860-node-exporter-full/>
 - <https://grafana.com/grafana/dashboards/14282-cadvisor-exporter/>
- TenPlus CPU and Memory Queries (Perses Open Source Dashboard)
 - https://demo.perses.dev/projects/cpu_and_memory_usage_queries_kubecon2024
 - PLEASE GO ADD MORE QUERIES – MAKE YOUR OWN DASHBOARD
- Stack Overflow Article
 - <https://stackoverflow.com/questions/34923788/prometheus-convert-cpu-user-seconds-to-cpu-usage>