



Observability Day

NORTH AMERICA

KubeCon + CloudNativeCon
Salt Lake City, Utah

Enhancing Asynchronous Communication Observability with OpenTelemetry

Liudmila Molkova, Microsoft

Shivanshu Raj Shrivastava, SigNoz

About the Speakers

Observability Day
NORTH AMERICA

November 12, 2024
Salt Lake City



Liudmila Molkova

- Principal Software Engineer
- Microsoft
- Working on Azure client libraries and developer experiences
- Member OTeL Technical Committee
- OTeL SemConv maintainer
- Twitter: [@ne_skazu](#)
- LinkedIn: [lmolkova](#)
- Github: [lmolkova](#)



Shivanshu Raj Shrivastava

- Founding Engineer
- SigNoz (YC W21)
- Building open source OTeL native observability products
- Member and contributor OpenTelemetry
- CNCF Ambassador
- Twitter: [@shivanshu1333](#)
- LinkedIn: [shivanshu1333](#)
- Github: [shivanshuraj1333](#)

Enhancing Asynchronous Communication Observability with OpenTelemetry

Liudmila Molkova, Microsoft

Shivanshu Raj Shrivastava, SigNoz

Async messaging: what can go wrong?

- **Message Loss**
- **High Latency and Throughput Issues**
- **Dead Letter Queues**
- **Monitoring Blind Spots**
- **Throttling and Quota Limits**
- **Message Duplication**
- **Message Ordering Issues**
- **Consumer Lag and Back-pressure**
- **Broker Failures**
- **Network Partitions and Latency**
- **Security and Authentication Issues**
- **Partition bloating**

Messaging queues are Complex!

Managing them is **pain** for cluster managers who are running queues at scale
:(

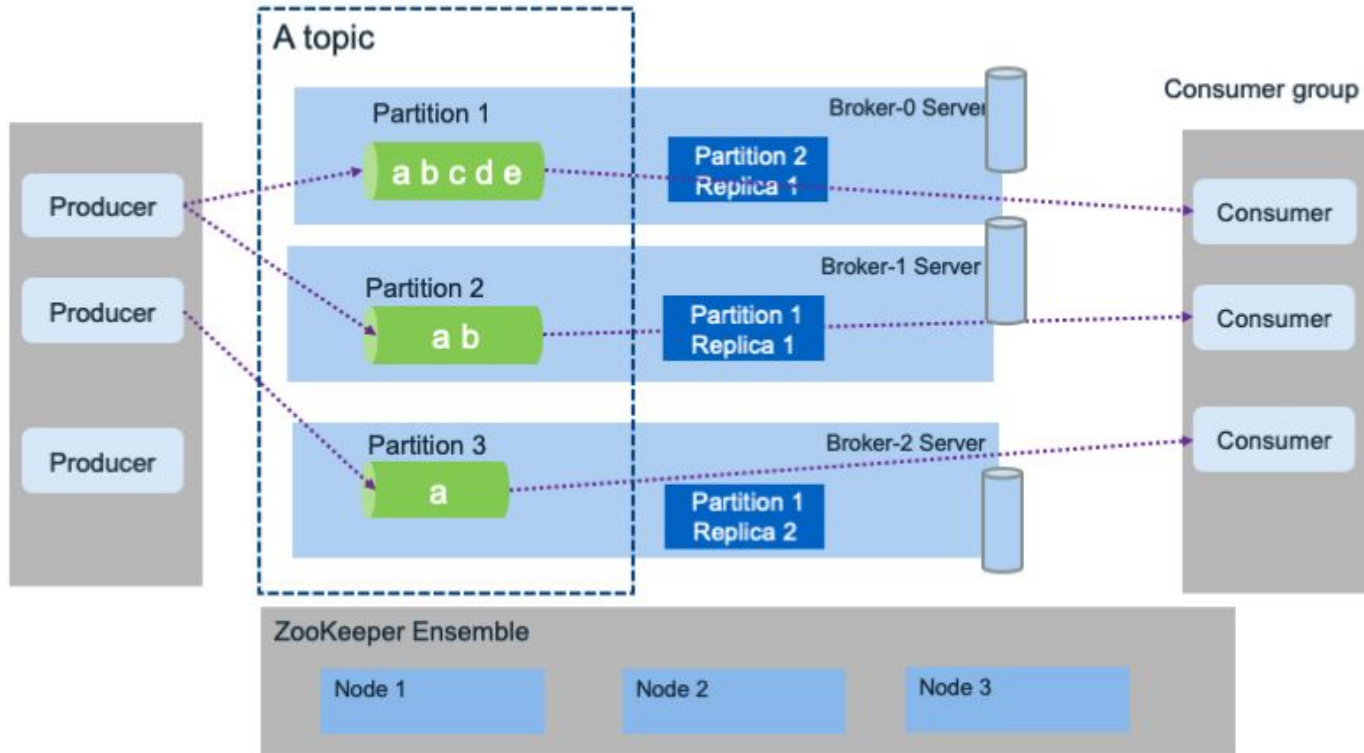
Developers often use **default** configurations, while fine tuning with **deeper observability** can help them **scale better**

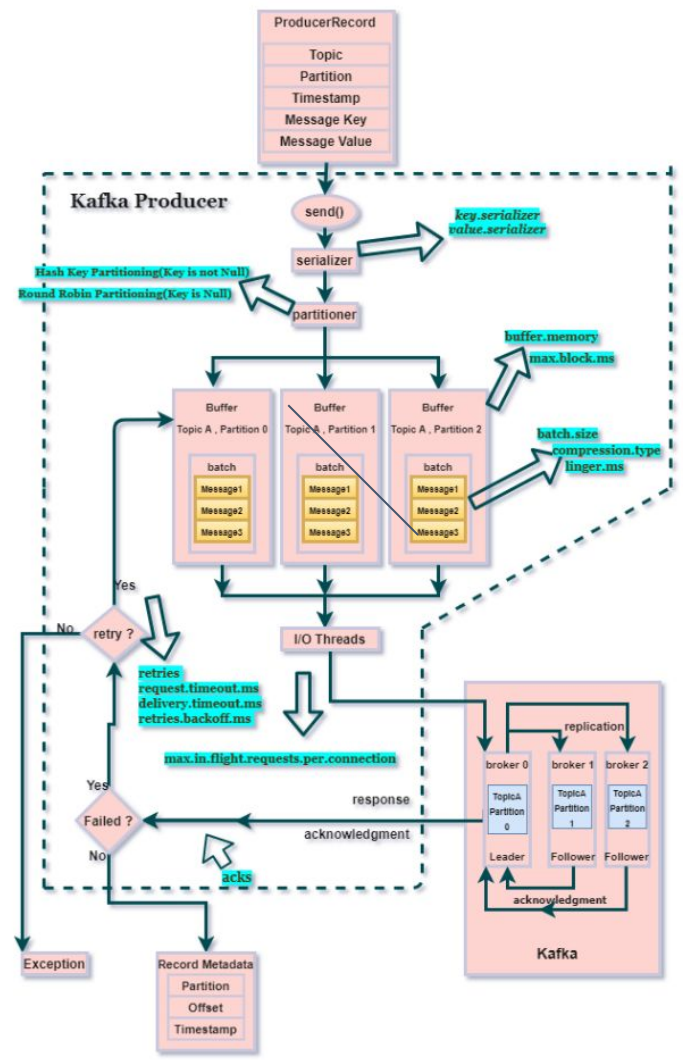
Knowing **when** and **how** things are breaking, can **help in fixing this faster** :)

Let's see how the magic happens inside a queue

**Our favourite
Kafka**

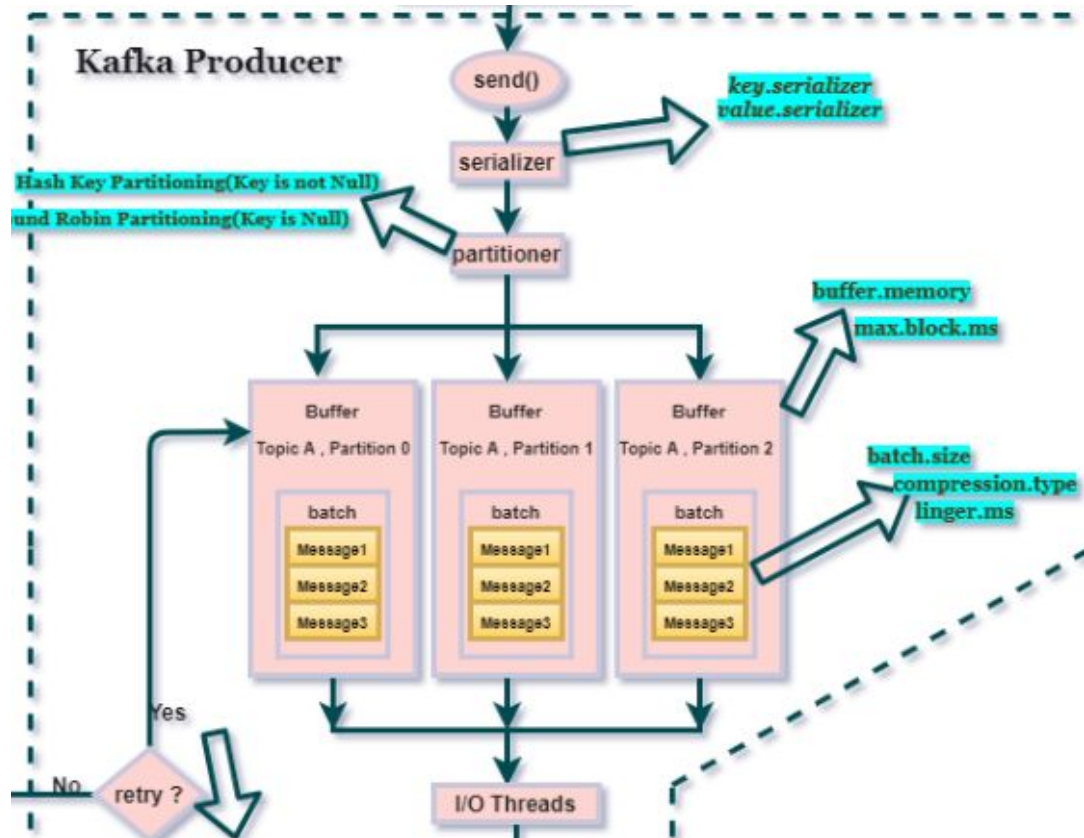
Messaging Queues





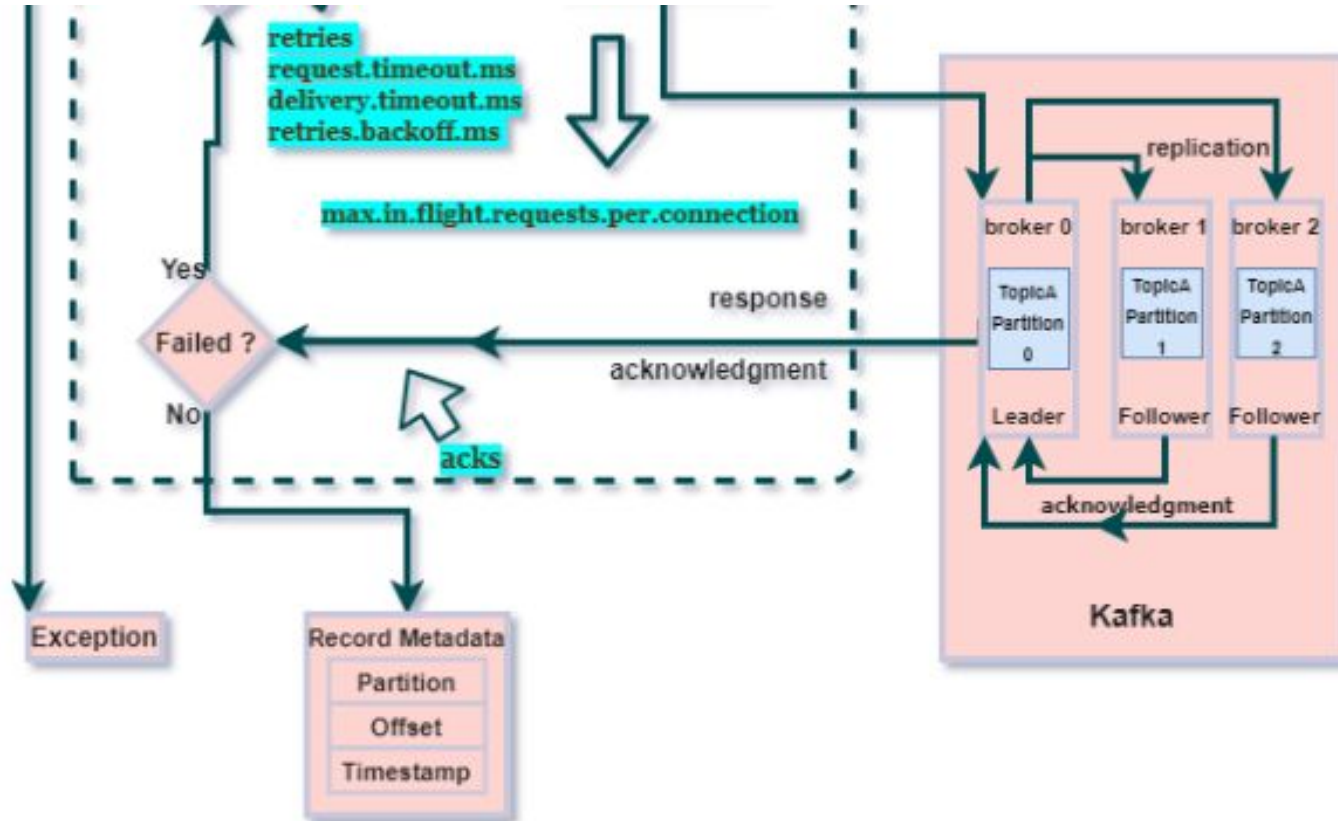
Complex!

Messaging Queues



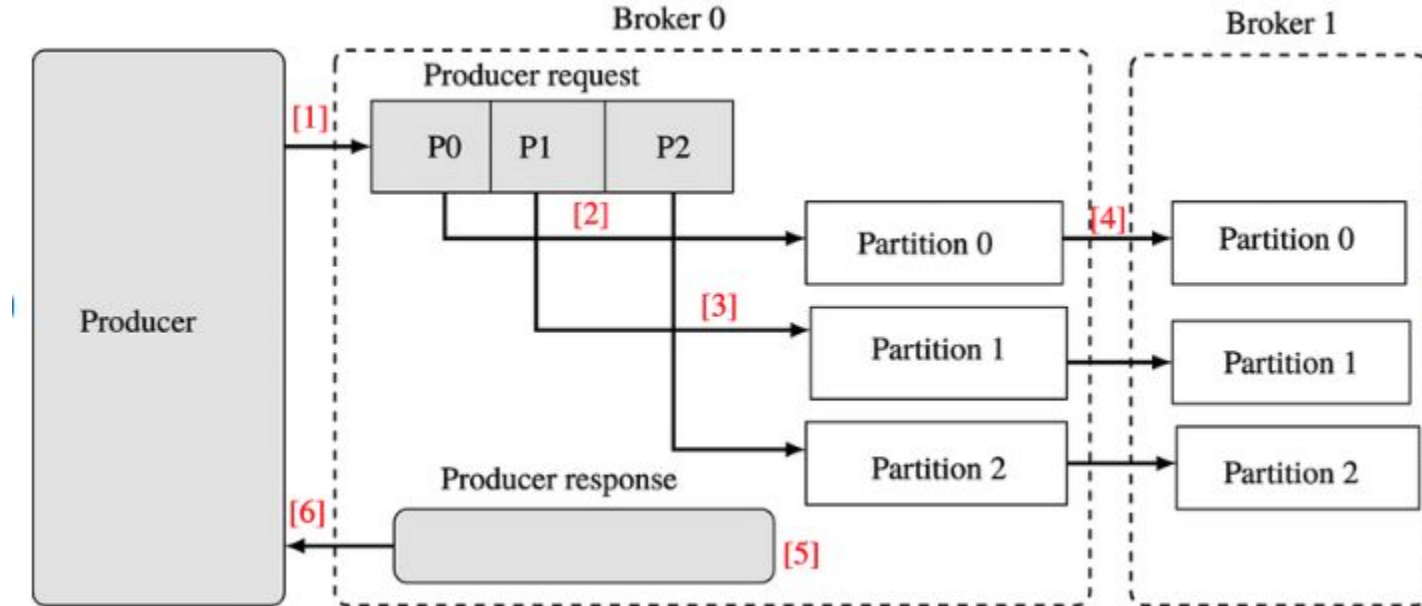
Complex!!

Messaging Queues



Complex!!!

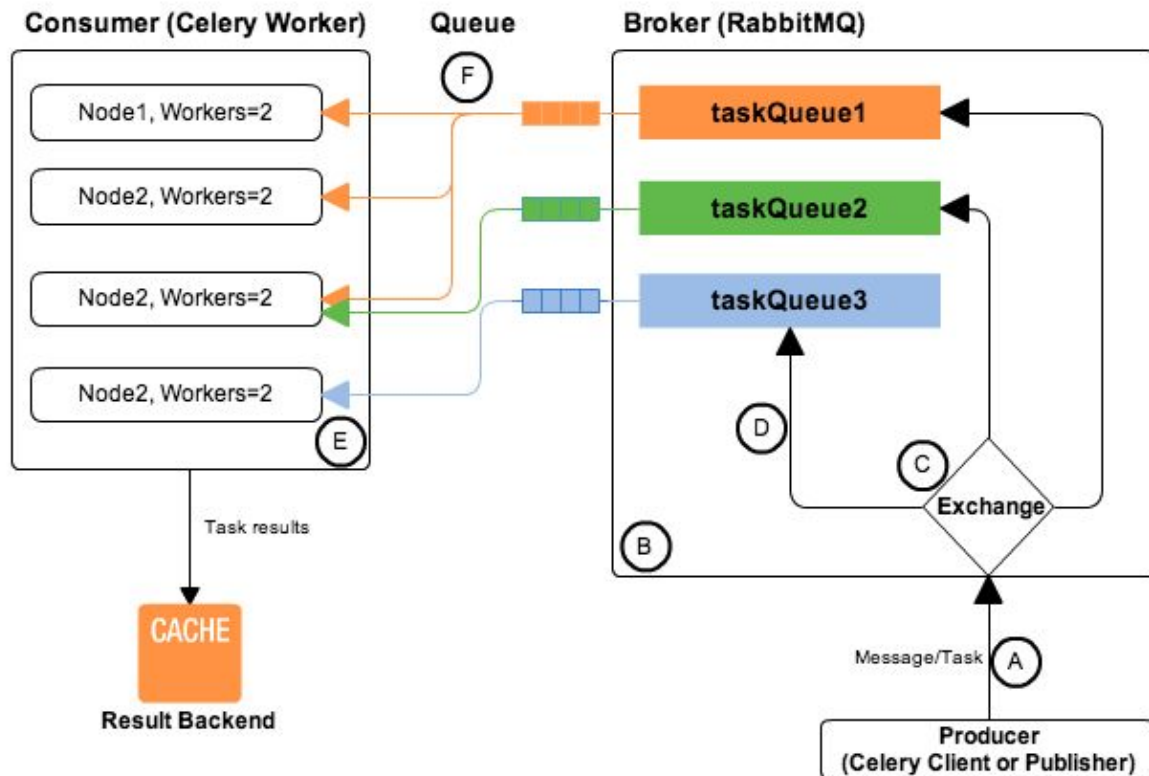
Messaging Queues



LGTM :)

Celery + RabbitMQ

Messaging Queues



There's a queue for every human :|

Messaging Queues

1. **RabbitMQ**
2. **Azure Event Hubs**
3. **Apache Kafka**
4. **ActiveMQ**
5. **Amazon MSK**
6. **Amazon SQS**
7. **Redis**
8. **ZeroMQ**
9. **IBM MQ**
10. **Google Cloud Pub/Sub**
11. **NATS**
12. **NSQ**
13. **Strimzi**

Some OSS, some vendor managed, some proprietary

TBH, many more....

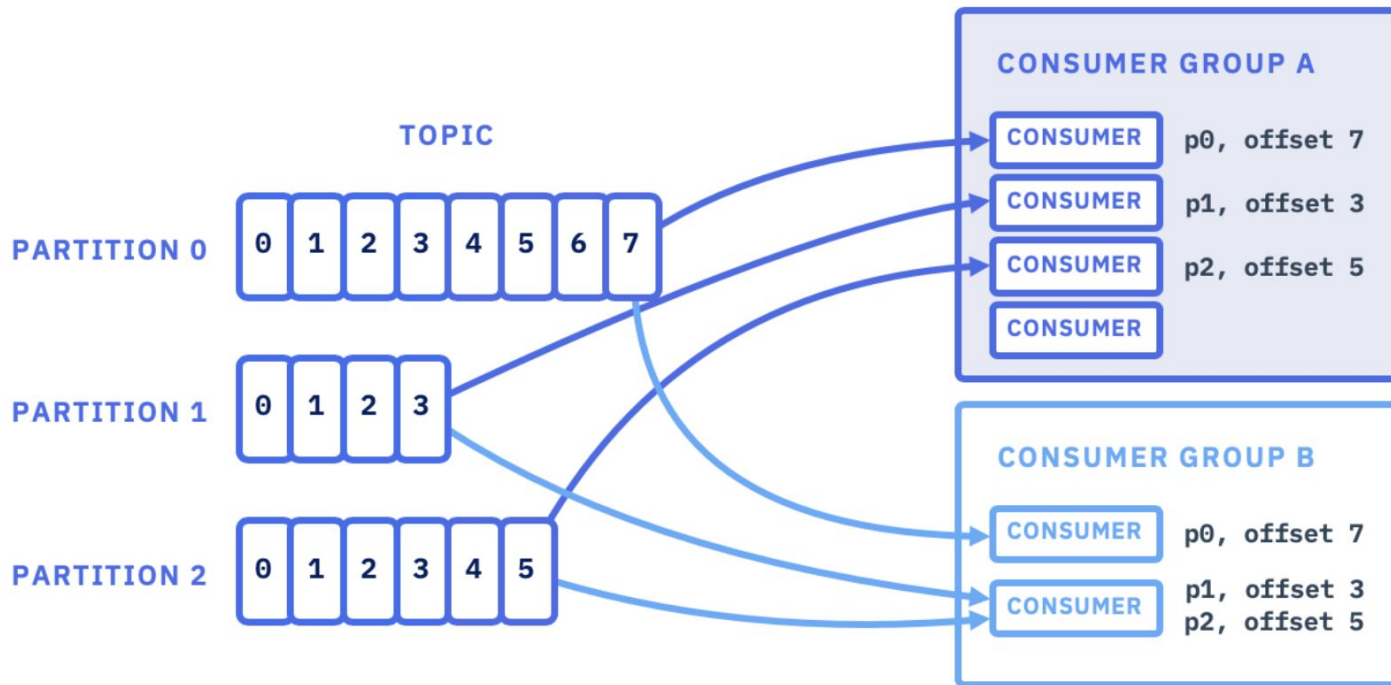
What you (don't) see on broker metrics

- **Kafka does not directly expose producer-side latency metrics per partition.**
 - No one ever knows why there is high producer latency.
 - There is no visibility if there's a problem with Kafka itself or producer.
 - Difficulty tracking producer throughput across different topics
 - Complex to monitor batch sizes and compression ratios
 - Complex to find the correlation between different partitions

- **End-to-End Latency (How much time it took for producer to send a message and a consumer to acknowledge it)**
 - Difficult to track processing time per message
 - No direct correlation between consumer performance and resource utilization
 - Complex to correlate broker metrics with client-side issues

- **Consumer Group Lag correlation with spans**
 - No visibility of which partition is causing the lag?
 - Is the lag caused by consumer application, network problems, or imbalances between producer and consumer performance.
 - Manual implementation needed for trace context propagation.
 - Limited visibility into consumer group rebalancing events

Common problems




- **Infrastructure Monitoring Gaps**

- Complex relationship between JVM metrics and Kafka performance
- Difficult to correlate network issues with message delivery problems
- Limited visibility into disk I/O patterns
- Difficult to track partition growth patterns

Noisy neighbour

Messaging Queues



Metrics -> Trace

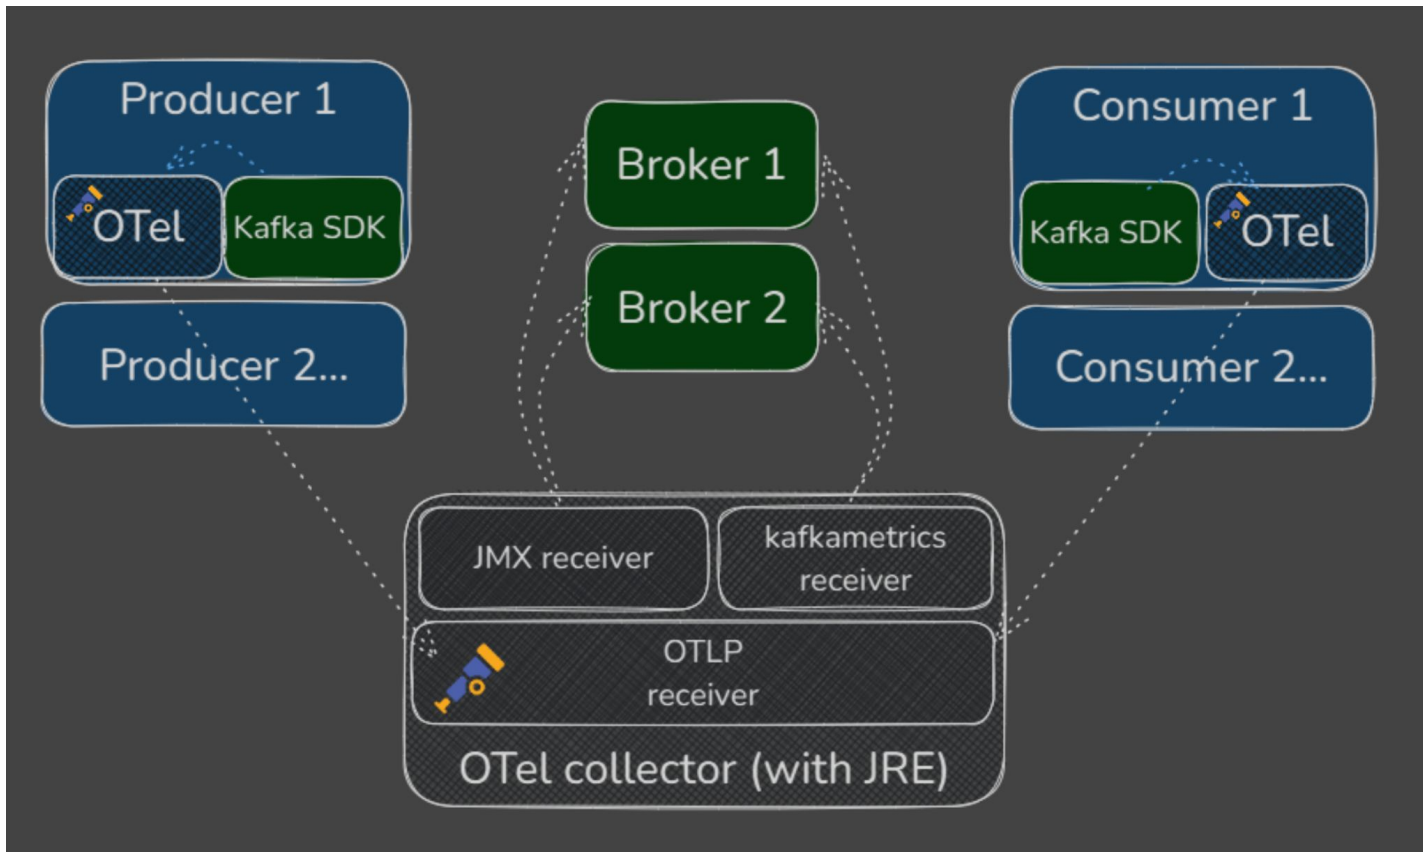
Trace -> Logs

Trace -> Metrics

**All these correlations are hard to achieve
natively in queues**

How to get more insights?

Kafka OTel setup



Sample collector config: <https://github.com/lmolkova/obs-day-messaging/blob/main/configs/otel-collector-config.yml>


- **Stable API** with ability to export telemetry anywhere
- **Instrumentations** for popular clients and frameworks
- **Semantic conventions** to unify telemetry format

Distributed tracing in messaging queues

- Useful in debugging individual flows and root cause issues
- Each operation is represented as span
- Trace context is propagated with the message through the messaging queues

producer: POST /orders 47d0aa9

Trace detail **10/29/2024 7:18:14.347 PM** Duration **62.35ms** Resources **2** Depth **4** Total spans **4** [View logs](#)



More context in span attributes

Name	Value
Name	orders_topic publish
Kind	Producer
messaging.client_id	kafka-producer-producer-1
messaging.destination.name	orders_topic
messaging.destination.partition.id	2
messaging.kafka.message.offset	79342
messaging.operation	publish
messaging.system	kafka

Check out <https://github.com/open-telemetry/semantic-conventions/blob/main/docs/messaging/messaging-spans.md>

- Useful to understand **overall state and health** of the system
- **Easy** to store and query
- Very **low performance overhead**
- Telemetry is **aggregated** on client and exported with **low-cardinality** dimensions

- **Operation duration** histograms
- Sent, received, processed **message count**
- **Extra dimensions** to filter and break things down
- More to come: lag, delivery count, etc
- Implemented in a **few instrumentations** so far

Check out <https://github.com/open-telemetry/semantic-conventions/blob/main/docs/messaging/messaging-metrics.md> for the details!

Name	Instrument Type	Unit (UCUM)	Description
<code>messaging.client.operation.duration</code>	Histogram	s	Duration of messaging operation initiated by a producer or consumer client.




Attribute	Type	Description	Examples
<code>error.type</code>	string	Describes a class of error the operation ended with.	<code>amqp:decode-error</code> ; <code>KAFKA_STORAGE_ERROR</code> ; <code>channel-error</code>
<code>messaging.consumer.group.name</code>	string	The name of the consumer group with which a consumer is associated.	<code>my-group</code> ; <code>indexer</code>
<code>messaging.destination.name</code>	string	The message destination name.	<code>MyQueue</code> ; <code>MyTopic</code>
<code>messaging.destination.partition.id</code>	string	The identifier of the partition messages are sent to or received from, unique within the <code>messaging.destination.name</code> .	<code>1</code>
<code>messaging.operation.name</code>	string	The system-specific name of the messaging operation.	<code>send</code> ; <code>receive</code> ; <code>ack</code>
<code>messaging.operation.type</code>	string	A string identifying the type of the messaging operation.	<code>create</code> ; <code>send</code> ; <code>receive</code>
<code>messaging.system</code>	string	The messaging system as identified by the client instrumentation.	<code>activemq</code> ; <code>aws_sqs</code> ; <code>eventgrid</code>
<code>server.address</code>	string	Server domain name if available without reverse DNS lookup; otherwise, IP address or Unix domain socket name.	<code>example.com</code> ; <code>10.1.2.80</code> ; <code>/tmp/my.sock</code>
<code>server.port</code>	int	Server port number.	<code>80</code> ; <code>8080</code> ; <code>443</code>


- `messaging.operation.name`
- `messaging.system`
- `messaging.batch.message_count`
- `messaging.consumer.group.name`
- `messaging.destination.anonymous`
- `messaging.destination.name`
- `messaging.destination.subscription.name`
- `messaging.destination.template`
- `messaging.destination.temporary`
- `messaging.operation.type`
- `messaging.client.id`
- `messaging.destination.partition.id`
- `messaging.message.conversation_id`
- `messaging.message.id`
- `messaging.message.body.size`
- `messaging.message.envelope.size`

- **Throughput:** records, requests, IO, etc
- **Latency:** request, IO, internal queues
- Connectivity, authentication, buffers, batching and more

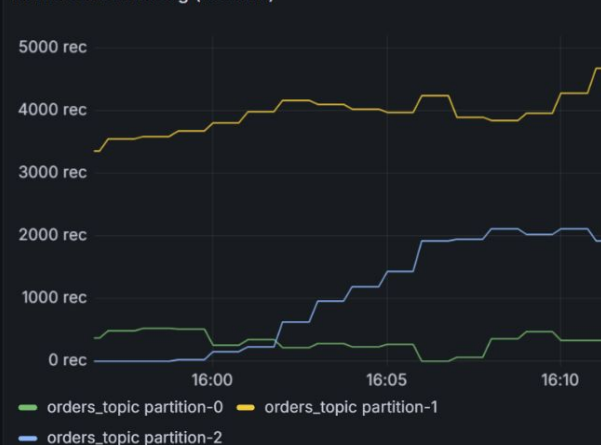
Kafka request latency, avg



Kafka throughput (records)

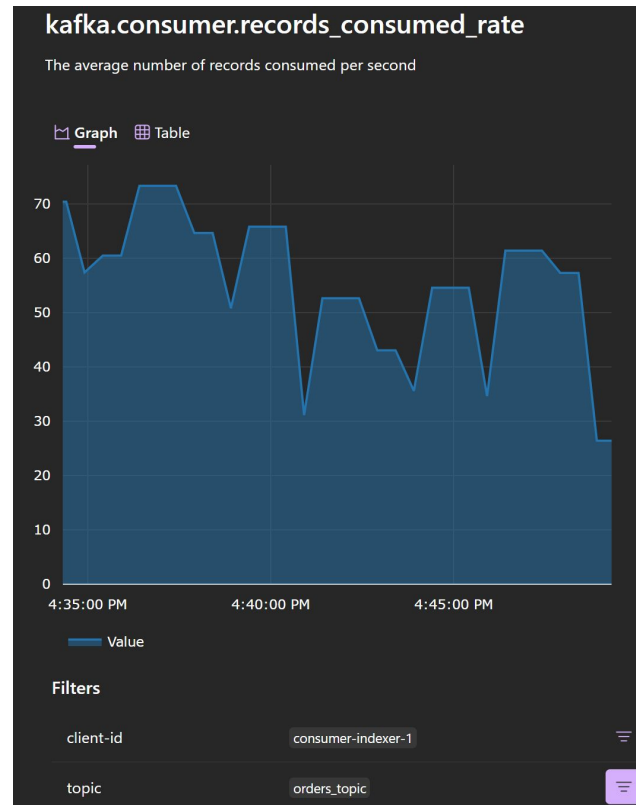


Kafka consumer lag (records)



- Client and language-specific
 - Instrumentations pull metrics from Kafka SDKs
 - Some metrics come through broker plugin ([KIP-714](#))
- **Don't provide enough context**
 - One common dimension: `client_id`
 - Some metrics include: `node_id`, `topic`, `partition`

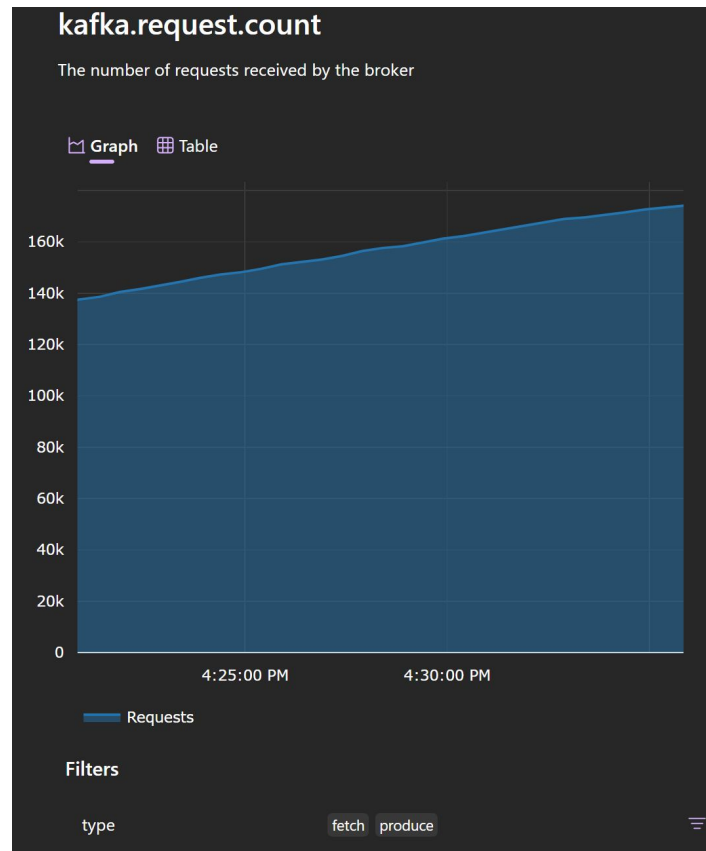
Existing Kafka metrics reported by OTel instrumentations: [Java](#), [.NET](#)



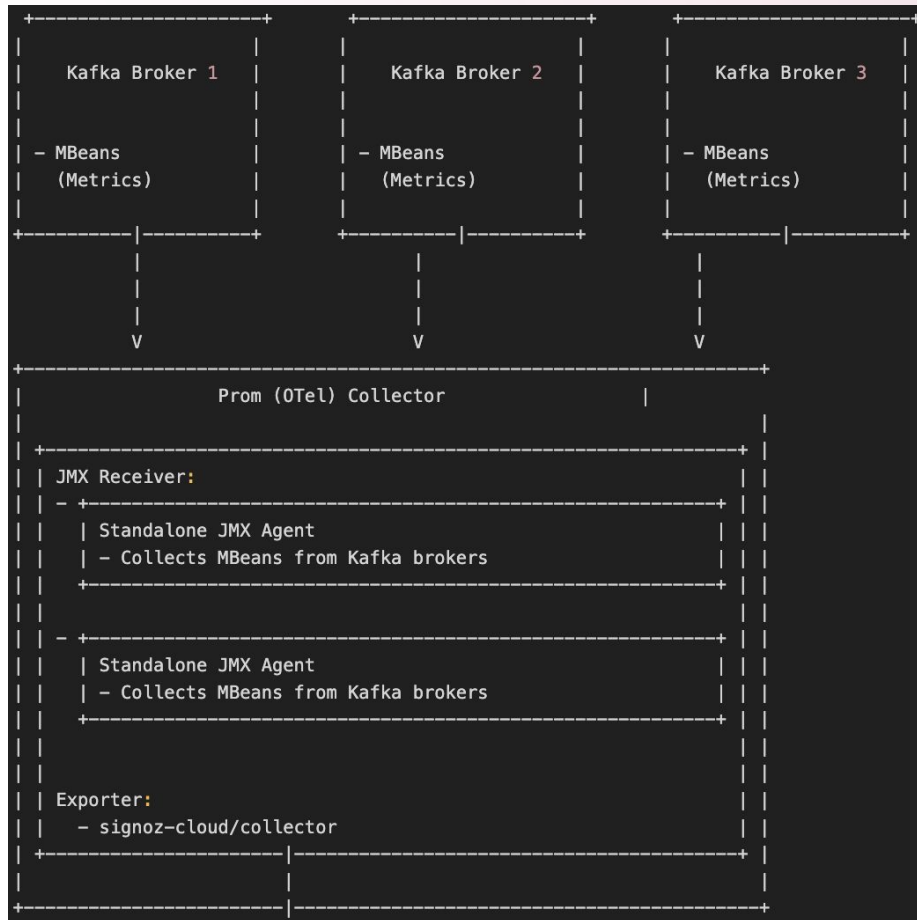
Kafka broker 'native' metrics

- Throughput, latency, elections, offsets, etc
- Available via [JMX](#) and [kafka-metrics](#) receivers in OTel collector
- **Still not enough context**

Existing Kafka metrics reported by OTel instrumentations: [Java](#), [.NET](#)

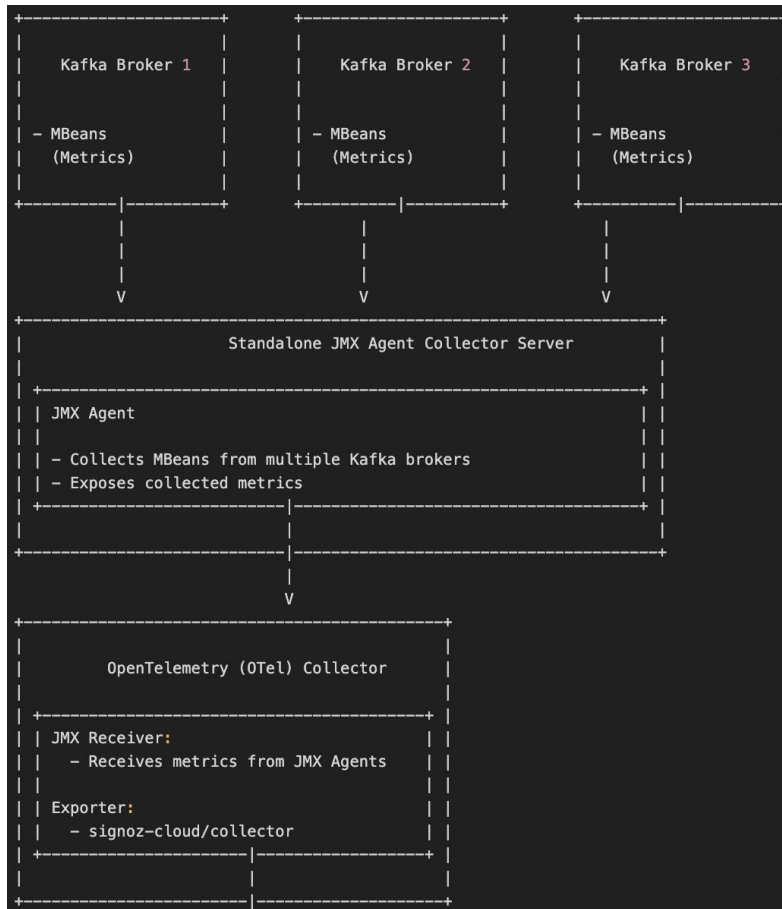


How to get deeper insights?

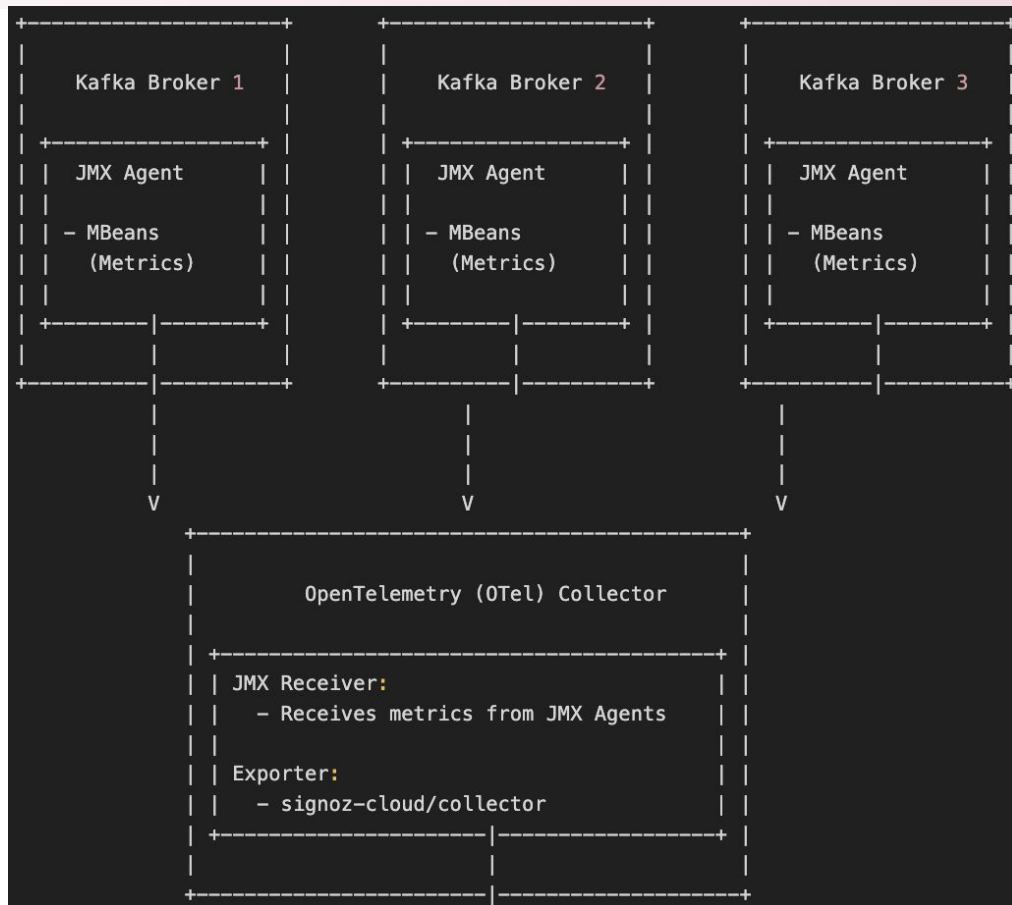


MBean
ObjectName: kafka.server:type=BrokerTopicMetrics, name=MessagesInPerSec
Attributes
<ul style="list-style-type: none">- Count: 123456- MeanRate: 456.78- OneMinuteRate: 78.90- FiveMinuteRate: 67.89- FifteenMinuteRate: 56.78
Operations
<ul style="list-style-type: none">- resetStatistics(): void- sampleStats(): CompositeData

How to get deeper insights?



How to get deeper insights?



Demo!

- Participate in OTEL Messaging Semantic Conventions
 - Define new metrics, add attributes or conventions for new systems
- Improve existing instrumentations
 - Available in [Java](#), [Python](#), [.NET](#), [node.js](#), and [other languages](#)
- Ask for native instrumentation in your favorite Kafka client library

Questions?

Thanks for joining!



November 12, 2024
Salt Lake City