

envoycon

NORTH AMERICA

Extending Envoy: A Guide to Custom Extensions with Envoy Gateway

Guy Daich, SAP

Huabing Zhao, Tetrade

Background: Gateway API


Gateway API: A Superior Alternative to Ingress API

- **API Features:** Offers a richer set of features compared to Ingress API
- **Gateway Standard:** Designed as a universal API standard to maximize compatibility across implementations
- **Extension Mechanism:** Introduces well-defined extension points, replacing the unstructured annotations in Ingress API
 - Policy Attachment: Enables attaching custom policies without altering the core API
 - Custom Filters: Supports rule-level custom traffic filters for granular control
 - Custom Backend: Routes traffic efficiently to various backends

Background: Envoy Gateway

Envoy Gateway: Empowering Gateway API with Envoy

- A Gateway controller fully compatible with Gateway API
- Simplifies deployment and management of Envoy as an API Gateway
- Handles traffic for both Kubernetes clusters and VM-based workloads
- Goes beyond the Gateway API with advanced features
 - Policy Attachment
 - ClientTrafficPolicy
 - BackendTrafficPolicy
 - SecurityPolicy
 - EnvoyPatchPolicy
 - EnvoyExtensionPolicy
 - HTTPRouteFilter
 - Backend
 - IP
 - HostName
 - Unix Domain Socket



Envoy Extension Policy

EnvoyExtensionPolicy allows loading custom extensions into Envoy to execute user-defined logic for request and response processing.

- Supported Extension types:
 - Wasm
 - Ext-Proc
 - More in future: Lua, Dynamic module ...
- Configuration Flexibility:
 - Supported attachment targets: Gateway, HTTPRoute
 - Multiple Ext-Proc/Wasm extensions can be defined per policy
 - **EnvoyProxy** CR can be used to determine filter ordering setting, allowing users to control when Wasm/Ext-Proc filters are invoked in the filter chain

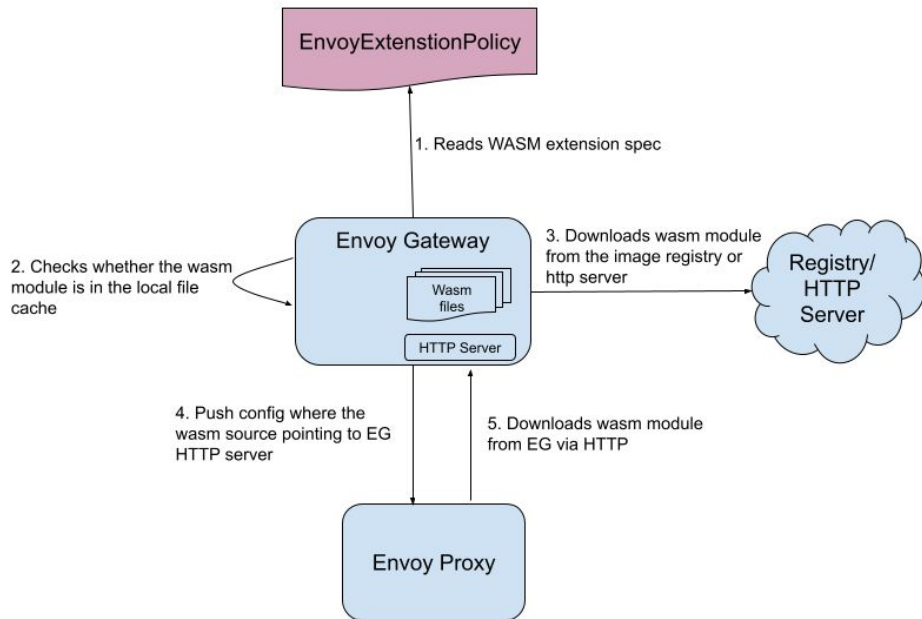
```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyProxy
metadata:
  name: custom-proxy-config
  namespace: envoy-gateway-system
spec:
  filterOrder:
    - name: envoy.filters.http.ext_proc
      after: envoy.filters.http.ratelimit
```

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyExtensionPolicy
metadata:
  namespace: envoy-gateway
  name: my-custom-extensions
spec:
  targetRef:
    group: gateway.networking.k8s.io
    kind: Gateway
    name: my-gateway
  wasm:
    - name: wasm-filter-1
      code:
        type: HTTP
        http:
          url: https://www.example.com/wasm-filter-1.wasm
          sha256: 746df05c8f3a0b07a46c0967cfbc5cbe5b9d48d0
    - name: wasm-filter-2
      code:
        type: Image
        image:
          url: oci://www.example.com/wasm-filter-2:v1.0.0
          pullSecretRef:
            name: my-pull-secret
            sha256: a1efca12ea51069abb123bf9c77889fcc2a31cc5
  extProc:
    - backendRefs:
        - name: my-ext-proc-svc
          port: 8000
```

Wasm Extension - OCI Image Support

Envoy Gateway supports OCI image as a remote Wasm code source.

- **Versioning:** Users can use the tag of the OCI image to manage the version of the Wasm module.
- **Security:** Users can use private registries to store the Wasm module securely.
- **Distribution:** Users can use the existing tools of the OCI registry to distribute the Wasm module.



EnvoyExtensionPolicy - Wasm Extension

OCI Image Wasm source

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyExtensionPolicy
metadata:
  name: wasm-test
spec:
  targetRefs:
  - group: gateway.networking.k8s.io
    kind: HTTPRoute
    name: backend
  wasm:
  - name: wasm-filter-1
    rootID: my_root_id
    code:
      type: Image
      image:
        url: zhaohuabing/testwasm:v0.0.1
  - name: wasm-filter-2
    rootID: "my-root-id"
    code:
      type: Image
      image:
        url: oci://my.private.registry/wasm-filter-2:v1.0.0
        pullSecretRef:
          name: my-pull-secret
          sha256: a1efca12ea51069abb123bf9c77889fcc2a31cc5483f
    config:
      parameter1: value1
      parameter2: value2
```

HTTP Wasm source

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyExtensionPolicy
metadata:
  name: wasm-test
spec:
  targetRefs:
  - group: gateway.networking.k8s.io
    kind: HTTPRoute
    name: backend
  wasm:
  - name: wasm-filter-1
    code:
      type: HTTP
      http:
        url: https://www.example.com/wasm-filter-1.wasm
        sha256: 746df05c8f3a0b07a46c0967cfbc5cbe5b9d48d0f79f
    config:
      parameter1:
        key1: value1
        key2: value2
      parameter2: value3
```

External Processing in Envoy Gateway

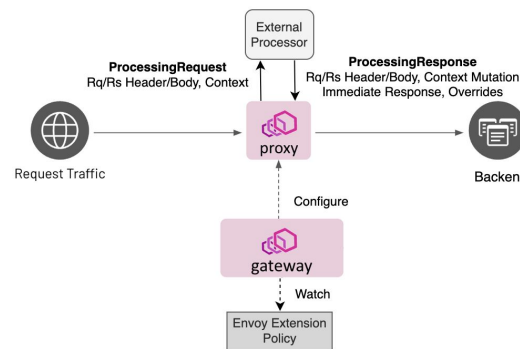
The **External Processing** extension is an HTTP filter that supports out-of-process extensibility using a callout to an **external gRPC service**.

The external gRPC service can inspect and mutate Headers and Bodies of the HTTP stream by registering to relevant **stream hooks**.

Ext-Proc Protocol:

- **Processing Request:** Created by Envoy and sent to external processor for a registered hook, containing one of HTTP Request/Response headers or body (chunks).
- **Processing Response:** Created by the grpc service in response to a Processing Request, containing mutation of HTTP request/response headers or body or Immediate Response.

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyExtensionPolicy
metadata:
  namespace: default
  name: policy-for-http-route
spec:
  targetRef:
    group: gateway.networking.k8s.io
    kind: HTTPRoute
    name: httproute-1
  extProc:
    - backendRefs:
      - Name: my-ext-proc-svc
        Port: 8000
      messageTimeout: 1s
      failedOpen: true
      processingMode:
        request:
          body: Buffered
        response:
          body: Buffered
```



Implications of an out-of-process extension:

- **Security:** Ensure identity of external processor and confidentiality/integrity of messages
- **Resilience:** communication to a different component should consider network instability, decoupled lifecycle of components, possible unavailability, etc.
- **Performance:** communication with a remote service, additional (un)marshaling, additional encryption

Envoy Gateway offers advanced integration options for the external grpc service:

- **Security:** Use **BackendTLSPolicy** and **EnvoyProxy** CRs to establish MTLS between Envoy Proxy and External Processing
- **Resilience:** Use Envoy Gateway's Extended **BackendRef** to configure a variety of behaviors for the External Processor connection pool, such as: Load Balancing, Circuit Breaking, HealthChecks (incl. grpc), Failover (primary, secondary), Timeouts, Buffer Sizes, KeepAlive and more
- **Performance:** Use **Backend** and **EnvoyProxy** CRs to integrate an External Processor that is exposed with Unix domain sockets in the same Pod (sidecar) or Host (decoupled deployment and scaling).

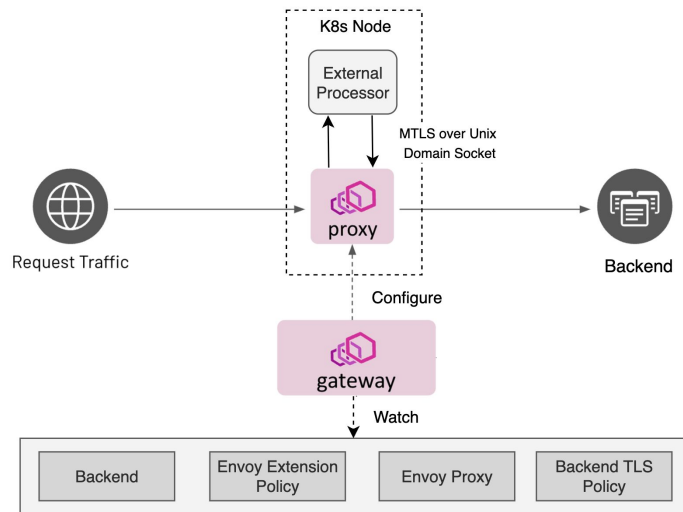
Advanced Deployment Patterns

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: Backend
metadata:
  name: grpc-uds-ext-proc
spec:
  endpoints:
    - unix:
        path: /var/run/ext-proc/extproc.sock
```

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyExtensionPolicy
metadata:
  name: policy-uds-ext-proc
spec:
  targetRefs:
    - group: gateway.networking.k8s.io
      kind: HTTPRoute
      name: http-with-extproc-uds-tls
  extProc:
    - backendRefs:
        - kind: Backend
          group: gateway.envoyproxy.io
          name: grpc-uds-ext-proc
  processingMode:
    request: {}
    response: {}
  backendSettings:
    healthCheck:
      active:
        healthyThreshold: 1
        type: GRPC
    circuitBreaker:
      maxConnections: 2048
  timeout:
    tcp:
      connectTimeout: 15s
```

```
apiVersion: gateway.networking.k8s.io/v1alpha3
kind: BackendTLSPolicy
metadata:
  name: policy-btls-uds-extproc
spec:
  targetRefs:
    - group: gateway.envoyproxy.io
      kind: Backend
      name: grpc-uds-ext-proc
  validation:
    caCertificateRefs:
      - name: grpc-ext-proc-ca
```

```
apiVersion: gateway.envoyproxy.io/v1alpha1
kind: EnvoyProxy
metadata:
  name: proxy-config
spec:
  provider:
    type: Kubernetes
    kubernetes:
      envoyDeployment:
        container:
          volumeMounts:
            - mountPath: /var/run/ext-proc
              name: socket-dir
      pod:
        volumes:
          - name: socket-dir
            hostPath:
              path: /var/run/ext-proc
              type: ""
      backendTLS:
        clientCertificateRef:
          kind: Secret
          name: example-client-cert
```



When to Use Wasm/ExtProc

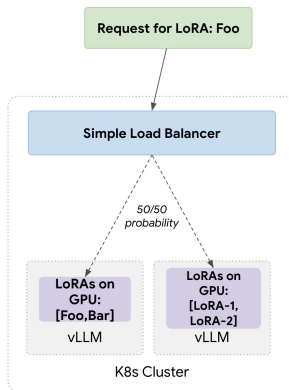

Wasm extensions excel at **lightweight, in-path data processing**, while External Process extensions can handle more **complex logic requiring external system interaction**.

Consider the below factors when choosing between them:

- **Performance:**
 - Wasm extensions offer superior performance since they run directly within the Envoy process.
 - External Process extensions rely on network calls, which can slightly reduce performance due to additional latency. Sidecar deployment can mitigate this to some degree.
- **Functionality:**
 - Wasm operates in a sandbox, restricting system calls and access to external resources.
 - External Process extensions have no such limitations and can be built in any language, with full access to system resources.
- **Deployment:**
 - Wasm extensions can be dynamically loaded by Envoy from an OCI registry or HTTP URL.
 - External Process extensions require managing a separate process, adding complexity to deployment.
- **Security:**
 - Wasm runs within Envoy, meaning bugs in the extension could impact Envoy's stability.
 - External Process extensions run independently, so failures won't directly affect Envoy's operation.
- **Scalability:**
 - Wasm extensions are embedded in Envoy and scale together with the Envoy instances.
 - External Process extensions can scale independently, with separate resource management.

Demo: K8s LLM Instance Gateway

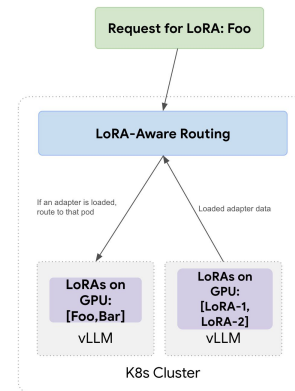
- **LLMs** can be **tuned** for specific tasks.
- For LLM operators, operating a fleet of single-purpose (tuned) LLM servers is **costly** and often leads to **underutilization of resources**.
- **Low-Rank Adaptation (LoRA)**, a popular tuning technique, can be used to run **multiple adapters with a shared model** in a single LLM Server instance, **improving resource utilization**.
- However, the number of **concurrently loaded adapters is limited**, and **naive round-robin** load balancing can lead to significant performance impacts, due to delays for **requests queued for an unloaded adapter**.
- A tailored Load Balancing algorithm has the potential to address these concerns.




**Generic
Load Balancing**

Demo: K8s LLM Instance Gateway

- The **K8s Serving WG** is working to improve K8s as a platform for LLM inference. Together with the **Gateway API WG**, they have explored a solution to this issue based on **Gateway extensions**.
- Gateways provide a **comprehensive view** of system-wide and **instance-specific LLM resource utilization and adapter state**. They are also independent of any specific LLM server implementation
- Gateways can implement **LoRA-aware load balancing** to select pods with the appropriate **LoRA adapter loaded, shortest queue, and lowest KV cache** utilization.
- Early results from PoCs showed that **throughput increased by 85%** and latency decreased significantly.

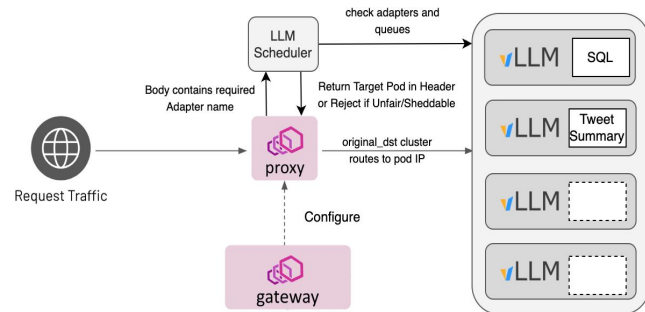


**LoRA-Aware
Load
Balancing**



Demo: K8s LLM Instance Gateway

- The LLM scheduler is an External Processing service which is registered to request body and header hooks.
 - Scheduler **monitors** the state of vLLM pods, consuming **metrics** such as Loaded adapters, queue sizes, KV Cache utilization.
 - Incoming requests are buffered and the desired LoRA adapter is **extracted from the body**.
 - The **optimal pod** for serving the request is selected and provided to Envoy via a dedicated header added with **header mutation**.
 - If load shedding or rate limiting is required an **immediate response** is returned.
- Upstream routing leverages the Envoy **original_dst** cluster, capable of dynamically routing to an address specified by the scheduler.



Demo: K8s LLM Instance Gateway

Learn more about LLM Instance Gateway and WG-Serving:


- LLM Instance [Github](#), [Proposal](#)
- WG-Serving: [Github](#), [KubeCon NA 24 Project Talk](#)

Get involved in Envoy Gateway AI related topics:

- Slack: [#envoy-ai-gateway](#)
- Envoy AI Gateway [Webinar](#), [KubeCon NA 24](#)

Demo

- K8s and Envoy Gateway manifests
- Implementation of an External Processing service
- Demonstrate LoRA-Aware routing, where requests are routed to vLLM pods that have the desired LoRA adapter loaded



Thank You

Q&A

Get involved in Envoy Gateway:

- <https://github.com/envoyproxy/gateway>
- <https://gateway.envoyproxy.io/docs/>

Session Feedback:

