

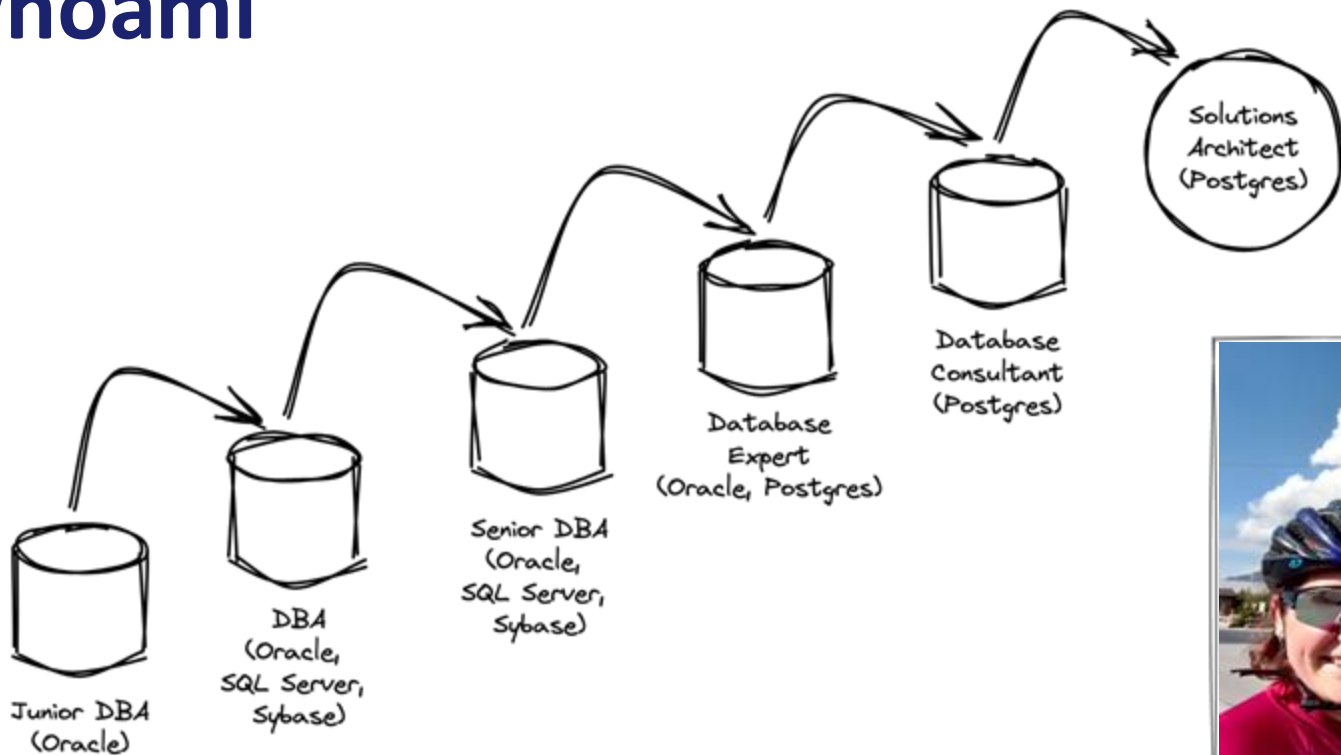


Kubernetes as Your (Virtual) DBA


Karen Jex | Senior Solutions Architect @ Crunchy Data

Data on Kubernetes Day @ KubeCon CloudNativeCon | Salt Lake City | Nov 2024

whoami



Introduction






Agenda

1. What Came Before?
2. What does a DBA do anyway?
3. Postgres Operators for Kubernetes
4. What do you want from an Operator?
5. How to get Started

Agenda


1. **What Came Before?**
2. What does a DBA do anyway?
3. Postgres Operators for Kubernetes
4. What do you want from an Operator?
5. How to get Started





Features of Containers

- Lightweight
- Scalable
- Portable
- Isolated
- Stateless
- Ephemeral





Herding Cats (Managing Containers)

Agenda

1. What Came Before?
2. **What does a DBA do anyway?**
3. Postgres Operators for Kubernetes
4. What do you want from an Operator?
5. How to get Started



What is a DBA?

“Database administrators **create, organise and look after computer systems that store data** for a company.” <https://nationalcareers.service.gov.uk/job-profiles/database-administrator>

“A database administrator, or DBA, is responsible for **maintaining, securing, and operating databases** and also ensures that **data is correctly stored and retrieved.**” <https://www.oracle.com/database/what-is-a-dba/>

“Database administrators **use specialist software to organise and maintain a secure database.**” <https://www.prospects.ac.uk/job-profiles/database-administrator>

“A Database Administrator, or Database Manager is responsible for **managing computer systems that store and organise data** for companies. Their duties include... **securing data and identifying areas for improvement with the infrastructure.**” <https://www.indeed.com/career/database-administrator>

DBA Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- DB Design/Data Modelling
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance

Agenda


1. What Came Before?
2. What does a DBA do anyway?
3. **Postgres Operators for Kubernetes**
4. What do you want from an Operator?
5. How to get Started



Kubernetes Features

- Provisioning & Deployment
- Configuration
- Scheduling
- Scaling up & down
- Self-healing
- Services
- Storage
- Resource allocation
- Load Balancing & Networking
- Security
- Stateful Sets
- Sidecars

Running Postgres on Kubernetes




Kubernetes Operators

“Operators are software extensions to Kubernetes that make use of **custom resources** to manage applications and their components. Operators follow Kubernetes principles, notably **the control loop**.”

<https://kubernetes.io/docs/concepts/extend-kubernetes/operator/>



Postgres Operators for Kubernetes




Agenda

1. What Came Before?
2. What does a DBA do anyway?
3. Postgres Operators for Kubernetes
4. **What do you want from an Operator?**
5. How to get Started



What do you want from an Operator?



~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~● DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance



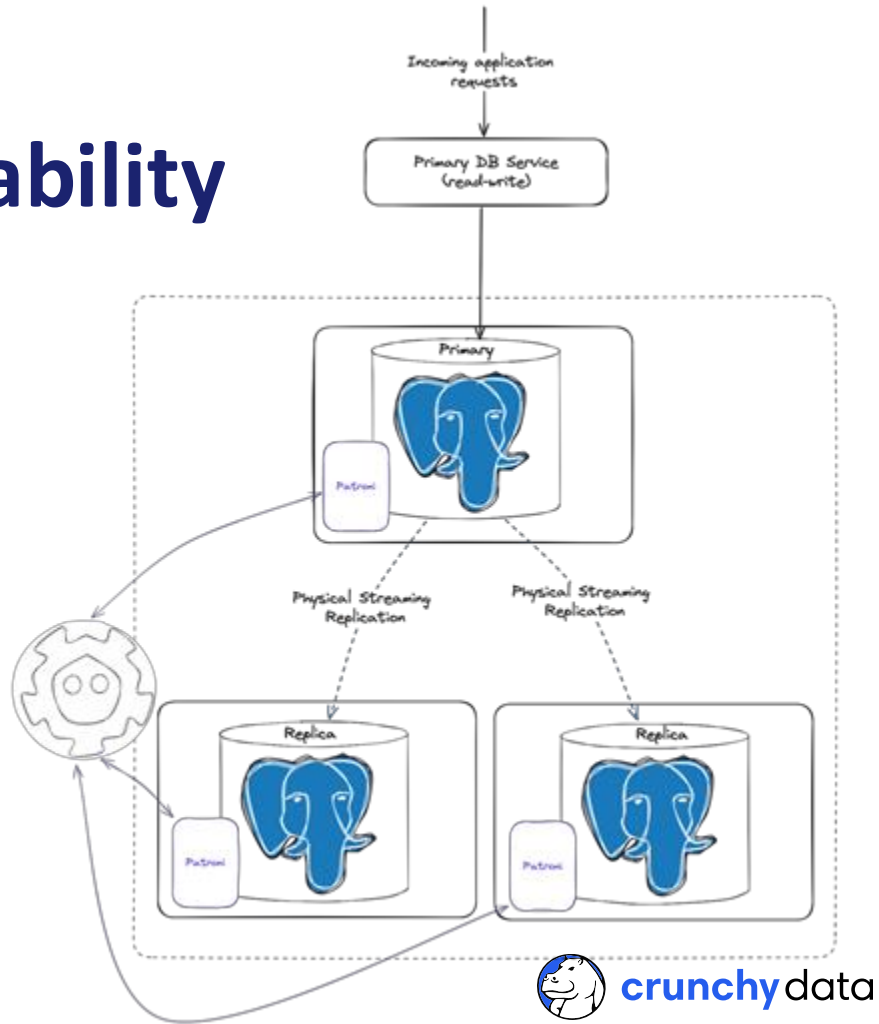
~~DBA~~ Operator Responsibilities

- **Database (High) Availability**
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance



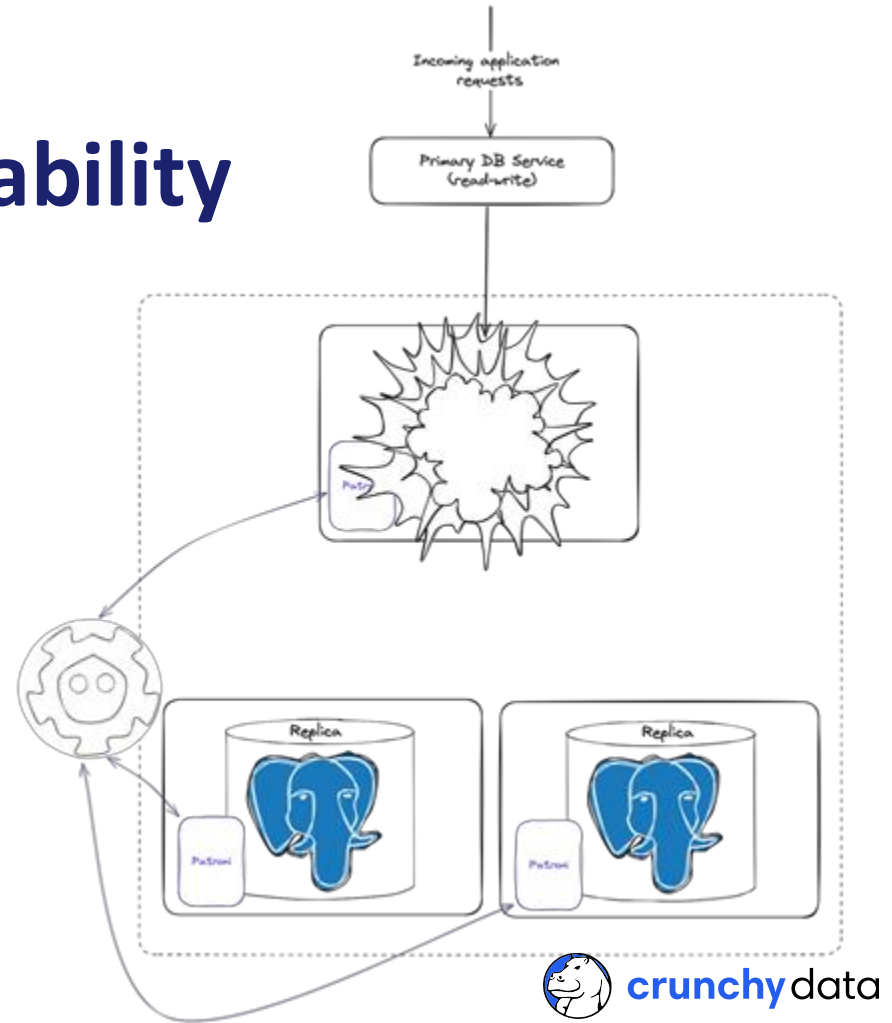
DBA Operator Responsibilities

Database (High) Availability



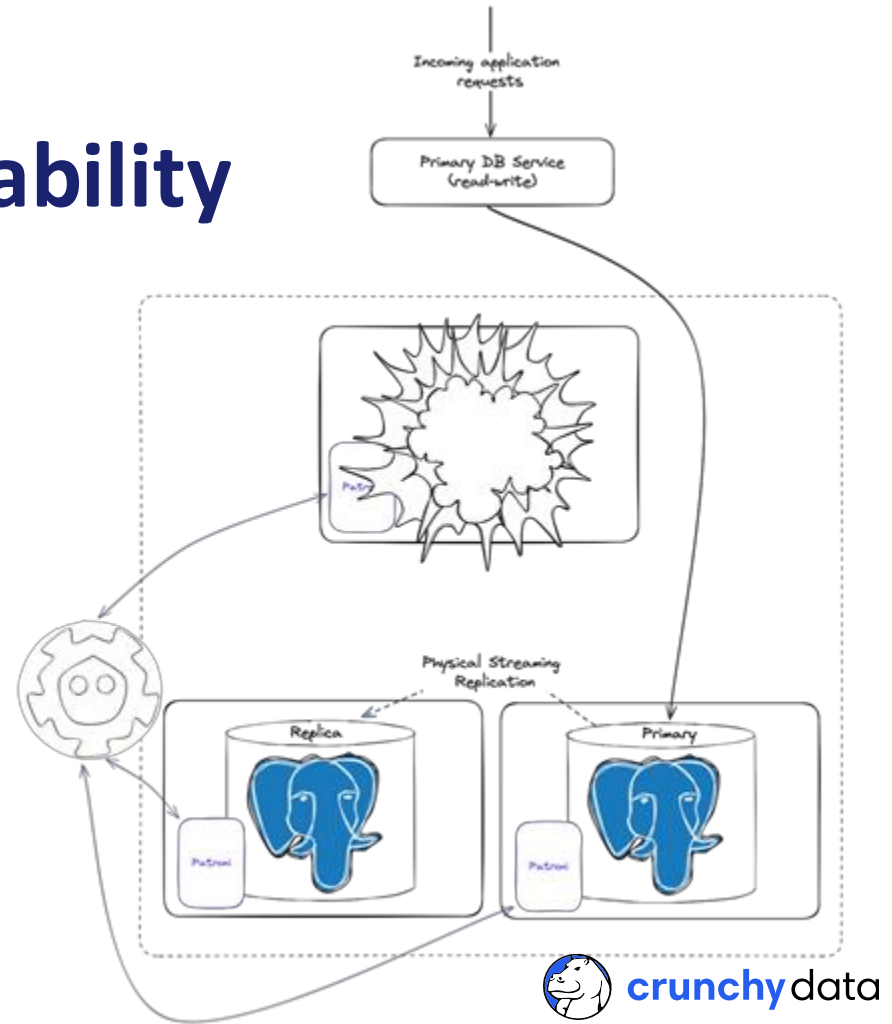
DBA Operator Responsibilities

Database (High) Availability



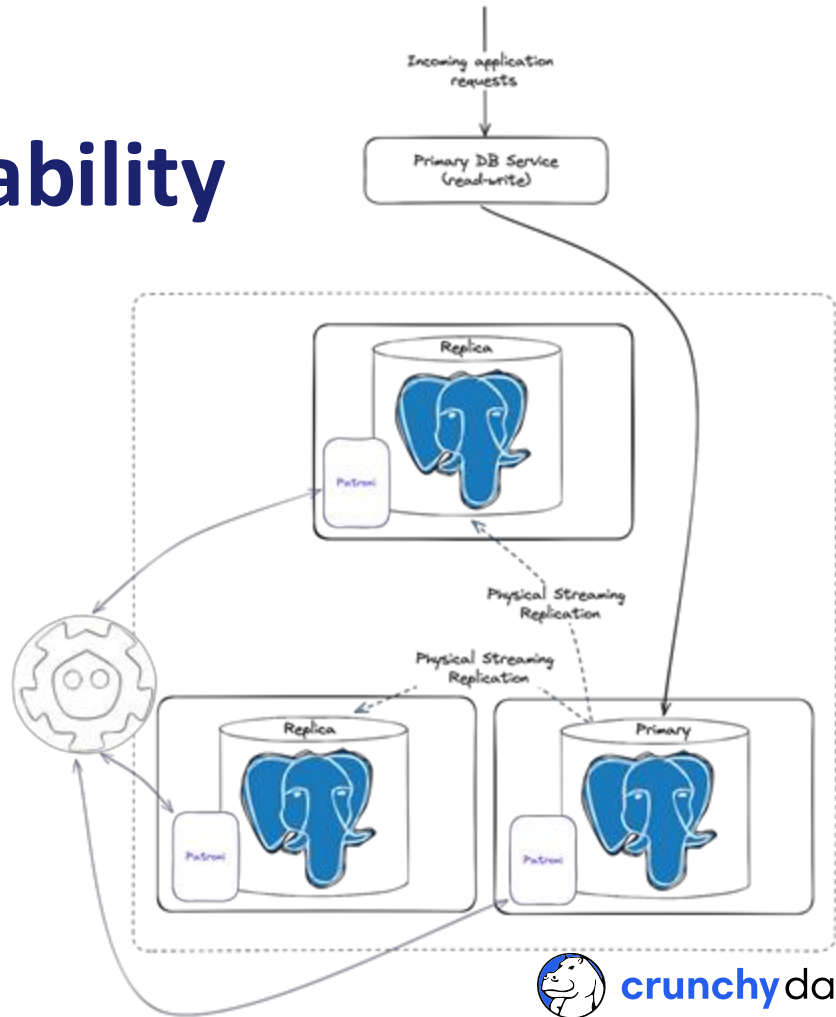
DBA Operator Responsibilities

Database (High) Availability



DBA Operator Responsibilities

Database (High) Availability




~~DBA~~ Operator Responsibilities

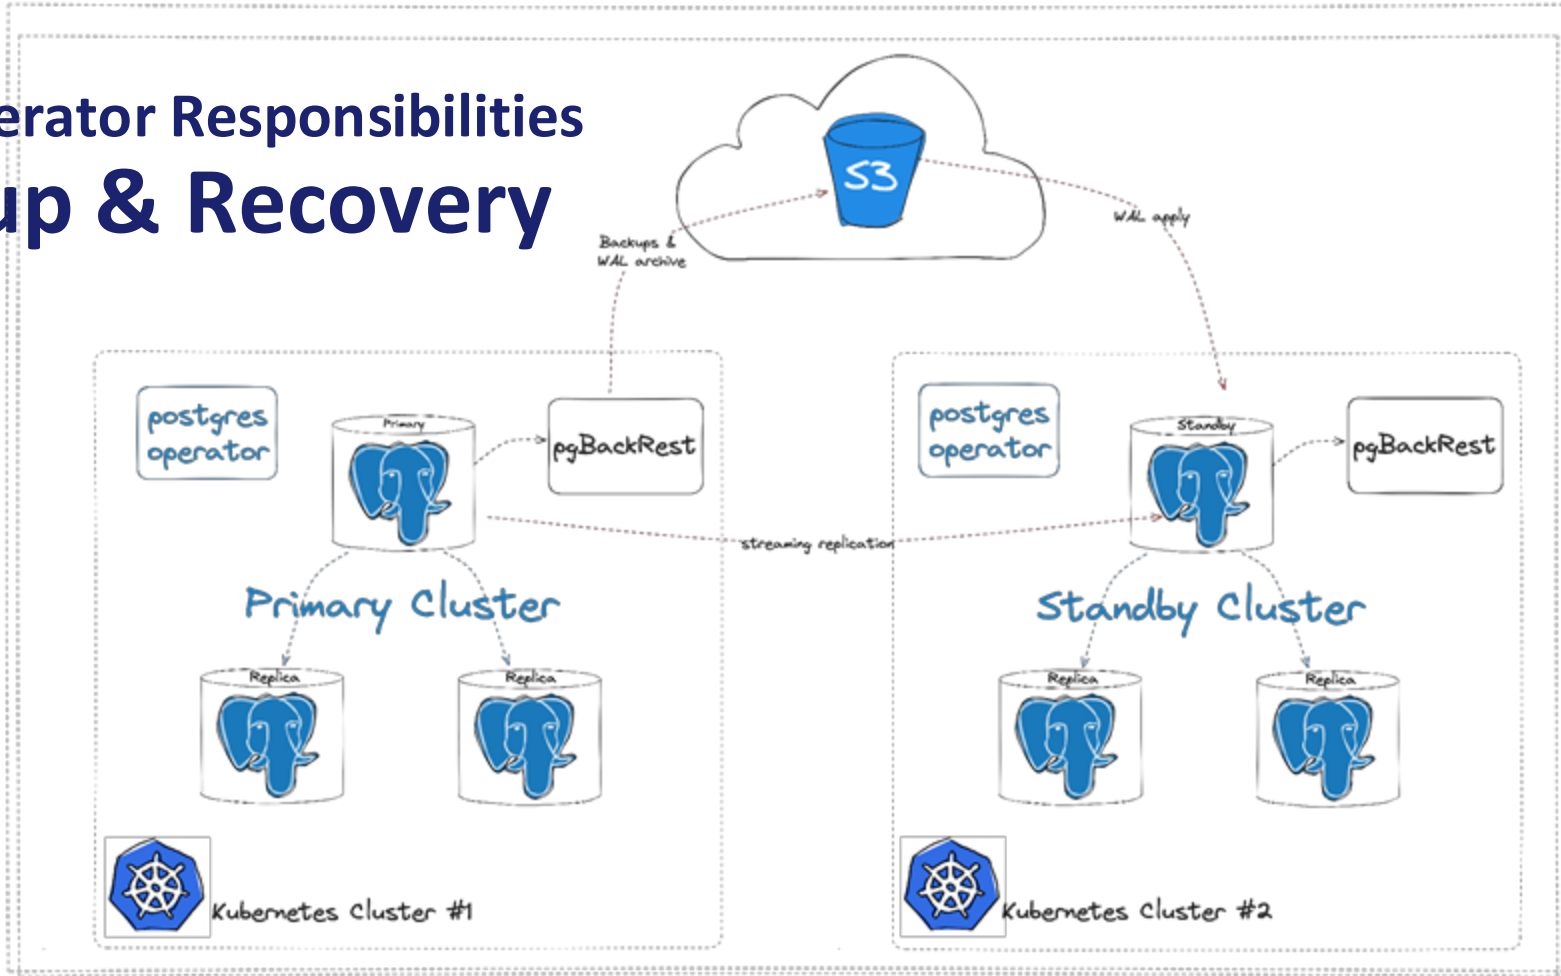
- Database (High) Availability
- **Backup and Recovery**
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance

DBA Operator Responsibilities

Backup & Recovery



DBA Operator Responsibilities Backup & Recovery



~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- **Security & Data Protection**
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance

DBA Operator Responsibilities

Security


- Database Access
- Password Encryption
- Secrets
- SSL/TLS
- Certificates

~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- **Monitoring**
- ~~● DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance

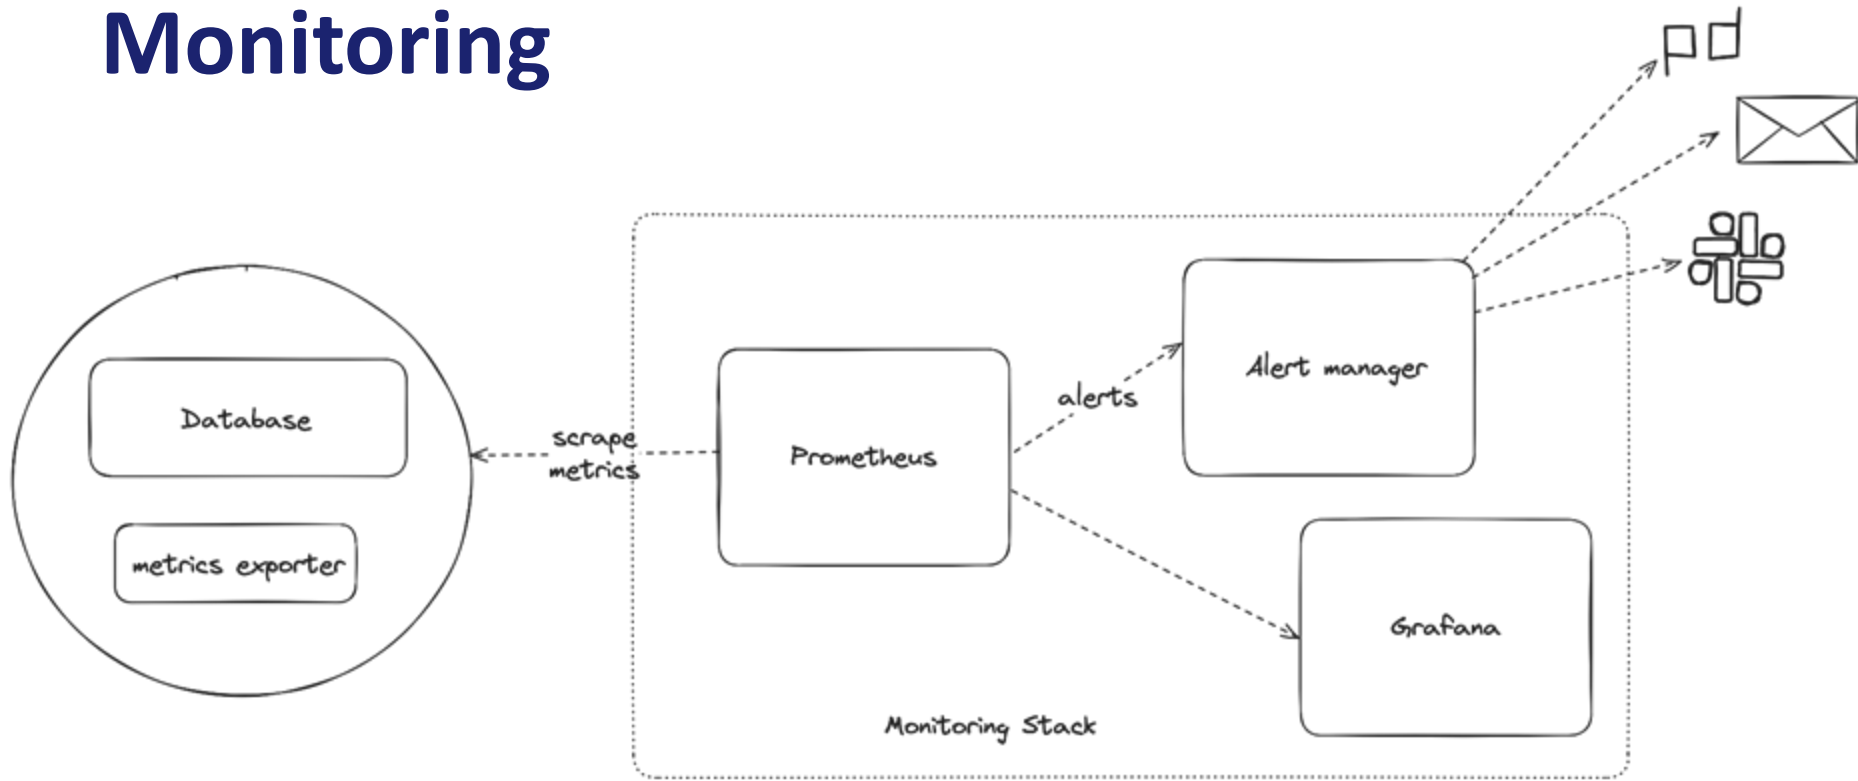
~~DBA~~ Operator Responsibilities

Monitoring



DBA Operator Responsibilities


Monitoring



~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- **Support/Troubleshooting**
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance

~~DBA~~ Operator Responsibilities Support/Troubleshooting




~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- **DB Software install/upgrade**
- Database Expertise
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance



~~DBA~~ Operator Responsibilities


Database Software Install



DBA Operator Responsibilities

Postgres (Minor) Upgrade


Ordered, automated rolling updates



~~DBA~~ Operator Responsibilities

Database Software (Major) Upgrade

- Requires planning & testing
- Automated by Operator



~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- **Database Expertise**
- Performance Tuning
- Capacity Planning
- Database Creation
- Database Maintenance


~~DBA~~ Operator Responsibilities

Database Expertise

- Operator: Built-in database expertise

Human expert:

- Strategic considerations
- Application needs
- Business requirements




~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- **Performance Tuning**
- Capacity Planning
- Database Creation
- Database Maintenance

~~DBA~~ Operator Responsibilities

Performance Tuning

- Configuration Parameters
- Connection pooling
- `pg_stat_statements`
- Logging
- Monitoring and alerting



~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- **Capacity Planning**
- Database Creation
- Database Maintenance

DBA Operator Responsibilities

Capacity Planning

```
dataVolumeClaimSpec:  
  accessModes:  
    - "ReadWriteOnce"  
  resources:  
    requests:  
      storage: 1Gi  
      storage: 10Gi
```


~~DBA~~ Operator Responsibilities

- Database (High) Availability
- Backup and Recovery
- Security & Data Protection
- Monitoring
- ~~DB Design/Data Modelling~~
- Support/Troubleshooting
- DB Software install/upgrade
- Database Expertise
- Performance Tuning
- Capacity Planning
- **Database Creation**
- **Database Maintenance**



~~DBA~~ Operator Responsibilities

Database Creation

```
users:  
- name: hippo  
  databases:  
    - dokday
```



~~DBA~~ Operator Responsibilities

Database Maintenance

- Taking backups
- Creating users
- Rebuilding Indexes
- Gathering statistics
- Log management

Not an exhaustive list!



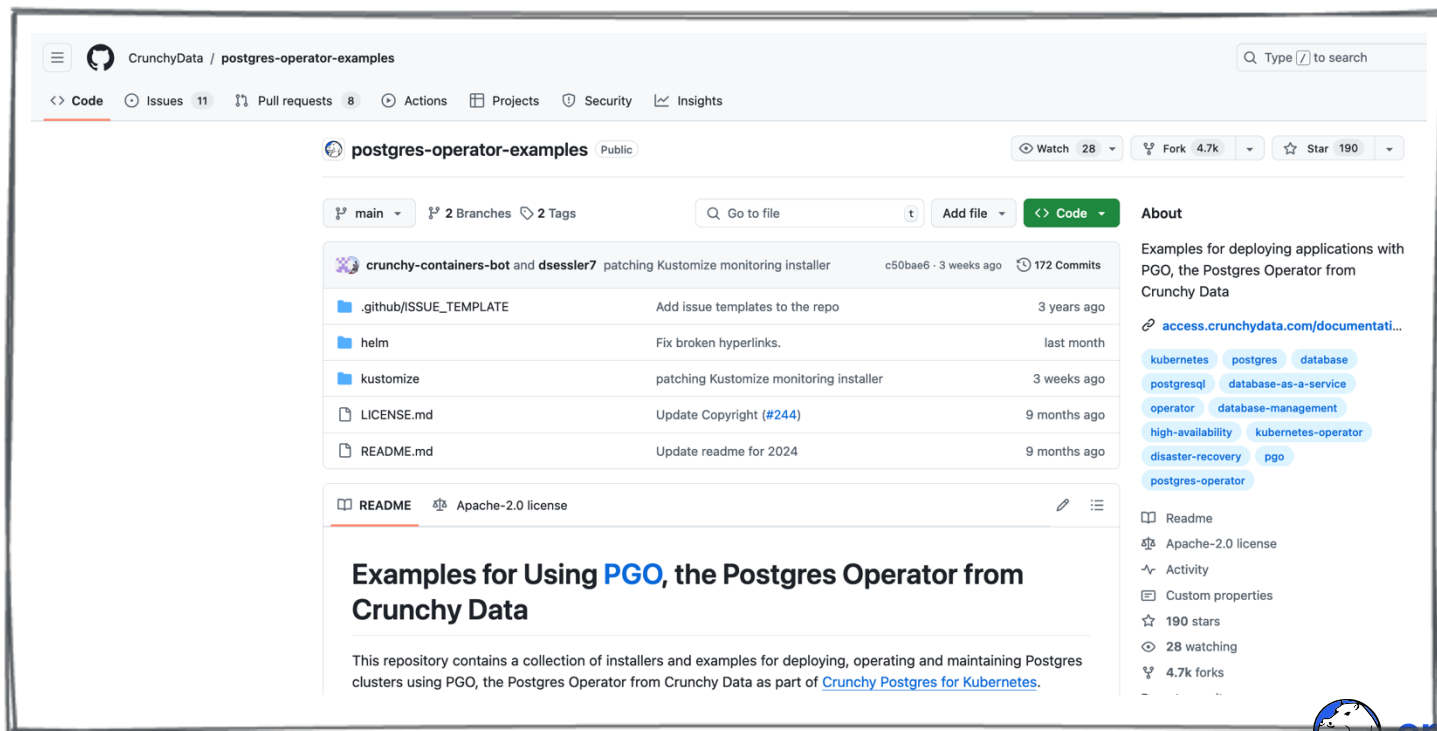
Agenda

1. What Came Before?
2. What does a DBA do anyway?
3. Postgres Operators for Kubernetes
4. What do you want from an Operator?
5. **How to get Started**



How to get Started

<https://github.com/CrunchyData/postgres-operator-examples>



The screenshot shows the GitHub repository page for `postgres-operator-examples` under the `CrunchyData` organization. The repository is public and has 28 watchers, 4.7k forks, and 190 stars. The main branch is selected, and there are 2 branches and 2 tags. The repository description states: "Examples for deploying applications with PGO, the Postgres Operator from Crunchy Data". The file list includes `.github/ISSUE_TEMPLATE`, `helm`, `kustomize`, `LICENSE.md`, and `README.md`. The `README.md` file is selected, showing the title "Examples for Using PGO, the Postgres Operator from Crunchy Data" and a description: "This repository contains a collection of installers and examples for deploying, operating and maintaining Postgres clusters using PGO, the Postgres Operator from Crunchy Data as part of [Crunchy Postgres for Kubernetes](#)." The right sidebar shows the repository's activity, including 190 stars, 28 watching, and 4.7k forks. The CrunchyData logo is visible in the bottom right corner.

How to get Started

Install the Operator

```
% kubectl apply --server-side -k kustomize/install/default
```

```
customresourcedefinition.apiextensions.k8s.io/pgadmins.postgres-  
operator.crunchydata.com serverside-applied  
customresourcedefinition.apiextensions.k8s.io/pgupgrades.postgres-  
operator.crunchydata.com serverside-applied  
customresourcedefinition.apiextensions.k8s.io/postgresclusters.postgres-  
operator.crunchydata.com serverside-applied  
serviceaccount/pgo serverside-applied  
clusterrole.rbac.authorization.k8s.io/postgres-operator serverside-applied  
clusterrolebinding.rbac.authorization.k8s.io/postgres-operator serverside-  
applied  
deployment.apps/pgo serverside-applied
```

postgres.yaml

```
apiVersion: postgres-operator.crunchydata.com/v1beta1
kind: PostgresCluster
metadata:
  name: kubecon
spec:
  postgresVersion: 16
  instances:
    - name: instance1
      replicas: 3
      dataVolumeClaimSpec:
        accessModes:
          - "ReadWriteOnce"
        resources:
          requests:
            storage: 1Gi
```



postgres.yaml (contd)

```
backups:  
  pgbackrest:  
    repos:  
      - name: repo1  
        volume:  
          volumeClaimSpec:  
            accessModes:  
              - "ReadWriteOnce"  
            resources:  
              requests:  
                storage: 1Gi
```



How to get Started

Create a Postgres Cluster

```
% kubectl apply -k kubecon
```

```
% kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubecon-backup-b7v4-krfkb	0/1	Completed	0	13m
kubecon-instance1-b9nq-0	4/4	Running	0	13m
kubecon-instance1-ljg8-0	4/4	Running	0	13m
kubecon-instance1-m6mx-0	4/4	Running	0	13m
kubecon-repo-host-0	2/2	Running	0	13m
pgadmin-611942b8-585c-4dea-886a-6774a1eccd7a-0	1/1	Running	0	45m
pgo-6fdd547449-6ghr9	1/1	Running	0	46m



Conclusions

Conclusions

Postgres Operator for Kubernetes

- Virtual database administrator
- PostgreSQL expertise and automation
- Robust, secure, scalable architecture
- Combined strength of PostgreSQL and Kubernetes
- Leaves you to do the Interesting Stuff!





Thank You!

Karen Jex | @karenhjex | karen.jex@crunchydata.com