




KubeCon



CloudNativeCon

North America 2024



CLOUD NATIVE &  
KUBERNETES

**AI DAY**

NORTH AMERICA

# LLM Powered Agents with Kubernetes

Hema Veeradhi  
Shrey Anand

Red Hat  
*Emerging Technologies (OCTO)*



# Agenda

- **Introduction to LLMs**
- **Introduction to RAG**
- **What are Agents**
  - Components of LLM agents
- **Types of Agents**
  - Router based
  - ReAct
- **Agent Use Case**
- **Demo**
- **Concluding Remarks**



KubeCon



CloudNativeCon

North America 2024



Hema Veeradhi  
Principal Data Scientist  
Emerging Technology | Red Hat



California, United States



hemaveeradhi

Hi! 🖐️



Shrey Anand  
Principal Data Scientist  
Emerging Technology | Red Hat



Toronto, Canada



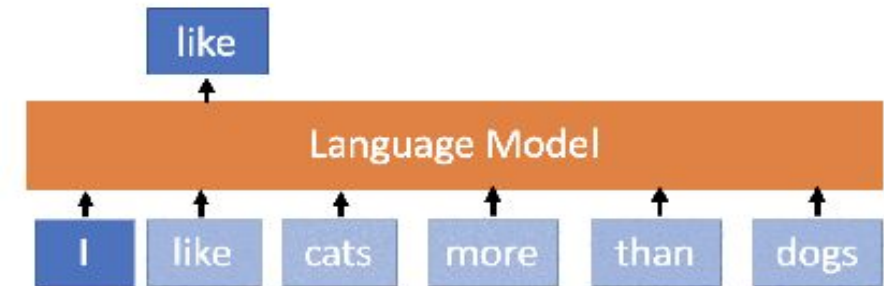
shrey-anand

# Introduction to LLMs

Large Language Models are **computational models that are capable of modeling and generating human language**. LLMs have the transformative ability to predict the likelihood of word sequences or generate new text based on a given input.

## What are LLMs good at?

- Text generation/code generation
- Chatbots and Conversational AI
- Information retrieval
- Sentiment analysis



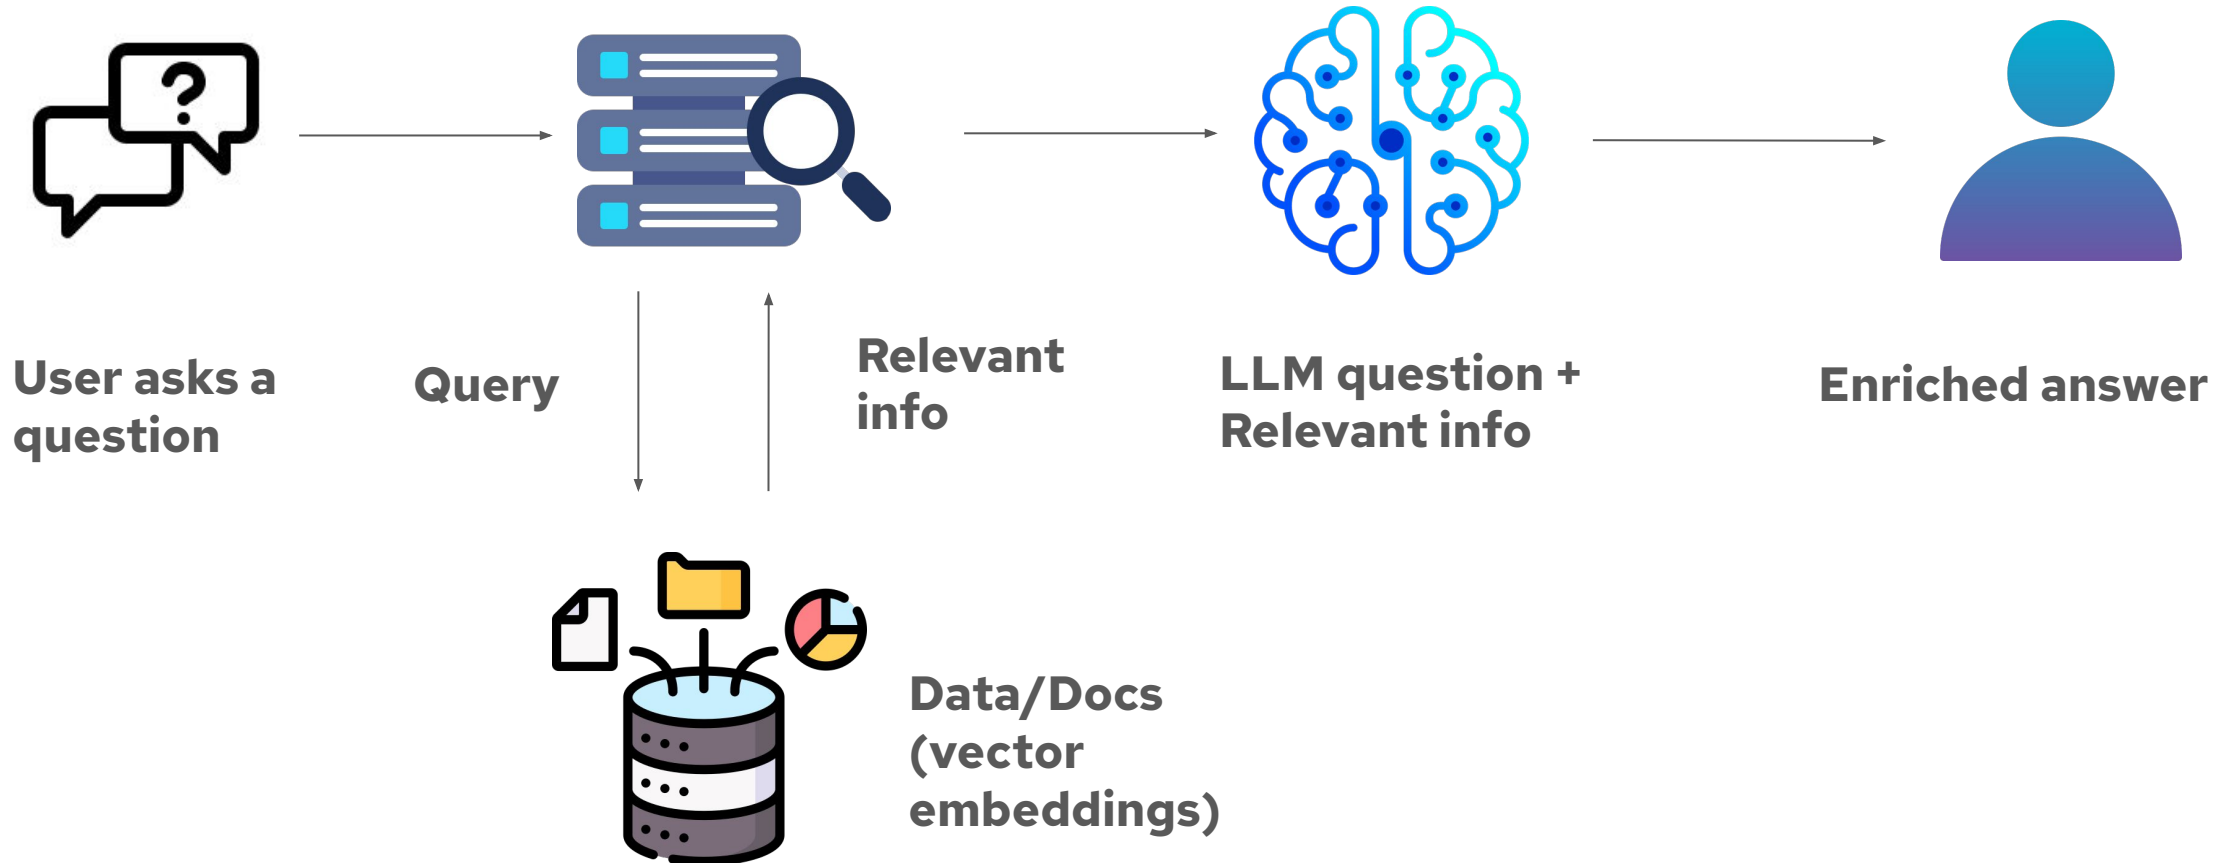
Input: n “tokens” -> I, like, cats, more, than

Output: 1 “token” -> dogs

# When can LLMs fail?

- **Complex reasoning tasks** - LLMs have limited reasoning capability, LLMs are good knowledge retrievers but not good reasoners
- **No Dynamicity** - LLMs are static and unable to access real-time information
- **Limited Knowledge (hallucination)** - While trained on vast data, LLMs lack up-to-date world knowledge

# RAG - Retrieval Augmented Generation





# LLM Agents

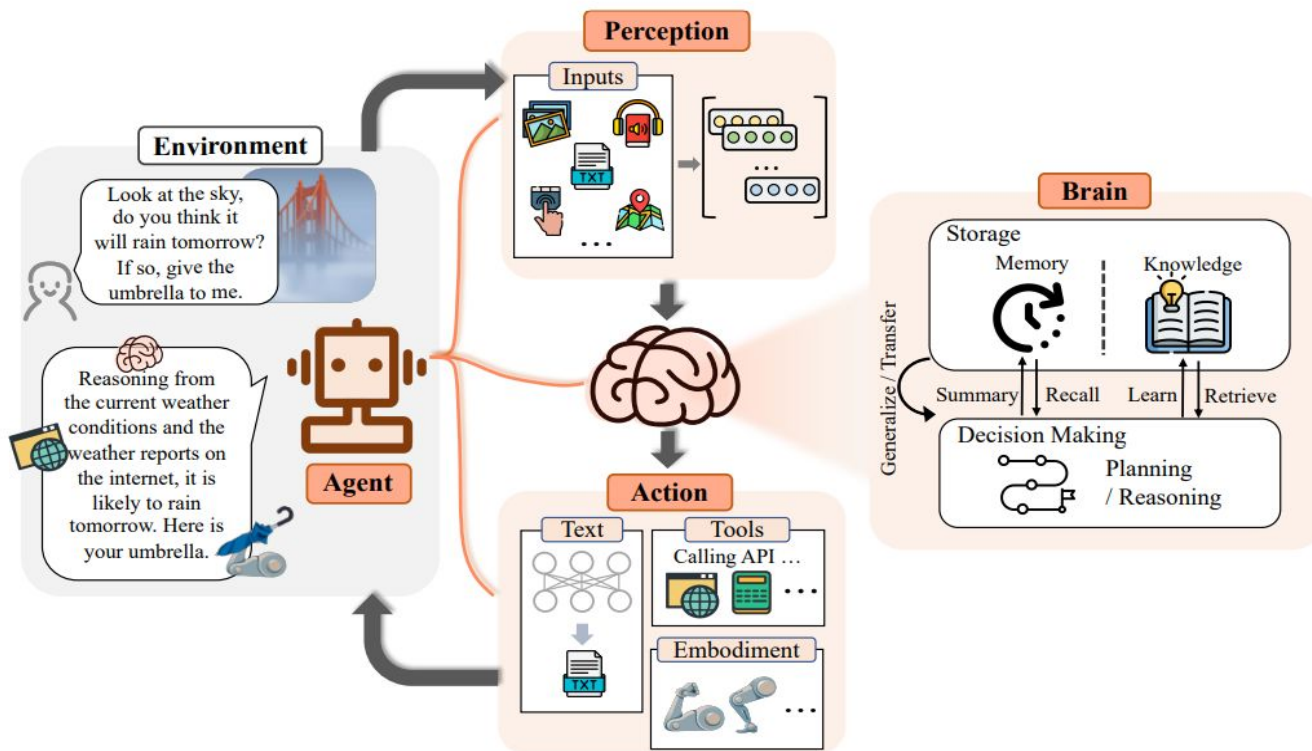
LLM Agents, also known as Large Language Model agents, **leverage LLMs to execute complex tasks by integrating them with essential components like planning and memory.**

**LLM Agent = LLM + Tools + State**

- **LLM** - Computational engine i.e “brain”
- **Tools** - Agents ability to interact with the external world
- **Memory/State** - Agent’s memory of previous messages and results from used tools



# Components of an LLM Agent



- **LLM** - Computational engine i.e. "brain"
- **Planning** - Chain of thought process (CoT) to create a plan for executing the tasks
- **Tools** - Executable functions, APIs, other services that complete various tasks
- **Memory** - Short term memory to retain the agent's train of thought and long term memory to retain context/conversation history
- **Actions** - Agents perform actions based on their environment and reasoning, adapting and solving tasks iteratively through feedback.



# Types of Agents

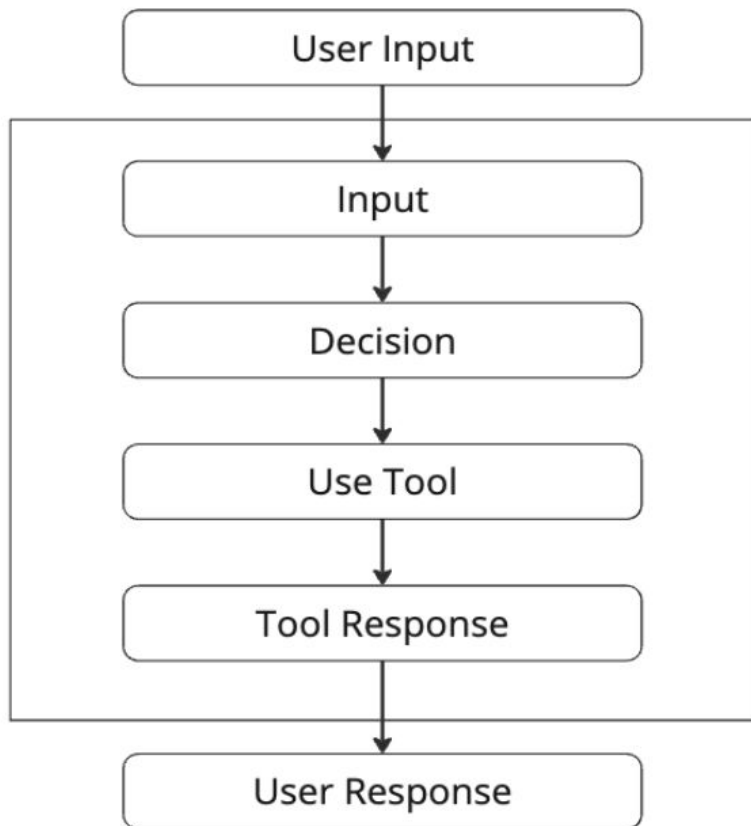
There are many different types of LLM agents to choose from, depending on the nature of your use case, including:

- React Agent
- Rule-based Agent
- Single/Multi Agent

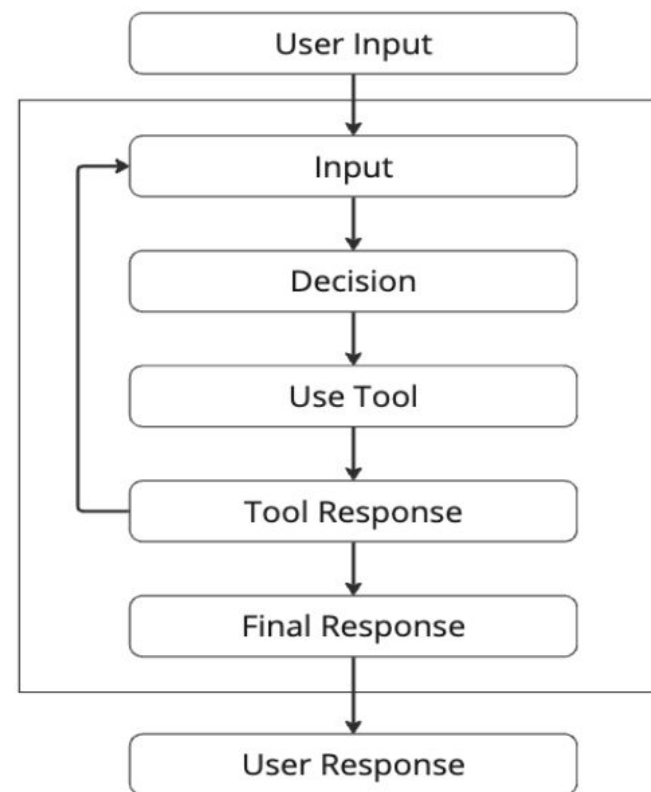
## Tools and frameworks to build LLM agents

- [LangChain](#)
- [Bee Agent](#)
- [OpenAgents](#)

# Router Agent



# ReAct Agent



# Use Case

- Imagine a **fictional cloud company CloudForge Dynamics that has several departments and is looking to have a unified chat assistant** for its employees to answer queries across these departments
- Let's say we have access to Products, HR, and Customer accounts documents
- How can we answer questions by gathering/combining information from all of these different sources?
  - A LLM agent can decide which department's sources should be used for answering the query




KubeCon



CloudNativeCon

North America 2024



**CLOUD NATIVE &  
KUBERNETES**

**AI DAY**

**NORTH AMERICA**

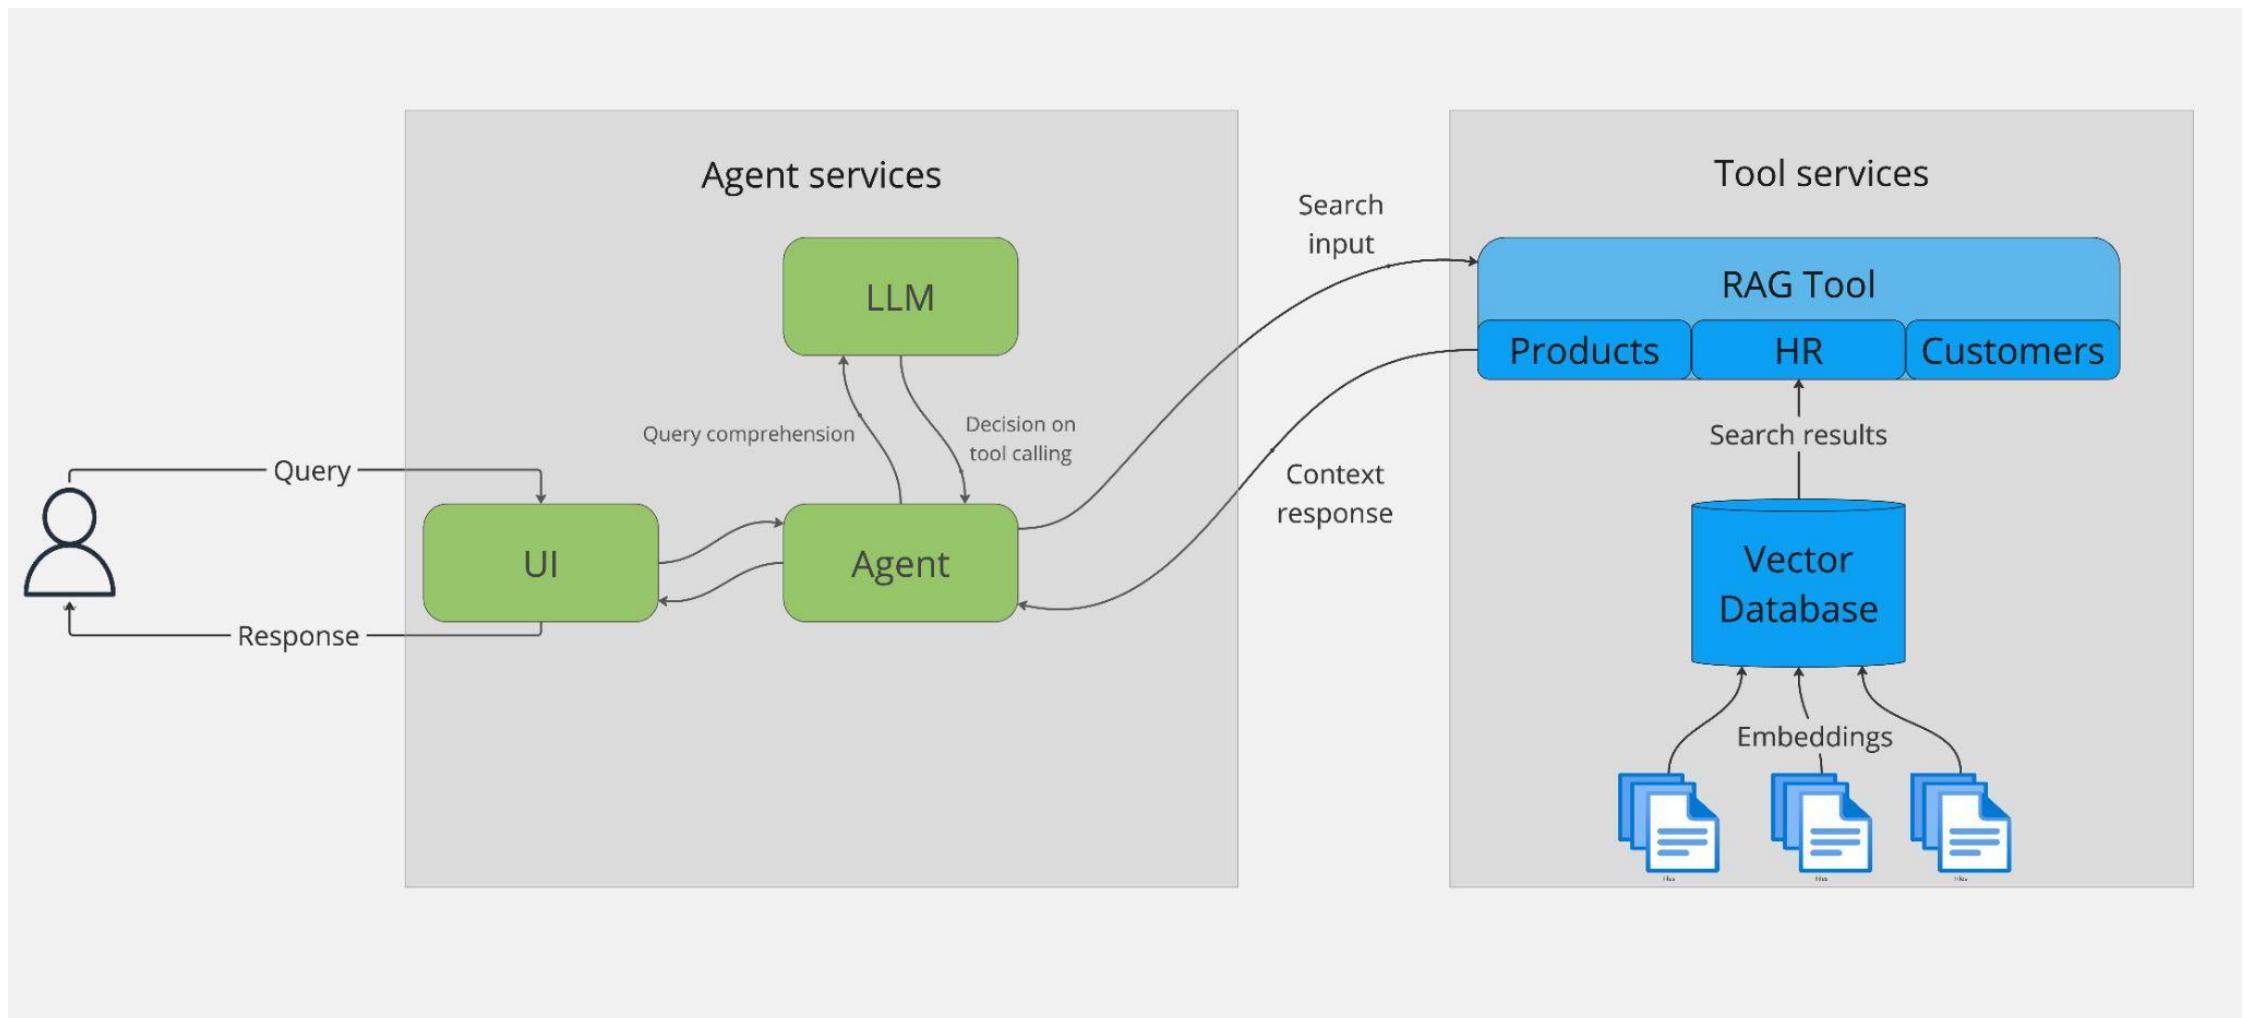
# DEMO



## Example Queries

- What does the CloudForge Migrate product do?
- What HR things should I do before the start date?
- Total Payments Received from FinNova Bank?

# Architecture



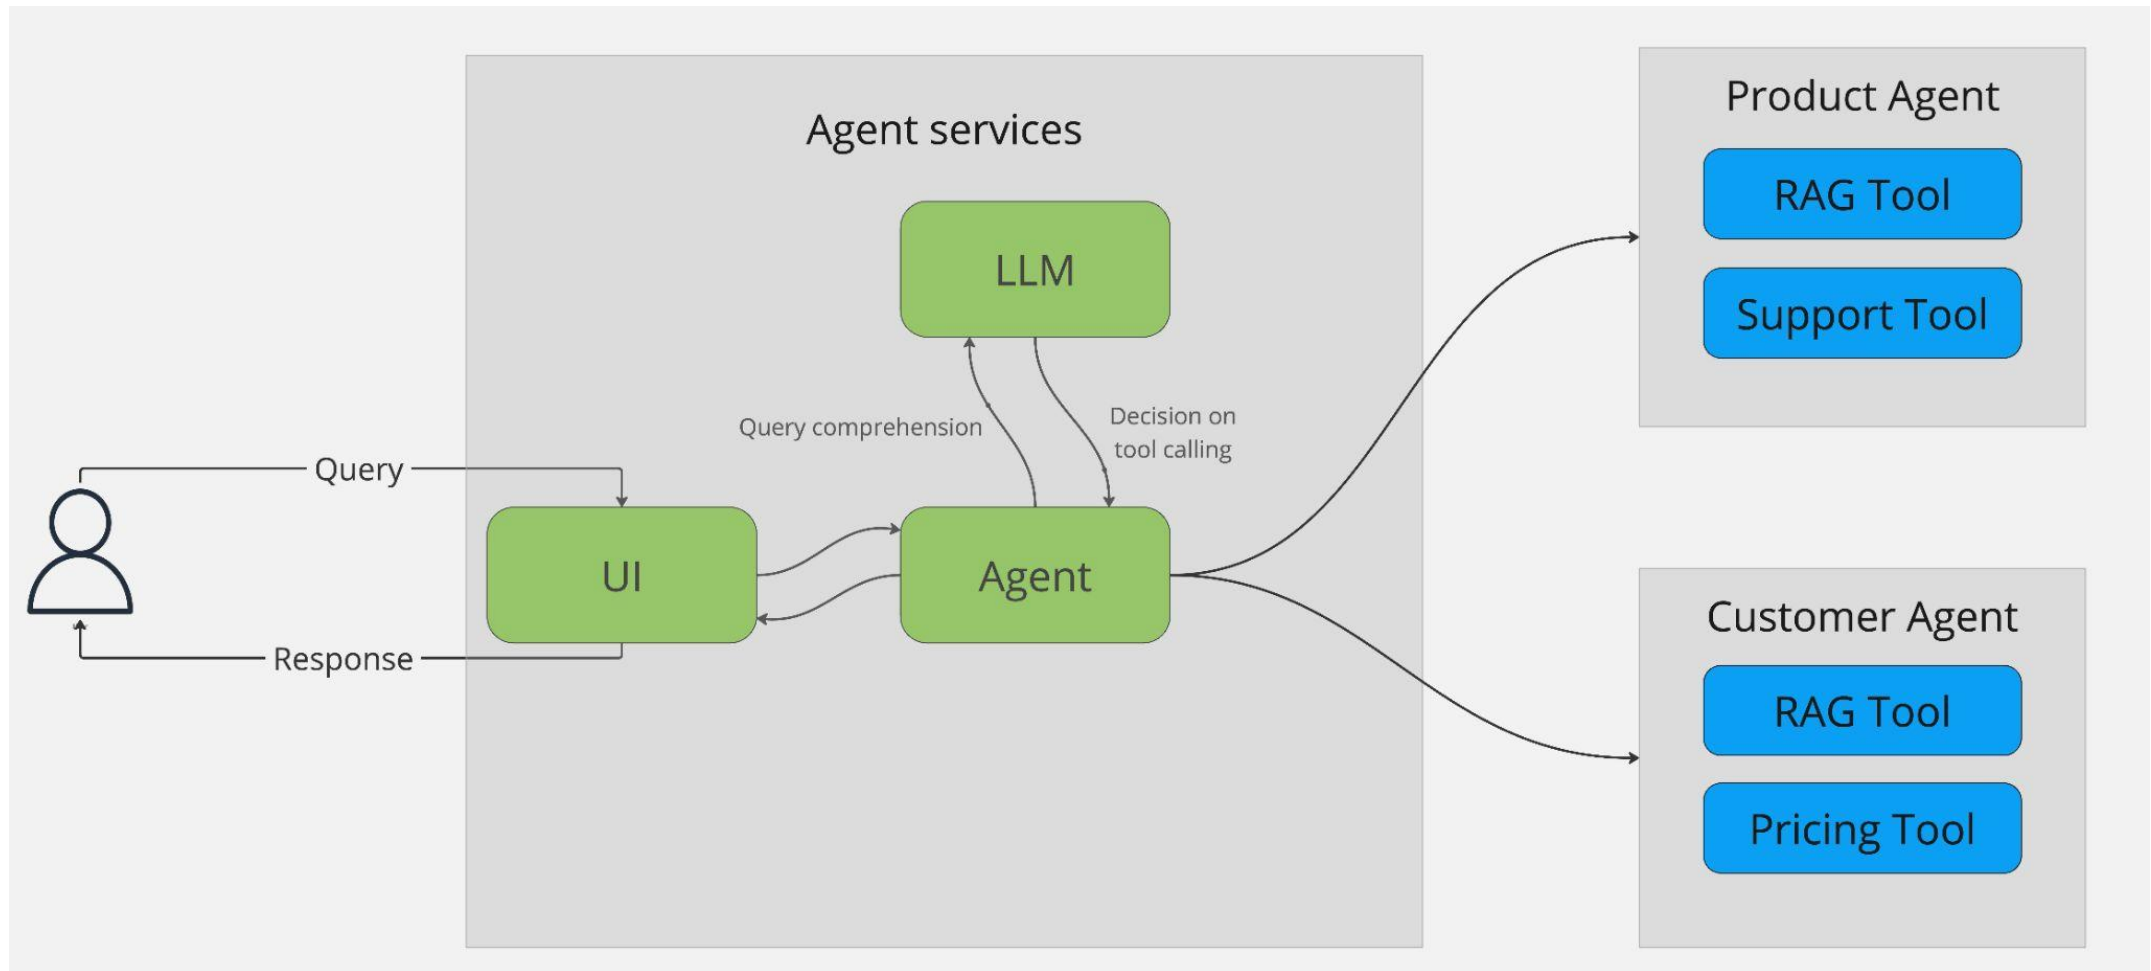
# Resources

- **GitHub Repository** - <https://github.com/redhat-et/llm-agents>
- **Agent Services** (UI, Agent, LLM deployment files)
  - **Streamlit UI** - <https://github.com/redhat-et/llm-agents/tree/main/streamlit/openshift>
  - **ReAct Agent** - [https://github.com/redhat-et/llm-agents/tree/main/react\\_agent/openshift](https://github.com/redhat-et/llm-agents/tree/main/react_agent/openshift)
- **RAG Tool Services**
  - **Milvus vector database** - [https://github.com/redhat-et/llm-agents/blob/main/rag/vector\\_db](https://github.com/redhat-et/llm-agents/blob/main/rag/vector_db)
  - **RAG service** - <https://github.com/redhat-et/llm-agents/blob/main/rag>





# Discussion





## Concluding Remarks

- LLMs are awesome, but they need RAG for grounding
- LLMs + RAG pattern is awesome, but it needs Agent's framework for planning and interacting with external sources
- Challenges with the agents approach:
  - Debugging is difficult as there can be many fault points such as selection of the right tool and input parameters
  - Latency of response is high if there are many API calls involved in creating the response
  - Load issues on external tools, and security concerns of limiting user access to data, with each tool having to handle access controls on its own.
- Further advancements are being made to develop more robust, scalable and reproducible open source agent frameworks (Eg: [CrewAI](#), [Bee Agent framework](#), [Arch](#))



**KubeCon**



**CloudNativeCon**

**North America 2024**



# Thank You!

## Q?



[hveeradh@redhat.com](mailto:hveeradh@redhat.com)  
[shanand@redhat.com](mailto:shanand@redhat.com)