

Observability Day

NORTH AMERICA

Perfect Match: Correlating Continuous Profiling with Distributed Tracing for Stronger Observability

Jonas Kunz & Christos Kalkanis

What is Profiling



Metrics



Logs



Traces



Profiles

- The **fourth** pillar of Observability
- Different sources of profiling data: On-CPU, Off-CPU, Memory
- Instrumentation vs **Sampling**
- Sampling CPU profilers interrupt CPU at regular intervals and record execution state
- An **On-CPU** profiler records stacktraces for every CPU core that gets interrupted
- **StackTrace**: Snapshot of function calls at a point in time
- StackTraces can span **kernel**, **native** application and **higher-level** runtime code
- **eBPF** can enable low-overhead continuous profiling in production without application instrumentation
- **Continuous** profiling offers deeper level of visibility, can expose **unknown-unknowns**

Benefits of Continuous Profiling

Optimize Compute
Efficiency

Quickly Identify
Performance
Bottlenecks

Save on Cloud Cost

Reduce CO2
Footprint

“If one saves 20% on 800 servers, and assumes 300W power consumption, that one-line code change is worth 160 metric tons of CO2 saved per year.”

Thomas Dullien

Where is my cloud spend going and where can we generate savings?

FinOps

There's an incident happening, what's going on and why is it so slow?

SRE

Continuous Profiling

How can I make my application faster and more efficient?

Developer

Low-Overhead Continuous Profiling

- Optimize.cloud launches low-overhead multi-runtime zero-instrumentation [profiler](#) in 2021
- Acquired by Elastic soon after
- Donated to OpenTelemetry in 2024
- Continued development and evolution



[opentelemetry-ebpf-profiler](#)

ARM64

x86-64



Unobtrusive. Frictionless Deployment

Powered by eBPF. Requires **zero-instrumentation**, no code changes or app restarts. **Gain faster ROI.**



Whole-System Visibility

Unlock unknown-unknowns - from the kernel through userspace into high-level code, across **multi-cloud workloads**.



Polyglot Visibility

C/C++, Rust & Go (without debug symbols on host)
PHP, Python, Java (or any JVM language), Ruby, DotNet, Perl & NodeJS.

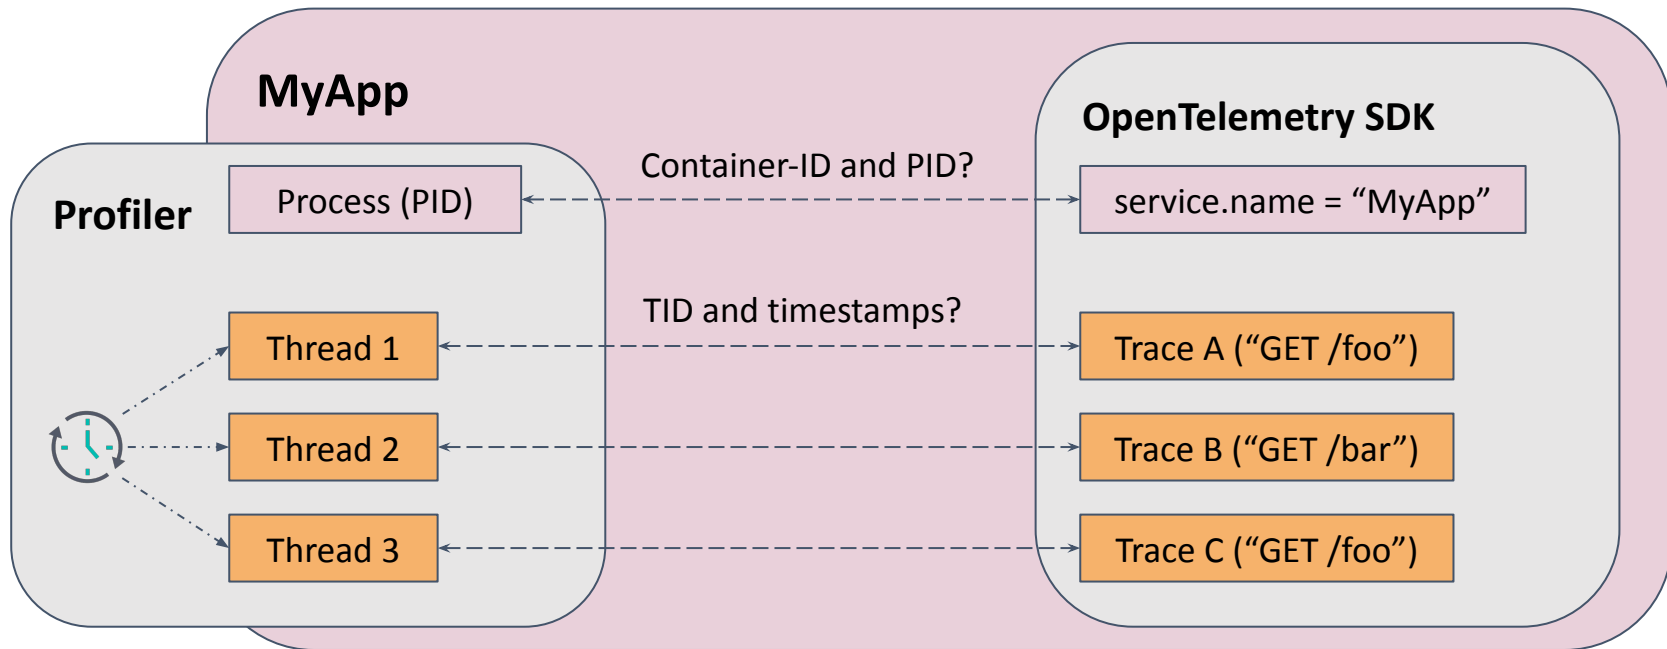


Extremely Low Overhead

Continuous profiling in **production with negligible overhead**.

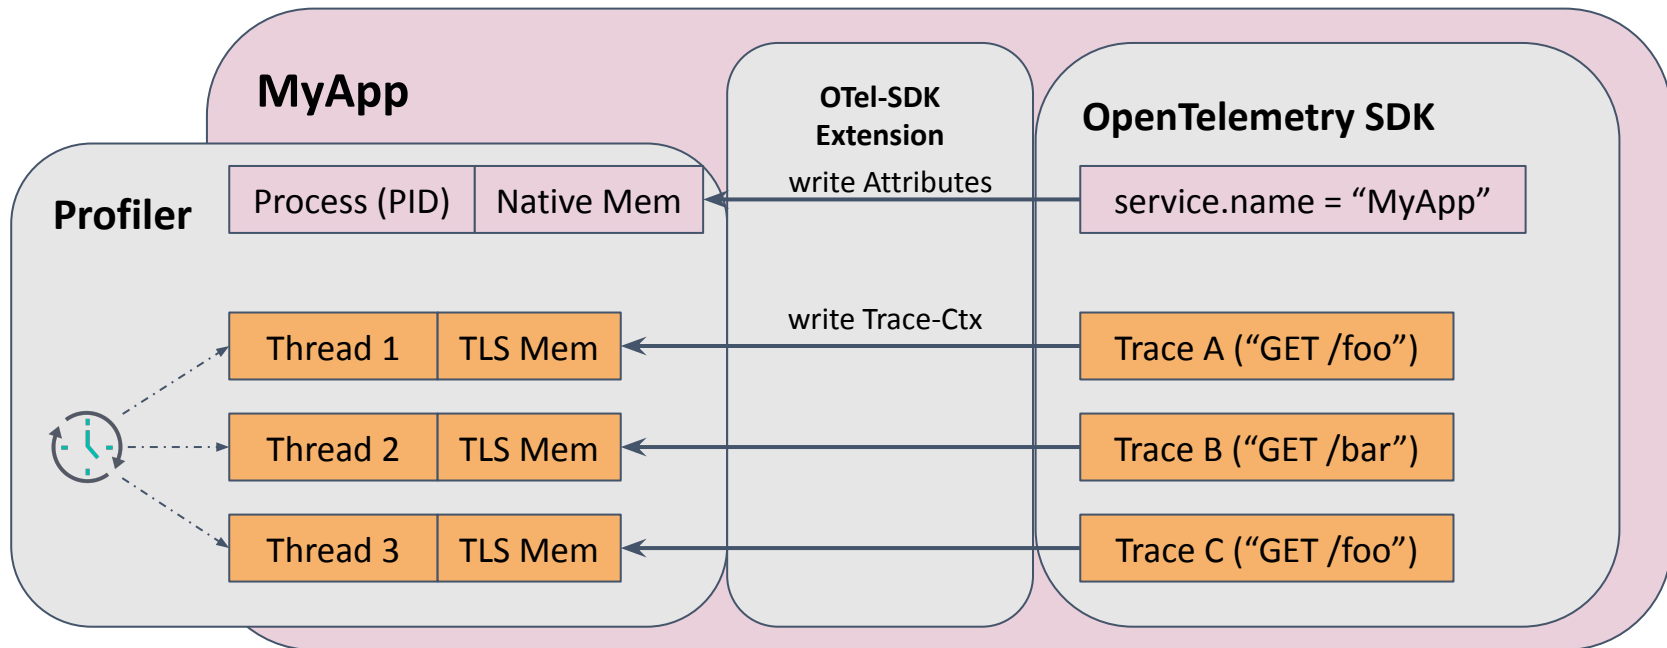
Typical case: < 1% system CPU, ~250MB of RAM

The Correlation Gap



- Root cause for CPU Usage of "MyApp"?
- Where does "GET /foo" spend its CPU-time?

The Correlation Gap Approach



- Root cause for CPU Usage of "MyApp"?
- Where does "GET /foo" spend its CPU-time?

Java SpanProcessor available on [GitHub](#):

```
Resource resource = Resource.builder()
    .put(ResourceAttributes.SERVICE_NAME, "my-service").build();

UniversalProfilingProcessor processor =
    UniversalProfilingProcessor.builder(exportingProcessor, resource).build();

SdkTracerProvider tracerProvider = SdkTracerProvider.builder()
    .addSpanProcessor(processor).build();
```

Loads a native JNI library which

- Populates a native, global var with resource-attributes
- Allocates native TLS
- Keeps the native TLS in sync on OTEL context changes



**The best telescopes
to see the world
closer**

Go Shopping



<https://opentelemetry.io/ecosystem/demo/>

- AdService: Java App for serving advertisements
 - Random Ads
 - Targeted Ads: (E.g. for “telescopes”, “binoculars”, ...)
 - Type recorded as span-attributes
- Problems:
 - Built-in: Background CPU-usage
 - Added:
 - Background allocations for GC-pressure
 - Code-path dependent CPU bottlenecks

Questions?