

# SkyRay: Seamlessly Extending KubeRay to Multi-Cluster Multi-Cloud Operation

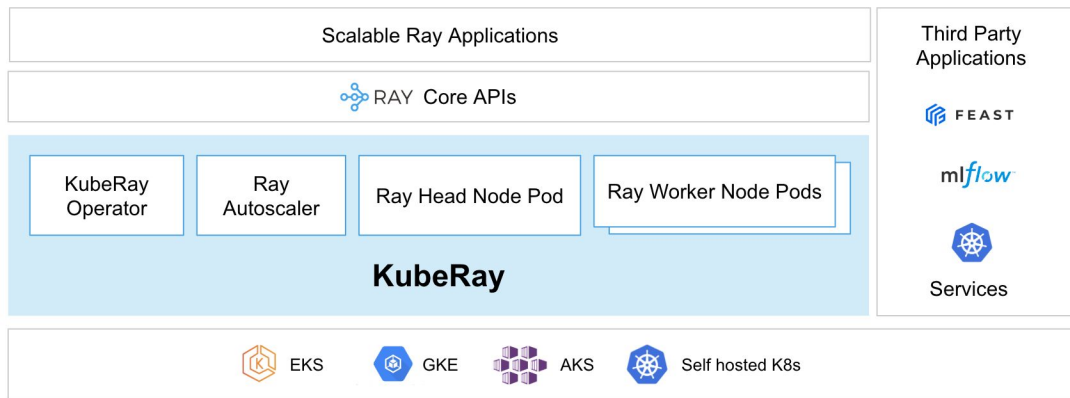
Anne Holler, Chief Scientist, Elotl  
Cloud Native & Kubernetes AI Day, 11/12/24

# SkyRay Overview: Ray and KubeRay Overview

**Ray**: unified framework that can scale ML/AI applications from laptop to cluster



**KubeRay**: creation, deletion, scaling of Ray clusters/jobs/services on K8s cluster



## KubeRay Structure

# SkyRay Overview: Multi-cluster Multi-cloud K8s Sky



Reasons to use **multiple K8s clusters** to group resources by characteristics include:

- **Service continuity**, e.g., limit impact of region outage
- **Workload purpose**, e.g., production (higher QoS) vs development (lower QoS)
- **Resource availability and cost**, e.g., per region GPU availability, per region pricing
- **Geo-location**, e.g., to satisfy workload location restrictions or to reduce response latency
- **Usage lifetime**, e.g., ML training or RAG ingestion may be periodic, ML serving is ongoing

Reasons to use K8s clusters from **multiple cloud vendors** include:

- **Avoids cloud vendor lock-in**
- **Supports customer cloud choice**
- **Allows differentiated costs**, e.g., fixed (on-premise, reserved) and dynamic (on-demand or spot)

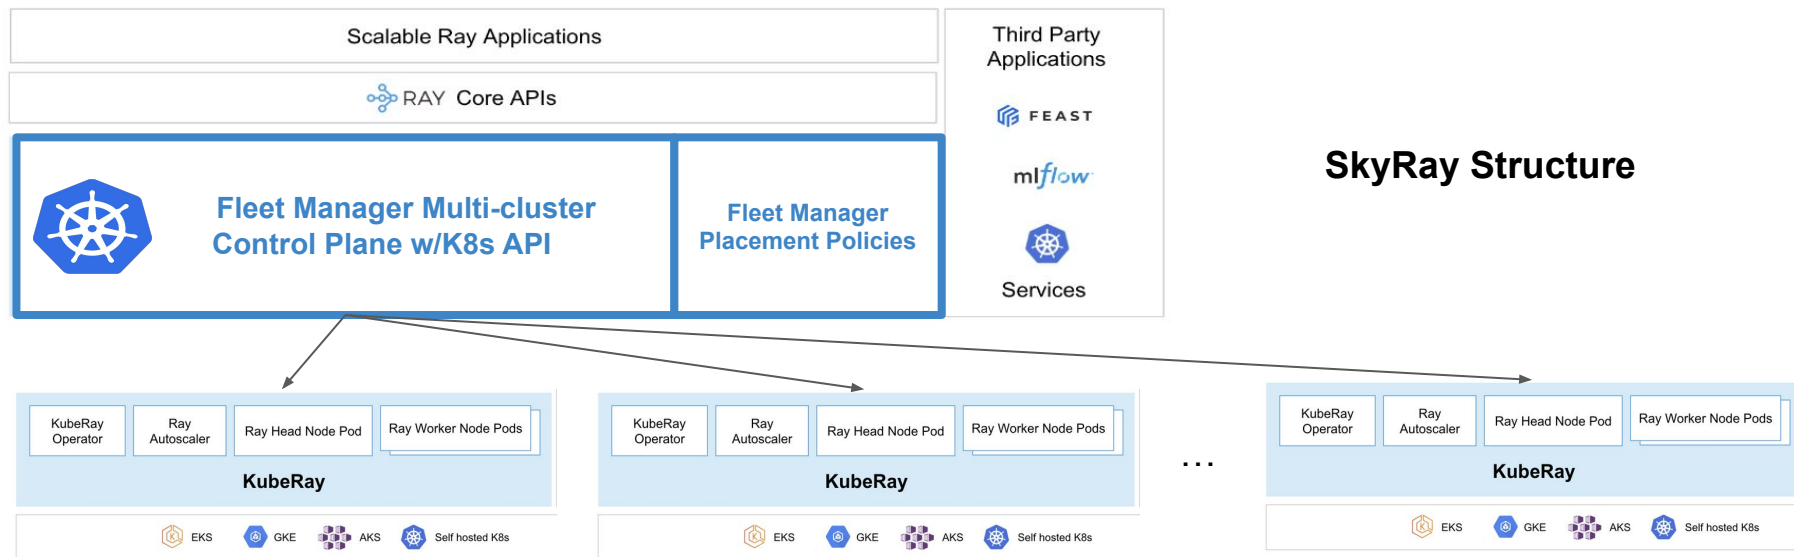
[From Cloud Computing To Sky Computing](#) [HotOS21]: need **commodity cloud compute layer**

- **K8s can represent a form of commodity compute layer for multi-cluster and multi-cloud use**
  - Requires using multi-cluster & multi-cloud K8s be nearly as **simple** as using single K8s cluster

# SkyRay Overview: SkyRay Vision

**SkyRay** seamlessly extends KubeRay from K8s cluster to multi-cluster multi-cloud operation

- **Achieved via policy-driven K8s fleet mgr (FM), presenting K8s API, interoperating w/KubeRay**
  - FM deploys KubeRay on each of a set of workload K8s clusters, as per FM policy
  - FM places KubeRay clusters/jobs/services on workload clusters, chosen as per FM policies
  - Some example K8s FMs include [Karmada](#), [Open Cluster Management](#), and [Nova](#)



# Talk Outline

SkyRay Overview

**SkyRay Operation**

SkyRay Example ML/AI Use Cases

Conclusion

# KubeRay Operation



**KubeRay**: K8s operator that simplifies deployment and management of Ray apps on K8s.

KubeRay supports three key Custom Resource Definitions:

- **RayCluster**
  - For creating a Ray cluster with the specified resources and attributes.
- **RayJob**
  - For creating a Ray cluster and submitting a job to it when the cluster is ready.
  - Can optionally delete the Ray cluster once the job finishes.
  - **Often used for ML/AI training.**
- **RayService**
  - For creating a Ray cluster and running a Ray Serve deployment graph.
  - Offers zero-downtime upgrades for Ray cluster, high availability, and [Ray Serve autoscaling](#).
  - **Often used for ML/AI serving.**

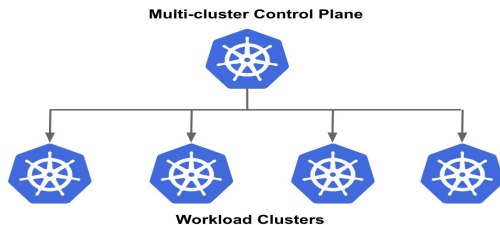
KubeRay deployments can optionally include the **Ray Autoscaler** for RayCluster scaling

- Automatically adds and removes workers from a RayCluster based on resource requests

# SkyRay Operation

**SkyRay** operates as follows:

- **Fleet Manager K8s control plane tracks K8s workload clusters' characteristics**
  - E.g., name, cloud provider, region, available capacity, labels, K8s version, cluster autoscaler, ...
- **Fleet Manager schedules KubeRay and its CRDs on all K8s workload clusters**
  - Specified by spread/duplicate Fleet Manager policy
- **All KubeRay cluster/job/service CR placement requests submitted to Fleet Manager K8s endpoint**
  - FM chooses workload cluster for requests based on cluster characteristics and placement policies



- Deploy workloads to one or more clusters from a central scheduler
- Aggregate view of workload topologies
- Orchestrate actions across workloads

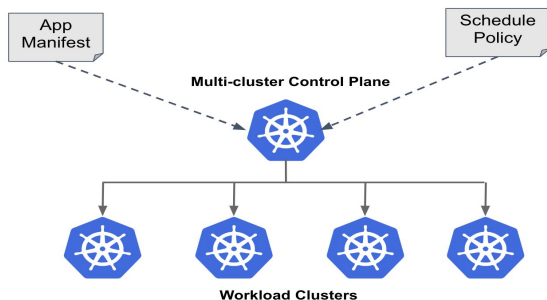
[Karmada](#), [Admiralty](#), [EloTi Nova](#) follow similar architecture.

# SkyRay K8s Fleet Manager Operation

For our SkyRay work, we used [Elotl Nova](#) K8s Fleet Manager, includes relevant capabilities:

- **Policies: spread w/duplicate** (or percent), **specified cluster**, **priority**, **available-capacity**
  - Nova workloads placed using an **available-capacity policy** are [gang-scheduled](#), meaning no part of a workload is scheduled until all of its components can be placed.
  - **ML/AI training jobs typically require gang-scheduling**, to allow all components of the distributed training task to make coordinated progress
- **Features: Just-in-time clusters [scale-to-0 or delete when idle, clone cluster]**, Overrides
  - JIT clusters can reduce costs or increase capacity for certain SkyRay scenarios

Decouple placement from workload definition

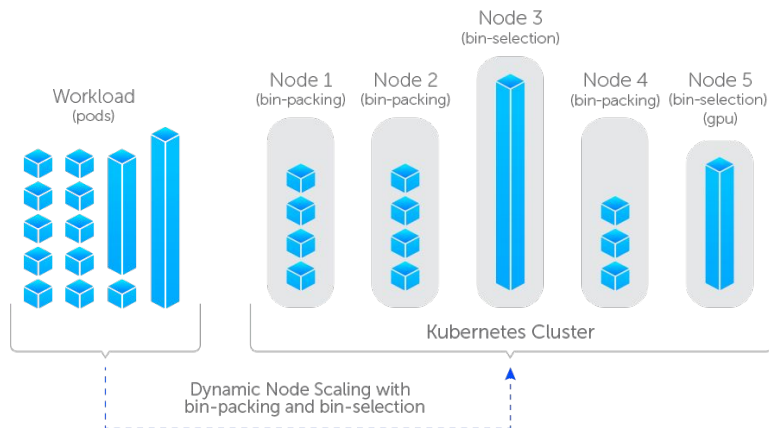




# Nova Fleet Manager Autoscaler-Aware Operation

When multiple clusters satisfy a workload's placement policy, Nova

- **Selects target cluster with existing sufficient available cluster resources, if any exists,**
- **Else selects target cluster with cluster autoscaler, either [K8s Cluster Autoscaler](#) or [Luna](#)**
  - **For our SkyRay work, we used the Luna smart cluster autoscaler**
    - Nova adds Luna's default pod placement label to workloads it schedules
  - **ML/AI workloads often require expensive GPU resources**
    - **Scaling GPU while needed from least expensive satisfactory type can reduce costs**



**Luna Operation**

# Talk Outline

SkyRay Overview

SkyRay Operation

**SkyRay Example ML/AI Use Cases**

Conclusion

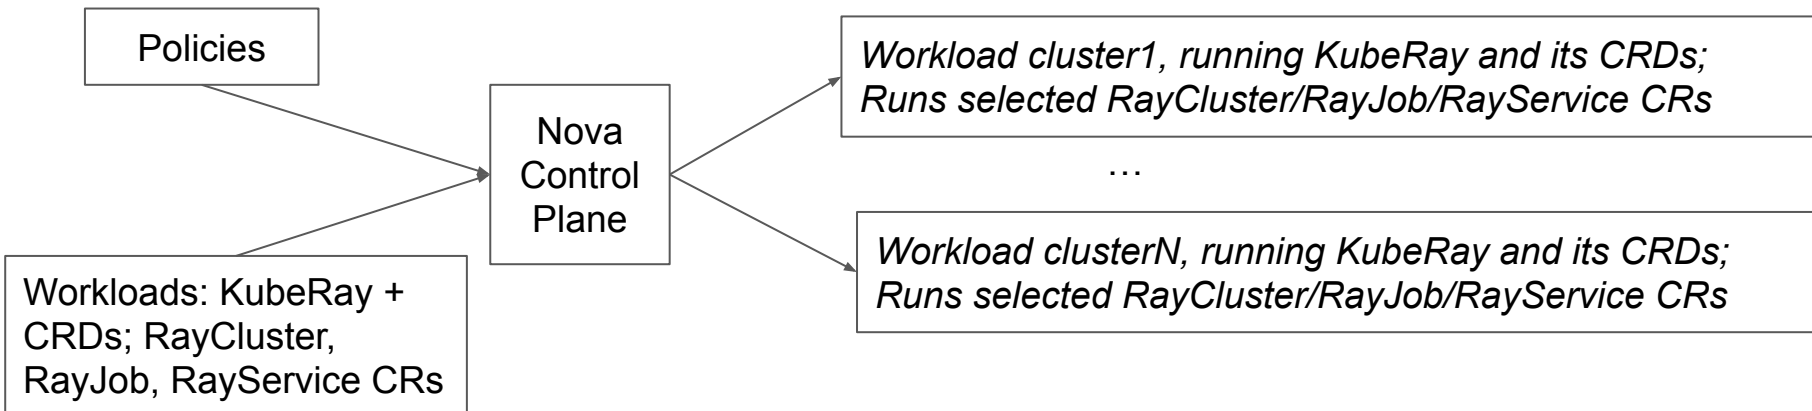
# SkyRay Examples Overview

## SkyRay recap

- Nova CP uses spread/duplicate policy to schedule KubeRay & CRDs to all workload clusters
- Nova CP schedules each RayCluster/RayJob/RayService CR to workload cluster
- KubeRay on workload cluster materializes head and worker pods for CR as usual

Scripts and K8s yaml for the examples available in github repo [elotl/skyray](https://github.com/elotl/skyray)

- Some examples include dynamic allocation via Nova JIT or via Ray + Luna autoscaling
- One includes Nova CP scheduling KubeRay + non-KubeRay objects for Compound AI



# SkyRay Example ML/AI Use Cases

**RayJobs Training production ML/AI models**

**RayJobs Training experimental ML/AI models**

**RayServices Serving production vs development ML/AI models**

**RayServices Serving LLM models multi-cloud by priority with suspend/resume JIT**

**RayService Serving LLM model, no downtime K8s upgrade using delete/create JIT**

**CompoundAI RAG+LLM Serving with RayService for LLM**



Image from KubeRay blog at AnyScale website

# Initial Setup for Scenario Examples

For simplicity of presentation, all scenario examples involve 2 workload clusters

Initial view from the Nova control Plane, after KubeRay spread/duplicate scheduled

```
$ kubectl get all --all-namespaces
```

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	service/kuberay-operator	ClusterIP	10.96.241.6	<none>	8080/TCP	91s
default	service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	6m50s

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kuberayns	deployment.apps/kuberay-operator	2/1	2	2	91s

Scenario examples expect SKYRAY\_PATH env var set to clone of github repo [elotl/skyray](https://github.com/elotl/skyray)

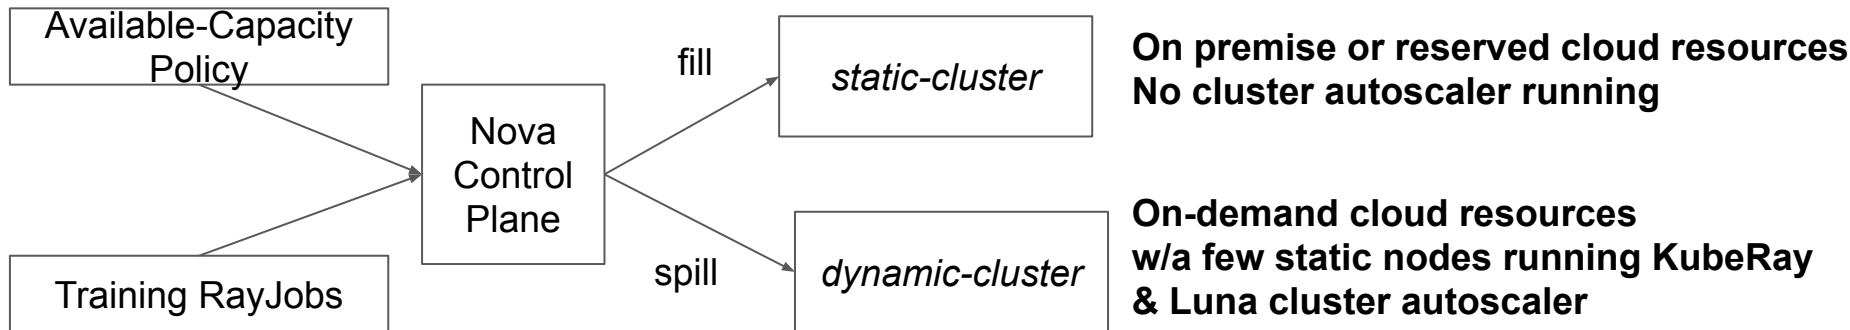
# Scenario: RayJobs Training Production Models

For RayJobs training production ML/AI models on GPUs, desired behavior is "fill and spill"

- Gang-scheduled on statically-allocated cluster if fit, else cluster w/dynamically-allocated resources
- Workloads' high value warrants on-demand resources; training so latency to scale resources ok

Nova [available-capacity placement policy](#) set to match both clusters; Nova gang-schedules

- Nova places training RayJobs on *static-cluster* preferentially since resources immediately available.
- When a training RayJob arrives that does not fully fit on *static-cluster*
  - Nova places it on *dynamic-cluster* and Luna adds resources to accommodate it



# RayJobs Training Production Models Example Overview

**As a proxy for a production training workload, we use Pytorch image train benchmark**

- Run as a RayJob deployed using KubeRay, adapted from the example [here](#).

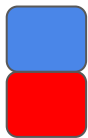


**Each RayJob's RayCluster is configured with CPU head and 2 single-GPU workers.**

- The configuration of the RayJob with its associated RayCluster is available [here](#).
- The RayJob available-capacity placement policy is [here](#).

**Three copies of the RayJob are deployed to the Nova Control plane.**

- Nova gang-schedules first 2 copies on static-cluster, since they fit there
- Nova gang-schedules third copy on dynamic-cluster, Luna scales up cluster resources
- Experiment run with nova control plane and workload clusters on [EKS](#)



static-cluster



dynamic-cluster

# RayJobs Training Production Models Example



## Deploy 1st training job instance in **rayjob1** namespace, scheduled on *static-cluster*

```
$ export RAYCLUSTER_NAMESPACE1=rayjob1
$ ${SKYRAY_PATH}/deploy-scripts/deploy-rayjob-train.sh ${SKYRAY_PATH} ${RAYCLUSTER_NAMESPACE1} ${AWS_ACCESS_KEY_ID}
${AWS_SECRET_ACCESS_KEY}
<snip>
$ export TARG_CLUSTER1=$(kubectl get rayjob.ray.io/rayjob-train -n ${RAYCLUSTER_NAMESPACE1} -L nova.elotl.co/target-cluster |
awk {'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER1}
static-cluster
```

## Deploy 2nd training job instance in **rayjob2** namespace, also scheduled on *static-cluster*

```
$ export RAYCLUSTER_NAMESPACE2=rayjob2
$ ${SKYRAY_PATH}/deploy-scripts/deploy-rayjob-train.sh ${SKYRAY_PATH} ${RAYCLUSTER_NAMESPACE2} ${AWS_ACCESS_KEY_ID}
${AWS_SECRET_ACCESS_KEY}
<snip>
$ export TARG_CLUSTER2=$(kubectl get rayjob.ray.io/rayjob-train -n ${RAYCLUSTER_NAMESPACE2} -L nova.elotl.co/target-cluster |
awk {'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER2}
static-cluster
```

## Deploy 3rd training job instance in **rayjob3** namespace, scheduled on *dynamic-cluster*

```
$ export RAYCLUSTER_NAMESPACE3=rayjob3
$ ${SKYRAY_PATH}/deploy-scripts/deploy-rayjob-train.sh ${SKYRAY_PATH} ${RAYCLUSTER_NAMESPACE3} ${AWS_ACCESS_KEY_ID}
${AWS_SECRET_ACCESS_KEY}
<snip>
$ export TARG_CLUSTER3=$(kubectl get rayjob.ray.io/rayjob-train -n ${RAYCLUSTER_NAMESPACE3} -L nova.elotl.co/target-cluster |
awk {'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER3}
dynamic-cluster
```





# “Fill and spill” achieved via available-capacity policy

## Multi-cluster deployment as simple as single cluster deployment; view jobs from Nova CP

```
$ kubectl get all --all-namespaces <SNIP>
```

NAMESPACE	NAME	JOB STATUS	DEPLOYMENT STATUS	START TIME	END TIME	AGE
rayjob1	rayjob.ray.io/rayjob-train	RUNNING	Running	2024-07-01T22:13:02Z		9m11s
rayjob2	rayjob.ray.io/rayjob-train	RUNNING	Running	2024-07-01T22:12:07Z		4m55s
rayjob3	rayjob.ray.io/rayjob-train		Initializing	2024-07-01T22:16:28Z		34s

## Luna scales up dynamic cluster accordingly

```
$ kubectl --context=dynamic-cluster get nodes -Lnode.kubernetes.io/instance-type
```

NAME	STATUS	ROLES	AGE	VERSION	INSTANCE-TYPE
ip-192-168-161-254.us-west-2.compute.internal	Ready	<none>	4m47s	v1.29.3-eks-ae9a62a	t3a.2xlarge
ip-192-168-182-75.us-west-2.compute.internal	Ready	<none>	55m	v1.29.3-eks-ae9a62a	t3a.small
ip-192-168-61-229.us-west-2.compute.internal	Ready	<none>	4m24s	v1.29.3-eks-ae9a62a	g4dn.2xlarge
ip-192-168-63-192.us-west-2.compute.internal	Ready	<none>	4m27s	v1.29.3-eks-ae9a62a	g4dn.2xlarge
ip-192-168-94-42.us-west-2.compute.internal	Ready	<none>	56d	v1.29.3-eks-ae9a62a	m5.large

## All 3 RayJobs run to completion

```
$ kubectl get all --all-namespaces <SNIP>
```

NAMESPACE	NAME	JOB STATUS	DEPL STATUS	START TIME	END TIME	AGE
rayjob1	rayjob.ray.io/rayjob-train	SUCCEEDED	Complete	2024-07-01T22:13:02Z	2024-07-01T22:26:30Z	22m
rayjob2	<a href="https://rayjob.ray.io/rayjob-train">rayjob.ray.io/rayjob-train</a>	SUCCEEDED	Complete	2024-07-01T22:12:07Z	2024-07-01T22:19:49Z	18m
rayjob3	rayjob.ray.io/rayjob-train	SUCCEEDED	Complete	2024-07-01T22:16:28Z	2024-07-01T22:30:27Z	14m

## Luna scales down the dynamic cluster, after rayjob3 removed

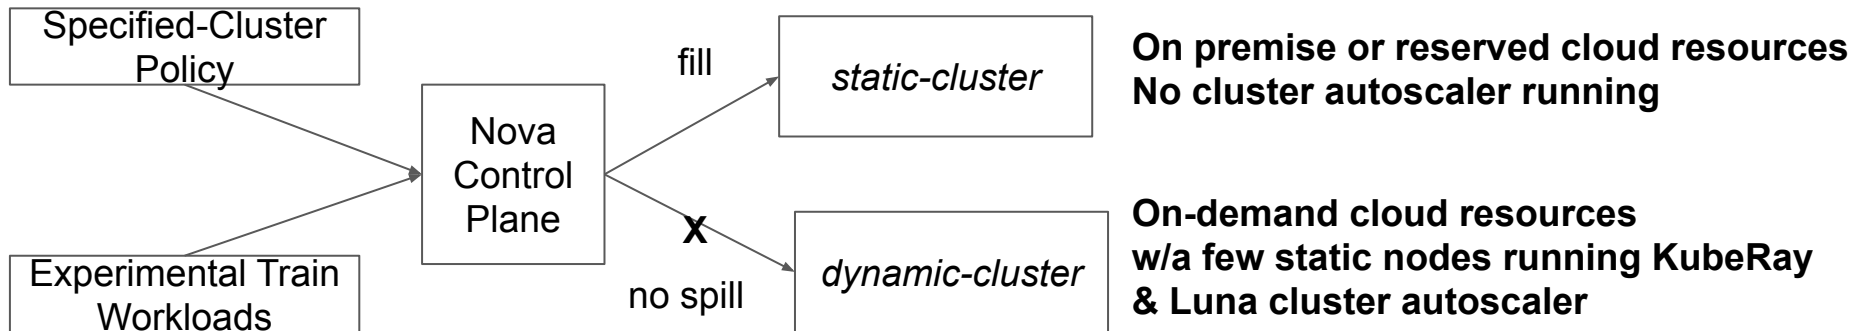
# Scenario: RayJobs Training Experimental Models

For training experimental ML/AI models on GPUs, desired behavior is **"fill, no spill"**

- **Workloads scheduled on statically-allocated on-premise or reserved cluster, sunk-cost**
- **Workloads have not proven value to warrant paying for on-demand cloud resources**

Nova [specified-cluster placement policy](#) is set to match only static-cluster

- Nova places all experimental training workloads on the cluster
- Any workloads that cannot be admitted are pending in the cluster.



# RayJobs Training Experimental Models Example Overview

**As proxy for experimental training workload, we again use Pytorch image train benchmark**

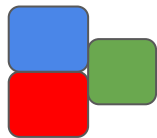
- Run as a RayJob deployed using KubeRay, adapted from the example [here](#)  PyTorch

**Each RayJob's RayCluster is again configured with a CPU head and 2 single-GPU workers.**

- The configuration of the RayJob with its associated RayCluster is available [here](#).
- The RayJob cluster-specific placement policy is [here](#).

**Three copies of the RayJob are deployed to the Nova Control plane.**

- Nova schedules all 3 on static-cluster
- Objects that do have sufficient resources to schedule stay pending



static-cluster

dynamic-cluster



# RayJobs Training Experimental Models

Deploy 3 training job instances in separate namespaces, all scheduled on *static-cluster*

```
$ export RAYCLUSTER_NAMESPACE=rayjob1
$ ${SKYRAY_PATH}/deploy-scripts/deploy-rayjob-train-static.sh ${SKYRAY_PATH} ${RAYCLUSTER_NAMESPACE} ${AWS_ACCESS_KEY_ID}
${AWS_SECRET_ACCESS_KEY}
<Snip>
$ export TARG_CLUSTER=$(kubectl get rayjob.ray.io/rayjob-train -n ${RAYCLUSTER_NAMESPACE} -L nova.elotl.co/target-cluster | awk
{'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER}
Static-cluster
```

```
$ export RAYCLUSTER_NAMESPACE=rayjob2
$ ${SKYRAY_PATH}/deploy-scripts/deploy-rayjob-train-static.sh ${SKYRAY_PATH} ${RAYCLUSTER_NAMESPACE} ${AWS_ACCESS_KEY_ID}
${AWS_SECRET_ACCESS_KEY}
<snip>
$ export TARG_CLUSTER=$(kubectl get rayjob.ray.io/rayjob-train -n ${RAYCLUSTER_NAMESPACE} -L nova.elotl.co/target-cluster | awk
{'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER}
static-cluster
```

```
$ export RAYCLUSTER_NAMESPACE=rayjob3
$ ${SKYRAY_PATH}/deploy-scripts/deploy-rayjob-train-static.sh ${SKYRAY_PATH} ${RAYCLUSTER_NAMESPACE} ${AWS_ACCESS_KEY_ID}
${AWS_SECRET_ACCESS_KEY}
<snip>
$ export TARG_CLUSTER=$(kubectl get rayjob.ray.io/rayjob-train -n ${RAYCLUSTER_NAMESPACE} -L nova.elotl.co/target-cluster | awk
{'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER}
Static-cluster
```

# “Fill, no spill” achieved via specified-cluster policy

Again, multi-cluster deployment as simple as single cluster deployment

In this case, *static-cluster* doesn’t have sufficient remaining resources to run 3rd RayJob.

- Its unschedulable pods stay pending until capacity freed up by removal of previous job(s).

```
$ kubectl get all --all-namespaces
. . .
NAMESPACE NAME                JOB STATUS DEPL STATUS START TIME                END TIME                AGE
rayjob1  rayjob.ray.io/rayjob-train SUCCEEDED Complete    2024-07-02T13:49:21Z    2024-07-02T13:56:49Z    8m5s
rayjob2  rayjob.ray.io/rayjob-train RUNNING    Running      2024-07-02T13:53:16Z
rayjob3  rayjob.ray.io/rayjob-train                Initializing 2024-07-02T13:54:47Z    2m39s
...
$ kubectl get all --all-namespaces
. . .
NAMESPACE NAME                JOB STATUS DEPL STATUS START TIME                END TIME                AGE
rayjob2  rayjob.ray.io/rayjob-train SUCCEEDED Complete    2024-07-02T13:53:16Z    2024-07-02T14:00:49Z    12m
rayjob3  rayjob.ray.io/rayjob-train RUNNING    Running      2024-07-02T13:54:47Z    11m
...
$ kubectl get all --all-namespaces
. . .
NAMESPACE NAME                JOB STATUS DEPL STATUS START TIME                END TIME                AGE
rayjob2  rayjob.ray.io/rayjob-train SUCCEEDED Complete    2024-07-02T13:53:16Z    2024-07-02T14:00:49Z    14m
rayjob3  rayjob.ray.io/rayjob-train SUCCEEDED Complete    2024-07-02T13:54:47Z    2024-07-02T14:07:53Z    13m
```

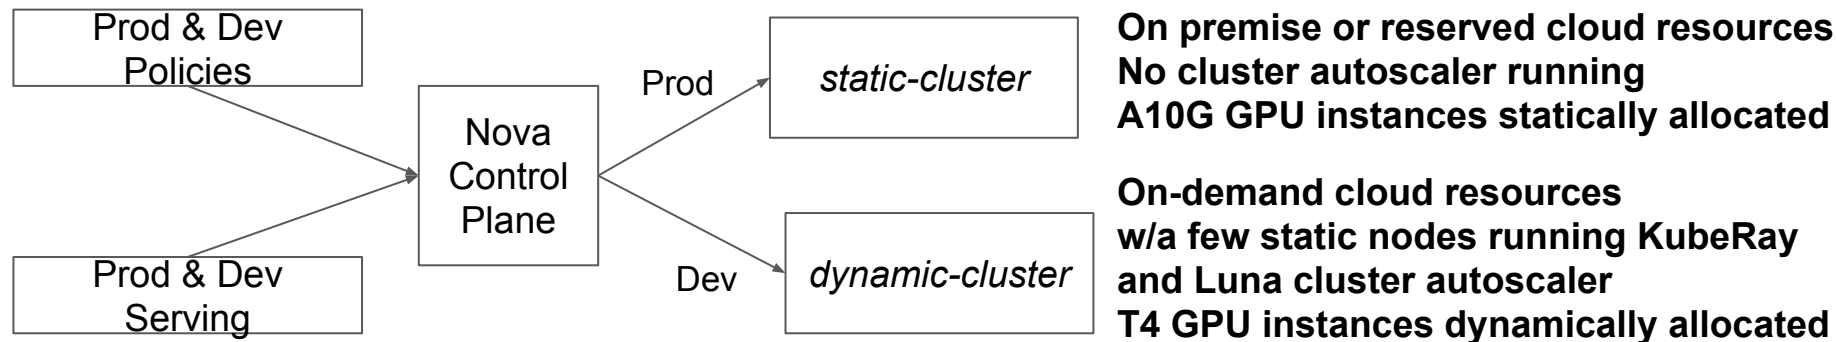
# Scenario: RayServices Serving Prod vs Dev AI/ML Models

For RayServices serving prod vs dev models, desired behavior is **"select the right cluster"**

- **Online prod serving placed on statically-allocated cluster configured for max load SLA**
  - **Low latency required**, serving on critical path of business app (e.g., predicting ride-share ETA)
- **Dev serving placed on cluster w/dynamically-allocated lower-cost resources**
  - **Low latency not required**

**Nova cluster selection policies ([prod](#), [dev](#)) set to match cluster label for workload class**

- **Label *production* on *static-cluster* cluster CR, *development* on *dynamic-cluster* cluster CR**
  - **Note: Cluster labels add layer of indirection that facilitates changing set of category clusters**



# RayServices Serving Prod vs Dev Models Example Overview

As a representative serving workload, we use the **text summarizer model service**

- Run as a **RayService** deployed using KubeRay, adapted from the example [here](#).

**Each RayService's RayCluster is configured with a CPU head and 1 single-GPU worker**

- The configuration of the RayService with its associated RayCluster is available [here](#)
- These are the RayService cluster selection policies for [production](#), [development](#)

**Two copies of the RayService are deployed to the Nova Control plane.**

- The RayService deployed to the production namespace is scheduled on *static-cluster*
- The RayService deployed to the development namespace is scheduled on *dynamic-cluster*



static-cluster



dynamic-cluster

# RayService Production Serving

Production namespace deployed to Nova CP; Placed w/ spread/duplicate policy

<SNIP>

## RayService deployed to Nova CP in production namespace; Placed on static-cluster

```
$ kubectl apply -f ${SKYRAY_PATH}/deploy-scripts/ray-service.text-summarizer.yaml --namespace=production
rayservice.ray.io/text-summarizer created
```

```
$ kubectl --context=static-cluster get all -n production
```

NAME	READY	STATUS	RESTARTS	AGE			
pod/text-summarizer-raycluster-ntcfh-head-tmnqr	1/1	Running	0	68m			
pod/text-summarizer-raycluster-ntcfh-worker-gpu-group-wft6f	1/1	Running	0	68m			
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE		
service/text-summarizer-head-svc	ClusterIP	10.100.6.157	<none>	10001/TCP,8265/TCP,6379/TCP,8080/TCP,8000/TCP	60m		
service/text-summarizer-raycluster-ntcfh-head-svc	ClusterIP	10.100.197.135	<none>	10001/TCP,8265/TCP,6379/TCP,8080/TCP,8000/TCP	68m		
service/text-summarizer-serve-svc	ClusterIP	10.100.205.162	<none>	8000/TCP	60m		
NAME	DESIRED WORKERS	AVAILABLE WORKERS	CPUS	MEMORY	GPUS	STATUS	AGE
raycluster.ray.io/text-summarizer-raycluster-ntcfh	1	1	5	20G	1	ready	68m
NAME	AGE						
rayservice.ray.io/text-summarizer	68m						

## Production Serving Validated:

```
$ kubectl --context=static-cluster port-forward svc/text-summarizer-serve-svc 8000 -n production
<SNIP>
```

```
$ python text_summarizer_req.py
```

Paris is the capital and most populous city of France. It has an estimated population of 2,175,601 residents as of 2018. The City of Paris is the centre of the French capital.



# RayService Development Serving

Development namespace deployed to Nova control plane; Placed w/ spread/duplicate policy

<SNIP>

**RayService** deployed to Nova CP in development namespace; Placed on *dynamic-cluster*

```
$ kubectl apply -f ${SKYRAY_PATH}/deploy-scripts/ray-service.text-summarizer.yaml --namespace=development
rayservice.ray.io/text-summarizer created
```

```
$ kubectl --context=dynamic-cluster get all -n development
```

NAME	READY	STATUS	RESTARTS	AGE
pod/text-summarizer-raycluster-2xnts-head-68bvm	1/1	Running	0	47m
pod/text-summarizer-raycluster-2xnts-worker-gpu-group-s8pbn	1/1	Running	0	47m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/text-summarizer-head-svc	ClusterIP	10.100.45.127	<none>	10001/TCP,8265/TCP,6379/TCP,8080/TCP,8000/TCP	37m
service/text-summarizer-raycluster-2xnts-head-svc	ClusterIP	10.100.46.227	<none>	10001/TCP,8265/TCP,6379/TCP,8080/TCP,8000/TCP	47m
service/text-summarizer-serve-svc	ClusterIP	10.100.209.7	<none>	8000/TCP	37m

NAME	DESIRED WORKERS	AVAILABLE WORKERS	CPUS	MEMORY	GPUS	STATUS	AGE
raycluster.ray.io/text-summarizer-raycluster-2xnts	1	1	5	20G	1	ready	47m

NAME	AGE
rayservice.ray.io/text-summarizer	47m

## Development Serving Validated

```
$ kubectl --context=dynamic-cluster port-forward svc/text-summarizer-serve-svc 8000 -n development
```

<SNIP>

```
$ python text_summarizer_req.py
```

Paris is the capital and most populous city of France. It has an estimated population of 2,175,601 residents as of 2018. The City of Paris is the centre of the French capital.



# “Select right cluster” achieved via label cluster-select policy

Again, multi-cluster deployment as simple as single cluster deployment

Prod on more costly statically-allocated GPU, dev on cheaper dynamically-allocated GPU

- ***static-cluster*** was configured w/ g5.xlarge, which includes an NVIDIA A10G GPU
- **Luna** allocates a g4dn.xlarge for ***dynamic-cluster***, which includes an NVIDIA T4 GPU
  - us-east per-hr on-demand price for g4dn.xlarge < per-hr reserved price for g5.xlarge
  - g4dn.xlarge is good choice for dev workload, which doesn't warrant more powerful GPU

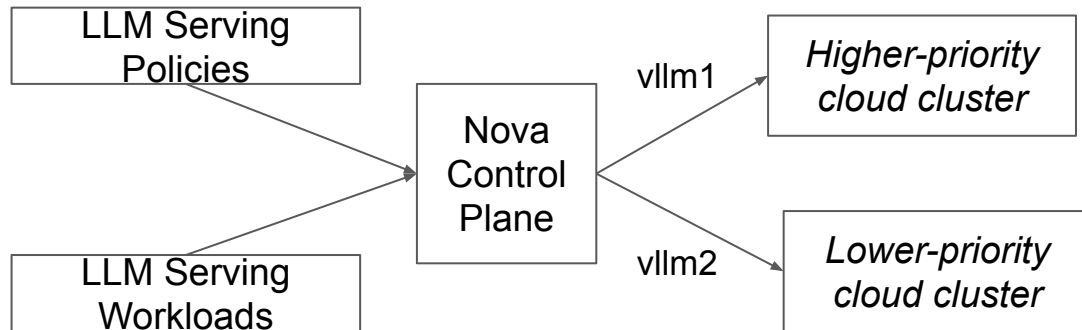
```
$ kubectl --context=dynamic-cluster get nodes -lnode.kubernetes.io/instance-type
```

NAME	STATUS	ROLES	AGE	VERSION	INSTANCE-TYPE
ip-192-168-164-97.us-west-2.compute.internal	Ready	<none>	8d	v1.29.3-eks-ae9a62a	t3a.small
ip-192-168-171-101.us-west-2.compute.internal	Ready	<none>	48m	v1.29.3-eks-ae9a62a	t3a.xlarge
ip-192-168-49-24.us-west-2.compute.internal	Ready	<none>	48m	v1.29.3-eks-ae9a62a	g4dn.xlarge
ip-192-168-94-42.us-west-2.compute.internal	Ready	<none>	64d	v1.29.3-eks-ae9a62a	m5.large

# Scenario: RayServices Serving LLM Models Multi-cloud Priority w/JIT

For serving LLM models w/ordered cloud preferences, desired "select cluster by cloud priority"

- Nova [cluster selection policy](#) use cloud in priority order, set as per some criteria, e.g.
  - Higher-priority may be lower cost, lower-priority may have more resources
- JIT w/Suspend/Resume standby; cluster idle->nodes scaled to 0, restored when active
  - Option NOVA\_NAMESPACES\_EXCLUDE\_ACTIVE set to Kuberay namespace *kuberayns*;  
So KubeRay running on workload clusters does not block cluster entering standby



## Presenting two cloud provider priority examples

- E1: Fill/spill w/GKE & EKS clusters, shows on-demand w/alternative on-demand
- E2: Fill/spill w/Rackspace & EKS clusters, shows allocated spot w/alternative on-demand

### GKE cluster on GCE cloud provider; nodes:

- 2x e2-medium: 2 CPUs, 4 GB memory
- 1x e2-standard-4: 4 CPUs, 16 GB memory
- 1x g2-standard-16: 16 CPUs, 64 GB memory, 1 L4 GPU; node taint key: nvidia.com/gpu; effect: NoSchedule

### EKS cluster on AWS cloud provider; nodes:

- 2x m5.large: 2 CPUs, 8 GB memory
- 1x t3a.xlarge: 4 CPUs, 16 GB memory
- 1x g6.4xlarge: 16 CPUs, 64 GB memory, 1 L4 GPU; node taint key: nvidia.com/gpu; effect: NoSchedule

### Rackspace Spot cluster on Openstack cloud provider managed by Platform9; nodes:

- 1x ao.2.24.128\_A30 (24 CPUs, 128 GB memory, 1 A30 GPU)

# E1: RayServices Serving LLM Models Multi-cloud Priority



EKS and GKE clusters initially idle and in standby, control plane costs \$0.10/hr  
KubeRay is pending with its *kuberayns* namespace set to not represent active use

```
$ kubectl get clusters
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY
eks-workload-cluster	1.30	cpu-us-west-2	us-west-2	us-west-2d	False	True	True
gke-workload-cluster	1.29	anne-nova-cp	us-central1	us-central1-c	False	True	True

Spread/dupe vllm1 namespace for LLM deploy; Set [cloud priority policy](#) for workloads

```
$ export RAYCLUSTER_NAMESPACE=vllm1
$ envsubst < ${SKYRAY_PATH}/deploy-scripts/namespace.yaml | kubectl apply -f -
$ envsubst < ${SKYRAY_PATH}/policies/rayservicensprioritypolicy.yaml | kubectl apply -f -
```

Deploy [RayService to serve LLM model mosaicml/mpt-7b-chat using vLLM](#) to vllm1 namespace;  
gke-workload-cluster exits standby and kuberay starts up rayservice head and worker on it

```
$ kubectl apply -f ${SKYRAY_PATH}/luna-llm-serve/ray-service.llm.yaml -n ${RAYCLUSTER_NAMESPACE}
$ kubectl wait --for=jsonpath='{.status.serviceStatus}'=Running rayservice.ray.io/llm-model-serve --namespace
${RAYCLUSTER_NAMESPACE} --timeout=20m
$ export TARG_CLUSTER=$(kubectl get rayservice.ray.io/llm-model-serve -n ${RAYCLUSTER_NAMESPACE} -L
nova.elotl.co/target-cluster | awk {'print $NF'} | tail -1)
$ echo ${TARG_CLUSTER}
gke-workload-cluster
```

# “Select cluster by cloud priority” achieved via ordered-selector policy

## Validate serving LLM model from vllm1 namespace via [query.py](#)

```
$ kubectl --context=${TARG_CLUSTER} port-forward service/llm-model-serve-serve-svc 8000:8000 -n ${RAYCLUSTER_NAMESPACE}
$ python query.py
```

Type your query here: what is a good christmas present for a 6 year old?

A good Christmas present for a 6-year-old could be a board game like Monopoly or a toy like a doll or a robot. Other ideas could be books, puzzles, or arts and crafts supplies. It's important to consider the child's interests and hobbies when choosing a gift.

## Perform same ops for LLM model in namespace vllm2, workload spilled to awakened EKS cluster

```
$ export RAYCLUSTER_NAMESPACE=vllm2 <SNIP>
$ echo ${TARG_CLUSTER}
eks-workload-cluster
```

## Validate serving LLM model from vllm2 namespace

...Type your query here: When is a good time to buy a stock?

As an AI language model, I don't have personal beliefs or opinions, but I can provide some general guidelines to help you make informed decisions when buying stocks. 1. Research the company's financial health and performance. Look at its past earnings reports, revenue growth, and debt-to-equity ratio. <SNIP>

## Both workload clusters active; delete namespace to remove model, cluster re-enters standby

# E2: RayServices Serving LLM Models Multi-cloud Priority



Rackspace & EKS clusters initially idle; KubeRay *kuberayns* namespace set to not be active use

```
$ kubectl get clusters
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY
elotl-llmexperiment	1.29		SJC3	nova	True	True	False
workload-cluster-1	1.30	gpu-cluster	us-west-2	us-west-2d	True	True	False

Spread/dupe vllm1 namespace for LLM deploy; Set [cloud priority policy](#) for workloads

```
$ export RAYCLUSTER_NAMESPACE=vllm1
```

```
$ envsubst < ${SKYRAY_PATH}/deploy-scripts/namespace.yaml | kubectl apply -f -
```

```
$ envsubst < ${SKYRAY_PATH}/policies/rayservicensprioritypolicy.yaml | kubectl apply -f -
```

Deploy [RayService to serve LLM model mosaicml/mpt-7b-chat using vLLM](#) to vllm1 namespace; elotl-llmexperiment cluster chosen and kuberay starts up rayservice head and worker on it

```
$ kubectl apply -f ${SKYRAY_PATH}/luna-llm-serve/ray-service.llm.yaml -n ${RAYCLUSTER_NAMESPACE}
```

```
$ kubectl wait --for=jsonpath='{.status.serviceStatus}'=Running rayservice.ray.io/llm-model-serve --namespace
```

```
${RAYCLUSTER_NAMESPACE} --timeout=20m
```

```
$ export TARG_CLUSTER=$(kubectl get rayservice.ray.io/llm-model-serve -n ${RAYCLUSTER_NAMESPACE} -L nova.elotl.co/target-cluster  
| awk {'print $NF'} | tail -1)
```

```
$ echo ${TARG_CLUSTER}
```

```
elotl-llmexperiment
```

# “Select cluster by cloud priority” achieved via ordered-selector policy

## Validate serving LLM model from vllm1 namespace via [query.py](#)

```
$ kubectl --context=${TARG_CLUSTER} port-forward service/llm-model-serve-serve-svc 8000:8000 -n ${RAYCLUSTER_NAMESPACE}
$ python query.py
```

Type your query here: what is the longest river in the world?

The longest river in the world is the Nile River, which flows through 11 countries in Africa. It has a length of 6,695 kilometers (4,130 miles) and a drainage basin of 2,130,000 square kilometers (827,000 square miles).

## Perform same ops for LLM model in namespace vllm2, workload spilled to EKS cluster

```
$ export RAYCLUSTER_NAMESPACE=vllm2 <SNIP>
$ echo ${TARG_CLUSTER}
eks-workload-cluster
```

## Validate serving LLM model from vllm2 namespace

Type your query here: what is zumba?

Zumba is a fitness program that combines dance and aerobic exercise. It was created by Alberto "Beto" Perez in 2002 and has since become a global phenomenon, with millions of people taking part in Zumba classes around the world. Zumba workouts are designed to be fun and energetic, and they typically involve a combination of fast-paced dancing, aerobic exercise, and strength training. Zumba classes are led by instructors who teach the moves and provide motivation and encouragement to participants.

## Both workload clusters active; again, can delete namespace to remove model



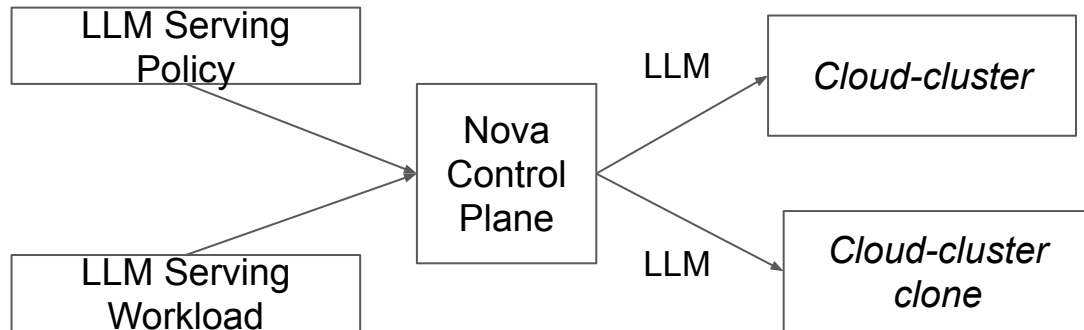
# Scenario: LLM Model Serving w/No downtime K8s upgrade

For upgrade K8s cluster serving LLM model, desired behavior “**clone-upgrade/spread/validate**”

- **RayService LLM serving on K8s cluster w/ spread/duplicate on labeled Nova cluster policy**
- **Nova JIT w/Delete/Create standby enabled**

**User adds clone of labeled Nova cluster CR w/updated K8s version, K8s and Nova names**

- **Nova JIT creates corresponding cluster, Nova schedules dupe LLM serving workload on it**
- **User validates LLM serving workload on upgraded K8s cluster, chooses when to switch DNS**
  - User then removes label from old Nova cluster CR, cluster becomes idle and JIT deletes it  
User can then remove old cluster CR



# LLM Model Serving w/No downtime K8s upgrade, Start

## Label applied to workload cluster

```
$ kubectl label cluster workload-cluster-1 cluster.type=rayservice
```

```
$ kubectl get clusters --context=nova --show-labels
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY	LABELS
workload-cluster-1	1.30	cpu-us-west-2	us-west-2	us-west-2d	True	True	False	cluster.type=rayservice, <a href="https://kubernetes.io/metadata.name=workload-cluster-1.nova.elotl.co/cluster.novacreated=false.nova.elotl.co/cluster.provider=aws.nova.elotl.co/cluster.region=us-west-2.nova.elotl.co/cluster.version=1.30.nova.elotl.co/cluster.zone=us-west-2d">kubernetes.io/metadata.name=workload-cluster-1.nova.elotl.co/cluster.novacreated=false.nova.elotl.co/cluster.provider=aws.nova.elotl.co/cluster.region=us-west-2.nova.elotl.co/cluster.version=1.30.nova.elotl.co/cluster.zone=us-west-2d</a>

## Spread/duplicate vllm1 ns for LLM deploy; Set [spread/duplicate/labeled policy](#) for ns workloads

```
$ export RAYCLUSTER_NAMESPACE=vllm1
```

```
$ envsubst < ${SKYRAY_PATH}/deploy-scripts/namespace.yaml | kubectl apply -f -
```

```
$ envsubst < ${SKYRAY_PATH}/policies/rayservicespreadypolicy.yaml | kubectl apply -f -
```

## Deploy [RayService to serve LLM model mosaicml/mpt-7b-chat using vLLM](#) to vllm1 namespace

```
$ kubectl apply -f ${SKYRAY_PATH}/luna-llm-serve/ray-service.llm.yaml -n ${RAYCLUSTER_NAMESPACE}
```

```
$ kubectl get rayservice.ray.io/llm-model-serve -n ${RAYCLUSTER_NAMESPACE} -o
```

```
jsonpath='{.metadata.annotations.nova\.elotl\.co/spread-onto}' | awk {'print $NF'}
```

```
workload-cluster-1::Duplicate
```

```
$ export TARG_CLUSTER=workload-cluster-1
```

# LLM Model Serving w/No downtime K8s upgrade, Con't

## Validate serving LLM model from vllm1 namespace via [query.py](#)

```
$ kubectl --context=${TARG_CLUSTER} port-forward service/llm-model-serve-serve-svc 8000:8000 -n ${RAYCLUSTER_NAMESPACE}
$ python query.py
Type your query here: what year did the titanic sink? The Titanic sank in April 1912.
```

## Get copy of Nova cluster CR

```
$ kubectl get cluster workload-cluster-1 -o yaml > workload-cluster-1.yaml
$ cp workload-cluster-1.yaml workload-cluster-1-upd.yaml
```

## Edit cluster CR copy to change the Nova cluster name, the K8s cluster name, and the K8s version

```
$ diff workload-cluster-1.yaml workload-cluster-1-upd.yaml
11c11
<     kubernetes.io/metadata.name: workload-cluster-1
-
>     kubernetes.io/metadata.name: workload-cluster-1-upd
17c17
<   name: workload-cluster-1
-
>   name: workload-cluster-1-upd
21,23c21,23
<   kubernetescluster: cpu-us-west-2
<   kubernetesversion: "1.30"
<   name: workload-cluster-1
-
>   kubernetescluster: cpu-us-west-2-upd
>   kubernetesversion: "1.31"
>   name: workload-cluster-1-upd
```

# “No Downtime K8s upgrade” via clone-upgrade/spread/validate

## Apply resulting updated CR copy

```
$ kubectl apply -f workload-cluster-1-upd.yaml
```

## New cluster CR label matches LLM spread/duplicate policy; Nova creates cluster for placement

```
$ kubectl get rayservice.ray.io/llm-model-serve -n ${RAYCLUSTER_NAMESPACE} -o
jsonpath='{.metadata.annotations.nova\.elotl\.co/spread-onto}' | awk {'print $NF'}
workload-cluster-1,workload-cluster-1-upd::Duplicate
$ kubectl get clusters
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY
workload-cluster-1	1.30	cpu-us-west-2	us-west-2	us-west-2d	True	False	False
workload-cluster-1-upd	1.31	cpu-us-west-2-upd	us-west-2	us-west-2d	True	False	False

## User validates serving LLM model from vllm1 namespace on the upgraded cluster via [query.py](#)

```
$ export TARG_CLUSTER=workload-cluster-1-upd
$ kubectl --context=${TARG_CLUSTER} port-forward service/llm-model-serve-serve-svc 8000:8000 -n ${RAYCLUSTER_NAMESPACE}
$ python query.py
Type your query here: how many days are in March? There are 31 days in March.
```

## User removes label from old cluster -> idle; JIT deletes after a while, user can then delete its CR

```
$ kubectl label cluster workload-cluster-1 cluster.type-
<SNIP>
```

```
$ kubectl get clusters
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY
workload-cluster-1-upd	1.31	cpu-us-west-2-upd	us-west-2	us-west-2d	True	False	False

# Scenario: RAG+LLM Serving w/RayService for LLM

For RAG+LLM serving w/RayService for LLM, desired behavior is **"select cluster for task"**

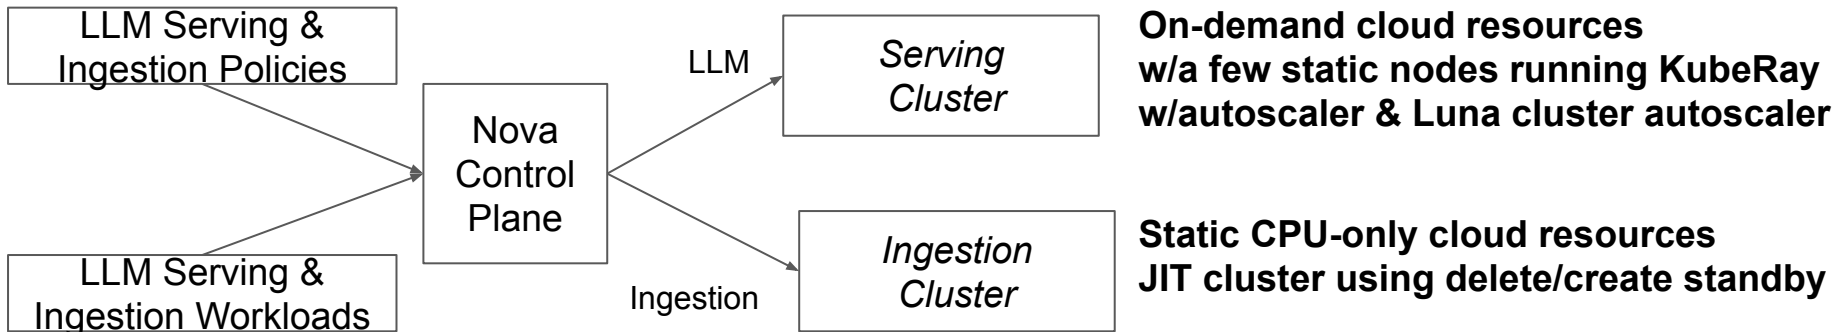
- **Nova cluster selection policy matches task namespace w/cluster for task**
- [CompoundAI](#) example that uses KubeRay and non-KubeRay components
- **Retrieval Augmented Generation ingests data into vectordb for context lookup during serving**

**Ingestion job placed on JIT ingestion cluster comprised of static CPU-only cloud resources**

- LangChain doc loaders/splitters, HuggingFace embeddings, FAISS vectorDB; pickled for S3 upload

**LLM serving and RAG+LLM serving placed on serving cluster, which is dynamic**

- LLM serving: RayService deployed by KubeRay w/autoscaler
- RAG+LLM serving: LangChain+FASTAPI+FAISS vectorDB S3 download/unpickle



# RayService Serving LLM model with RAG, Start

## Workload clusters set up

```
$ kubectl get clusters
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY
ingestion-cluster	1.30	cpu-cluster	us-west-2	us-west-2b	True	True	False
workload-cluster-1	1.30	gpu-cluster	us-west-2	us-west-2d	True	True	False

## EKS workload-cluster-1 static nodes

- 2x m5.large: 2 CPUs, 8 GB memory

## EKS ingestion-cluster static nodes

- 2x m5.large: 2 CPUs, 8 GB memory

## workload-cluster-1 labeled for rayservice serving and for model rag serving

```
$ kubectl label cluster workload-cluster-1 cluster.type=rayservice
```

```
$ kubectl label cluster workload-cluster-1 nova.elotl.co/cluster.modelragllmserve=true
```

# RayService Serving LLM model with RAG, Continue

## Place LLM; KubeRay creates head, Ray Autoscaler creates worker, Luna allocates nodes

```
$ export RAYCLUSTER_NAMESPACE=serving
$ envsubst < ${SKYRAY_PATH}/deploy-scripts/namespace.yaml | kubectl apply -f -
$ envsubst < ${SKYRAY_PATH}/policies/rayservicespreadpolicy.yaml | kubectl apply -f -
$ kubectl apply -f ${SKYRAY_PATH}/luna-llm-serve/ray-service.llm.mpt-7b-chat.autoscale.yaml -n ${RAYCLUSTER_NAMESPACE}
```

## Validate LLM model running on workload-cluster-1

```
$ kubectl --context=workload-cluster-1 port-forward service/llm-model-serve-serve-svc 8000:8000 -n serving
$ python query.py
Type your query here: what is the longest bridge in the world?
The longest bridge in the world is the Akashi-Kaikyo Bridge in Japan. It is a suspension bridge that spans a total length of 20.8 kilometers (12.6 miles) and connects the city of Kobe in Japan to the island of Honshu.
```

## Run ingestion job on ingestion-cluster

```
$ export MODEL_INGESTION_CLUSTER=ingestion-cluster
$ export MODEL_NAMESPACE=serving
$ envsubst < ${GENAI_PATH}/demo/llm.vdb.service/specificclusterpolicy.yaml | kubectl apply -f -
# export AWS_ACCESS_KEY_ID=<censored>, AWS_SECRET_ACCESS_KEY=<censored>, VECTOR_DB_INPUT_TYPE=text-docs.
# VECTOR_DB_INPUT_ARG=mini-rag-wikipedia-input, VECTOR_DB_S3_BUCKET=servi-faiss-vectordbs.
VECTOR_DB_S3_FILE=anne-rag-wikipedia.pkl
$ envsubst < ${GENAI_PATH}/demo/llm.vdb.service/createvdb.yaml | kubectl apply -n ${MODEL_NAMESPACE} -f -
```

## Place RAG+LLM serving service on workload-cluster-1

```
$ envsubst < ${GENAI_PATH}/demo/llm.rag.service/spreadacrossclusterset.yaml | kubectl apply -f -
$ export MODEL_LLM_SERVER_URL=http://llm-model-serve-serve-svc.serving.svc.cluster.local:8000
$ envsubst < ${GENAI_PATH}/demo/llm.rag.service/chat-serveragllmpluslb.yaml | kubectl apply -n ${MODEL_NAMESPACE} -f -
```

# "Select cluster for task" via Nova policies and JIT, Ray & Luna autoscaler

Run RAG+LLM query, RAG data was [rag-mini-wikipedia](https://rag-mini-wikipedia.com)

```
$ curl -X GET
```

```
"http://a8524e9bf88034df4b01095a0949752c-549500553.us-west-2.elb.amazonaws.com/answer/what%20are%20the%20two%20types%20of%20elephants%20in%20Africa?"
```

```
{"question":"what are the two types of elephants in Africa","answer":"The two types of elephants in Africa are the African Forest Elephant and the African Savanna Elephant."}
```

And note that the idle ingestion cluster entered standby and was deleted from the cloud

```
$ kubectl get clusters
```

NAME	K8S-VERSION	K8S-CLUSTER	REGION	ZONE	READY	IDLE	STANDBY
ingestion-cluster	1.30	cpu-cluster	us-west-2	us-west-2b	False	True	True
workload-cluster-1	1.30	gpu-cluster	us-west-2	us-west-2d	True	False	False

As usual, when done, cleanup is easy via namespace deletion

```
$ kubectl delete ns serving
```

Again, multi-cluster deployment as simple as single cluster deployment



# Talk Outline

SkyRay Overview

SkyRay Operation

SkyRay Example ML/AI Use Cases

**Conclusion**

# Conclusion

## SkyRay seamlessly extends Kuberay for multi-cluster scheduling scenarios

- **Reduces launch-time** of critical ML/AI workloads by scheduling on available GPUs
- **Increases efficiency** by directing experimental ML/AI jobs to sunk-cost clusters
- **Manages costs** via policies that select GPUs with the desired price/performance
- **Enhances robustness** by supporting prioritized cluster use from multiple cloud vendors
- **Facilitates K8 cluster maintenance** with no workload downtime K8s upgrade
- **Handles scheduling compound AI jobs** with KubeRay models and with non-Ray components

## We want SkyRay to handle your additional desired scheduling outcomes as well!

- And BTW Nova's flexible scheduling policies have been applied to various domains, including Managing LLM+RAG deployments, Performing cloud-agnostic gitops, Handling DR failover

If you'd like to try SkyRay, please download [free trial versions](#) of Nova and Luna.

**ELOT**

Thanks! Questions?

