



Why does continuous profiling matter to developers?

KubeCon NA Salt Lake City 2024
Co-located event: AppDeveloperCon



Who are we?

Jonas Kunz

Observability SW Engineer @ Elastic

OpenTelemetry-Java Contributor

Jonas Kunz @ CNCF-Slack

Mauricio Salatino

@Diagrid @Daprdev

Application Development WG co-chair

<https://salaboy.com>



Application Development WG Survey

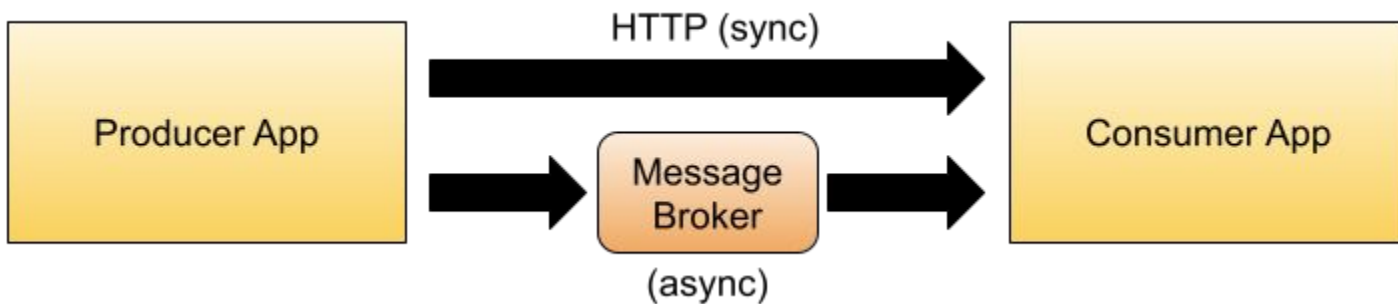
ADD QR CODE



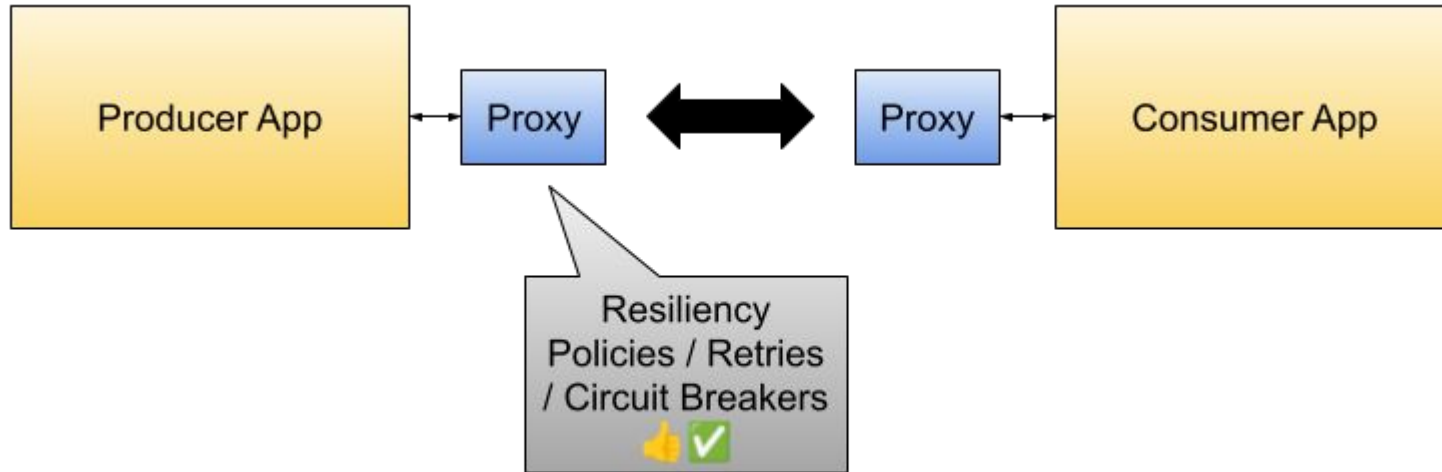
Agenda

- Building Cloud Native Resilient and Observable applications
- Pillars of Observability
- (Continuous) Profiling
- Next steps

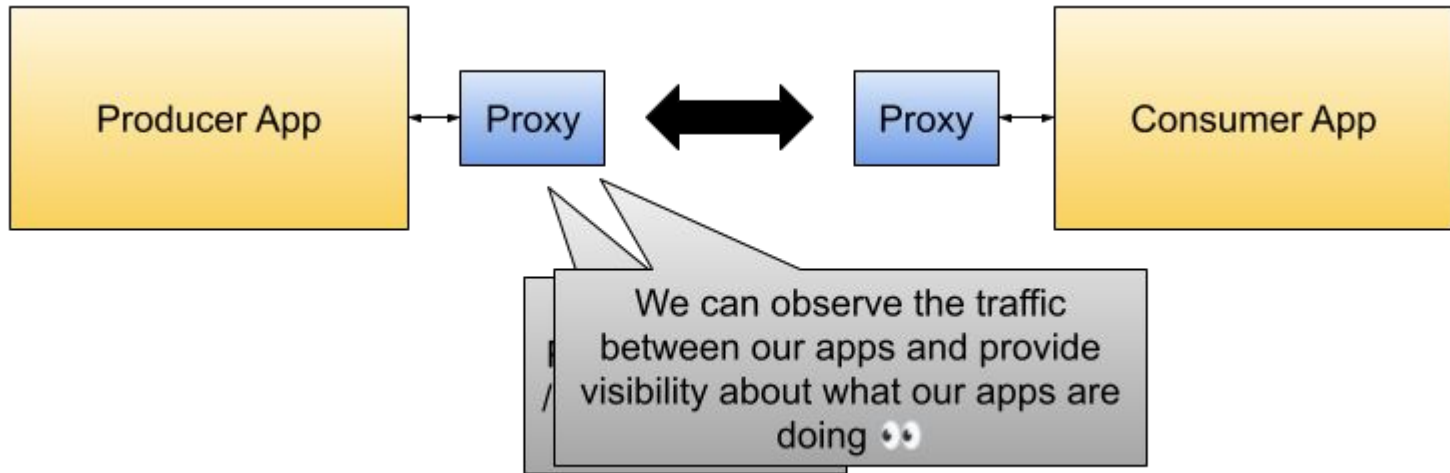
Distributed applications 101



Cloud Native distributed applications



Resilient and observable distributed applications



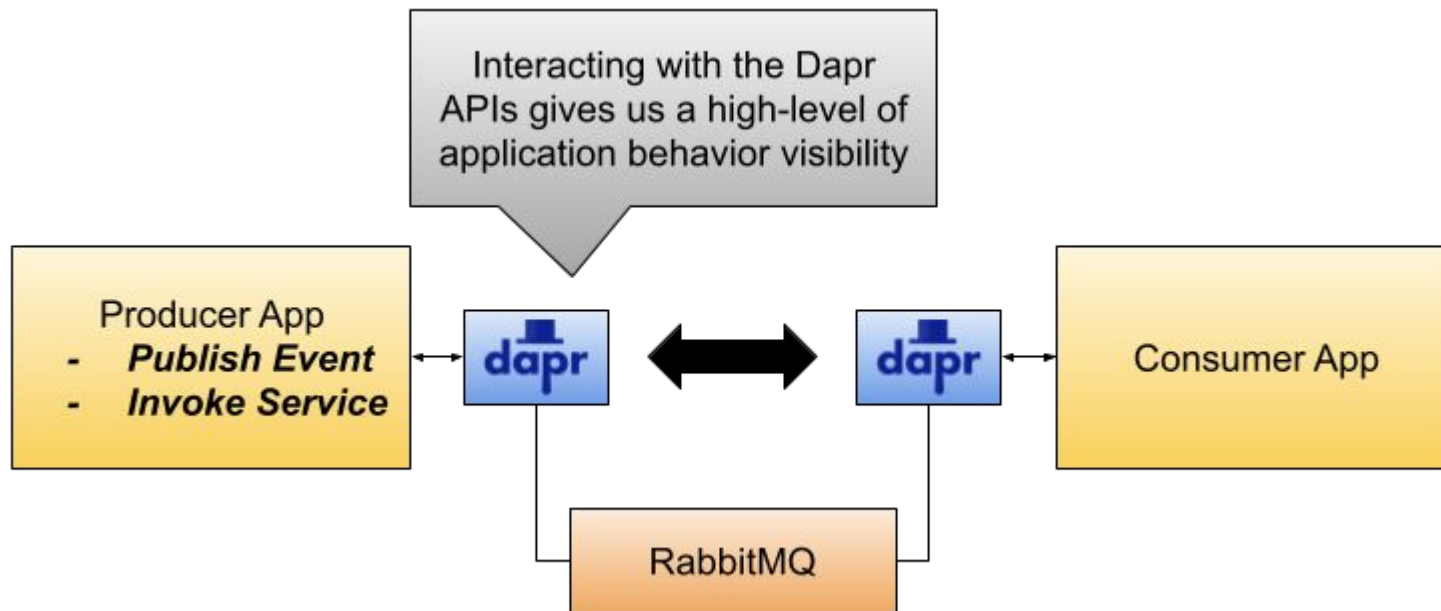


Dapr: enabling developers with APIs



- Stands for **D**istributed **a**pplication **R**untime
- Uses a proxy to expose application-level APIs to solve common distributed challenges
- We used two APIs for this examples:
 - Service to Service invocation
 - PubSub for async communications
- All APIs cover cross-functional concerns such as resiliency, observability and security

With Dapr




Demo #1

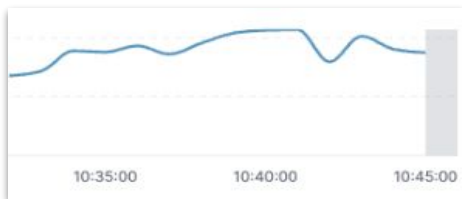
Let's use our apps!



The 3 Pillars of Observability



10:54:50.259	info ▾	GetSupportedCurren
10:54:50.259		Currencies fetched, response
10:54:50.259		Processing supported currenci
10:54:50.197	info ▾	GetSupportedCurren



Logs

Traces

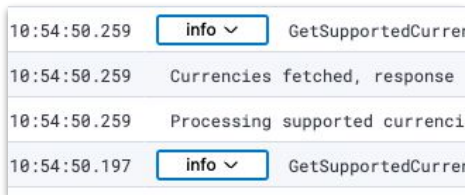
HTTP 2xx GET /cart 15 ms	
render route (pages) /cart	4.3 ms
resolve page components	185 µs



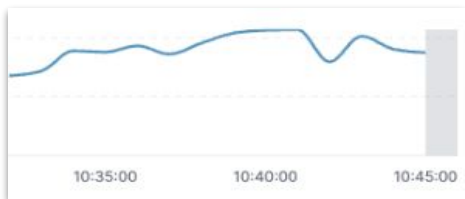
since 2019

Metrics

The 3 4 Pillars of Observability

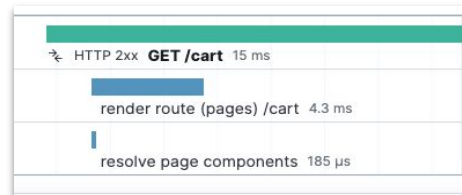


10:54:50.259	info	GetSupportedCurren
10:54:50.259		Currencies fetched, response
10:54:50.259		Processing supported currenci
10:54:50.197	info	GetSupportedCurren



Logs

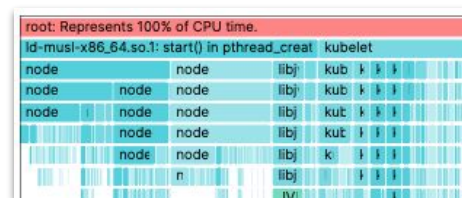
Traces



HTTP 2xx	GET /cart	15 ms
	render route (pages) /cart	4.3 ms
	resolve page components	185 µs

Metrics

Profiles
since 2022



root: Represents 100% of CPU time.			
id-musi-x86_64.so.1: start() in pthread_creat	kubelet		
node	node	libj	kub
node	node	libj	kub
node	node	libj	kut
node	node	libj	kut
node	node	libj	k
n	libj		

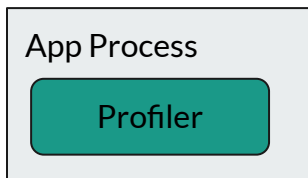


since 2019

Profiling

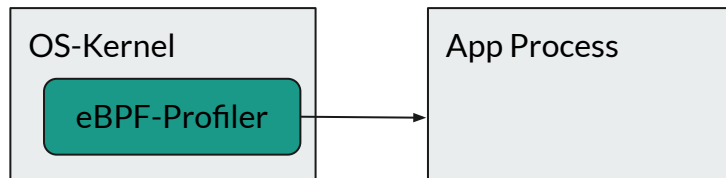
- Measuring where and how an application spends its time **without** having to modify/instrument it
- “Time” can be many things
 - CPU-Time, Wall-Clock, IO-time, ...
- Profiling sees the world from OS-perspective (Threads, processes, OS-resources)

In-Process Profiling



E.g. Linux Perf, Java Flight Recorder

eBPF Profiling



E.g. opentelemetry-ebpf-profiler

Continuous-Profiling in Production

- Optimize.cloud launches low-overhead multi-runtime zero-instrumentation profiler in 2021
- Acquired by Elastic soon after
- Donated to OpenTelemetry in 2024
- Continued development and evolution



[opentelemetry-ebpf-profiler](https://github.com/open-telemetry/opentelemetry-ebpf-profiler)



Whole-System Visibility

Unlock unknown-unknowns - from the kernel through userspace into high-level code, across multi-cloud workloads.



Polyglot Visibility

C/C++, Rust & Go (without debug symbols on host) PHP, Python, Java (or any JVM language), Ruby, DotNet, Perl & NodeJS.

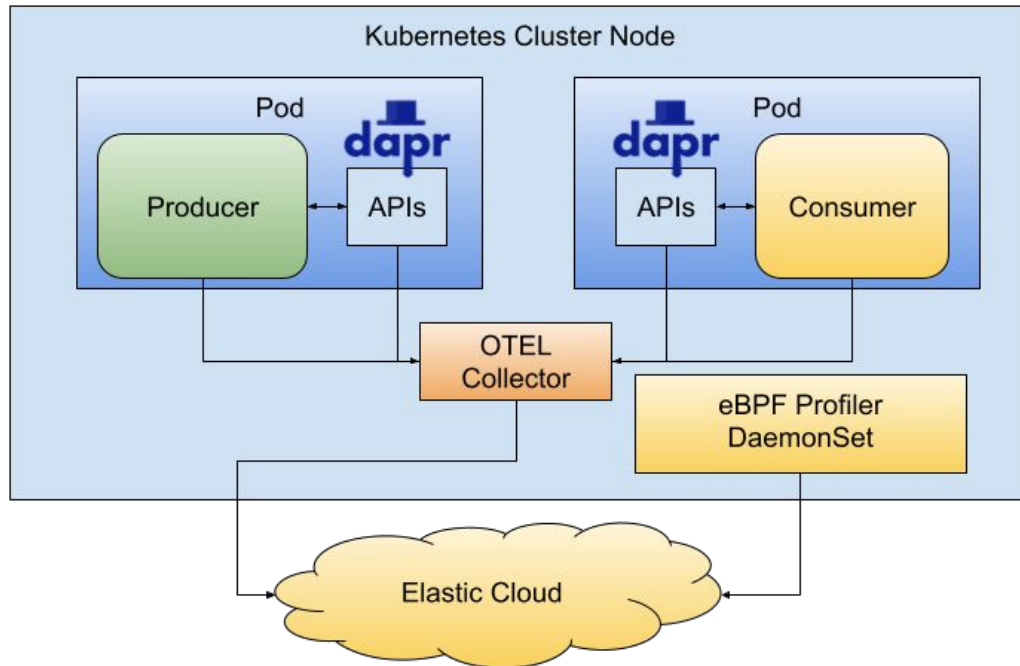


Extremely Low Overhead

Continuous profiling in production with negligible overhead.

Typical case: < 1% CPU, ~250MB of RAM

Recap deployment



Demo #2

Profiling in production





What's ahead

- Stabilization of the profiling signal (OTLP)
- Stabilization of profiling support in the OTEL-collector
- Standardization of trace - profiling correlation
- More than CPU-profiling (e.g. IO, page-faults,etc)

... and much more!

Thanks!

Jonas Kunz (Jonas Kunz @ CNCF-Slack)

Mauricio Salatino @diagridio @daprdev @salaboy

