
Is Everyone O-KEDA?

“Exciting” lessons learned in our journey
to use KEDA Pod Autoscaling

Brian Davis

Principal Software Engineer, Red Canary

X  @brianthedavis





PHASE 1 PHASE 2 PHASE 3

Implement
KEDA


?

Profit

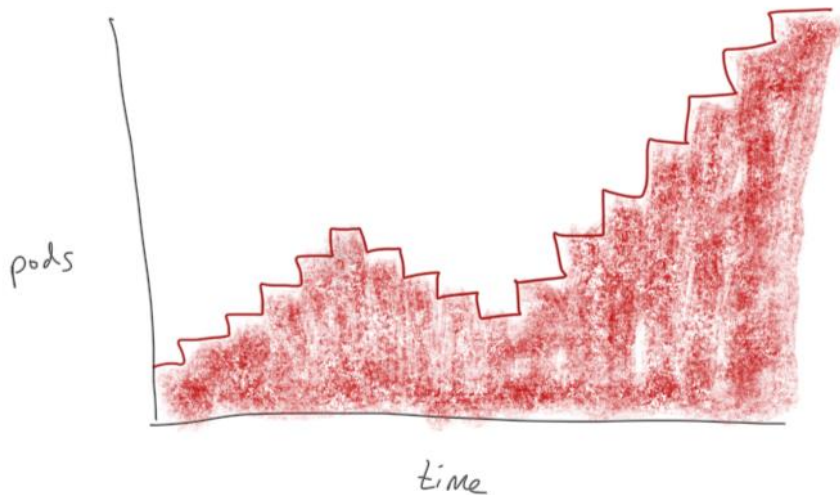
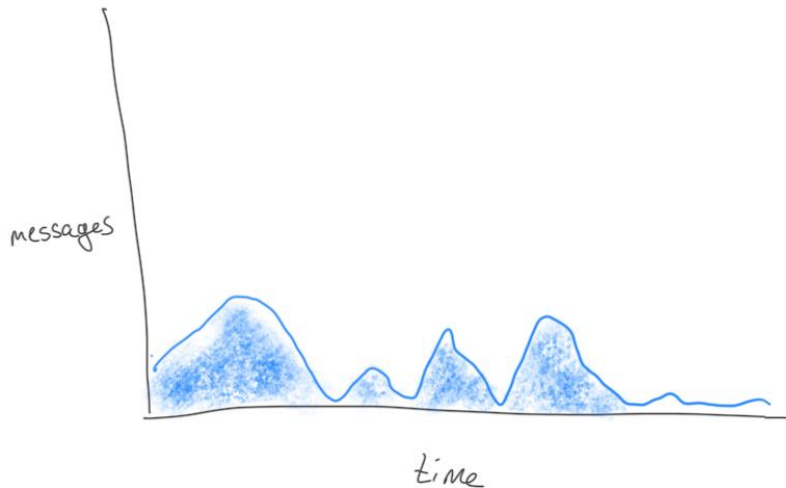


Five Phases of adopting KEDA

1. Accepting your current auto scaling is no good
2. Figuring out how to pronounce it
3. Implement it
4. Trough of despair
5. Profit





Our old scaling



Zero to N Scaling:

One of the coolest features of KEDA is its ability to scale your application back to zero when there are no events to process. This can lead to significant cost savings.

The new hotness

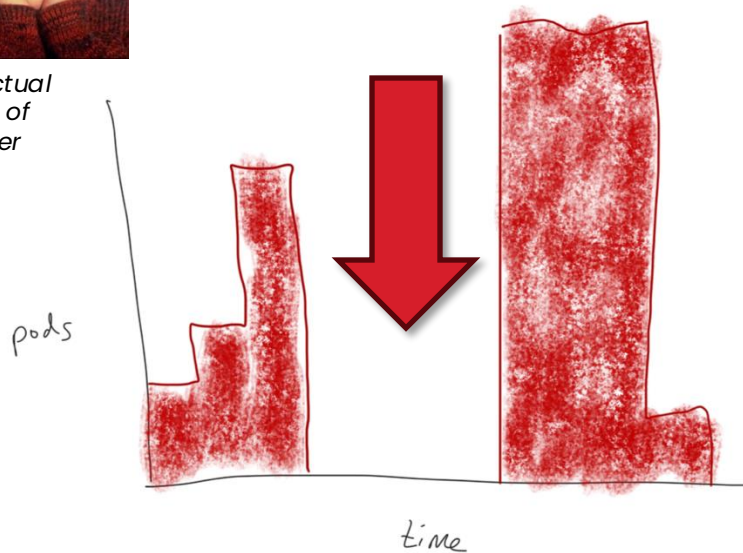



...until



**Not actual
picture of
engineer*

Engineer gets paged at 3am




Zero to N Scaling:

One of the coolest features of KEDA is its ability to scale your application back to zero when there are no events to process **OR THE K8S METRICS SERVER GETS EVICTED.**



"I'll take scalability for 1000, Alex"

- AWS STS Rate Limiting
- AWS S3 Rate Limiting





Tuning KEDA

Make messages wait around for *a little bit* to smooth out the peaks and valleys and waste less resources



The night club theory

You want a line out the nightclub – otherwise people won't see how cool it really is....



Are you O-KEDA?

KEDA is AWESOME!

...but

be prepared for unexpected issues
when you implement a new auto scaler

X  @brianthedavis

