

# Building Reliable Cross-Cloud Kubernetes Clusters on Spot Instances with Drafter and PVM

Felicitas Pojtinger  
@pojntfx



# Chapters

Commoditization

Silo

Architect

PVM

Conduit

Drafter





# Felicitas Pojtinger

Fediverse: @pojntfx@mastodon.social

Bluesky: @pojntfx.mastodon.social.ap.brid.gy

Github: @pojntfx

LinkedIn: in/pojntfx

Web: felicitas.pojtinger.com



[ABOUT SITE](#)[ABOUT ME](#)[NEW ESSAYS](#)[NEW LINKS](#)[PATREON](#)

# LAWS OF TECH: COMMODITIZE YOUR COMPLEMENT

[AI economics](#), [economics](#), [insight.porn](#), [Google](#)

*A classic pattern in technology economics, identified by Joel Spolsky, is layers of the stack attempting to become monopolies while turning other layers into perfectly-competitive markets which are commoditized, in order to harvest most of the consumer surplus; discussion and examples.*

2018-03-17 – 2022-01-11 · [finished](#) · [certainty: highly likely](#) · [importance: 5](#)

[backlinks](#) <sup>↗</sup> · [bibliography](#) <sup>≡</sup>

1 “Smart Companies Try To Commoditize Their Products’ Complements”

2 “Open Source As a Strategic Weapon”

3 Generalizing

4 Examples

5 See Also

6 External Links

7 Appendix

7.1 Information Rules

8 Footnotes

9 Backlinks

10 Bibliography

Joel Spolsky in 2002<sup>22ya</sup> identified a major pattern in technology business & economics: the pattern of “commoditizing your complement”, an alternative to vertical integration, where companies seek to secure a chokepoint or quasi-monopoly in products composed of many necessary & sufficient layers by dominating one layer while fostering so much competition in another layer above or below its layer that no competing monopolist can emerge, prices are driven down to marginal costs elsewhere in the stack, total price drops & increases demand, and the majority of the consumer surplus of the final product can be diverted to the quasi-monopolist. No matter how valuable the original may be and how much one could charge for it, it can be more valuable to make it



<https://gwer.net/complement>

# **"Smart Companies Try To Commoditize Their Products' Complements"**



LoopHole Labs

<https://gwern.net/complement>

Cars → Electricity/Gas

Computers →  
Software

Shipping  
company →  
Rail/Road





LoopHole Labs

- ✦ ✦ **Headline:** [Netscape<sup>w</sup>](#) [Open Sources<sup>w</sup>](#) [Their Web Browser<sup>w</sup>](#).
- ✦ **Myth:** They're doing this to get free source code contributions from New Zealand.
- ✦ **Reality:** They're doing this to commoditize the web browser. They *egy from day one*. Have a look at [the very first Netscape press release](#). Netscape gave away the browser so they could [sell](#) [servers](#) and servers are classic complements. The cheaper the browser, the more servers. This was never as true as it was in October 1994<sup>30ya</sup>. ..

- ✦ ✦ **Headline:** [Transmeta<sup>w</sup>](#) [Hires<sup>w</sup>](#) [Linus<sup>w</sup>](#), [Pays Him To Hack on<sup>w</sup>](#) [Linux<sup>w</sup>](#).
- ✦ **Myth:** They just did it to get publicity. Would you have heard of Transmeta otherwise?
- ✦ **Reality:** Transmeta is a CPU company. The natural complement of a CPU is an operating system. Transmeta wants OSs to be a commodity.
- ✦ ✦ **Headline:** [Sun<sup>w</sup>](#) [and HP<sup>w</sup>](#) [Pay<sup>w</sup>](#) [Ximian<sup>w</sup>](#) [To Hack on<sup>w</sup>](#) [Gnome<sup>w</sup>](#).
- ✦ **Myth:** Sun and HP are supporting free software because they like [Bazaars](#), [not Cathedrals<sup>w</sup>](#).
- ✦ **Reality:** Sun and HP are hardware companies. They make boxen. In order to make money on the desktop, they need for [windowing systems<sup>w</sup>](#), which are a complement of desktop computers, to be a commodity. Why don't they take the money they're paying Ximian and use it to develop a proprietary windowing system? They tried this (Sun had [NeWS<sup>w</sup>](#) and HP had [New Wave<sup>w</sup>](#)), but these are really hardware companies at heart with pretty crude software skills, and they need windowing systems to be a *cheap commodity*, not a proprietary advantage which they have to pay for. So they hired the nice guys at Ximian to do this for the same reason that Sun bought [Star Office<sup>w</sup>](#) and open sourced it: to commoditize software and make more money on hardware.

<https://gwer.net/complement>

**If you produce software, what is  
your complement?**



**Loophole Labs**





**HETZNER**

**How would commoditized compute  
look like?**



**Loophole Labs**

Can run  
**everything**

Can run  
**everywhere**




Can it run **everything**?



# VMs





VT-d (kvm\_intel)



AMD-V (kvm\_amd)



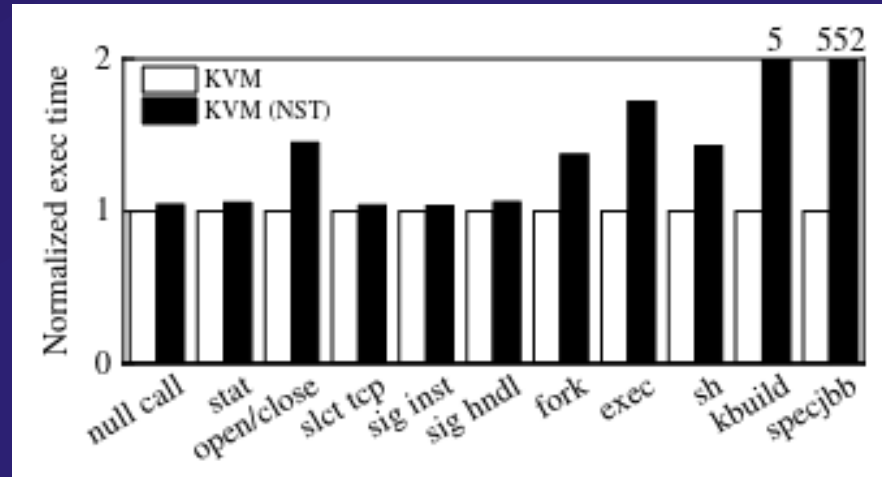


**HETZNER**





```
pojntfx@fels-dell-xps-13-plus:~  
pojntfx@fels-dell-xps-13-plus:~$ sudo modprobe kvm_intel nested=1  
pojntfx@fels-dell-xps-13-plus:~$ cat /sys/module/kvm_intel/parameters/nested  
Y  
pojntfx@fels-dell-xps-13-plus:~$
```



<https://huilucs.github.io/pubs/pvm.pdf>

# Component 1: **PVM**



linux-kernel.vger.kernel.org archive mirror

[help](#) / [color](#) / [mirror](#) / [Atom feed](#)

\* [RFC PATCH 00/73] KVM: x86/PVM: Introduce a new hypervisor

@ 2024-02-26 14:35 Lai Jiangshan

2024-02-26 14:35 ` [RFC PATCH 01/73] KVM: Documentation: Add the specification for PVM Lai Jiangshan  
` (74 more replies)

0 siblings, 75 replies; 82+ messages in thread

From: Lai Jiangshan @ 2024-02-26 14:35 UTC ([permalink](#) / [raw](#))

To: [linux-kernel](#)

Cc: Lai Jiangshan, Linus Torvalds, Peter Zijlstra,  
Sean Christopherson, Thomas Gleixner, Borislav Petkov,  
Ingo Molnar, [kvm](#), Paolo Bonzini, x86, Kees Cook, Juergen Gross,  
Hou Wenlong

From: Lai Jiangshan <[jiangshan.ljs@antgroup.com](mailto:jiangshan.ljs@antgroup.com)>


This RFC series proposes a new virtualization framework built upon the KVM hypervisor that does not require hardware-assisted virtualization techniques. PVM (Pagetable-based virtual machine) is implemented as a new vendor for KVM x86, which is compatible with the KVM virtualization software stack, such as Kata Containers, a secure container technique in a cloud-native environment.

The work also led to a paper being accepted at SOSP 2023 [[sosp-2023-acm](#)] [[sosp-2023-pdf](#)], and Lai delivered a presentation at the symposium in Germany in October 2023 [[sosp-2023-slides](#)]:

PVM: Efficient Shadow Paging for Deploying Secure Containers in Cloud-native Environment

PVM has been adopted by Alibaba Cloud and Ant Group in production to host tens of thousands of secure containers daily, and it has also been adopted by the Openanolis community.

[https://lore.kernel.org/lkml/  
CABgObfaSGOt4AKRF5WEJt2fGMj\\_hLXd7J2x2etce2ymv  
T4HkpA@mail.gmail.com/T/](https://lore.kernel.org/lkml/CABgObfaSGOt4AKRF5WEJt2fGMj_hLXd7J2x2etce2ymvT4HkpA@mail.gmail.com/T/)

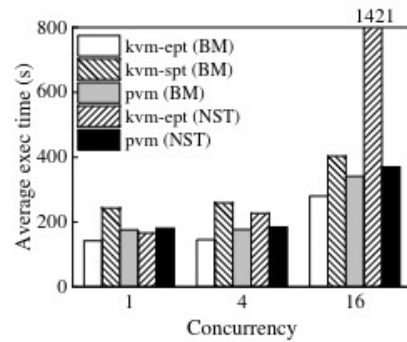


<https://huilucs.github.io/pubs/pvm.pdf>


Configurations	kvm (BM)	pvm (BM)	kvm (NST)	pvm (NST)
Hypercall	0.46/0.46	0.54/0.54	7.43/7.87	0.48/0.48
Exception	1.66/1.65	1.67/1.65	9.20/9.01	2.21/2.2
MSR access	0.87/0.87	2.53/2.51	8.18/8.47	2.88/2.86
CPUID	0.54/0.54	0.60/0.59	7.10/7.16	0.51/0.51
PIO	3.79/3.39	4.91/4.54	29.34/28.27	12.94/12.03




<https://huilucs.github.io/pubs/pvm.pdf>




(a) Kbuild (lower is better)



(b) Blogbench (higher is better)



(c) Specjbb2005 (higher is better)



(d) Fluidanimate (lower is better)



<https://huilucs.github.io/pubs/pvm.pdf>

# Can it run **everywhere**?

Spot Instances are available at a discount of up to **90% off** compared to On-Demand pricing. To compare the current Spot prices against standard On-Demand rates, visit the [Spot Instance Advisor](#).

A *Spot Instance interruption notice* is a warning that is issued **two minutes before Amazon EC2 stops or terminates your Spot Instance**. If you specify hibernation as the interruption behavior, you receive an interruption notice, but you do not receive a two-minute warning because the hibernation process begins immediately.



# Live Migration

## Component 2: **Firecracker**



[BENEFITS](#) [HOW IT WORKS](#) [FAQS](#) [LEARN MORE](#)

## Secure and fast microVMs for serverless computing

[VIEW ON GITHUB](#)

Firecracker is an open source virtualization technology that is purpose-built for creating and managing secure, multi-tenant container and function-based services.

Firecracker enables you to deploy workloads in lightweight virtual machines, called microVMs, which provide enhanced security and workload isolation over traditional VMs, while enabling the speed and resource efficiency of containers. Firecracker was developed at Amazon Web Services to improve the customer experience of services like [AWS Lambda](#) and [AWS Fargate](#).

Firecracker is a virtual machine monitor (VMM) that uses the Linux Kernel-based Virtual Machine (KVM) to create and manage microVMs. Firecracker has a minimalist design. It excludes unnecessary devices and guest functionality to reduce the memory footprint and attack surface area of each microVM. This improves security, decreases the startup time, and increases hardware utilization. Firecracker is generally available on [64-bit Intel, AMD and Arm CPUs with support for](#)

hub.io...



<https://firecracker-microvm.github.io/>



CPU template	CPU vendor	CPU model
C3	Intel	any
T2	Intel	any
T2A	AMD	Milan
T2CL	Intel	Cascade Lake or newer
T2S	Intel	any
V1N1	ARM	Neoverse V1

[https://github.com/firecracker-microvm/  
firecracker/blob/main/docs/cpu\\_templates/cpu-  
templates.md](https://github.com/firecracker-microvm/firecracker/blob/main/docs/cpu_templates/cpu_templates.md)



```
/snapshot/create:
  put:
    summary: Creates a full or diff snapshot. Post-boot only.
    description:
      Creates a snapshot of the microVM state. The microVM should be
      in the 'Paused' state.
    operationId: createSnapshot
    parameters:
      - name: body
        in: body
        description: The configuration used for creating a snapshot.
        required: true
        schema:
          $ref: "#/definitions/SnapshotCreateParams"
    responses:
      204:
        description: Snapshot created
      400:
        description: Snapshot cannot be created due to bad input
        schema:
          $ref: "#/definitions/Error"
    default:
      description: Internal server error
      schema:
        $ref: "#/definitions/Error"
```

```
/snapshot/load:
  put:
    summary: Loads a snapshot. Pre-boot only.
    description:
      Loads the microVM state from a snapshot.
      Only accepted on a fresh Firecracker process (before configuring
      any resource other than the Logger and Metrics).
    operationId: loadSnapshot
    parameters:
      - name: body
        in: body
        description: The configuration used for loading a snapshot.
        required: true
        schema:
          $ref: "#/definitions/SnapshotLoadParams"
    responses:
      204:
        description: Snapshot loaded
      400:
        description: Snapshot cannot be loaded due to bad input
        schema:
          $ref: "#/definitions/Error"
    default:
      description: Internal server error
      schema:
        $ref: "#/definitions/Error"
```

src/firecracker/swagger/firecracker.yaml



```
src/vmm/src/arch/x86_64/msr.rs

@@ -49,7 +49,7 @@ const APIC_BASE_MSR: u32 = 0x800;
49 49 /// Number of APIC MSR indexes
50 50 const APIC_MSR_INDEXES: u32 = 0x400;
51 51
52 - /// Custom MSRs fall in the range 0x4b564d00-0x4b564dff
52 + /// Custom KVM MSRs fall in the range 0x4b564d00-0x4b564def (0x4b564df0-0x4b564dff is reserved for PVM)
53 53 const MSR_KVM_WALL_CLOCK_NEW: u32 = 0x4b56_4d00;
54 54 const MSR_KVM_SYSTEM_TIME_NEW: u32 = 0x4b56_4d01;
55 55 const MSR_KVM_ASYNC_PF_EN: u32 = 0x4b56_4d02;
@@ -58,6 +58,6 @@ const MSR_KVM_PV_EOI_EN: u32 = 0x4b56_4d04;
58 58 const MSR_KVM_POLL_CONTROL: u32 = 0x4b56_4d05;
59 59 const MSR_KVM_ASYNC_PF_INT: u32 = 0x4b56_4d06;
60 60
61 + /// Custom PVM MSRs fall in the range 0x4b564df0-0x4b564dff
62 + const MSR_PVM_LINEAR_ADDRESS_RANGE: u32 = 0x4b56_4df0;
63 + const MSR_PVM_VCPU_STRUCT: u32 = 0x4b56_4df1;
64 + const MSR_PVM_SUPERVISOR_RSP: u32 = 0x4b56_4df2;
65 + const MSR_PVM_SUPERVISOR_REDZONE: u32 = 0x4b56_4df3;
66 + const MSR_PVM_EVENT_ENTRY: u32 = 0x4b56_4df4;
67 + const MSR_PVM_RETU_RIP: u32 = 0x4b56_4df5;
68 + const MSR_PVM_RET_S_RIP: u32 = 0x4b56_4df6;
69 + const MSR_PVM_SWITCH_CR3: u32 = 0x4b56_4df7;
70 +
61 71 /// Taken from arch/x86/include/asm/msr-index.h
62 72 /// Spectre mitigations control MSR
63 73 pub const MSR_IA32_SPEC_CTRL: u32 = 0x0000_0048;
@@ -237,6 +237,6 @@ static SERIALIZABLE_MSR_RANGES: &[MsRRange] = &[
237 247 MSR_RANGE!(MSR_KVM_POLL_CONTROL),
238 248 MSR_RANGE!(MSR_KVM_ASYNC_PF_INT),
239 249 MSR_RANGE!(MSR_IA32_TSX_CTRL),
250 + MSR_RANGE!(MSR_PVM_LINEAR_ADDRESS_RANGE),
251 + MSR_RANGE!(MSR_PVM_VCPU_STRUCT),
252 + MSR_RANGE!(MSR_PVM_SUPERVISOR_RSP),
253 + MSR_RANGE!(MSR_PVM_SUPERVISOR_REDZONE),
254 + MSR_RANGE!(MSR_PVM_EVENT_ENTRY),
255 + MSR_RANGE!(MSR_PVM_RETU_RIP),
256 + MSR_RANGE!(MSR_PVM_RET_S_RIP),
257 + MSR_RANGE!(MSR_PVM_SWITCH_CR3),
240 258 ];
241 259
242 260 /// Specifies whether a particular MSR should be included in vcpu serialization.
```

<https://github.com/loopholelabs/firecracker/pull/15/files#diff-646931758f35a261e2f848a0970552a6da048816eae0393b597d5830d857fba1>

```
SnapshotType::Msync | SnapshotType::MsyncAndState => {
  mark_queues_as_dirty(vmm);

  vmm.guest_memory().msync().map_err(MemoryMsync)
}
```



7 src/firecracker/swagger/firecracker.yaml		
@@ -1203,6 +1203,8 @@ definitions:		
1203	1203	enum:
1204	1204	- Full
1205	1205	- Diff
1206	+	- Msync
1207	+	- MsyncAndState
1206	1208	description:
1207	1209	Type of snapshot to create. It is optional and by default, a full
1208	1210	snapshot is created.
@@ -1238,6 +1240,11 @@ definitions:		
1238	1240	type: boolean
1239	1241	description:
1240	1242	When set to true, the vm is also resumed if the snapshot load is successful.
1243	+	shared:
1244	+	type: boolean
1245	+	description: When set to true and the guest memory backend is a file,
1246	+	changes to the memory are asynchronously written back to the
1247	+	backend as the VM is running.
1241	1248	
1242	1249	TokenBucket:
1243	1250	type: object

<https://github.com/loopholelabs/firecracker/pull/15/files#diff-646931758f35a261e2f848a0970552a6da048816eae0393b597d5830d857fba1>

# Component 3: **Silo**





[README](#) [Code of conduct](#) [AGPL-3.0 license](#)

# Silo

A Storage Primitive Designed for Live Migration

license **AGPL-3.0** [Loophole Labs](#) 262 members go version **>=1.2.1** [reference](#)

## Overview

Silo is a storage primitive designed to support live migration. One of the core functionalities within Silo is the ability to migrate/sync storage to various `backends` while it is still in use (without affecting performance).

## Sources

All storage sources within Silo implement `storage.StorageProvider`. You can find some example sources at [pkg/storage/sources](#).

## Expose

If you wish to expose a Silo storage device to an external consumer, one way would be to use the NBD kernel driver. See [pkg/expose/sources](#).

## Block orders

When you wish to move storage from one place to another, you'll need to specify an order. This can be dynamically changing. For example, there is a volatility monitor which can be used to migrate storage from least volatile to most volatile. Also you may wish to prioritize certain blocks for example if the destination is

<https://github.com/loopholelabs/silo>



```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo
~/Projects/silo

pojntfx@fels-dell-xps-13-plus:~/Projects/silo$ silo serve --help
Start up serve

Usage:
  silo serve [flags]

Flags:
  -a, --addr string    Address to serve from (default ":5170")
  -c, --conf string    Configuration file (default "silo.conf")
  -C, --continuous    Continuous sync
  -h, --help           help for serve
  -o, --order          Any order (faster)
  -p, --progress       Show progress
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$
```

```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo
~/Projects/silo

pojntfx@fels-dell-xps-13-plus:~/Projects/silo$ silo connect --help
Connect to a Silo instance, and stream available devices.

Usage:
  silo connect [flags]

Flags:
  -a, --addr string    Address to serve from (default "localhost:5170")
  -e, --expose         Expose as an nbd devices
  -h, --help           help for connect
  -m, --mount          Mount the nbd devices
  -p, --progress       Show progress
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$
```



```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo
~/Projects/silo


pojntfx@fels-dell-xps-13-plus:~/Projects/silo$ silo sync --help
Continuous sync to s3

Usage:
  silo sync [flags]

Flags:
  -a, --access string      S3 access
  -l, --blocksize int      S3 block size (default 1048576)
  -b, --bucket string      S3 bucket
  -c, --conf string        Configuration file (default "silo.conf")
      --dirtylimit int     Dirty block limit per period (default 16)
      --dirtymaxage duration Dirty block max age (default 1s)
      --dirtyminchanged int Dirty block min subblock changes (default 4)
      --dirtyperiod duration Dirty block check period (default 100ms)
  -d, --dirtyshift int     Dirty tracker block shift (default 10)
  -y, --dummy              Dummy destination
  -e, --endpoint string    S3 endpoint
  -h, --help               help for sync
  -r, --replay             Replay existing binlog(s)
  -s, --secret string      S3 secret
  -t, --timelimit duration Sync time limit (default 30s)
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$
```



Overview Status HowTo FAQ



# virtiofs

## Overview

Virtiofs is a shared file system that lets virtual machines access a directory tree on the host. Unlike existing approaches, it is designed to offer local file system semantics and performance.

Virtiofs was started at Red Hat and is being developed in the Linux, QEMU, FUSE, and Kata Containers open source communities.

See the [design document](#) for a more in-depth explanation of virtiofs.

## Status

Available in mainline since [Linux](#) 5.4, [QEMU](#) 5.0, [libvirt](#) 6.2, and [Kata Containers](#) 1.7.

The new [virtiofsd-rs](#) Rust daemon is receiving the most attention for new feature development.

## Community

Chat: [#virtiofs on Matrix](#)

Mailing list: [virtio-fs@lists.linux.dev](#) ([list info](#))

Community call: Bi-weekly on Wednesdays via [video conference](#) or [phone](#) (meeting ID 318831955). [Meeting times and agenda](#).

## HowTo

- [Sharing files with virtiofs using libvirt](#)
- [Installing virtiofs drivers on Windows](#)
- [Kata Containers with virtiofs](#)
- [Bootimg from virtiofs](#)

More documentation: [QEMU with virtiofs](#)



<https://virtio-fs.gitlab.io/>


# Component 4: **Drafter**



Loophole Labs






 README  AGPL-3.0 license



## Drafter

A Compute Primitive Designed for Live Migration

license **AGPL-3.0**  **LoopHole Labs** 262 members  hydrun CI **failing** go version **>= 1.21**  **reference**

### Overview

Drafter is a compute primitive with live migration support.

It enables you to:

- **Snapshot, package, and distribute stateful VMs:** With an opinionated packaging format and simple developer tools, managing, packaging, and distributing VMs becomes as straightforward as working with containers.
- **Run OCI images as VMs:** In addition to running almost any Linux distribution (Alpine Linux, Fedora, Debian, Ubuntu etc.), Drafter can also run OCI images as VMs without the overhead of a nested Docker daemon or full CRI implementation. It uses a dynamic disk configuration system, an optional custom Buildroot-based OS to start the OCI image, and a familiar Docker-like networking configuration.
- **Easily live migrate VMs between heterogeneous nodes with no downtime:** Drafter leverages a [custom optimized Firecracker fork](#) and [patches to PVM](#) to enable live migration of VMs between heterogeneous nodes, data centers and cloud providers without hardware virtualization support, even across continents. With a [customizable hybrid pre- and post-copy strategy](#), migrations typically take below 100ms within the same data center and around 500ms for Europe ↔ North America migrations over the public internet, depending on the application.

<https://github.com/loopholelabs/drafter>



```
resumedPeer, err := migratedPeer.Resume(  
    goroutineManager.Context(),  
  
    *resumeTimeout,  
    *rescueTimeout,  
  
    struct{}{},  
    ipc.AgentServerAcceptHooks[ipc.AgentServerRemote[struct{}], struct{}]{},  
  
    runner.SnapshotLoadConfiguration{  
        ExperimentalMapPrivate: *experimentalMapPrivate,  
  
        ExperimentalMapPrivateStateOutput: *experimentalMapPrivateStateOutput,  
        ExperimentalMapPrivateMemoryOutput: *experimentalMapPrivateMemoryOutput,  
    },  
)  
  
if err != nil {  
    panic(err)  
}  
  
defer func() {  
    defer goroutineManager.CreateForegroundPanicCollector()()  
  
    if err := resumedPeer.Close(); err != nil {  
        panic(err)  
    }  
}()  
}
```

cmd/drafter-peer/main.go

```

type ProxyClientRemote[G any] struct {
    GuestService G

    Dial func(ctx context.Context, connID, raddr string) error
    Write func(ctx context.Context, connID string, p []byte) (int, error)
    Close func(ctx context.Context, connID string) error
}

```

```

proxyService := proxy.NewProxyClient(
    ctx, // This will continue to live after any of the individual RPCs complete

    struct{}{},

    logger.SubLogger("ProxyClient"),

    func(ctx context.Context, network, address string) (io.ReadWriteCloser, error) {
        return (&net.Dialer{}).DialContext(ctx, network, address)
    },
)

```

```

agentClient := ipc.NewAgentClient(
    mtuService,

    func(ctx context.Context) error {
        logger.Info().Msg("Running pre-suspend command")

        if strings.TrimSpace(*beforeSuspendCmd) != "" {
            cmd := exec.CommandContext(ctx, *shellCmd, "-c", *beforeSuspendCmd)
            cmd.Stdout = os.Stdout
            cmd.Stderr = os.Stderr

            if err := cmd.Run(); err != nil {
                return err
            }
        }

        logger.Info().Msg("Running pre-suspend CRI service handler")

        return guest.CRIServiceBeforeSuspend(ctx, criService)
    },
    func(ctx context.Context) error {
        logger.Info().Msg("Running after-resume command")

        if strings.TrimSpace(*afterResumeCmd) != "" {
            cmd := exec.CommandContext(ctx, *shellCmd, "-c", *afterResumeCmd)
            cmd.Stdout = os.Stdout
            cmd.Stderr = os.Stderr

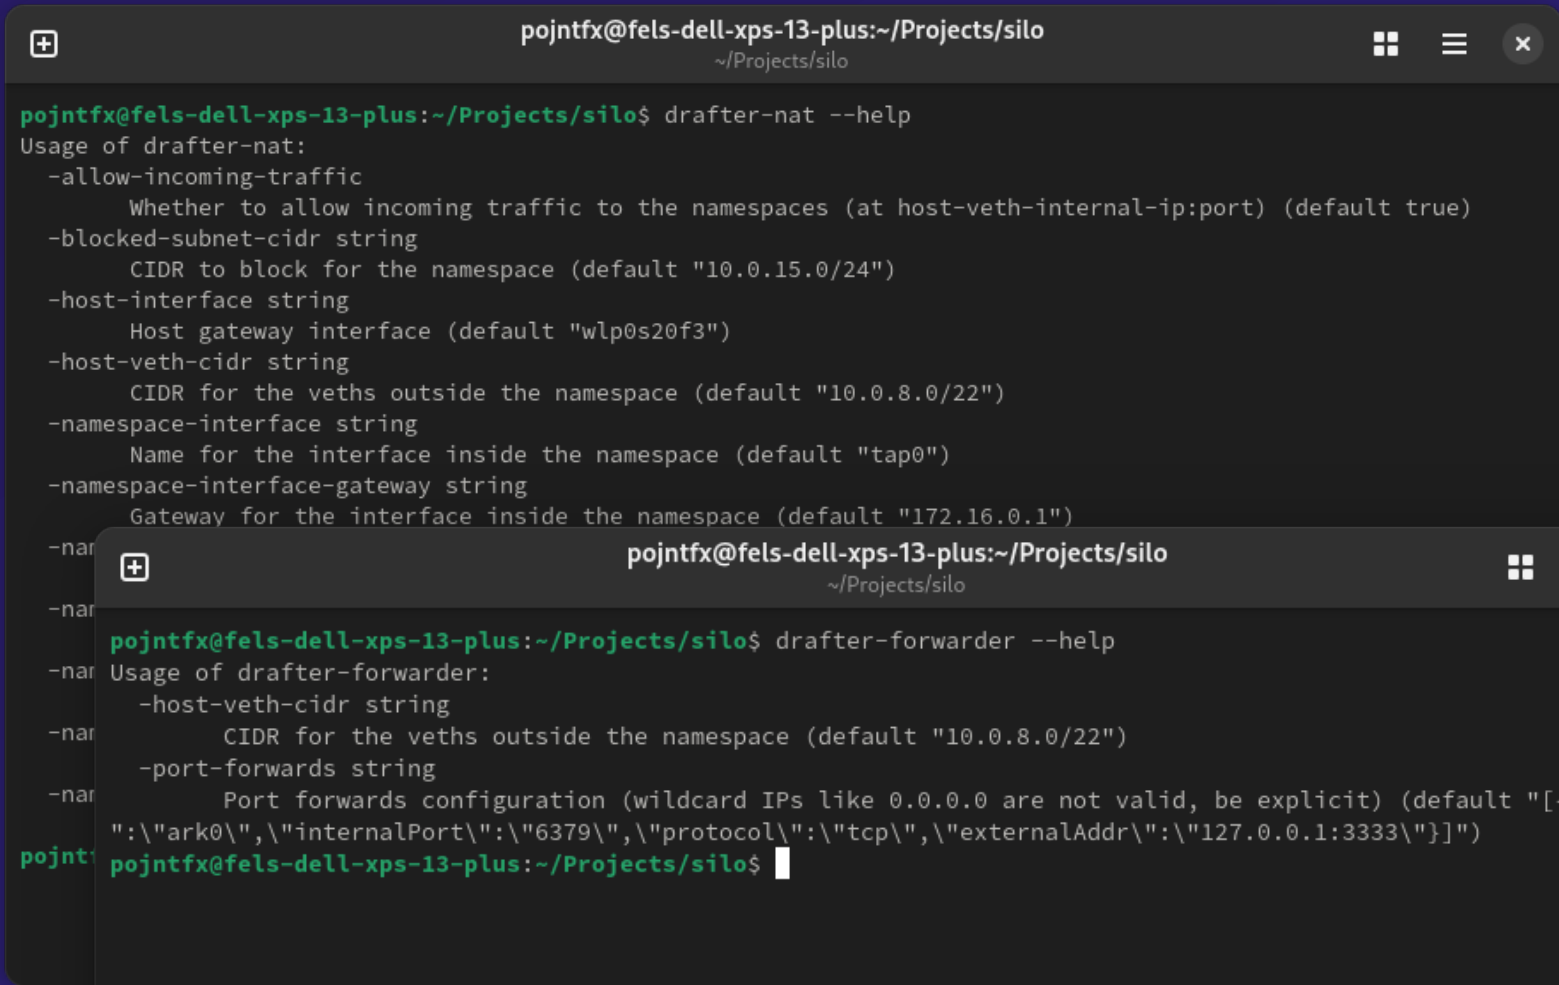
            if err := cmd.Run(); err != nil {
                return err
            }
        }

        return nil
    },
)

```









pojntfx@fels-dell-xps-13-plus:~/Projects/silo

~/Projects/silo



```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$ drafter-liveness --help
```

Usage of drafter-liveness:

-vsock-port int

VSocket port (default 25)

-vsock-timeout duration

VSocket dial timeout (default 1m0s)

```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$
```



pojntfx@fels-dell-xps-13-plus:~/Projects/silo

~/Projects/silo



```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$ drafter-agent --help
```

Usage of drafter-agent:

-after-resume-cmd string

Command to run after the VM has been resumed (leave empty to disable)

-before-suspend-cmd string

Command to run before the VM is suspended (leave empty to disable)

-shell-cmd string

Shell to use to run the before suspend and after resume commands (default "sh")

-vsock-port uint

VSocket port (default 26)

-vsock-timeout duration

VSocket dial timeout (default 1m0s)

```
pojntfx@fels-dell-xps-13-plus:~/Projects/silo$
```





```
pojntfx@fels-dell-xps-13-plus:~$ drafter-snapshotter --help
Usage of drafter-snapshotter:
  -agent-vsock-port int
    Agent VSocket port (default 26)
  -boot-args string
    Boot/kernel arguments (default "console=ttyS0 panic=1 pci=off modules=ext4 rootfstype=ext4 root=/dev/vda i8042.noaux i8042.nomux i8042.nopnp i8042.dumbkbd rootflags=rw printk.devkmsg=on printk_ratelimit=0 printk_ratelimit_burst=0 clocksource=tsc random.trust_cpu=on")
  -cgroup-version int
    Cgroup version to use for Jailer (default 2)
  -chroot-base-dir string
    chroot base directory (default "out/vms")
  -cpu-count int
    CPU count (default 1)
  -cpu-template string
    Firecracker CPU template (see https://github.com/firecracker-microvm/firecracker/blob/main/docs/cpu_templates/cpu-templates.md#static-cpu-templates for the options) (default "None")
  -devices string
    Devices configuration (default "[{\"name\":\"state\",\"input\":\"\",\"output\":\"out/package/state.bin\"},{\"name\":\"memory\",\"input\":\"\",\"output\":\"out/package/memory.bin\"},{\"name\":\"kernel\",\"input\":\"out/blueprint/vmlinux\",\"output\":\"out/package/vmlinux\"},{\"name\":\"disk\",\"input\":\"out/blueprint/rootfs.ext4\",\"output\":\"out/package/rootfs.ext4\"},{\"name\":\"config\",\"input\":\"\",\"output\":\"out/package/config.json\"},{\"name\":\"oci\",\"input\":\"out/blueprint/oci.ext4\",\"output\":\"out/package/oci.ext4\"}]")
  -enable-input
    Whether to enable VM stdin
  -enable-output
    Whether to enable VM stdout and stderr (default true)
  -firecracker-bin string
    Firecracker binary (default "firecracker")
  -pid int
```



pojntfx@fels-dell-xps-13-plus:~



```
pojntfx@fels-dell-xps-13-plus:~$ drafter-packager --help
```

```
Usage of drafter-packager:
```

```
-devices string
```

```
    Devices configuration (default "[{\"name\":\"state\",\"path\":\"out/package/state.bin\"},{\"name\":\"memory\",\"path\":\"out/package/memory.bin\"},{\"name\":\"kernel\",\"path\":\"out/package/vmlinux\"},{\"name\":\"disk\",\"path\":\"out/package/rootfs.ext4\"},{\"name\":\"config\",\"path\":\"out/package/config.json\"},{\"name\":\"oci\",\"path\":\"out/blueprint/oci.ext4\"}]")
```

```
-extract
```

```
    Whether to extract or archive
```

```
-package-path string
```

```
    Path to package file (default "out/app.tar.zst")
```

```
pojntfx@fels-dell-xps-13-plus:~$
```





```
pojntfx@fels-dell-xps-13-plus:~  
pojntfx@fels-dell-xps-13-plus:~$ drafter-runner --help  
Usage of drafter-runner:  
-cgroup-version int  
    Cgroup version to use for Jailer (default 2)  
-chroot-base-dir string  
    chroot base directory (default "out/vms")  
-devices string  
    Devices configuration (default "[{\"name\":\"state\",\"path\":\"out/package/state.bin\",\"shared\":false},{\"name\":\"memory\",\"path\":\"out/package/memory.bin\",\"shared\":false},{\"name\":\"kernel\",\"path\":\"out/package/vmlinux\",\"shared\":false},{\"name\":\"disk\",\"path\":\"out/package/rootfs.ext4\",\"shared\":false},{\"name\":\"config\",\"path\":\"out/package/config.json\",\"shared\":false},{\"name\":\"oci\",\"path\":\"out/blueprint/oci.ext4\",\"shared\":false}]")  
-enable-input  
    Whether to enable VM stdin  
-enable-output  
    Whether to enable VM stdout and stderr (default true)  
-experimental-map-private  
    (Experimental) Whether to use MAP_PRIVATE for memory and state devices  
-experimental-map-private-memory-output string  
    (Experimental) Path to write the local changes to the shared memory to (leave empty to write back to device directly) (ignored unless --experimental-map-private)  
-experimental-map-private-state-output string  
    (Experimental) Path to write the local changes to the shared state to (leave empty to write back to device directly) (ignored unless --experimental-map-private)  
-firecracker-bin string  
    Firecracker binary (default "firecracker")  
-gid int  
    Group ID for the Firecracker process  
-jailer-bin string  
    Jailer binary (from Firecracker) (default "iail")
```



```
pojntfx@fels-dell-xps-13-plus:~  
~  
pojntfx@fels-dell-xps-13-plus:~$ drafter-registry --help  
Usage of drafter-registry:  
-concurrency int  
    Number of concurrent workers to use in migrations (default 1024)  
-devices string  
    Devices configuration (default "[{"name":"state","base":"out/package/state.bin","overlay":"out/overlay/state.bin","state":"out/state/state.bin","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"memory","base":"out/package/memory.bin","overlay":"out/overlay/memory.bin","state":"out/state/memory.bin","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"kernel","base":"out/package/vmlinux","overlay":"out/overlay/vmlinux","state":"out/state/vmlinux","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"disk","base":"out/package/rootfs.ext4","overlay":"out/overlay/rootfs.ext4","state":"out/state/rootfs.ext4","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"config","base":"out/package/config.json","overlay":"out/overlay/config.json","state":"out/state/config.json","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"oci","base":"out/package/oci.ext4","overlay":"out/overlay/oci.ext4","state":"out/state/oci.ext4","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}])"  
-laddr string  
    Local address to listen on (leave empty to disable) (default "localhost:1337")  
-raddr string  
    Remote address to connect to (leave empty to disable) (default "localhost:1337")  
pojntfx@fels-dell-xps-13-plus:~$  
~  
pojntfx@fels-dell-xps-13-plus:~$ drafter-mounter --help  
Usage of drafter-mounter:  
-concurrency int  
    Number of concurrent workers to use in migrations (default 1024)  
-devices string  
    Devices configuration (default "[{"name":"state","base":"out/package/state.bin","overlay":"out/overlay/state.bin","state":"out/state/state.bin","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"memory","base":"out/package/memory.bin","overlay":"out/overlay/memory.bin","state":"out/state/memory.bin","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"kernel","base":"out/package/vmlinux","overlay":"out/overlay/vmlinux","state":"out/state/vmlinux","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"disk","base":"out/package/rootfs.ext4","overlay":"out/overlay/rootfs.ext4","state":"out/state/rootfs.ext4","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"config","base":"out/package/config.json","overlay":"out/overlay/config.json","state":"out/state/config.json","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}, {"name":"oci","base":"out/package/oci.ext4","overlay":"out/overlay/oci.ext4","state":"out/state/oci.ext4","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true}])"  
-laddr string  
    Local address to listen on (leave empty to disable) (default "localhost:1337")  
-raddr string  
    Remote address to connect to (leave empty to disable) (default "localhost:1337")  
pojntfx@fels-dell-xps-13-plus:~$
```



```
pojntfx@fels-dell-xps-13-plus:~  
pojntfx@fels-dell-xps-13-plus:~$ drafter-peer --help  
Usage of drafter-peer:  
-cgroup-version int  
    Cgroup version to use for Jailer (default 2)  
-chroot-base-dir string  
    chroot base directory (default "out/vms")  
-concurrency int  
    Number of concurrent workers to use in migrations (default 1024)  
-devices string  
    Devices configuration (default "[{"name":"state","base":"out/package/state.bin","overlay":"out/overlay/state.bin","state":"out/state/state.bin","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true,"shared":false}, {"name":"memory","base":"out/package/memory.bin","overlay":"out/overlay/memory.bin","state":"out/state/memory.bin","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true,"shared":false}, {"name":"kernel","base":"out/package/vmlinux","overlay":"out/overlay/vmlinux","state":"out/state/vmlinux","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true,"shared":false}, {"name":"disk","base":"out/package/rootfs.ext4","overlay":"out/overlay/rootfs.ext4","state":"out/state/rootfs.ext4","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true,"shared":false}, {"name":"config","base":"out/package/config.json","overlay":"out/overlay/config.json","state":"out/state/config.json","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true,"shared":false}, {"name":"oci","base":"out/package/oci.ext4","overlay":"out/overlay/oci.ext4","state":"out/state/oci.ext4","blockSize":65536,"expiry":1000000000,"maxDirtyBlocks":200,"minCycles":5,"maxCycles":20,"cycleThrottle":500000000,"makeMigratable":true,"shared":false}])")  
-enable-input  
    Whether to enable VM stdin  
-enable-output  
    Whether to enable VM stdout and stderr (default true)
```

Demo 1:




# Live Migrating Valkey/Redis between Cloud Providers with Drafter





Loophole Labs



# Component 5: **Conduit**









Bringing it all together:  **Architect**

Demo 2:

**Live Migrating a Kubernetes  
Cluster between Cloud Providers  
with Architect**



LoopHole Labs

Can run  
**everything**

Can run  
**everywhere**





```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: architect
handler: architect
```





```
apiVersion: v1
kind: Pod
metadata:
  name: xonotic-pod
  namespace: default
  labels:
    architect.run/migratable-id: my-migratable-xonotic-pod
  annotations:
    architect.run/ingress-port-mapping-xonotic: "26000:26000"
spec:
  runtimeClassName: architect
  nodeName: ip-172-31-52-200.us-west-2.compute.internal
  # nodeName: ip-172-31-61-246.us-west-2.compute.internal
  containers:
    - name: xonotic-container
      image: docker.io/loopholelabs/xonotic-demo:latest-antilag # Xonotic configured with `g_antilag 0`
      imagePullPolicy: Always
      resources: {}
      volumeMounts: []
      ports:
        - containerPort: 26000
          protocol: UDP
          hostPort: 26000
```



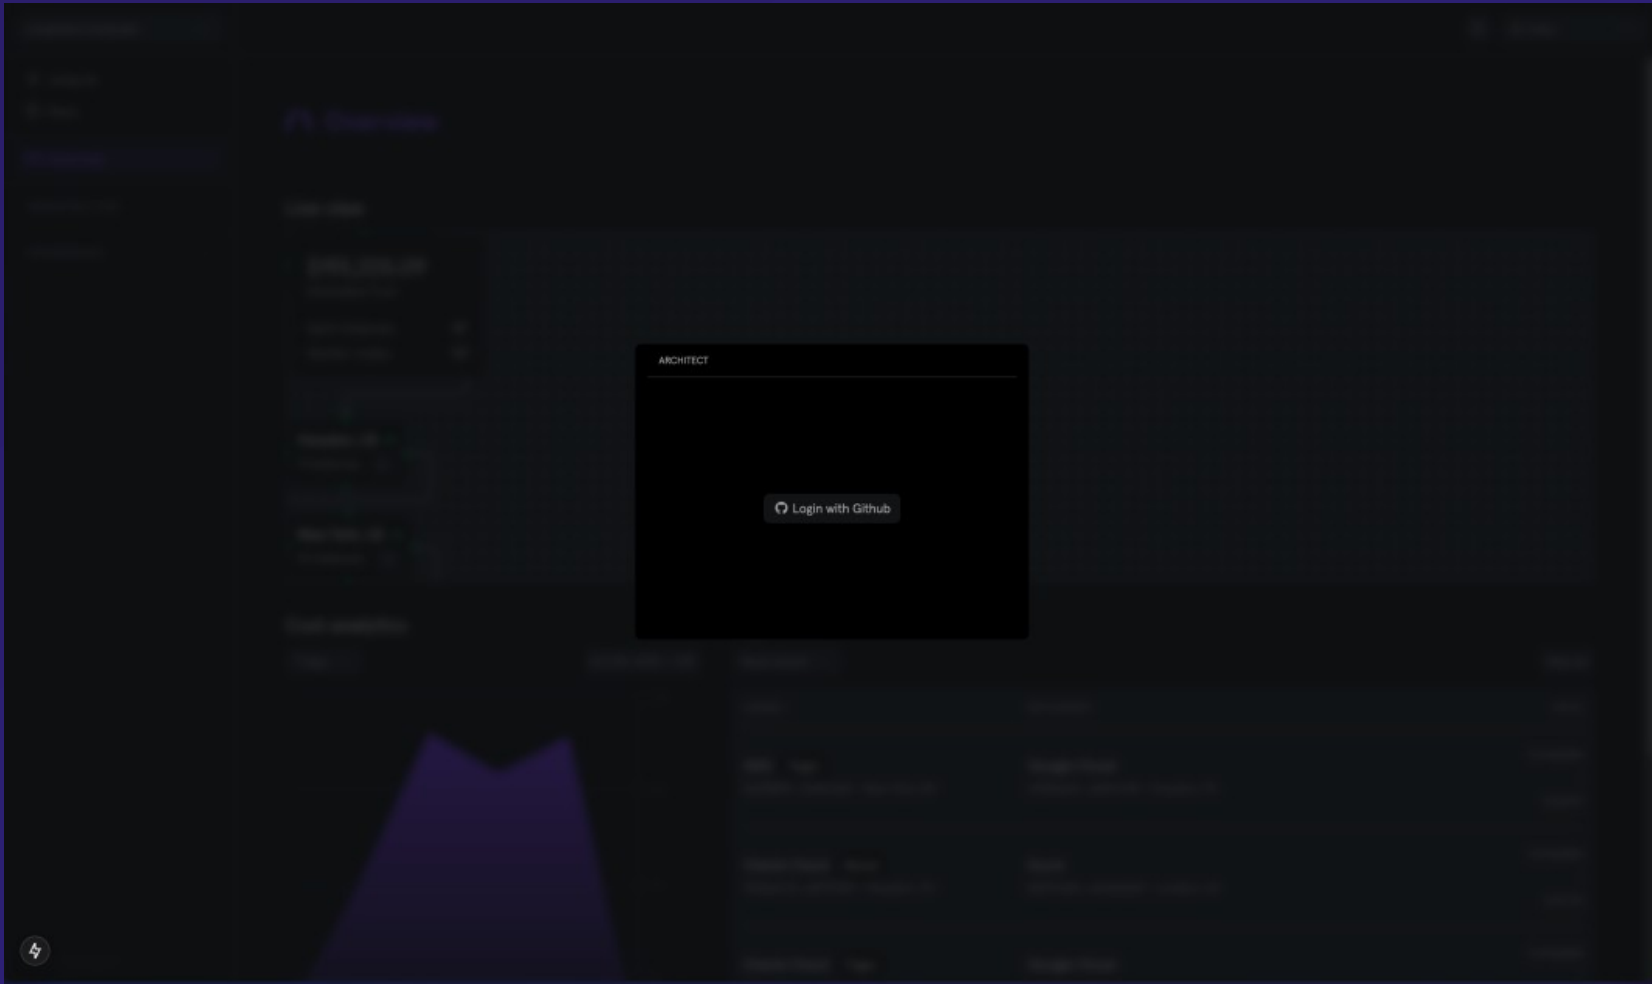
```
apiVersion: v1
kind: Pod
metadata:
  name: valkey-pod
  namespace: default
  labels:
    architect.run/migratable-id: my-migratable-valkey-pod
  annotations:
    architect.run/ingress-port-mapping-valkey: "6379:6379"
spec:
  runtimeClassName: architect
  nodeName: ip-172-31-52-200.us-west-2.compute.internal
  # nodeName: ip-172-31-61-246.us-west-2.compute.internal
  containers:
    - name: valkey-container
      image: quay.io/panquest/bitnami-valkey:7.2.6-debian-12-r3-1.8.2
      imagePullPolicy: Always
      env:
        - name: ALLOW_EMPTY_PASSWORD
          value: "yes"
      resources: {}
      volumeMounts: []
      ports:
        - containerPort: 6379
          protocol: TCP
          hostPort: 6379
```

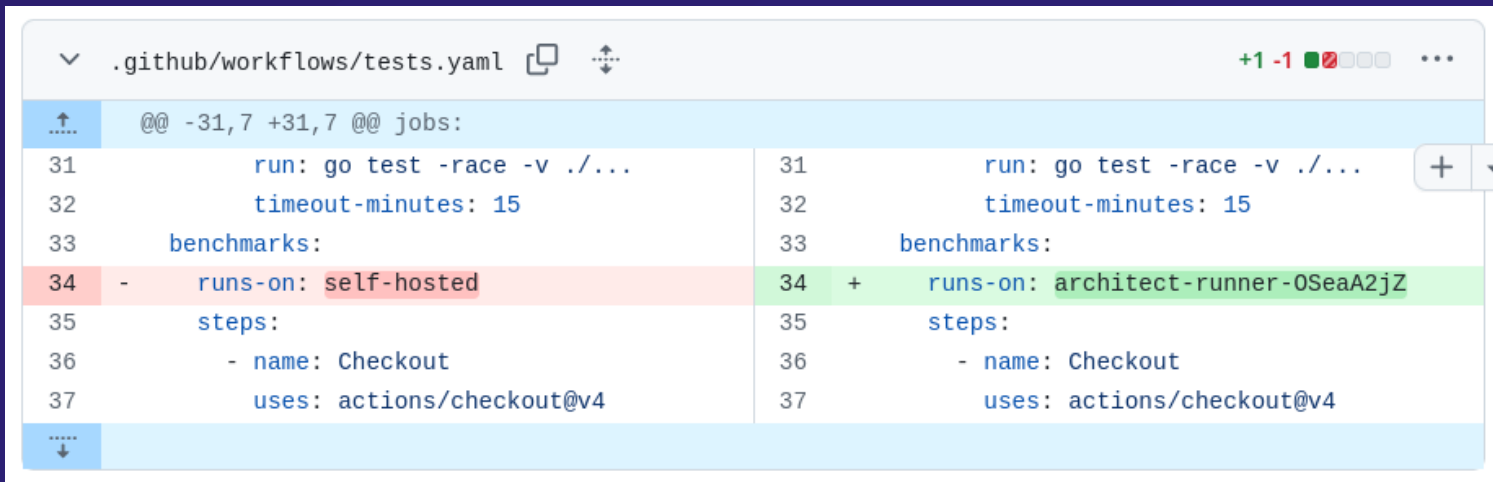
Demo 3:

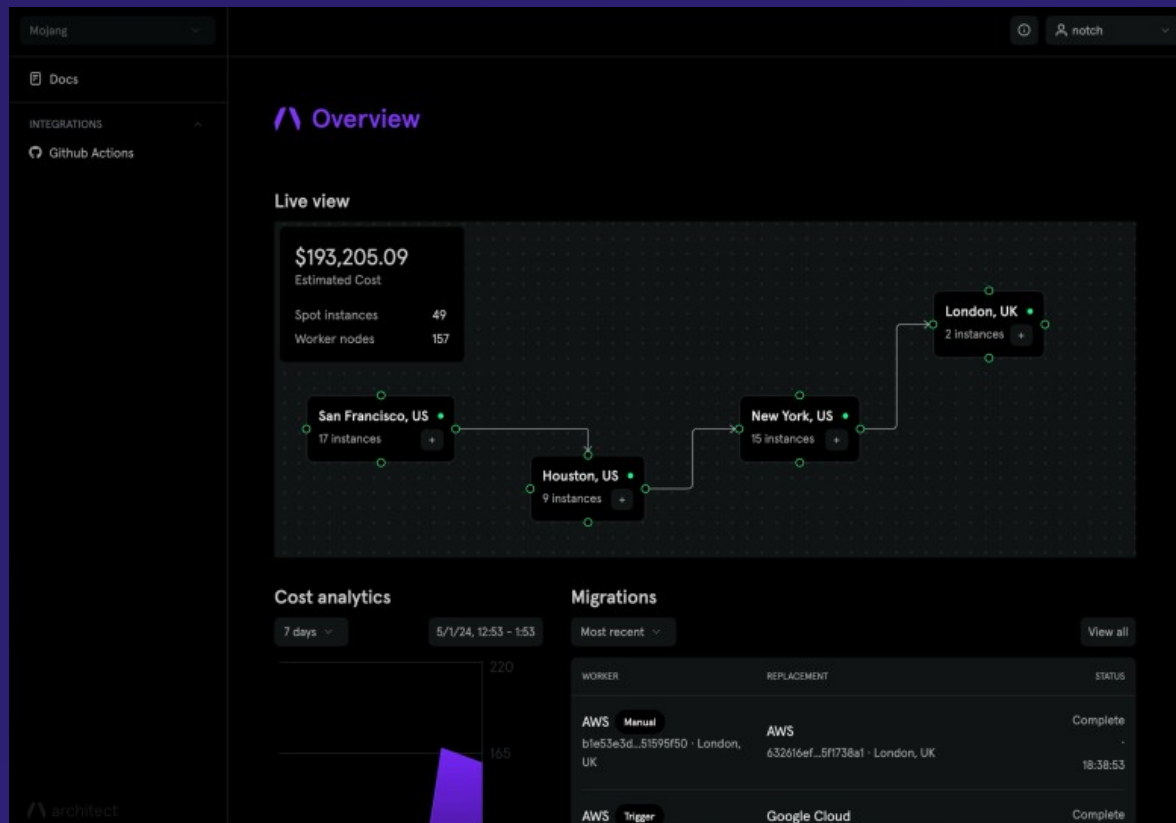
# GitHub Actions CI/CD on Spot Instances with Architect

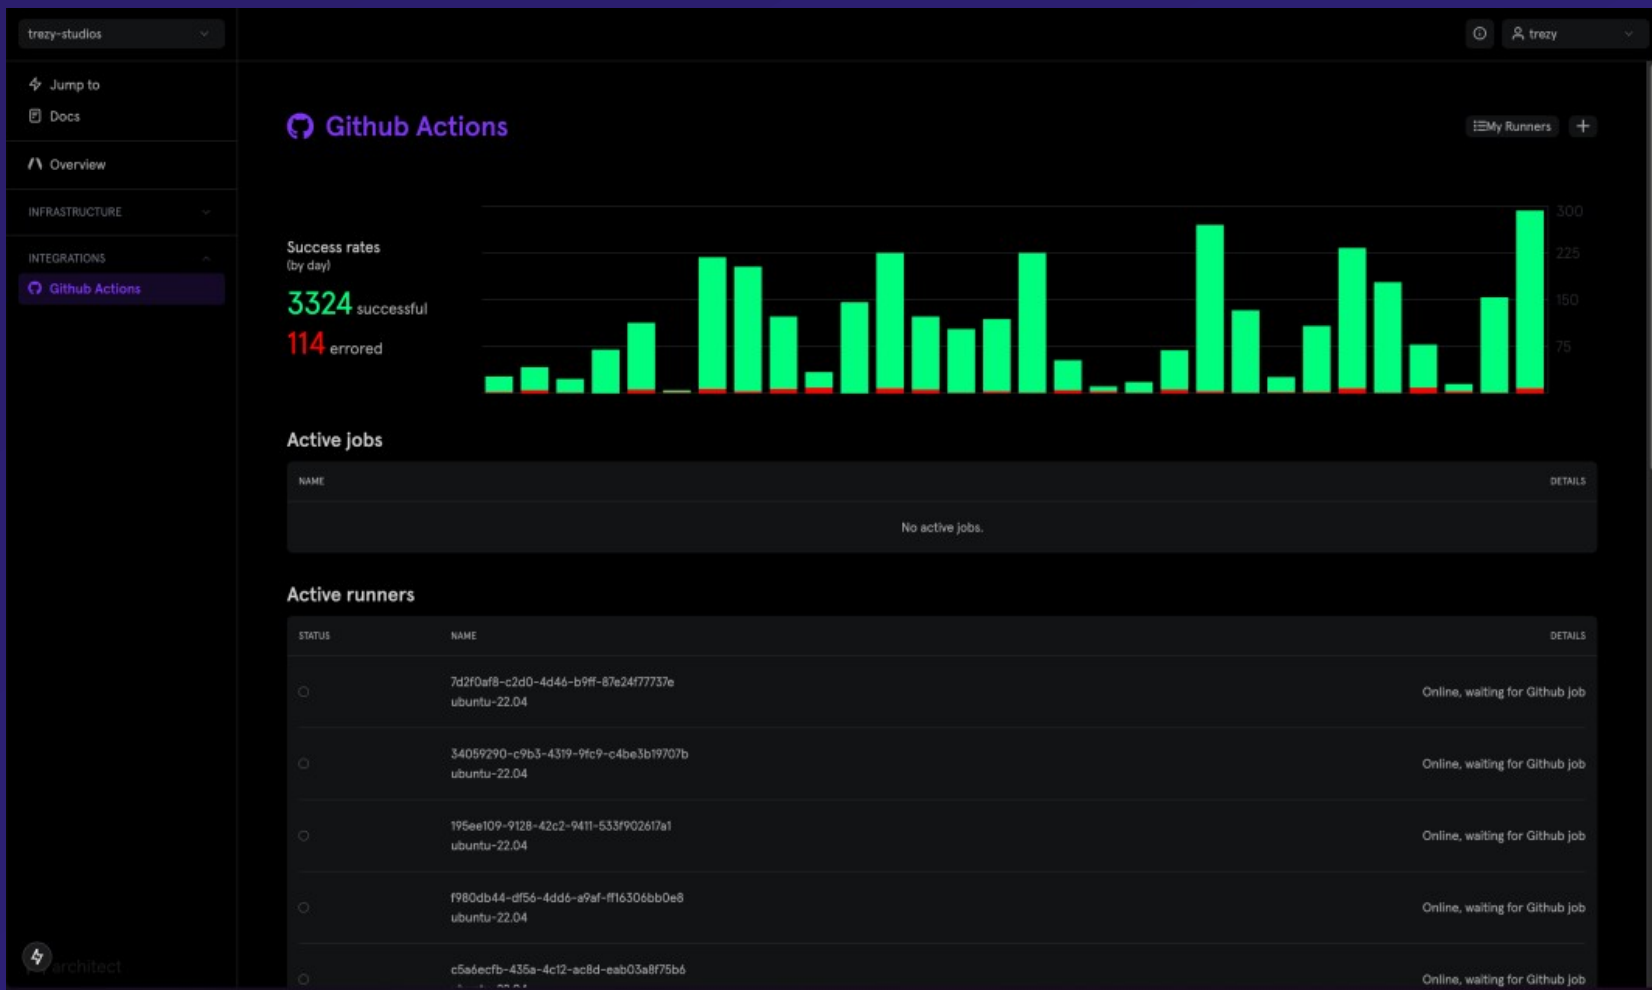


Loophole Labs











# Recap

# Recap

Commoditized Compute

PVM

Silo

Architect

Conduit

Drafter



# Links and Resources

<https://loophole.sh/kubecon2024>





# Felicitas Pojtinger

Fediverse: @pojntfx@mastodon.social

Bluesky: @pojntfx.mastodon.social.ap.brid.gy

Github: @pojntfx

LinkedIn: in/pojntfx

Web: felicitas.pojtinger.com



Introducing Scale Functions

# Changing the Way Developers Think About Networking

Modern application delivery for developers and DevOps teams. From Open-Source to Enterprise.

Stay in the Loop

We care about the protection of your data. Read our [Privacy Policy](#).

[Check Out Our Blog](#) →

Used By Developers at the world's best companies

Google



Dgraph

Berkeley  
UNIVERSITY OF CALIFORNIA



# Use Spot Instances for Any Workload

Architect lets you take advantage of Spot Instances for any workload on any Cloud – saving you 90% on your compute costs for even the most demanding tasks like Data Pipelines, Machine Learning, and CI/CD.

[JOIN WAITLIST](#)


Game Servers



Spot



Data Pipelines



Spot



Video Encoding



## WHAT ARE SPOT INSTANCES?

### Spot Instances are 90% Cheaper than On-Demand Instances

Spot Instances are unused cloud capacity available at steep discounts. They offer the same performance as On-Demand Instances at a fraction of the cost – and they don't require monthly commits.

On Demand



<https://architect.run/>