



KubeCon



CloudNativeCon

North America 2024

# Does My K8s Application Need CPR?

Performance Evaluation of a Multi-Cluster  
Workload Management Application

**Braulio Dumba & Ezra Silvera, IBM Research**

# Agenda

- Motivation
- KubeStellar: Manage workloads across multi-clusters
- Testing framework & Experiments examples
- Lessons learned & Conclusion

- Using multiple k8s clusters becomes common
  - Isolation: environments (dev, prod) , teams, etc..
  - Compliance with enterprise security or data governance requirements
  - Access to heterogeneous resources (GPUs, special HW, etc.)
- Multiple solutions already exist (e.g., KubeStellar, Karmada, OCM, etc.)
- Performance is a key factor for **managing workloads** across clusters
  - Users expect timely application deployment & status updates
  - Critical for adoption & success

# Performance Evaluation Challenges

- Multi-cluster workload management applications are complex
  - Solutions include many components in different locations
  - Require multiple infrastructure configurations
- Most of the popular performance evaluation tools target single cluster
  - Might need to extend/enhance existing tools
- Scale-testing: need to provision many **clusters** → resource consuming
  - Emulation may not be enough for all test cases
- Such solutions offer a large number of configurable settings
  - Significantly increases the number of test options

# KubeStellar: Background

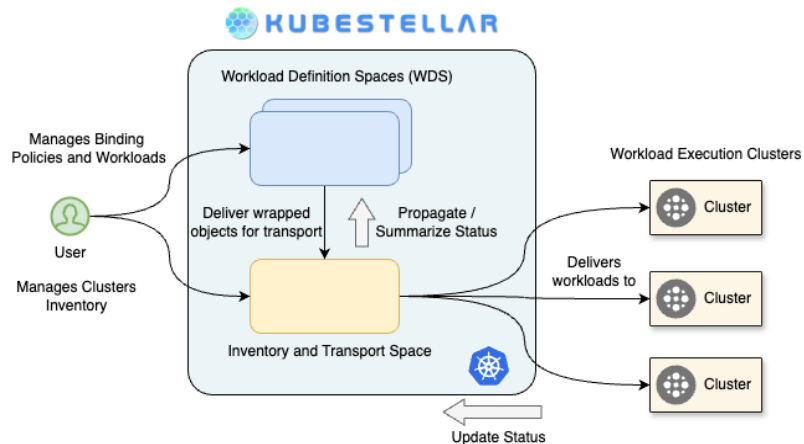
- A CNCF Sandbox project
- Deploy, configure and manage workloads across multiple K8s clusters
  - Use “native” K8s interfaces
  - Policy-based placement, workload customization, status summarization
- Find more info: <https://kubestellar.io/>



# KubeStellar: Architecture Overview

- Based on "spaces" abstraction for isolation and multi-tenancy
  - Space: behaves like a regular k8s cluster. Exposes k8s API end-point
- Pluggable transport framework
  - Propagate resources from WDS to the execution cluster(s)
  - Currently using OCM (<https://open-cluster-management.io/>)
- Customize resources deployed to WEC
- Collect & summarize statuses from WECs

**WEC:** Workload Execution cluster  
**WDS:** Workload Distribution Space,  
**ITS:** Inventory and Transport Space



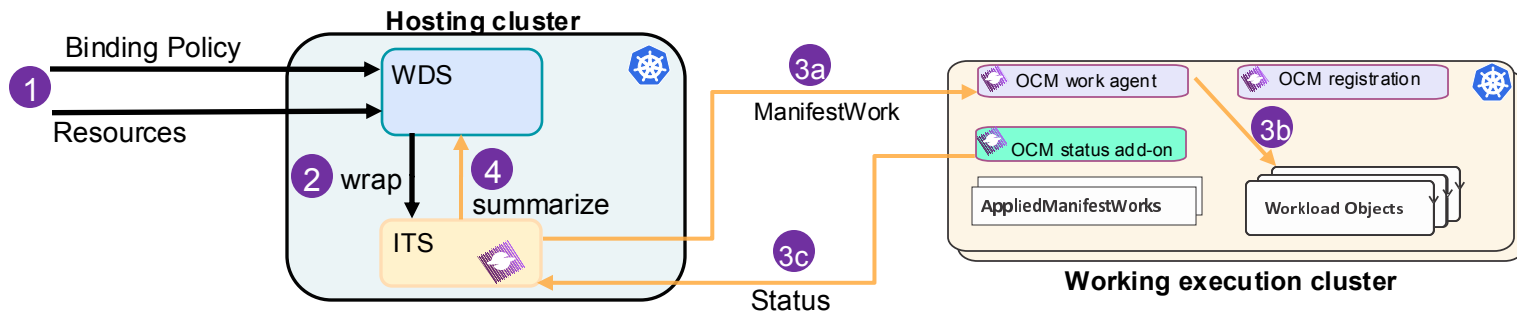
# KubeStellar: Binding Policy (BP)

- Defines “what goes where”
- Associates a subset of workload objects in the WDS with a subset of WECs
  - Defined by the user in the WDS
  - Performance implications
- Allows customization of how workload objects are down-synced

```
apiVersion:
control.kubestellar.io/v1alpha1
kind: BindingPolicy
metadata:
  name: nginx
spec:
  clusterSelectors:
  - matchLabels:
    location-group: edge
  downsync:
  - objectSelectors:
    - matchLabels:
      app.kubernetes.io/name: nginx
  .....
```

# KubeStellar Example: Deploy Workloads

1. User:
  - Deploys workload/resources into WDS
  - Defines the binding policy (desired placement)
2. KS: Transport controller pushes resources into the inventory space (ITS)
  - a. Wraps resources into a container object (e.g., ManifestWork)
3. Transport mechanism (OCM)
  - a. Distributes wrapped resources into the WEC
  - b. Unwraps & deploys resources in the WEC
  - c. Status is returned through the KS transport status plugin
4. KS: Propagates status into WDS (user facing objects)

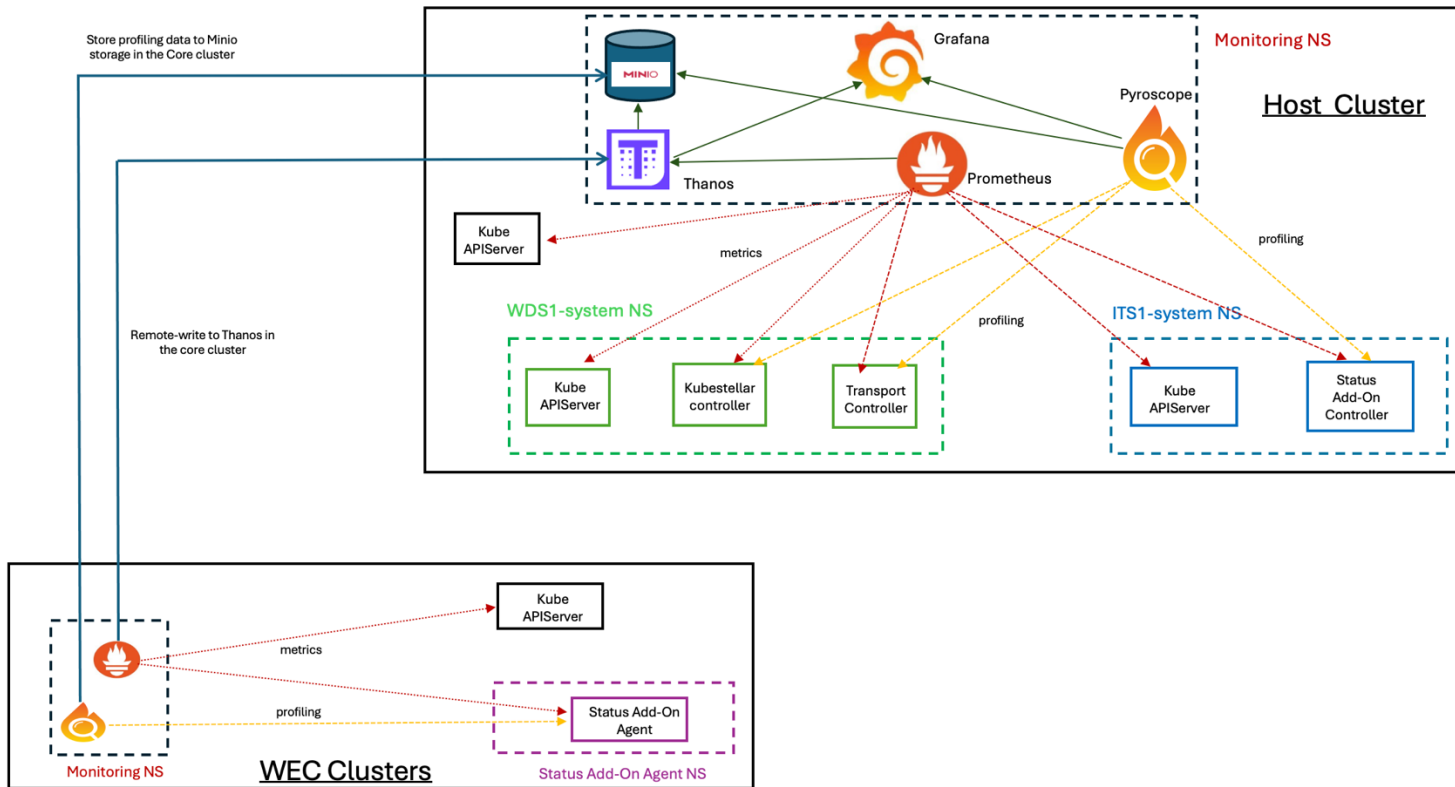




## Leverage existing tools when possible

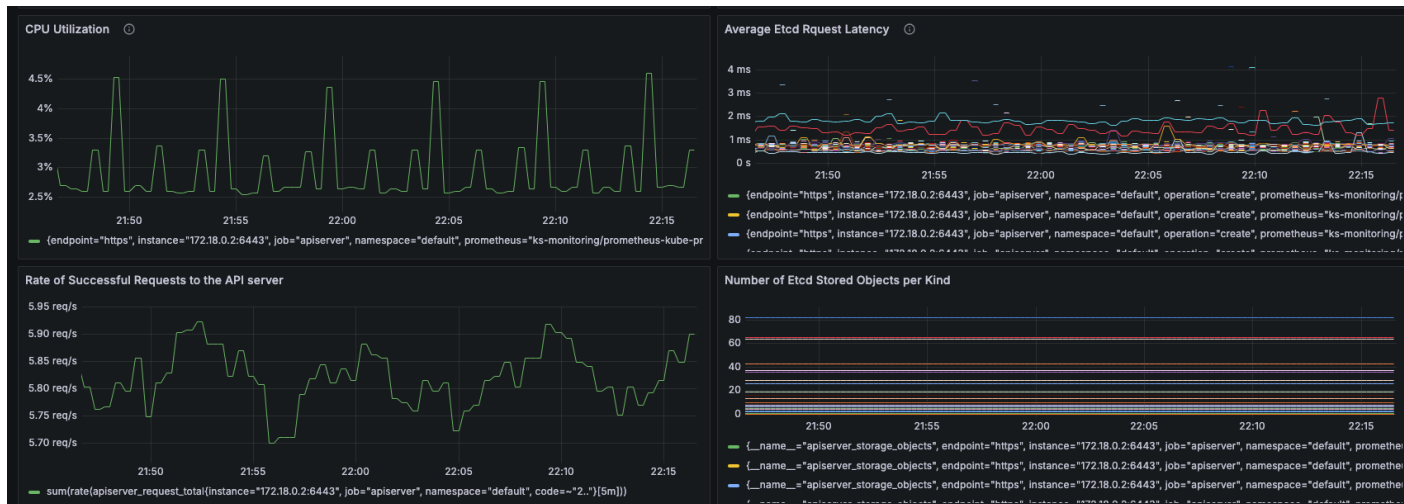
- **ClusterLoader2:** Kubernetes cluster performance test tool
- **Prometheus/Thanos, Pyroscope & Grafana:** Monitoring, profiling, visualization
- **Kind:** Running Kubernetes clusters using Docker container "nodes"
- **Kwok:** Kubernetes WithOut Kubelet, simulating any number of nodes and maintain pods on those nodes
- **Kube-burner:** Performance and scale test orchestration toolset. Mimic production workloads used to stress Kubernetes clusters

# Framework Architecture: Monitoring Setup



# Framework Architecture: Metrics Collection

- Targets: KS control plane (e.g., WDS & ITS API servers, transport, etc.)
- Single pane of glass to monitor KS hub and remote clusters



**Grafana dashboards to monitor API server, APF and KS controllers**

Used two profiles from Kube-burner reference workloads

## Workload benchmark for plain **Kubernetes** environments

Cluster-density profile
<b>1 deployments</b> , with two pod replicas (pause), mounting 2 secrets, 2 config maps
<b>3 services</b> , the first service points to the TCP/8080 port of the deployments
<b>10 secrets</b> containing a 2048-character random string
<b>10 configmaps</b>

## Workload benchmark for **OpenShift** environments


Cluster-density-ms profile
<b>1 image stream</b>
<b>4 deployments</b> , each with two pod replicas (pause), mounting 4 secrets, 4 config maps, and 1 downward API volume each
<b>2 services</b> , each pointing to the TCP/8080 and TCP/8443 ports of the first and second deployments
<b>1 edge</b> route pointing to the first service
<b>20 secrets</b> containing a 2048-character random string
<b>10 configmaps</b> containing a 2048-character random string

# Workload Generation

- Poisson Workload generator function:
  - Introduce some randomness in the workload
- Extended ClusterLoader2 (CL2)
  - Support poisson distribution tuning set\*
  - Add KubeStellar provider\*\*

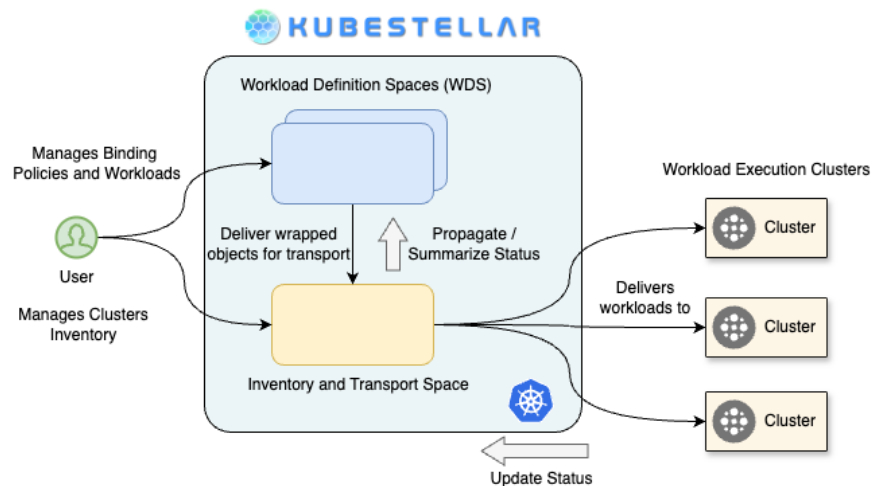
\* Merged PR: <https://github.com/kubernetes/perf-tests/pull/2633>

\*\* Merged PR: <https://github.com/kubernetes/perf-tests/pull/2632>



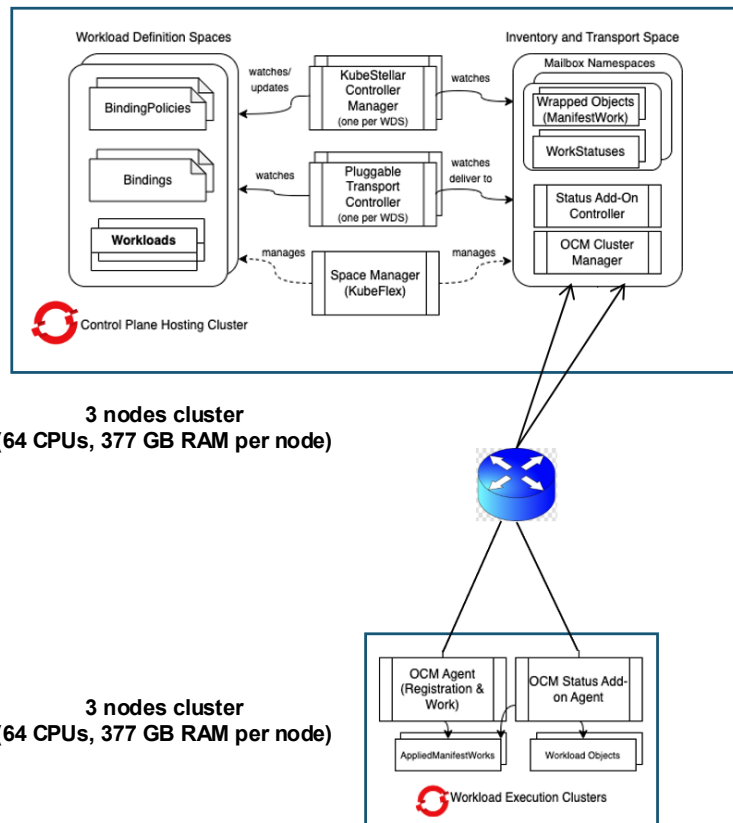
Selected KubeStellar performance parameters:

- **Number of binding Policies**
- **Number of workload execution clusters (WEC)**
- Number of workload description space (with shared ITS)
- Number of workload description space (with dedicated ITS)
- Workload size per binding policy



# Experiment 1: Number of Binding Policies

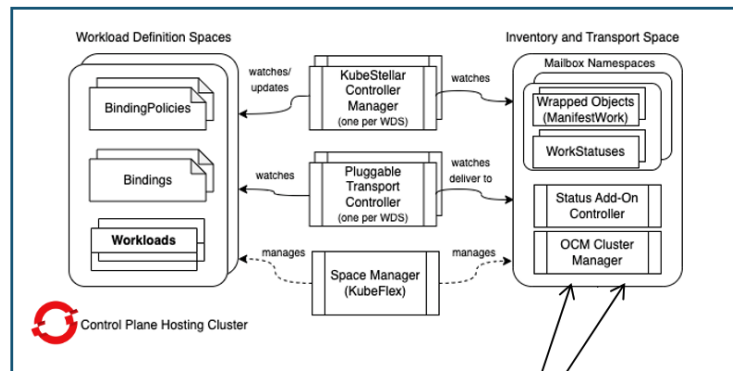
- Setup:
  - 2 OCP clusters
- Workload generator function
  - clusterloader2 **Poisson** tuningSet (alpha=0.5)



# Experiment 1: Number of Binding Policies

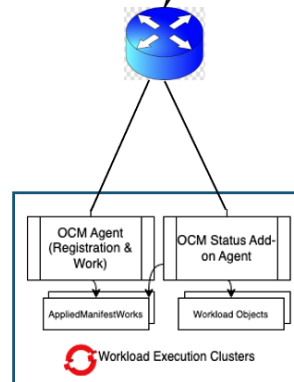
- Setup:
  - 2 OCP clusters
- Workload:
  - cluster-density-ms

Object type	Total #
BindingPolicies	150
Namespaces	150
Image stream	150
Deployments	600
Secrets	3000
Services	300
Configmaps	1500
routes	150



**3 nodes cluster**  
(64 CPUs, 377 GB RAM per node)

**3 nodes cluster**  
(64 CPUs, 377 GB RAM per node)






# Experiment 1: Number of Binding Policies - Results

Issue detected: controllers fight (component: KS transport controllers)

**Before fixes**



**After fixes**




# Experiment 2: Long Running Test

- Setup:
  - 2 OCP clusters
- Long running measurements: 1 binding Policy
  - Workload: pod that sleeps for 20 seconds
  - Workload generator function: clusterloader **RSteppedLoad** tuningSet (burstSize=1 and stepDelay=60 sec)
  - Custom controller: deletes a pod after reaching the completed state – another pod is created after 1 minute with a different name
  - Experiment duration (48 hours): 10/13 - 10/15


# Experiment 2: Long Running Test - Results

Issue detected: memory leak (component: status-agent controller)

## Before fixes



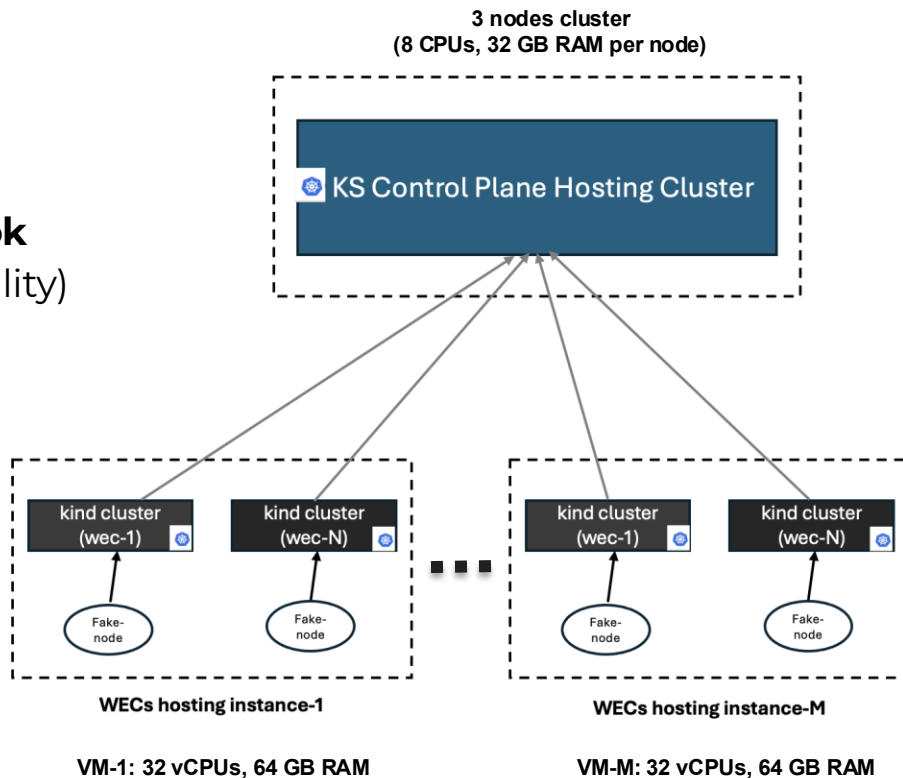
## After fixes



# Experiment 3: Number of WECs


- Setup:

- Ubuntu 24.04 VMs ( $M = 5$ )
- **100 kind clusters** ( $N = 20$ )
- Emulation of nodes & pods with **kwok**
- Using Ansible automation (repeatability)




# Experiment 3: Number of WECs (Cont.)

- Measurements: E2E & down-sync latencies, resource utilization, etc.
- Workload: kwok fake pod



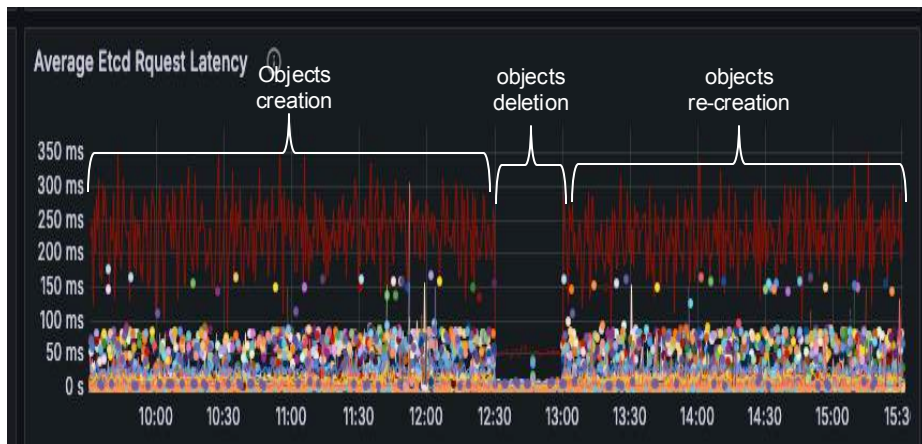
**1-1 Deployment**  
(1 binding Policy per cluster)



**1-Many Deployment**  
(1 binding Policy for all clusters)

# Experiment 3: Number of WECs - Results

## Resource Utilization: ITS API Server



- A single pane of glass to analyze performance is extremely helpful
- Define performance variable & stay focused on them
- Long tests are a must – some issues can't be detected without it
- Profiling is your friend
- Leverage existing open-source tools when possible
- Performance analysis should be used to help building usage guidelines/benchmark

# So, Does My k8s Application Need CPR?



**No!** I can see and hear the **heartbeats** even for my **Multi-Cluster** management app