

# Misadventures In Large Scale Cluster Performance

Shane Corbett, AWS  
Dima Ilchenko, Fortinet

# Scaling

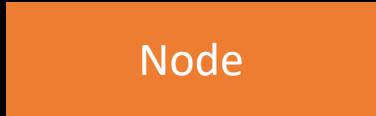
**5,000 nodes**

**150,000 total pods**

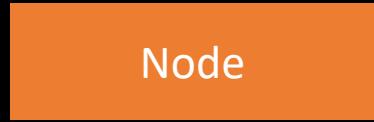
**300,000 total containers**

**110 pods per node**

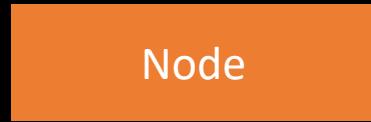
# Why Not Nodes?



Node

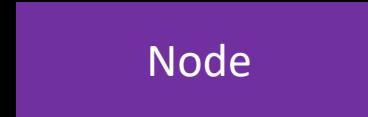


Node

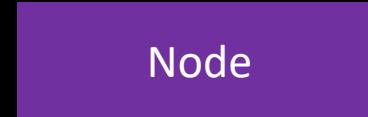


Node

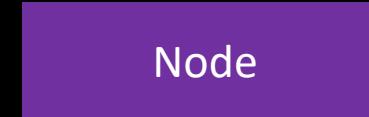
500 Nodes  
**Crash**



Node



Node



Node

6,000 Nodes  
**FINE**

# Why Not Nodes?

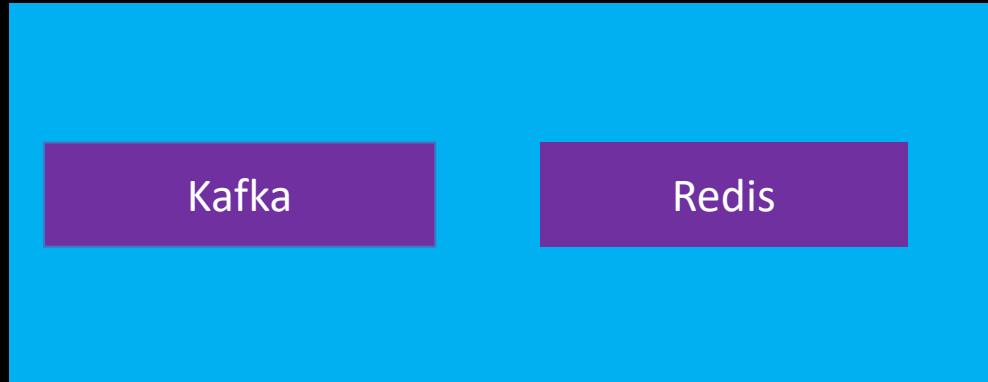
16 vCPU Node

VS.

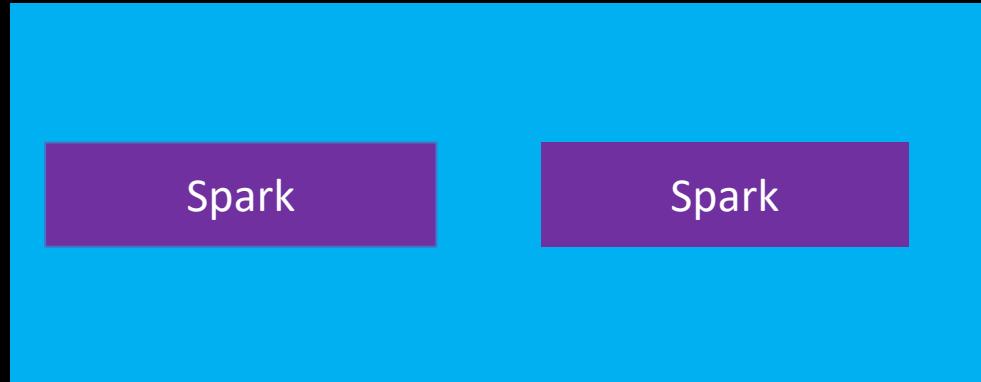
96 vCPU Node

# Why Not Pods?

**10K Pods**  
**5 Day Pod lifetime**



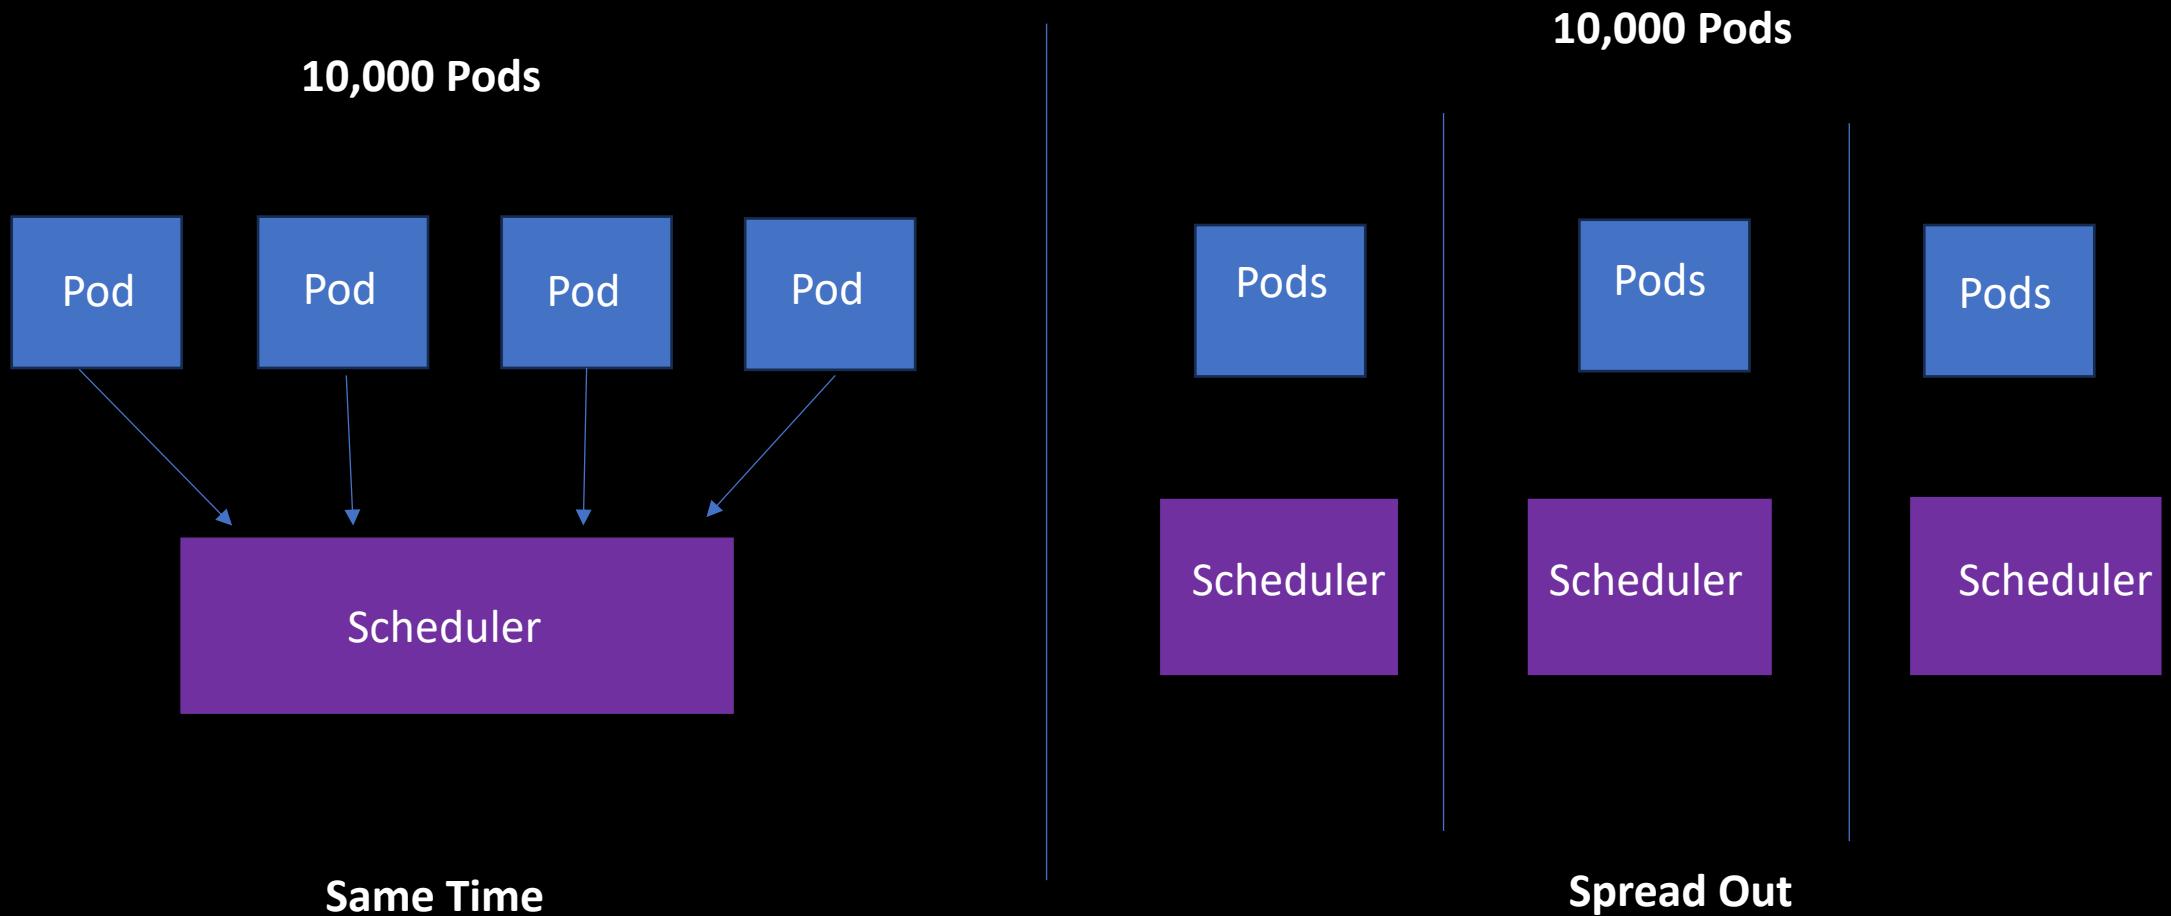
**10K Pods**  
**15 Min Pod lifetime**



# Node and Pods



# Slow Down



# Volume of Change



# API Server

CONTAINERS

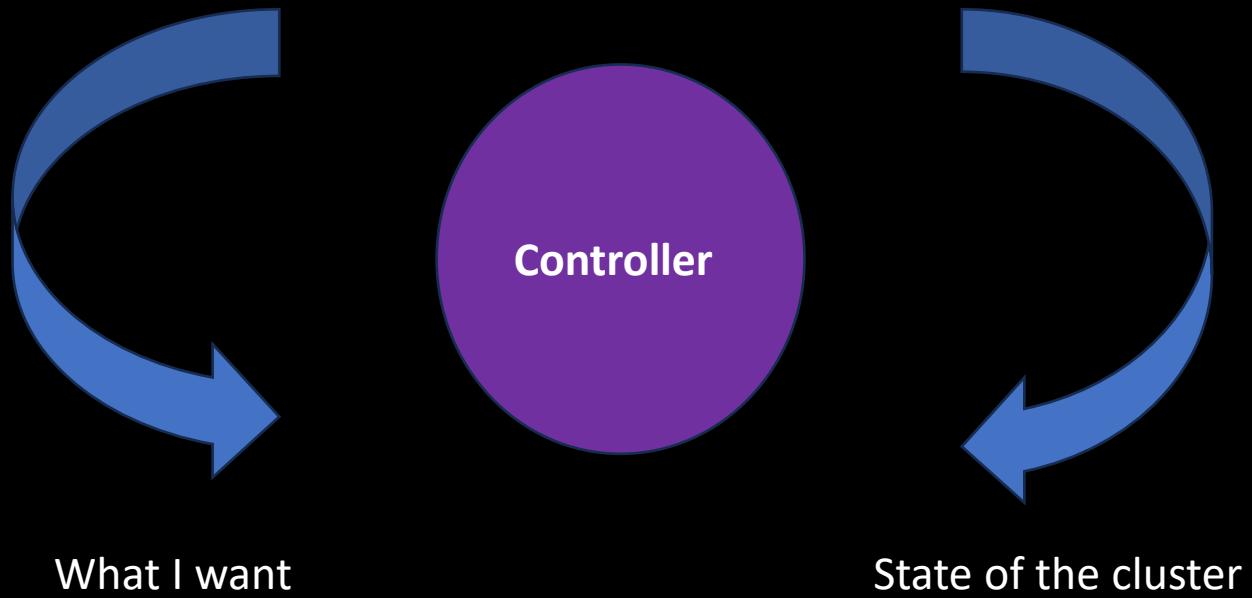
## Troubleshooting Amazon EKS API Servers with Prometheus

[Read the blog post ›](#)

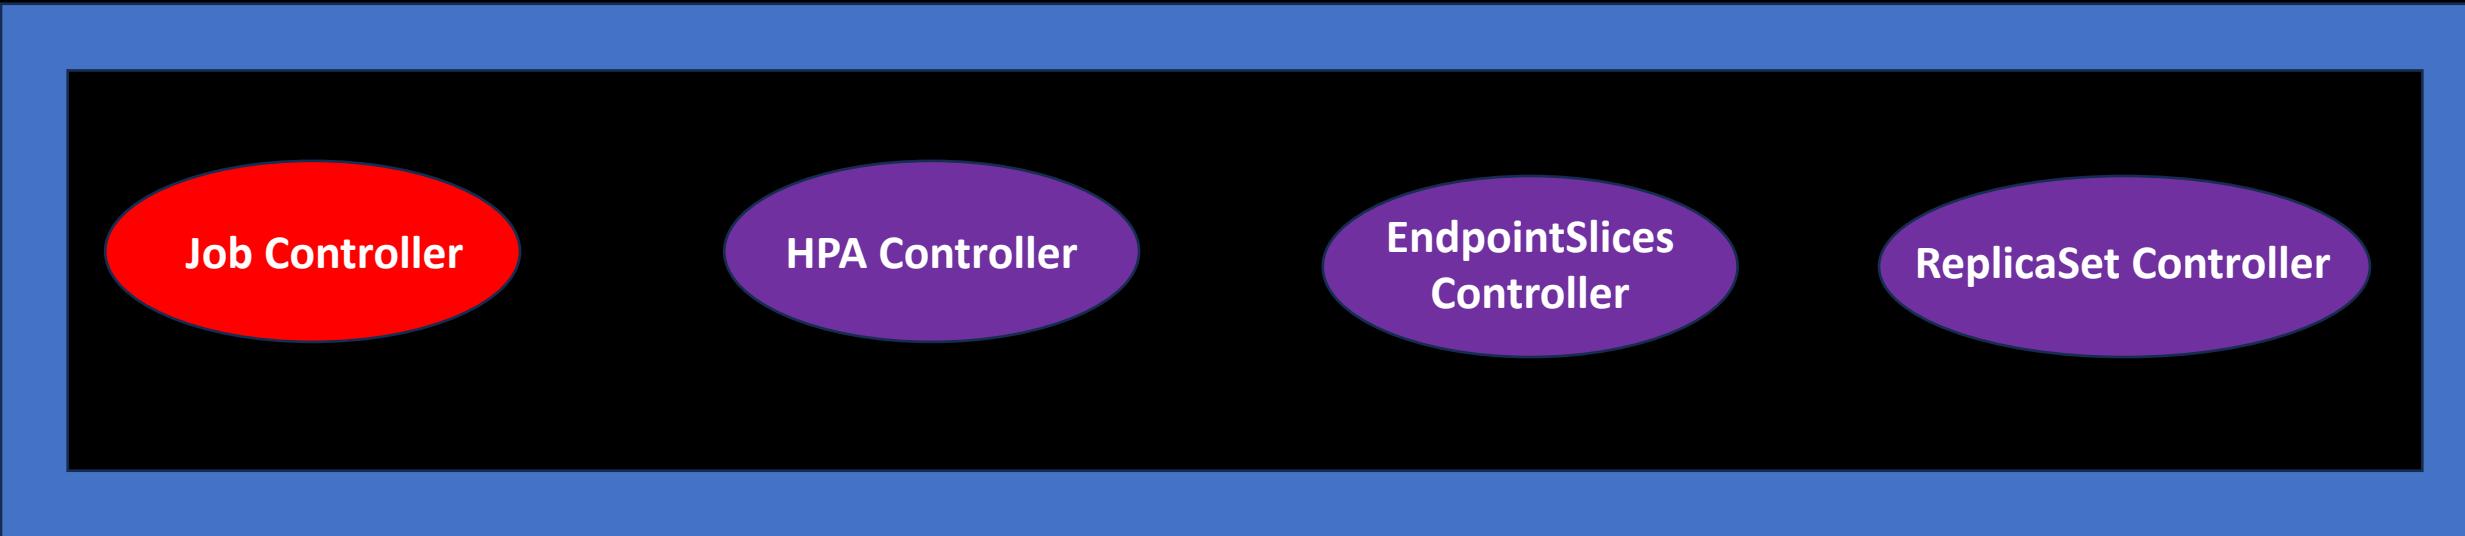
<https://aws.amazon.com/blogs/containers/troubleshooting-amazon-eks-api-servers-with-prometheus/>

# Controller

Volume of Change



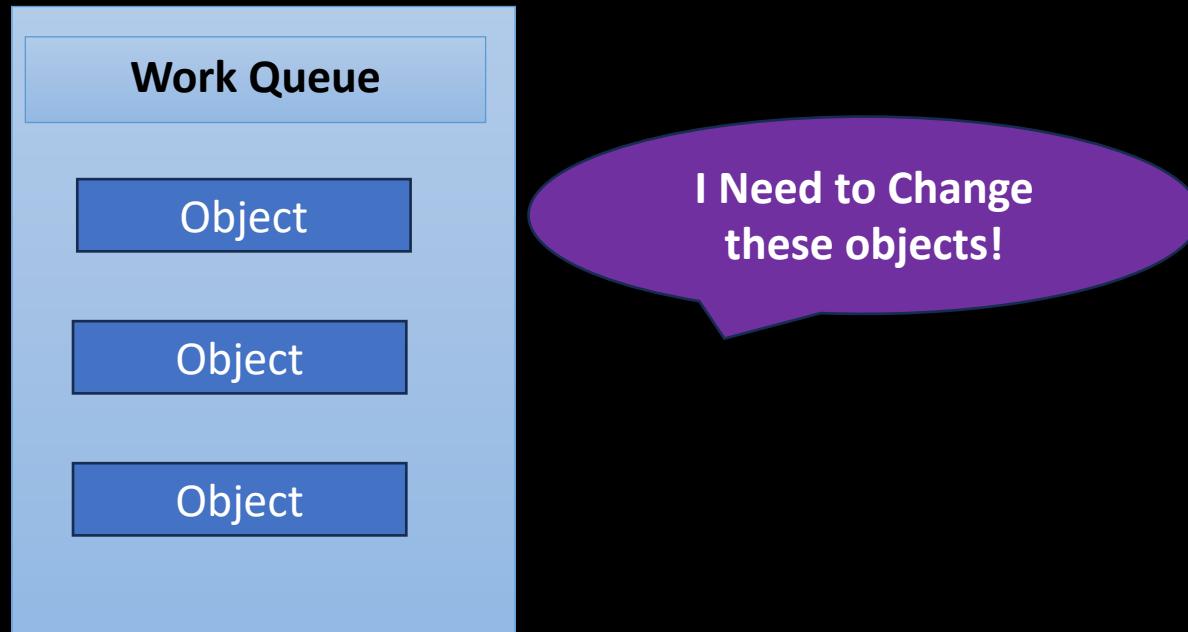
# Kube Controller Manager



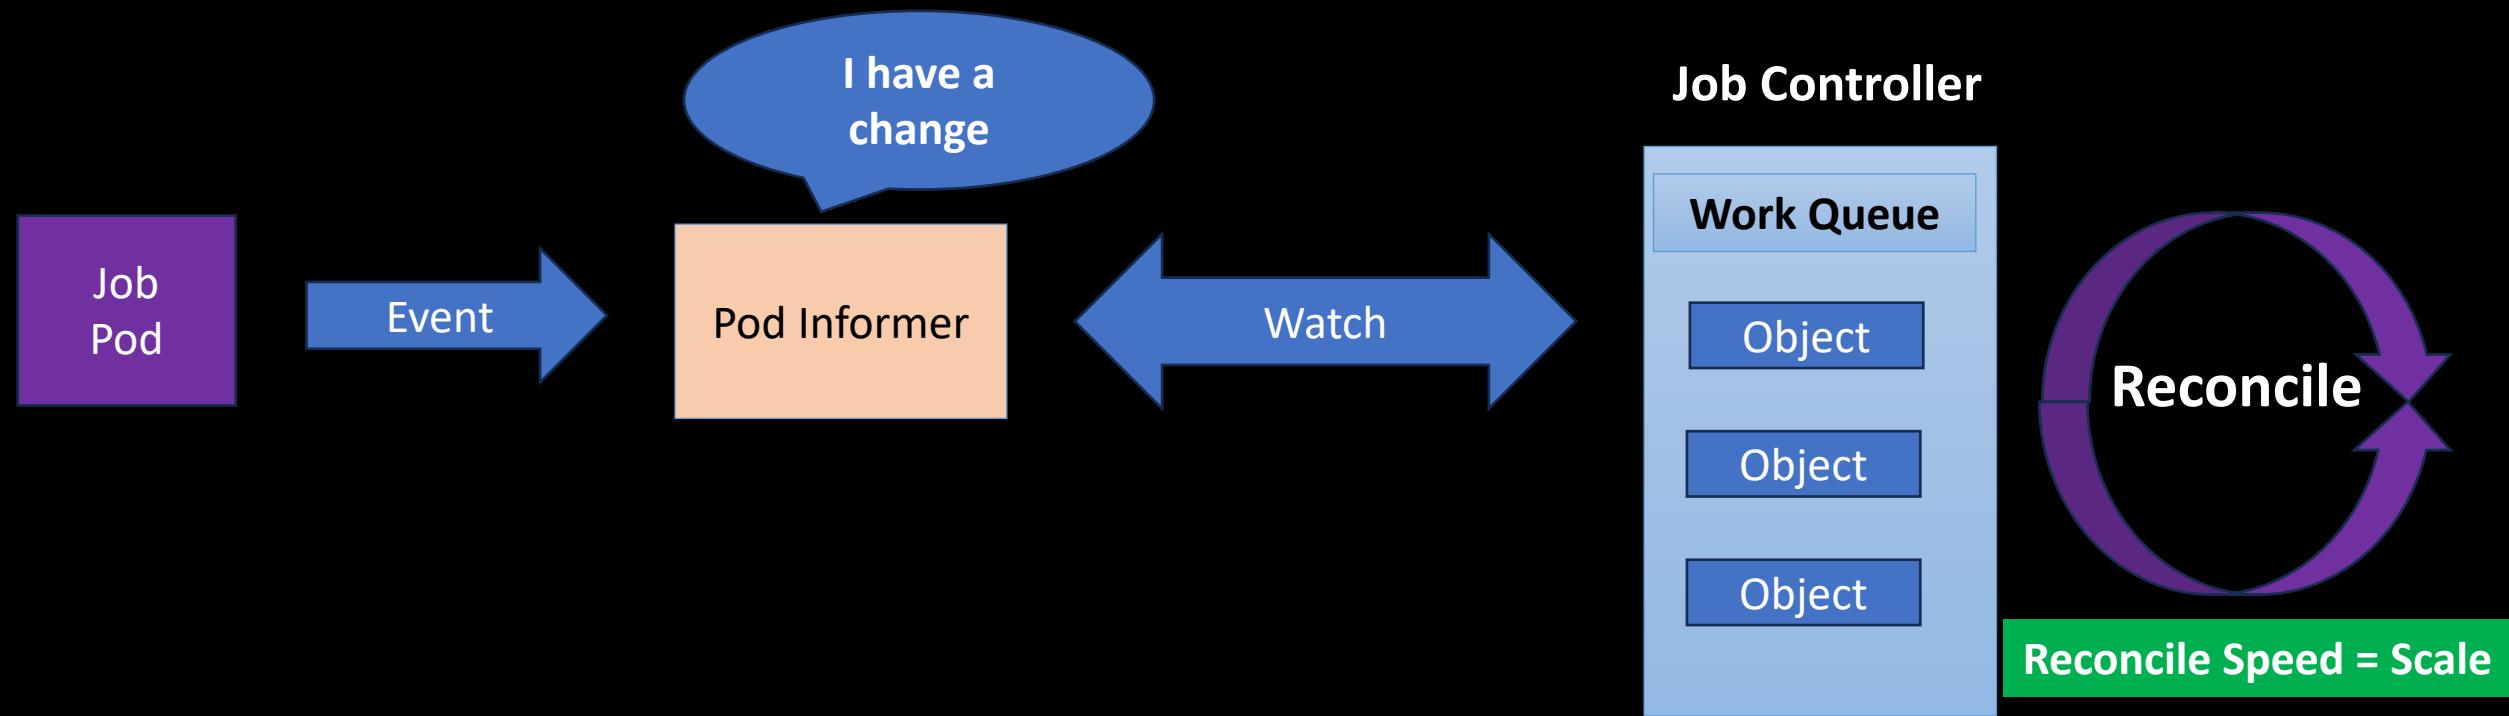
Think in Controller Volume

# Looks like a Queue

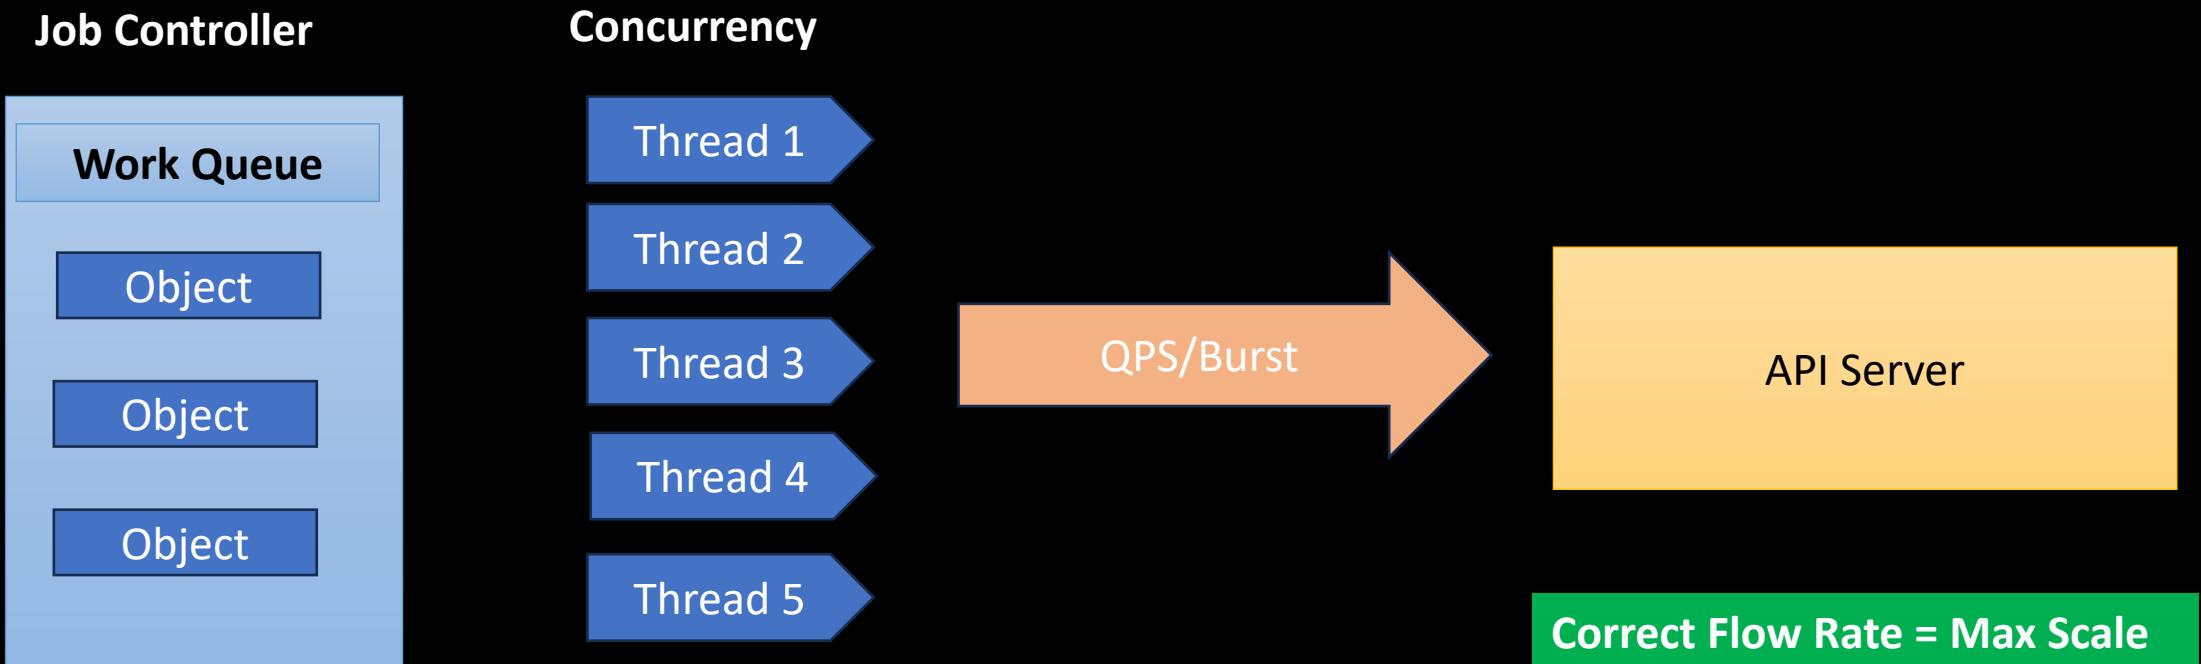
Job Controller



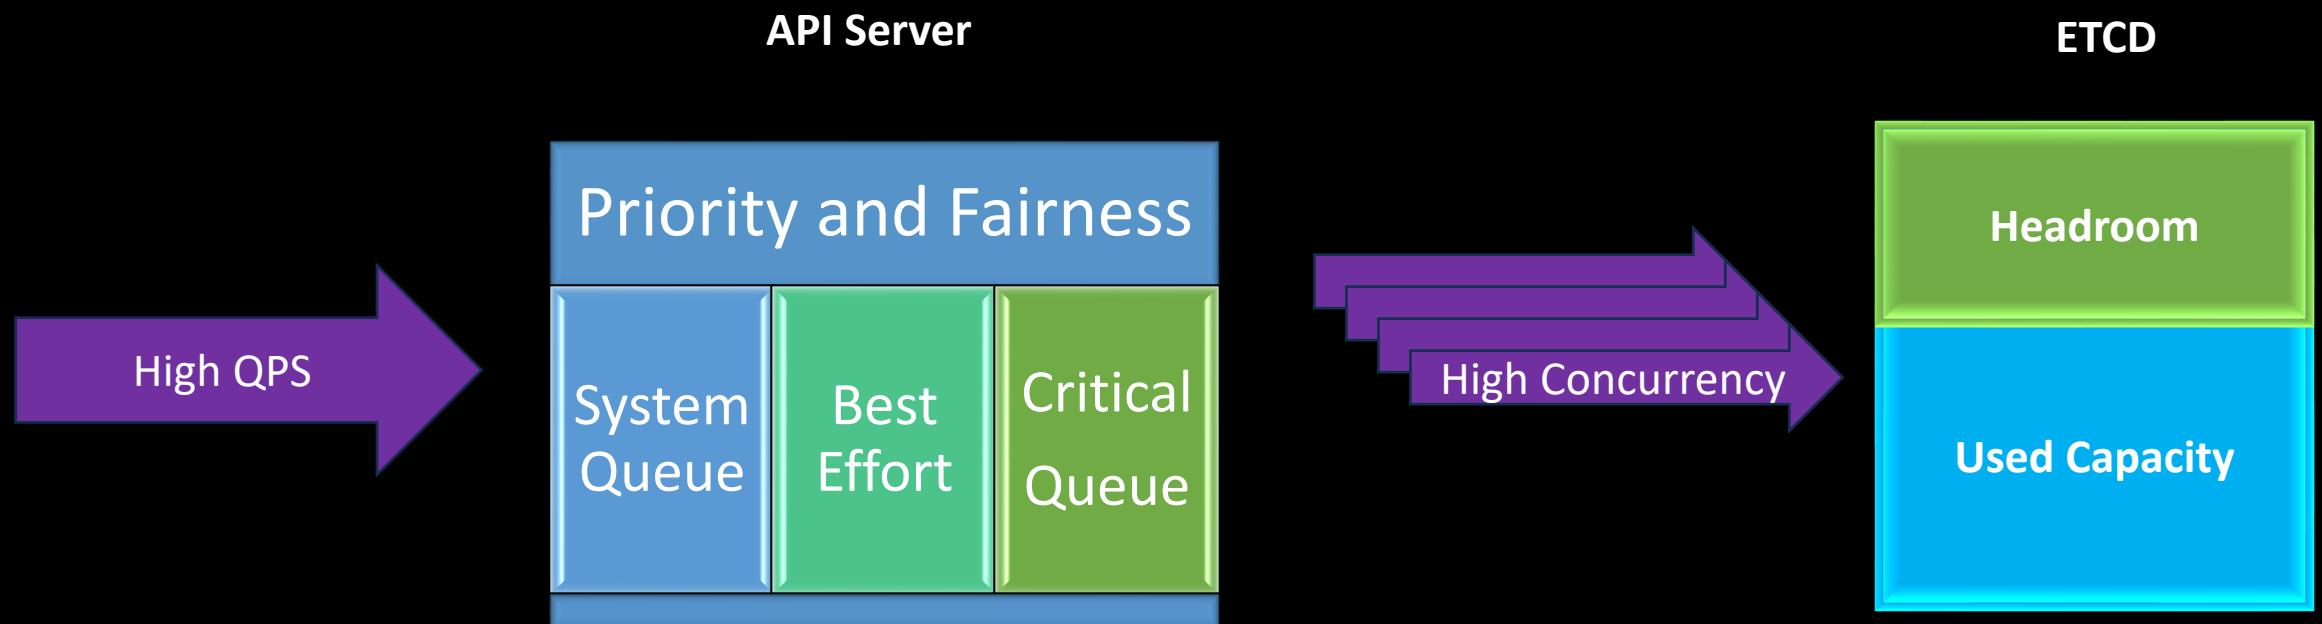
# Kube Controller Manager



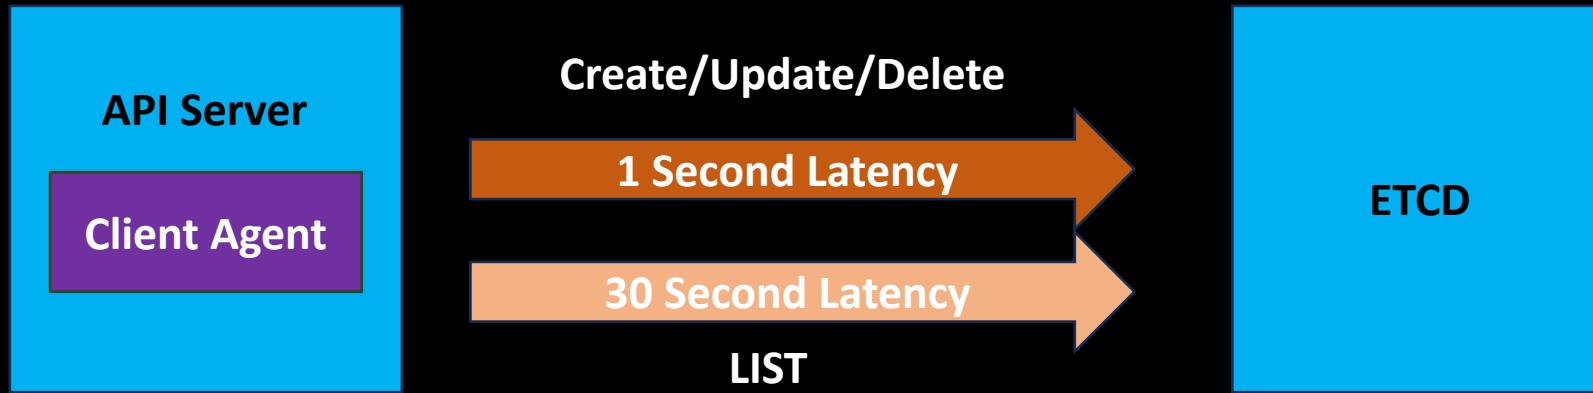
# Controllers



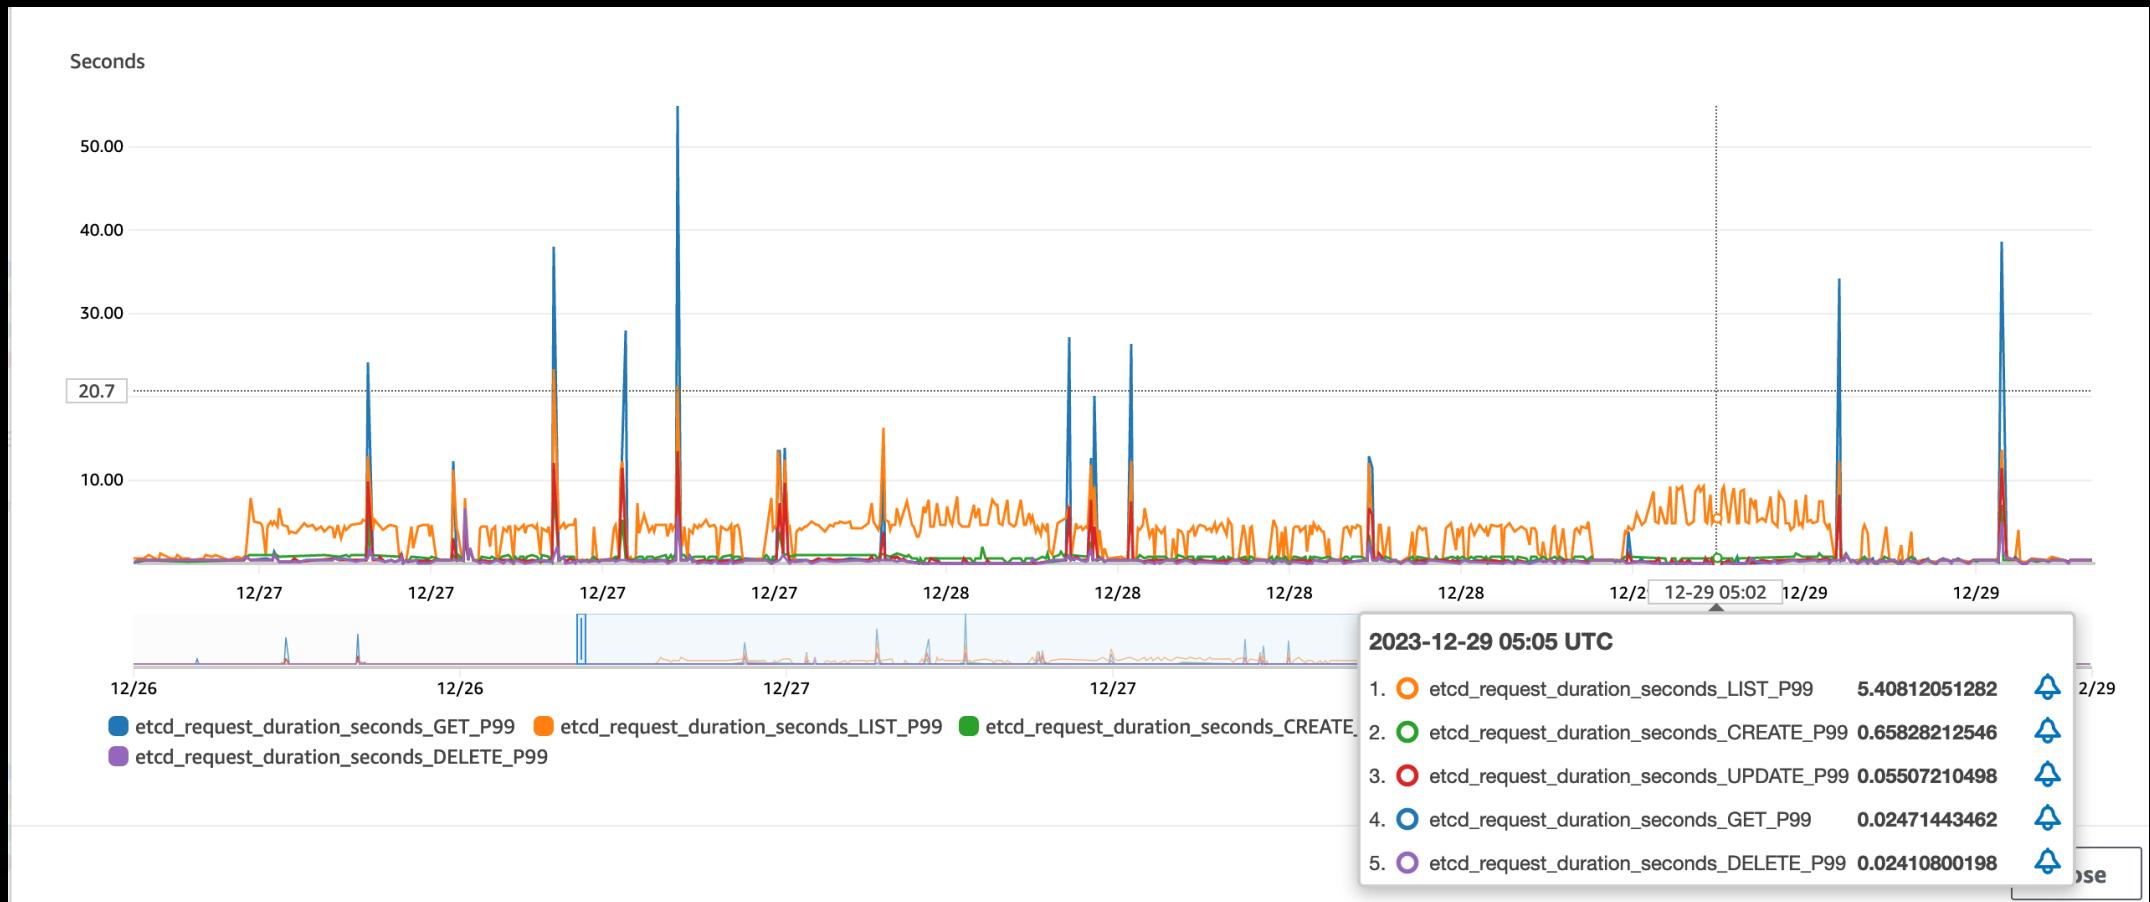
# The Right Value?



# What's The Target

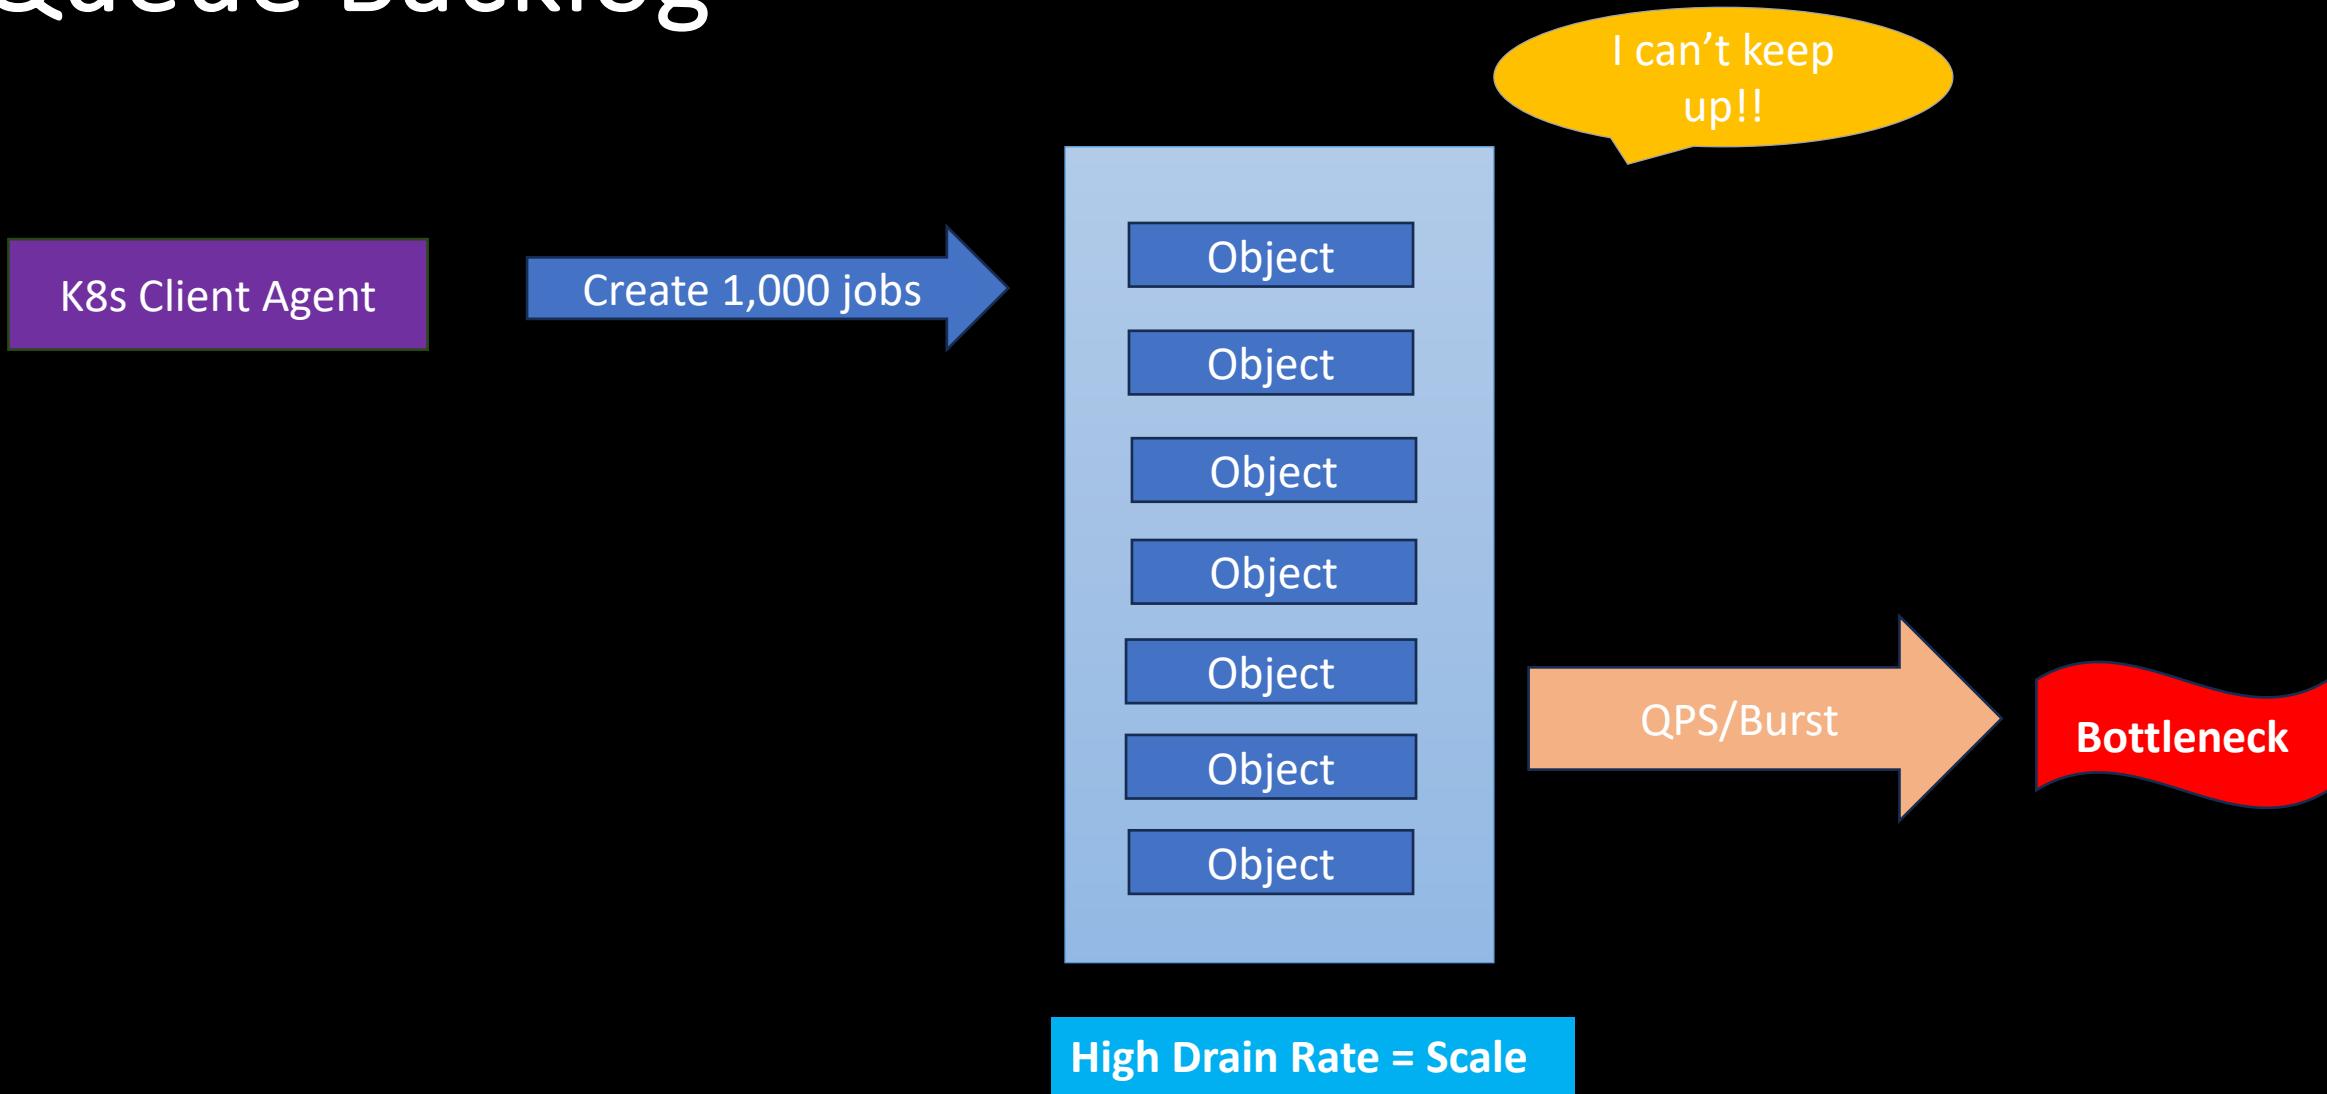


# API to ETCD Latency

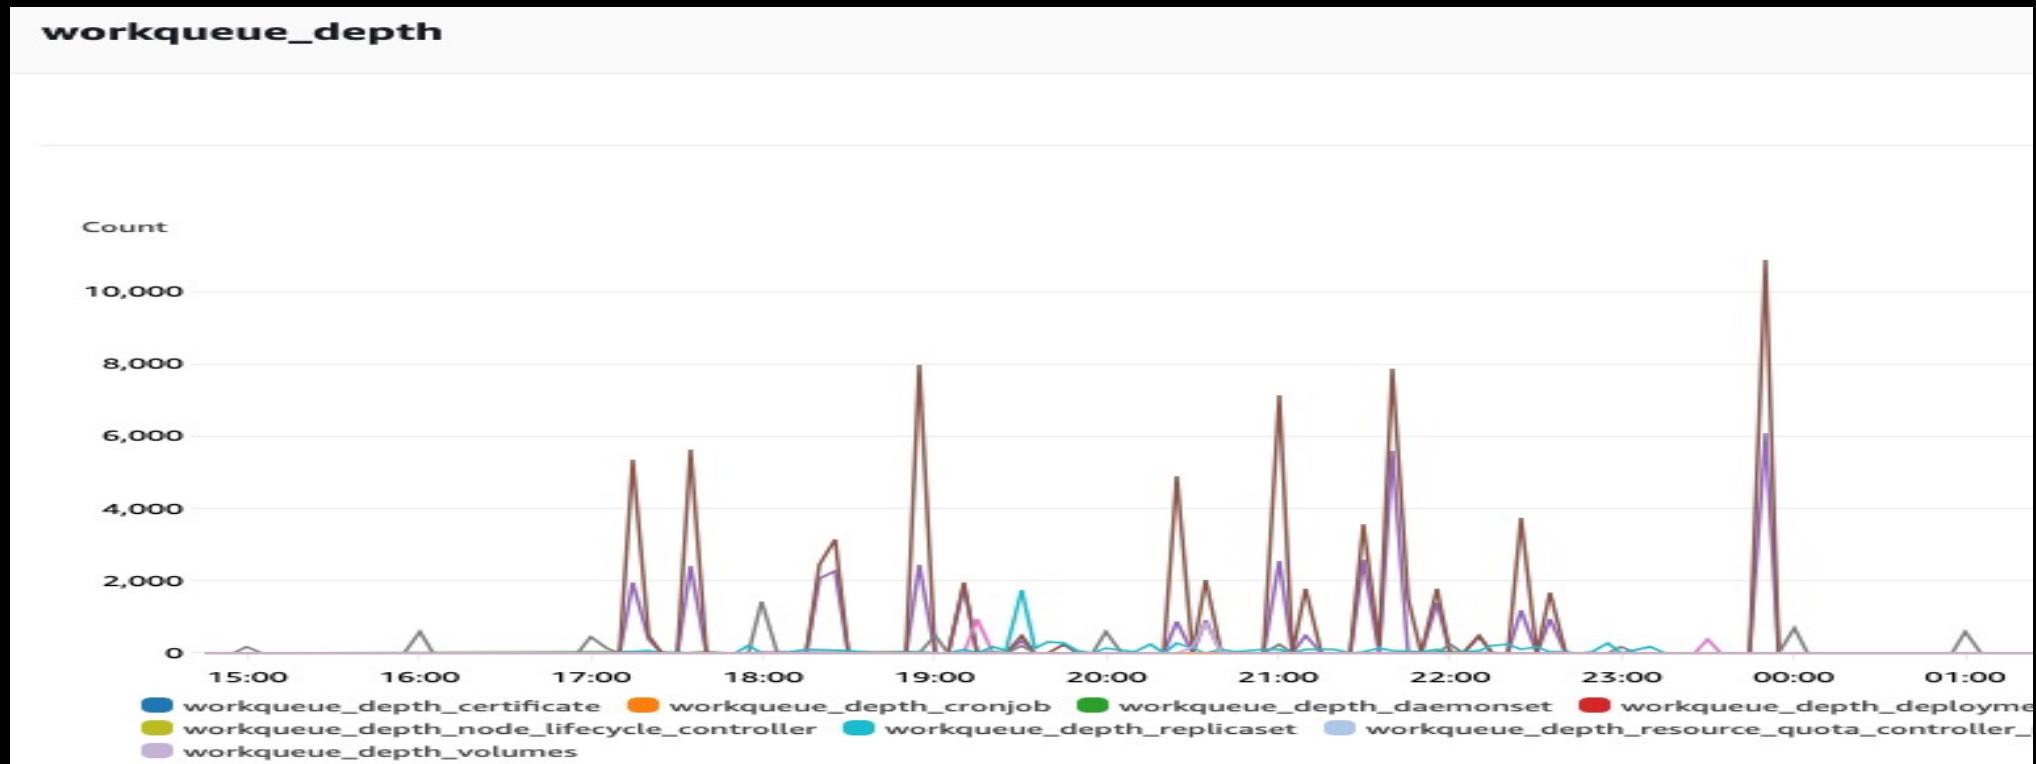


Which Controller is Sad?  
(and why)

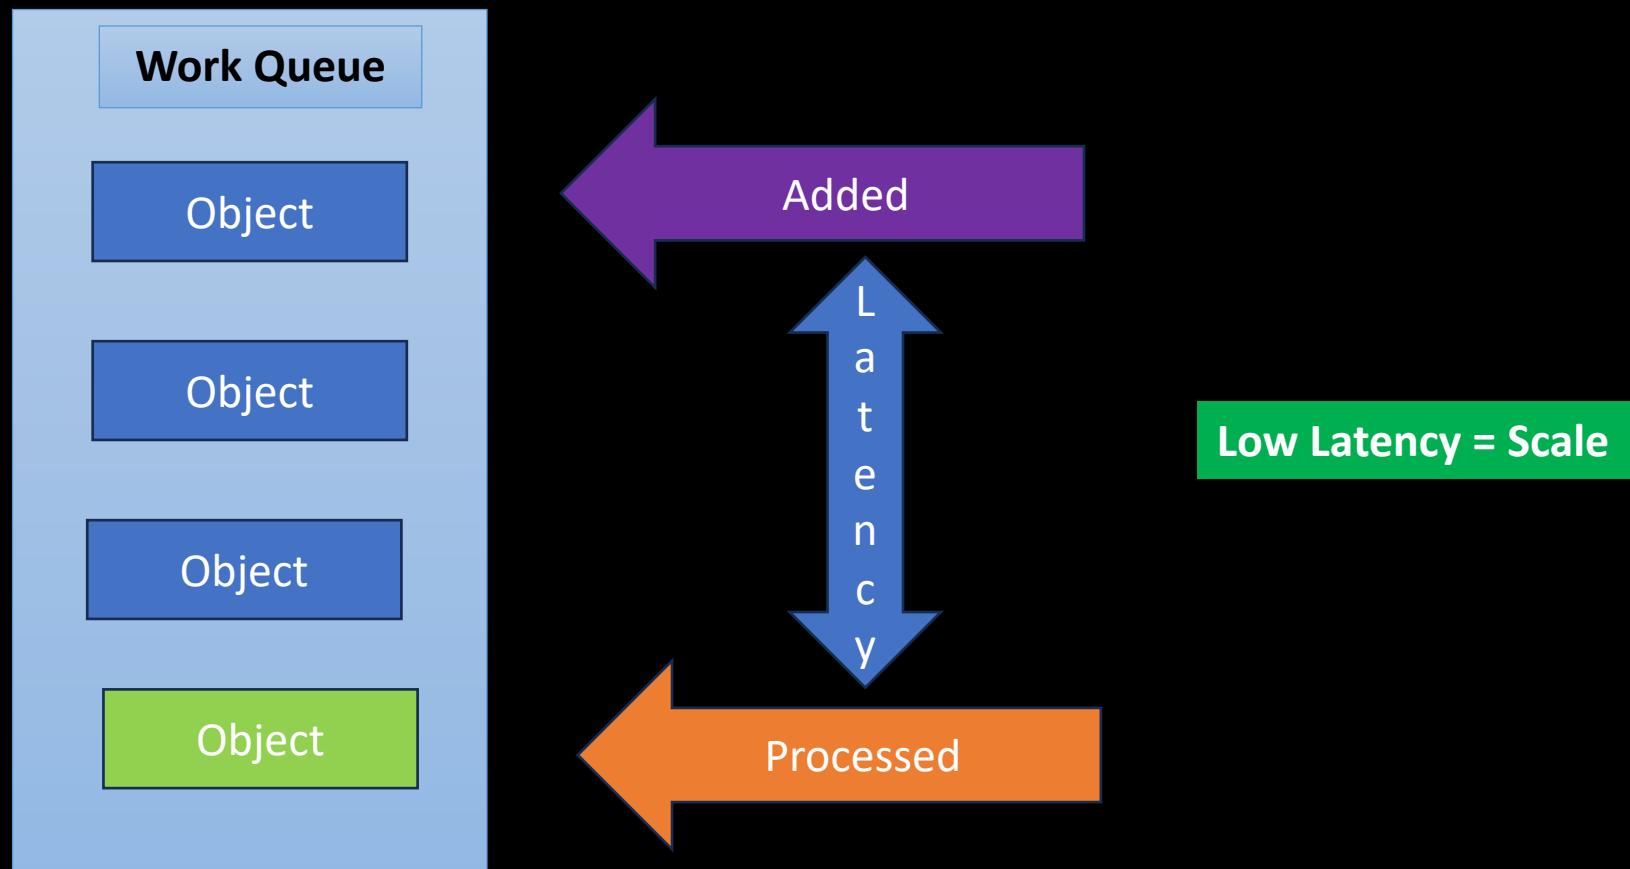
# Queue Backlog



# Work Queue Depth



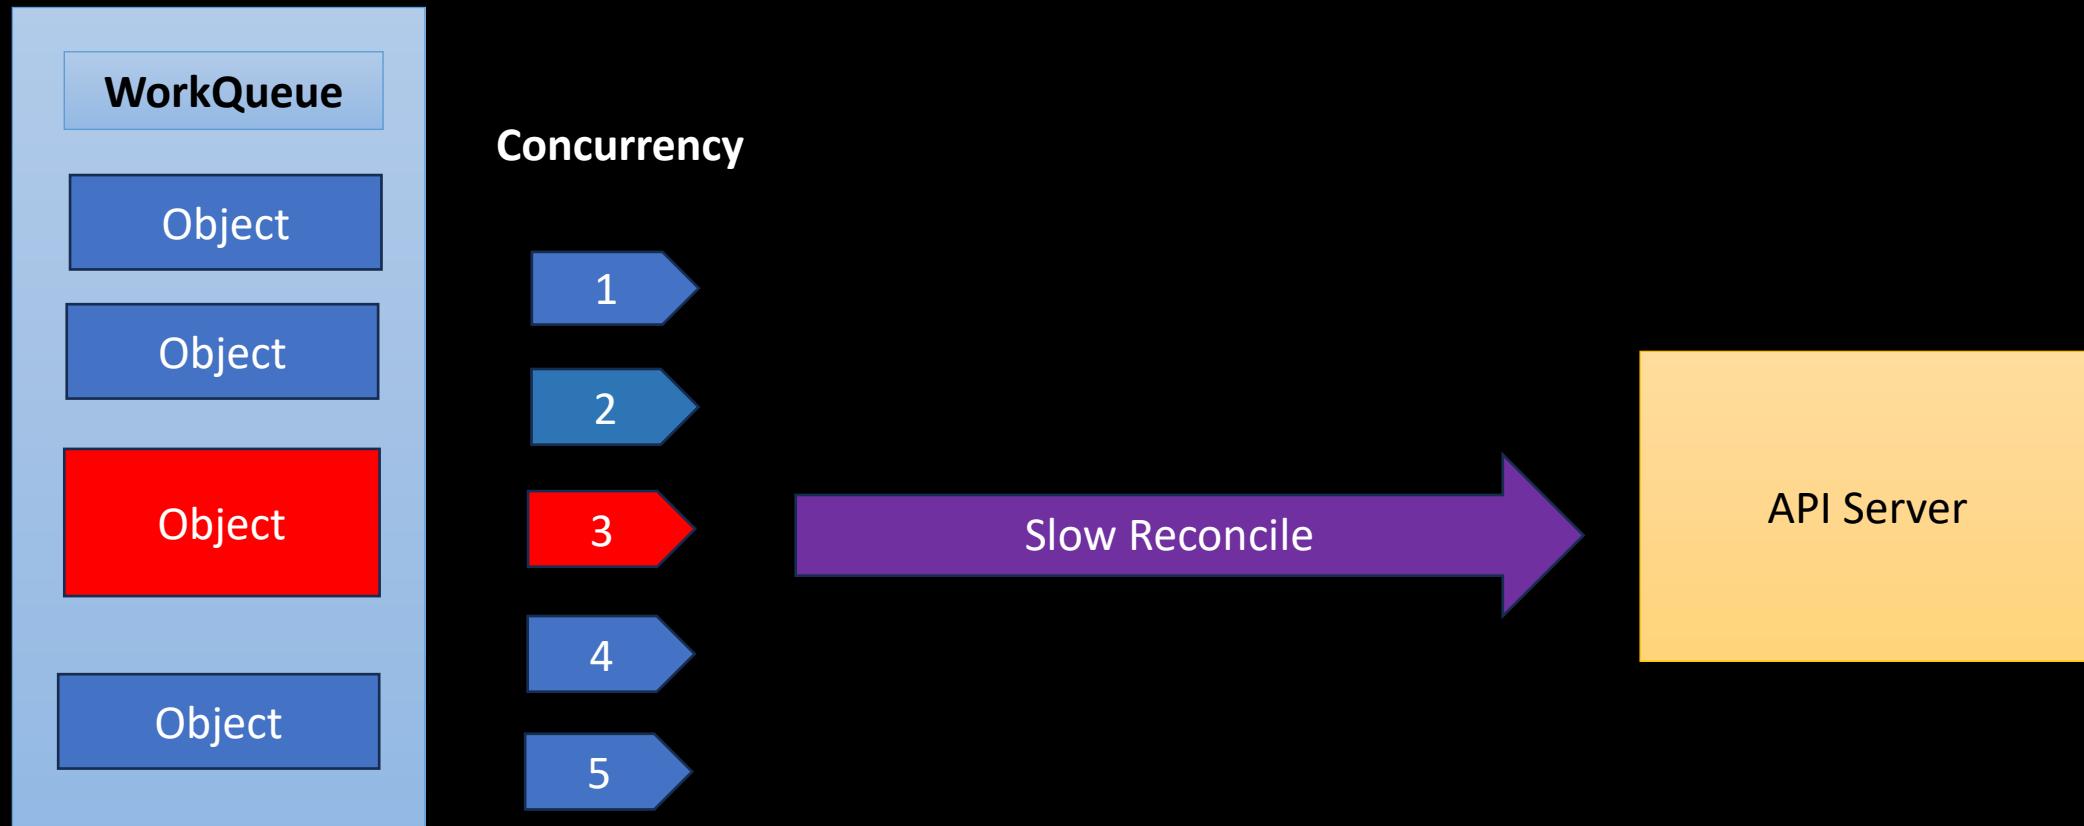
# Request Latency



# Trusting the Metrics



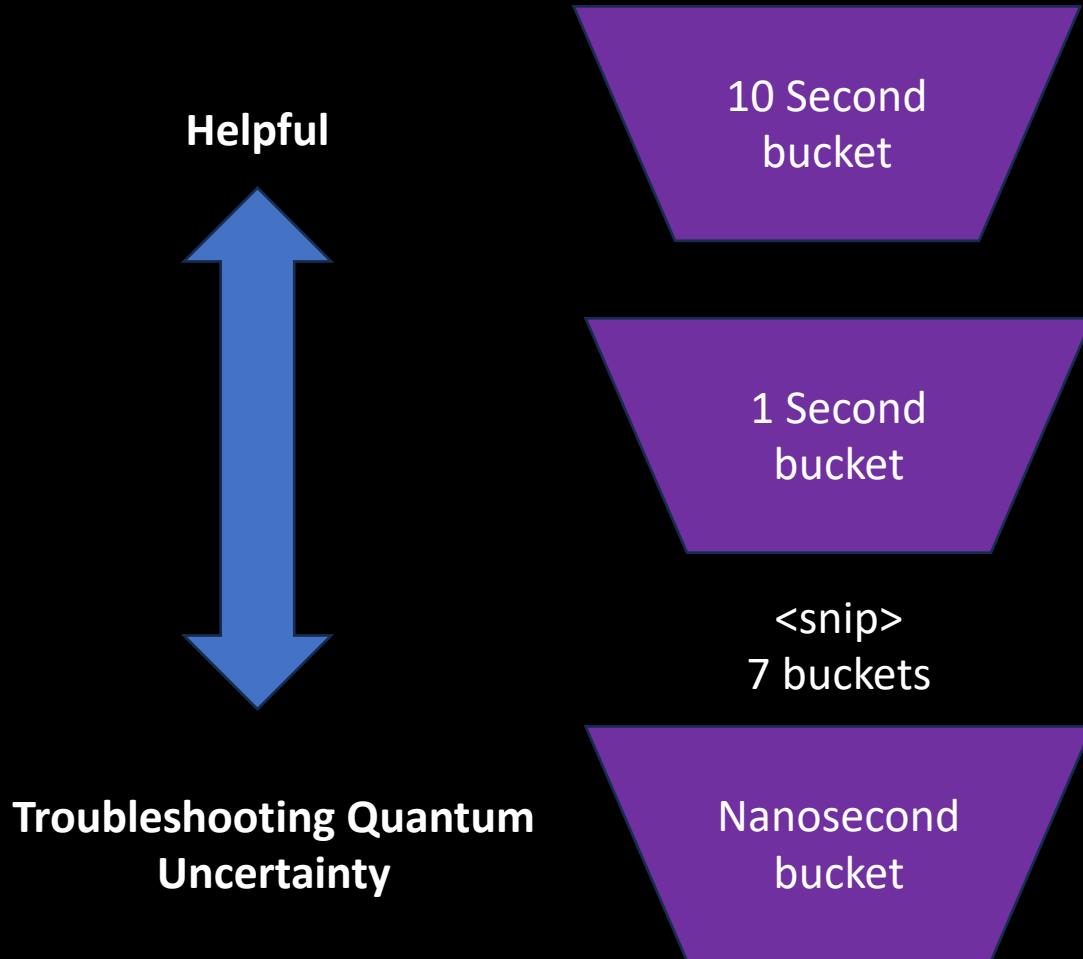
# Slow Requests



# Long Work Duration

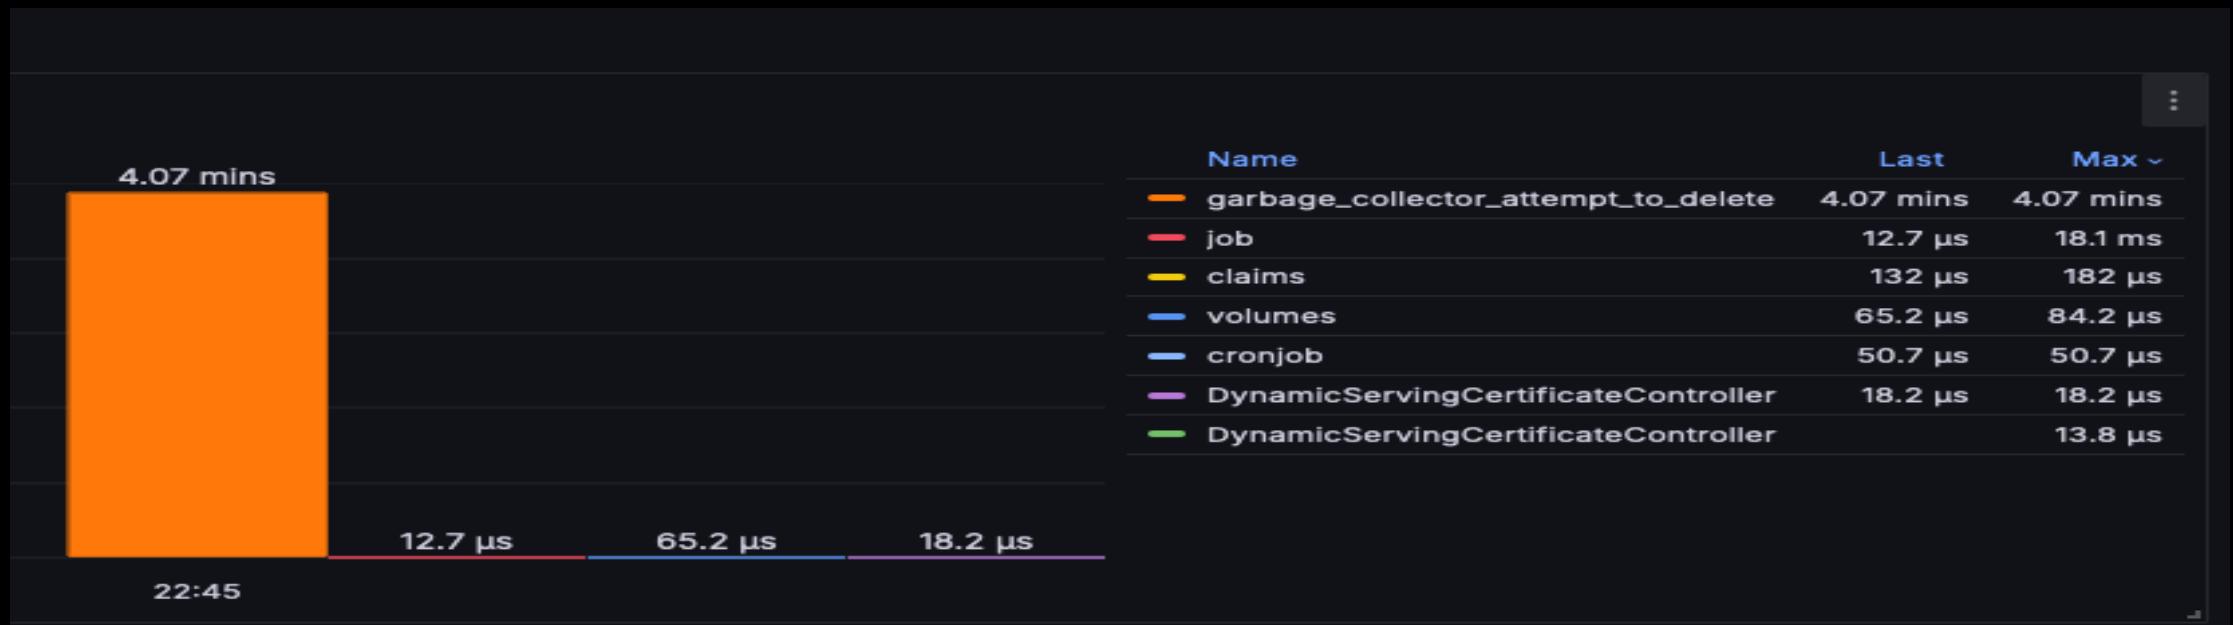


# Betrayed by Buckets



# 10 Seconds to 4 Minutes

sum/ count = AVG



# What is Slow?

#	requestURI	userAgent	objectRef.re...	objectRe...	duration_in_sec
▶ 1	/api/v1/pods	datasources-agent/v0.0.0 (linux/amd64) kubernetes/\$Format	pods		44.2675
▶ 2	/api/v1/pods	datasources-agent/v0.0.0 (linux/amd64) kubernetes/\$Format	pods		42.8898
▶ 3	/api/v1/pods	datasources-agent/v0.0.0 (linux/amd64) kubernetes/\$Format	pods		35.7841
▶ 4	/api/v1/pods	datasources-agent/v0.0.0 (linux/amd64) kubernetes/\$Format	pods		35.6
▶ 5	/api/v1/pods	datasources-agent/v0.0.0 (linux/amd64) kubernetes/\$Format	pods		35.062

## Monitoring guide

[https://aws.github.io/aws-eks-best-practices/scalability/docs/kcp\\_monitoring/](https://aws.github.io/aws-eks-best-practices/scalability/docs/kcp_monitoring/)

## Queries

<https://github.com/aws-samples/specialist-tam-container-dashboards/blob/main/troubleshooting-queries/CloudWatch-Logs-Troubleshooting-Queries.md#calculate-the-average-response-time-for-each-request-uri>

# The Unexpected

@timestamp	1717628797993
annotations.apiserver.latency.k8s.io/etcd	8.51405529s
annotations.apiserver.latency.k8s.io/response-write	4.82036605s
annotations.apiserver.latency.k8s.io/serialize-response-object	34.796534043s
annotations.apiserver.latency.k8s.io/total	1m0.002500633s
annotations.apiserver.latency.k8s.io/transform-response-object	2.518μs
annotations.authorization.k8s.io/decision	allow

Response-Write – Slow Client (Lack of Network, CPU)  
Serialize-response-object – JSON Serialization (Protobuf)

# Mentors



Shyam Jeedigunta



Chao Chen



Nathan Herz

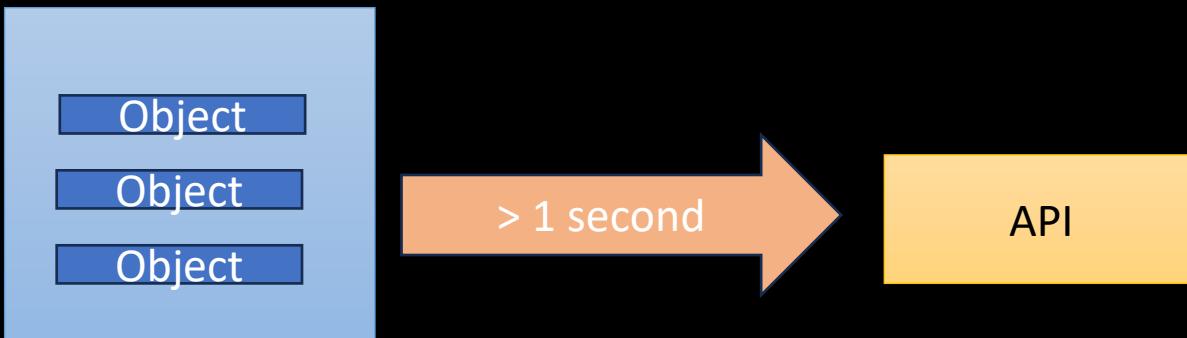


Harish Kuna

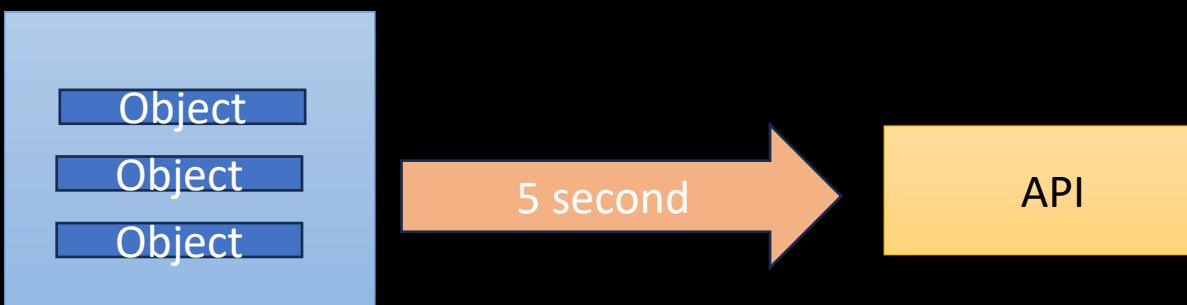
# Latency Effects

The max number of objects!

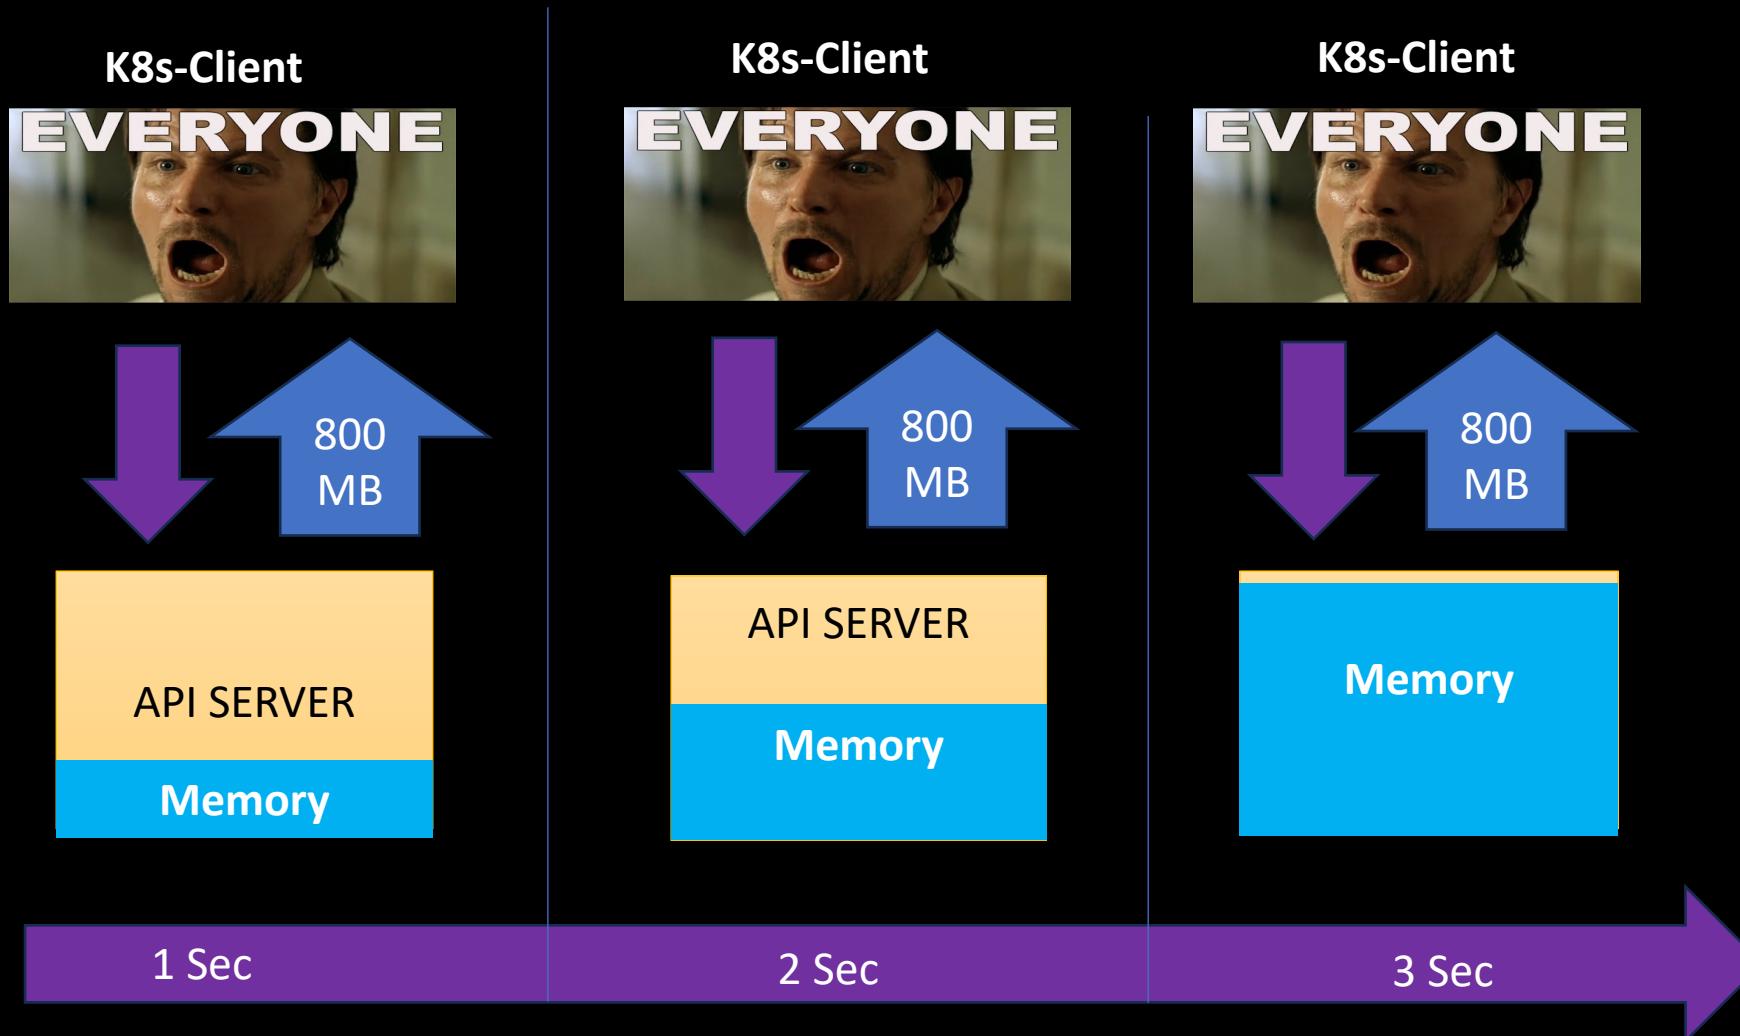
1000 Objects



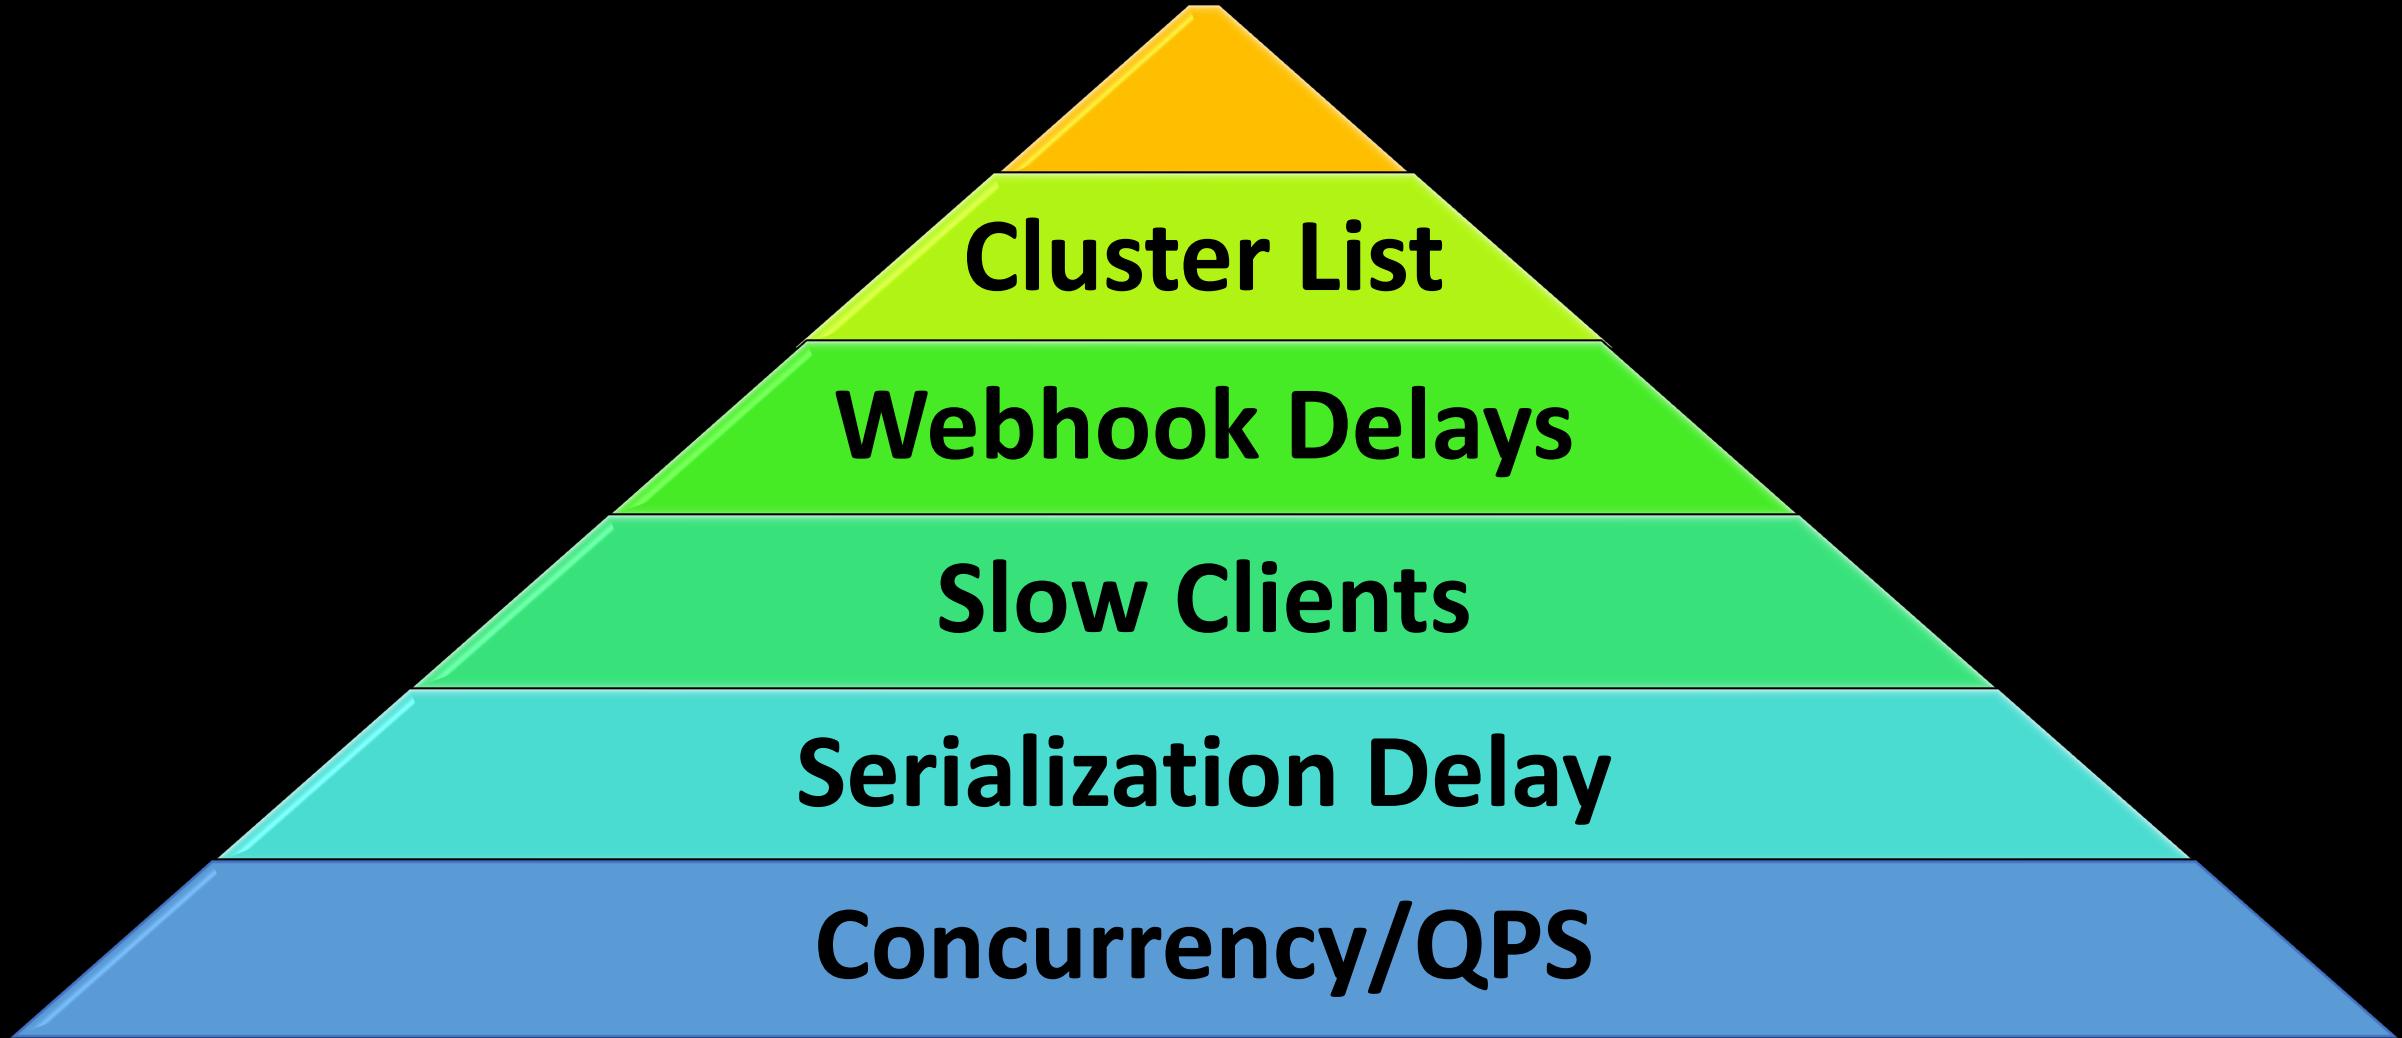
100 Objects



# Workloads Cause Latency



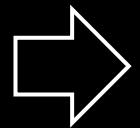
# Pyramid of Problems



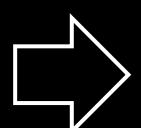
# Nº1 Cluster Reduction Trick

hosting providers hate him

Custom  
kube-scheduler



NodeResourcesFit:  
MostAllocated

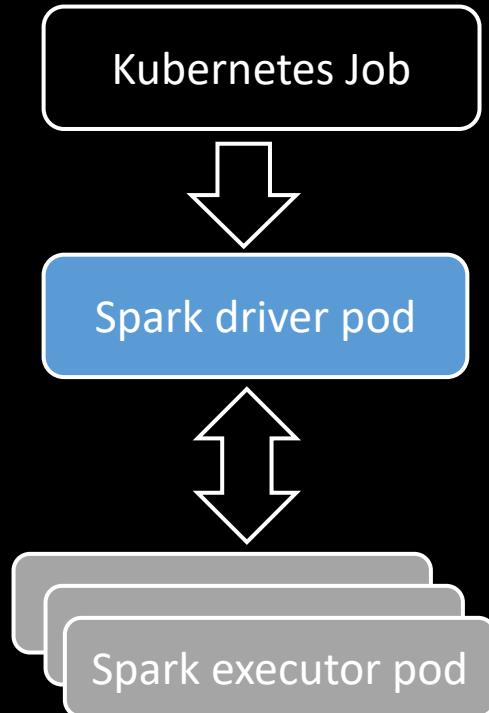


50%  
Cluster Size

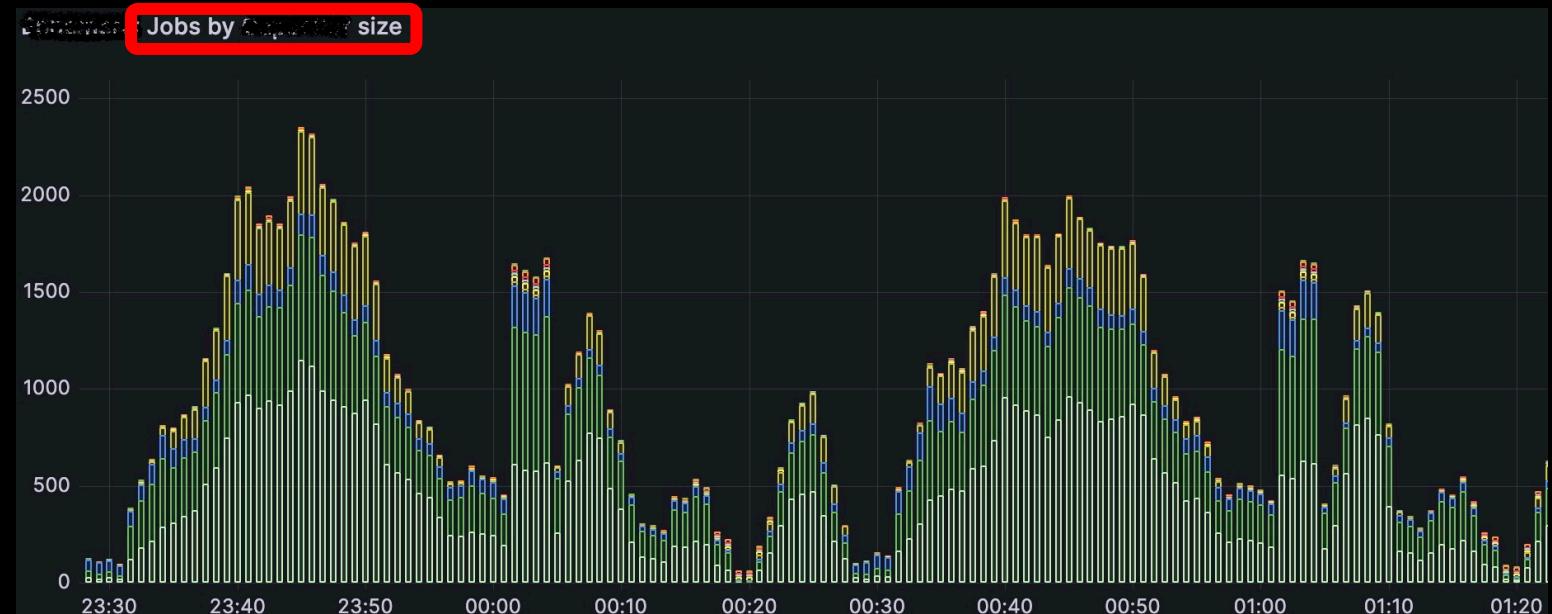


# The Mischiefous Workload

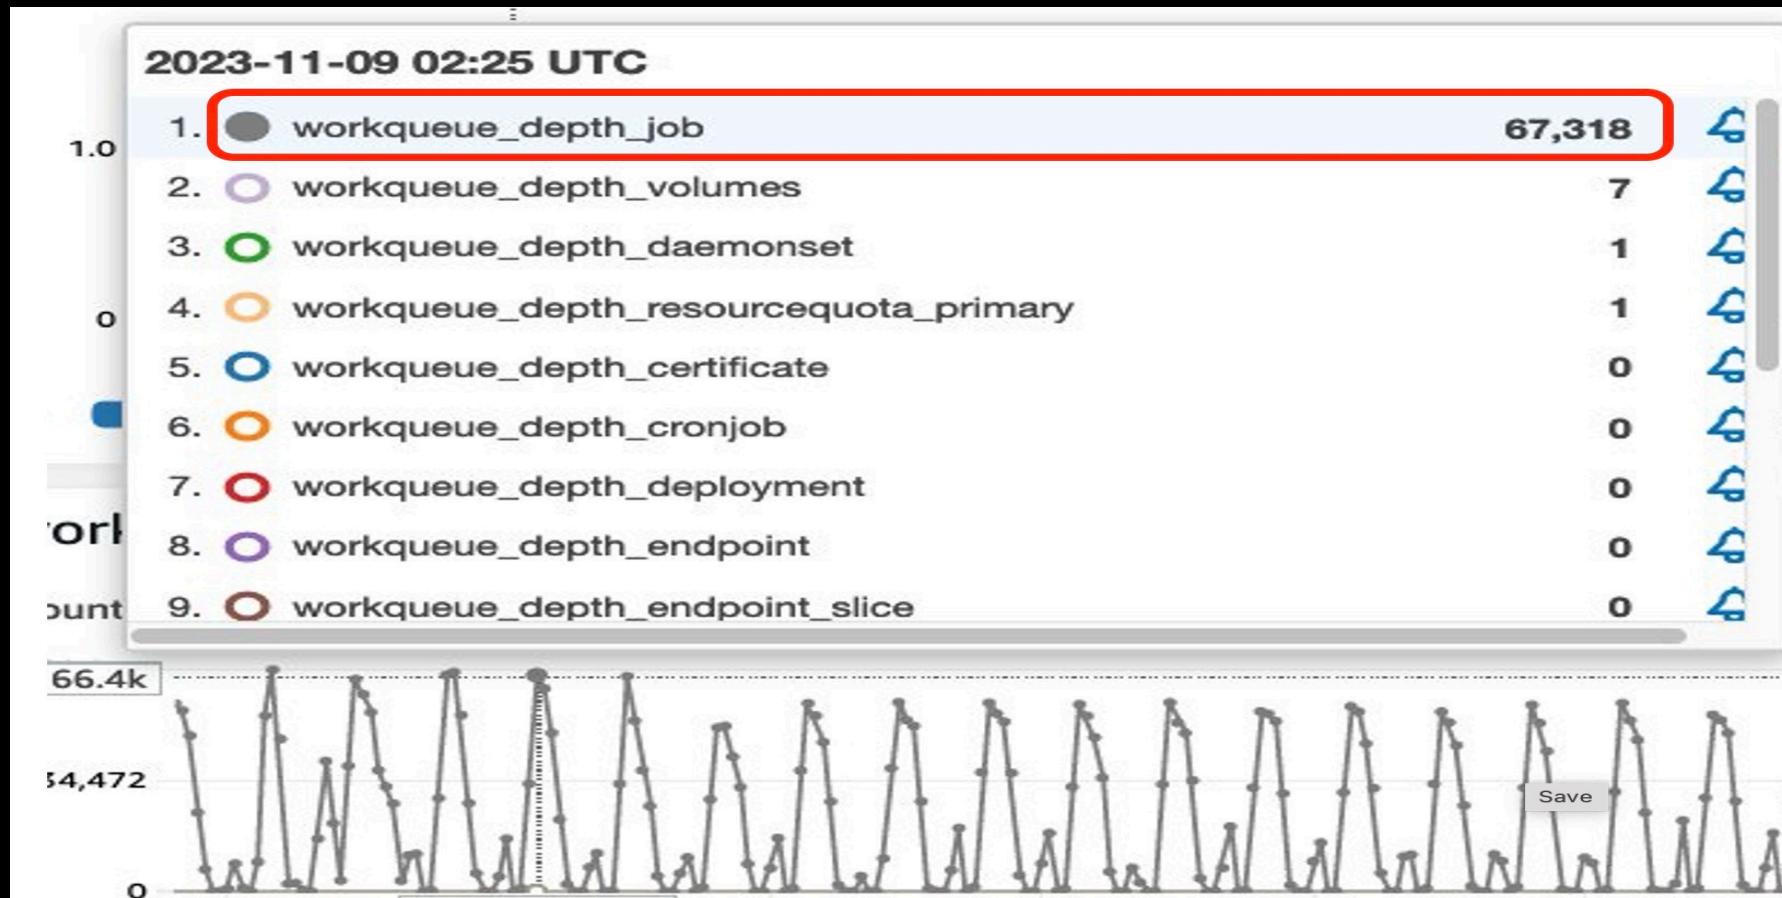
Spark Jobs anatomy



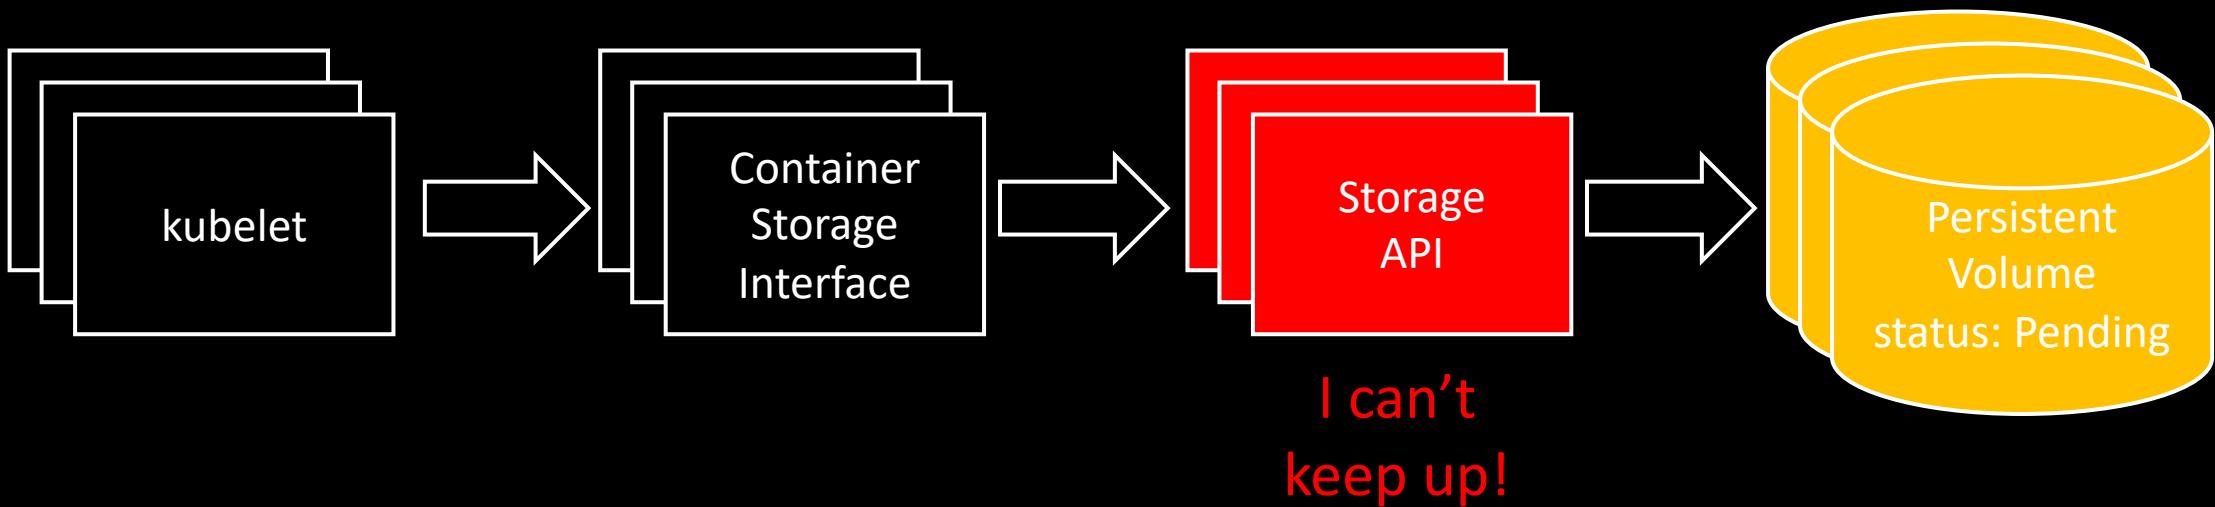
Spark Jobs scale



# Control Plane Catastrophe



# 1,000 Persistent Volumes **(PER SECOND)**



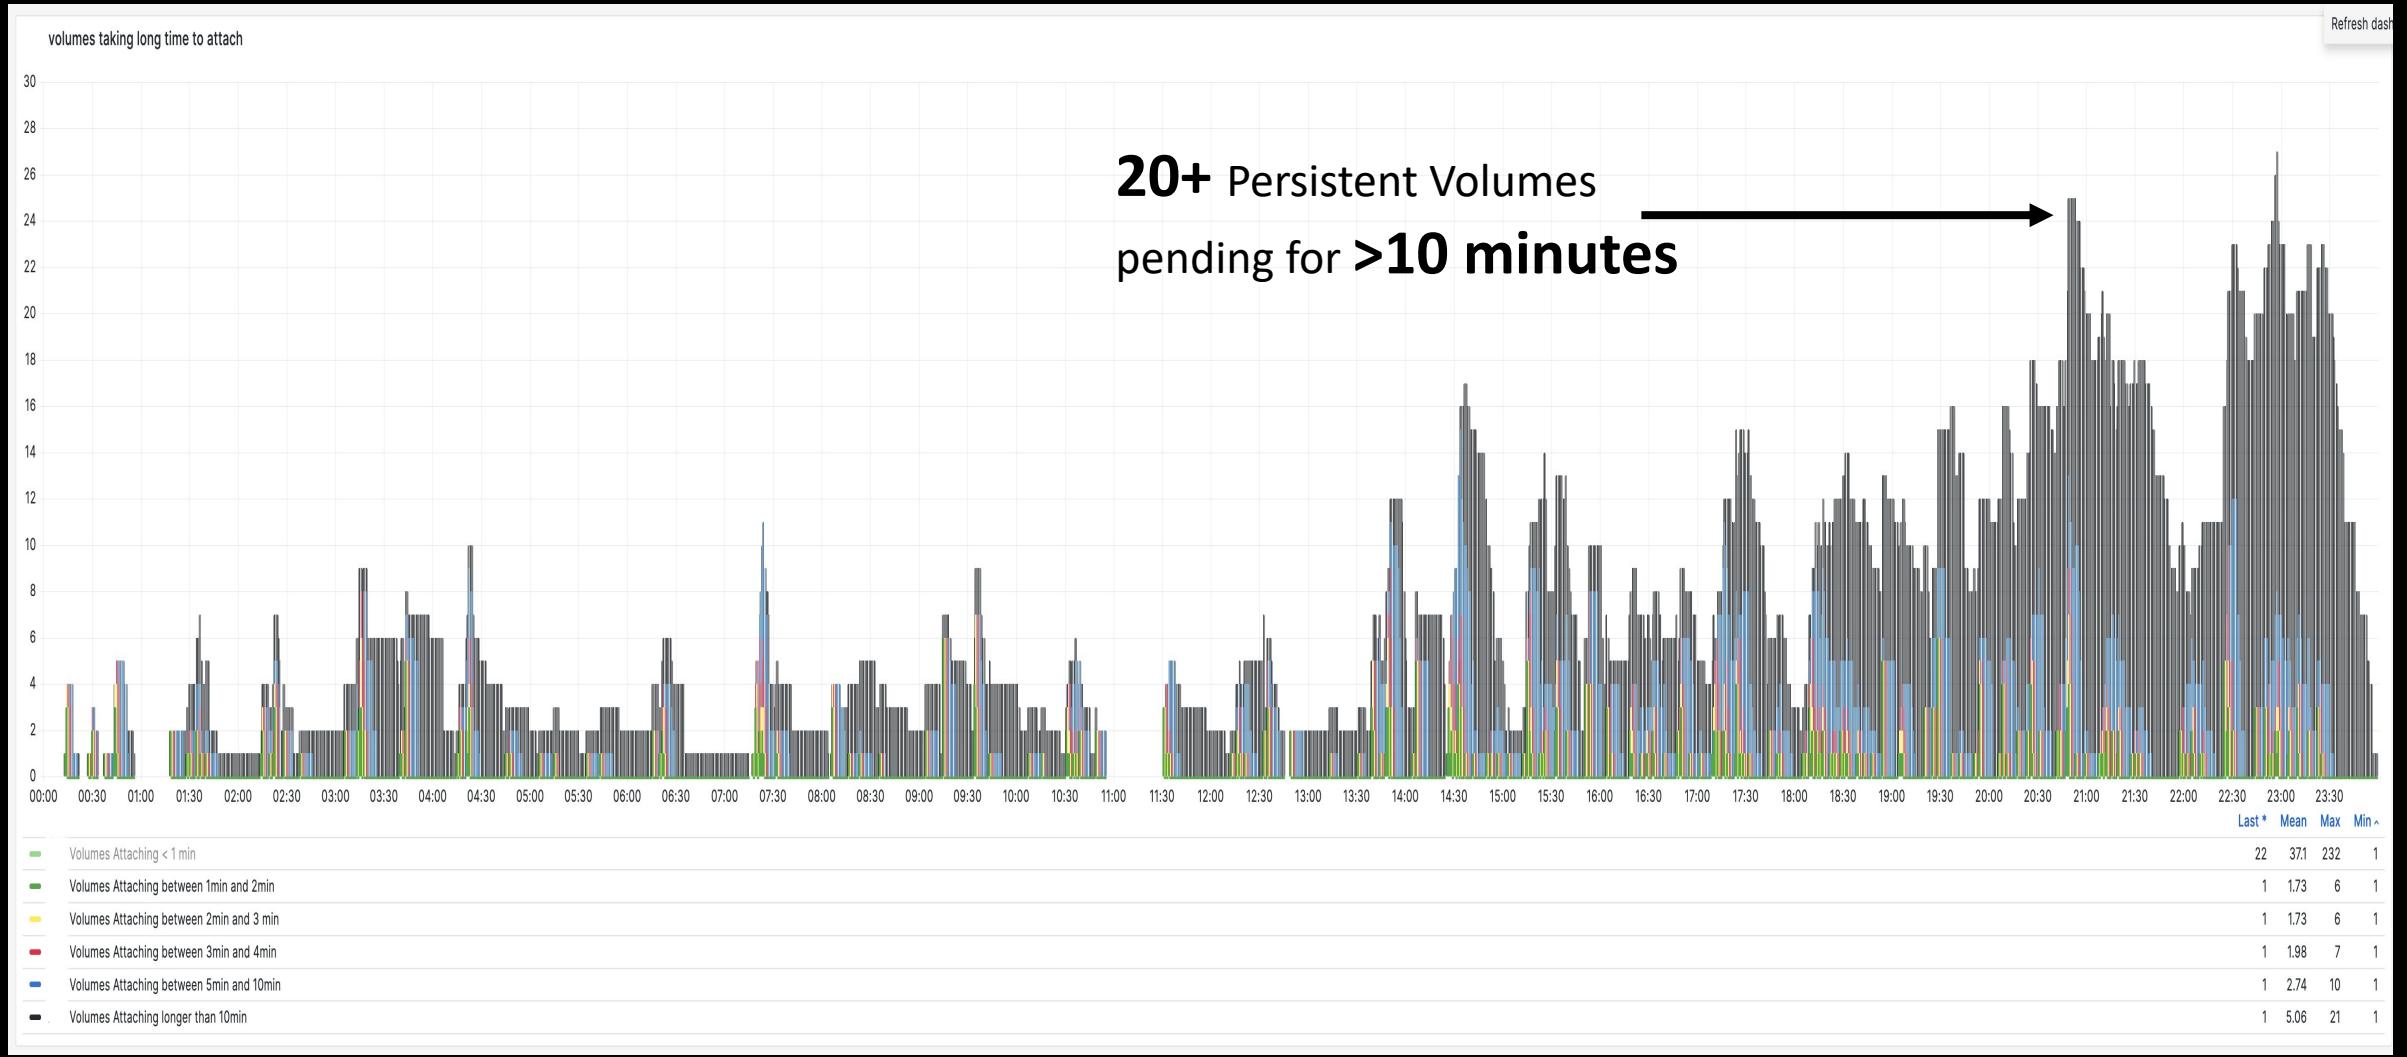
# CSI Problem: Logs

```
kcm created an event at [REDACTED] -  
AttachVolume.Attach failed for volume "pvc-567770b1-b2f9-  
4807-9707-21aab54c37b8" : timed out waiting for external-  
attacher of [REDACTED] CSI driver to attach volume  
[REDACTED]
```

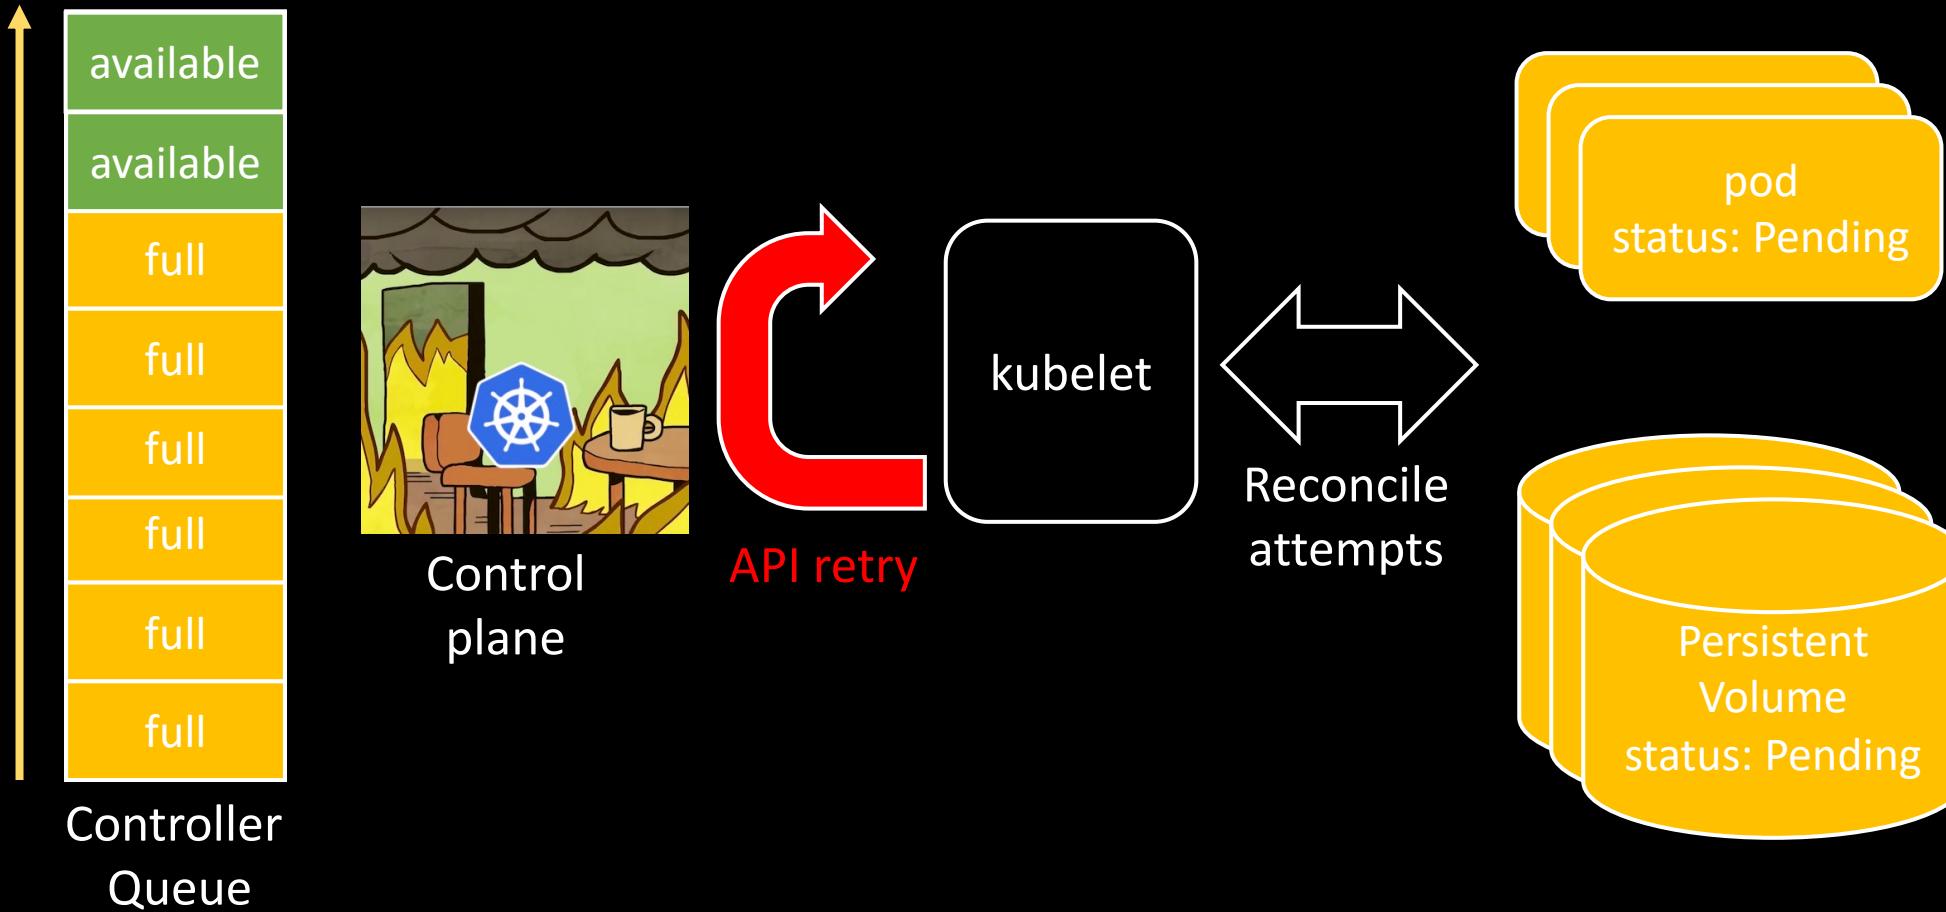
Always look at the logs!



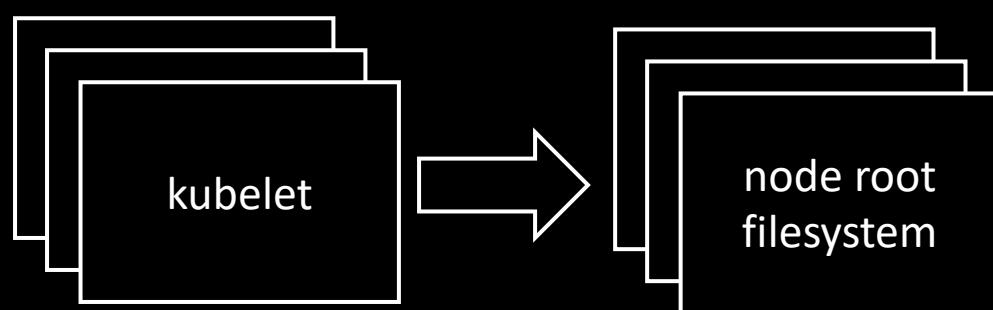
# CSI Problem: Not Very Scary Graph



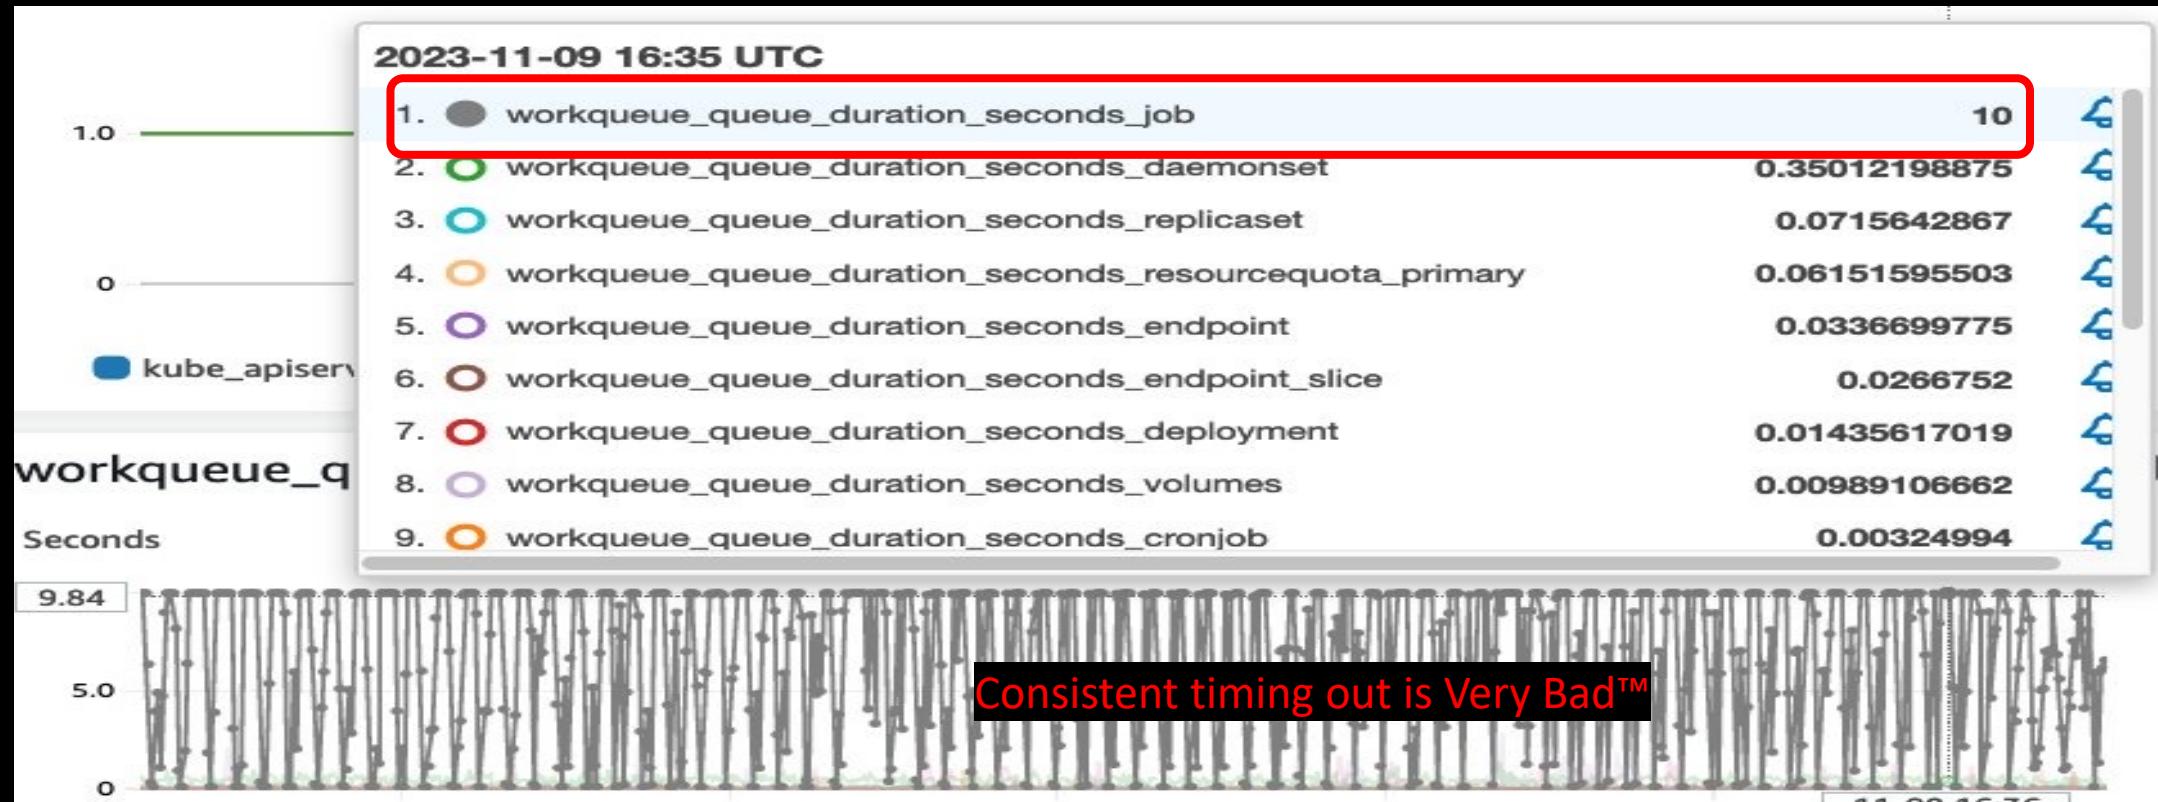
# (Not Just) CSI Problem



# No CSI – No Problem



# No CSI – No Problem Different Problem



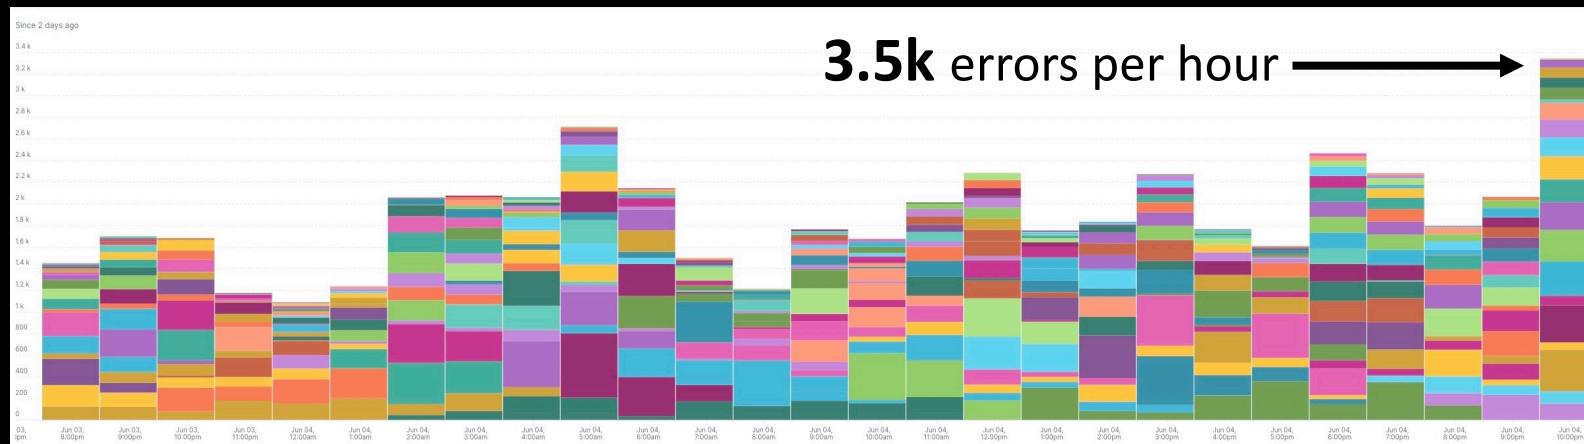
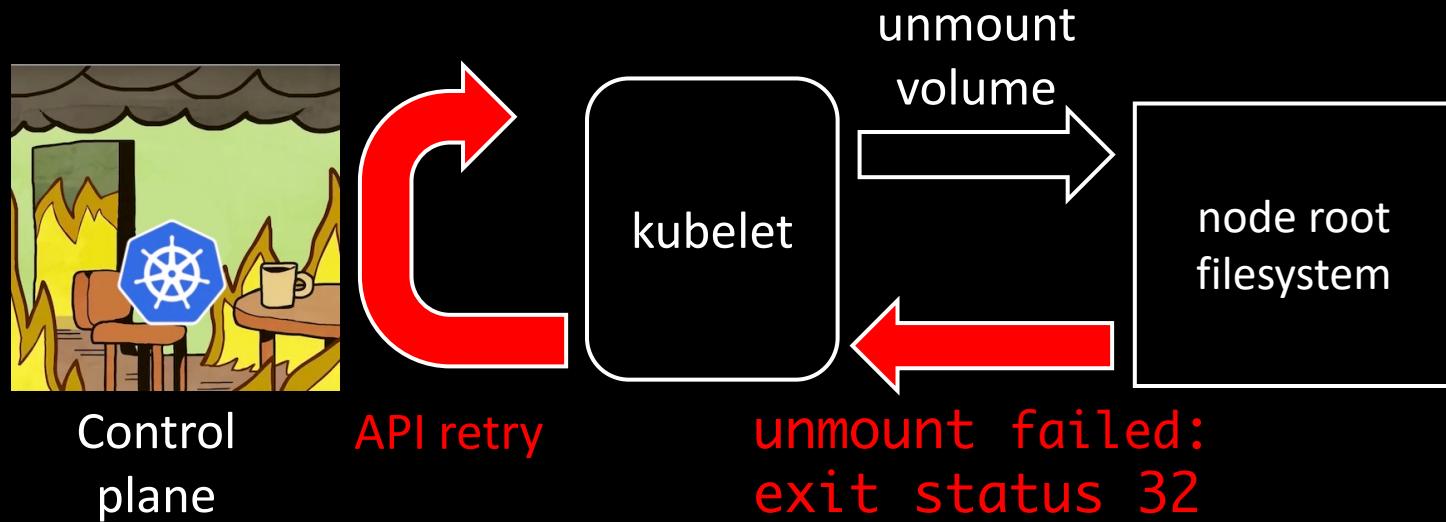
# Error 32

```
Error: UnmountVolume.TearDown failed for volume "kube-api-access-bc5sr" (UniqueName: "kubernetes.io/projected/693f4a5e-f879-47af-baae-6235cccd667b-kube-api-access-bc5sr") pod "693f4a5e-f879-47af-baae-6235cccd667b" (UID: "693f4a5e-f879-47af-baae-6235cccd667b") :  
unmount failed: exit status 32
```

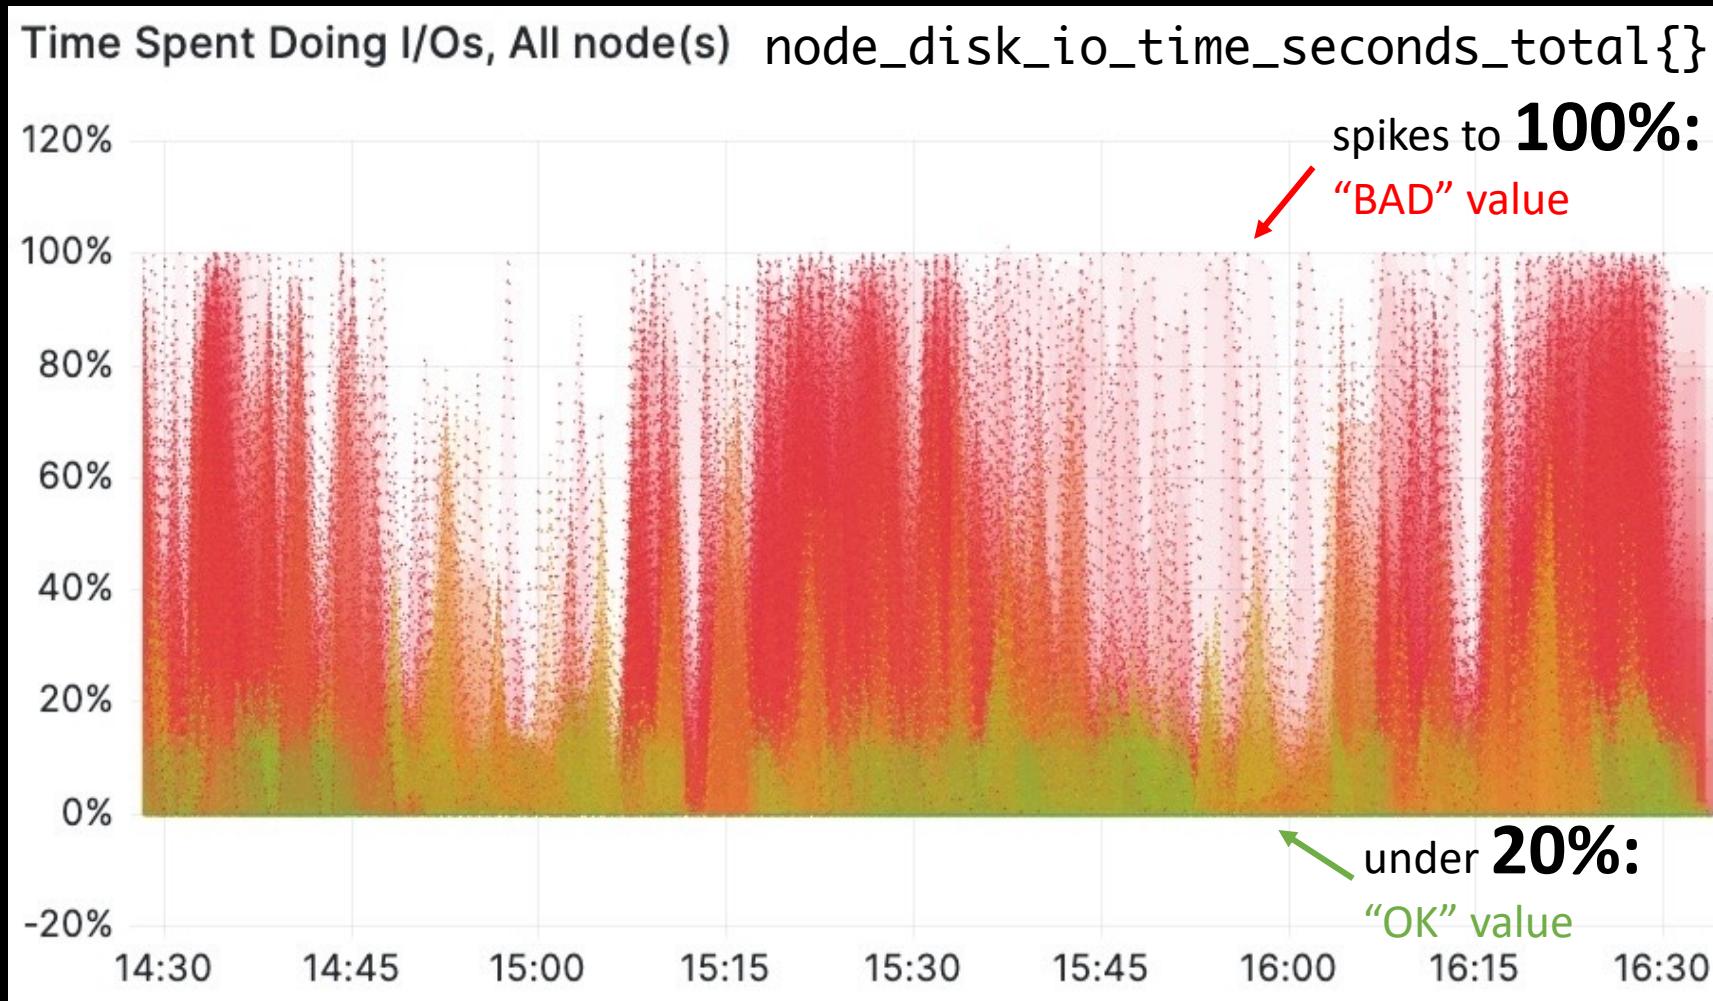
Always look at the logs!



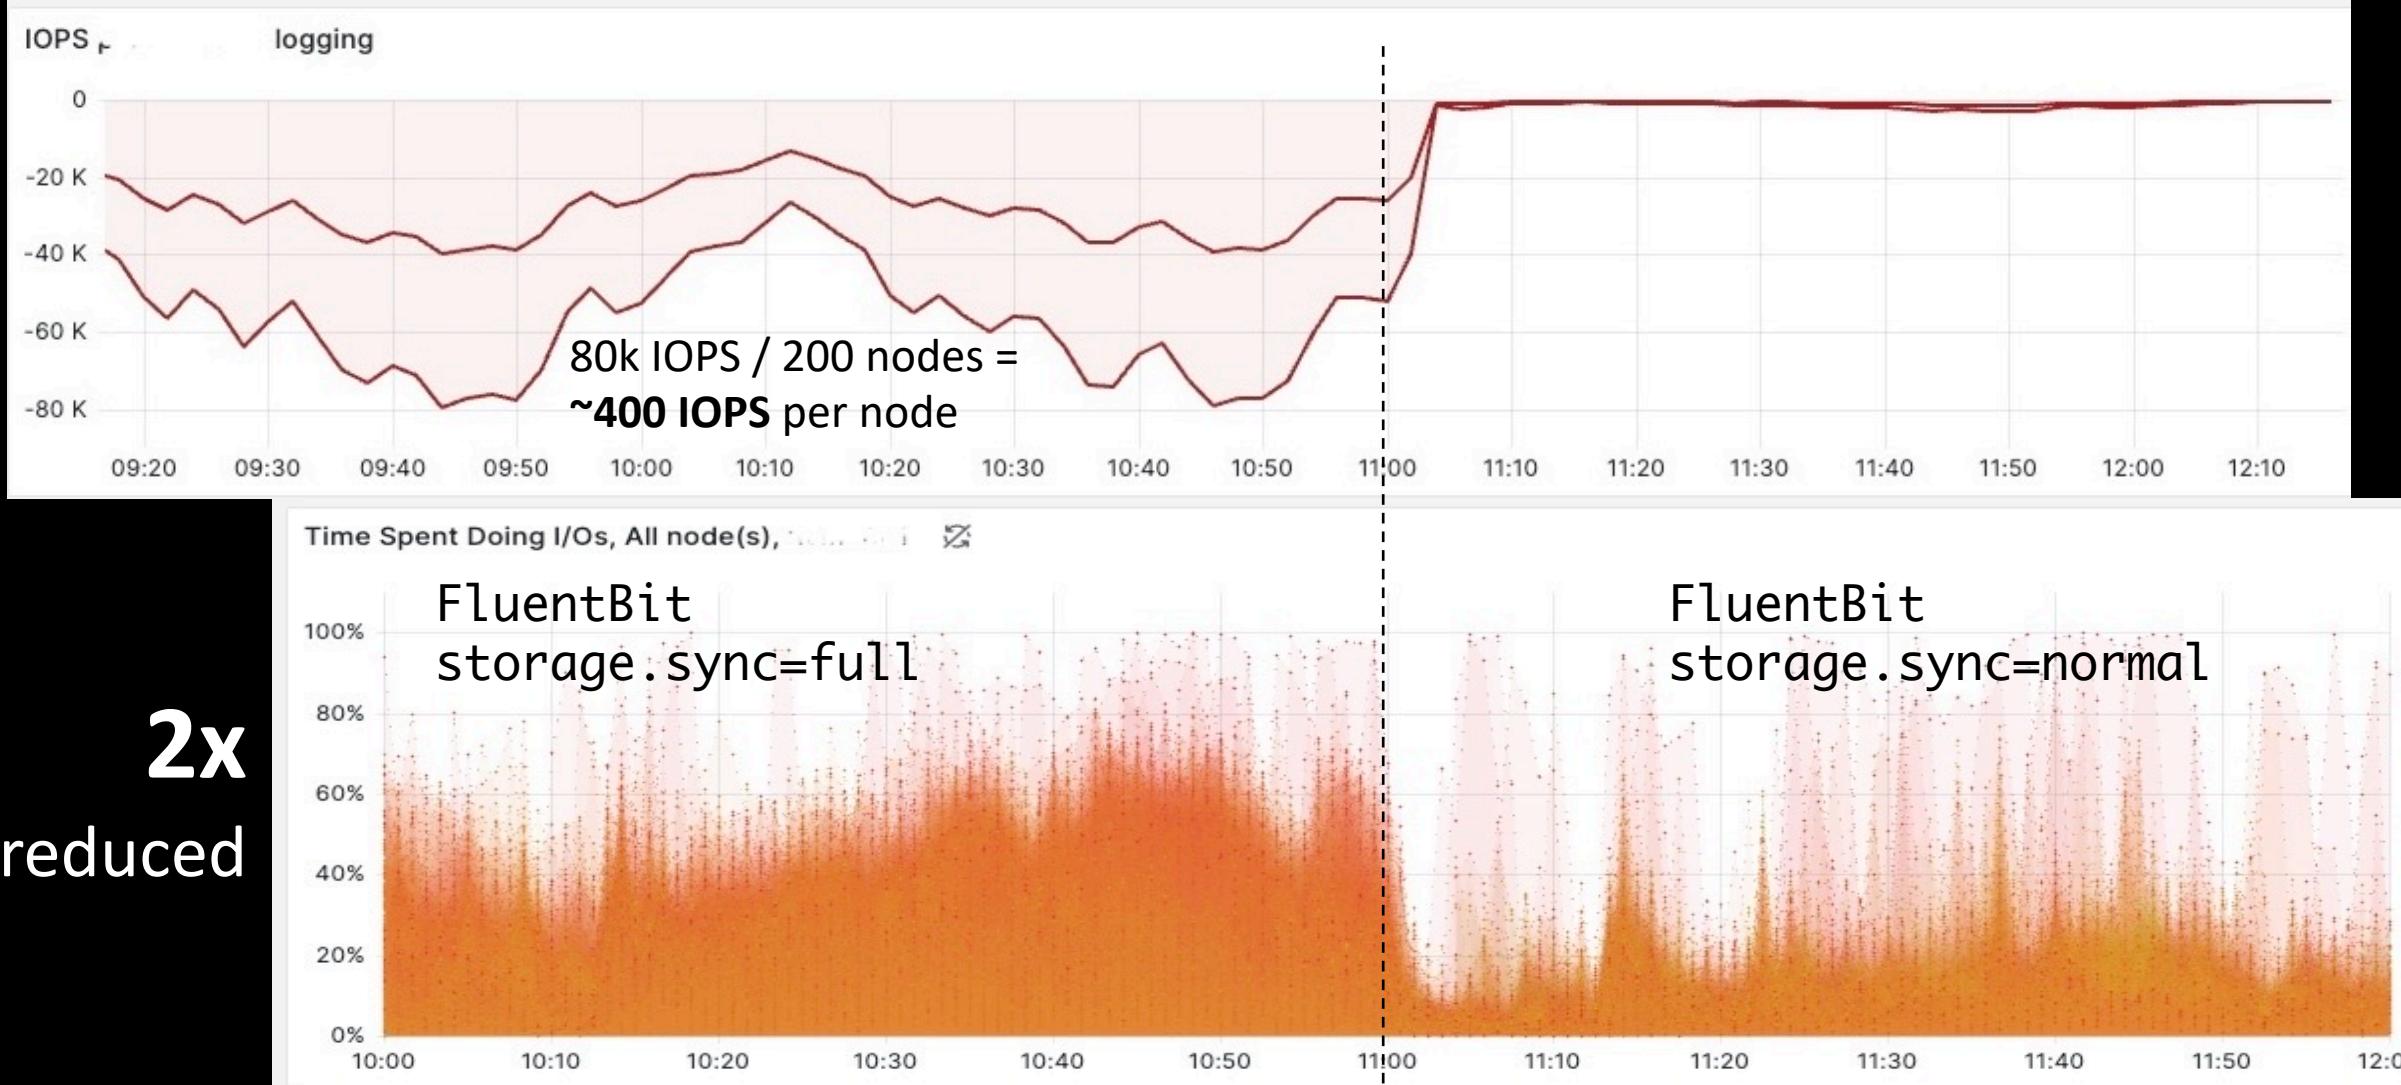
# The Good, the Bad, and the Control Plane

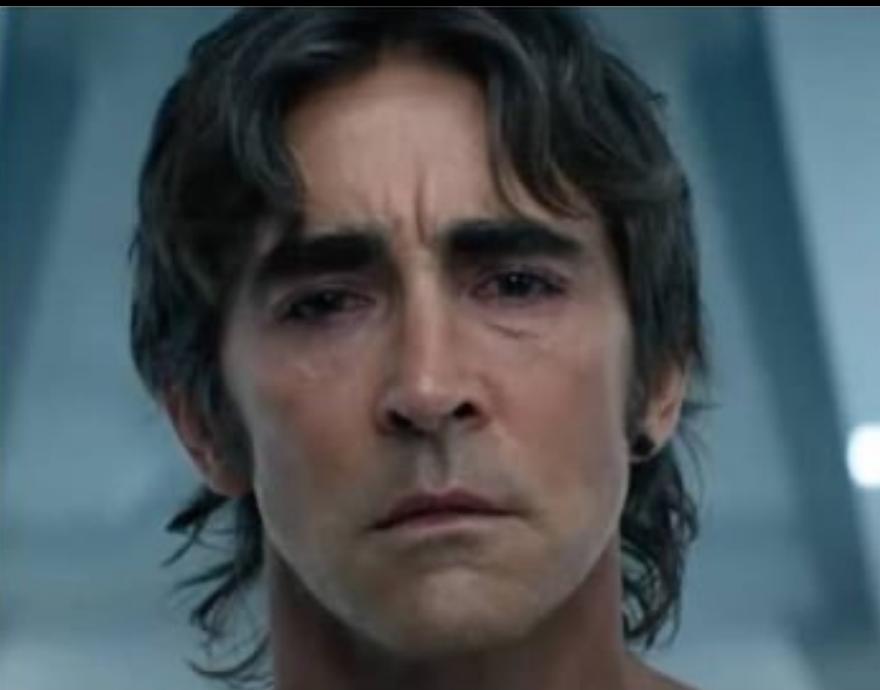


# Storage: Red Alert



# The Truth Is Out There At The Node Level





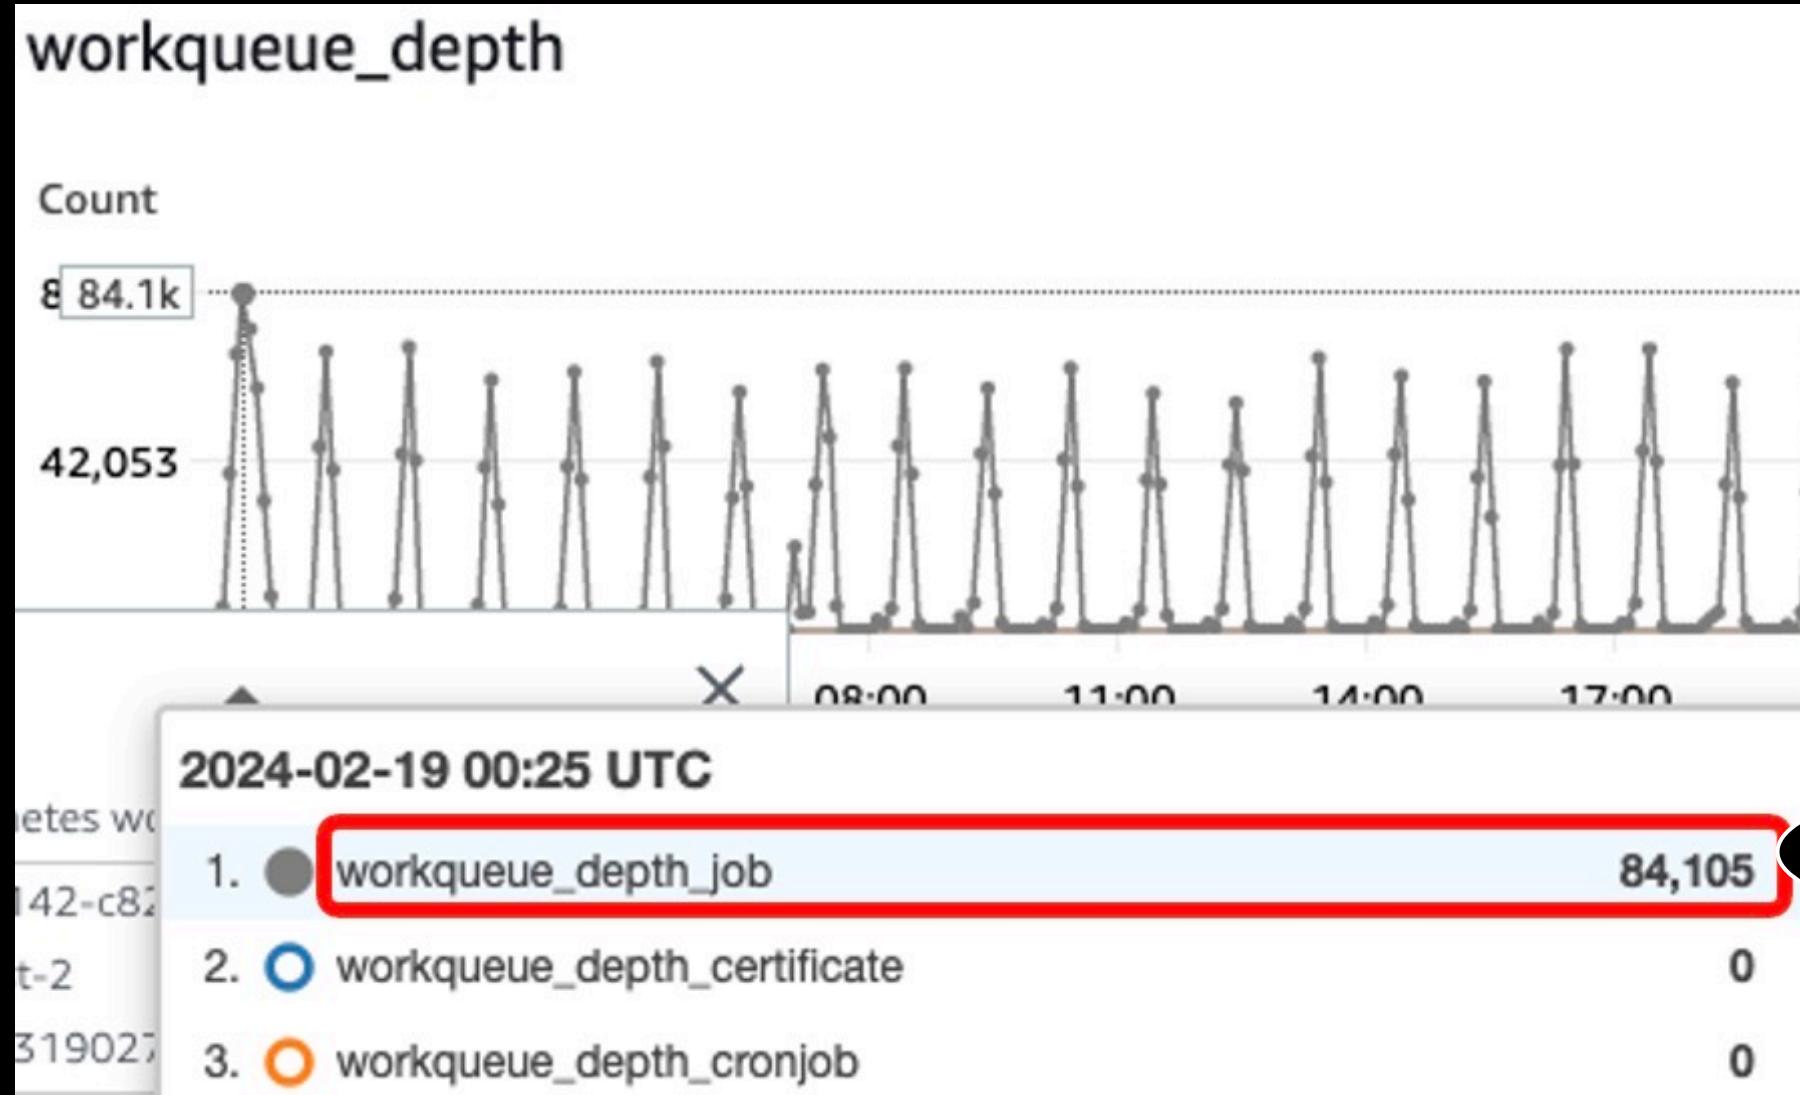
# Storage Goes Brrrrr

unmount failed:  
exit status 32



Control  
plane

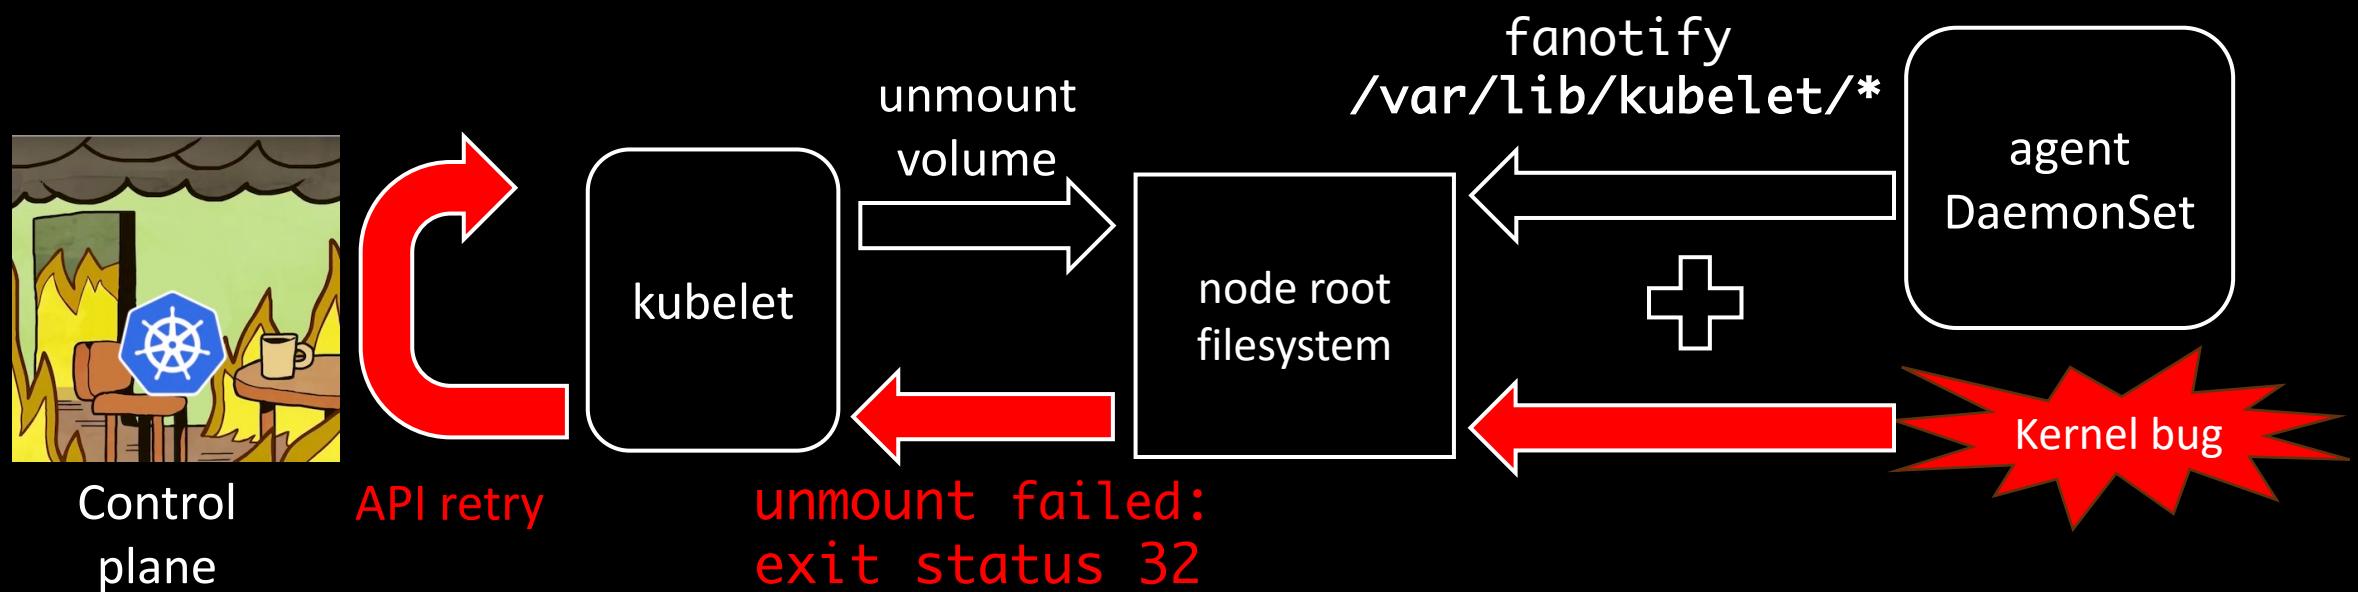
# Control Plane: New Low



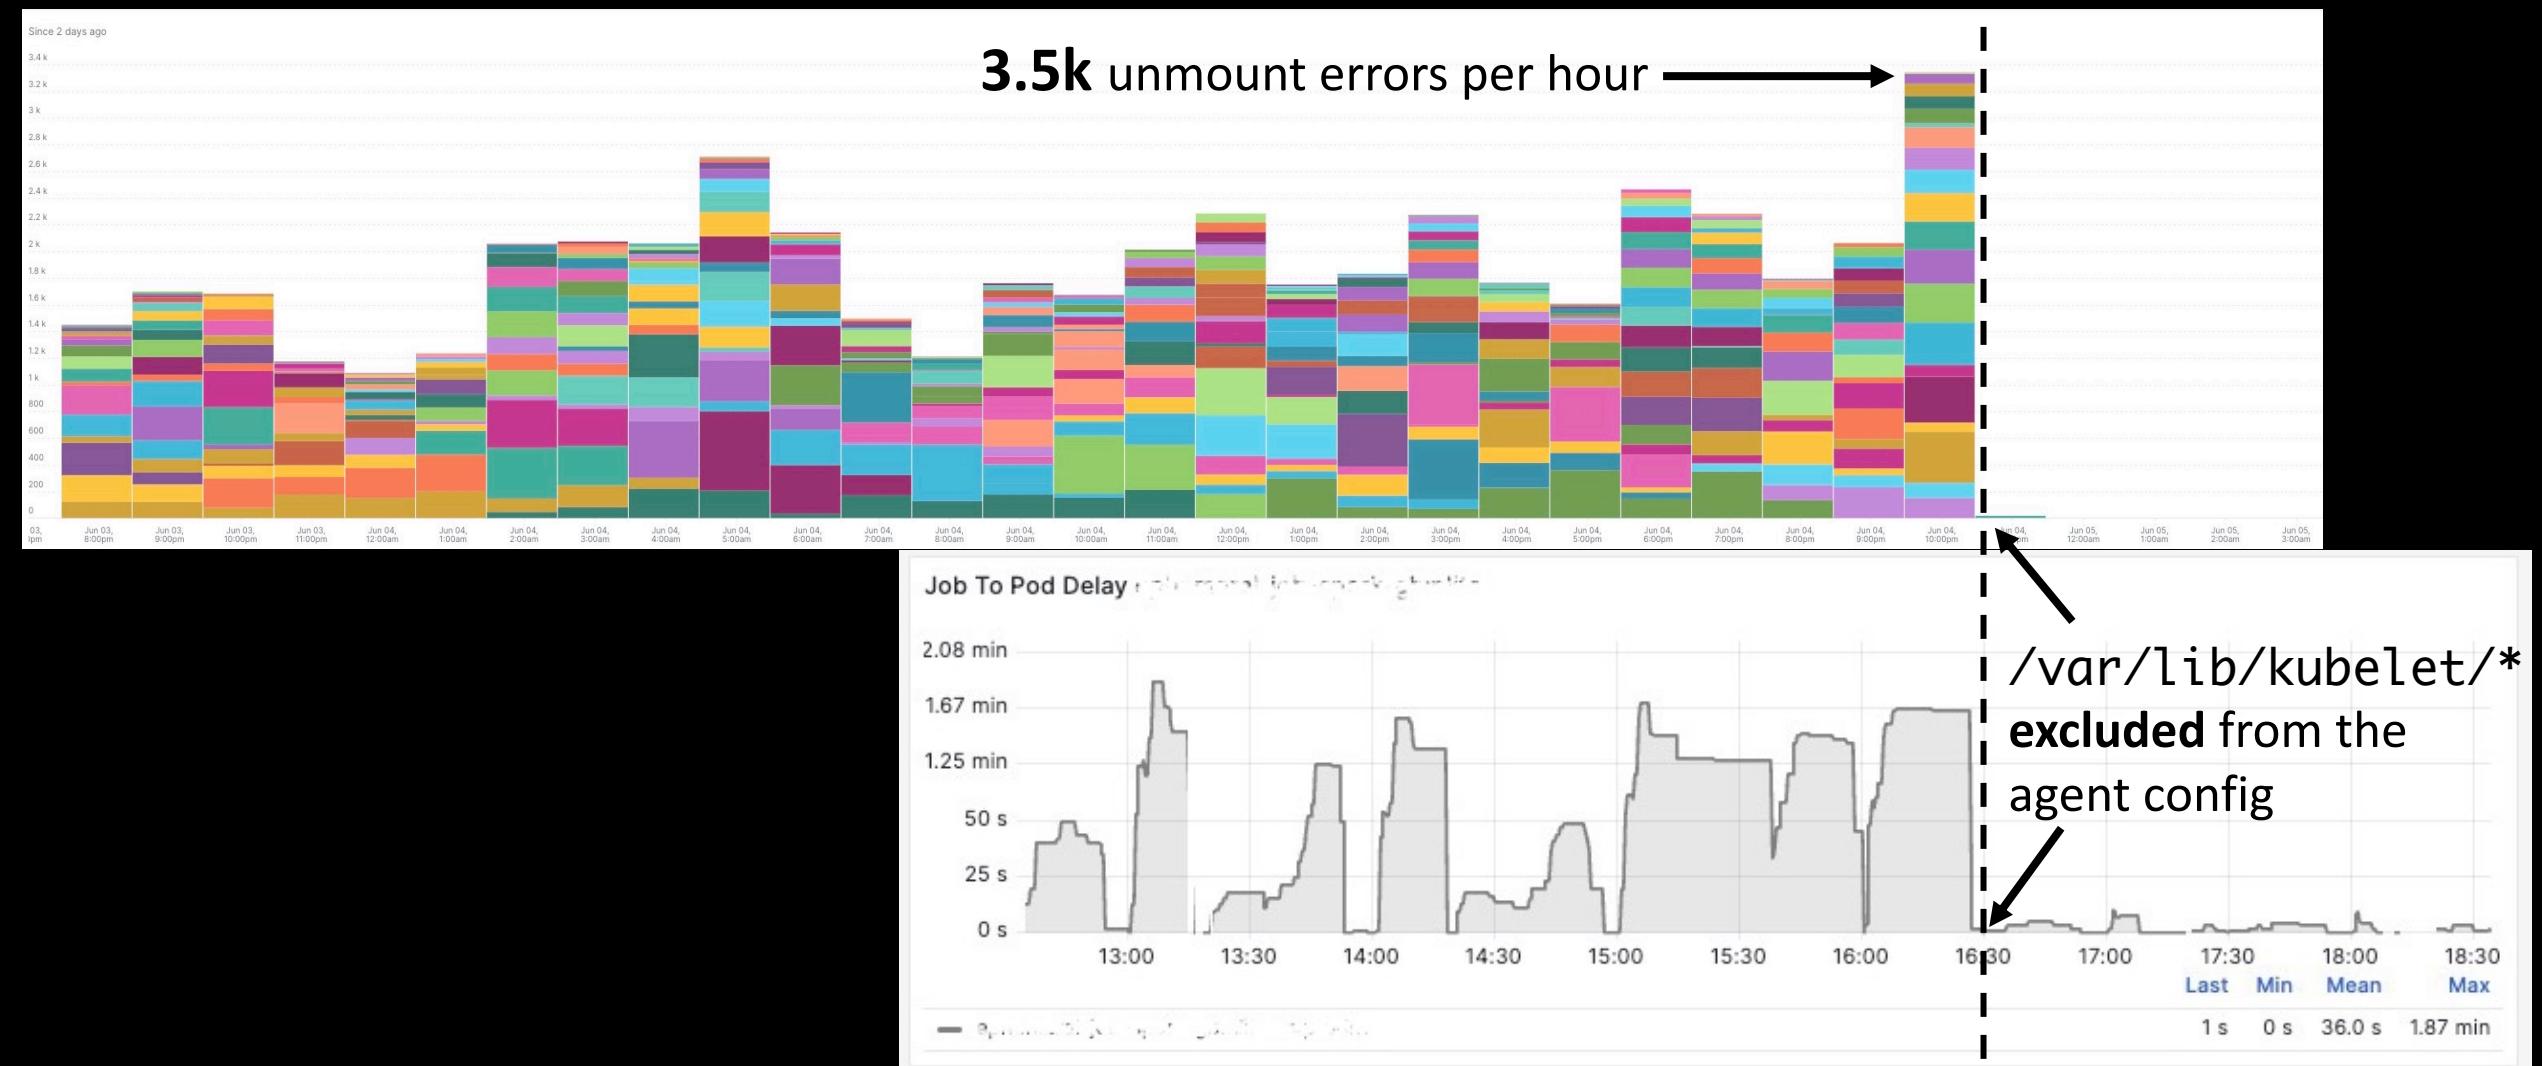
# Job to Pod Delay



# It Cannot be a DaemonSet



# It was a DaemonSet



# Custom kube-scheduler: Efficiency

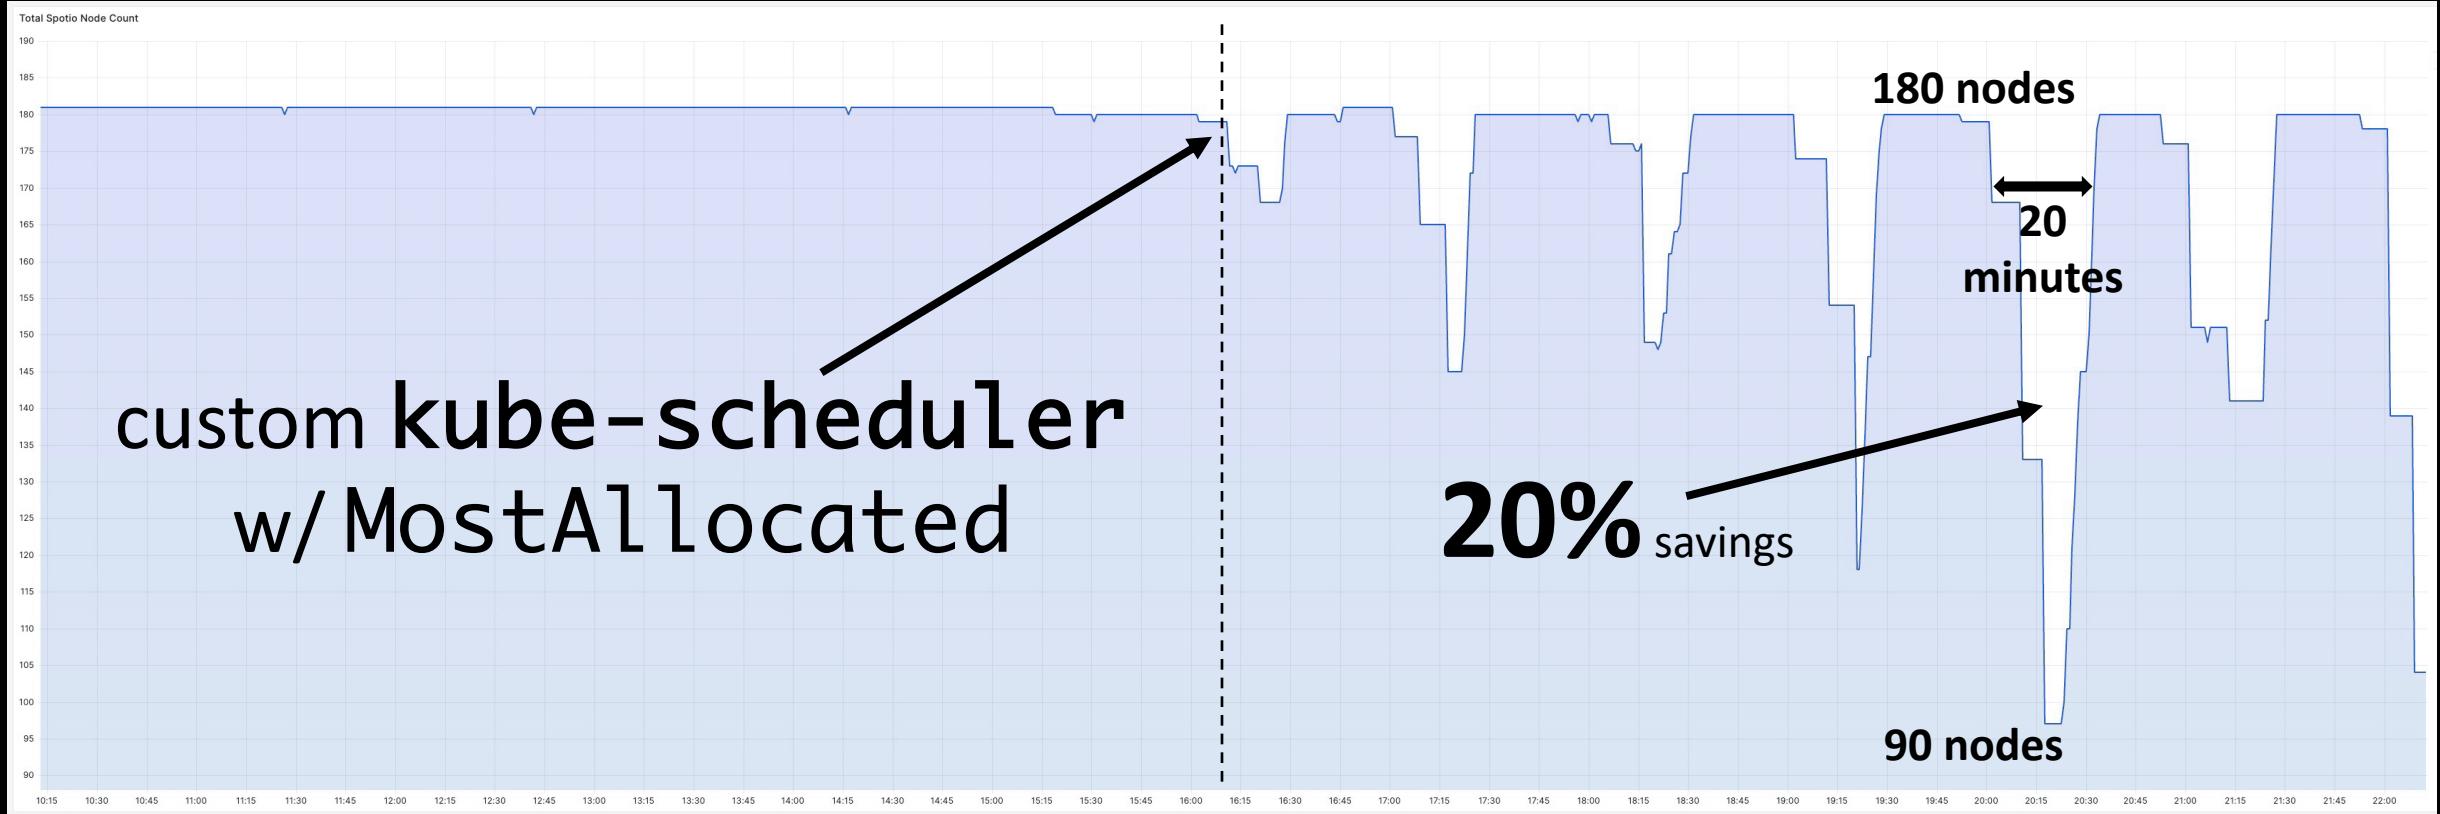
LeastAllocated (default)



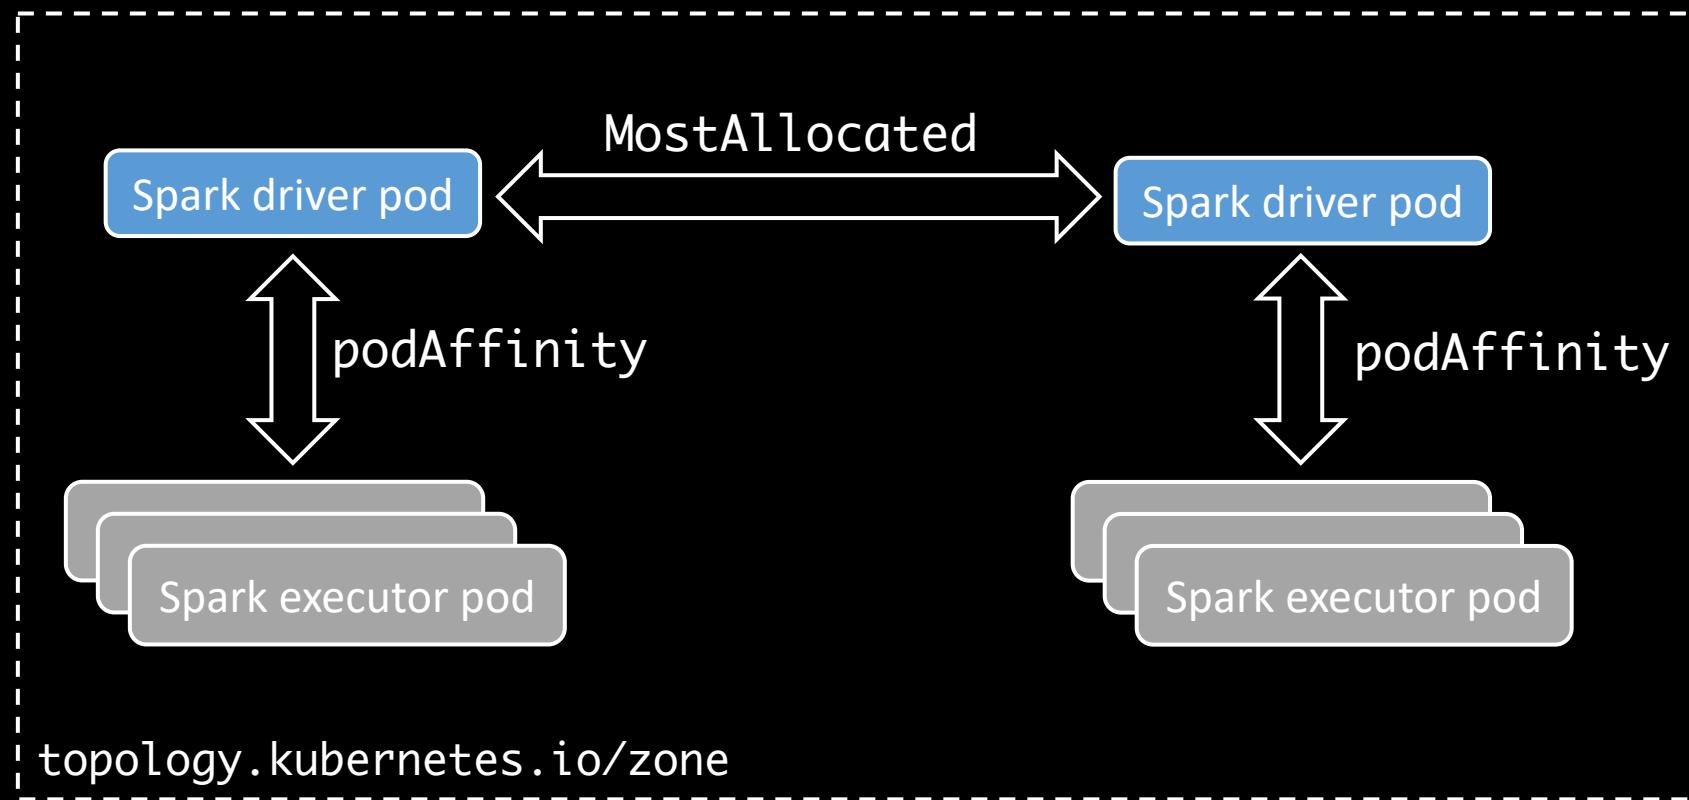
MostAllocated



# Custom kube-scheduler: Taste of Victory



# podAffinity + kube-scheduler: a Problem



# Very Default, Very Demure

**cluster-wide default**

*kube-scheduler*

**defaultConstraints:**

- maxSkew: 5

topologyKey: "topology.kubernetes.io/zone"

Suggestion, allows scheduling

**whenUnsatisfiable: ScheduleAnyway**

# Enforced Pod Spread

Match number of pods  
across zones

```
Spark driver pod
topologySpreadConstraints:
  - maxSkew: 1
    minDomains: 3
    topologyKey: topology.kubernetes.io/zone
  → whenUnsatisfiable: DoNotSchedule
    labelSelector:
      matchLabels:
        → spark-pod-kind: driver
      matchLabelKeys:
        - tier
```

Enforced, **blocks scheduling**

Spread drivers, executors follow

# No Tricks, Just Teamwork



Neil  
Chao

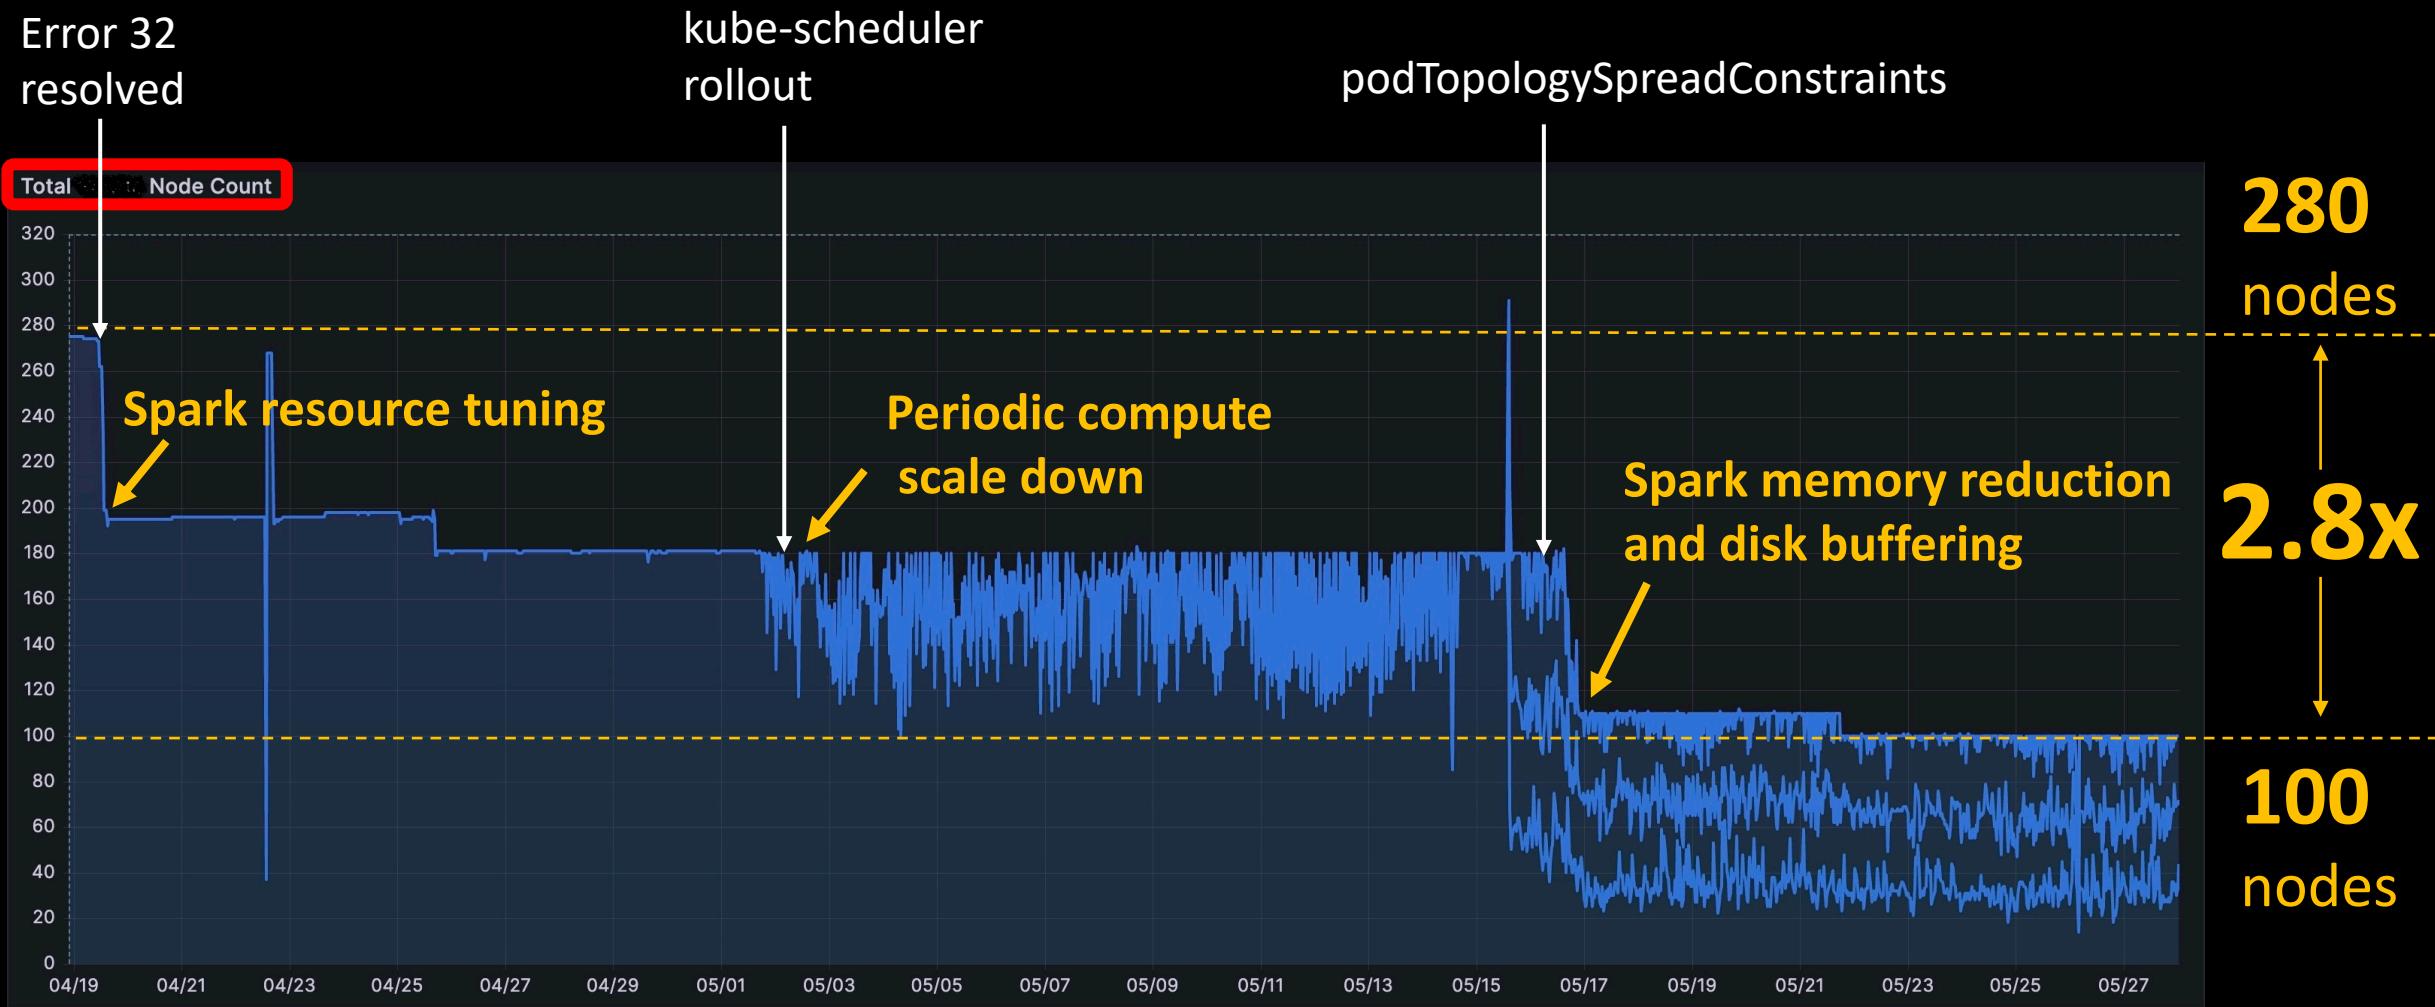


Chonghan  
Chen



Craig  
Skeinfell

# No Tricks, Just Deliberate Optimization

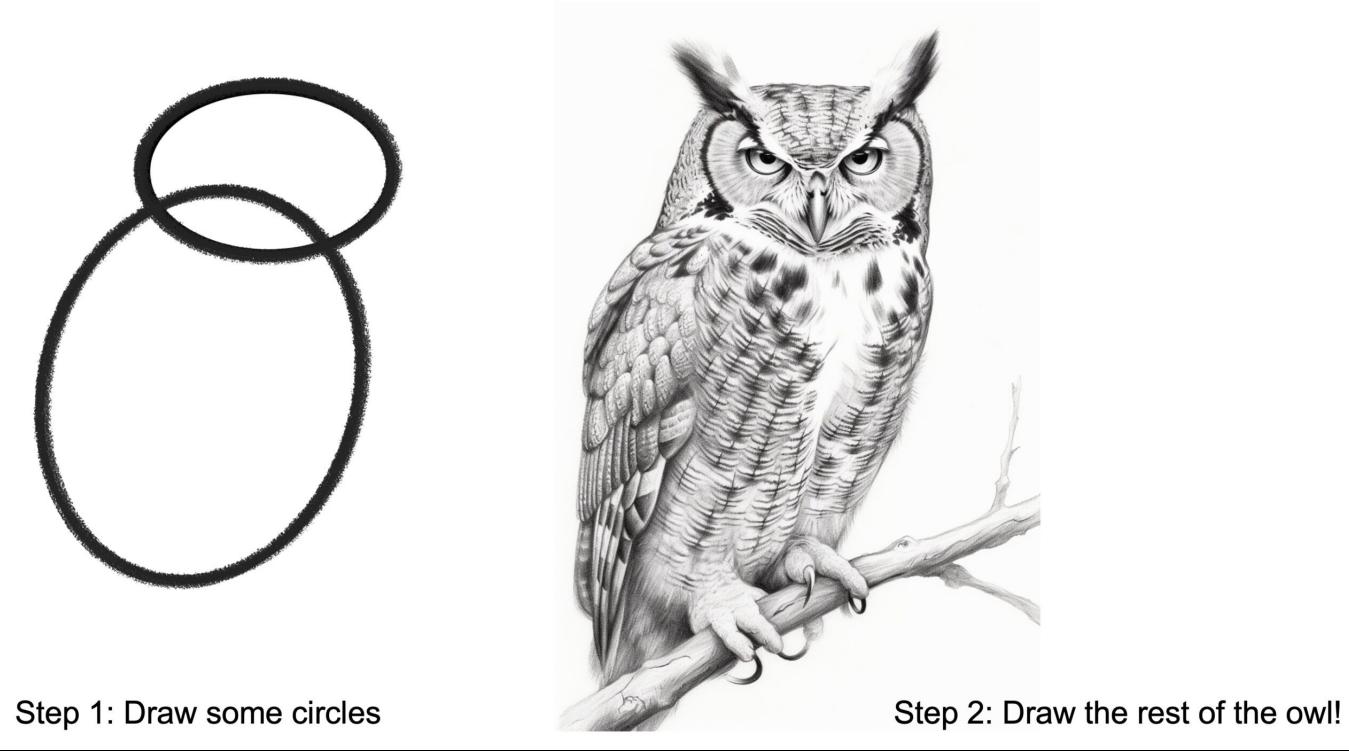


# The Secret of Scale

- Find the right flow control settings
- Find chaos causing clients
- Optimize the workload



# Optimize the workload



## Guides

[https://aws.github.io/aws-eks-best-practices/scalability/docs/node\\_efficiency/](https://aws.github.io/aws-eks-best-practices/scalability/docs/node_efficiency/)  
[https://aws.github.io/aws-eks-best-practices/scalability/docs/scaling\\_theory/](https://aws.github.io/aws-eks-best-practices/scalability/docs/scaling_theory/)

## Talk

[https://www.youtube.com/watch?v=NqtfDy\\_KAqg](https://www.youtube.com/watch?v=NqtfDy_KAqg)

Want More?

