



KubeCon



CloudNativeCon

North America 2024

# Optimizing Load Balancing and Autoscaling for LLM Inference on Kubernetes

David Gray

Software Engineer – Red Hat Performance and Scale

## Hi, I'm David!

- Based out of Toronto, ON, Canada
- Past: k8s operator development for out-of-tree kernel driver enablement and node tuning
- Present: LLM Inference performance on k8s

Part of the Performance and Scale for AI Platforms (PSAP) team at Red Hat

- Making AI applications run better and faster on Linux, Containers, and Kubernetes

[Explore the Red Hat Performance and Scale Blogs](#)



## What we'll discuss today

- Background
- Deploying LLMs on Kubernetes:  
vanilla and KServe
- Load-balancing for LLM Inference
- Pod autoscaling for LLM Inference
- Conclusion and future ideas
- Q&A



KubeCon

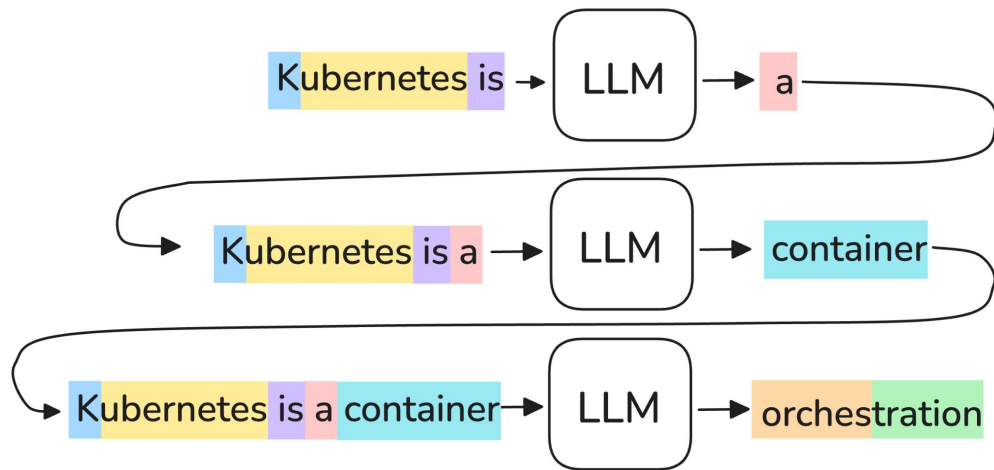


CloudNativeCon

North America 2024

# Background

# Background and motivation



Why optimize inference performance?

- Cost-efficiency
- Energy efficiency
- Trend towards inference with long sequences:
  - RAG
  - Chain-of-thought

# Model servers (inference engines / runtimes)

Examples:

- [vLLM](#)
- [text-generation-inference](#) 🙌
- [TensorRT-LLM](#)
- [SGLang](#)

Software for loading and running models

- HTTP / gRPC server
- Optimized accelerator kernels
- Batch processing
- Performance optimizations

# Measuring LLM inference performance



KubeCon



CloudNativeCon

North America 2024

## Latency

- Time-to-first-token (TTFT)
- Inter-token latency (ITL)
- Response time: depends on the number of tokens

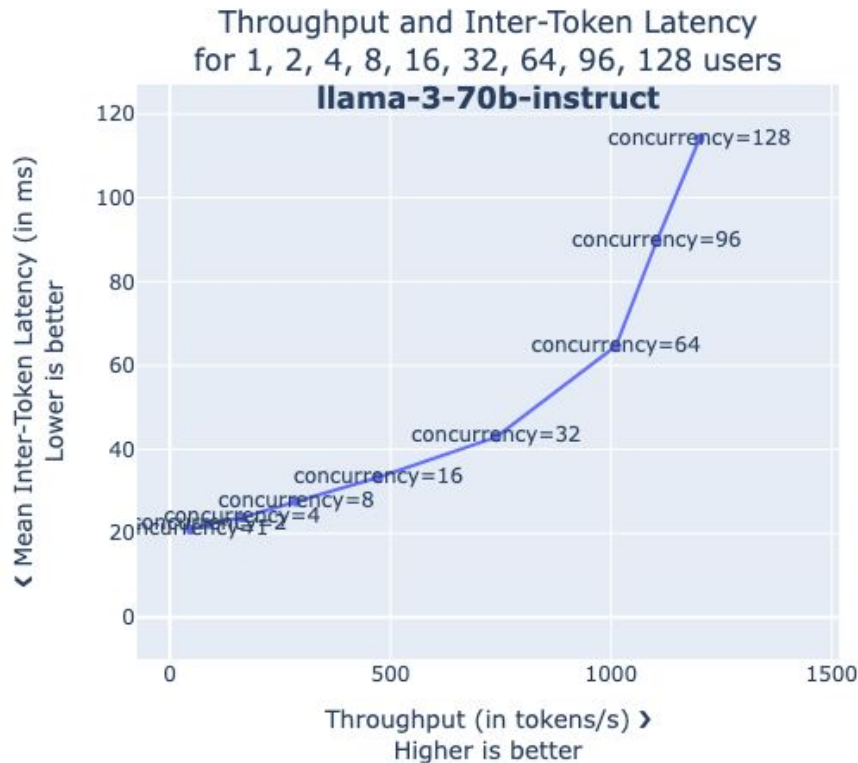
## Throughput

- Tokens / second
- Across all concurrent requests

## Benchmarking tools

- [llm-load-test](#)
- [fmpref](#)

Example data:





KubeCon



CloudNativeCon

North America 2024

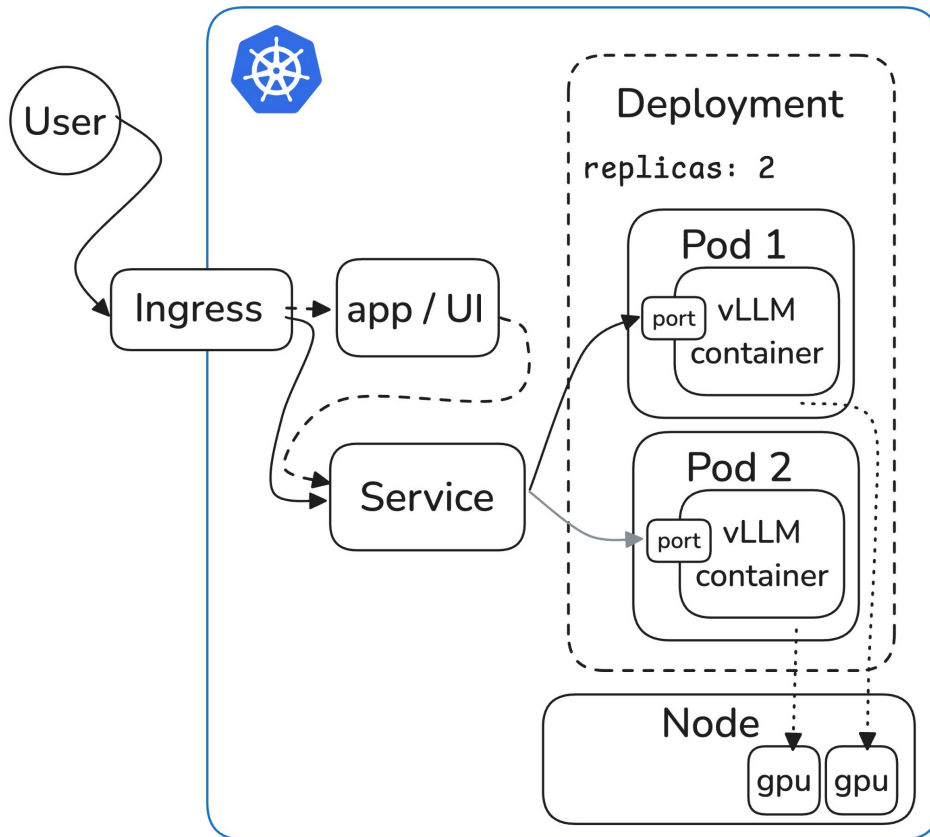
# LLM Inference on Kubernetes



# LLM Inference on Kubernetes

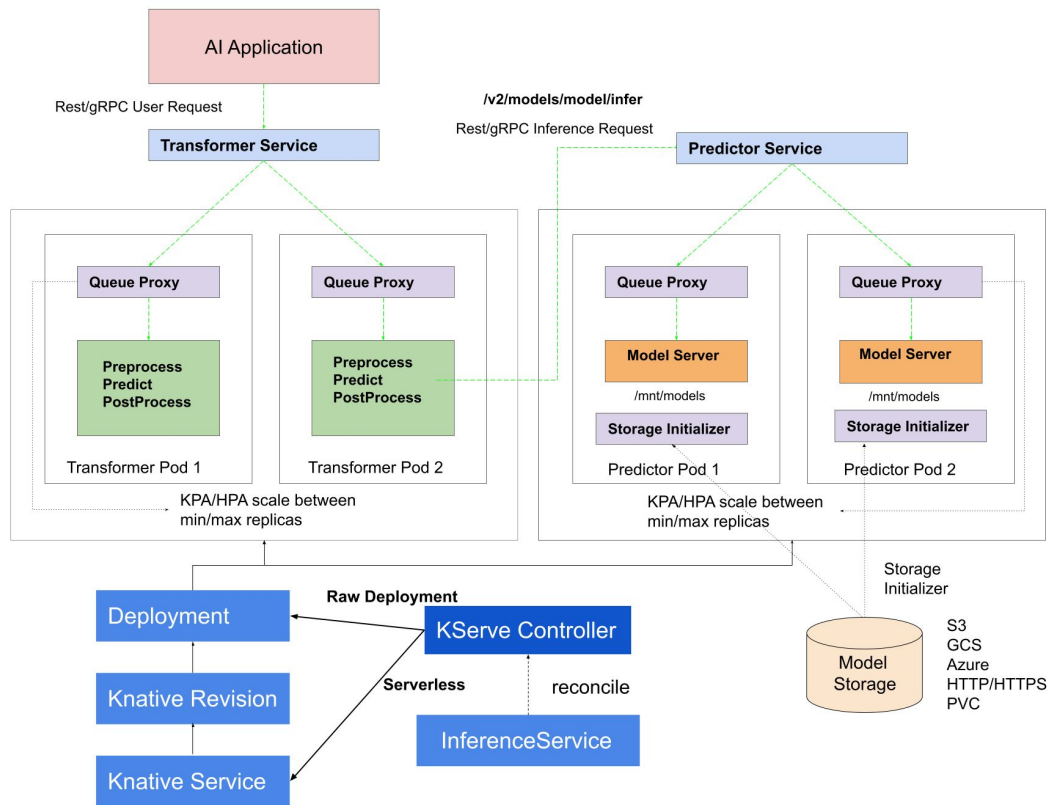
Deploying an LLM model on k8s the “vanilla way”

- Deployment + Service
- Part of another application, or behind a preprocessing app.



# KServe Introduction

- CRDs
  - ServingRuntime
  - InferenceService
- Deployment modes
  - RawDeployment
  - Knative
- Knative load balancing
  - Least-requests based
- Knative Pod Autoscaler (KPA)



[KServe control plane documentation](#)

For more info, see [Yuan Tang's talk on KServe features](#)



KubeCon



CloudNativeCon

North America 2024

# Load balancing for LLM inference

# Load balancing: Test methodology

Model: Llama-3.1-8b

Hardware: AWS g5.48xlarge

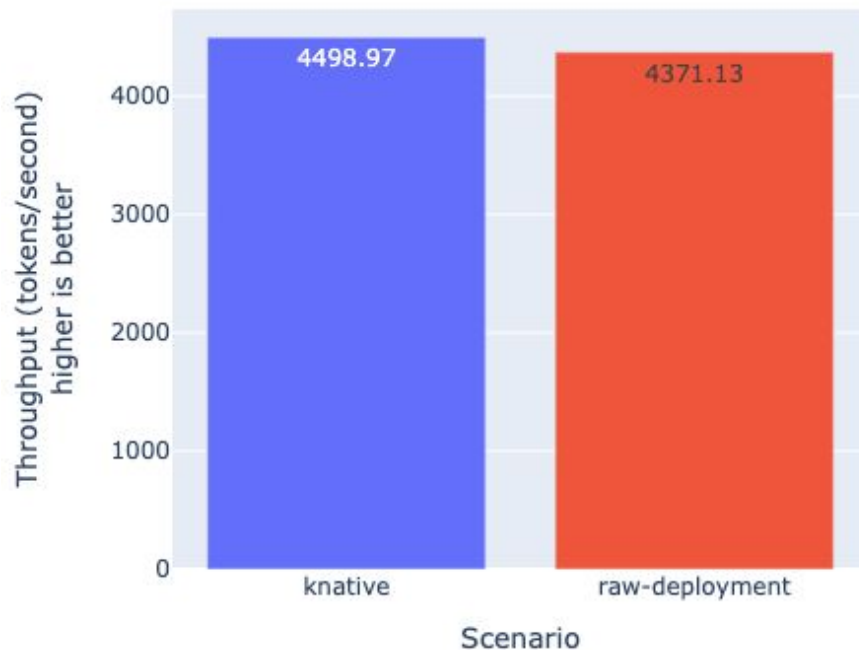
8 replicas, each on 1xA10G GPU (24GB)

llm-load-test parameters

- Dataset: subset of [OpenOrca](#)
- Input tokens:
  - 16 - 1600 tokens
  - mean: ~540 tokens
- Output tokens:
  - 16 - 1600 tokens
  - mean: ~415
- RPS=12

## Throughput by load balance strategy rps=12, 8 replicas

llama-3.1-8b on g5.48xlarge



# Load balancing: Knative vs RawDeployment



KubeCon

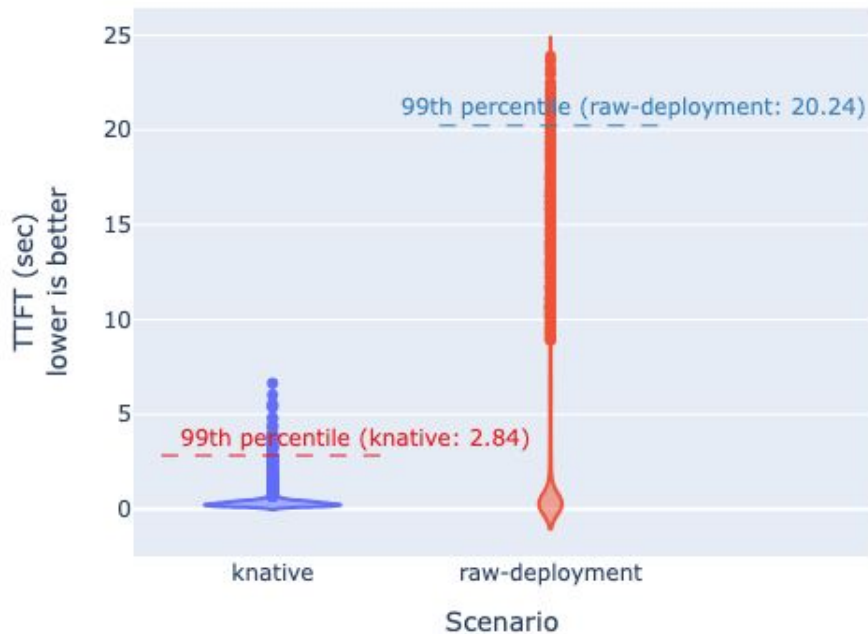


CloudNativeCon

North America 2024

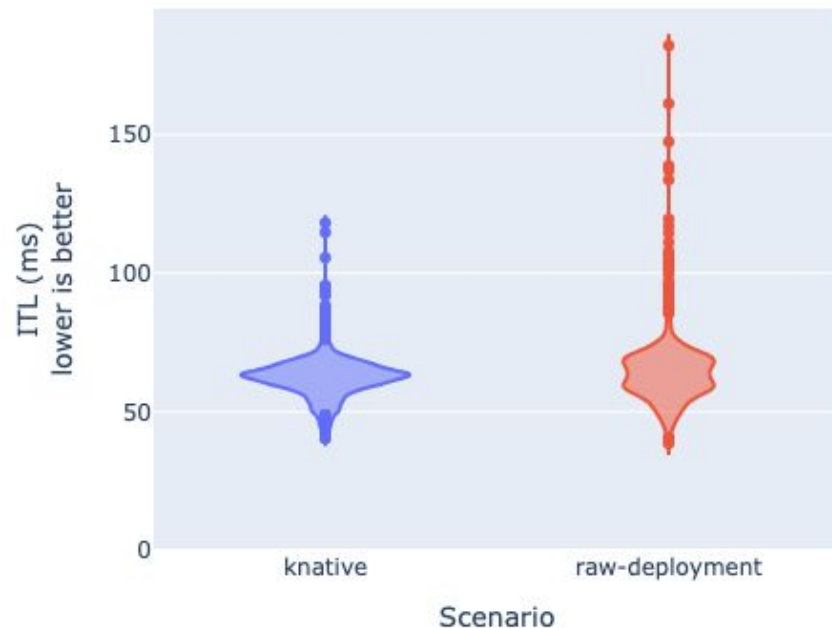
TTFT by load balance strategy rps=12, 8 replicas

llama-3.1-8b on g5.48xlarge



ITL by load balance strategy rps=12, 8 replicas


llama-3.1-8b on g5.48xlarge



# Load balancing: Knative vs RawDeployment

Large request queues on some replicas

- Max queue size (num\_requests\_waiting):
  - RawDeployment: **36**
  - Knative: **5**



vLLM metric: num\_requests\_waiting  
(raw-deployment)



# Custom load balancing

Issues with requests-based load balancing

- Maximum batch size is dynamic
- Batch size is limited by KV-cache
- KV-cache usage depends on:
  - Number of requests
  - Sequence lengths of requests
  - Model size

Some replicas near 100% KV-cache utilization while others < 80%

vLLM metric: GPU KV-cache usage %:



# Custom load balancing results

## Custom strategy

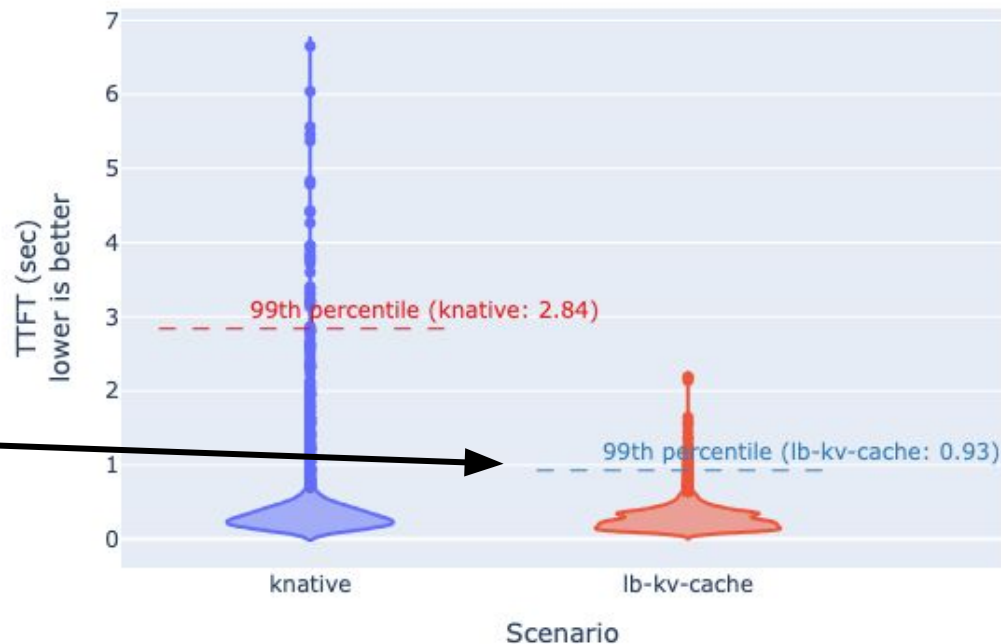
- Scrape KV-cache usage %
- Pick 2 replicas randomly
- Send to lower of the 2

## Implemented client-side (PoC)

99%ile TTFT under 1 second!

## TTFT by load balance strategy rps=12, 8 replicas

llama-3.1-8b on g5.48xlarge





# Custom load balancing: higher load

At RPS=13, Knative beats the custom strategy


The problem under high load

- Maximum num\_requests\_waiting:
  - Knative: **26**
  - Custom: **36**
- KV-cache usage >98% for all replicas

A better solution should prioritize queue size before KV-cache usage %

## TTFT by load balance strategy rps=13, 8 replicas

llama-3.1-8b on g5.48xlarge



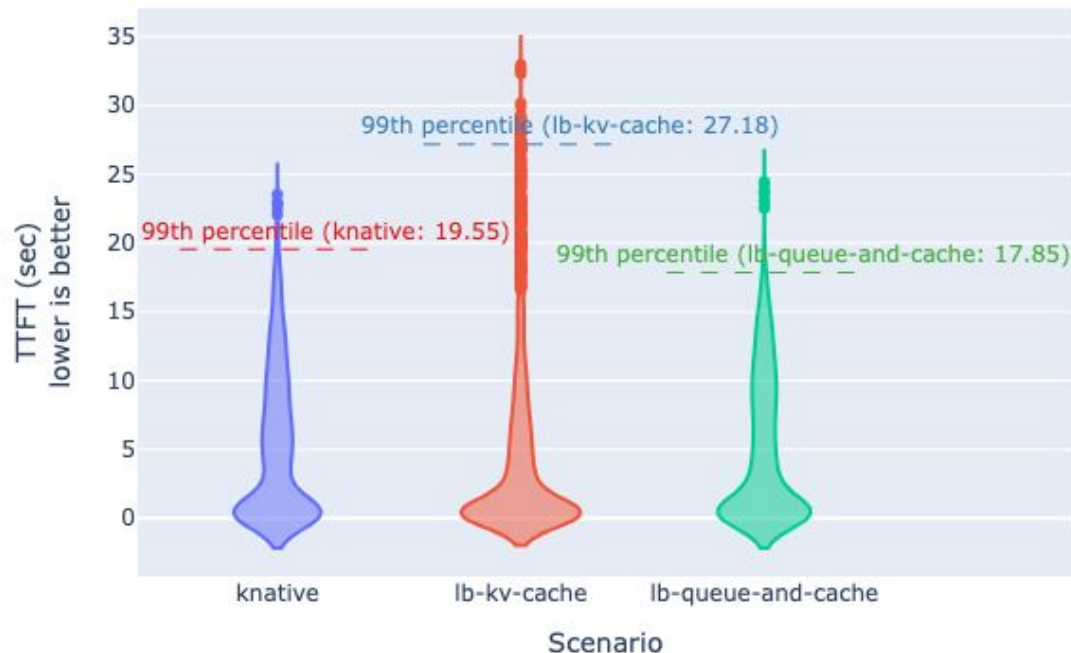
# Custom load balancing: better strategy

## Better strategy

- Scrape num\_requests\_waiting and KV-cache usage %
- Pick 2 replicas randomly
- Send to the one with lower num\_requests\_waiting
  - if tied, send to lower KV-cache usage %

## TTFT by load balance strategy rps=13, 8 replicas

llama-3.1-8b on g5.48xlarge



# An aside: Replicas vs tensor parallelism

Options to get more throughput

- More replicas
- tensor-parallelism (TP)
- pipeline-parallelism

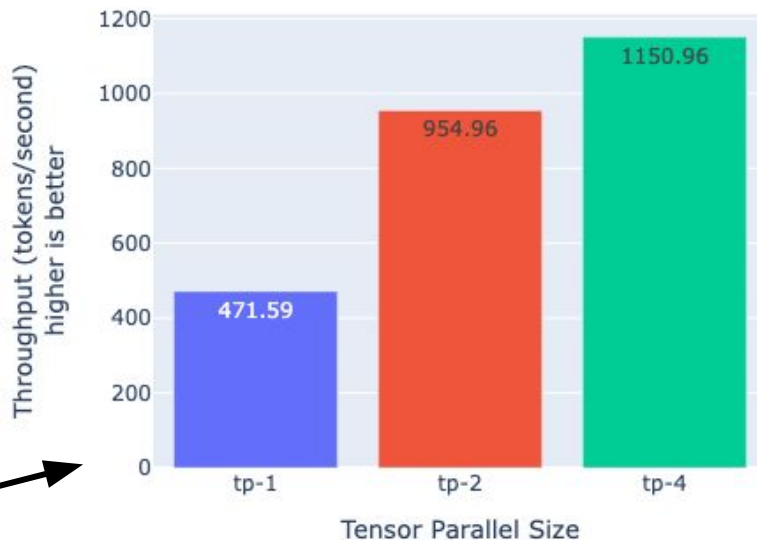
When to use tensor-parallelism

- Model doesn't fit in GPU memory
- To improve per-token latency\*
- Need more GPU memory for KV-cache
  - To enable higher batch sizes or longer sequences
  - Example: 20B model in 48GB GPU

\* TP performance depends on GPU interconnect

Throughput with 95% ITL < 100 ms, varying num. GPUs

Higher is better, granite-20b-instruct, L40S GPUs





KubeCon



CloudNativeCon

North America 2024

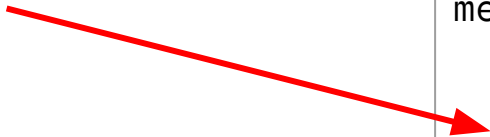
# Autoscaling for LLM Inference

# Pod autoscaling for LLM inference

How to know when to scale up

- KServe / Knative:
  - Concurrency
  - RPS
  - CPU/Memory
- Inference-engine metrics:
  - KV-cache utilization %
  - Queue depth  
(num\_requests\_waiting)
  - ITL, TTFT

KServe feature requests for KEDA  
integration: [#3561](#), [#4007](#)



```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  annotations:
    # Soft limit to trigger scaleup
    autoscaling.knative.dev/target: "30"
    ...
  name: llama-3-1-8b-isvc
  namespace: models
spec:
  predictor:
    # Hard limit enforced by queue-proxy
    containerConcurrency: 50
    ...
```

# Pod autoscaling

Delays before model is ready

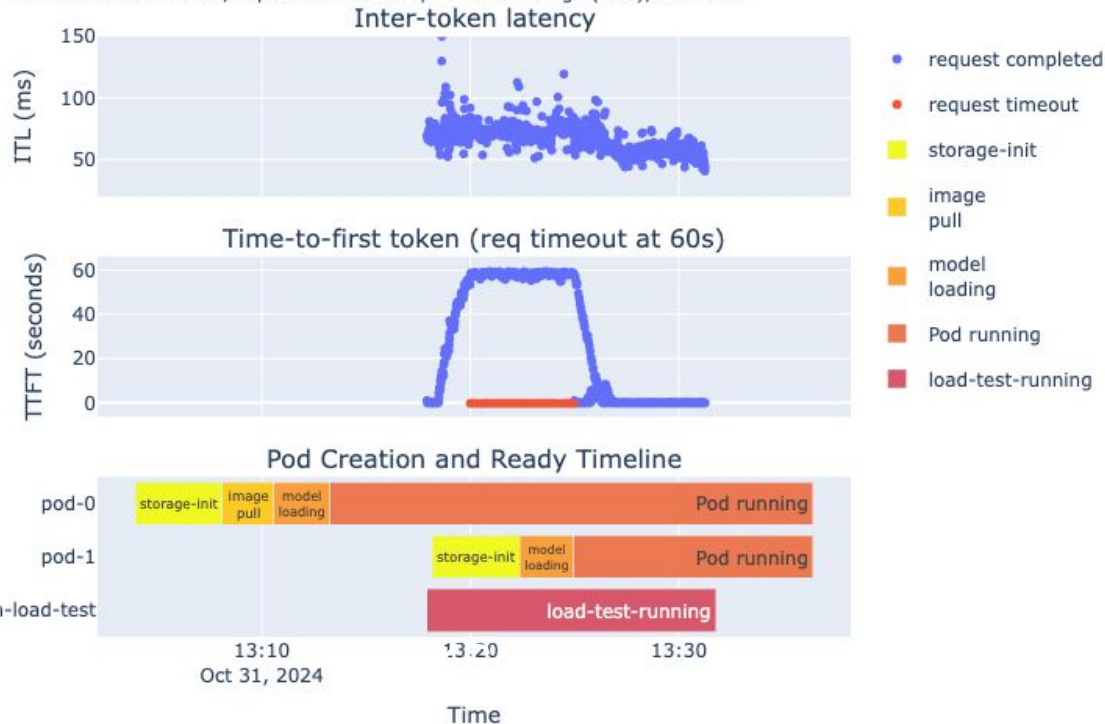
- Pull runtime image
- Download model files
- Load model from storage into GPU memory
- Graph warm-up

Solutions

- Cache models locally
- Use faster local storage
- Store image + model files somewhere with fast upload speeds

Latency under load while scaling model replicas from 1 to 2

Llama-3.1-8B-instruct, copied from S3 to ephemeral storage (EBS), RPS=2.5




# Autoscaling optimization: modelcars

## KServe Modelcars

- Load model files from containers
- Leverages k8s image caching
  - Simplifies orchestration compared to local storage
- Issues with large models (>20GB)
- Storage speed (r/w) matters
  - Ex: gp3 to io1 EBS volume

## Latency under load while scaling model replicas from 1 to 2

Llama-3.1-8B-instruct, loaded from modelcar (io1, 64k IOPS), RPS=2.5





KubeCon



CloudNativeCon

North America 2024

# Conclusion



## Summary

- Requests-based load balancing is well suited for LLM inference.
  - Custom load balancing can further improve balancing, reducing 99% TTFT
- Autoscaling performance is dependent on fast storage
- Modelcars simplifies orchestration for caching model files locally

## Future ideas

- Autoscaling + load-balancing experiments with other metrics
- KServe support for autoscaling on custom metrics (KServe issues [#3561](#), [#4007](#))
- [fastsafetensors](#) to load models with GPU Direct Storage
- More enhancements to load testing tools: [join WG-Serving](#) for more on this!



KubeCon



CloudNativeCon

North America 2024

# Thank you!

Don't hesitate to reach out

[dagray@redhat.com](mailto:dagray@redhat.com)



[dagrayvid on GitHub](#)



[David Whyte-Gray on LinkedIn](#)