



KubeCon



CloudNativeCon

North America 2024


The OTTL Cookbook

Tyler Helmuth, Honeycomb
Evan Bradley, Dynatrace

Agenda

1. Quick intro to the OpenTelemetry Collector and OTTL
2. Solve a few scenarios with OTTL
3. Audience scenarios

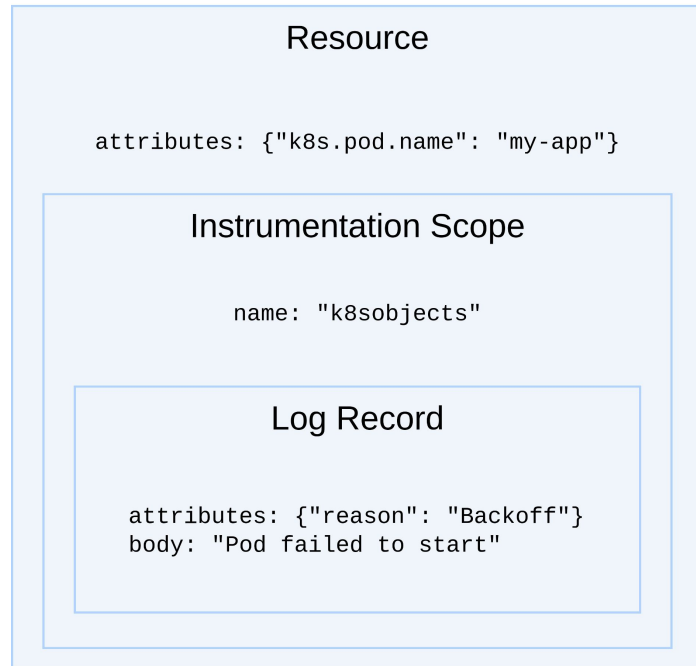
Quick Collector Intro



OTTL is a Domain Specific Language (DSL) tailor-made for interacting with telemetry in the OTel Collector.

- Provides access to all OTLP fields.
- Enables expressing complex transformations using a simple syntax.

Log



Use Case: Parse Unstructured Logs

We're receiving unstructured MySQL logs

Recipes:

- Set OTLP log fields
- Parse into a structured format
- Conditionally set values
- Update an OTLP resource

Log example

```
# Time: 2024-10-10T19:34:04.232967Z  
# Query_time: 0.776920  
SELECT * FROM my_table
```

Recipe: Set a Field

```
set(severity_number, SEVERITY_NUMBER_INFO)  
set(severity_text, "INFO")
```

severity_number: 0
severity_text:

severity_number: 9
severity_text: INFO

Recipe: Parse Unstructured Log

```
merge_maps(attributes, ExtractPatterns(body, "Query_time: (?P<query_time>[0-9\\.]+)"), "upsert")
```

attributes: {}
severity_number: 0
severity_text:

attributes:
 query_time: "0.776920"
severity_number: 9
severity_text: INFO

Recipe: Conditionally Set Values

```
set(attributes["slow"], true) where Int(attributes["query_time"]) > 0.7
```

```
attributes:  
  query_time: "0.776920"  
severity_number: 0  
severity_text:
```



```
attributes:  
  query_time: "0.776920"  
  slow: true  
severity_number: 9  
severity_text: INFO
```


Recipe: Update a Resource



KubeCon



CloudNativeCon

North America 2024

Context: resource

```
set(attributes["log.group.name"], "/db/instance/production-mysql-mycompany")
```

attributes: {}

attributes:
log.group.name: "/db/instance/..."

```
log_statements:
  - context: resource
    statements:
      - set(attributes["log.group.name"], "/db/instance/production-mysql-mycompany")

  - context: log
    statements:
      - set(severity_number, SEVERITY_NUMBER_INFO)
      - set(severity_text, "INFO")
      - merge_maps(attributes, ExtractPatterns(body, "<big regex here"), "upsert")
      - set(attributes["slow"], true) where attributes["query_time"] > 0.7
```

Use Case: Manipulating JSON

We're receiving JSON strings in our log bodies and need to process them.

Recipes:

- Parse a JSON log
- Work with timestamps
- Handle trace/span IDs
- Manipulate strings

Formatted JSON body

```
{
  "object": {
    "timestamp": 2024-10-29T22:17:30Z,
    "trace_id": "959295b148d6a8c71728ef6a3eaaef80",
    "span_id": "405b51fdec8afd00",
    "brand": "otel",
    "item": "telescope",
    "log": "stuff is broke"
  },
  "version": "v1"
}
```

Recipe: Parse JSON

```
merge_maps(cache, ParseJSON(body)["object"], "upsert")
```

cache: {}



```
cache:
  epoch_timestamp: 1730239218,
  trace_id: 959295b1,
  span_id: 405b51fd,
  name: my-app,
  env: prod,
  log: network error
```

Recipe: Parse JSON

```
set(body, cache["log"])
```

body: long json string

body: network error

Recipe: Work with Timestamps



KubeCon



CloudNativeCon

North America 2024

```
set(time, Time(cache["timestamp"], "%Y-%m-%dT%H:%M:%SZ"))
```


time: 0000000000



time: 1730240250

Recipe: Handle Trace/Span IDs

```
set(trace_id.string, cache["trace_id"])  
set(span_id.string, cache["span_id"])
```



Recipe: Manipulate Strings



KubeCon



CloudNativeCon

North America 2024

```
set(attributes["item-id"], ConvertCase(Concat([cache["brand"], cache["item"]], "-"), "upper"))
```

attributes: {}



attributes:
item-name: OTEL-TELESCOPE


```
log_statements:
- context: log
  statements:
    - merge_maps(cache, ParseJSON(body)["object"], "upsert")
    - set(body, cache["log"])
    - set(time, Time(cache["timestamp"], "%Y-%m-%dT%H:%M:%SZ"))
    - set(trace_id.string, cache["trace_id"])
    - set(span_id.string, cache["span_id"])
    - set(attributes["item-id"], ConvertCase(Concat([body["brand"], body["item"]], "-"), "upper"))
```

Use Case: Metrics to Stable SemConv

Need to coalesce different HTTP semantic conventions version

Recipe:

- Reuse conditions
- Rename an attribute
- Scaling a metric
- Rename a metric

```
name: http.client.duration
unit: ms
datapoints:
  - attributes:
      http.method: GET
      buckets: 5, 10, 25
```



```
name:
  http.client.request.duration
unit: s
datapoints:
  - attributes:
      http.request.method: GET
      buckets: 0.005, 0.01, 0.025
```

Recipe: Reuse Conditions



KubeCon



CloudNativeCon

North America 2024

```
metric_statements:
  - context: datapoint
    conditions:
      - metric.name == "http.client.duration"
      - metric.name == "http.server.duration"
    statements:
      - [...]
```

Recipe: Rename an Attribute



KubeCon



CloudNativeCon

North America 2024

```
set(attributes["http.request.method"], attributes["http.method"])  
delete_key(attributes, "http.method")
```

attributes:
http.method: GET



attributes:
http.request.method: GET

Recipe: Scale a Metric

`scale_metric(0.001, "s")` where `unit == "ms"`

`unit: ms`
`datapoints:`
 `buckets: 5, 10, 25`



`unit: s`
`datapoints:`
 `buckets: 0.005, 0.01, 0.025`

Recipe: Rename a Metric

```
set(name, "http.client.request.duration") where name == "http.client.duration"
```

```
set(name, "http.server.request.duration") where name == "http.server.duration"
```

name: http.client.duration



name: http.client.request.duration

```
metric_statements:
```

```
- context: datapoint
```

```
  conditions:
```

```
    - metric.name == "http.client.duration"
```

```
    - metric.name == "http.server.duration"
```

```
  statements:
```

```
    - set(attributes["http.request.method"], attributes["http.method"])
```

```
    - delete_key(attributes, "http.method")
```

```
- context: metric
```

```
  conditions:
```

```
    - name == "http.client.duration"
```

```
    - name == "http.server.duration"
```

```
  statements:
```

```
    - scale_metric(0.001, "s") where unit == "ms"
```

```
    - set(name, "http.client.request.duration") where name == "http.client.duration"
```

```
    - set(name, "http.server.request.duration") where name == "http.server.duration"
```

Use Case: Converting Prometheus

We are scraping the Prometheus endpoint of our message queue

Recipes:

- Dynamically rename a metric
- Aggregate a metric using the transform processor
- Restructure a metrics payload

Recipe: Dynamically Rename a Metric

Context: metric

```
replace_pattern(name, "mq_queue_(.+)", "mq.queue.$$1")
```

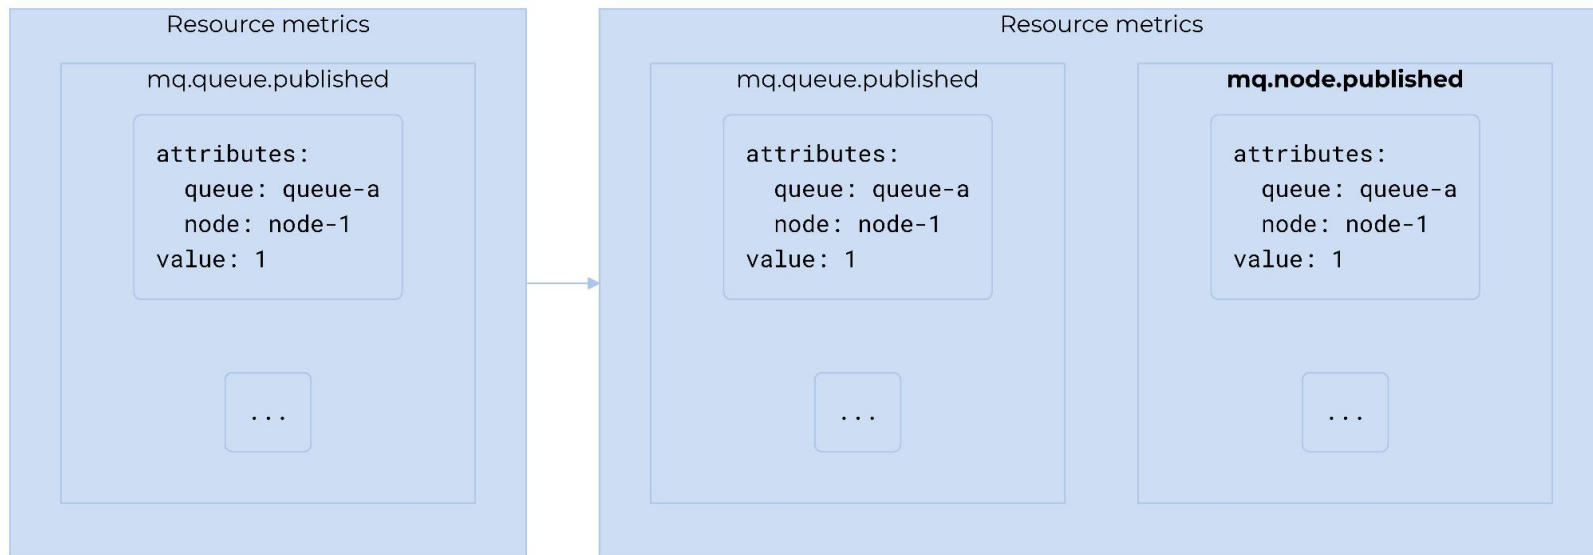
name: "mq_queue_published_total"

name: "mq.queue.published_total"

Recipe: Aggregate a Metric

Context: metric

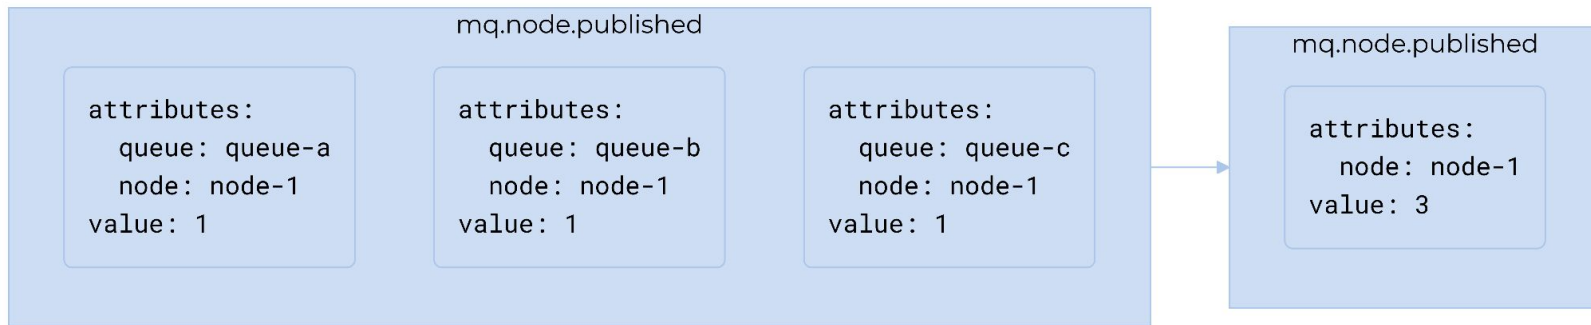
```
copy_metric(name="mq.node.published") where name == "mq.queue.published"
```



Recipe: Aggregate a Metric

Context: metric

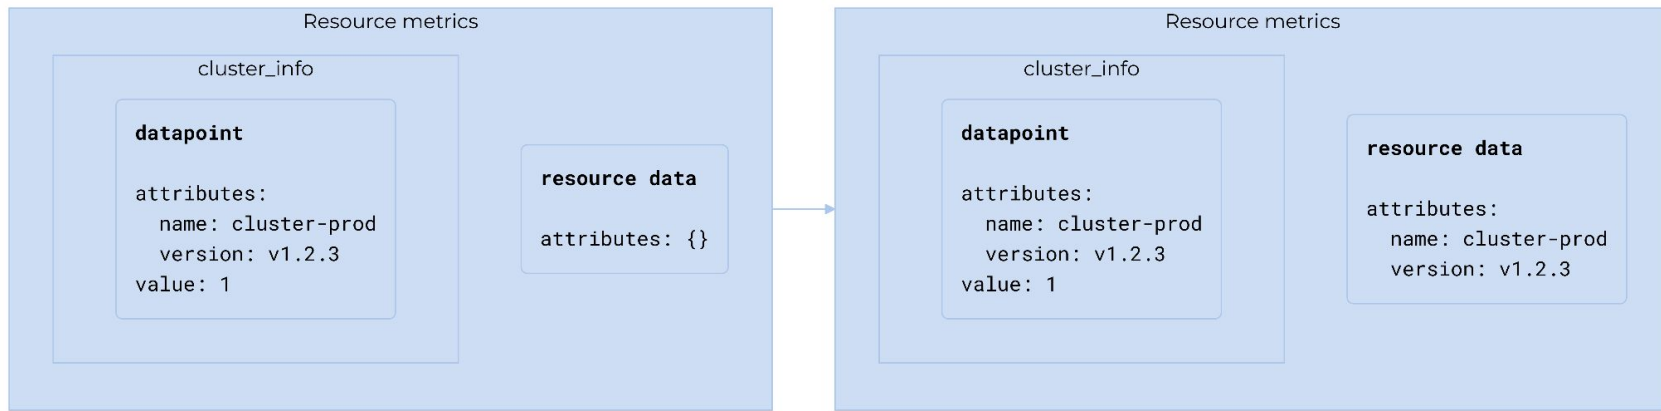
`aggregate_on_attributes("sum", ["node"]) where name == "mq.node.published"`



Recipe: Restructure a Metrics Payload

Context: datapoint

```
merge_maps(resource.attributes, attributes, "upsert") where meric.name == "mq_cluster_info"
```



Recipe: Restructure a Metrics Payload



KubeCon



CloudNativeCon

North America 2024

```
filter:
  metrics:
    metric:
      - name == "mq_cluster_info"
```

transform:

metric_statements:

- context: metric

statements:

- replace_pattern(name, "mq_queue_(.+)", "mq.queue.\$\$1")

- copy_metric(name="mq.node.publish") where name == "mq.queue.published_total"

- aggregate_on_attributes("sum", ["node"]) where name == "mq.node.publish"

- context: datapoint

statements:

- merge_maps(resource.attributes, attributes, "upsert") where metric.name == "mq_cluster_info"

filter:

error_mode: ignore

metrics:

metric:

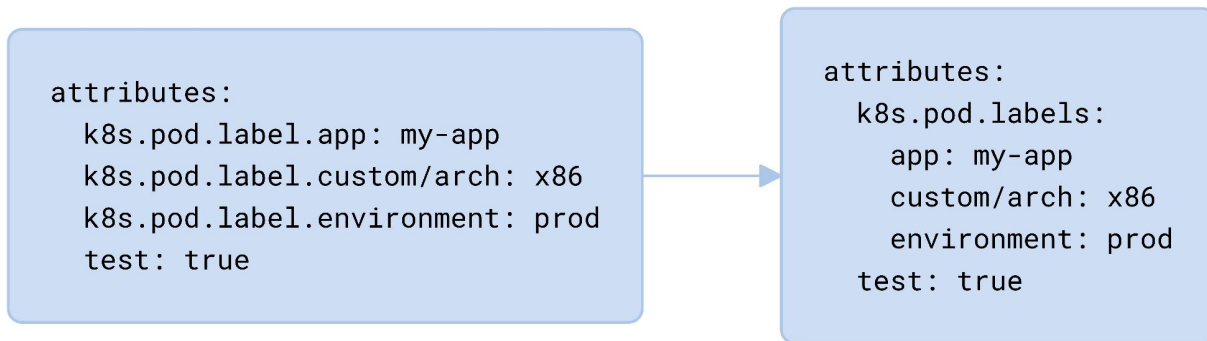
- name == "mq_cluster_info"

Use Case: Clean Up Span Attributes

Our spans have some attributes that all share the same known prefix.

Recipe:

- Grouping related attributes



Recipe: Group Related Attributes

`set(cache, attributes)`

```
attributes:  
  k8s.pod.label.app: my-app  
  k8s.pod.label.custom/arch: x86  
  k8s.pod.label.environment: prod  
  test: true  
cache:
```



```
attributes:  
  k8s.pod.label.app: my-app  
  k8s.pod.label.custom/arch: x86  
  k8s.pod.label.environment: prod  
  test: true  
cache:  
  k8s.pod.label.app: my-app  
  k8s.pod.label.custom/arch: x86  
  k8s.pod.label.environment: prod  
  test: true
```


Recipe: Group Related Attributes

```
delete_matching_keys(attributes, "^k8s\\.pod\\.label\\.")  
keep_matching_keys(cache, "^k8s\\.pod\\.label\\.")
```

```
attributes:  
  k8s.pod.label.app: my-app  
  k8s.pod.label.custom/arch: x86  
  k8s.pod.label.environment: prod  
  test: true  
cache:  
  k8s.pod.label.app: my-app  
  k8s.pod.label.custom/arch: x86  
  k8s.pod.label.environment: prod  
  test: true
```



```
attributes:  
  test: true  
cache:  
  k8s.pod.label.app: my-app  
  k8s.pod.label.custom/arch: x86  
  k8s.pod.label.environment: prod
```

Recipe: Group Related Attributes

```
replace_all_patterns(cache, "key", "^k8s\\.pod\\.label\\.\"", "")
```

cache:

```
k8s.pod.label.app: my-app  
k8s.pod.label.custom/arch: x86  
k8s.pod.label.environment: prod
```

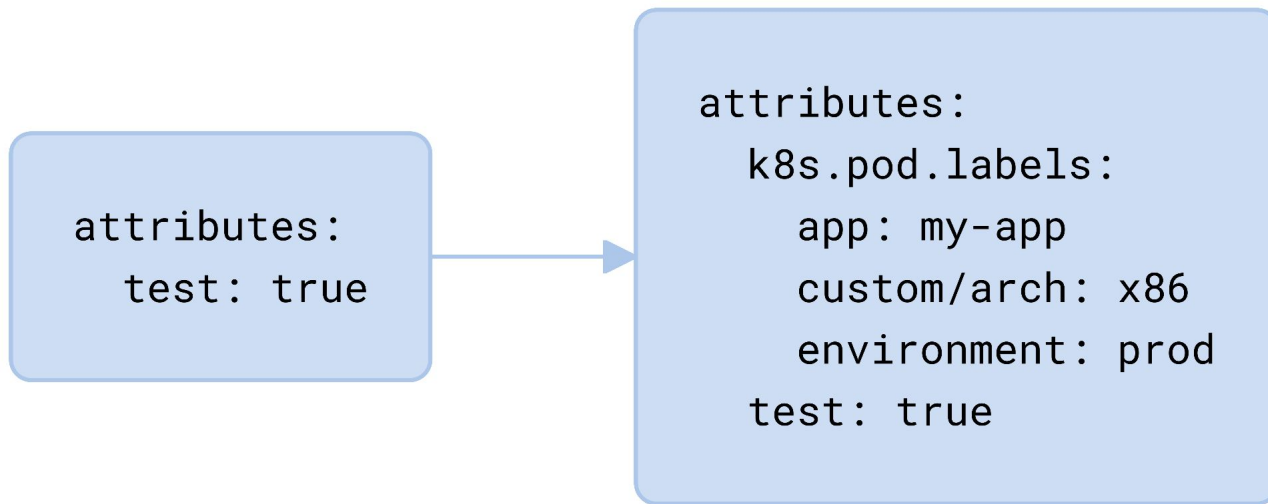


cache:

```
app: my-app  
custom/arch: x86  
environment: prod
```

Recipe: Group Related Attributes

```
set(attributes["k8s.pod.labels"], cache)
```



```
trace_statements:
```

```
- context: resource
```

```
  statements:
```

- set(cache, attributes)
- delete_matching_keys(attributes, "^k8s\\.pod\\.label\\.")
- keep_matching_keys(cache, "^k8s\\.pod\\.label\\.")
- replace_all_patterns(cache, "key", "^k8s\\.pod\\.label\\.\"", "")
- set(attributes["k8s.pod.labels"], cache)

If you have a use case in mind we'll try to solve it using OTTL

Restrictions:

- No stateful transformations
- No cross-signal transformations
- No Profile transformations - OTTL can't handle them yet.

We reserve the right to say “interesting problem, I need more time to solve that”

We don't need the “why” only the “what”.

Suggested templates for suggesting a scenario:


- **I need my data to** <blank> **instead of** <blank>
- **I need to drop data when** <blank>
- **I need to** <blank> **using** <blank>

Statements will be uploaded to the session page.

Want to Know More?

OTTL

<https://bit.ly/ottl-docs>



Transform processor

<https://bit.ly/tp-docs>

