



KubeCon



CloudNativeCon

North America 2024

# ***Using OpenTelemetry for Deep Observability Within Messaging Queues***

Shivanshu Raj Shrivastava, SigNoz  
Ekansh Gupta, SigNoz



KubeCon



CloudNativeCon

North America 2024

# About the speakers



Ekansh Gupta

Software Engineer, SigNoz (YC  
W21)



Shivanshu Raj Shrivastava

Founding Engineer SigNoz (YC W21)  
CNCF ambassador  
Member & contributor opentelemetry



KubeCon



CloudNativeCon

North America 2024

# Talk Outline



SigNoz

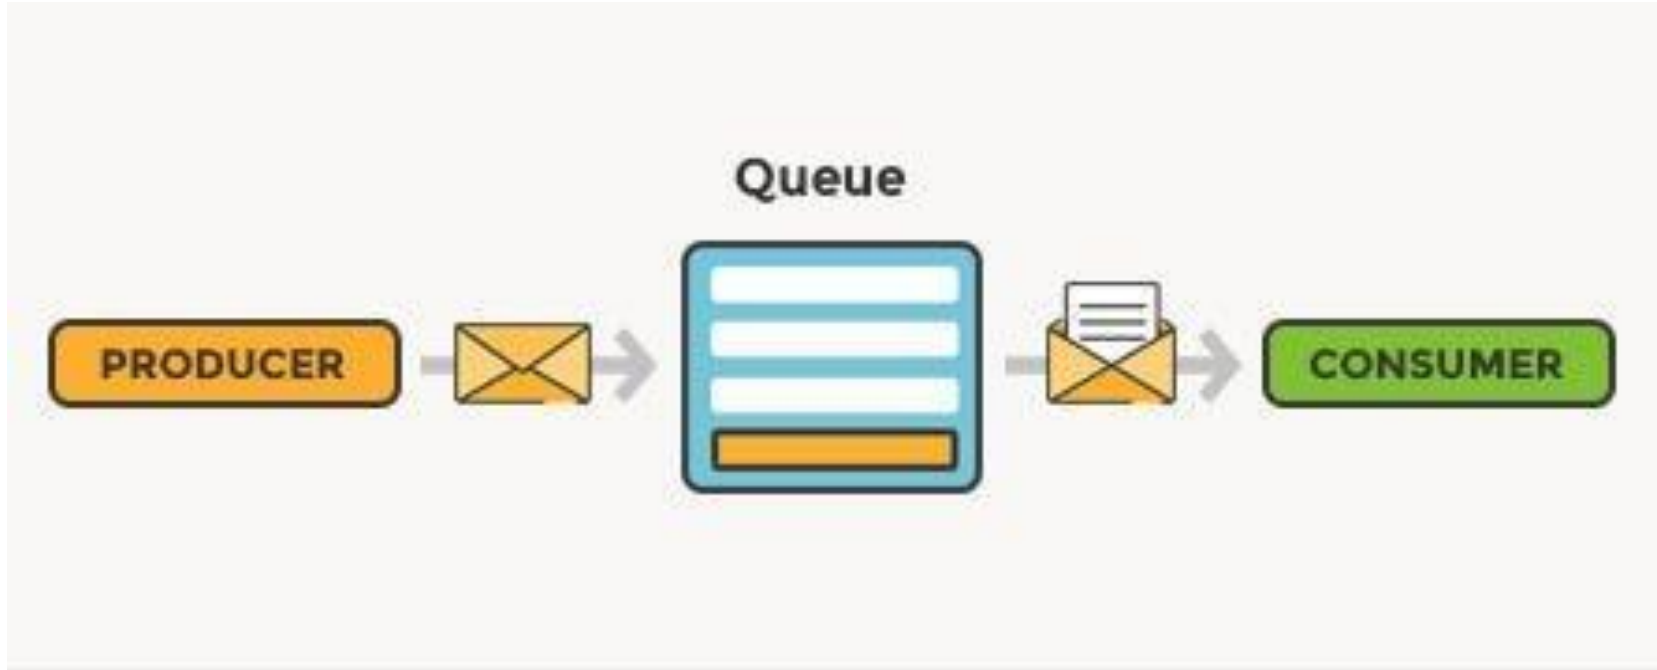
- **What is Messaging Queue and where all we use it?**
- **What are the common problems that we face while using Messaging Queues such as Kafka / RabbitMQ**
- **Messaging queues are complex**
- **Lets deep dive into Kafka**
- **What we don't see with Kafka Metrics**
- **The all powerful OTel**
- **Correlations and deeper insights into Messaging Queues.**

# **What are Messaging Queues?**

## **And**

# **Where do we use it?**

# What are MQs and where do we use it?



General Understanding of a Messaging Queue

Source: CloudAMQP

# What are MQs and where do we use it?




KubeCon



CloudNativeCon

North America 2024



# What are MQs and where do we use it?

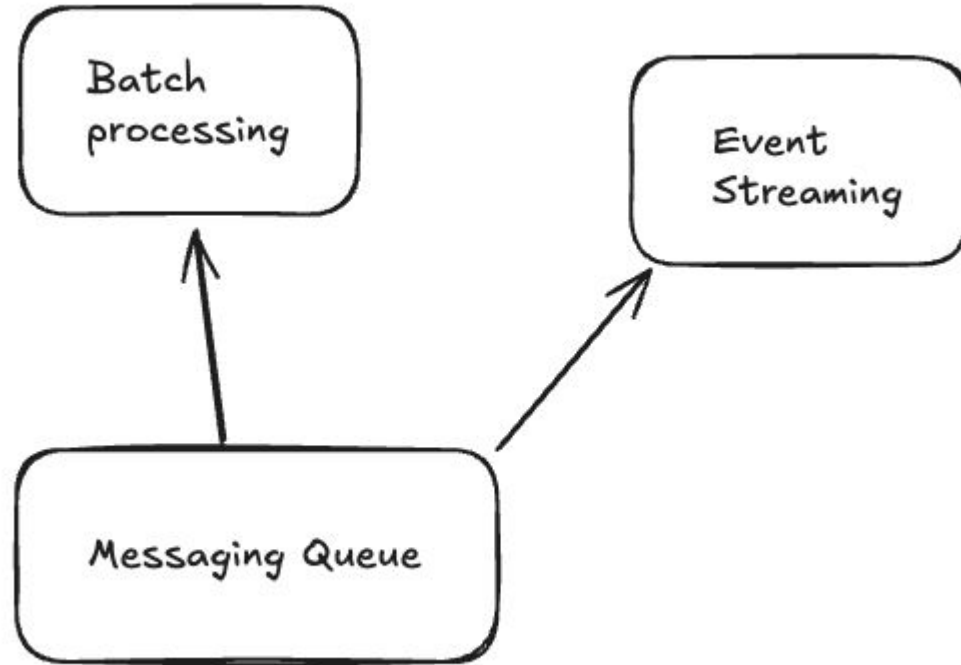


KubeCon



CloudNativeCon

North America 2024





# What are MQs and where do we use it?

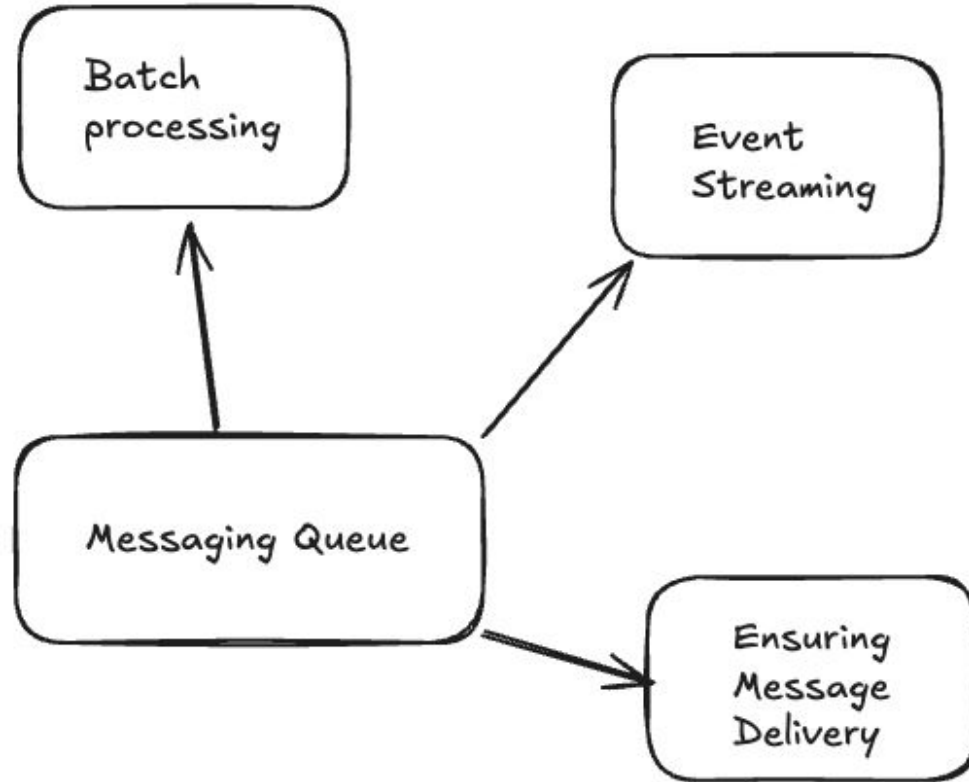


KubeCon

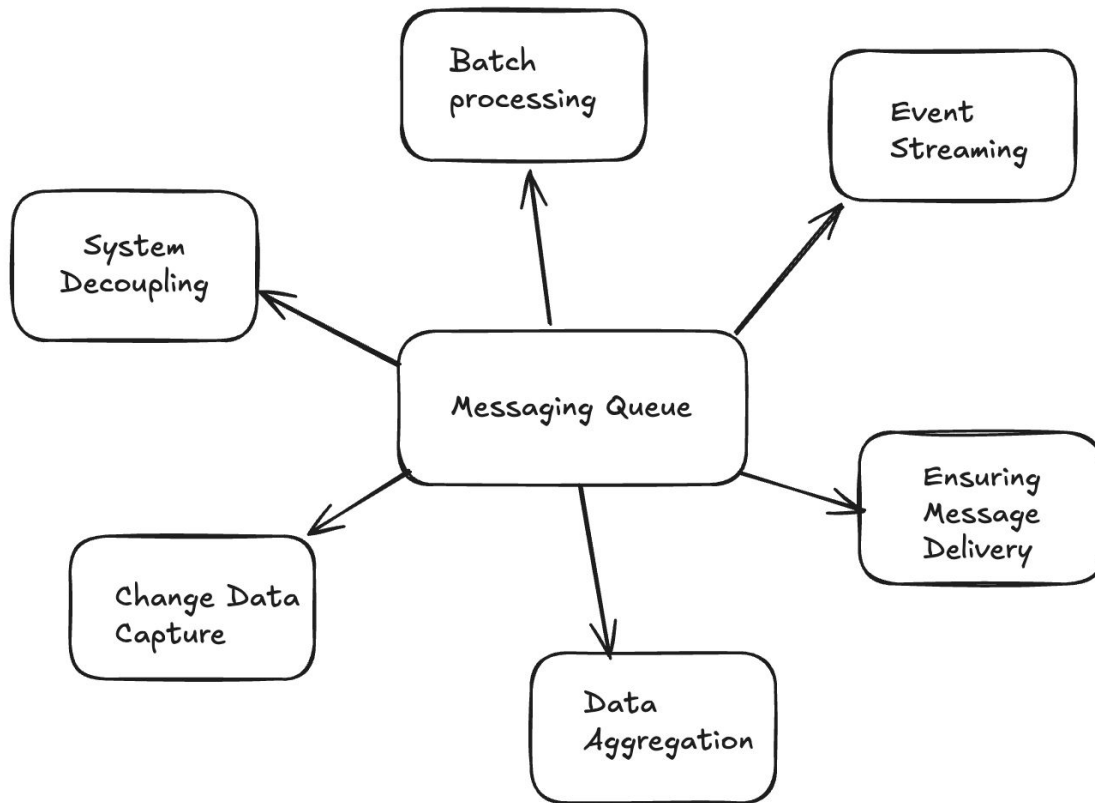


CloudNativeCon

North America 2024



# What are MQs and where do we use it?



# Messaging Queues:

## What can Actually go wrong?

# What can go Wrong?

**High Latency and  
Throughput Issues**

**No Visibility about  
Latency**

**Throttling Issues**

**Not Enough Metrics**

**Broker Failures**

**Consumer Lag and  
Back Pressure**

**Message Ordering Issues**

**Dead Letter Queues**



KubeCon



CloudNativeCon

North America 2024



imgflip.com



SigNoz



KubeCon



CloudNativeCon

North America 2024

# Messaging Queues are Complex



SigNoz

Developers often use **default** configurations,  
while fine tuning with **deeper observability**  
can help them **scale better**

Knowing **when** and **how** things are  
breaking, can **help in fixing this faster** :)



KubeCon



CloudNativeCon



North America 2024

# Let's Dive deeper into Messaging Queue



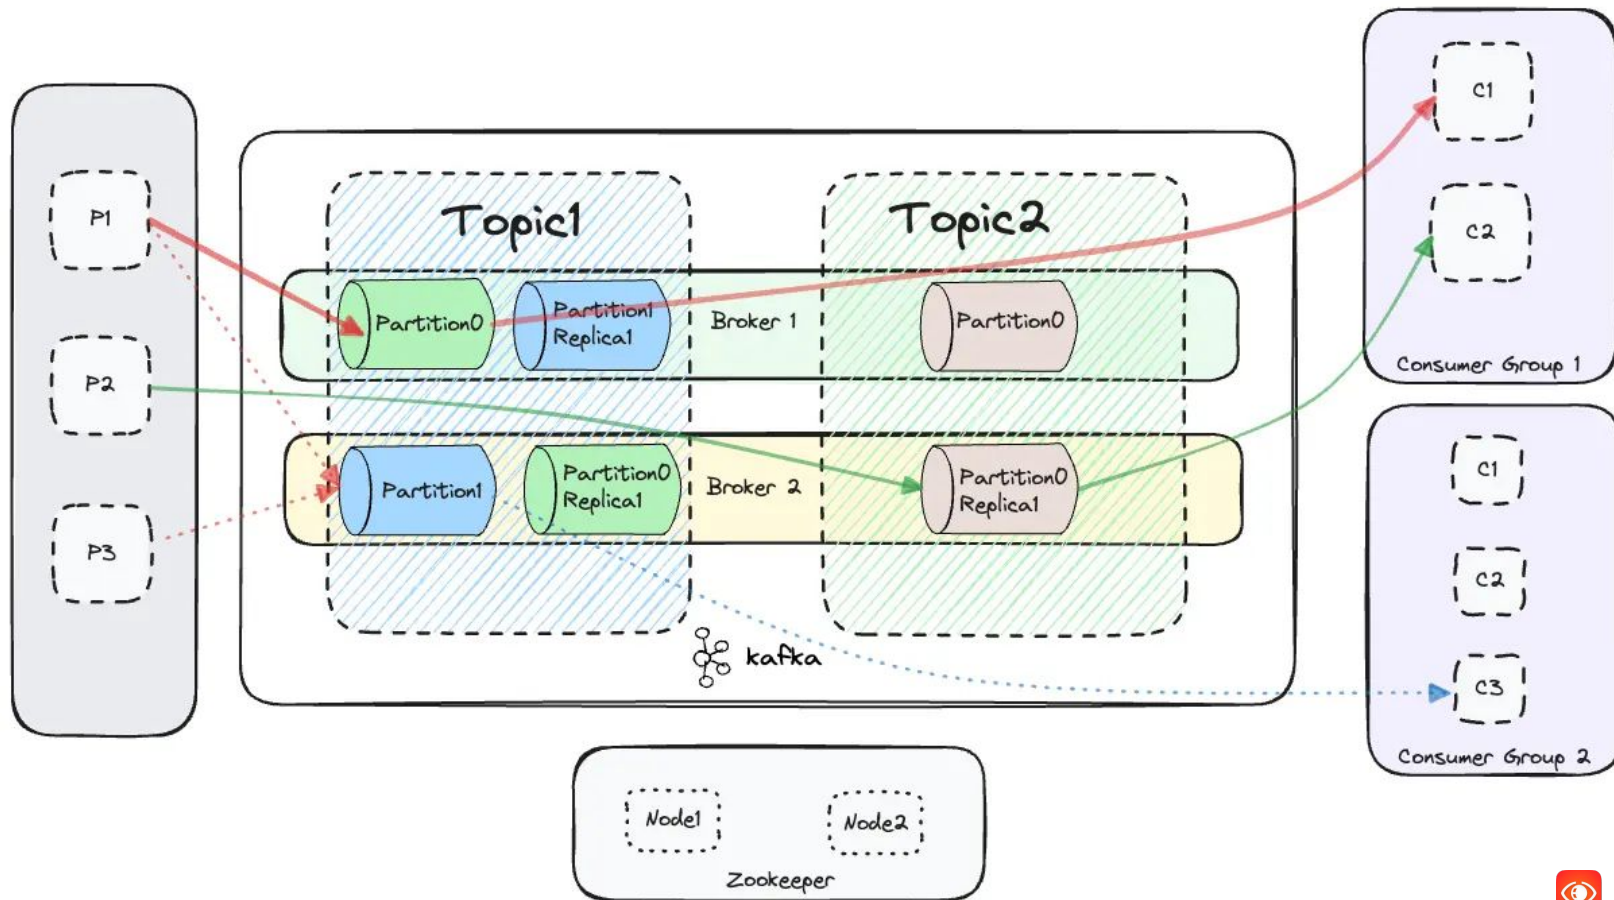
SigNoz

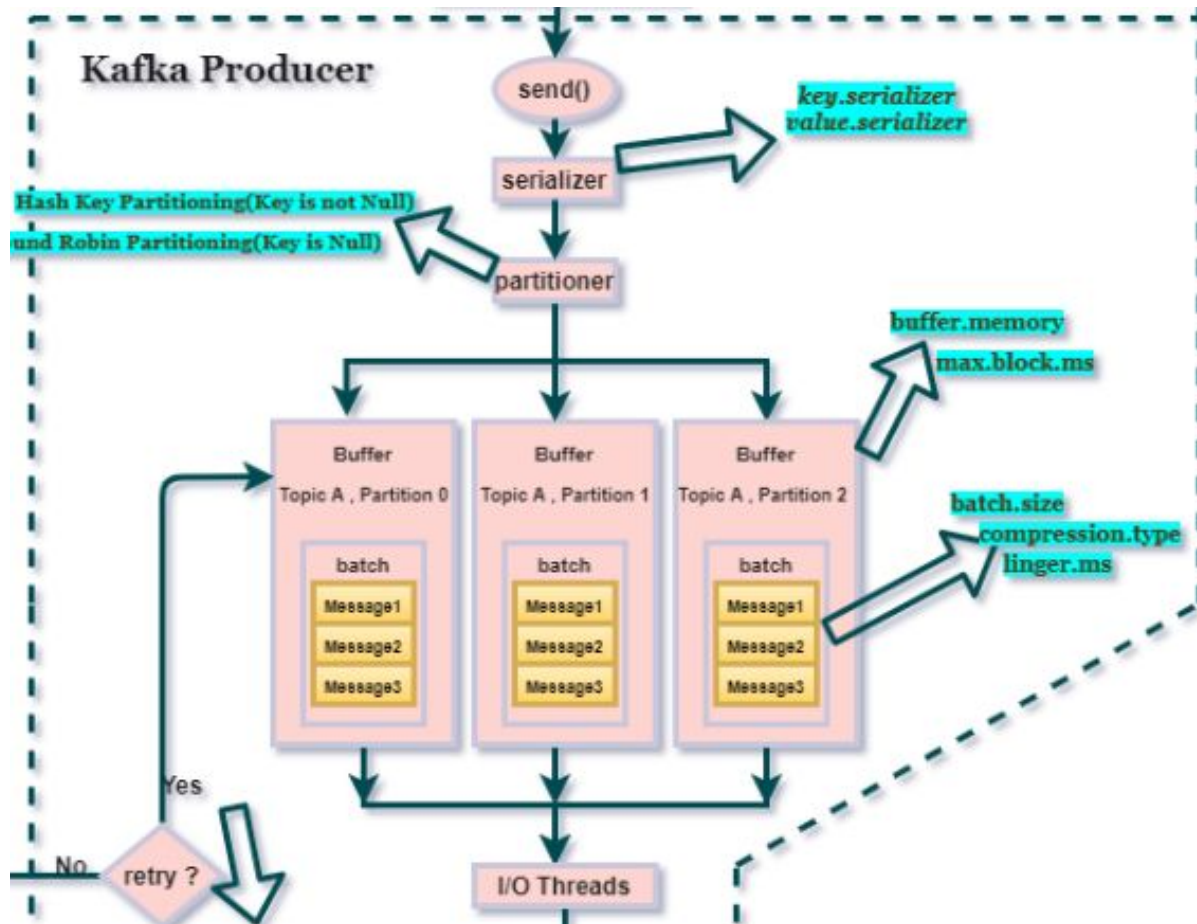


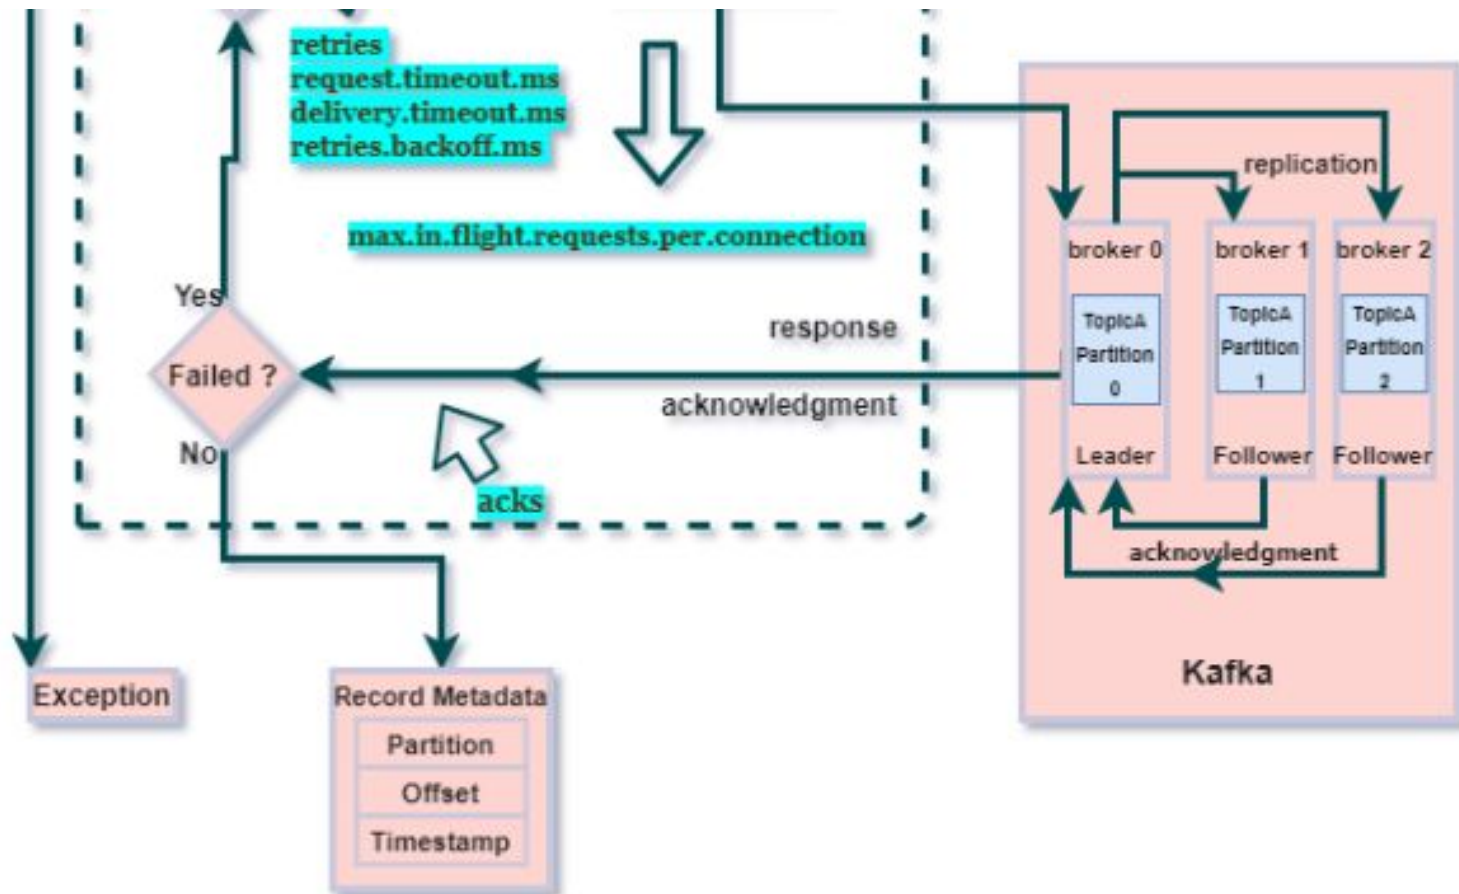


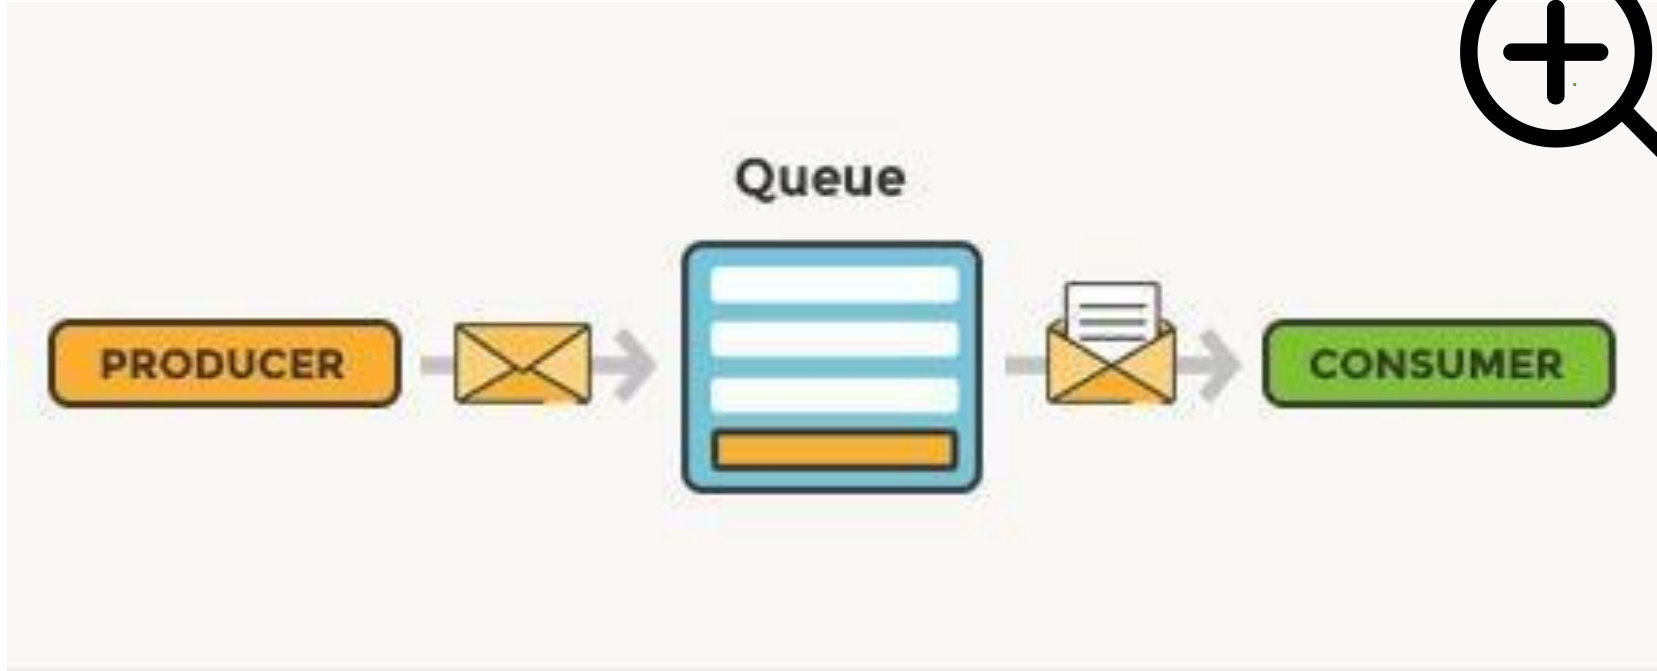

General Understanding of a Messaging Queue

Source: CloudAMQP





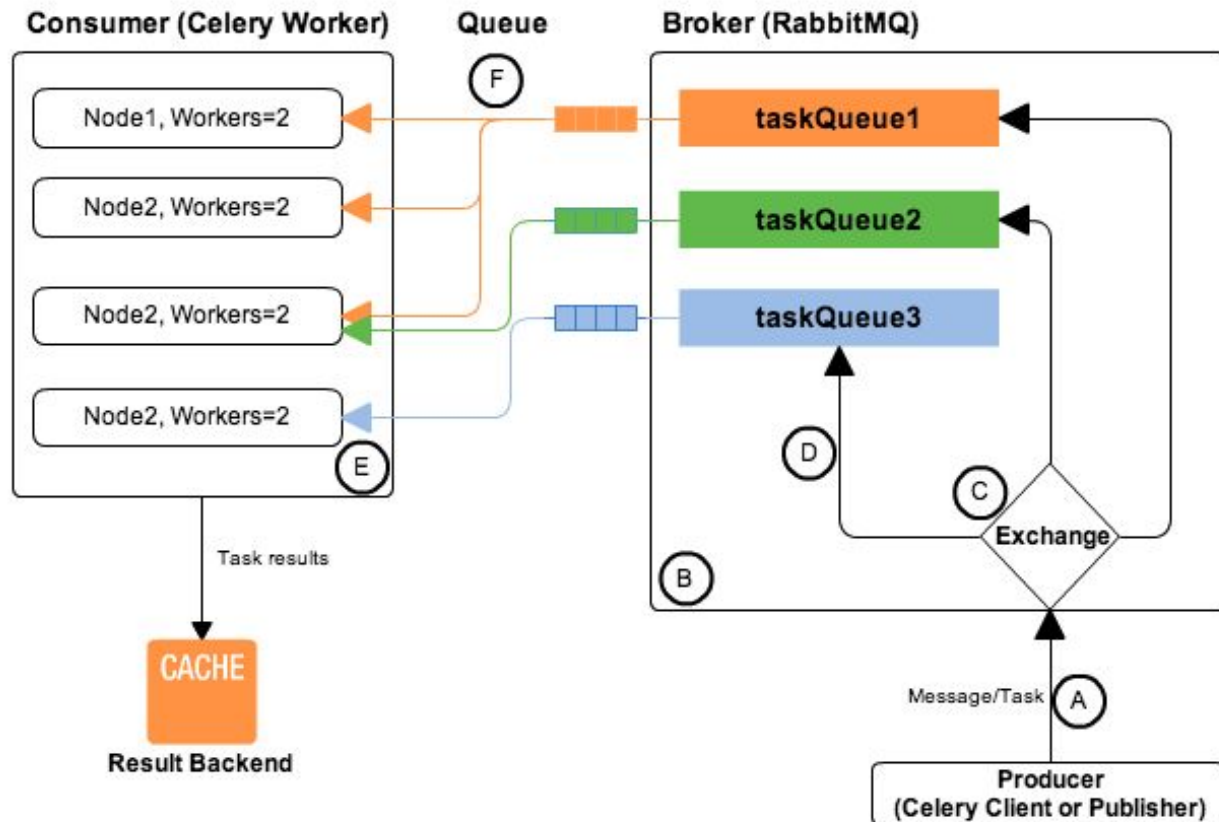




General Understanding of a Messaging Queue

Source: CloudAMQP

# RabbitMQ + Celery





KubeCon



CloudNativeCon

North America 2024

# So what we don't see with Broker Metrics?



SigNoz

- **Kafka does not directly expose producer-side latency metrics per partition.**
  - No one ever knows why there is high producer latency.
  - There is no visibility if there's a problem with Kafka itself or producer.
  - Difficulty tracking producer throughput across different topics
  - Complex to monitor batch sizes and compression ratios
  - Complex to find the correlation between different partitions



- **End-to-End Latency (How much time it took for producer to send a message and a consumer to acknowledge it)**
  - Difficult to track processing time per message
  - No direct correlation between consumer performance and resource utilization
  - Complex to correlate broker metrics with client-side issues

- **Consumer Group Lag correlation with spans**
  - No visibility of which partition is causing the lag?
  - Is the lag caused by consumer application, network problems, or imbalances between producer and consumer performance.
  - Manual implementation needed for trace context propagation.
  - Limited visibility into consumer group rebalancing events

- **Infrastructure Monitoring Gaps**

- Complex relationship between JVM metrics and Kafka performance
- Difficult to correlate network issues with message delivery problems
- Limited visibility into disk I/O patterns
- Difficult to track partition growth patterns

**Metrics -> Trace**

**Traces -> Logs**

**Traces -> Metrics**

**All these correlations are hard  
to achieve natively in messaging queues**



KubeCon



CloudNativeCon

North America 2024

# What is the answer to this Problem?



**SigNoz**



KubeCon



CloudNativeCon

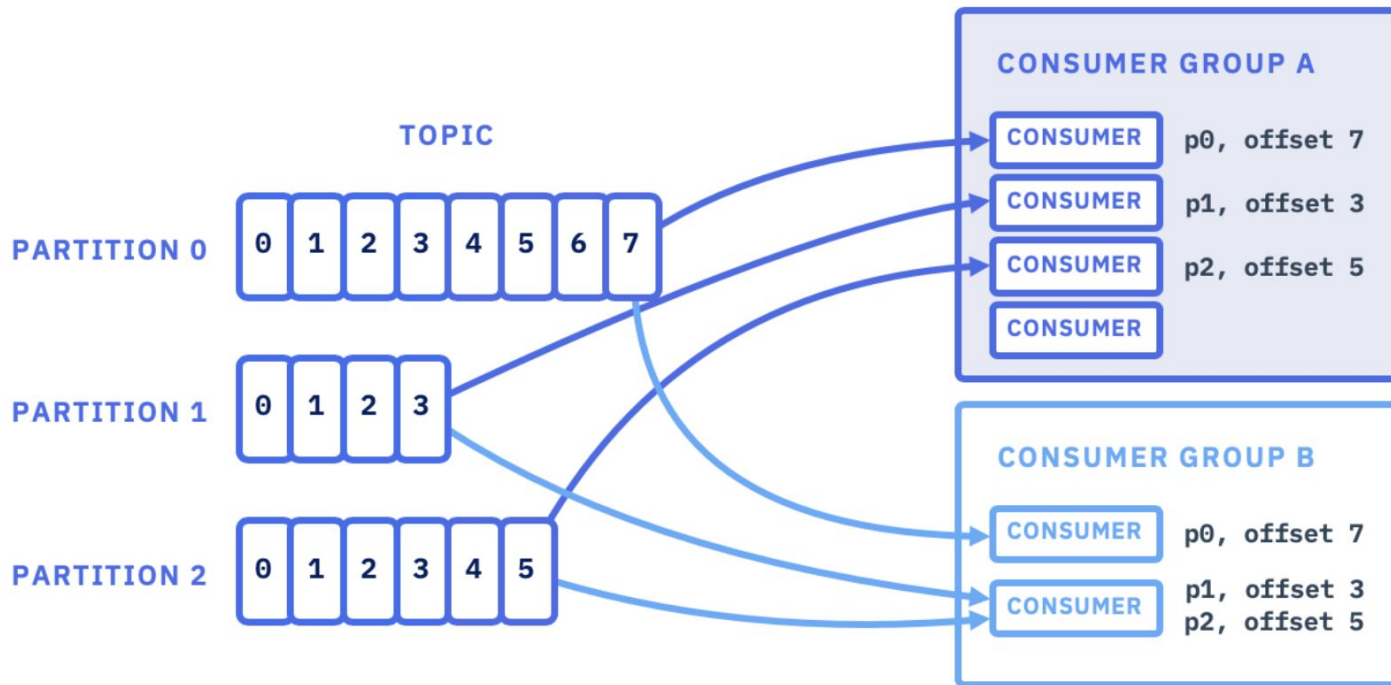
North America 2024



SigNoz

- **Consumer Group Lag correlation with spans**
  - No visibility of which partition is causing the lag?
  - Is the lag caused by consumer application, network problems, or imbalances between producer and consumer performance.
  - Manual implementation needed for trace context propagation.
  - Limited visibility into consumer group rebalancing events

# Common problems







KubeCon




CloudNativeCon

North America 2024

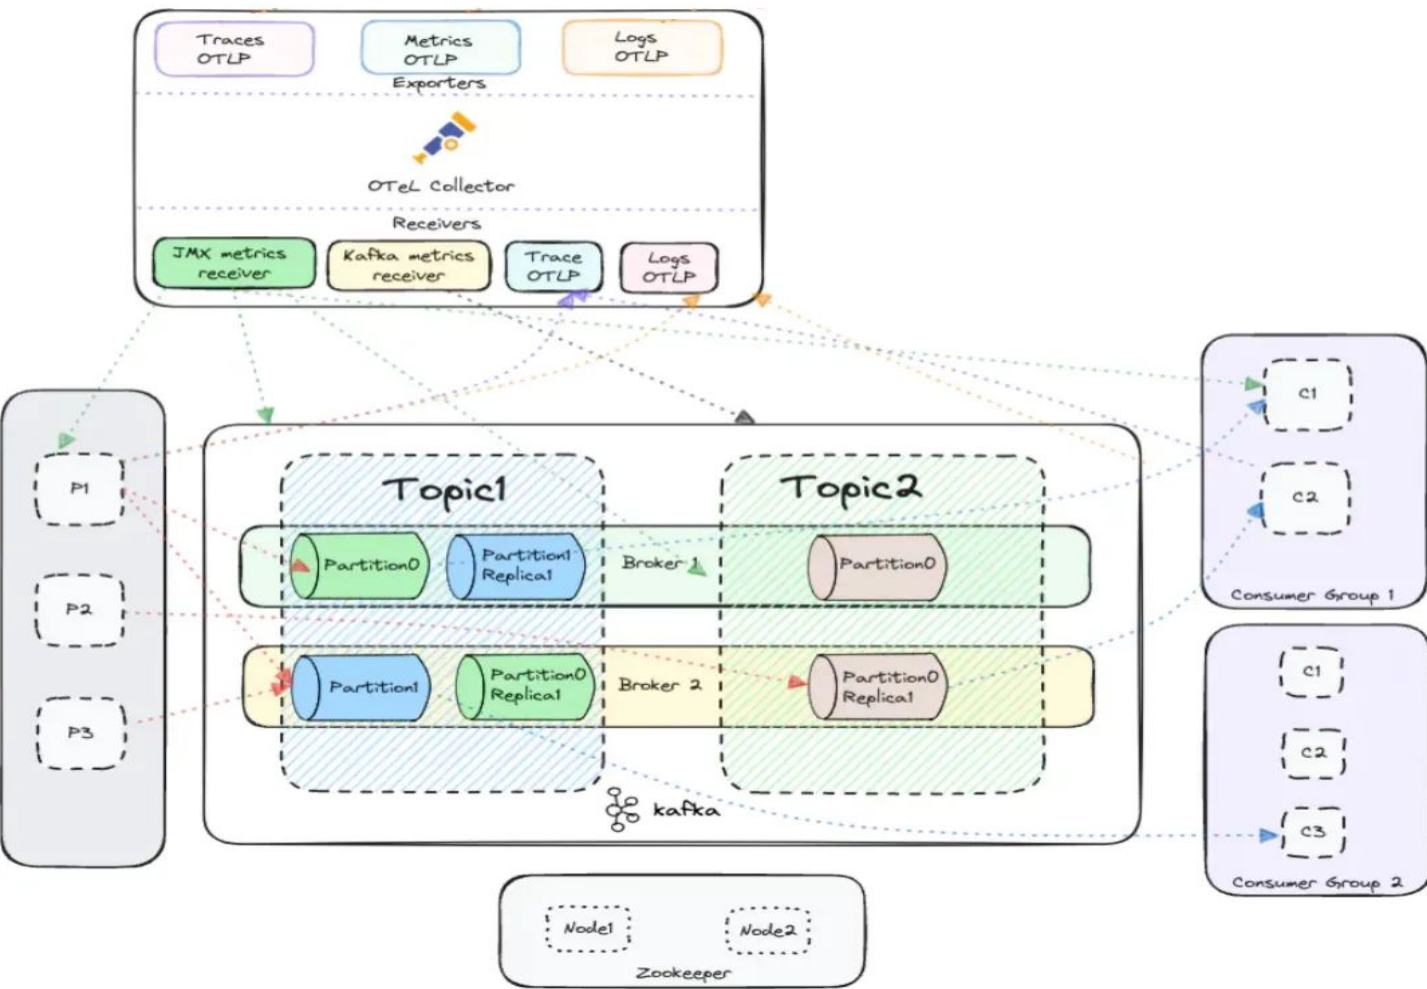
# Noisy Neighbour



SigNoz



- **Stable API** with ability to export telemetry anywhere
- **Instrumentations** for popular clients and frameworks
- **Semantic conventions** to unify telemetry format



```
receivers:
  otlp:
    protocols:
      grpc:
        # OTeL receiver endpoint (grpc)
        endpoint: 127.0.0.1:4317
      http:
        # OTeL receiver endpoint (http)
        endpoint: 127.0.0.1:4318
  kafkametrics:
    brokers:
      # Kafka brokers endpoints
      - 127.0.0.1:9092
      - 127.0.0.1:9093
    protocol_version: 2.0.0
  scrapers:
    - brokers
    - topics
    - consumers
```

jmx:

# configure the path where you installed opentelemetry-jmx-metrics receiver

jar\_path: \${PWD}/opentelemetry-jmx-metrics.jar

endpoint: service:jmx:rmi:///jndi/rmi://127.0.0.1:2020/jmxrmi

target\_system: jvm,kafka,kafka-consumer,kafka-producer

collection\_interval: 10s

log\_level: info

resource\_attributes:

broker.name: broker1

jmx/2:

# configure the path where you installed opentelemetry-jmx-metrics receiver

jar\_path: \${PWD}/opentelemetry-jmx-metrics.jar

endpoint: service:jmx:rmi:///jndi/rmi://127.0.0.1:2021/jmxrmi

target\_system: jvm,kafka,kafka-consumer,kafka-producer

collection\_interval: 10s

log\_level: info

resource\_attributes:

broker.name: broker2

```
exporters:
  otlp:
    # configure the grpc endpoint for signoz otel collector
    endpoint: "127.0.0.1:65118"
    tls:
      insecure: true
  debug:
    verbosity: detailed

service:
  pipelines:
    traces:
      receivers: [otlp]
      exporters: [otlp]
    logs:
      receivers: [otlp]
      exporters: [otlp]
    metrics:
      receivers: [otlp, kafkametrics]
      exporters: [otlp, debug]
```

# Distributed Tracing

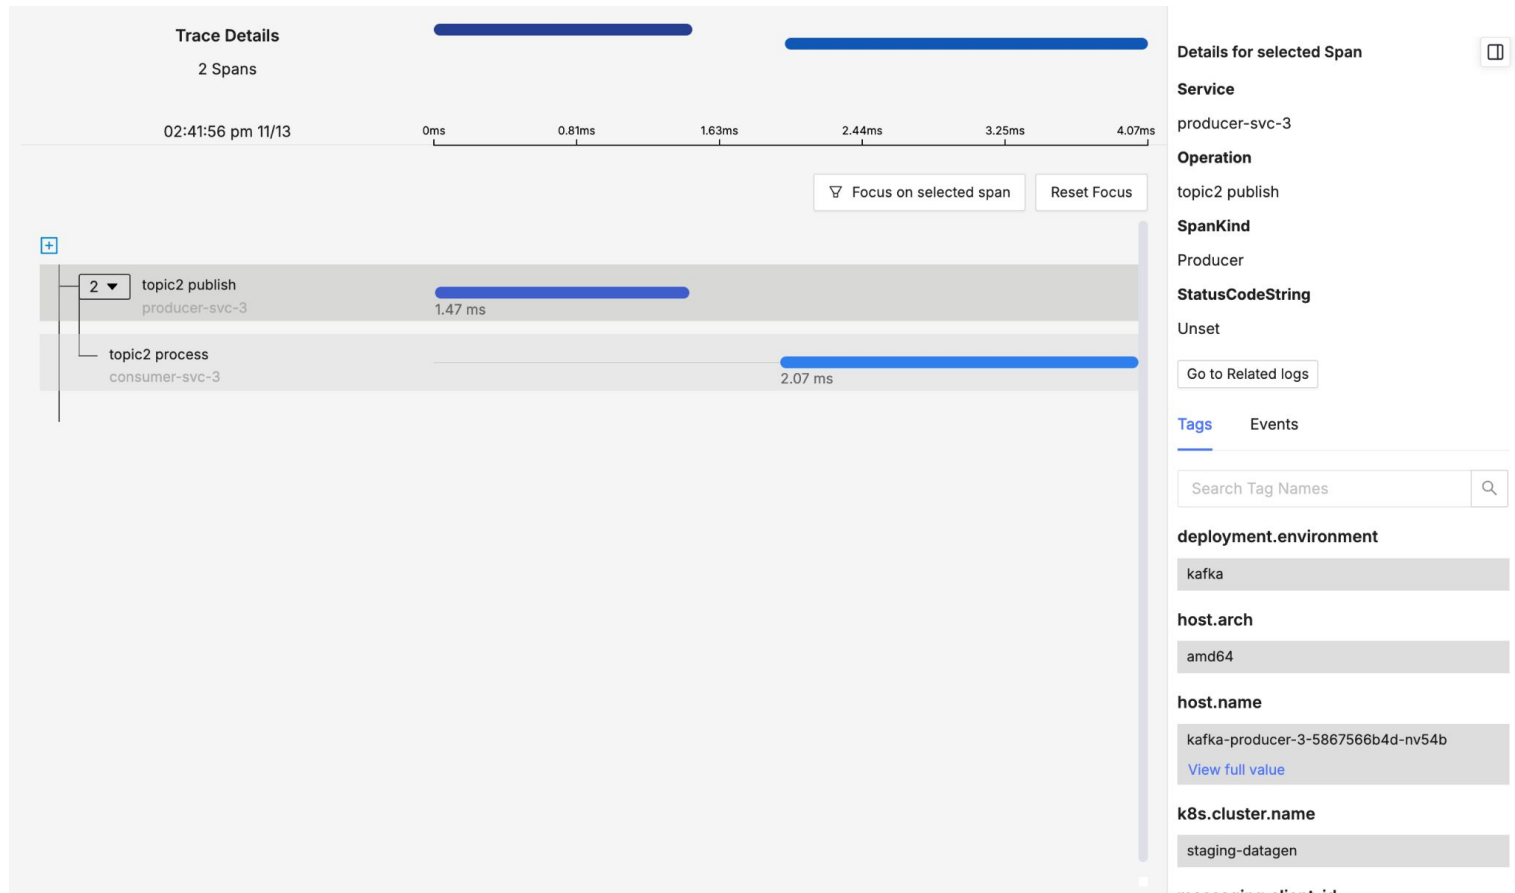


KubeCon



CloudNativeCon

North America 2024



SigNoz



```
InstrumentationScope io.opentelemetry.kafka-clients-0.11 2.4.0-alpha
```

```
Span #0
```

```
Trace ID      : d373467a29492793345307e726fb1404
Parent ID     : ee72a4415b8d6574
ID            : 9acfab51e1b25826
Name          : topic1 process
Kind          : Consumer
Start time    : 2024-05-21 15:16:47.166954 +0000 UTC
End time      : 2024-05-21 15:16:47.167983166 +0000 UTC
Status code   : Unset
Status message :
```

```
Attributes:
```

```
-> messaging.kafka.message.offset: Int(4)
-> messaging.kafka.consumer.group: Str(my-consumer-group)
-> messaging.kafka.message.key: Str(MY_KEY)
-> messaging.system: Str(kafka)
-> messaging.destination.partition.id: Str(0)
-> kafka.record.queue_time_ms: Int(30)
-> messaging.destination.name: Str(topic1)
-> messaging.operation: Str(process)
-> thread.id: Int(39)
-> messaging.client_id: Str(consumer-my-consumer-group-1)
-> thread.name: Str(Thread-2)
-> messaging.message.body.size: Int(10)
```

# Kafka Client-Native Metrics

Name	Instrument Type	Unit (UCUM)	Description
<code>messaging.client.operation.duration</code>	Histogram	s	Duration of messaging operation initiated by a producer or consumer client.

Attribute	Type	Description	Examples
<code>error.type</code>	string	Describes a class of error the operation ended with.	<code>amqp:decode-error</code> ; <code>KAFKA_STORAGE_ERROR</code> ; <code>channel-error</code>
<code>messaging.consumer.group.name</code>	string	The name of the consumer group with which a consumer is associated.	<code>my-group</code> ; <code>indexer</code>
<code>messaging.destination.name</code>	string	The message destination name.	<code>MyQueue</code> ; <code>MyTopic</code>
<code>messaging.destination.partition.id</code>	string	The identifier of the partition messages are sent to or received from, unique within the <code>messaging.destination.name</code> .	<code>1</code>
<code>messaging.operation.name</code>	string	The system-specific name of the messaging operation.	<code>send</code> ; <code>receive</code> ; <code>ack</code>
<code>messaging.operation.type</code>	string	A string identifying the type of the messaging operation.	<code>create</code> ; <code>send</code> ; <code>receive</code>
<code>messaging.system</code>	string	The messaging system as identified by the client instrumentation.	<code>activemq</code> ; <code>aws_sqs</code> ; <code>eventgrid</code>
<code>server.address</code>	string	Server domain name if available without reverse DNS lookup; otherwise, IP address or Unix domain socket name.	<code>example.com</code> ; <code>10.1.2.80</code> ; <code>/tmp/my.sock</code>
<code>server.port</code>	int	Server port number.	<code>80</code> ; <code>8080</code> ; <code>443</code>

- `messaging.operation.name`
- `messaging.system`
- `messaging.batch.message_count`
- `messaging.consumer.group.name`
- `messaging.destination.anonymous`
- `messaging.destination.name`
- `messaging.destination.subscription.name`
- `messaging.destination.template`
- `messaging.destination.temporary`
- `messaging.operation.type`
- `messaging.client.id`
- `messaging.destination.partition.id`
- `messaging.message.conversation_id`
- `messaging.message.id`
- `messaging.message.body.size`
- `messaging.message.envelope.size`

- Correlation with metrics based on Span Attributes
- See the time a message spent in kafka
- Can provide deeper visibility into your async systems

# How to collect broker metrics collection?

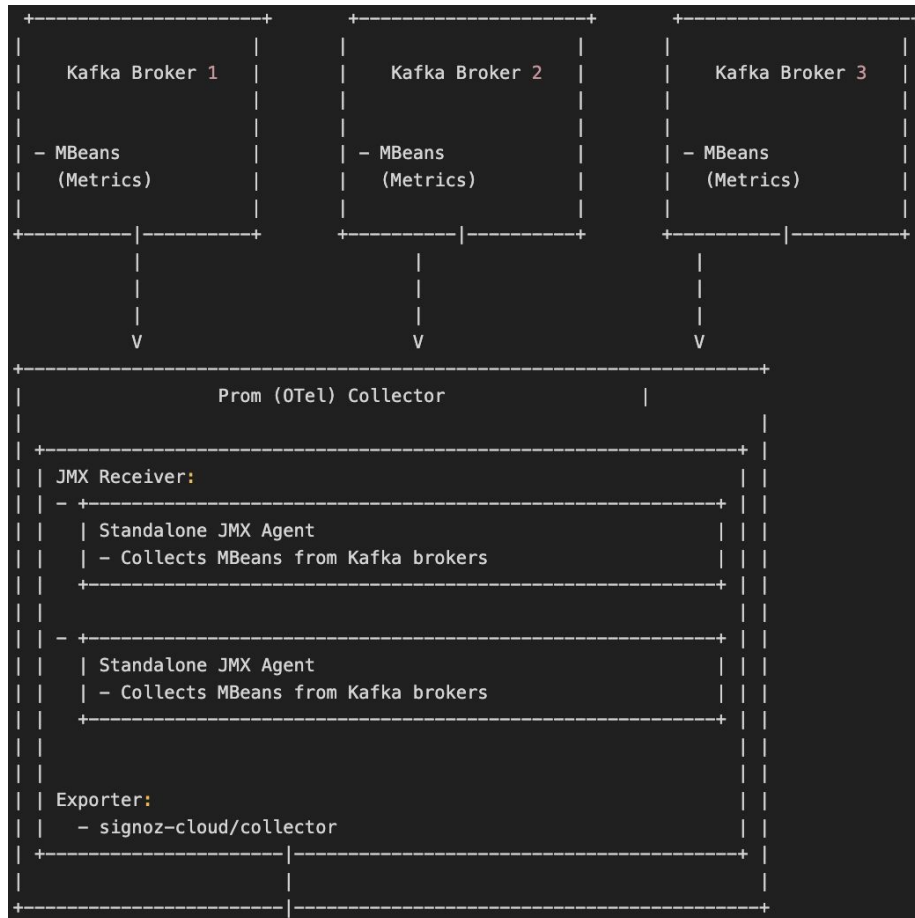


KubeCon



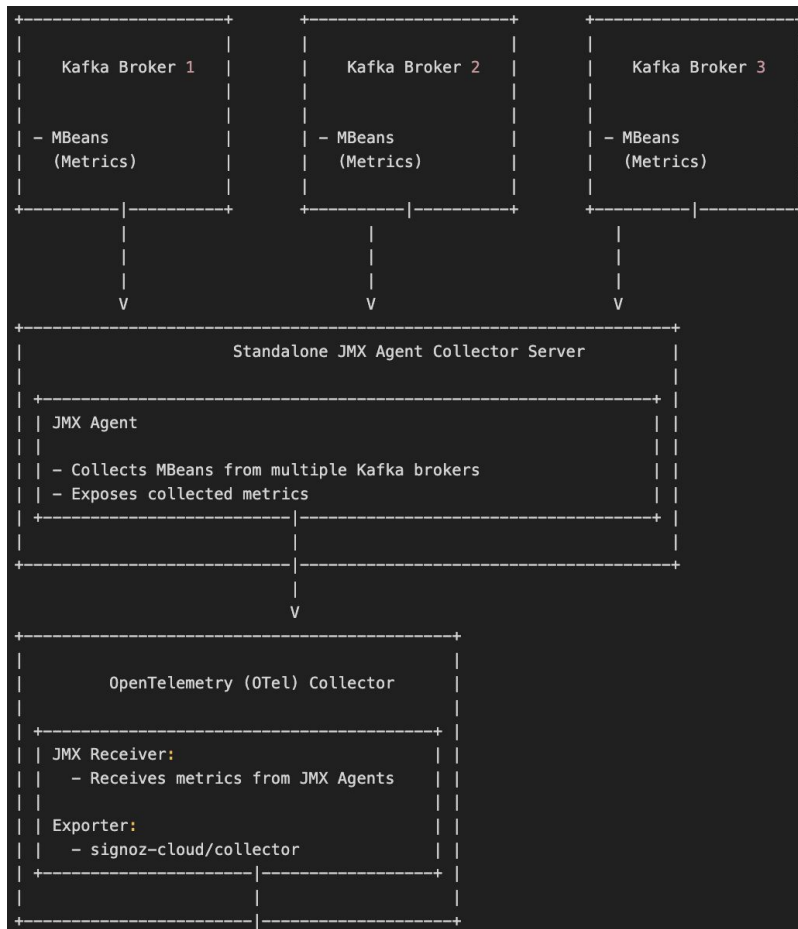
CloudNativeCon

North America 2024



MBean
ObjectName: kafka.server:type=BrokerTopicMetrics, name=MessagesInPerSec
Attributes
<ul style="list-style-type: none"><li>Count: 123456</li><li>MeanRate: 456.78</li><li>OneMinuteRate: 78.90</li><li>FiveMinuteRate: 67.89</li><li>FifteenMinuteRate: 56.78</li></ul>
Operations
<ul style="list-style-type: none"><li>resetStatistics(): void</li><li>sampleStats(): CompositeData</li></ul>

# How to collect broker metrics collection?



# How to collect broker metrics collection?

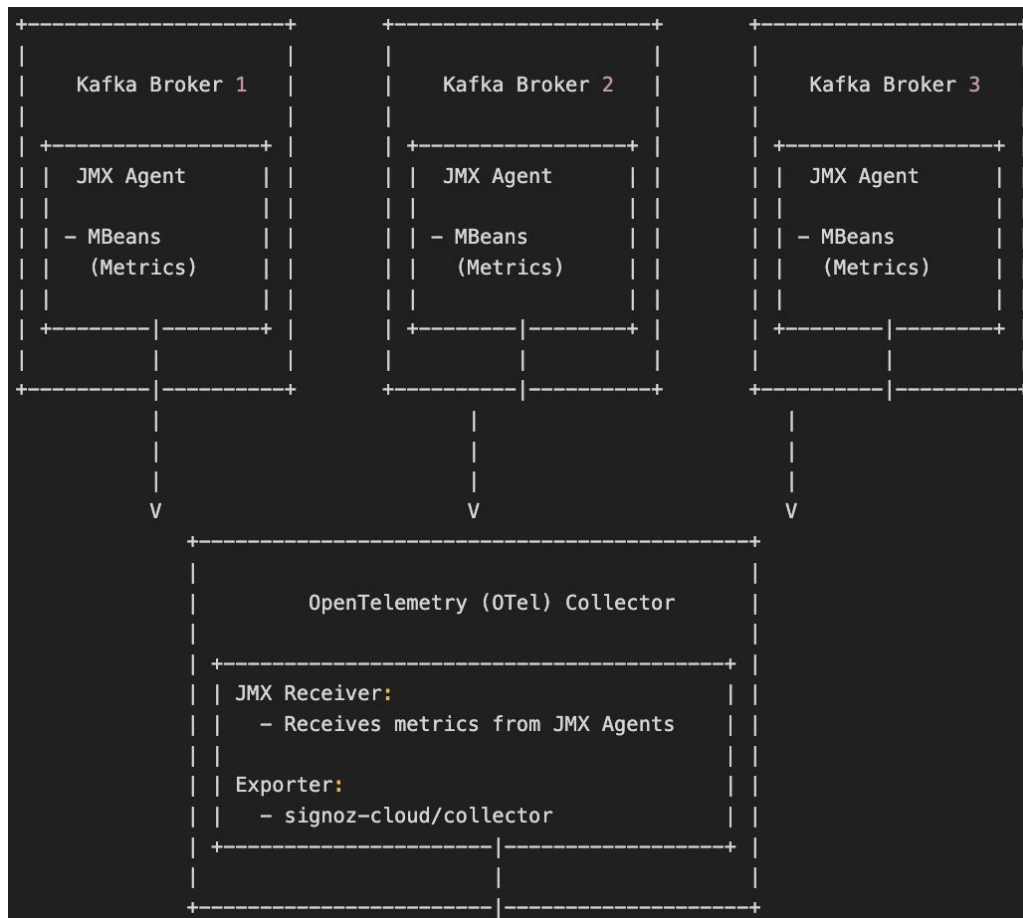


KubeCon



CloudNativeCon

North America 2024





KubeCon



CloudNativeCon



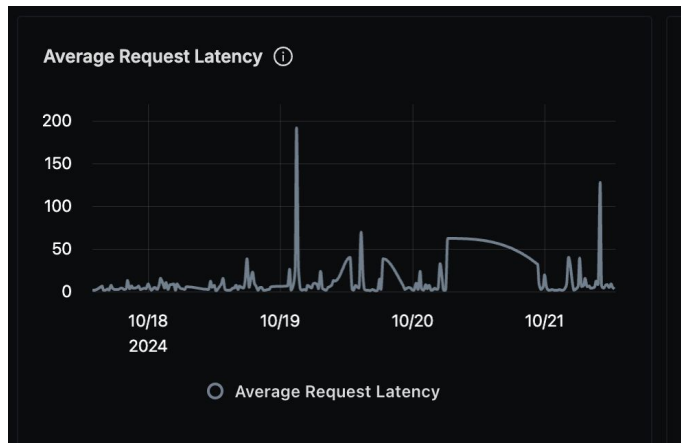
North America 2024

# What's missing with the metrics?



SigNoz

# Kafka Broker-Native Metrics







# Demo





KubeCon



CloudNativeCon

North America 2024

# Try out the correlation!

<https://signoz.io/docs/messaging-queues/kafka/>

<https://github.com/SigNoz/kafka-opentelemetry-instrumentation>



**SigNoz**

- Participate in OTel Messaging Semantic Conventions

- Define new metrics, add attributes or conventions for new systems

- Improve existing instrumentations

- Available in [Java](#), [Python](#), [.NET](#), [node.js](#), and [other languages](#)

- Ask for native instrumentation in your favorite Kafka client library



KubeCon




CloudNativeCon

North America 2024

# Thanks for joining!

## Questions?



We appreciate your feedback!