

# Engineering a Kubernetes Operator

## Lessons Learned Versions 1 to 5

**Andrew L'Ecuyer**  
Sr. Director of Kubernetes Engineering

**crunchy**data

# About Me



- Sr. Director of Kubernetes Engineering at Crunchy Data
- Manage the Engineering Team responsible for PGO, Crunchy Data's Postgres Operator
- I have experienced firsthand the **explosion of operators in recent years**

**Email:** [andrew.lecuyer@crunchydata.com](mailto:andrew.lecuyer@crunchydata.com)

**Discord:** andrewlecuyer

**GitHub:** andrewlecuyer

# Crunchy Data: Postgres Anywhere

## BARE METAL, VMs, CLOUD

### Crunchy Postgres

Crunchy Certified PostgreSQL is production ready Postgres.

#### INCLUDES:

- ✓ Backups
- ✓ Disaster recovery
- ✓ High availability
- ✓ Monitoring
- ✓ Automation
- ✓ Self managed

## KUBERNETES

### Crunchy Postgres for Kubernetes

Cloud Native Postgres on Kubernetes powered by Crunchy Postgres Operator.

#### INCLUDES:

- ✓ Simple provisioning
- ✓ Backups and DR included
- ✓ High availability
- ✓ Seamless upgrades
- ✓ Scale from 1 to thousands of databases
- ✓ Self managed

## FULLY MANAGED CLOUD

### Crunchy Bridge

The fully managed Postgres option on your choice of Cloud provider.

#### INCLUDES:

- ✓ AWS, Azure or GCP
- ✓ Continuous protection
- ✓ Backups
- ✓ Point in Time Recovery
- ✓ Pay for what you use
- ✓ The developer experience you want

# Outline



Insights & lessons-learned from Crunchy Data's journey building the first five versions of a Kubernetes Operator for Postgres



Will specifically focus on: High Availability (HA), Disaster Recovery (DR) and Upgrades



Will highlight important evolutions within Kubernetes (and Postgres) that have empowered operator development in recent years



As you will see, there has never been a better time to build an operator!

# Kubernetes Operator

Operators manage complexity.  
Adding HA, DR and seamless upgrades to Postgres isn't easy!

Operators bring new user communities into Kubernetes by making Kubernetes accessible

Operator

Current State

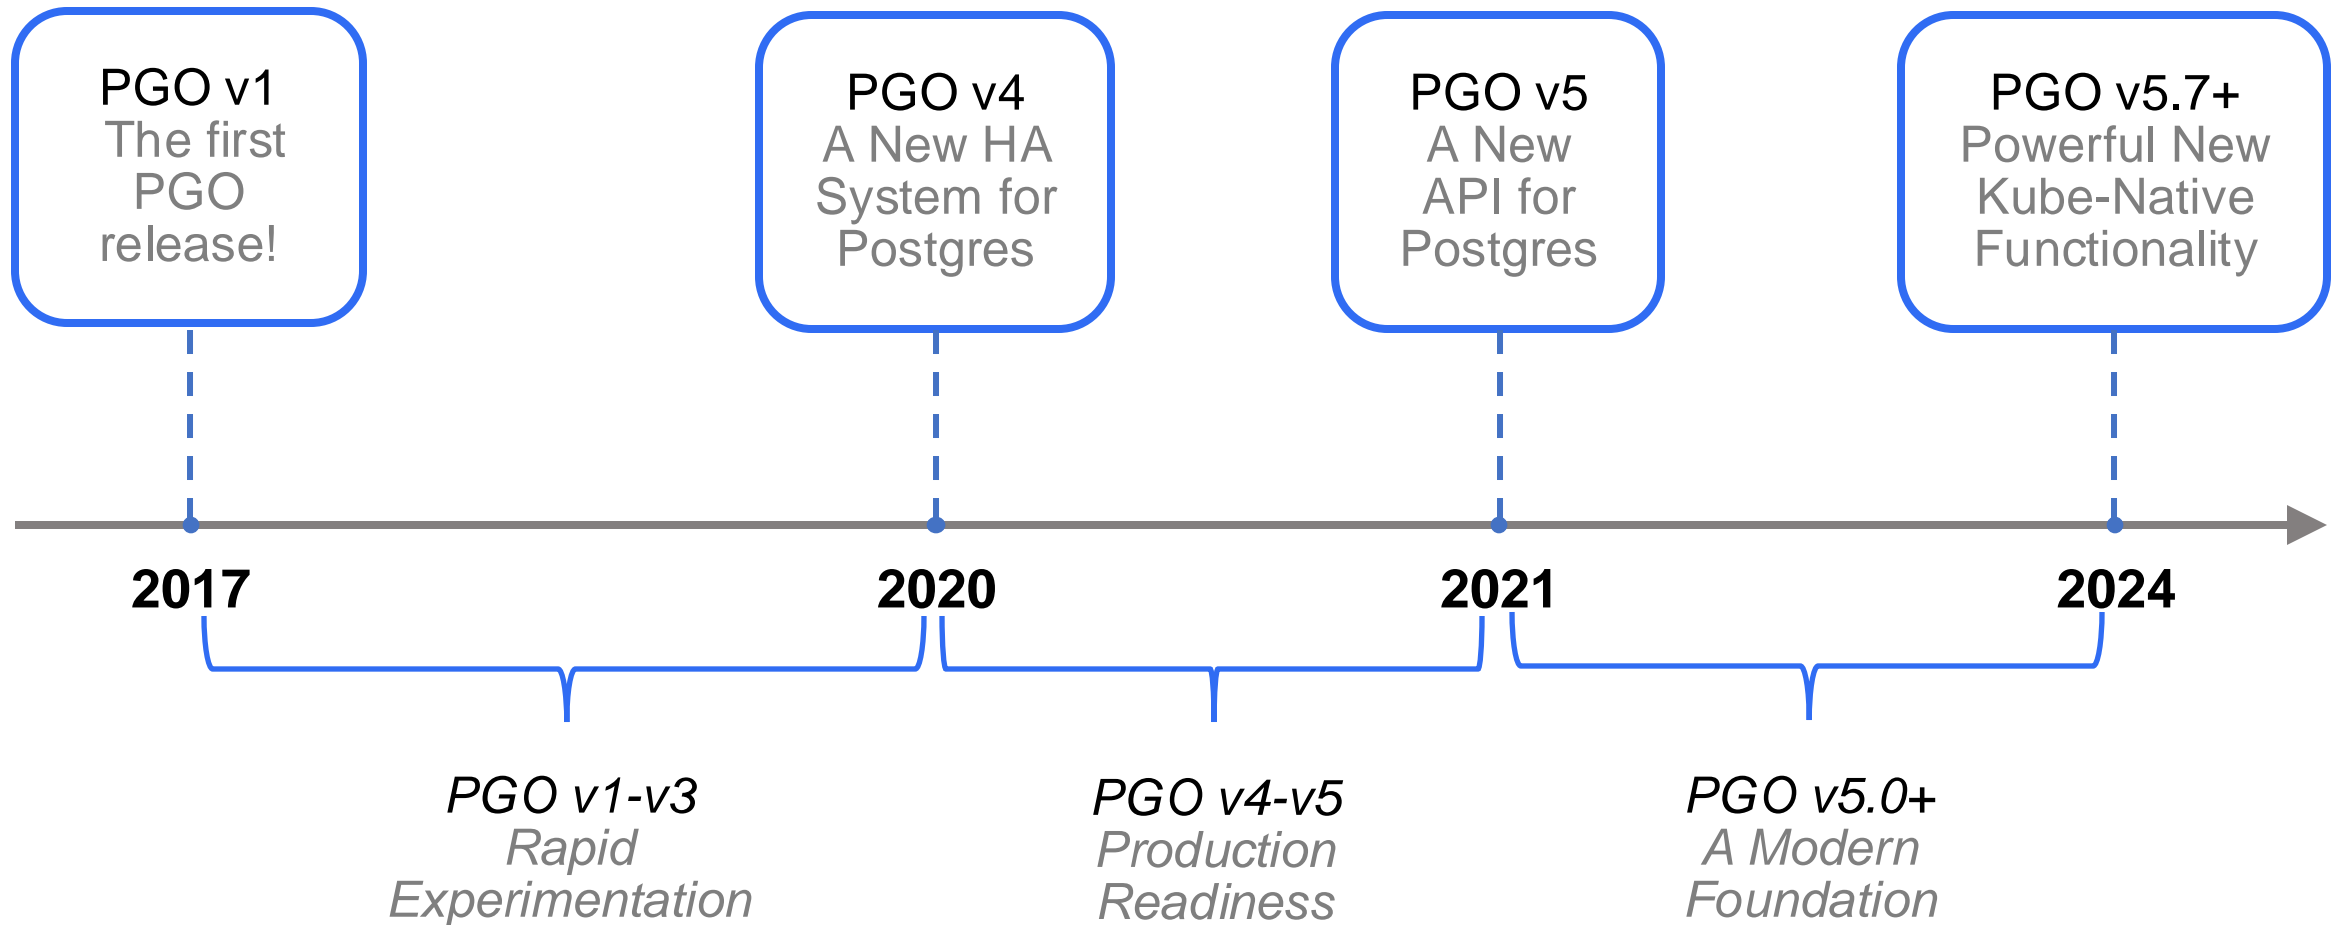
Observe

Diff

Act

Desired State

# PGO, the Postgres Operator from Crunchy Data



# Kubernetes Landscape in 2017

2014

- June 7 – Initial release
- July 21 – Kubernetes v.1.0 gets released
- Feb 23 – First release of Helm
- July 21 – Kubernetes v1.0 is released and CNCF is established
- Nov 3 – OpenShift joins the project
- Sep 29 – Pokémon GO! Case study

March 27, 2017  
PGO 1.0 released

March 28, 2017  
Kubernetes v1.6 released

Nov 16 – Helm 2.0.0 is released

Mar 27 – PGO 1.0 released

March 28 – Kubernetes v1.6 release moves dynamic storage provisioning to stable

2024

## Kubernetes Tooling

- Helm still on version 2
- Kustomize has not been released

## Kubernetes API

- The new StatefulSets API is in beta, after being renamed from the “PetSet” API in Kubernetes v1.5
- StorageClass and dynamic volume provisioning were promoted to stable in Kubernetes v1.6

## Operator Tooling

- Kubebuilder, Operator SDK and Controller Runtime projects do not exist
- Primary focus is getting applications and services up and running



# High Availability

**crunchy**data



# High Availability: Design Considerations

Use an existing HA solution for Postgres that is Kubernetes-ready?

Vs.

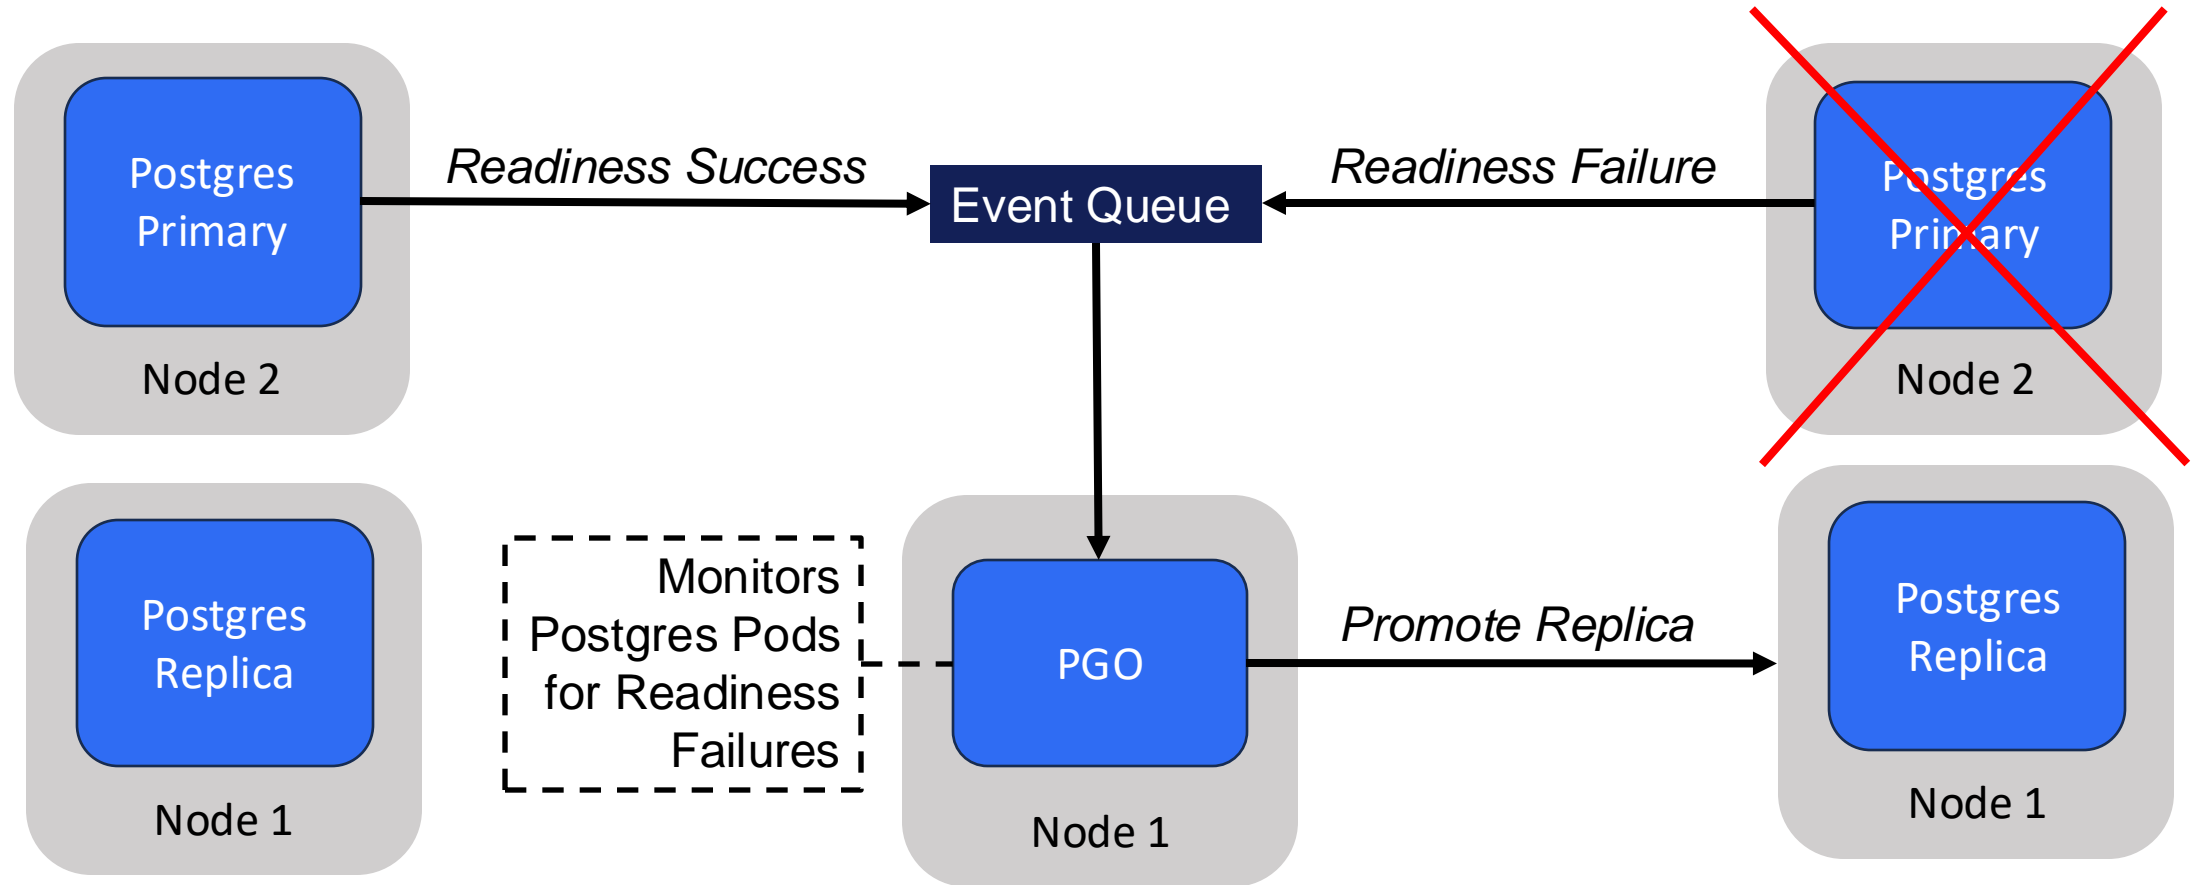
Or build a custom HA solution for Postgres using the operator?

Ideally both Postgres and the operator should be highly-available

...

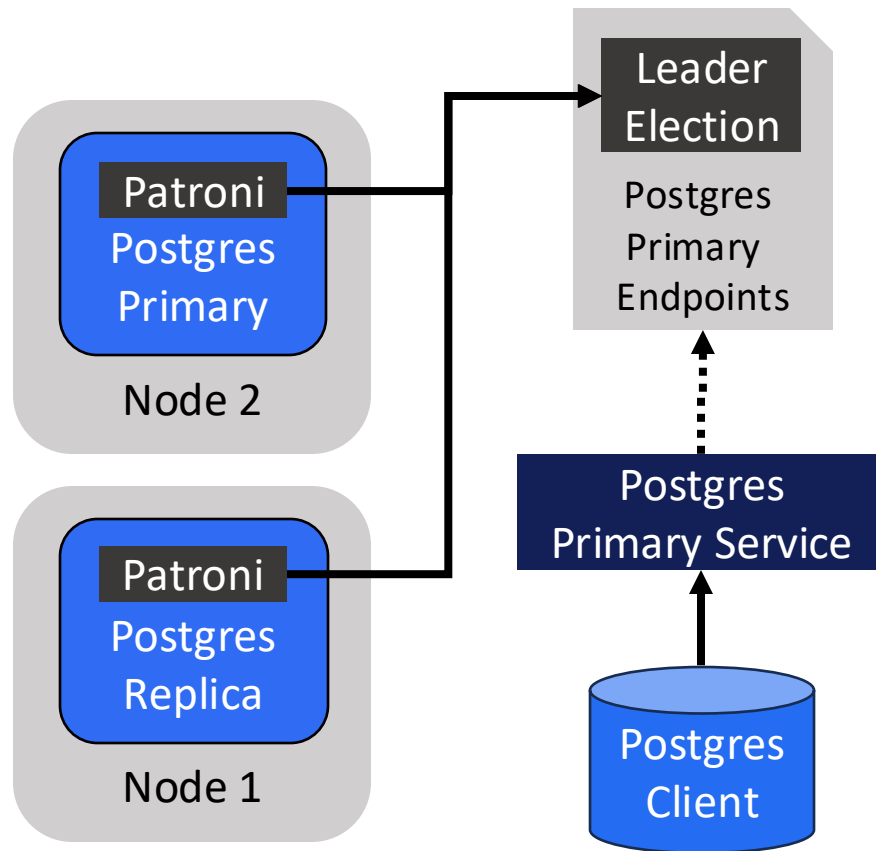
However, database availability is the top priority!

# Postgres High Availability: PGO Versions 1-3

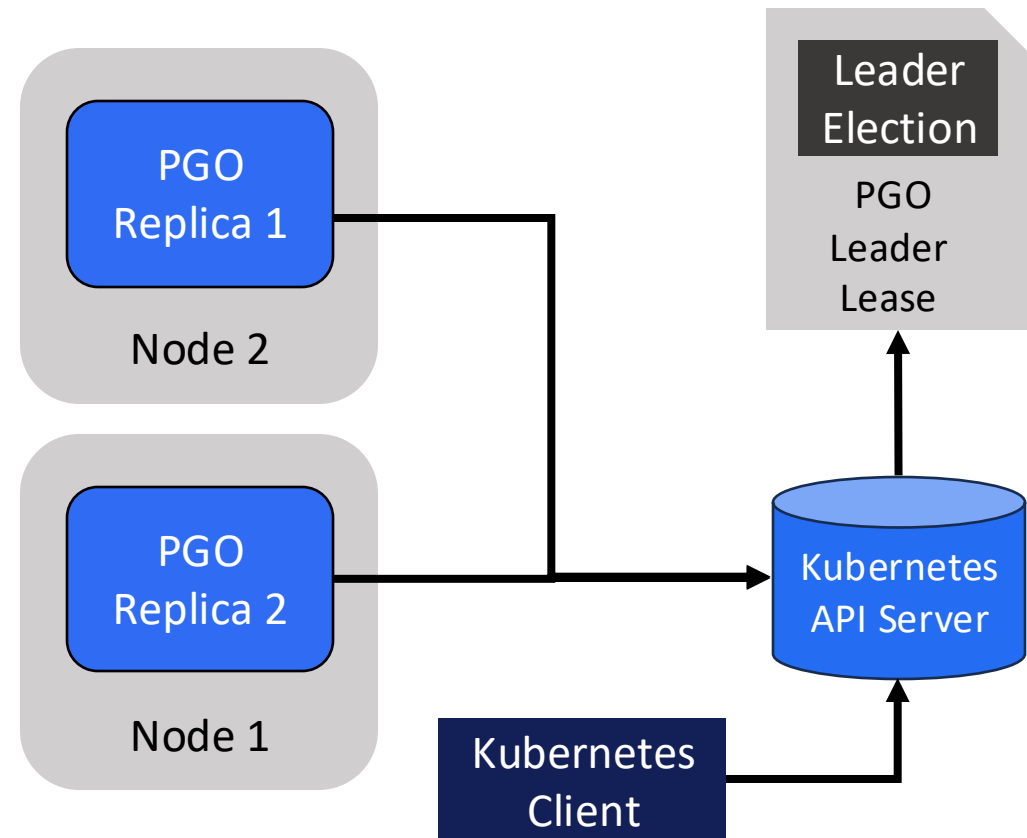


# High Availability: Current Solution

## Postgres



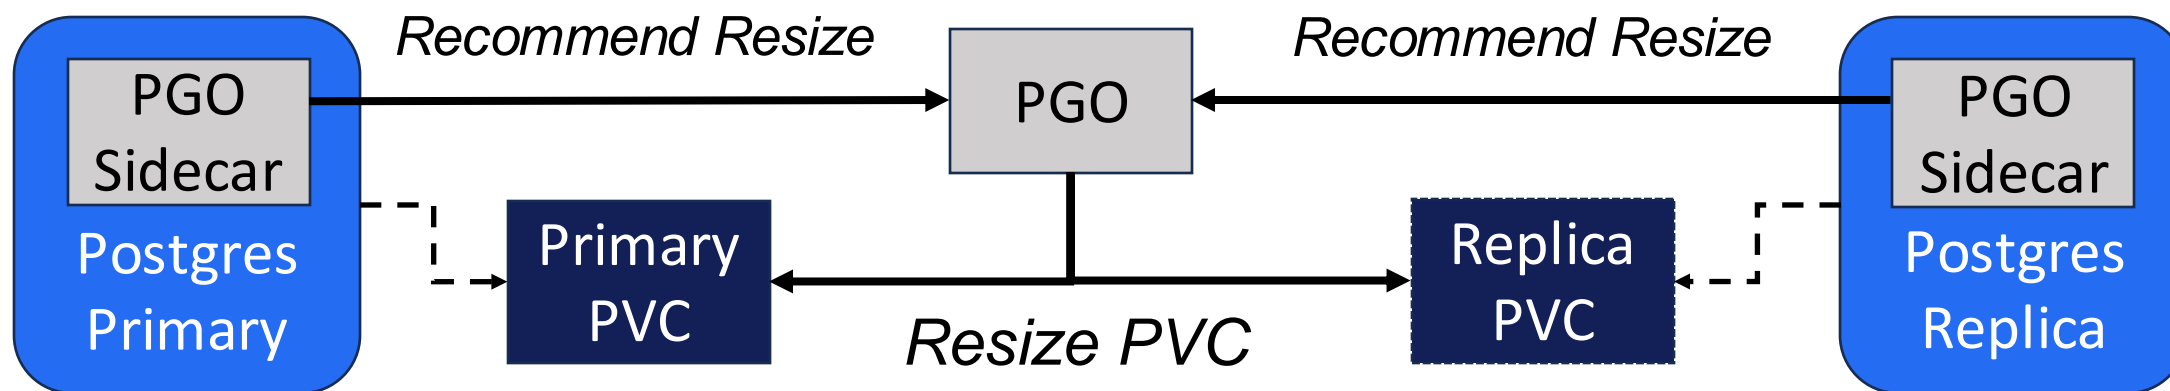
## PGO



> Postgres High Availability (HA) demo here

> [Postgres Operator High Availability \(HA\) demo here](#)

# Auto-Grow Evolution



Run a PGO sidecar in each Postgres instance Pod to determine when we're running out of storage space. Then, use Kubernetes PVC volume expansion to allow automatic resizing of PVC's without a rolling update

# Upgrades

**crunchy**data

# Upgrades: Design Considerations

Fully automate any/all upgrades?


Vs.

Or require manual intervention in certain places?



# Upgrades: Solution

A safe, fully-automated rolling update strategy for most upgrades & changes



A user-initiated, semi-automated strategy for Postgres major version upgrades

Steps:

1. Take a full backup
2. Create a PGUpgrade resource
3. Shutdown & annotate the cluster
4. Wait for upgrade to complete
5. Start the cluster
6. Complete post-upgrade tasks

> Rolling update demo here

> Postgres major version upgrade demo here

# Disaster Recovery

crunchydata

# Disaster Recovery: Design Considerations

Build a custom solution on top of existing DR tooling?

Vs.

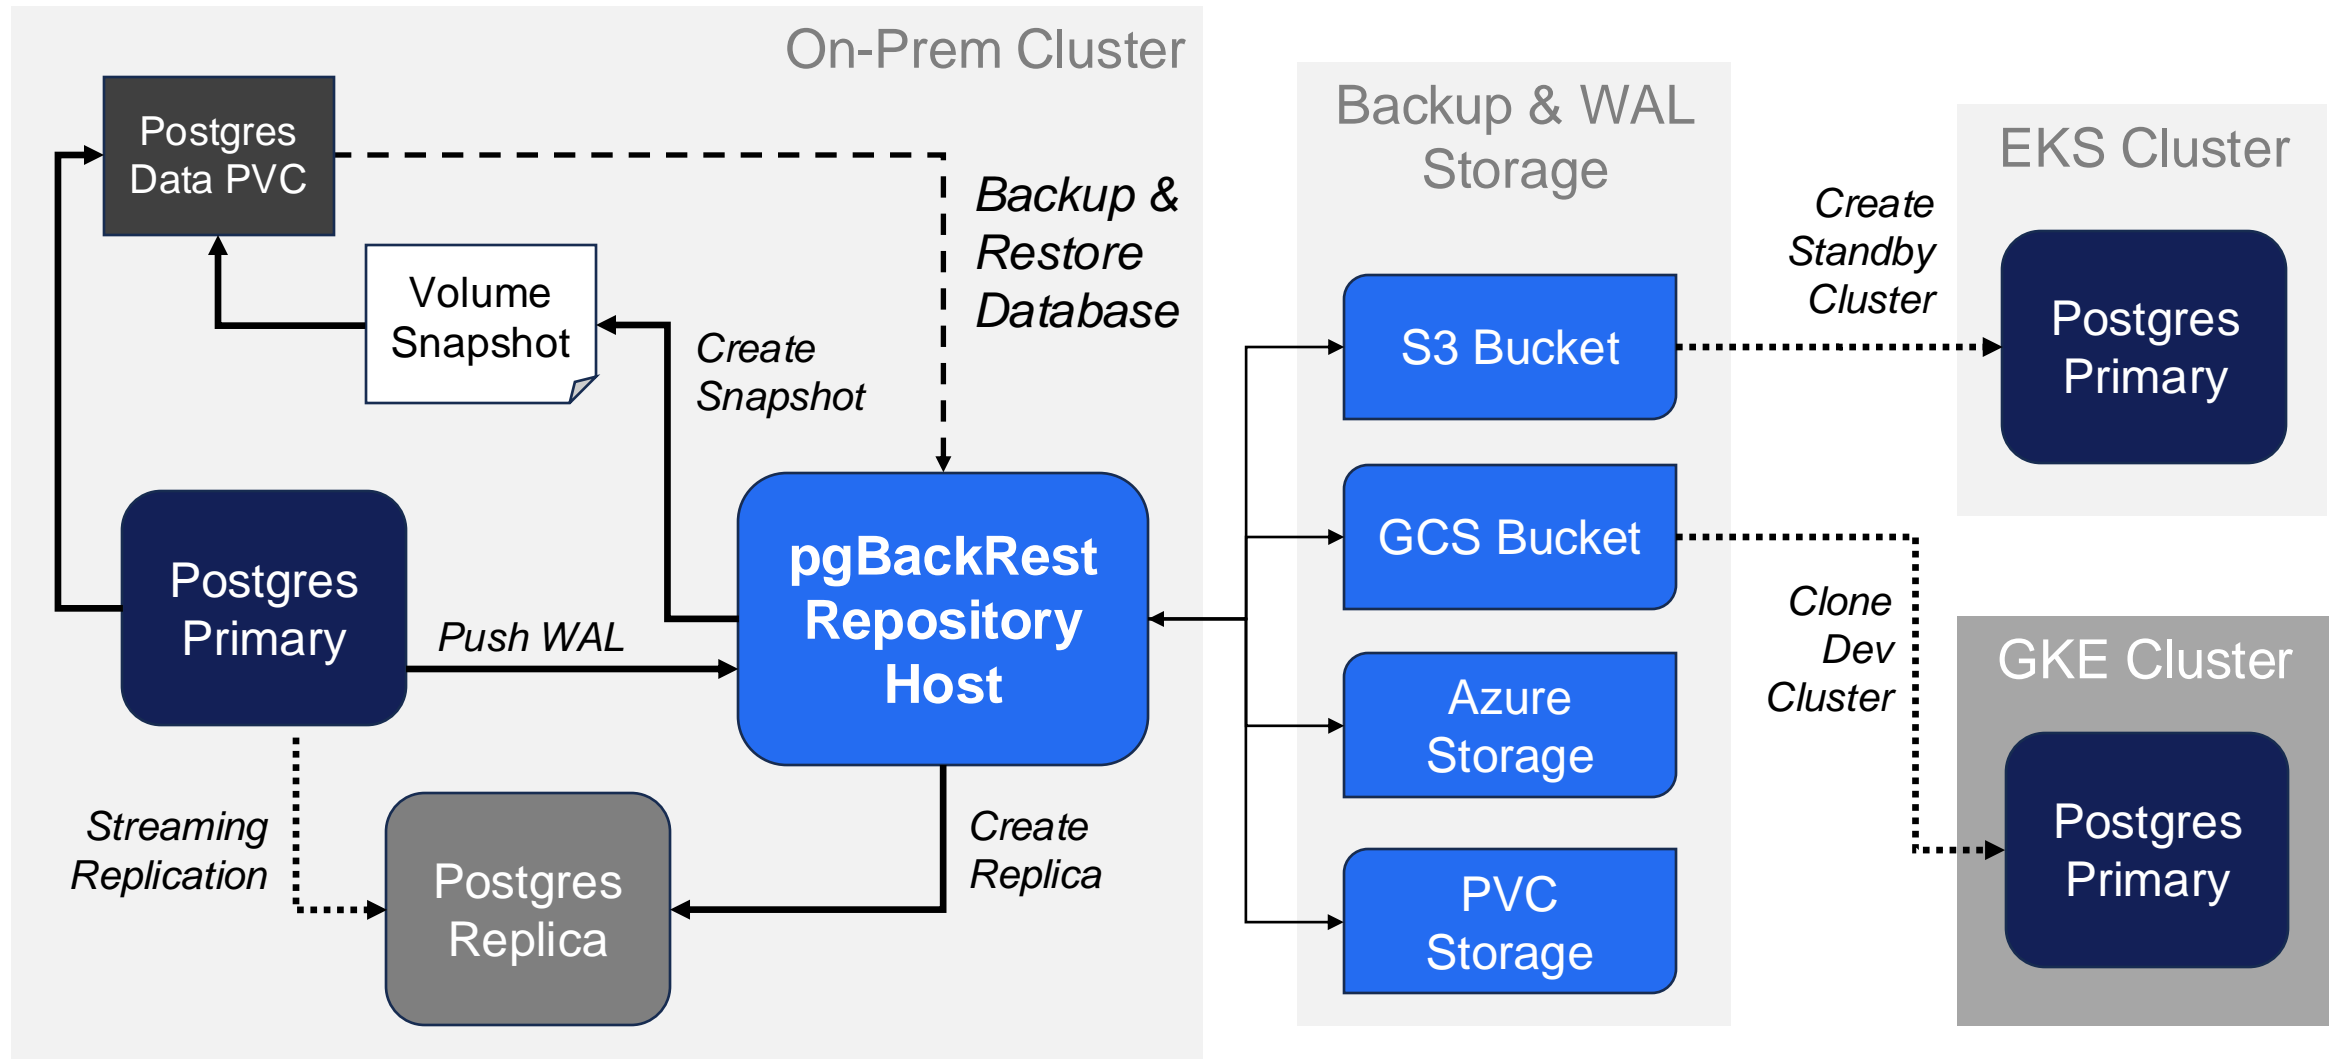
Or work with the Postgres community to better align DR tooling with Kubernetes?

Stick with Postgres-native solutions for Disaster Recovery?

Vs.

Or use Kubernetes-native solutions for DR?

# Disaster Recovery: Solution



> Disaster Recovery (DR) demo here

# Summary

**crunchy** data



# HA: Solution & Lessons Learned



**PGO Solution Summary:** Patroni for Postgres High Availability, controller-runtime for operator/PGO High Availability and Kubernetes volume expansion for auto-growing disks.



- Fight the “Not Invented Here” syndrome, and embrace existing solutions within the community
- A decentralized architecture allows us to scale
- Prevention is better than preparedness

# Upgrades: Solution & Lessons Learned



**PGO Solution Summary:** A safe rolling update strategy for config changes, minor Postgres upgrades and PGO upgrades, and an orchestratable solution for Postgres major version upgrades



- Manage risk associated with upgrade automation, and only automate when risks can be mitigated
- When we can't automate, ensure we can orchestrate
- Use status, conditions & events for upgrade visibility, and to empower engineers to safely perform upgrades

# DR: Solution & Lessons Learned



**PGO Solution Summary:** pgBackRest for multi-cloud backup/restore functionality & data mobility, and Volume Snapshots to improve restore performance



- Focus on recovery rather than backups
- A robust DR solution can enable data mobility
- Use Postgres-native solutions to *safely* utilize Kubernetes-native solutions

# Conclusion

**crunchy**data

# Should You Build an Operator?

- **A great solution for Postgres!** Manages the complexity of deploying & managing a production-ready Postgres cluster
- More knowledge is available than ever before (documentation, blogs, books, etc.) for operator development
- Multiple mature operator frameworks to help you get started
- **Provides practical knowledge & skills for contributing back to Kubernetes**

# Kubernetes Landscape in 2024

2014

- June 7 – Initial release
- July 21 – Kubernetes v.1.0 gets released
- Feb 23 – First release of Helm
- July 21 – Kubernetes v1.0 is released and CNCF is established
- Nov 3 – OpenShift joins the project
- Sep 29 – Pokémon GO! Case study

August 13, 2024 – Helm 2.0.0 is released  
Kubernetes v1.31 released

March 18, 2024 – PGO 1.0 released  
PGO v5.7 released  
March 28 – Kubernetes v1.6 release moves dynamic storage provisioning to stable

2024

## Kubernetes Tooling

- Helm now on version 3
- Kustomize now included in kubectl

## Kubernetes API

- Stateful deployments are first-class within the Kubernetes platform, with a StatefulSet API that is both stable & mature

## Operator Tooling

- Kubebuilder, Operator SDK & controller-runtime projects stable & mature
- Focus on advanced needs, e.g. multi-cluster, security, supply chain, and more



