



KubeCon



CloudNativeCon

North America 2024

GitOps at Production Scale with Flux

Leigh Capili

Senior DevRel Engineer, Flox

Priyanka “Pinky” Ravi

Platform Tech Advocate, G-Research





KubeCon



CloudNativeCon

North America 2024



CLOUD NATIVE
COMPUTING FOUNDATION

What is Flux?

- A foundation for Continuous and Progressive Delivery solutions for Kubernetes
- A flexible toolkit for you to build your own GitOps platform



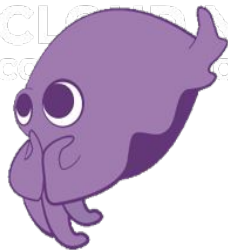


Benefits of Flux

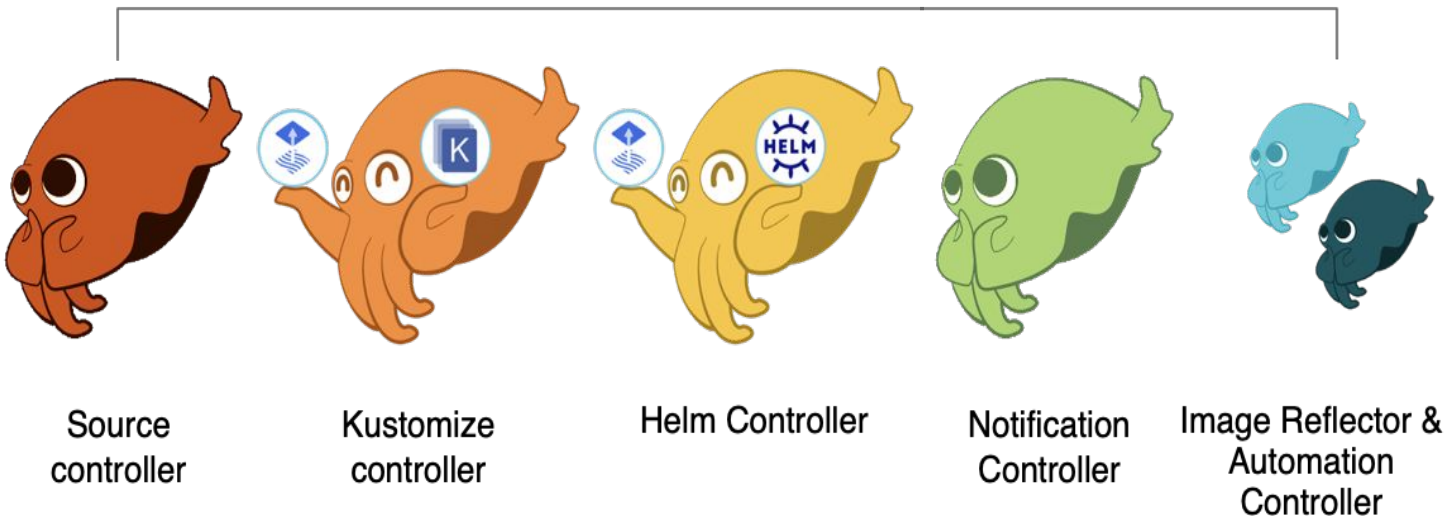
- Reduces developer burden
- Extensible
- Comes with out of the box support for Kustomize and Helm
- Designed For Kubernetes
- Security at its core



Overview of Flux



Flux is a set of
Kubernetes Controllers





KubeCon



CloudNativeCon

North America 2024

Flux is easy to manage



```
> flux bootstrap
```

The bootstrap sub-commands push the Flux manifests to a Git repository and deploy Flux on the cluster.

Usage:

```
flux bootstrap [command]
```

Available Commands:

bitbucket-server	Deploy Flux on a cluster connected to a Bitbucket Server repository
git	Deploy Flux on a cluster connected to a Git repository
gitea	Deploy Flux on a cluster connected to a Gitea repository
github	Deploy Flux on a cluster connected to a GitHub repository
gitlab	Deploy Flux on a cluster connected to a GitLab repository



Flux Operator

The Flux Operator manages the lifecycle of the CNCF Flux project.

Flux UIs




Enter Filter reset

chatbot-uiLogsDescribe


Pods

Running

Dependencies


 chatbot-ui

Address

chatbot-ui.default.svc.cluster.local
http://127.0.0.1:3000  (port-forward)

Port-forward command

Sync



 Applied about 3 hours ago -> 3bc14c1b (flux-system/builtin-apps)

expressjs-test-appLogsDescribe


Pods

Running

Dependencies


 expressjs-test-app  expressjs-test-app

Address

expressjs-test-app.default.svc.cluster.local
http://127.0.0.1:8080  (port-forward)

Port-forward command

Sync

 Applied about 3 hours ago -> 3bc14c1b (flux-system/builtin-apps)

notification-controllerLogsDescribe

Pods




Dependencies

Address

notification-controller.default.svc.cluster.local

Port-forward command

Sync

SOURCES:  Ready(2/2) KUSTOMIZATIONS:  Reconciling(2/3) HELM-RELEASES:  Ready(3/3) ↑

Flux UIs



Headlamp

Cluster: iq-new-changes-hl-cluster

Home

Cluster

Workloads

Storage

Network

Security

Configuration

Custom Resources

Map (beta)

Flux

Kustomizations

HelmReleases

Sources

Image Automations

Notifications

Flux Runtime

Image Repositories

Image Policies

Image Update Automations

Name	Namespace	Status	Insecure	Secret Ref	Image	Interval	Schedule	Service Account	Timeout	Age
podinfo	default	Ready	False	-	ghcr.io/stefanprodan/podinfo	5h		-		2w
podinfo	flux-system	Ready	False	-	ghcr.io/stefanprodan/podinfo	5m		-		3w

Name	Namespace	Status	Policy	Age
podinfo	default	Ready	("semver":("range":"5.1.x"))	2w
podinfo-policy	default	Ready	("semver":("range":"5.0.x"))	2w

Name	Namespace	Status	Git	Interval	Update	Age
podinfo-update	default	Failed	("commit":("author":("email":"fluxcdbot@users.noreply.github.com","name":"fluxcdbot")),("push":("branch":"main"))	30m	("strategy":"Setters")	2w



CLOUD NATIVE
COMPUTING FOUNDATION



**Flux adapts to your
organizational structure and
growth velocity**

CD growth factors

Microservices proliferation

Apps made out of many moving pieces each with its own release cadence and delivery workflow

SLAs & High availability & Disaster Recovery

Multi-AZ, Multi-region app instances

Security constraints

Multi-tenancy, Network isolation, GDPR

Business expansion

New apps are developed, while legacy apps still require maintenance





CLOUD NATIVE
COMPUTING FOUNDATION



**We'll need to deploy
thousands of workloads
over hundreds of clusters**



Can Flux do it?



Flux scaling strategies

**Source Optimizations
Controller Fine Tuning**



Vertical Scaling



Horizontal Scaling & Sharding





Optimizations & Fine Tuning

Source optimizations

- Migrate from Helm HTTP/S repositories to OCI
- Use dedicated Git repositories for Flux
- Split the Kubernetes resources into multiple Flux Kustomizations to take advantage of concurrent reconciliations

Controller fine tuning

- Use persistent storage for Flux internal artifacts
- Use tmpfs for in-memory kustomize builds





Vertical scaling

Resource limits

- Increase controller CPU and Memory limits
- Increase the controller reconciliation concurrency level based on CPU limits

Rate limits

- Monitor Kubernetes API requests rate limits error rate and adjust the reconciliation concurrency level

<https://fluxcd.io/flux/installation/configuration/vertical-scaling>



Flux “Mean Time To Production” benchmark



Starting with Flux v2.2, we have set up benchmarks that measure Mean Time To Production (MTTP). The MTTP benchmark measures the time it takes for Flux to deploy application changes into production.

Specs:

GitHub hosted-runner
ubuntu-latest-16-cores

Kubernetes KinD
v1.28.0 / 3 nodes

Flux source-controller
1CPU / 1Gi / concurrency 10

Flux kustomize-controller
1CPU / 1Gi / concurrency 20

Flux helm-controller
2CPU / 1Gi / concurrency 10

Objects	Type	Flux component	Duration	Max Memory
500	OCIRepository	source-controller	45s	65Mi
500	Kustomization	kustomize-controller	2m2s	72Mi
500	HelmChart	source-controller	45s	68Mi
500	HelmRelease	helm-controller	2m55s	350Mi
1000	OCIRepository	source-controller	1m30s	67Mi
1000	Kustomization	kustomize-controller	4m15s	112Mi
1000	HelmChart	source-controller	1m30s	110Mi
1000	HelmRelease	helm-controller	8m2s	620Mi



Horizontal scaling

To enable horizontal scaling, each controller can be deployed multiple times with a unique label selector which is used as the sharding key.

To spread the load between controller instances, the Flux resources can be assigned to a particular instance using the `sharding.fluxcd.io/key` label.

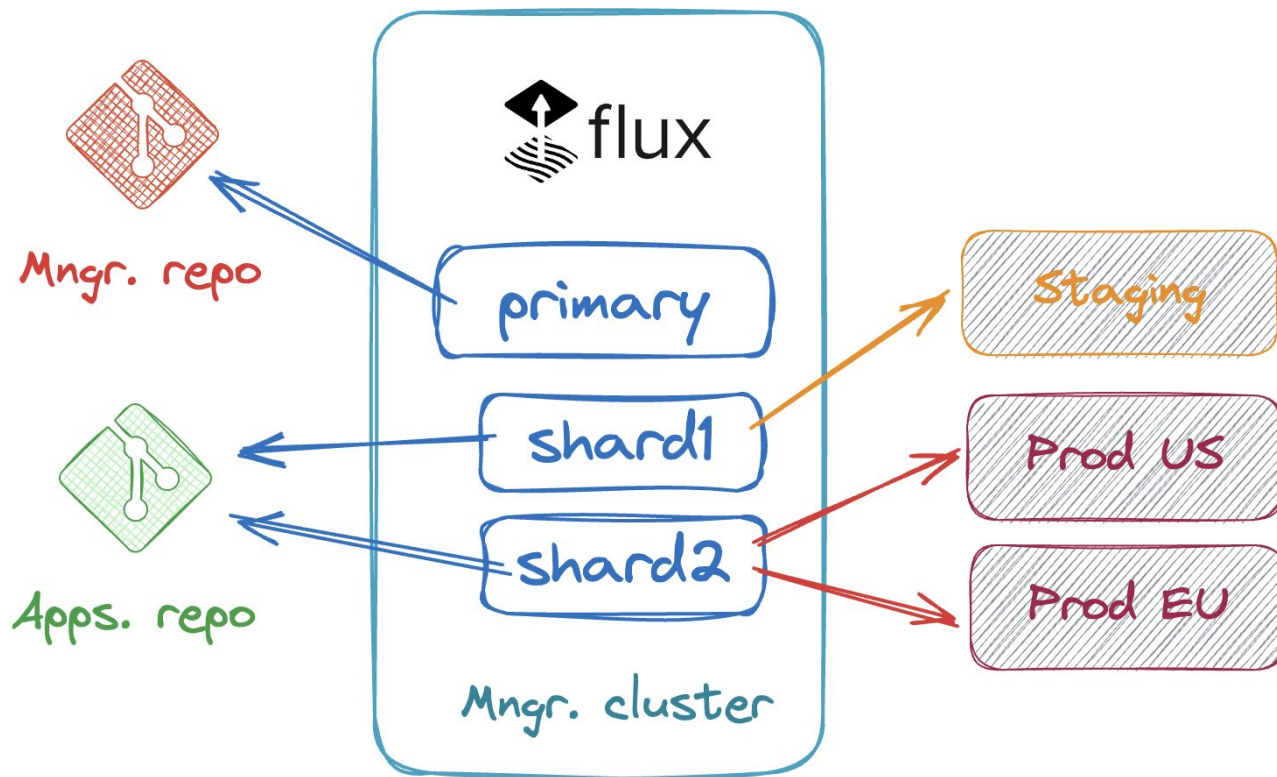
Sharding strategies

- Multi-tenant (instance per tenant)
- Multi-cluster (instance per cluster group)

<https://fluxcd.io/flux/installation/configuration/sharding>



Flux sharding for multi-cluster delivery





CLOUD NATIVE
COMPUTING FOUNDATION



Flux Updates 2024

<https://fluxcd.io/roadmap>

General Availability

Flux v2.3 (Q2 2024)

Promoted to GA the Flux Helm APIs and the Flux Helm functionalities.

- HelmRepository
- HelmChart
- HelmRelease

Flux v2.4 (Q3 2024)

Promoted to GA the S3-compatible storage API.

- Bucket





Helm OCI improvements

Flux v2.3 came with support for HelmReleases to refer to OCIRepositories, as an alternative to HelmRepositories/HelmCharts.

- Better UX when debugging Helm releases
- Allow reusing the same chart between releases
- Allow pinning Helm charts by OCI digests
- Allow automating Helm upgrades using semver ranges scoped to release candidates only
- Provenance verification with Cosign and Notary Notation





Flux Operator and OpenShift Compatibility

Flux can be installed on Red Hat OpenShift cluster directly from OperatorHub or from the RedHat production-ready catalog.

The Flux Operator, a new component in the Flux CD ecosystem **developed by ControlPlane** that automates the lifecycle management of Flux components and streamlines the GitOps workflows for Kubernetes clusters.





Terraform provider improvements

The Flux Terraform provider has undergone a major refactoring and now supports

- air-gapped bootstrap
- drift detection and correction for Flux components
- the ability to upgrade and restore the Flux controllers in-cluster

Starting with this release, the provider is fully compatible with OpenTofu.





Azure DevOps OIDC Authentication

Starting with Flux v2.4, you can configure source-controller and image-automation-controller to authenticate against Azure DevOps repositories using AKS Workload Identity.

Big thanks to **Dipti** from **Microsoft** for contributing this feature and for helping the Flux team maintaining the Azure integration on the long-run.





GitHub App OIDC Authentication (WIP)

In a future release, Flux will support OIDC-based authentication against Git repositories using GitHub App credentials.

This feature is **work-in-progress** and will be available early next year.





CLOUD NATIVE
COMPUTING FOUNDATION



v2.5 (Q1 2025)

Status: In progress

The primary goal of this milestone is to make a generally available release for the Flux image automation APIs.

<https://fluxcd.io/roadmap>

Get involved with Flux!

We invite the community to help us shape the future of Flux.

We have in place an **RFC process** for new features and enhancements proposals. We are keen to work with the community on RFCs and drive the project forward in a sustainable manner.

We want to **enable community members to take full ownership** of Flux features and share the responsibility of feature stability and longer-term maintenance.

<https://github.com/fluxcd/flux2/tree/main/rfcs>

