





North America 2024

Microsegment Your Network Like Mastercard with AdminNetworkPolicy





North America 2024



Surya Seetharaman
Principal Software Engineer @ Red Hat
SIG-Network-Policy-API Contributor
@@tssurya



John Zaiss
Principal Engineer @ Mastercard
20+ years at Mastercard
6+ years working in Kubernetes



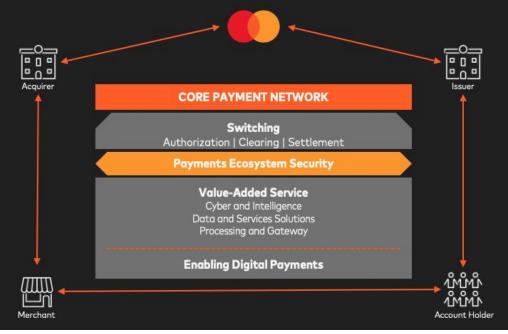
Daniel Ruggeri
Distinguished Engineer @ Mastercard
Web, Cloud, Network, Infra, Automation, and
FOSS nerd

Mastercard - Overview



Mastercard - Overview

Mastercard is a technology company in the global payments industry. We connect consumers, financial institutions, merchants, governments, digital partners, businesses and other organizations worldwide by enabling electronic payments and making those payment transactions safe, simple, smart and accessible.



The Mastercard Journey to Kubernetes



The Mastercard private cloud ecosystem lacked a multifunctional enterprise Kubernetes platform essential for effectively managing and executing diverse container workloads at scale. We've been on this journey for quite some time and have designed a solution, which we have internally branded as the Mastercard Kubernetes Service (MKS).













Reliability







North America 2024

Microsegmentation in Mastercard

The Declarative Vision



Application Team's Single Manifest

Here is my code

def myApp():
// application logic
// business rules

Resources I need

- 4 CPU cores
- 8GB RAM
- 50GB Storage

Connections I need

- database access
- API endpoints
- Message queue

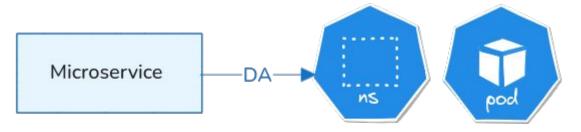
Declares Requirements Infrastructure

?

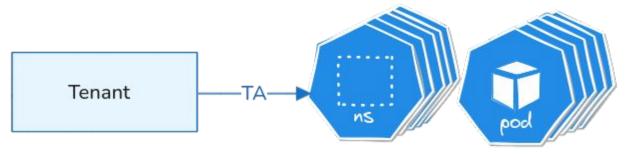
Implementation Details are a blackbox to the App Team. They don't know about K8s Resource details like Network Policies...

Converting declarative request to Kubernetes resources





Each microservice (deployable asset) is 1 pod in 1 namespace

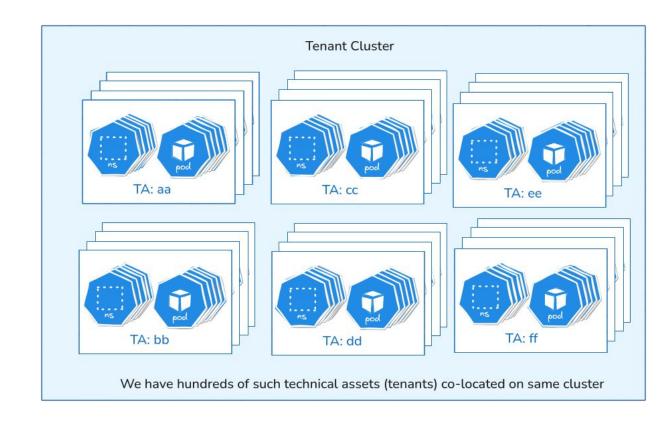


Each tenant (technical asset) has approximately 5 microservices (deployable assets)

Requirement for Microsegmentation

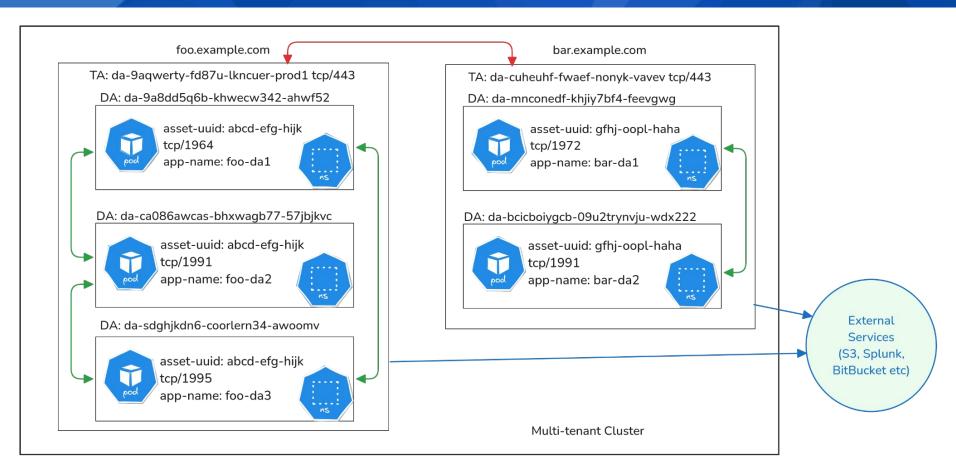


- Mastercard's default policy: "If it can be microsegmented, it SHOULD be... but if it is in the application data plane, it MUST be."
- Need to adhere to PCI
 Data Security Standard



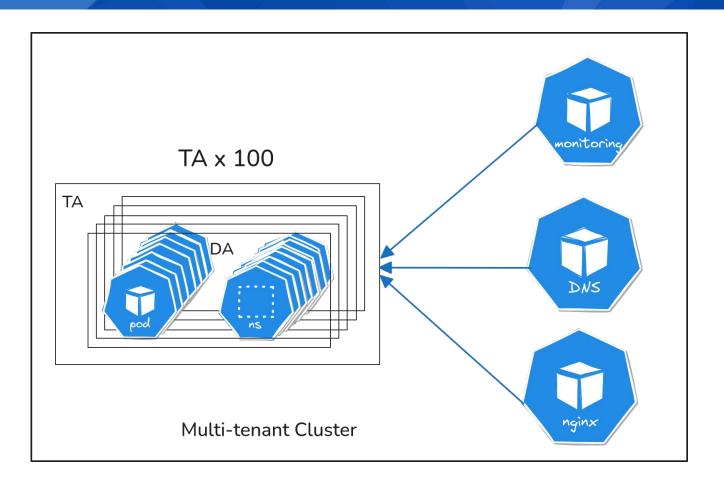
Wiring up the declarative connectivity requirements





Wiring up the infrastructure connectivity requirements

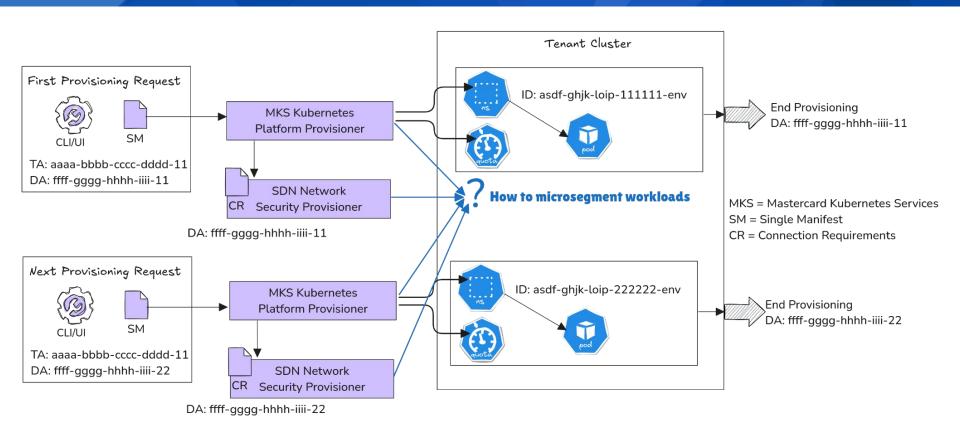




The platform team has a set of common connections we want to enforce on all TAs

Provision->Build->Deploy Workflow in Infrastructure





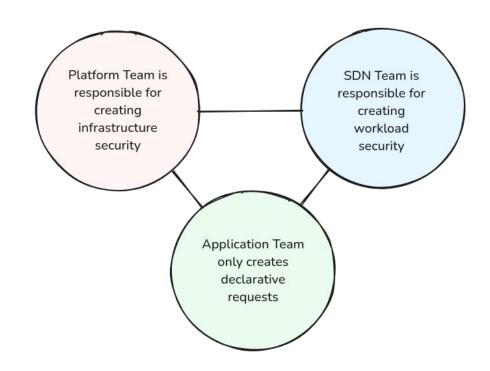
Why not use Network Policies?



 Need for a simple placeholder policy that met compliance regulations even before the workloads (ns+pod) are created

 SDN team needs to be proactive to App team requests (not reactive)

 Avoid duplication of Network Policy in every namespace



Why not use Network Policies?



Dynamic nature of DAs with SMs

Multiple updates or multiple policies would need to be created when additional DAs get

added into the TA Tenant Cluster ID: asdf-ghjk-loip-111111-env First Provisioning Request **End Provisioning** MKS Kubernetes DA: ffff-gggg-hhhh-iiii-11 Platform Provisioner SM TA: aaaa-bbbb-cccc-dddd-11 DA: ffff-gggg-hhhh-iiii-11 SDN Network MKS = Mastercard Kubernetes Services CR Security Provisioner SM = Single Manifest DA: ffff-gggg-hhhh-iiii-11 CR = Connection Requirements Next Provisioning Request ID: asdf-ghjk-loip-222222-env MKS Kubernetes **End Provisioning** Platform Provisioner DA: ffff-gggg-hhhh-iiii-22 SM TA: aaaa-bbbb-cccc-dddd-11 DA: ffff-gggg-hhhh-iiii-22 SDN Network CR Security Provisioner DA: ffff-gggg-hhhh-iiii-22

We wanted a Kubernetes upstream solution for this



- Around the same time in Kubernetes the sig-network-policy-API subproject was working on Admin Network Policy API
- Two cluster scoped CRDs:
 - AdminNetworkPolicy: Cannot be overridden by developer network policies
 - BaselineAdminNetworkPolicy: CAN be overridden by developer network policies
- Lives in <u>kubernetes-sigs repo</u>: Small, healthy community which participates in building new features around these APIs and maintains them
 - Current status: Alpha, on the journey to Beta



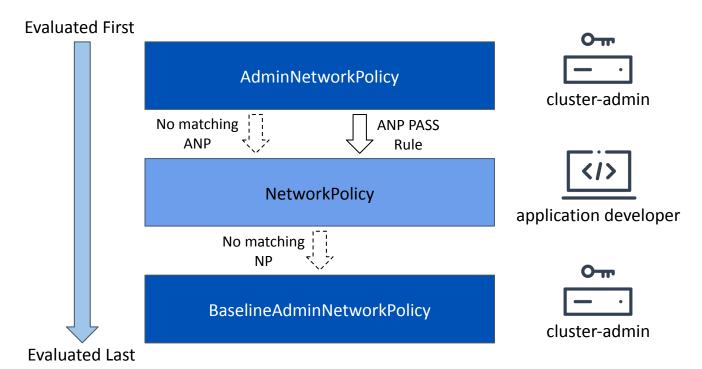


North America 2024

AdminNetworkPolicy



For the first iteration Mastercard does not plan to include network policies





```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: AdminNetworkPolicy
metadata:
name: sensitive-workload-control
spec:
priority: 34 -
subject:
  namespaces:
    matchLabels:
      tenant: true
ingress:
 - name: "allow-from-monitoring" # rule0
  action: "Allow"
  from:
  namespaces:
      matchLabels:
        kubernetes.io/metadata.name: monitoring
  ports:
  - portNumber:
      protocol: TCP
      port: 7564
   - namedPort: "scrape"
<...>
```

> 0 means highest priority and 1000 is the lowest



```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: AdminNetworkPolicy
metadata:
name: sensitive-workload-control
spec:
priority: 34 -
                                                        > 0 means highest priority and 1000 is the lowest
subject:
  namespaces:
                                                        namespaces on which this policy applies
    matchLabels:
      tenant: true
ingress:
 - name: "allow-from-monitoring" # rule0
  action: "Allow"
  from:
  namespaces:
      matchLabels:
         kubernetes.io/metadata.name: monitoring
  ports:
  - portNumber:
      protocol: TCP
      port: 7564
   - namedPort: "scrape"
<...>
```

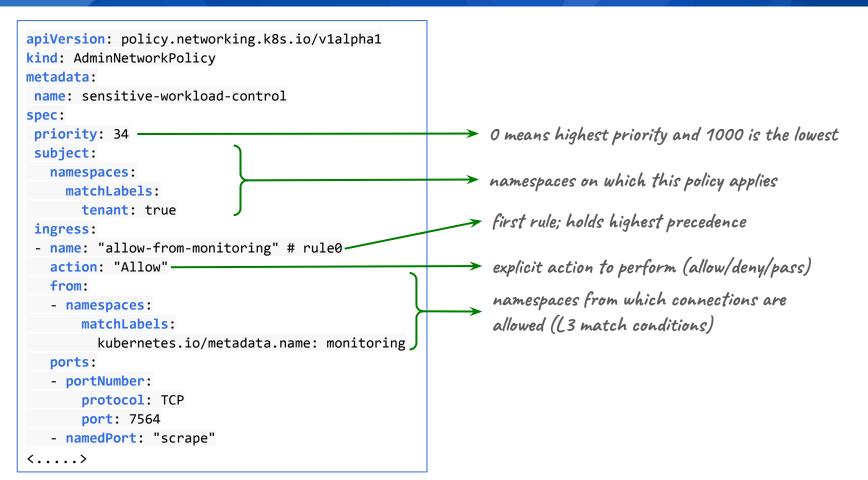


```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: AdminNetworkPolicy
metadata:
name: sensitive-workload-control
spec:
priority: 34 -
                                                         > 0 means highest priority and 1000 is the lowest
subject:
  namespaces:
                                                            namespaces on which this policy applies
    matchLabels:
       tenant: true
                                                            first rule; holds highest precedence
ingress:
 - name: "allow-from-monitoring" # rule0-
   action: "Allow"
  from:
   namespaces:
       matchLabels:
         kubernetes.io/metadata.name: monitoring
  ports:
   - portNumber:
       protocol: TCP
       port: 7564
   - namedPort: "scrape"
<...>
```

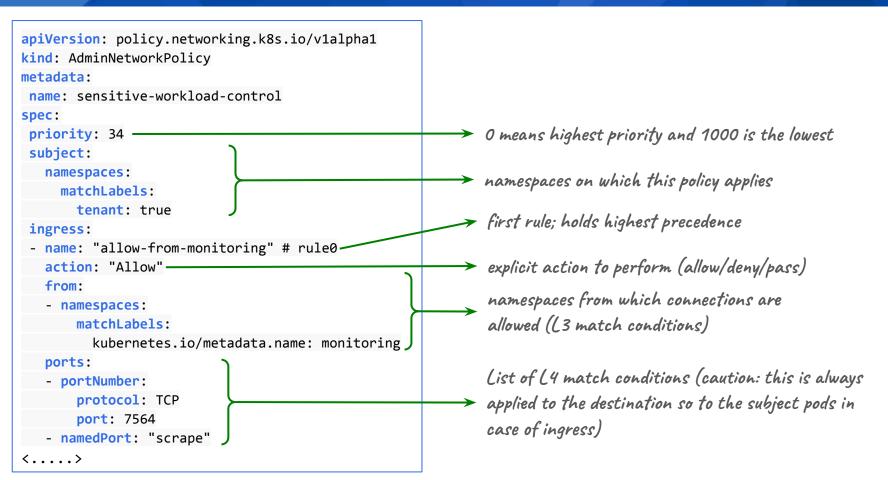


```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: AdminNetworkPolicy
metadata:
name: sensitive-workload-control
spec:
priority: 34 -
                                                         O means highest priority and 1000 is the lowest
subject:
  namespaces:
                                                           namespaces on which this policy applies
    matchLabels:
       tenant: true
                                                           first rule; holds highest precedence
ingress:
 - name: "allow-from-monitoring" # rule0
                                                         explicit action to perform (allow/deny/pass)
  action: "Allow" -
  from:
  namespaces:
       matchLabels:
         kubernetes.io/metadata.name: monitoring
  ports:
   - portNumber:
       protocol: TCP
       port: 7564
   - namedPort: "scrape"
<...>
```

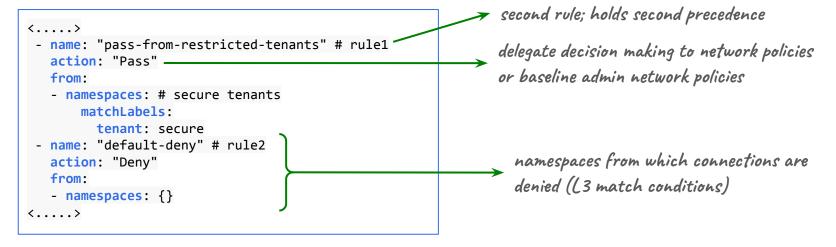












Admin Network Policy API Nodes + Networks Peer



```
<....>
egress:
- name: "allow-to-kapi-server" # rule0
   action: "Allow"
  to:
   - nodes:
                                                                  nodes to which connections are allowed
      matchExpressions:
                                                                  (L3 match conditions)
       - key: node-role.kubernetes.io/control-plane
         operator: Exists
   ports:
   - portNumber:
       protocol: TCP
       port: 6443
- name: "default-deny" # rule1
   action: "Deny"
  to:
                                                                 networks to which connections are denied
   - networks:
     - 0.0.0.0/0
                                                                 (L3 match conditions)
<....>
```

Baseline Admin Network Policy API

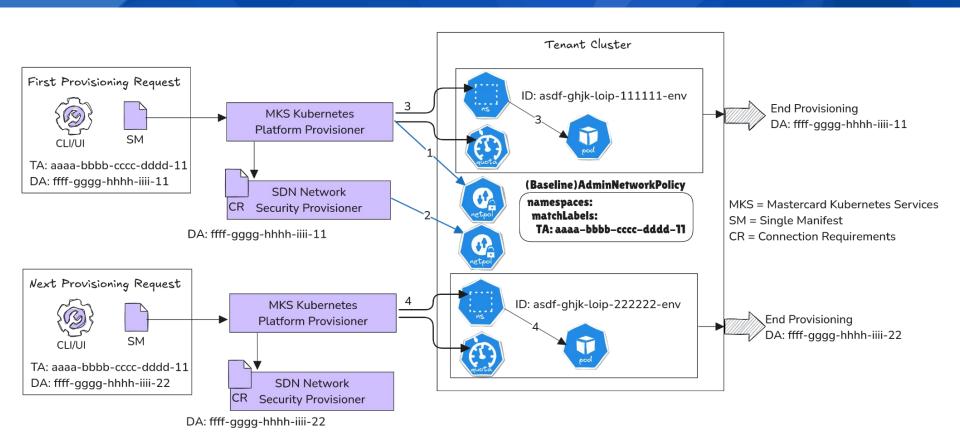


```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: BaselineAdminNetworkPolicy
metadata:
name: default
spec:
subject:
   namespaces:
     matchLabels:
         tenant: true
                                                                  namespaces from which connections are
ingress:
- name: "default-denv"
                                                                  denied (L3 match conditions)
   action: "Denv"
  from:
   - namespaces: {}
                                                                 namespaces to which connections are
egress:
                                                                 denied (L3 match conditions)
- name: "default-denv"
   action: "Deny"
  to:
   - namespaces: {}
```

BANP has same API fields as ANP except its singleton thus has no priority; no PASS action

Pre Declaring ANPs that select Technical Assets









North America 2024

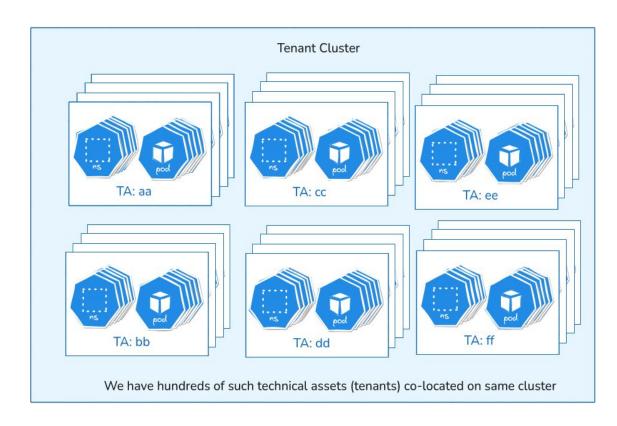
DEMO Time!



DEMO's Architecture Diagram: Setup



- 1 Multi-Tenant Cluster
- 10 Tenants (TAs)
- 5 Deployable Assets (DA) per tenant
- 1 Namespace per DA
- 1 BANP deployed
- 1 Platform ANP (DNS)
- 10 Tenant Scoped ANP's







North America 2024

ANP Caveats



One has to explicitly DENY the traffic, there are no implicit DENY's like with Network Policies

```
- name: "default-deny" # rule2
  action: "Deny"
  from:
   ...
```

For ingress, at max a "deny-all-traffic" means denying all traffic from all namespaces in your cluster. Unlike network policies there is no way to ask for "deny-everything"

```
- name: "default-deny" # rule2
  action: "Deny"
  from:
    - namespaces: {}
```



Caution when using empty selectors, specially in your subject - so you can potentially lock the cluster up

Specifying namespace and pod selectors to match on subjects and peers does NOT include host-networked pods

```
subject:
  namespaces: {}
```

```
to:
    pods:
    namespaceSelector:
    matchLabels:
        tenant: true
    podSelector:
    matchLabels:
    env: prod
```



nodes and networks peers are only available in Egress rules (subject to change in future depending on this NPEP)

Hence no need to worry about allowing ingress traffic from host network like is required for kubelet health checks when network policies are applied

```
to:
    - nodes:
    - networks:
```



Use "allow" and "deny" when you are sure you want to at any cost make that happen. Once you put these you cannot have finer controls using NetworkPolicies

```
metadata:
name: sensitive-workload-control
spec:
 priority: 30
 subject:
   namespaces:
     matchLabels:
       tenant: true
 ingress:
 - name: "pass-from-intra-cluster"
   action: "Pass"
   from:
   - namespaces:
       matchLabels:
         name: production
```



Don't create multiple admin network policies at the same priority that have overlapping rule definitions - in this case only of those rules will apply - which one gets applied is "implementation defined"

```
metadata:
 name: platform-team-controls
spec:
 priority: 44
 subject:
  namespaces:
     matchLabels:
       tenant: true
 ingress:
 - name: "allow-from-monitoring"
   action: "Allow"
  from:
   namespaces:
       matchLabels:
         name: monitoring
```

```
metadata:
name: sdn-team-controls
spec:
 priority: 44
 subject:
   namespaces:
     matchLabels:
       tenant: true
 ingress:
 - name: "deny-from-monitoring"
   action: "Deny"
   from:
   - namespaces:
       matchLabels:
         name: monitoring
```

BaselineAdminNetworkPolicy API Usage Caveats



Given BANP is a singleton use it wisely setting default deny's for all your workload
namespaces -> forces tenants to add explicit
allows and think about security

```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: BaselineAdminNetworkPolicy
metadata:
name: default
spec:
 subject:
   namespaces:
     matchLabels:
         tenant: true
 ingress:
 - name: "default-deny"
   action: "Deny"
  from:
   - namespaces: {}
 egress:
 - name: "default-deny"
   action: "Deny"
  to:
   - namespaces: {}
```





North America 2024

Lessons Learnt && What's Next?

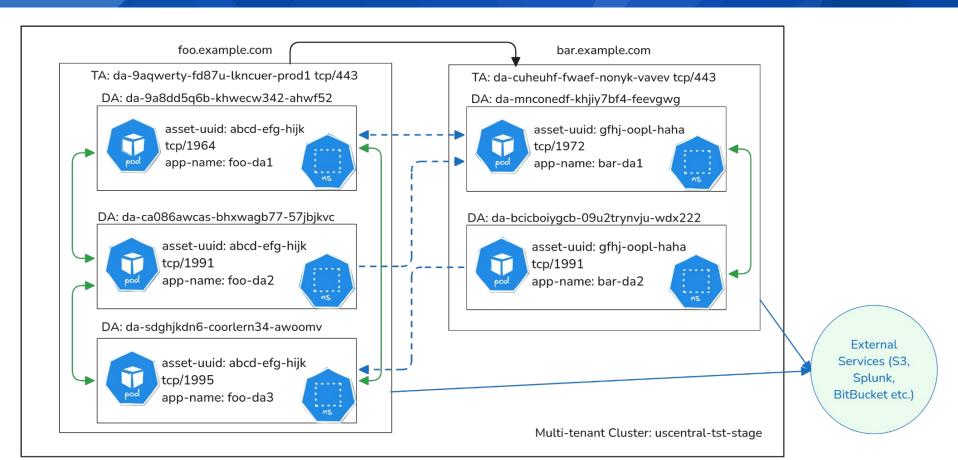
Lessons Learned



- Microsegmentation Controls: Designing connections requires collaboration across multiple teams, including the platform team.
- PCI DSS Compliance: This is more than microsegmentation is about becoming secure by default.
- Collaborative Design: Adopting open standards and collaborating with the community, specifically the sig-network-policy-api WG.
- **Start Simple**: Begin with a minimalistic approach to reduce complexity and yet still leave room to build in finer controls as needed.

Future with AdminNetworkPolicy + NetworkPolicy









North America 2024

Questions?

Feedback on the talk is welcome!->









KubeCon CloudNativeCon

North America 2024

Leveraging Priority Field between Platform & SDN Teams



- Platform team creates the ANPs for defining connections to infrastructure components from the microservices
 - Can use a range of 20-40

- SDN team is responsible for defining the connections between microservices
 - Can use a range of 40-60

This way even upon accidental overlap platform team's rules will be evaluated first

DEMO



- We have 10 tenants each with 5 Deployable Assets
- Tenant AA & Tenant BB for the demo
- We need a BANP {} {} to isolate tenants
- After BANP no inter or intra tenant connectivity
 - TAAA -> TAAA Backend(s) fails
- Define ANP's for each Tenant (Allow) →
 - TAAA -> TAAA Backend(s) still fails, because DNS is missing
 - Deploy an Infra ANP for the cluster with DNS ANP
 - TAAA -> TAAA Backend(s) SUCCESS!
 - TAAA -> TABB Backend(s) BLOCKED!
- Kapi server still doesn't work because we have BANP denying it