


# What agent to trust with your k8s:



HENRIK REXED  
CLOUDNATIVE ADVOCATE

# Henrik Rexed

---




## Cloud Native Advocate

- 15+ years of Performance engineering
- Owner of : IsitObservable



Producer of : Perfbytes






## DISCLAIMER

---

- No Spartans, Eagles or Shields were harmed in the making of this presentation
- The intention behind this talk track is not to assign blame to any CNCF project.
- This session is made to help the community in choosing their runtime security agent.










# KubeOlympics 2024





# KubeOlympics 2024



# The Athletes



**Falco**



**Tetragon**



**Kubearmor**



**Tracee**


# If you stay with me you will ...

- What we expect from a runtime security agent
  - Compare each solution in :
    - Desing Experience
    - The Components required
    - The Observability
- See various Benchmarking results
- Recommendation on which agents needs to used under specific conditions





# THE RULES







# Performance




# Observability



# Components




# Filter



# Capture


# **STAGE 1: COMPONENTS**






# Components

- A security agent usually relies on ebpf to collect the kernel or user event of our environment
- In k8S , the security agent has usually at least :
  - A daemonset deploying the ebpf probe on each node of our cluster




## Components required


- Falco




- Tetragon




- KubeArmor



- Tracee




## **STAGE 2: CAPTURE & FILTER**



# Capture

- Be able to capture events from our k8S cluster related to :
  - Process launched/ finished
  - File Access Read/Write
  - Networking details
  - Change of privileges
  - ...etc.
- Get the details attached to the event:
  - Path: executable path, file path...
  - Process details: pid, process name ,
  - User: userid, groupid
  - Capabilities
  - Container: image , name
  - K8s : pod name, namespace, ..etc





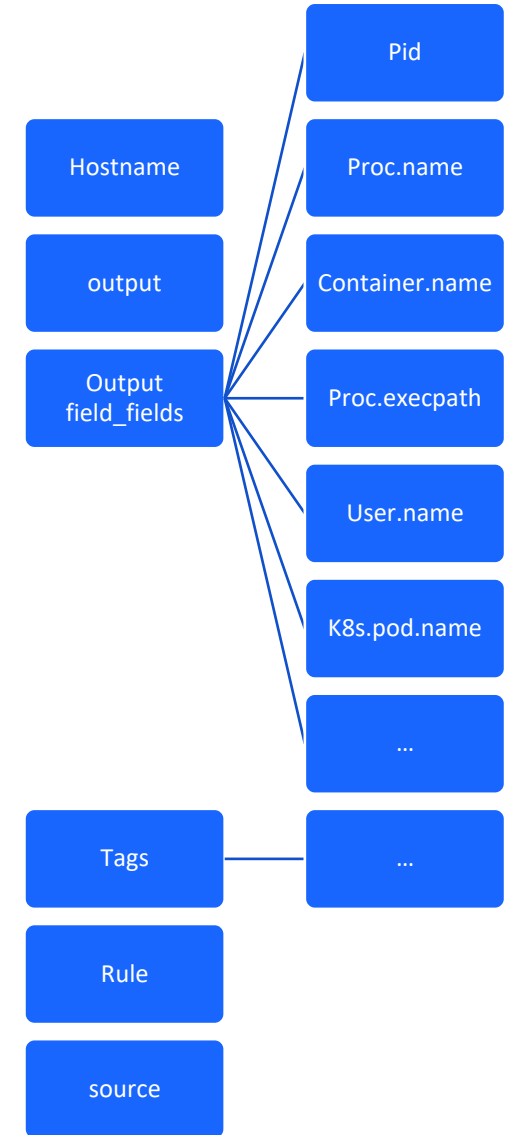
# Filter

- Filtering means having the option to create a policy filtering the event on :
  - Process properties
  - User details
  - K8S metadata
  - File details
  - Syscall/ kernel functions or event
- Have Predefined Rules
- Be able to react



# Falco Capture

- Falco is rule engine producing logs if the falco event is matching a Falco rule
- Falco receive event from :
  - Falco kernel agent
  - External plugins
- Falco provides a SDK allowing us to build our custom Plug
- Falco agents captures :
  - Process details
  - Syscall
  - Tracepoints
  - File






## Falco rule

- Falco starts by loading the rule files that would be used to generate the events. By default, Falco provide a default sets of rule.
- Falco provides “fields” helping us to define our filtering rule based on:
  - Event information
  - Proces
  - file,
  - Syscalls
- Having a global rule allow to to reuse filtering rules or conditions with macros

```
macro: access_file
condition: evt.type=open

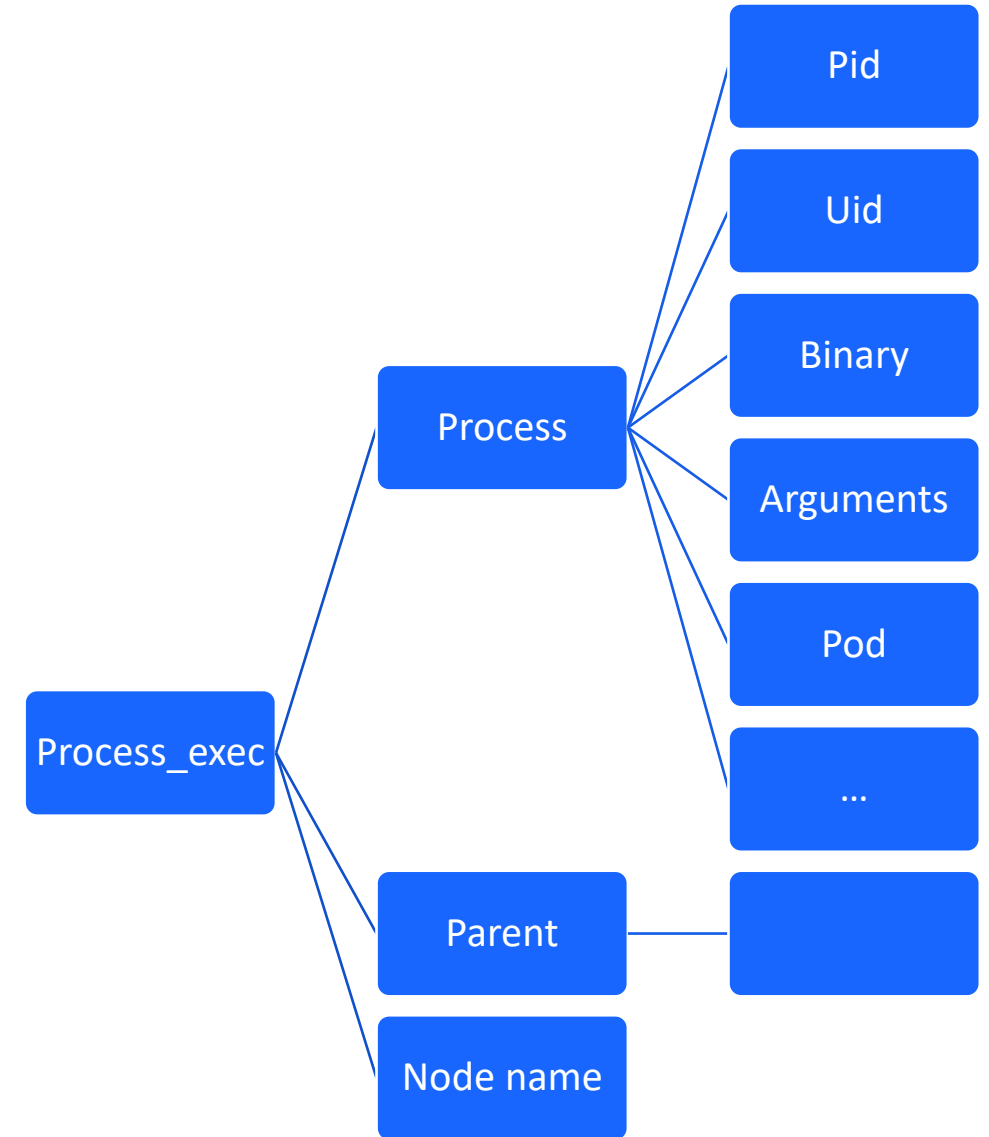
- rule: program_accesses_file
  desc: track whenever a set of programs opens a file
  condition: (access_file) and proc.name in (cat, ls)
  output: a tracked program opened a file
  (user=%user.name command=%proc.cmdline file=%fd.name)
  priority: INFO

- rule: test_rule
  desc: test rule description
  condition: evt.type = close
  output: user=%user.name command=%proc.cmdline
file=%fd.name
  priority: INFO
  enabled: false
```



# Tetragon capture


- Tetragon will by default produce events related to :
  - Process execution
  - Process exit
- If you deploy Policies, it will produce events matching our policy rule
- Policy required to defined a hook point to our system using:
  - Kprobe
  - Tracepoints
  - Uprobe
  - lsm



# Tetragon policy

- Building a policy means defining creating a TracingPolicy defining:
  - the right hookpoint
  - The index argument we would like to extract
  - The data type of the argument
- Filtering in Tetragon means applying a selector on :
  - The Args extracted
  - The process, file ...etc details
  - The user
  - The k8S metadata
- TracingPolicy is also allowing us to define how to react on the event ( block, Audit....etc)

```
apiVersion: cilium.io/v1alpha1
kind: TracingPolicy
metadata:
  name: "k8s-api-calls"
spec:
  kprobes:
    - call: "tcp_connect"
      syscall: false
      args:
        - index: 0
          type: "sock"
  selectors:
    - matchArgs:
        - index: 0
          operator: "DAddr"
          values:
            - "10.43.0.1"
            - "10.1.0.0/20"
    matchBinaries:
      - operator: "NotIn"
        values:
          - "/usr/bin/rancher"
          - "/usr/bin/dumb-init"
```



# Kubearmor capture

- KubeArmor will only produces events maching:
  - The Policiy deployed
  - Or the K8S objects having the right annotations
- KubeArmor will capture :
  - Process
  - File
  - Network
  - Capatbilities
  - Syscalll
- KubeArmor Operator provides more feature to manage with the help of annotations
  - The type of events to report ( file, process, network)
  - The default posture on how to react

Clustername

Hostname

Namesapce

Podname

Labels

ParentProcessName

ProcessName

HostPPID

HostPID

PPID

Source

Operation

Data

...




# KubeArmor Policy

- KubeArmor policies are defined by building a :
  - KubeArmorClusterPolicy
  - KubeArmorPolicy
  - KubeArmorHostPolicy
- The policy defines :
  - The Tags
  - The message of the event
  - The Selector to filter to a specific namespace or workload
  - The rule for the event type ( process, file, network..Etc)
  - And the action : Block, Audit, Allow

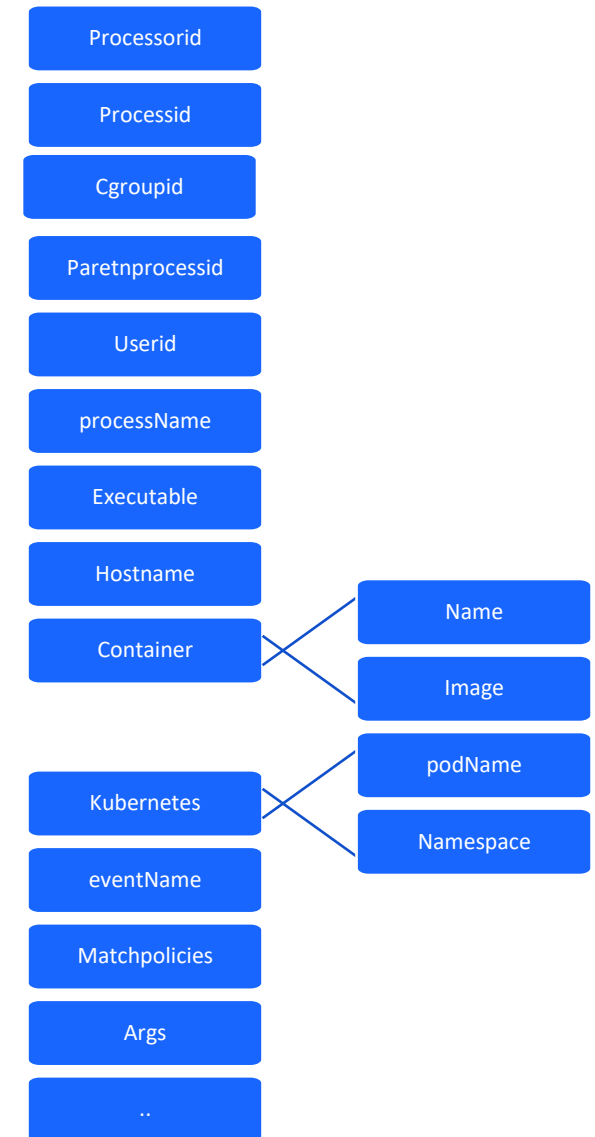
```
apiVersion: security.kubearmor.com/v1
kind: KubeArmorClusterPolicy
metadata:
  name: ksp-nist-remote-access-and-control
```

```
spec:
  tags: ["NIST","system","ksp","AC-17(1)"]
  message: "warning! someone tried to access and control"
  selector:
    matchExpressions:
      - key: namespace
        operator: NotIn
        values:
          - kube-system
          - istio-system
  process:
    severity: 4
    matchPaths:
      - path: /usr/bin/ssh
      - path: /etc/ssh
    action: Audit
```



# Tracee capture

- Tracee will only produce events matching the policies deployed
- Tracee will capture :
  - Security event
  - Network details
  - Anything using a given list of syscalls ( process, file...Etc)
- The type of events detected will rely on a set of “signatures” defining
  - What to capture
  - What to decode
- Tracee provides a “sdk” helping us to build our own signature






# Tracee Policy

- Tracee policies are defined by configuring a :
  - Policy
- Tracee provides a default policy based on the security signature
- The policy defines :
  - The scope
  - The set of rules defined by a event name and filters


```
apiVersion: tracee.aquasec.com/v1beta1
kind: Policy
metadata:
  name: sample-data-filter
  annotations:
    description: sample data filter
spec:
  scope:
    - global
  rules:
    - event: vfs_read
      filters:
        - data.pathname=/etc/*
        - data.pathname=/etc/fstab
        - data.pathname=/etc/crontab
        - data.pathname=/etc/hosts
        - data.pathname=/etc/hosts.allow
```



# Summary

				
Capture	X	X	X	X
Filter	X	X	X	X
Default Policy	X			X
Ability to react		X	X	
Customizable	X	X		X

# **STAGE 3: OBSERVABILITY**







# Observability

- Extend observability with the event produced by :
  - Collecting the logs produced
  - Create parsing rule
- Limit the number of fields exported
- Limit the events by enabling throttling
- Be able to report health metrics related to :
  - The policy deployed
  - The various components required to run the runtime agent

# Falco


- Falco events would be available in the Falco agent logs
- The event structure highly depends on the rule created and the syscall used.
- To simplify the log collection we usually rely on FalcoSidekick that provides a large number of integrations
- Falco provides a metric server reporting :
  - Metrics related to the rule
  - Actual health of Falco agent







# Tetragon

- Tetragon expose all the events directly in the logs of the agent
- Tetragon expose 2 Prometheus exporter:
  - Agent: sharing details on the policies deployed and the health
  - Operator: Sharing the health of the agent




# Kubearmor

- Kubearmor will not produce events in the logs from the agents
- To collect events you will either need to :
  - Enable the logging option on the Relay Server
  - Use the Kubearmor receiver that will collect the logs from the agent. This receiver is currently only compatible with v0.96 of the collector
- Kubearmor is not providing any metrics








# Tracee


- Tracee is producing the event directly at the logs of the agent (default settings)
- Tracee can also push the logs using :
  - the FluentForward protocol
  - A webhook endpoint
- The Tracee agent expose Prometheus metrics related to the events produced and the errors.



# Summary

				
Collecting events	X	X	X	X
Extend observability	X	X	X	X
Rule metrics	X	X		X
Heath metrics	X	X		X


# **STAGE 4: PERFORMANCE**



## Kpi that we want to measure


- Measure the latency added in our application


Latency



- Cpu usage
- Memory usage

Resource  
usage





# What type of policies we would apply

- Using the default policies

Falco



- k8S api calls
- Service account files
- Sensitive files
- Egress communication
- Install tools
- Network activities
- Process spawned

Tetragon



- Block unauthorized binaries
- Sensitive files
- Audit write in sensitive folder
- Suspicious Network tools
- External access


Kubearmor



- Sensitive files
- New containers created
- Process spawned
- Security signature
- Read write access from non root user

Tracee





## The various tests executed



No Agent



Falco



Tetragon with default events



Tetragon with default events & policies



KubeArmor with policies and no events



KubeArmor with policies and events




KubeArmor no policies but with events



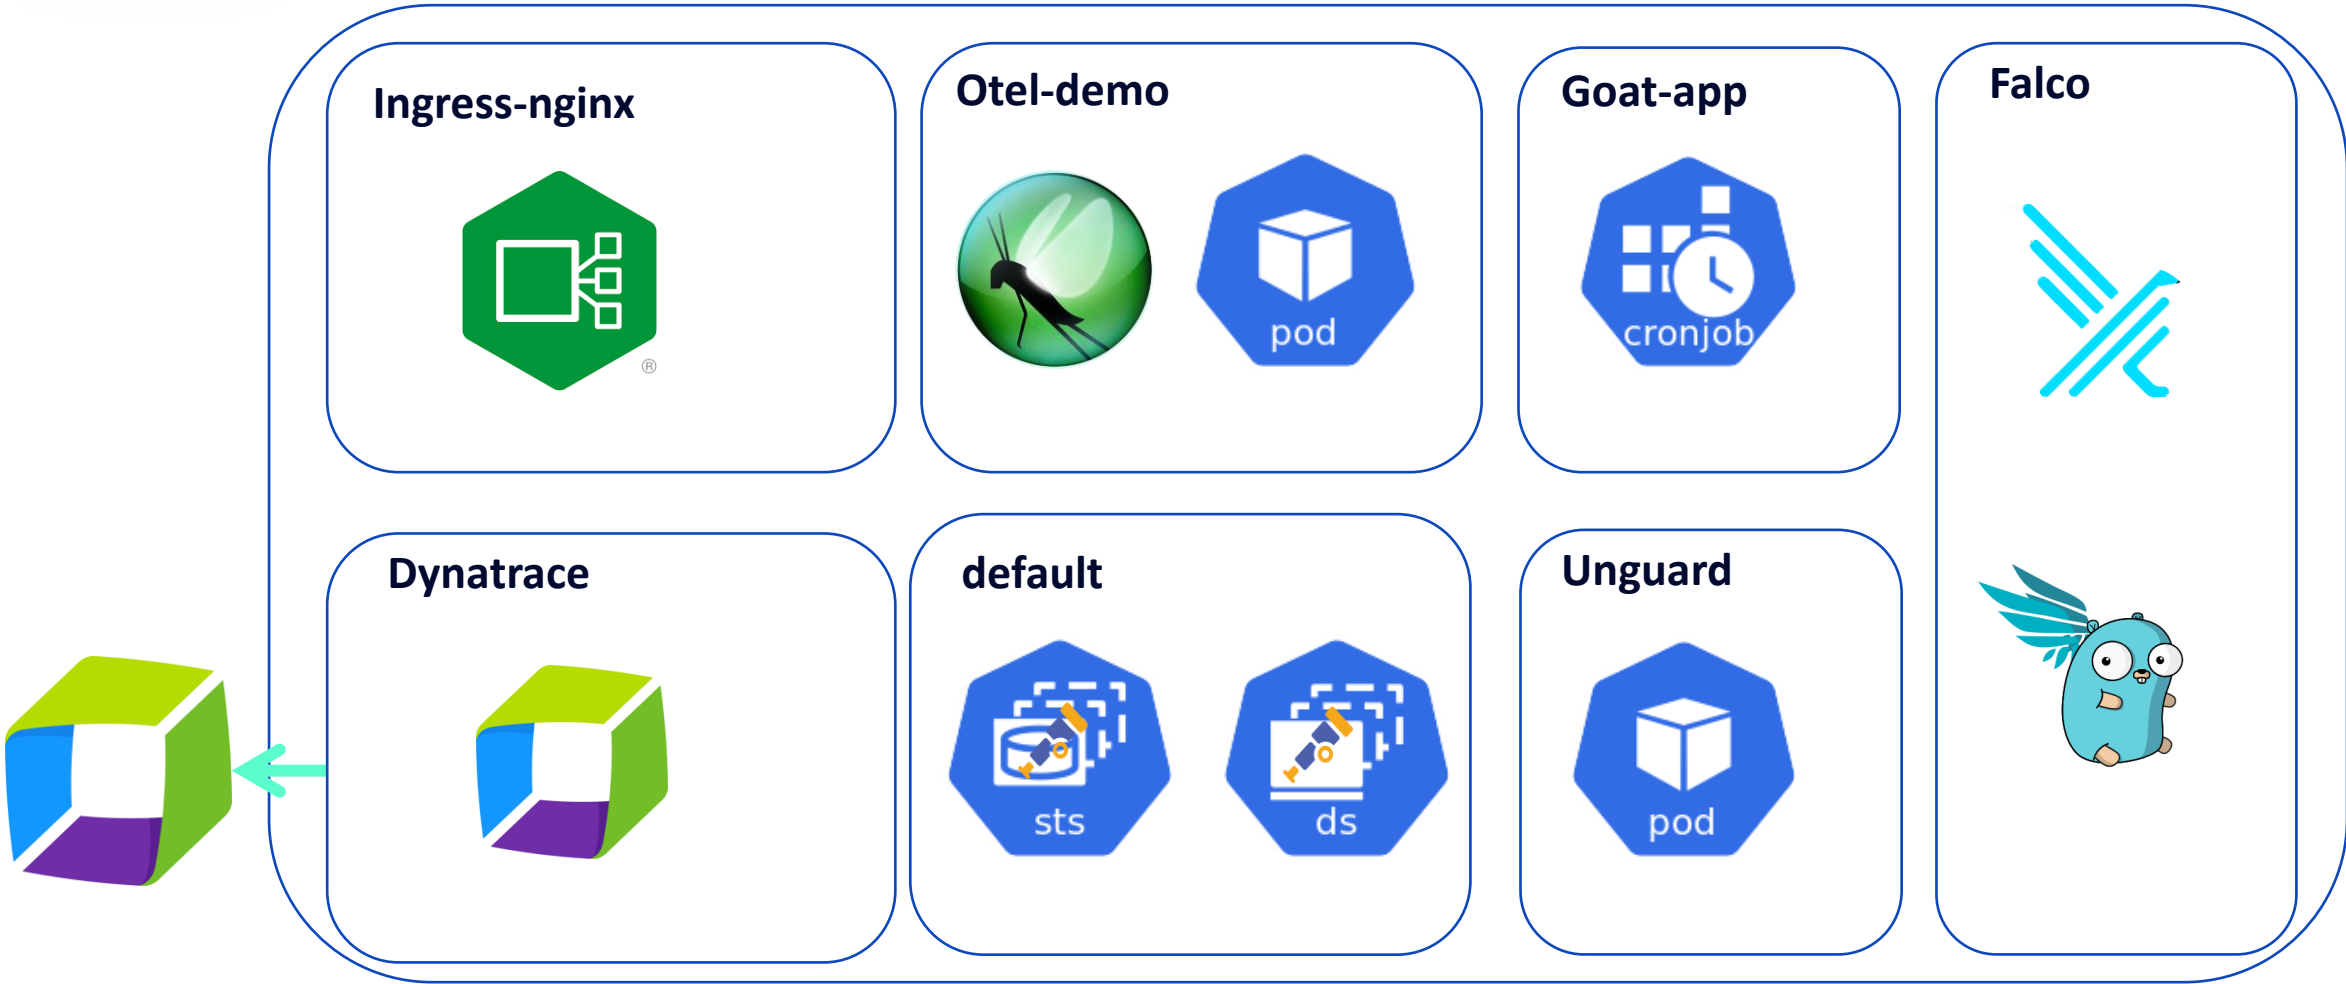
Tracee with default policy




Tracee with default policy and custom policies



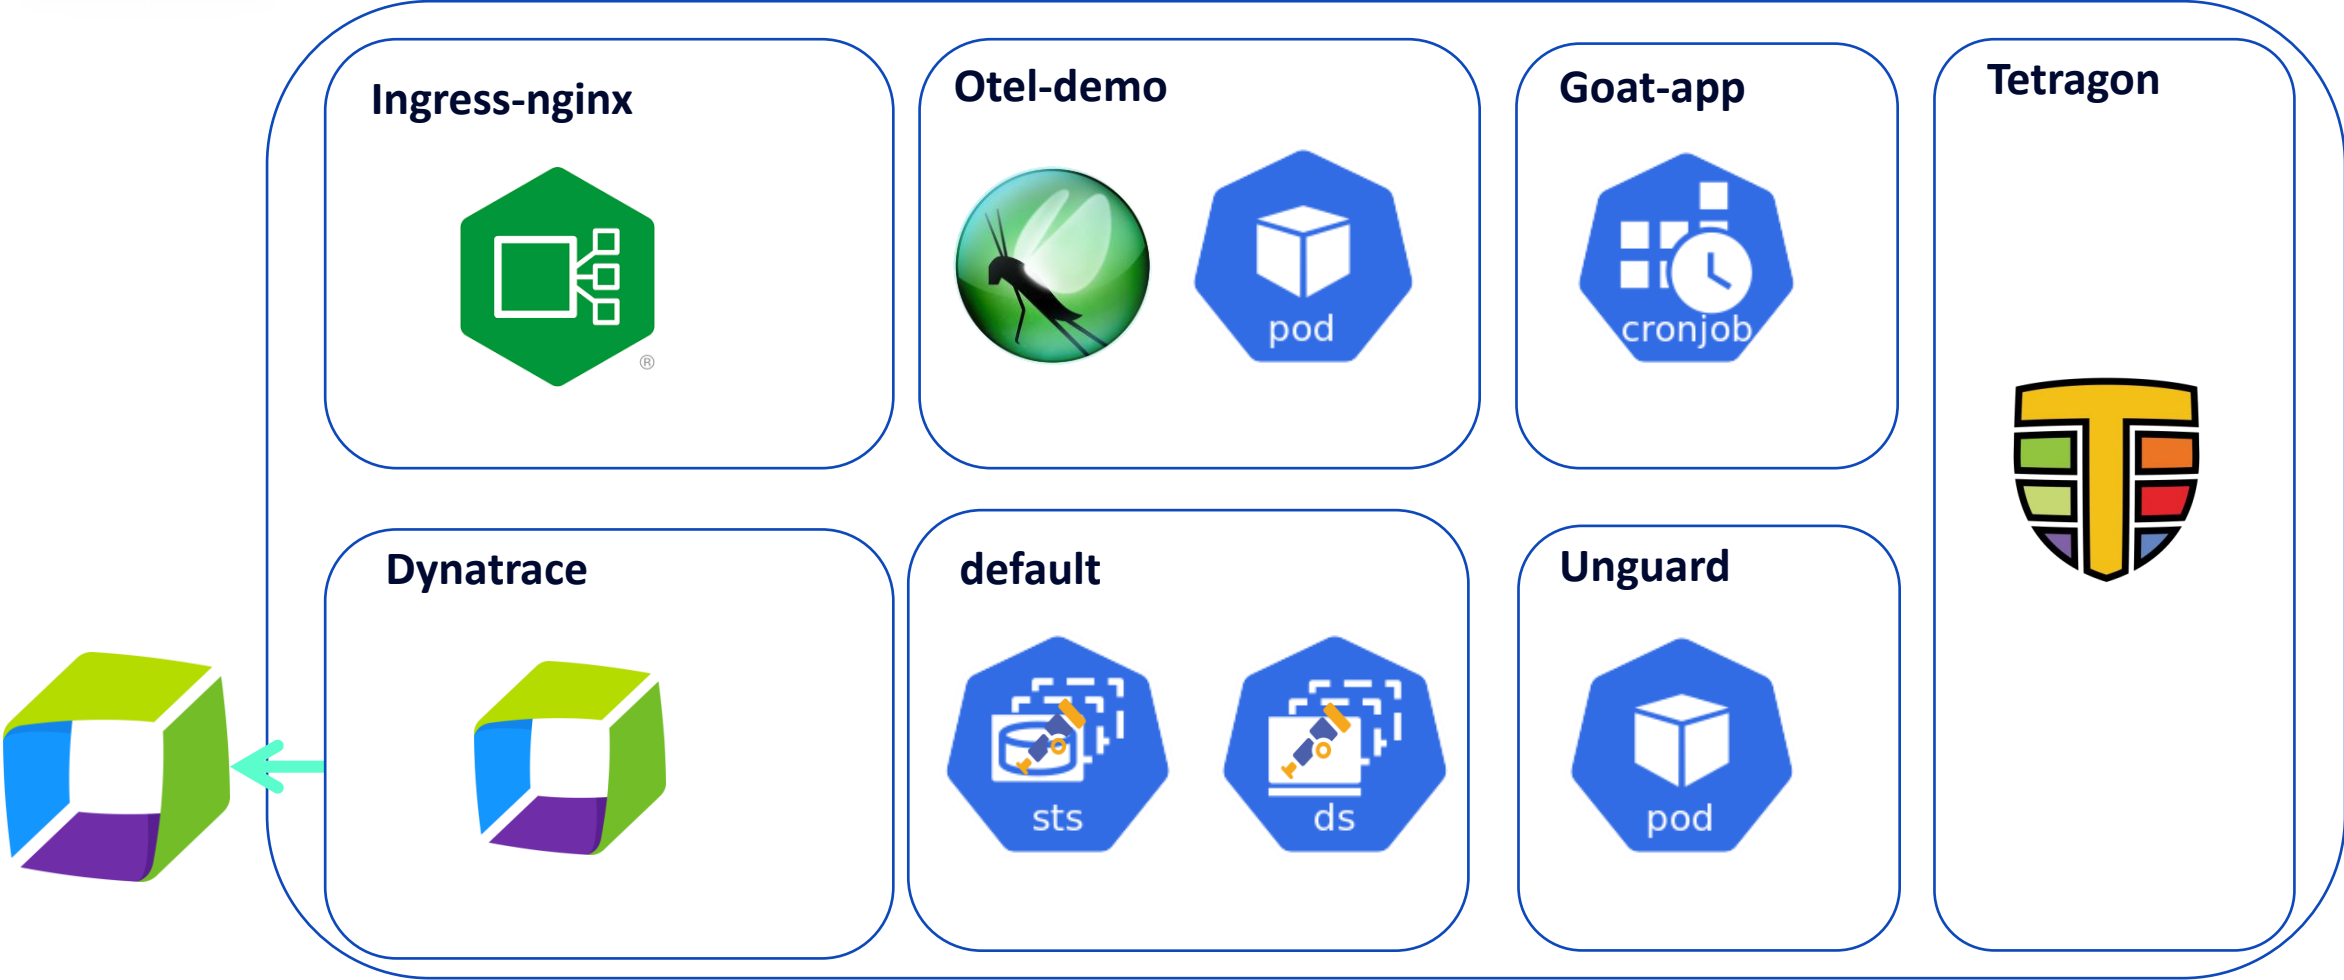
# Falco Architecture






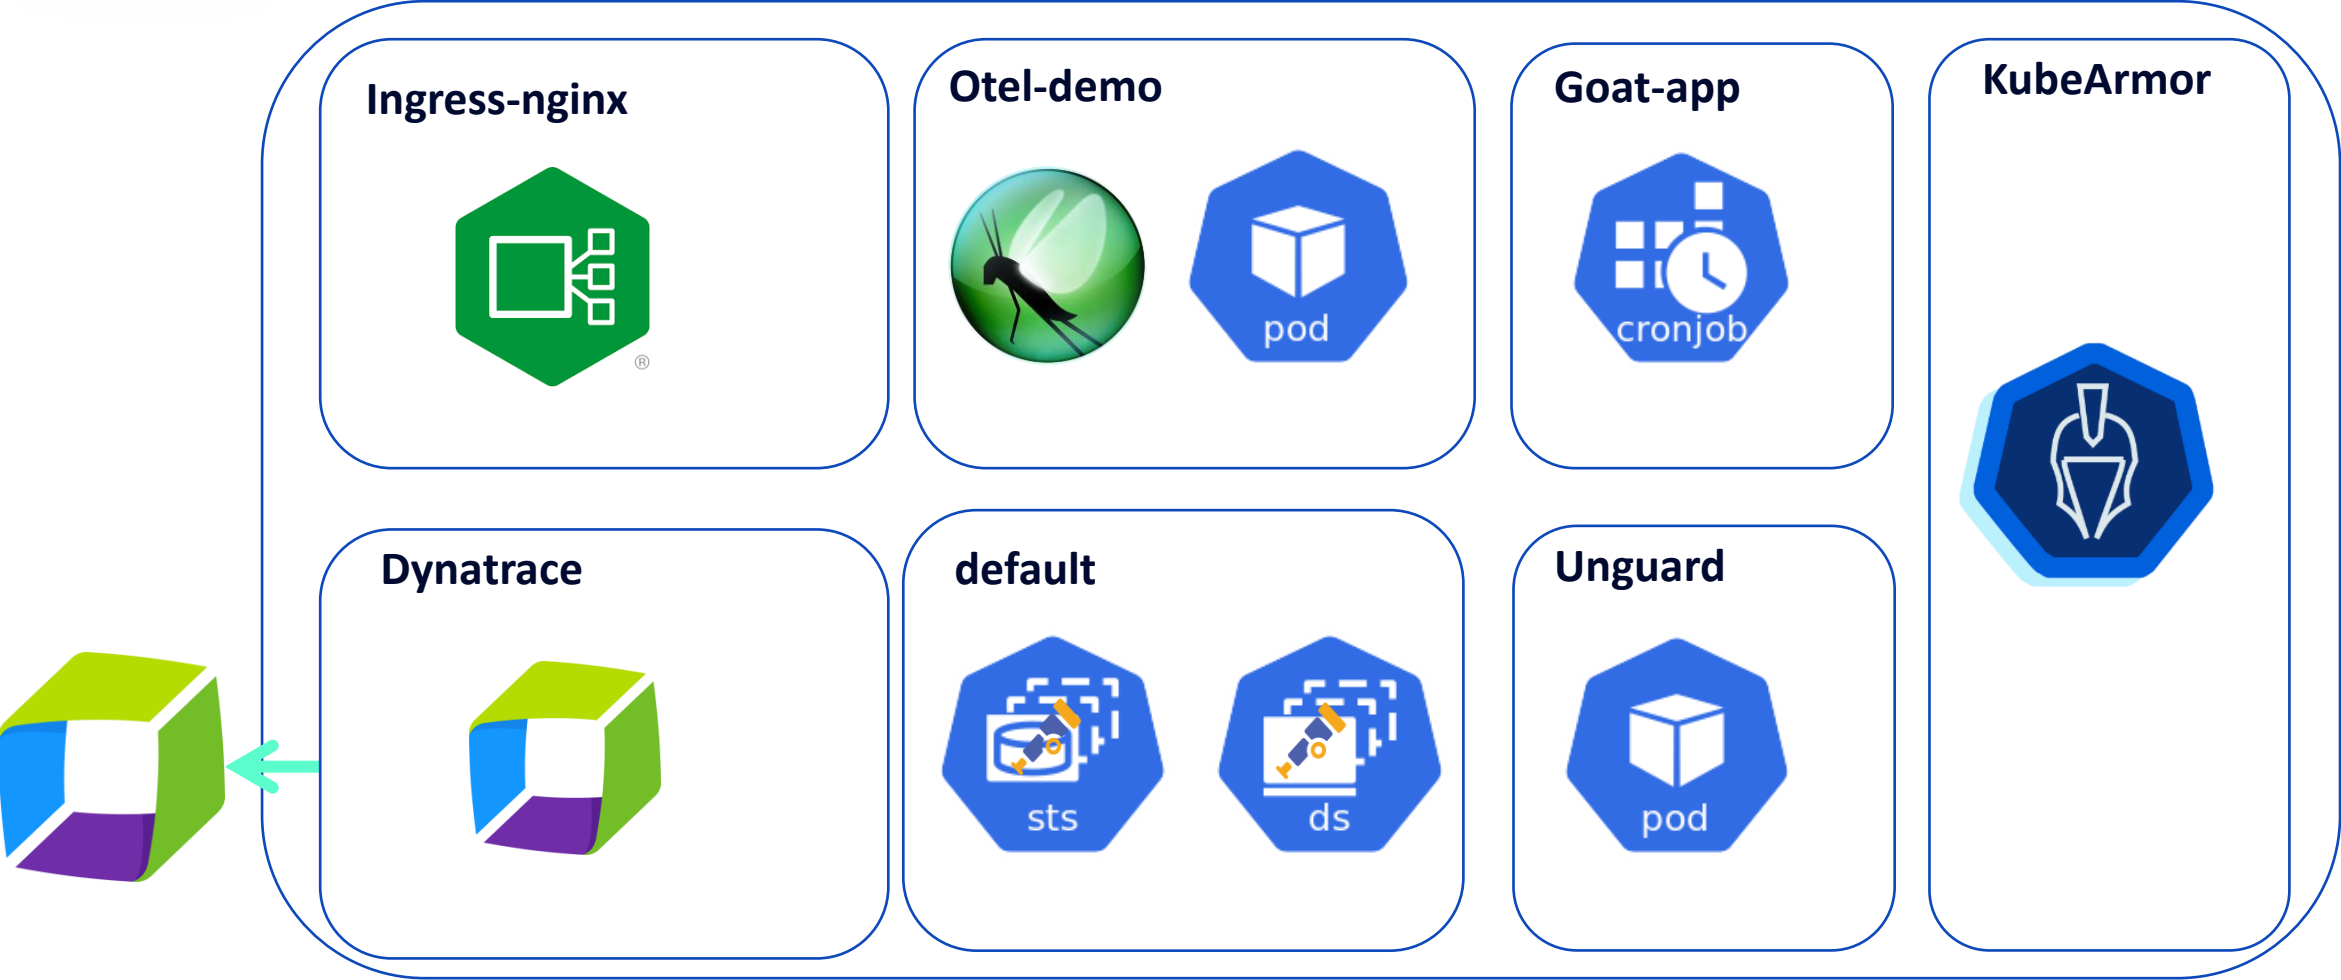



# Tetragon Architecture



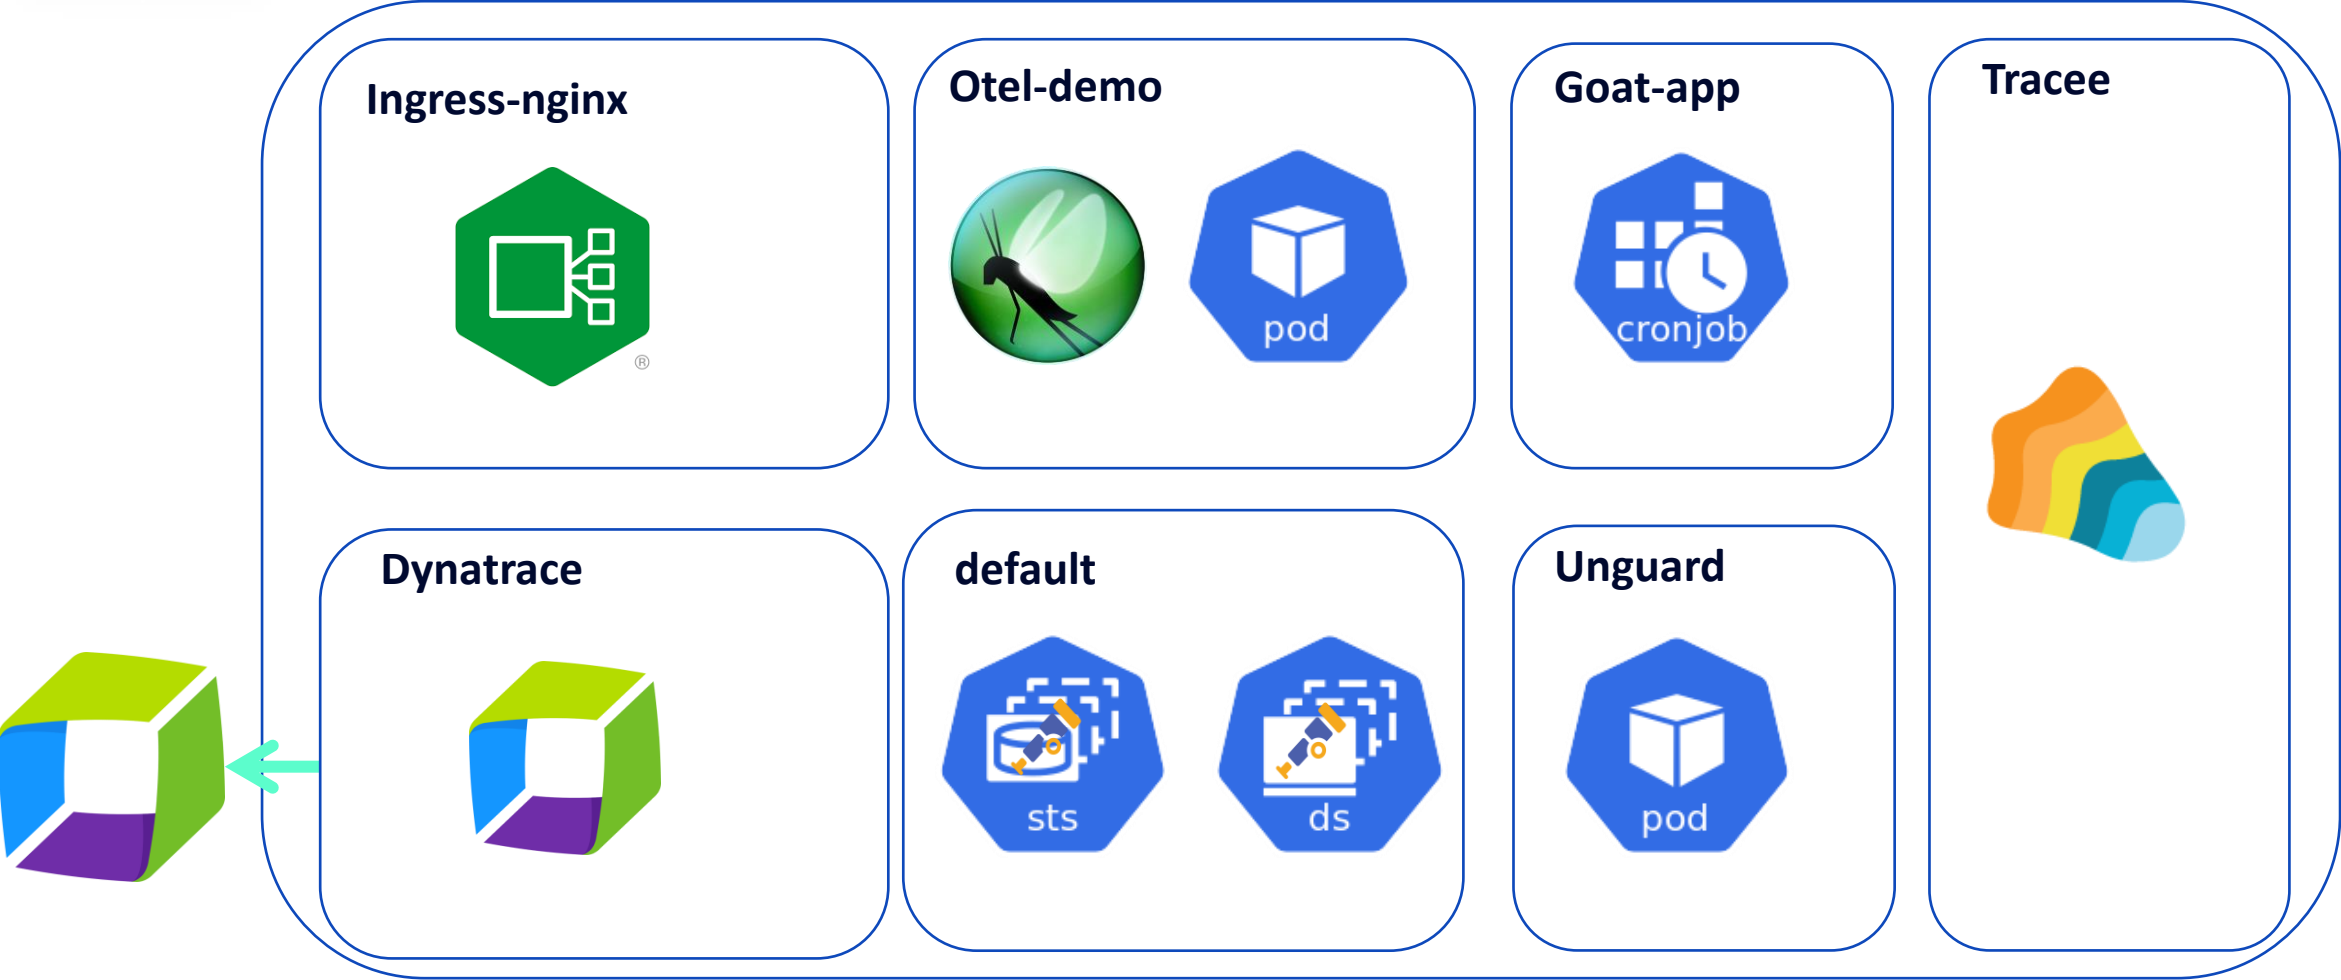



# KubeArmor Architecture






# Tracee Architecture








# Constant Load with no policies



Falco	+10,84ms
Tetragon	+ 5,03ms
KubeArmor	+15,73ms
Tracee	+110,79ms



# Constant Load with policies






<b>Falco</b>	+ 10,84ms
--------------	-----------

<b>Tetragon</b>	+26,13ms
-----------------	----------


<b>KubeArmor</b>	+176,62ms
------------------	-----------

<b>Tracee</b>	+114,23ms
---------------	-----------








# Constant Load with policies & events







Falco	+ 10,84ms
Tetragon	+26,13ms
KubeArmor	+477,34ms
Tracee	+114,23ms






# Summary

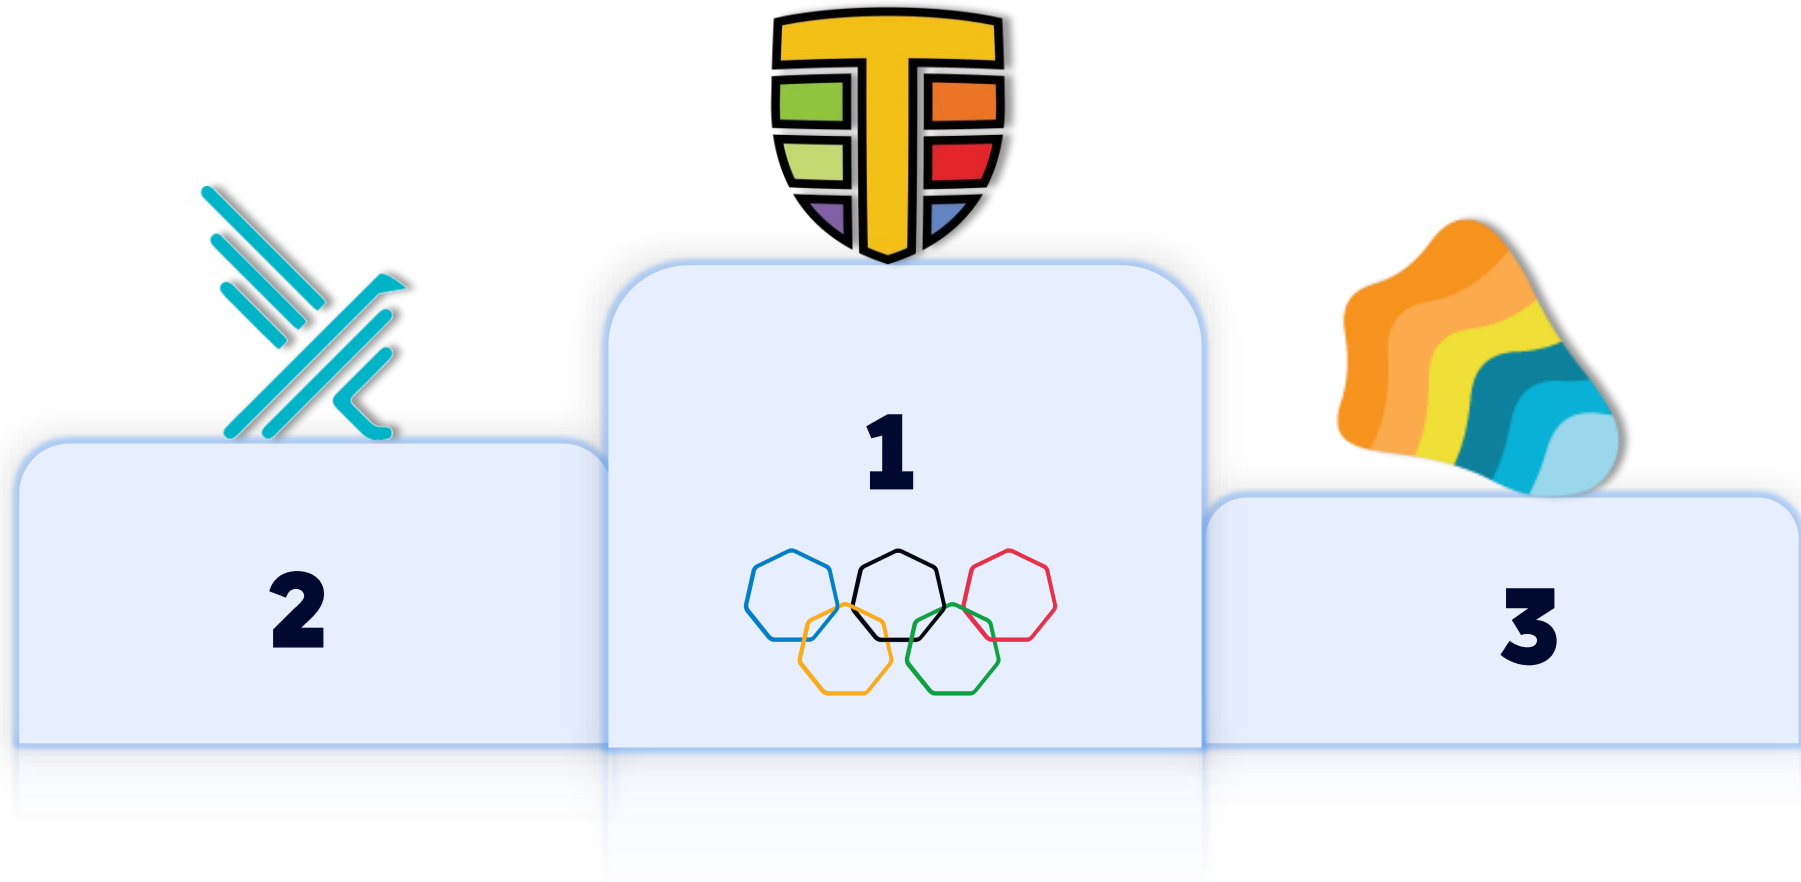
				
Latency	× × ×	× × ×	×	×
CPU	× ×	× × ×	×	×
memory	× ×	× ×	× × ×	×


# Conclusion





# Conclusion







# Is it observable

- Looking for educational content on Observability , Checkout the YouTube Channel :

## Is It Observable





**Thank You**

