KubeCon | CloudNativeCon

North America 2024

Best Practices for Deploying LLM Inference, RAG and Fine Tuning Pipelines in K8s

Shiva Krishna Merla, NVIDIA
Meenakshi Kaushik, NVIDIA

Agenda

- Overview
- Inference and Fine-Tuning
- Looking forward
- Conclusion

Overview

Generative AI Fuels Rise of Sovereign Clouds

Data Sovereignty and Compliance

Tailored AI Infrastructure

Enhanced Security and Privacy

Enterprise Benefits from Customer LLM

Custom Models Provide Insightful Responses

Base Models Generate Generic Responses

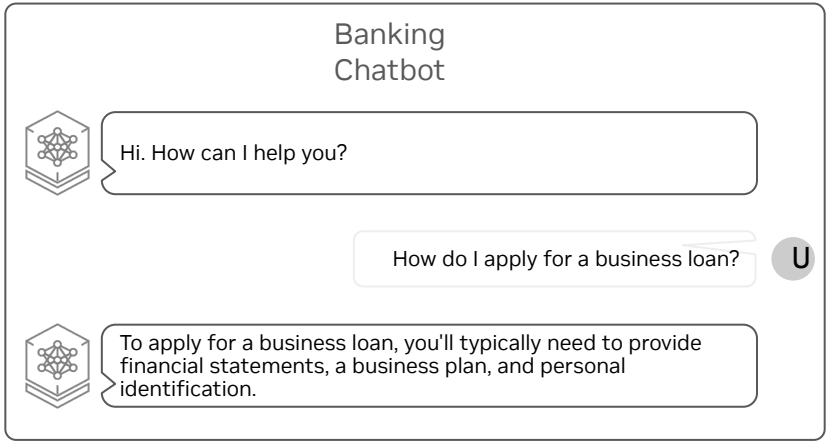

Custom Models Provide Business-Specific Answers

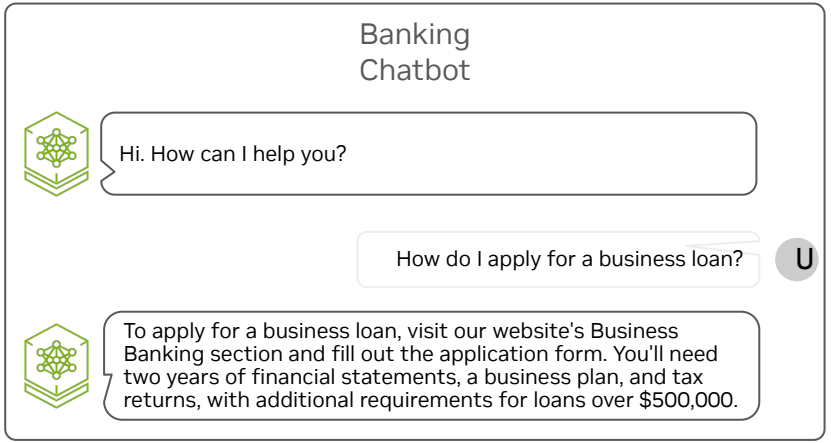

Inference

Inference Servers
Fine Tuning Jobs
Dependencies

Inference, RAG

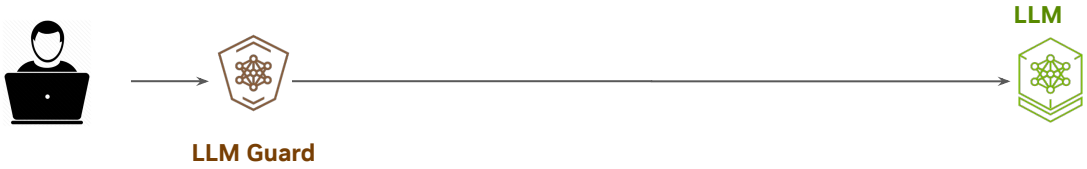

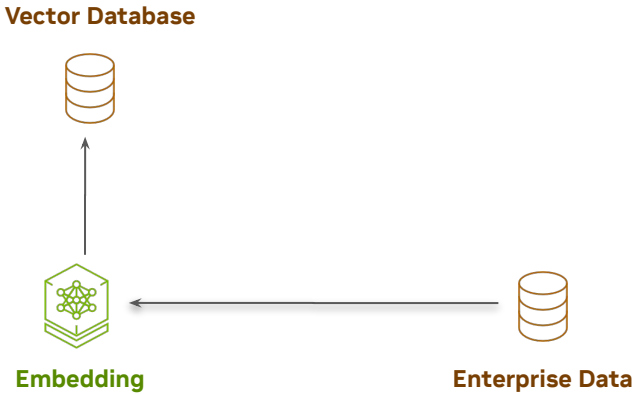

Inference Servers
Fine Tuning Jobs
Dependencies

Inference, RAG

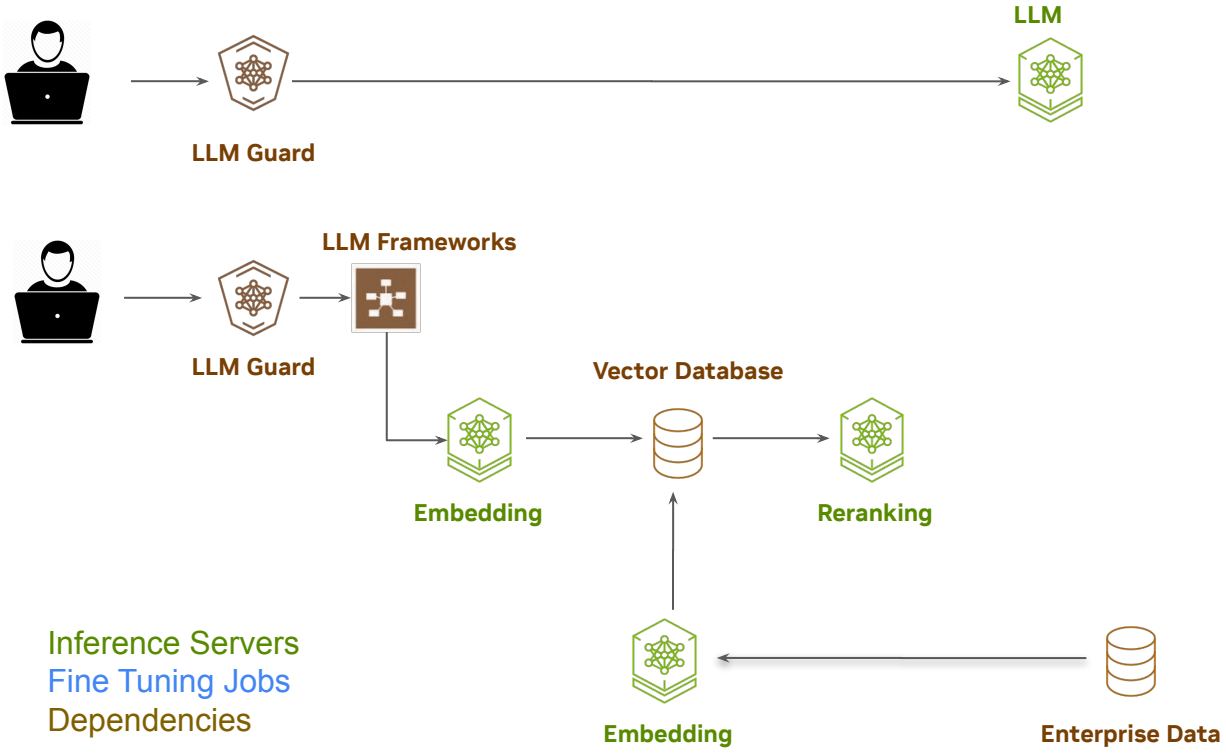

Inference, RAG

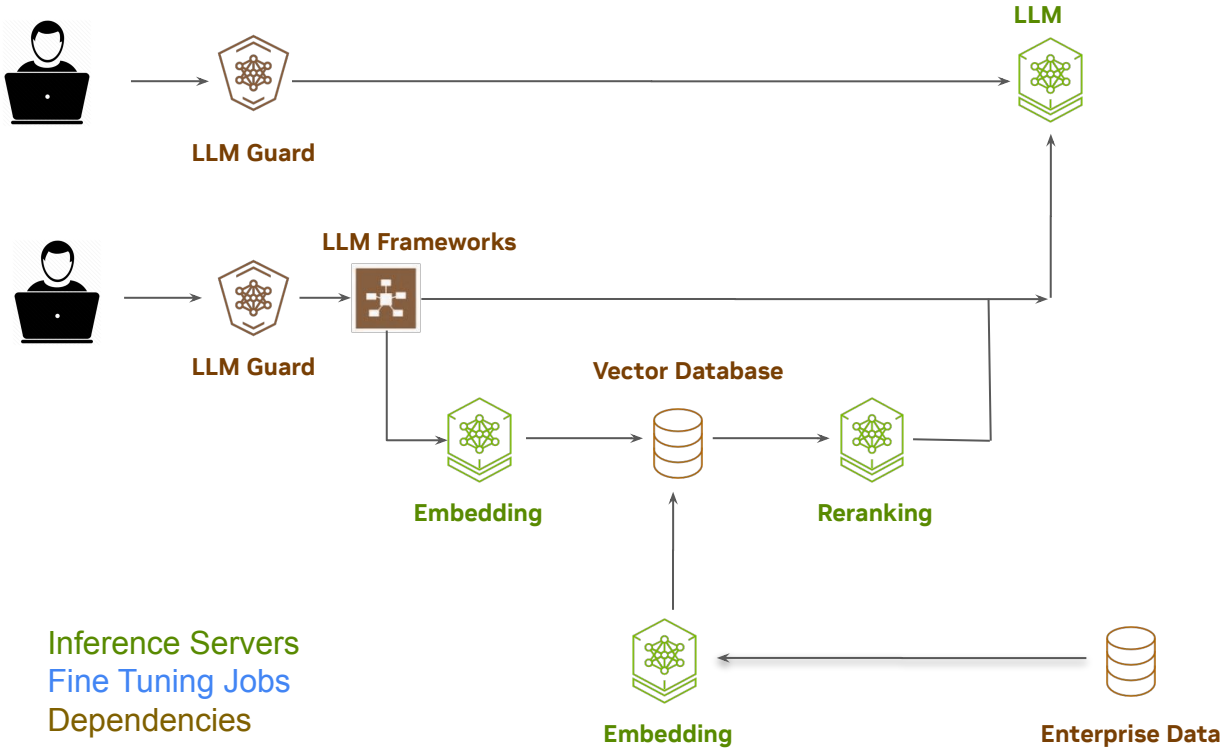

Inference Servers
Fine Tuning Jobs
Dependencies

Inference, RAG and Fine Tuning Pipelines

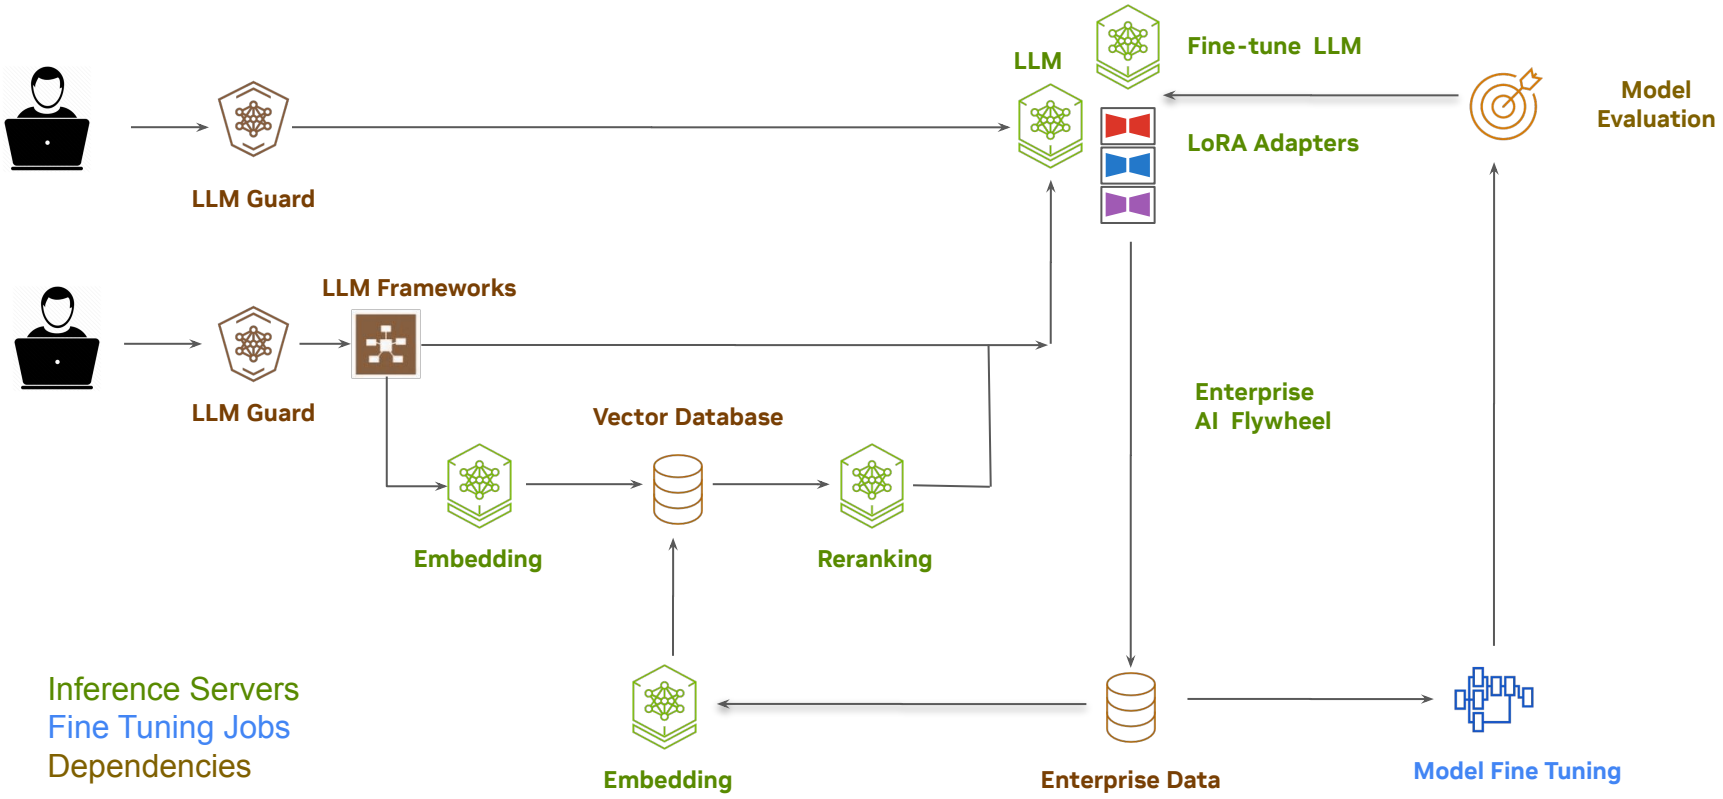

Current Landscape

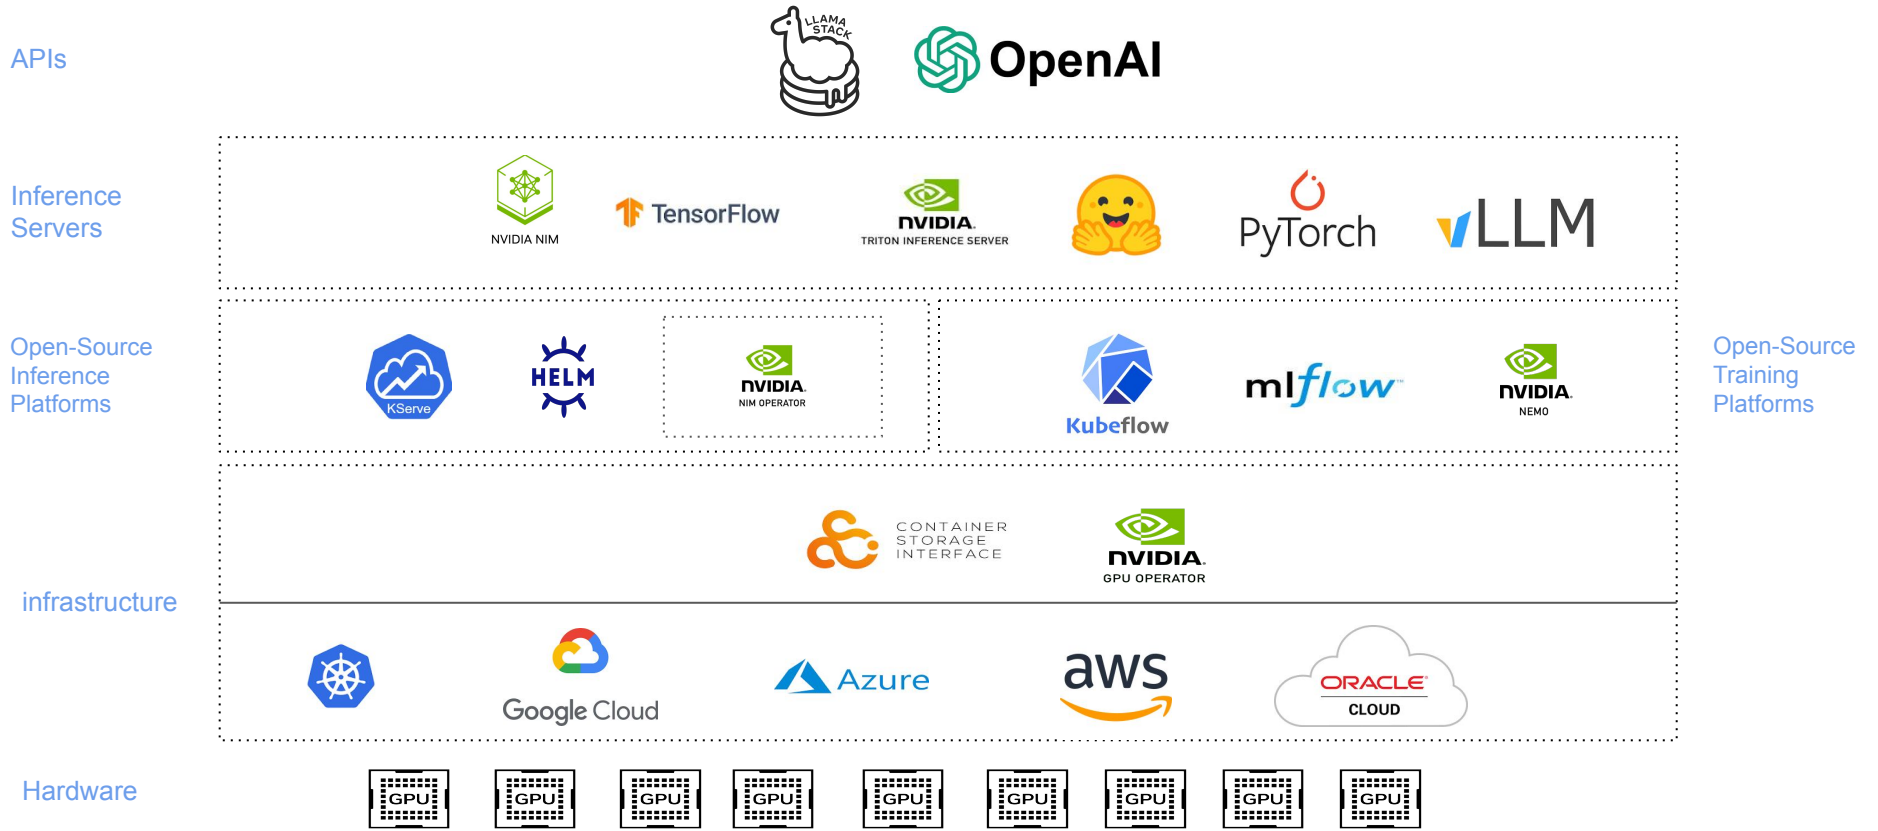

Inference and Fine Tuning

An Example Inference Server

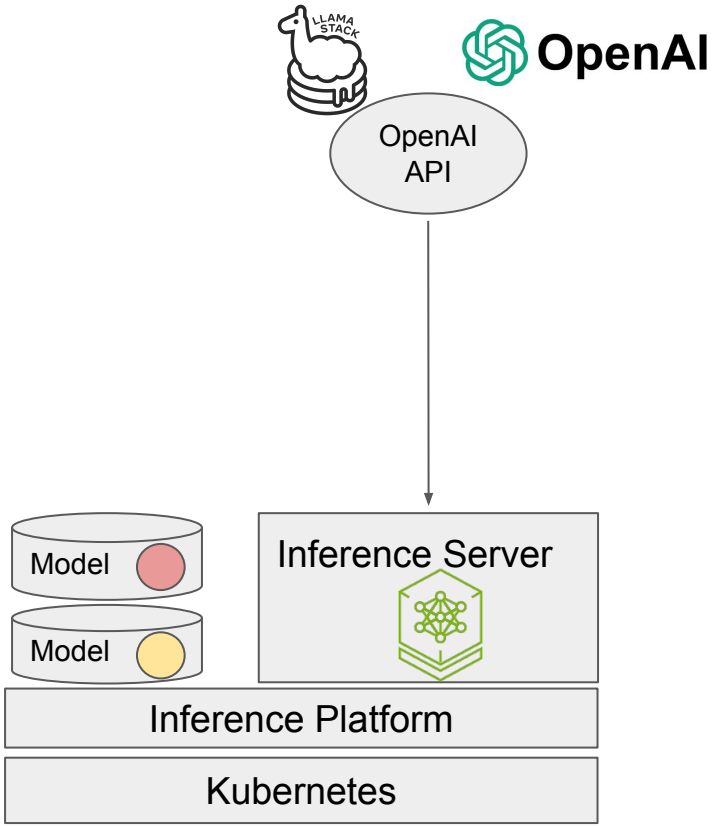

An Example Inference Server

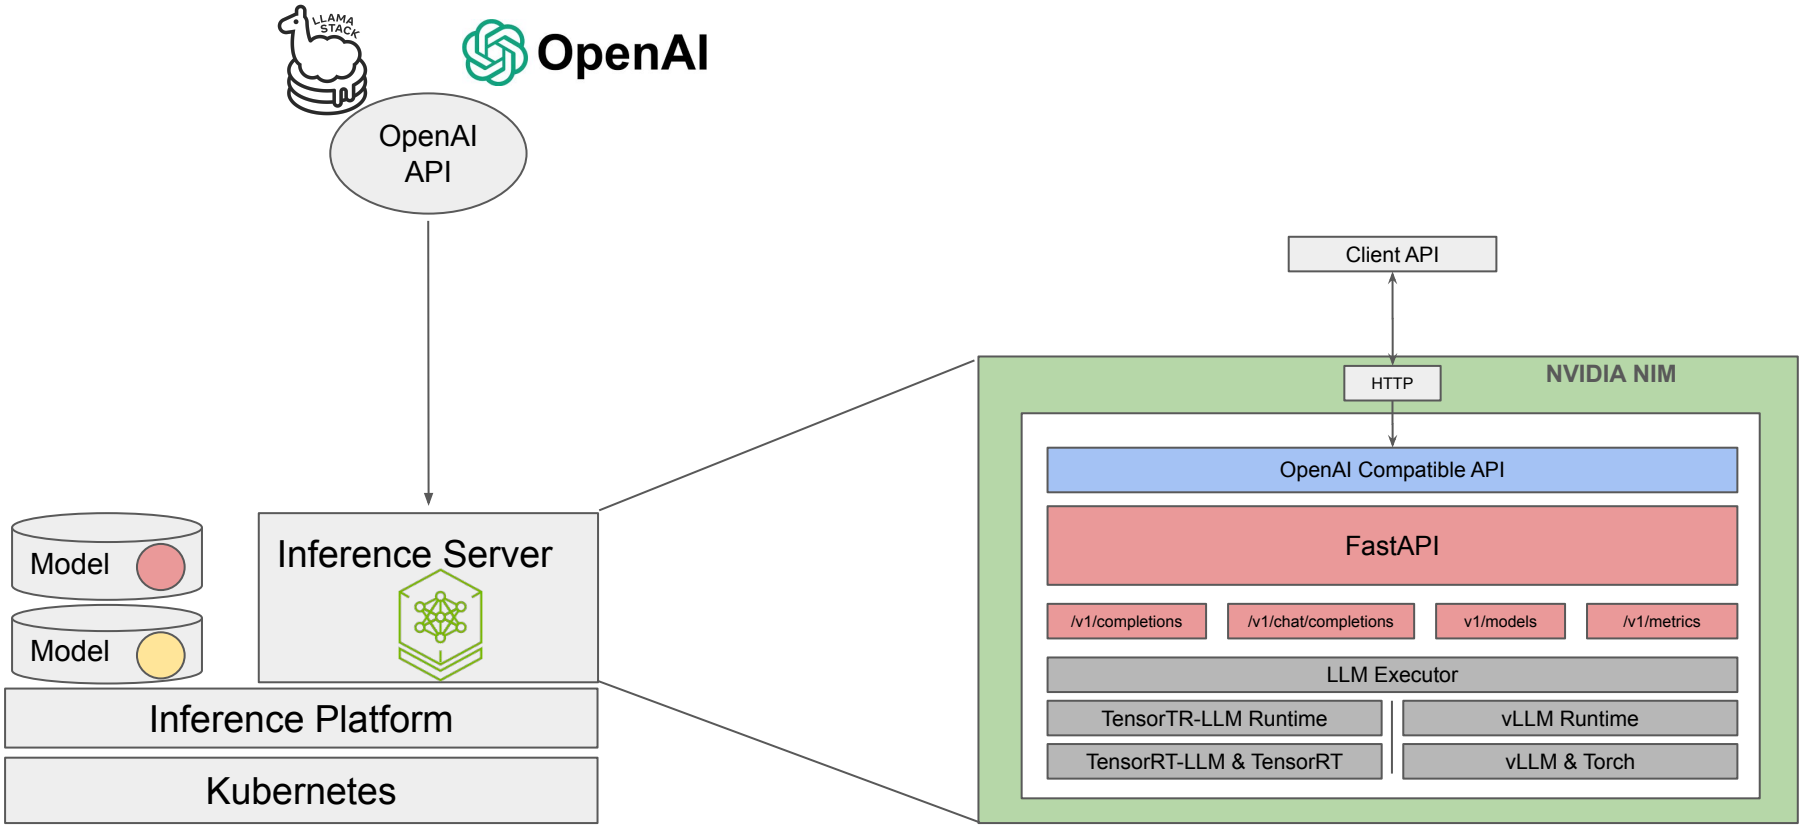

Typical Model Serving Pipeline

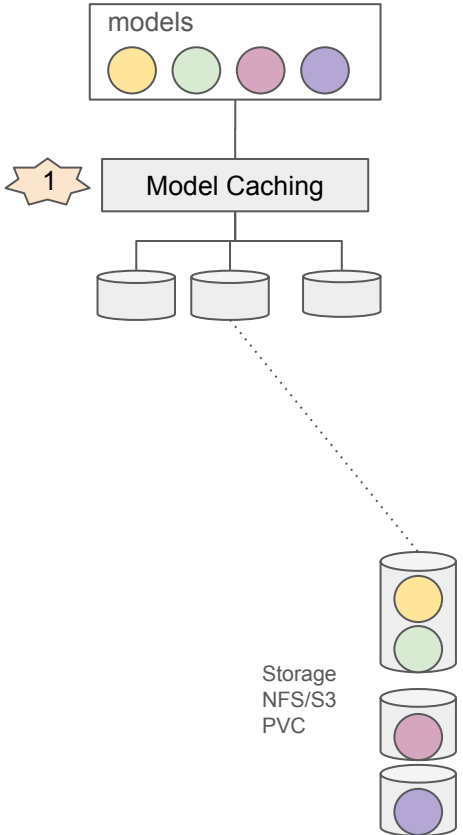

Typical Model Serving Pipeline

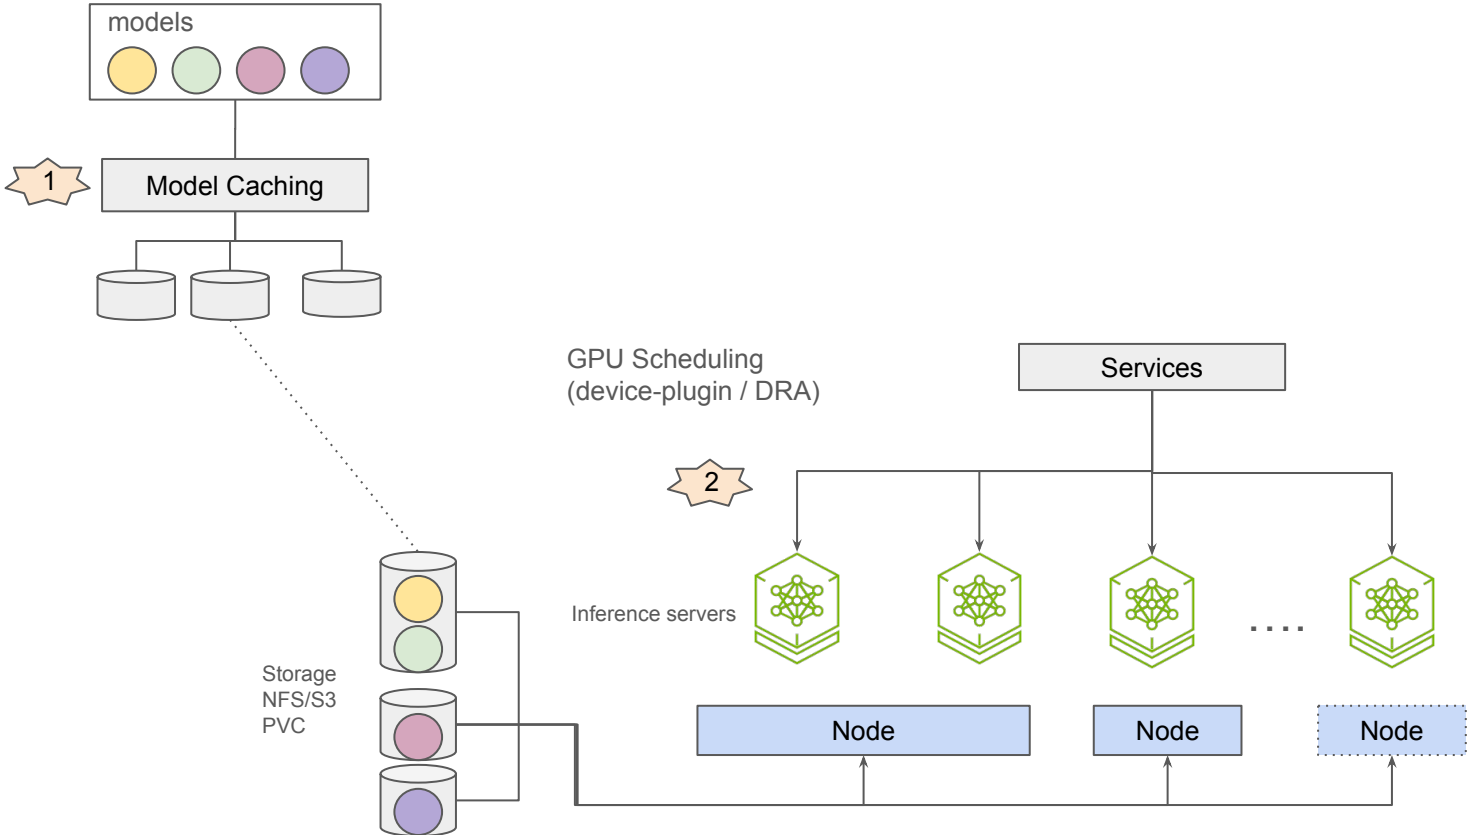

Typical Model Serving Pipeline

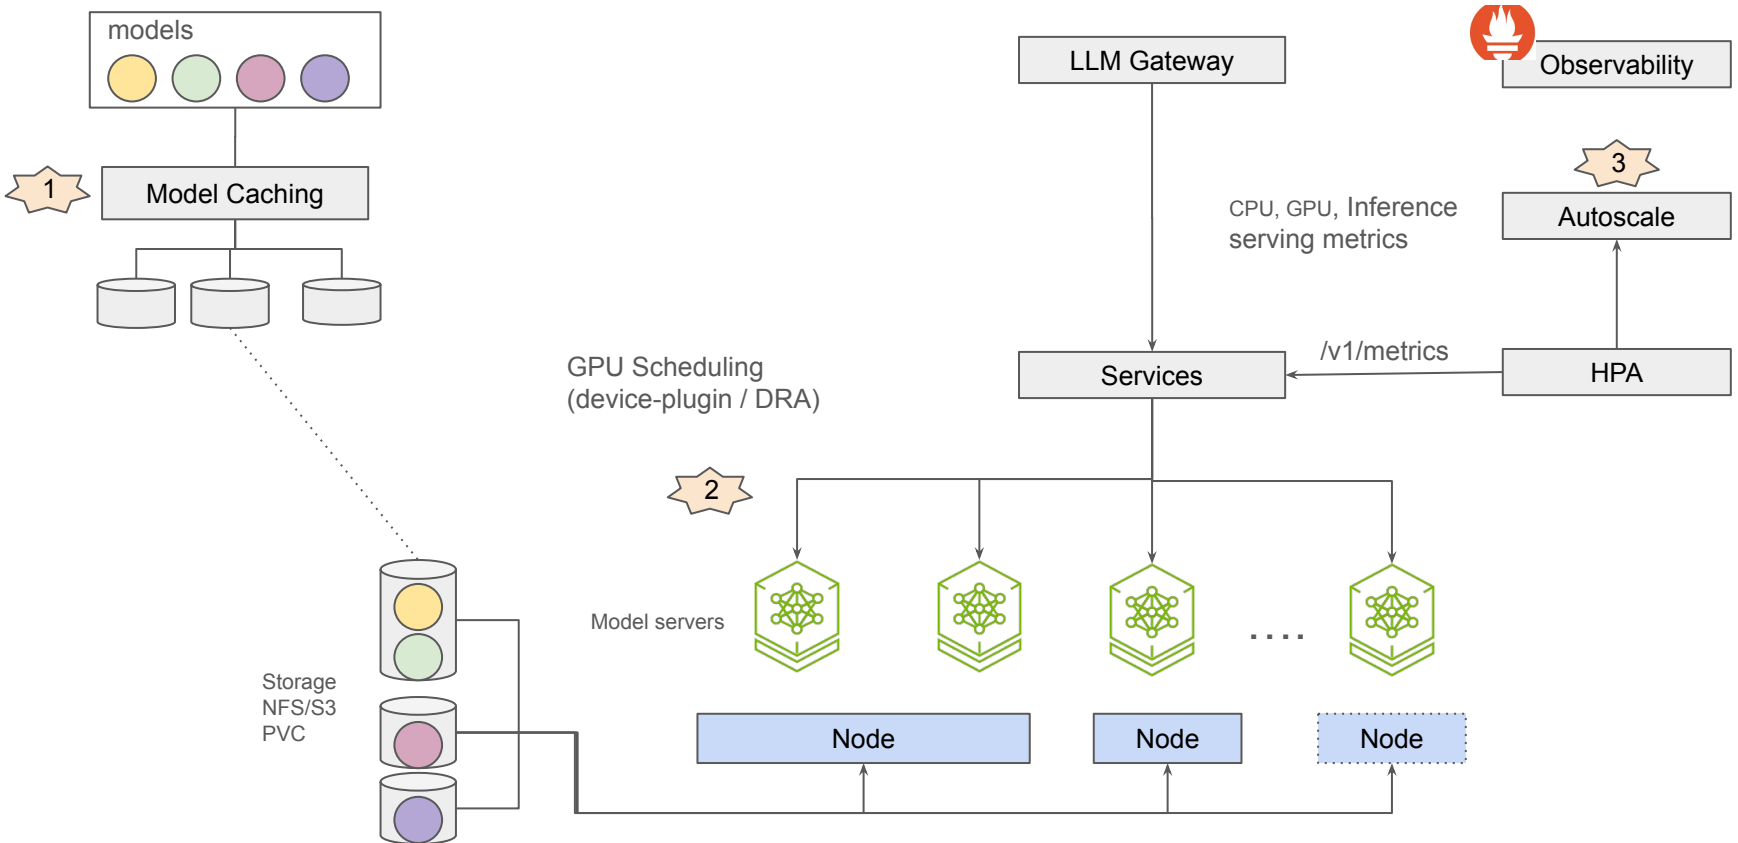

Typical Fine Tuning Pipeline

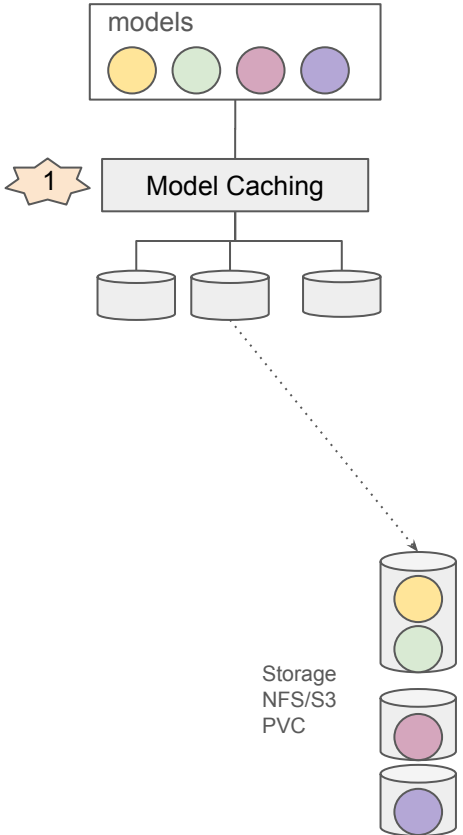

Typical Fine Tuning Pipeline

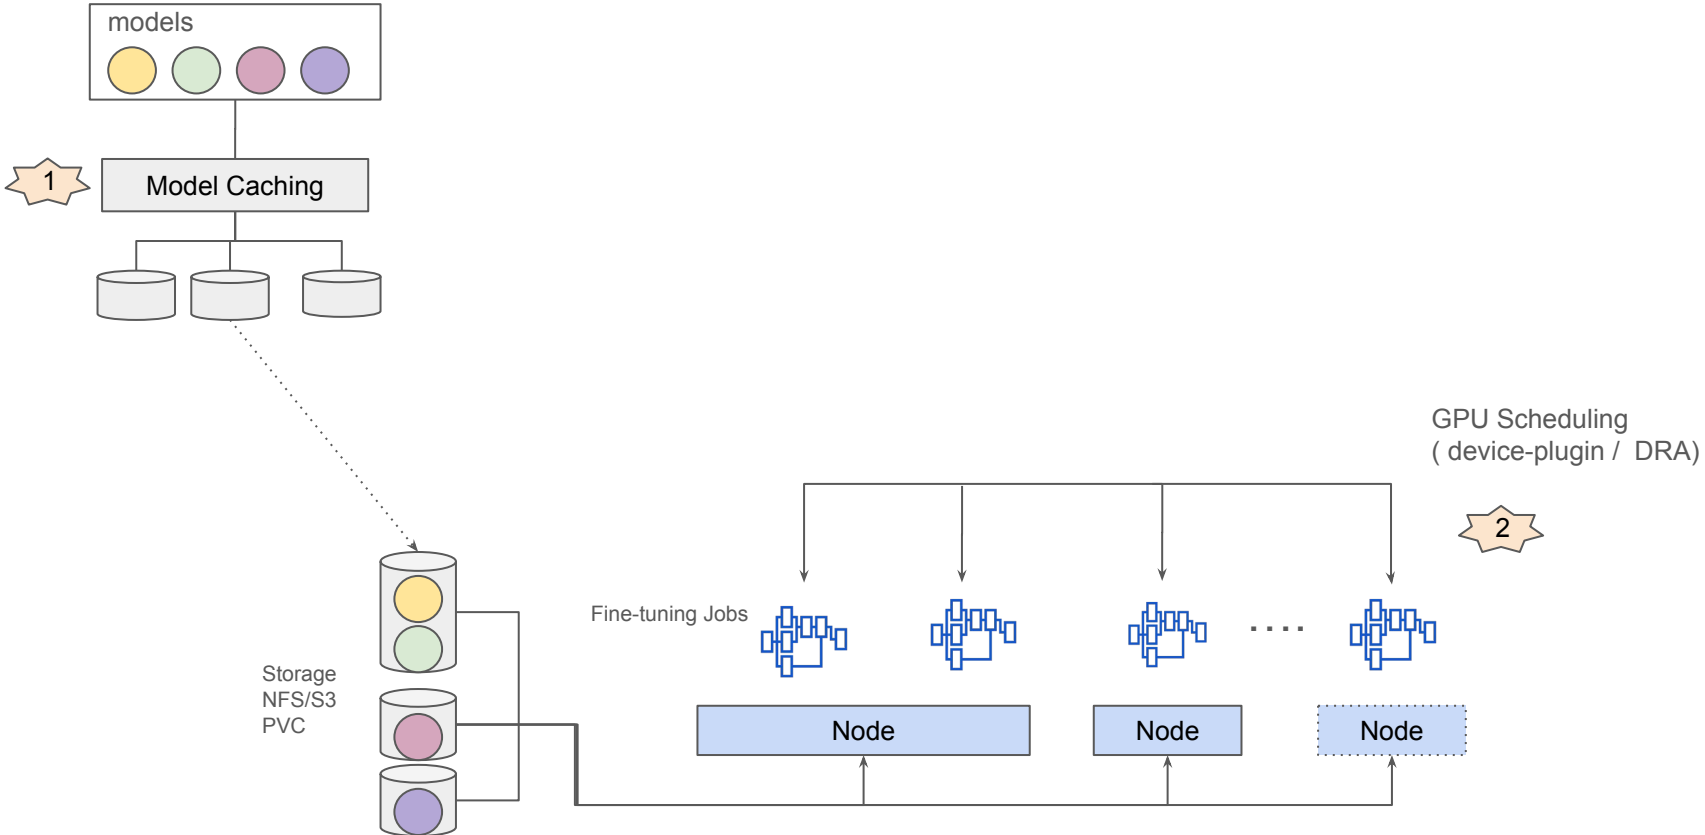

Typical Fine Tuning Pipeline

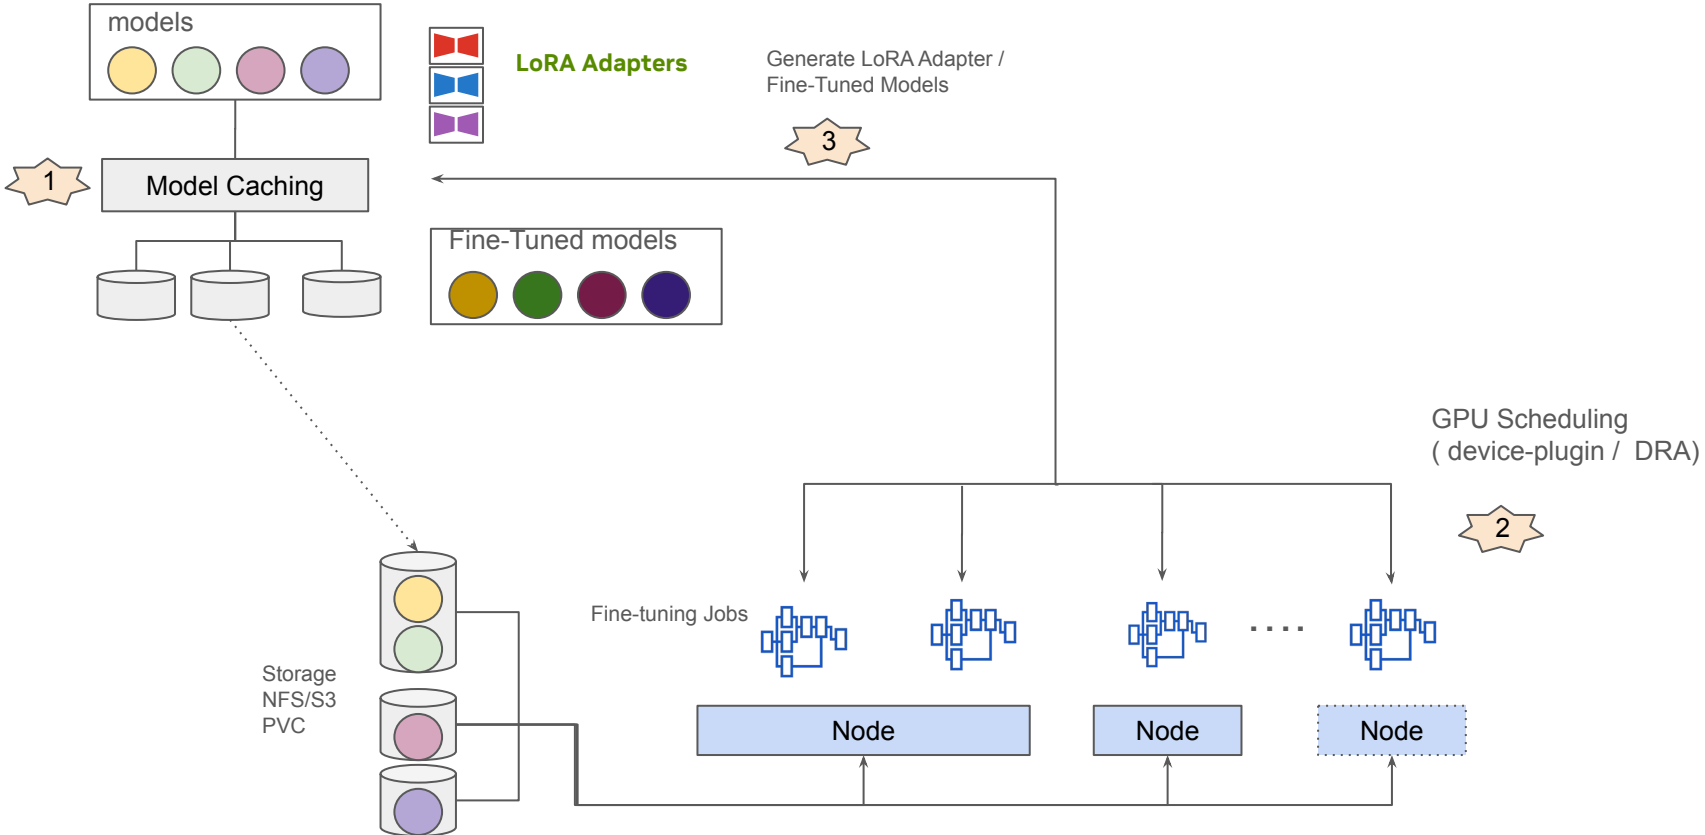

Model Management

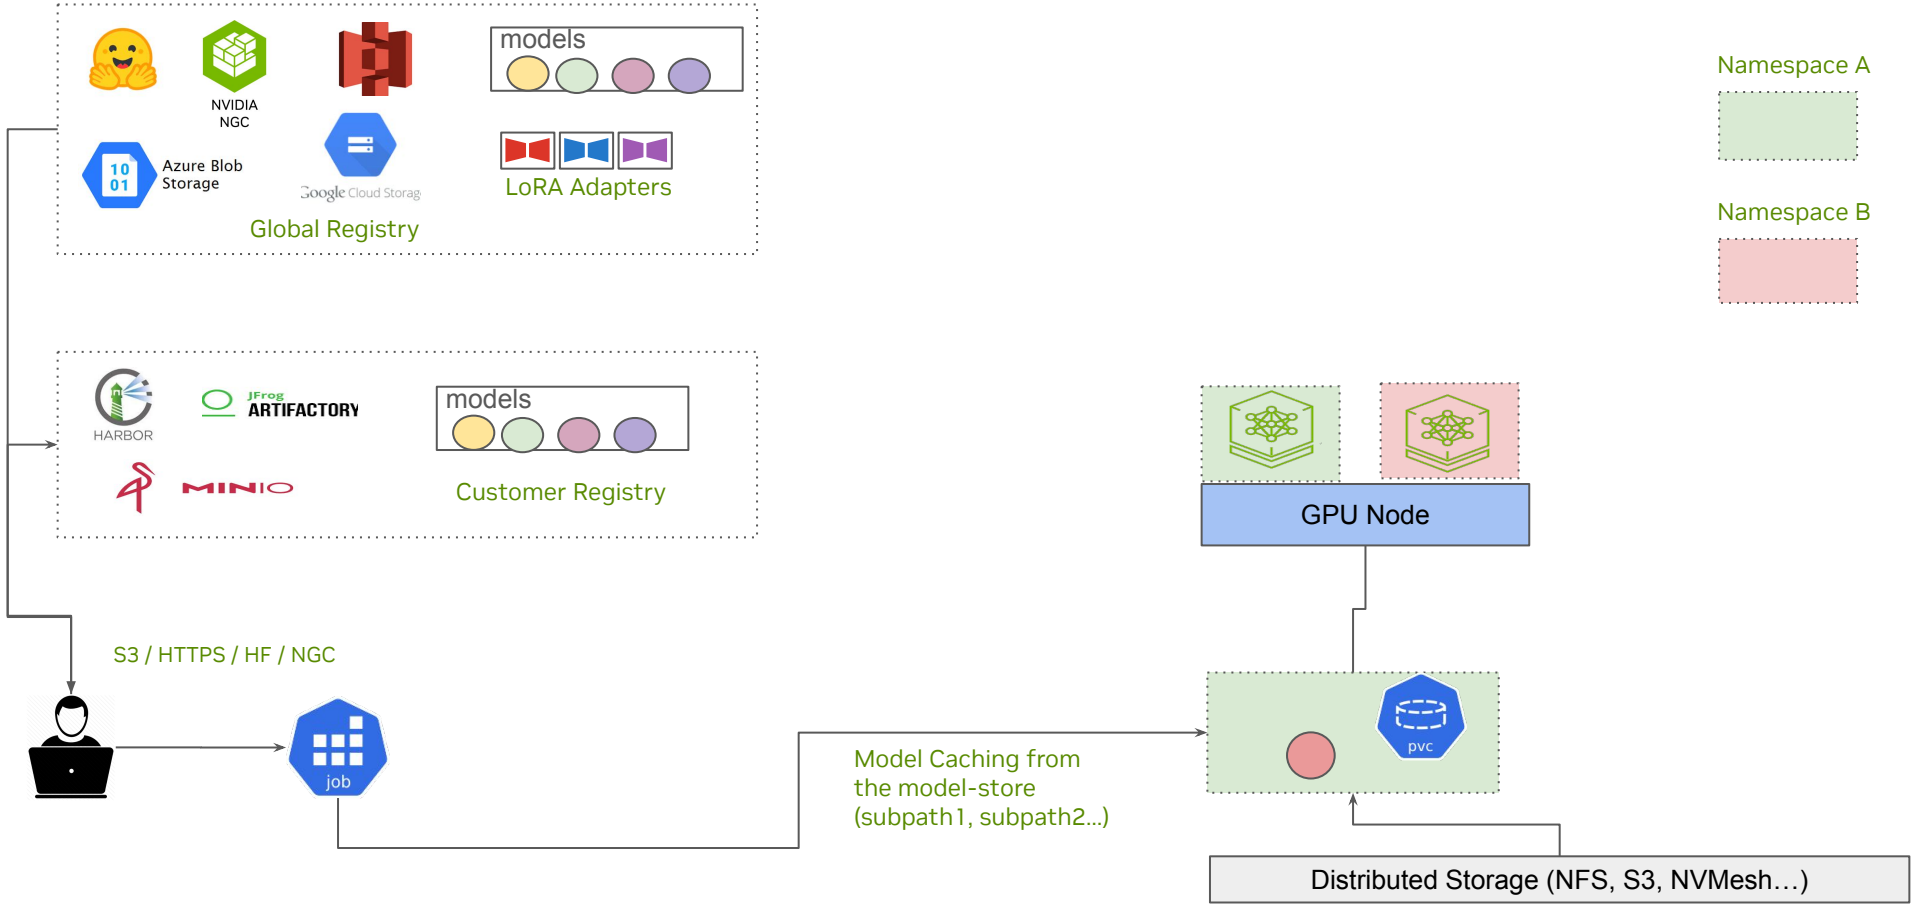

Model Management

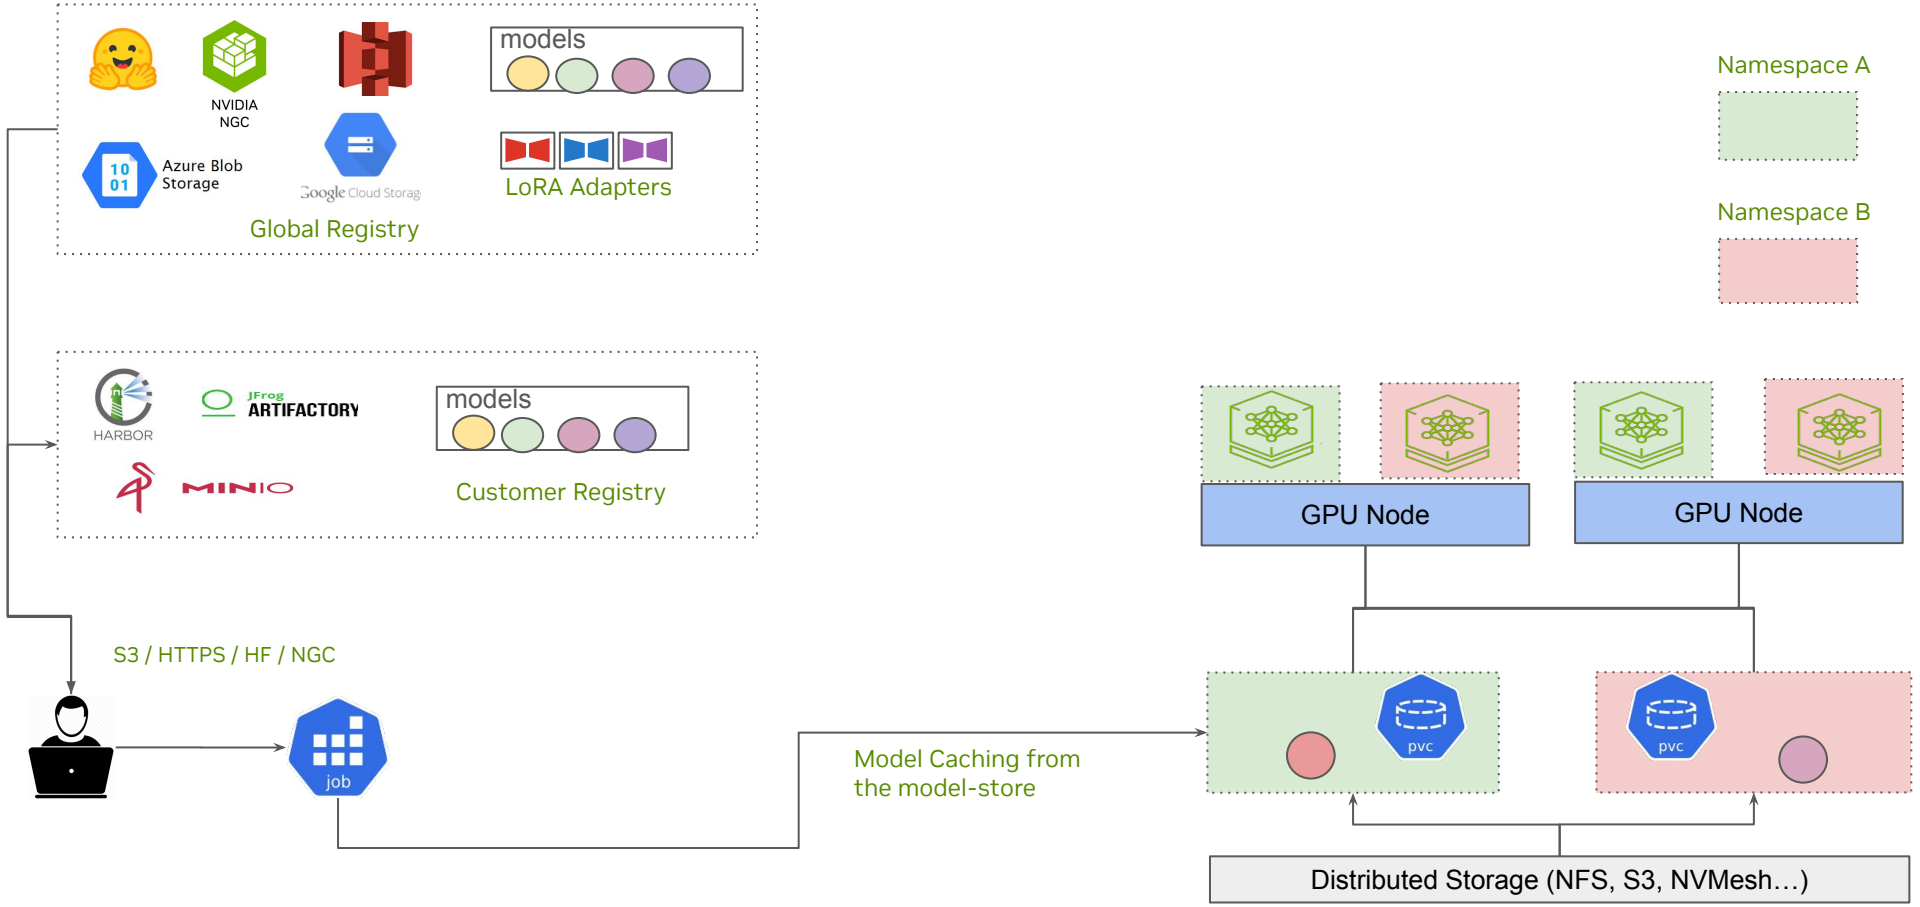

Model Management

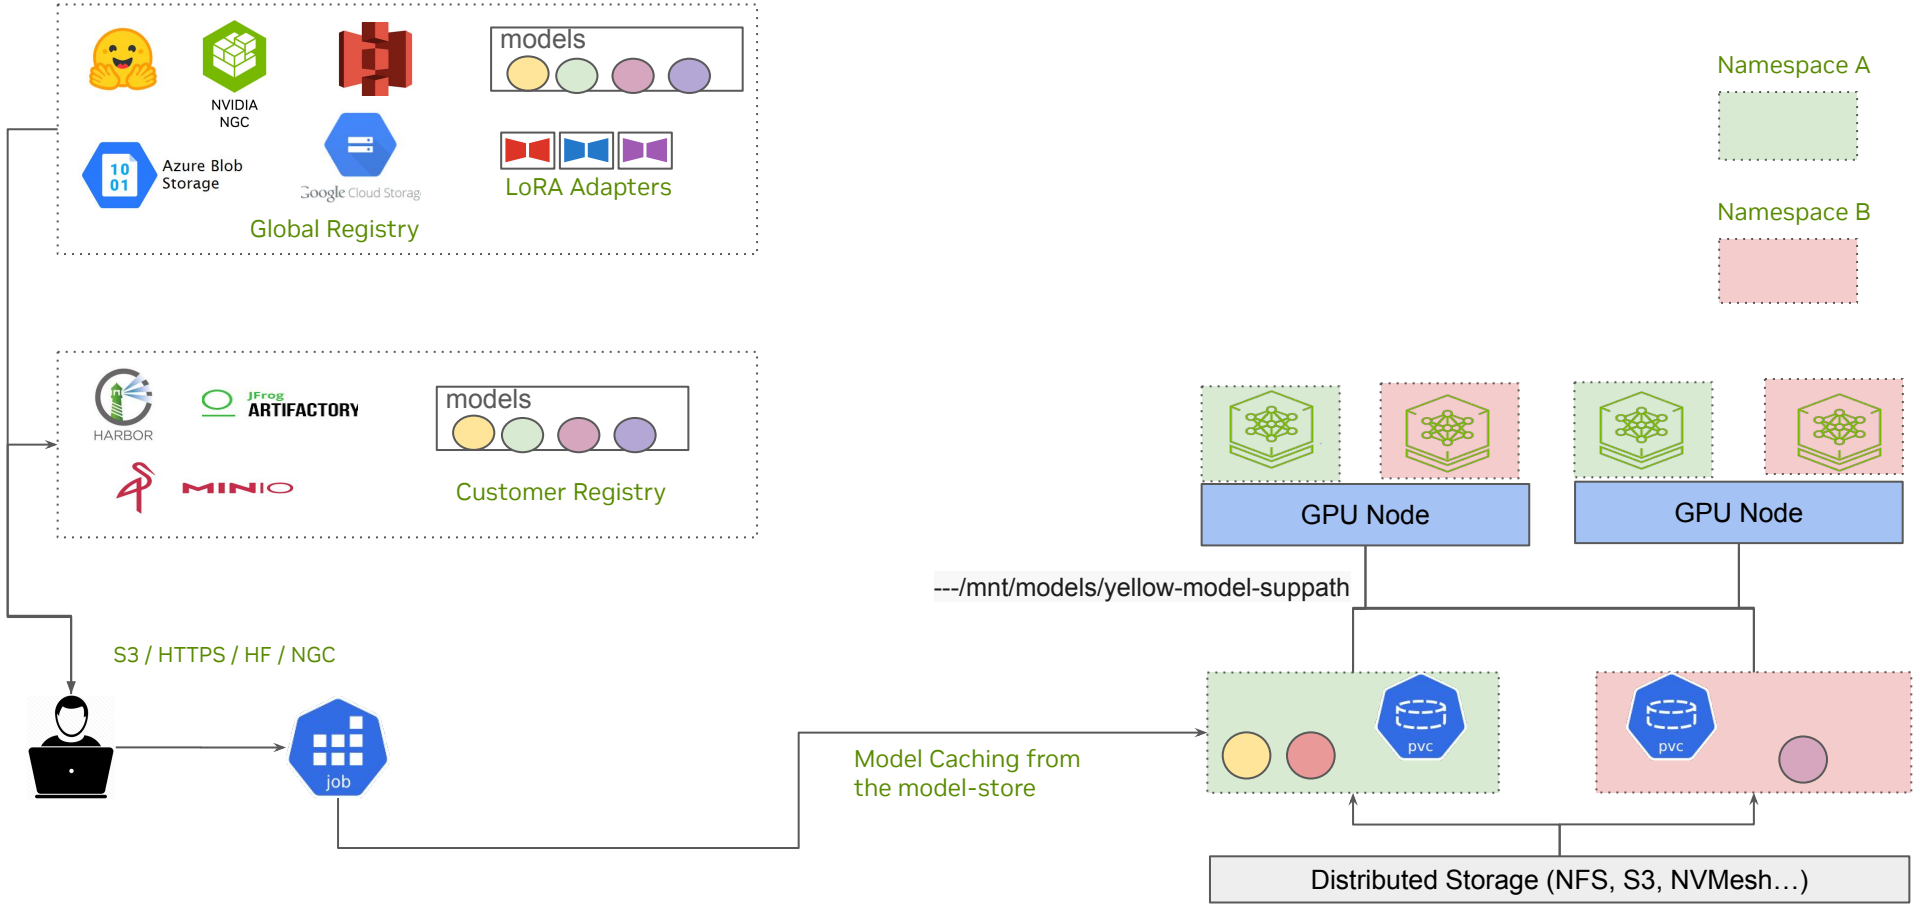

Model Management

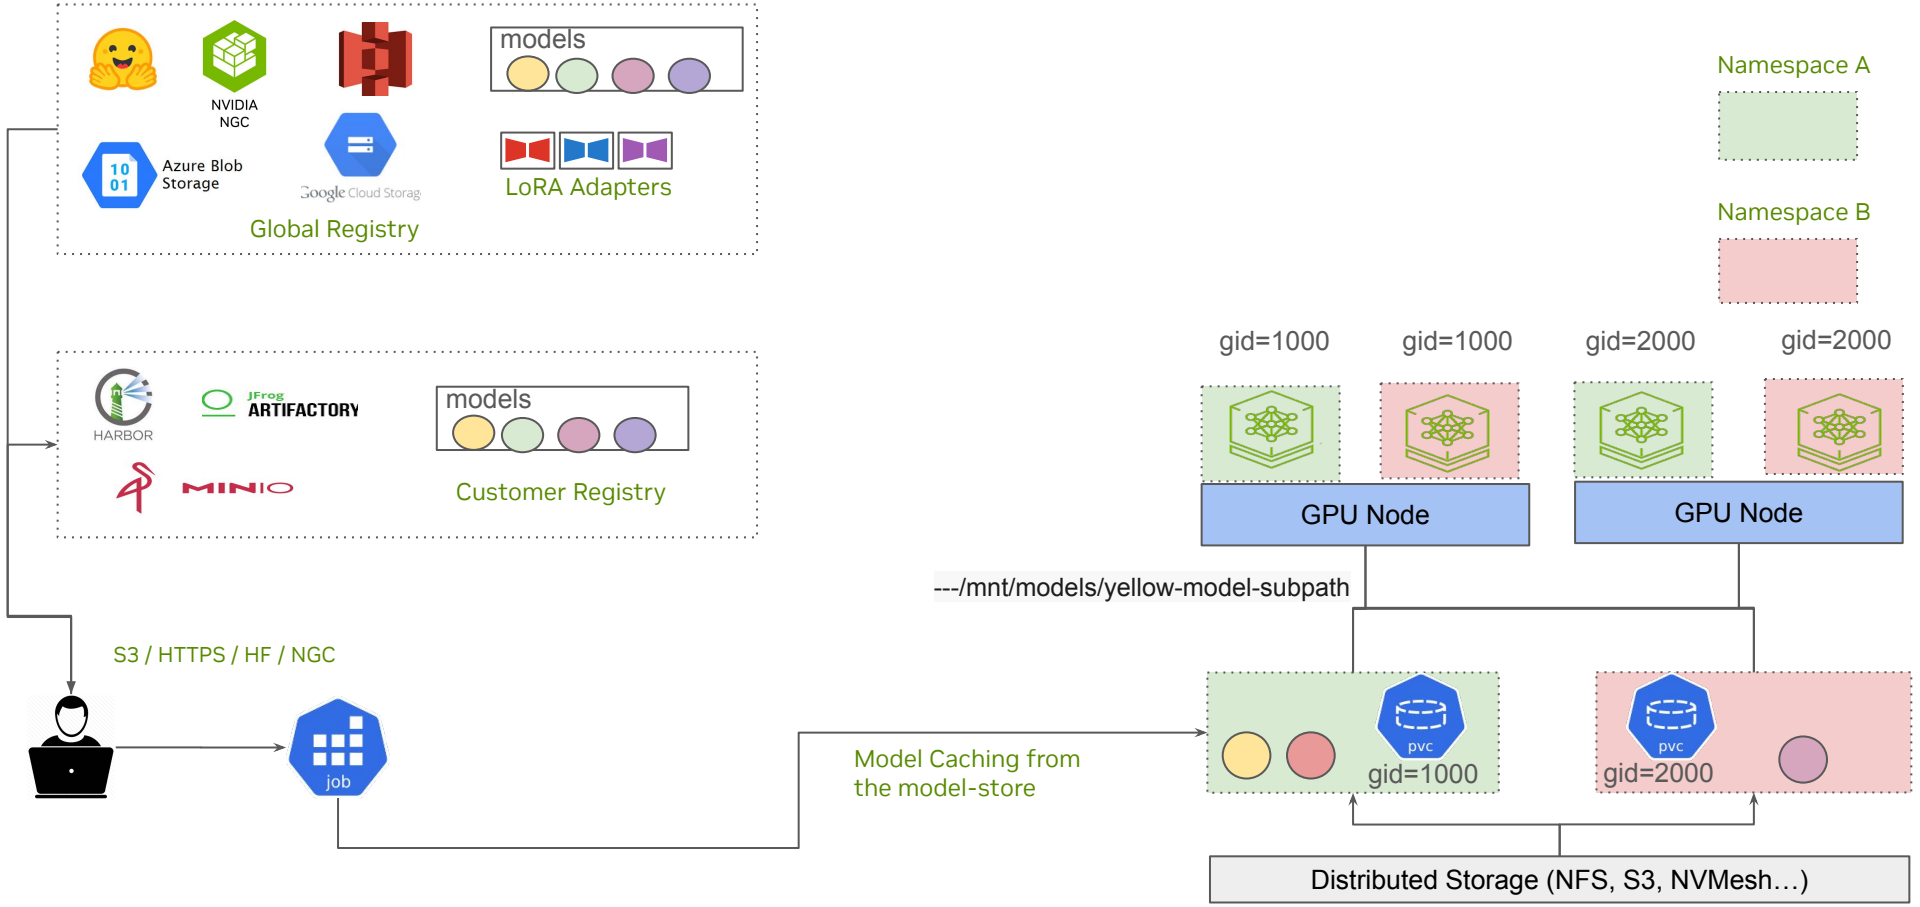

GPU Scheduling

Allocate GPUs to Inference workloads in a Kubernetes Cluster

```
Model Selection
```

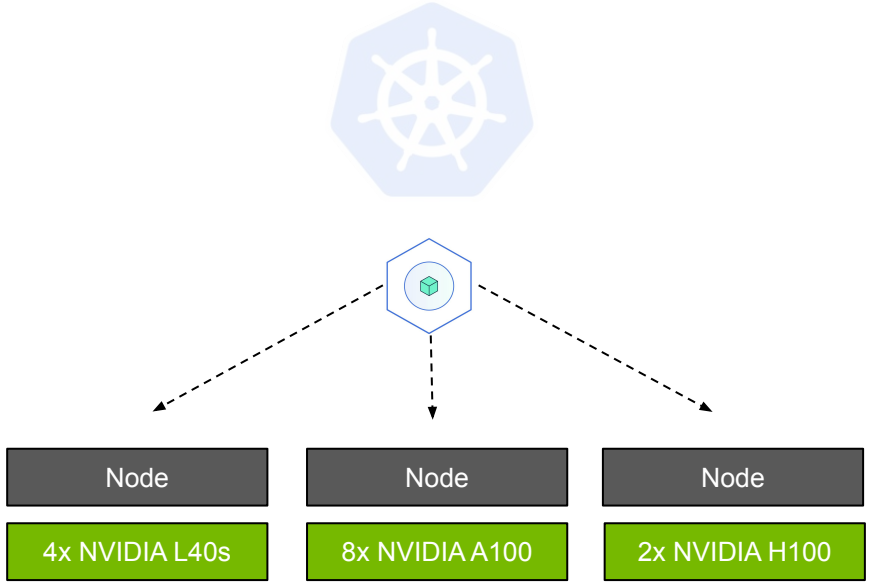

GPU Scheduling

Allocate GPUs to Inference workloads in a Kubernetes Cluster

```
Model Selection -> GPU Scheduling
```

```
apiVersion: v1
kind: Pod
metadata:
  name: model-server
spec:
  containers:
    - name: my-model-server
      image:
nvcr.io/nim/meta/llama3-70b-instruct:1.0.3
      resources:
        limits:
          nvidia.com/gpu: 1
```

- **Device Plugin**

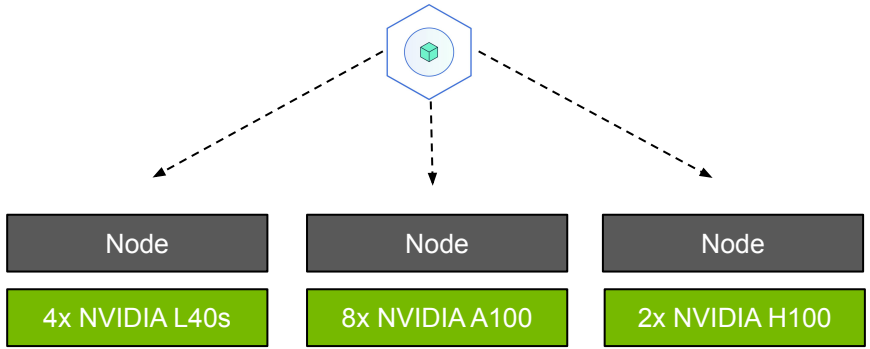

https://github.com/NVIDIA/gpu-operator

GPU Scheduling

Allocate GPUs to Inference workloads in a Kubernetes Cluster

```
Model Selection -> GPU Scheduling -> GPU Allocation
```

```
apiVersion: v1
kind: Pod
metadata:
  name: model-server
spec:
  containers:
    - name: my-model-server
      image:
nvcr.io/nim/meta/llama3-70b-instruct:1.0.3
      resources:
        limits:
          nvidia.com/gpu: 1
  nodeSelector:
    nvidia.com/gpu.product: H100-PCIE-80GB
    nvidia.com/cuda.runtime: 12.7
    nvidia.com/cuda.driver: 565.57.01
```

- Device Plugin
- GPU Feature Discovery

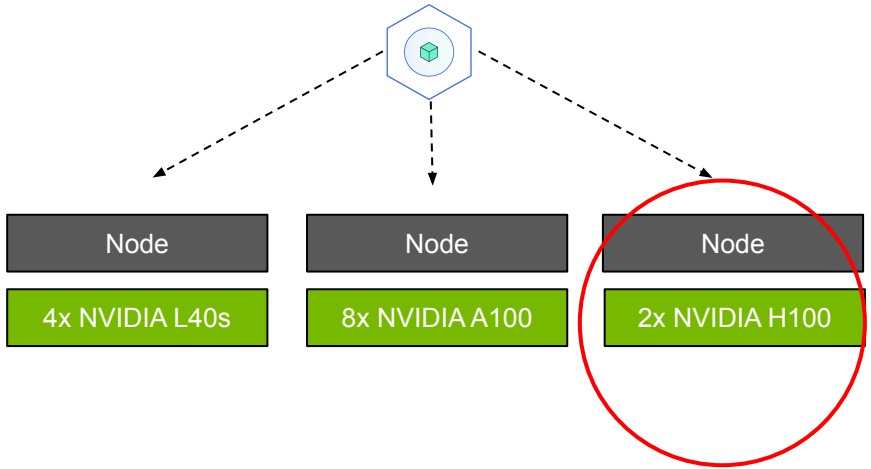

https://github.com/NVIDIA/gpu-operator

GPU Scheduling - With DRA


```yaml
apiVersion: v1
kind: Pod
metadata:
  name: my-model-server
spec:
  containers:
  - name: model-server-ctr
    image:
nvcr.io/nim/meta/llama3-70b-instru
ct:1.0.3
    command: ["nvidia-smi", "-L"]
    resources:
      limits:
        nvidia.com/gpu: 1
```

```yaml
apiVersion: resource.k8s.io/v1alpha2
kind: ResourceClaimTemplate
metadata:
  name: gpu-template
spec:
   devices:
     requests:
       - name: h100
         deviceClassName: gpu.nvidia.com
         selectors:
         - cel:
             expression: |
device.attributes['gpu.nvidia.com'].productName.lowerA
scii().matches('^.*h100.*$')
         count: 1
```

```yaml
---
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
  - name: ctr
    image: nvidia/cuda
    command: ["nvidia-smi" "-L"]
    resources:
      claims:
      - name: gpu
  resourceClaims:
  - name: gpu
    source:
      resourceClaimTemplateName: gpu-template
```

Observability and Autoscaling

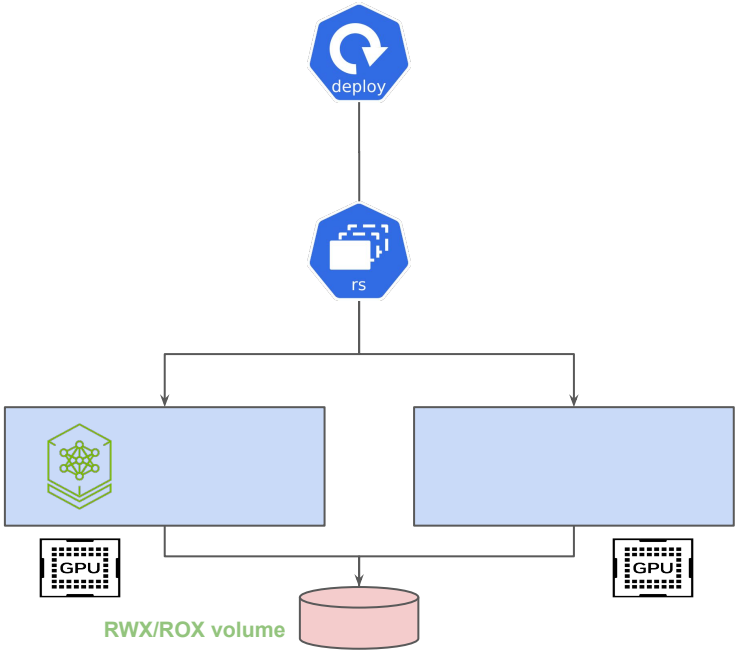

Observability and Autoscaling

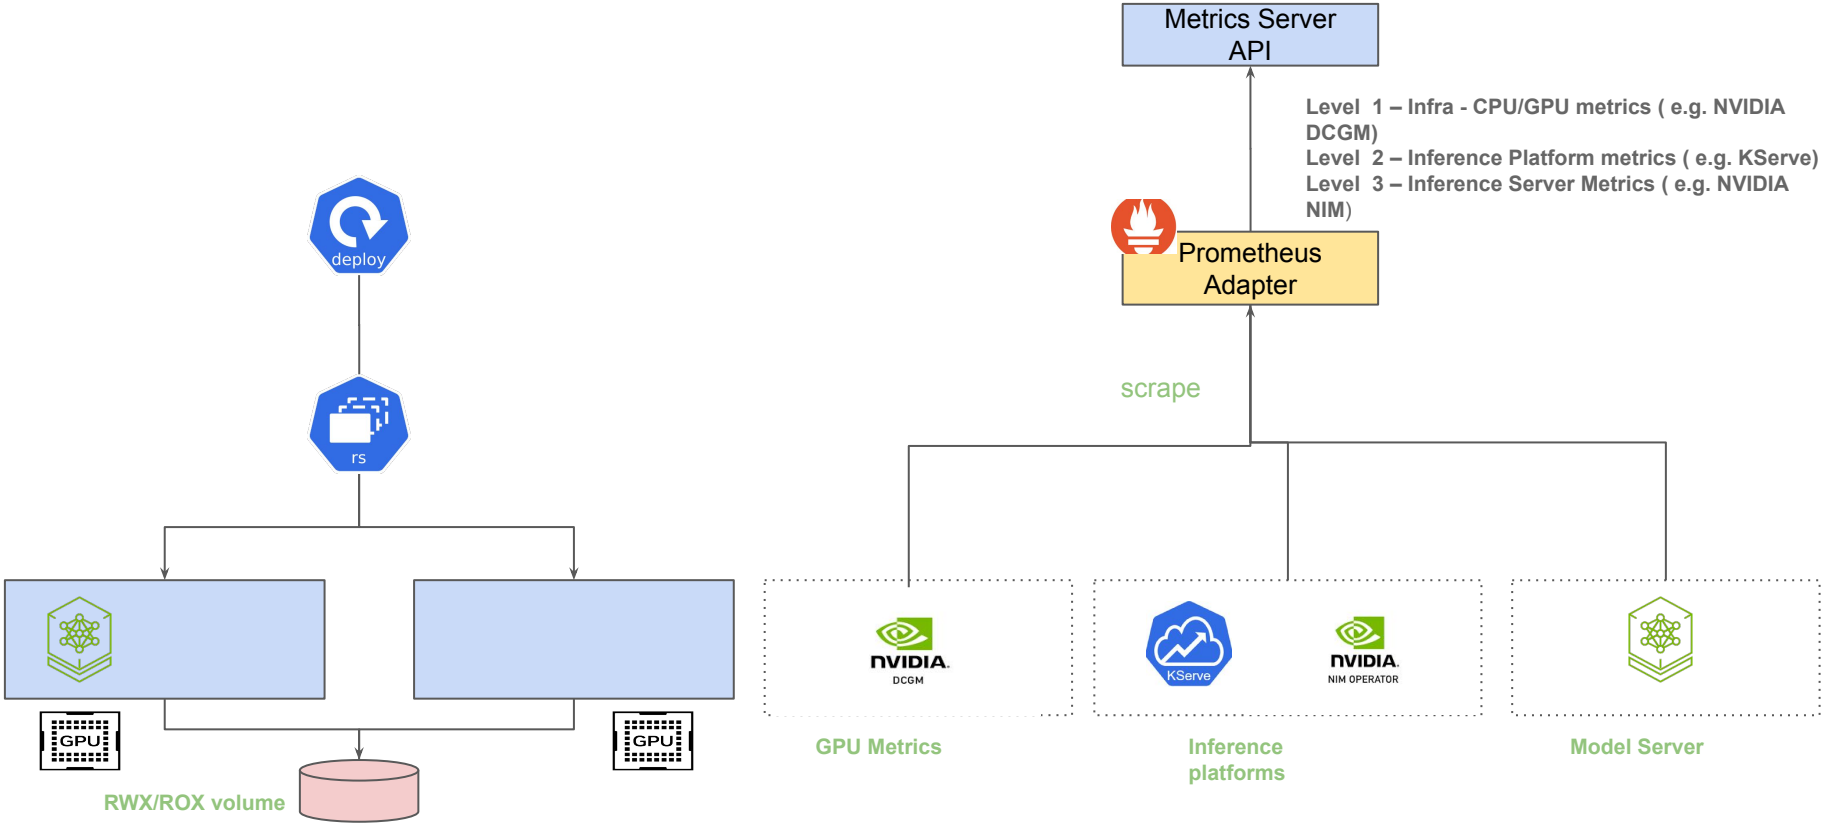

Observability and Autoscaling

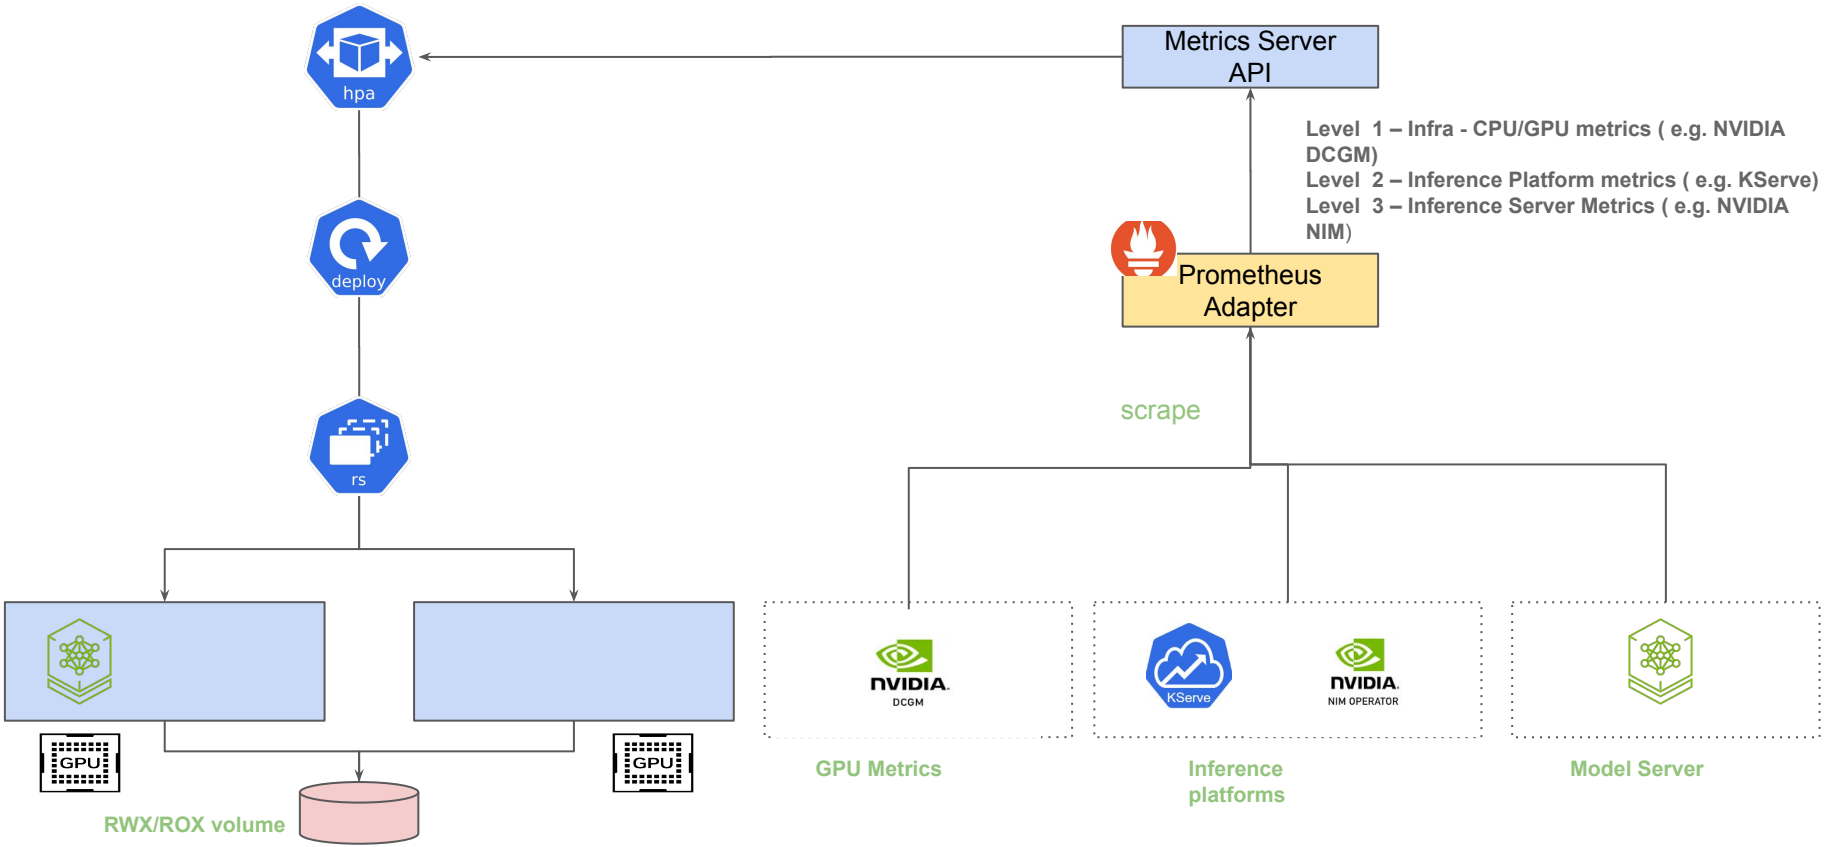

Observability and Autoscaling

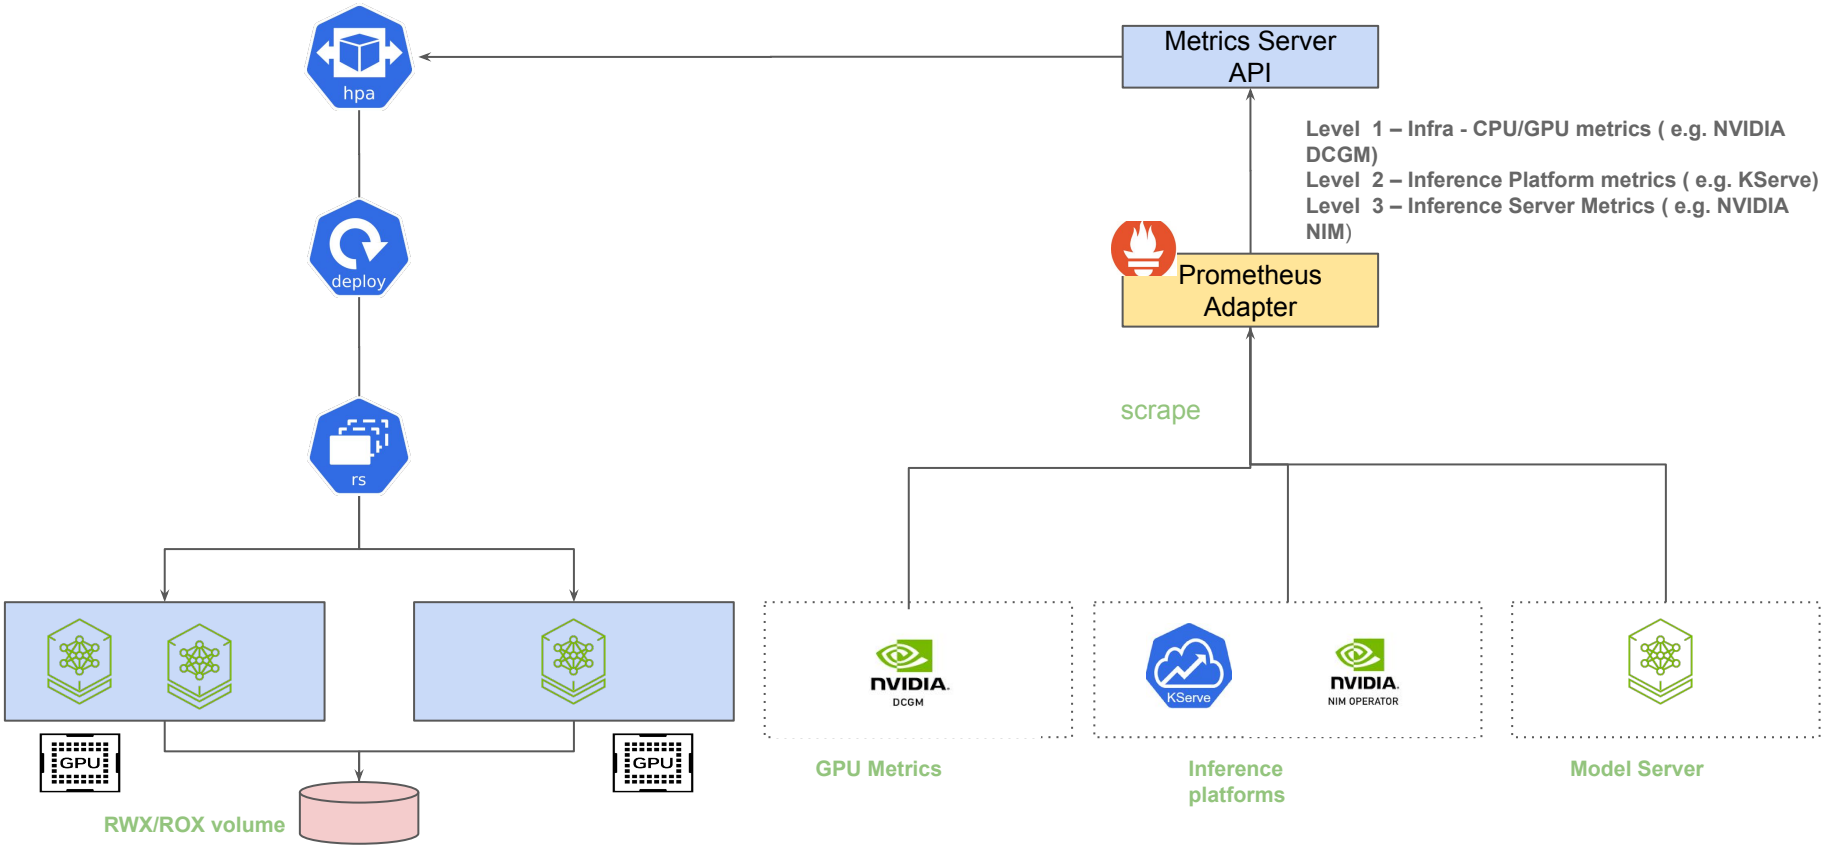

Observability and Autoscaling

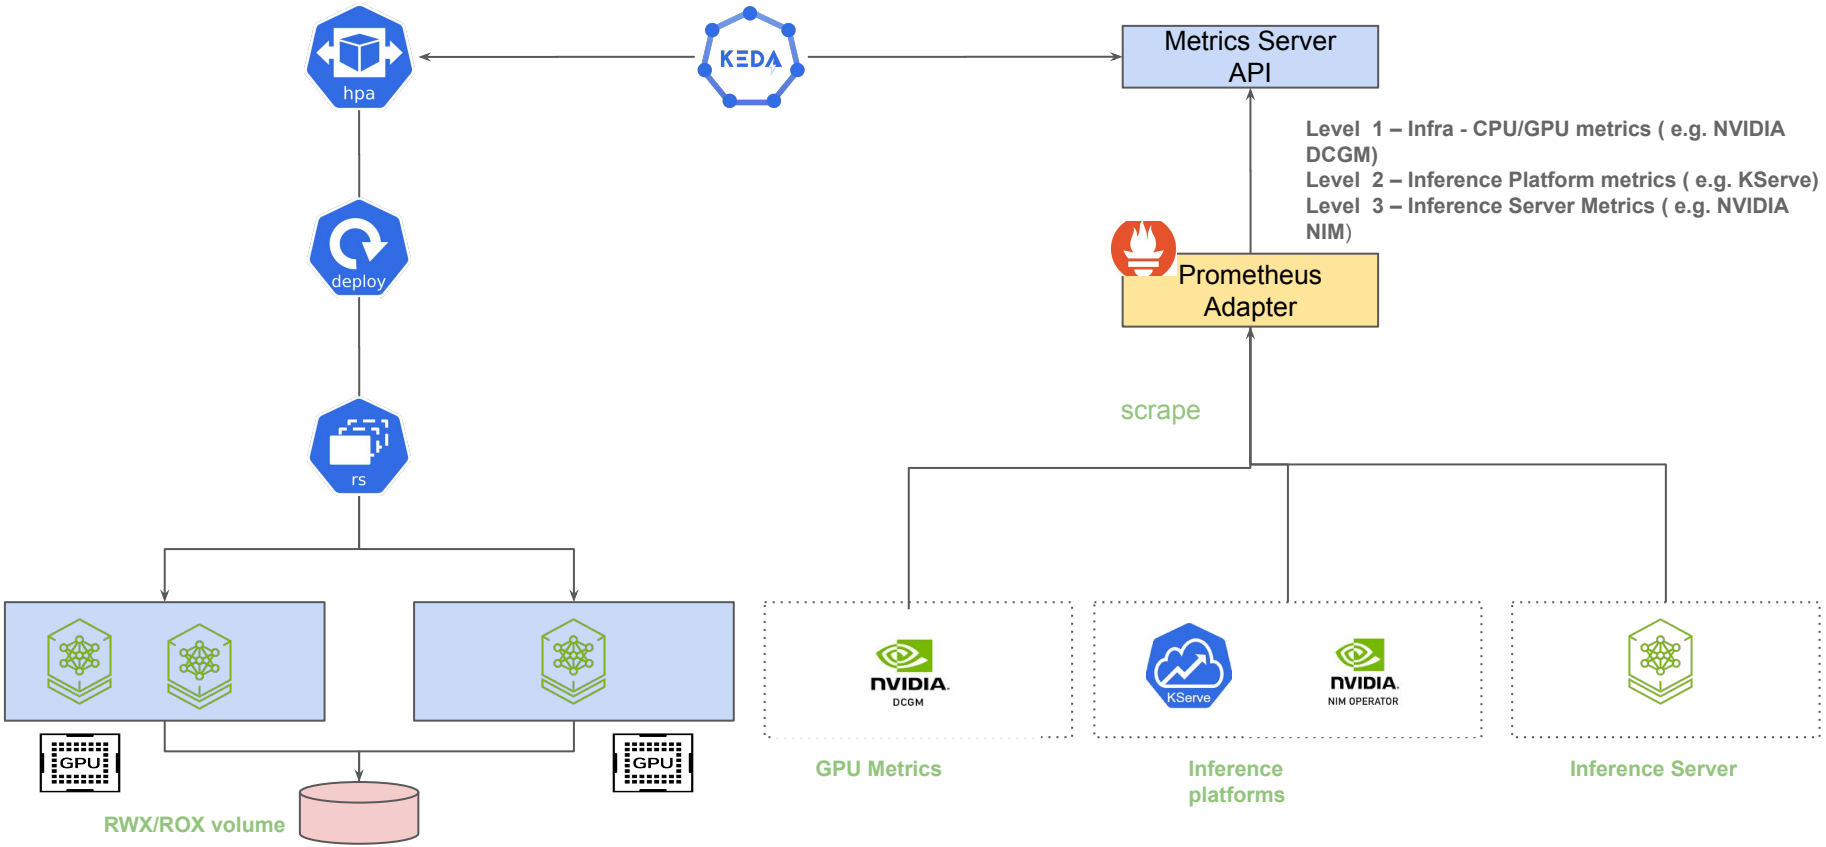

Model and Multi-LoRA Serving

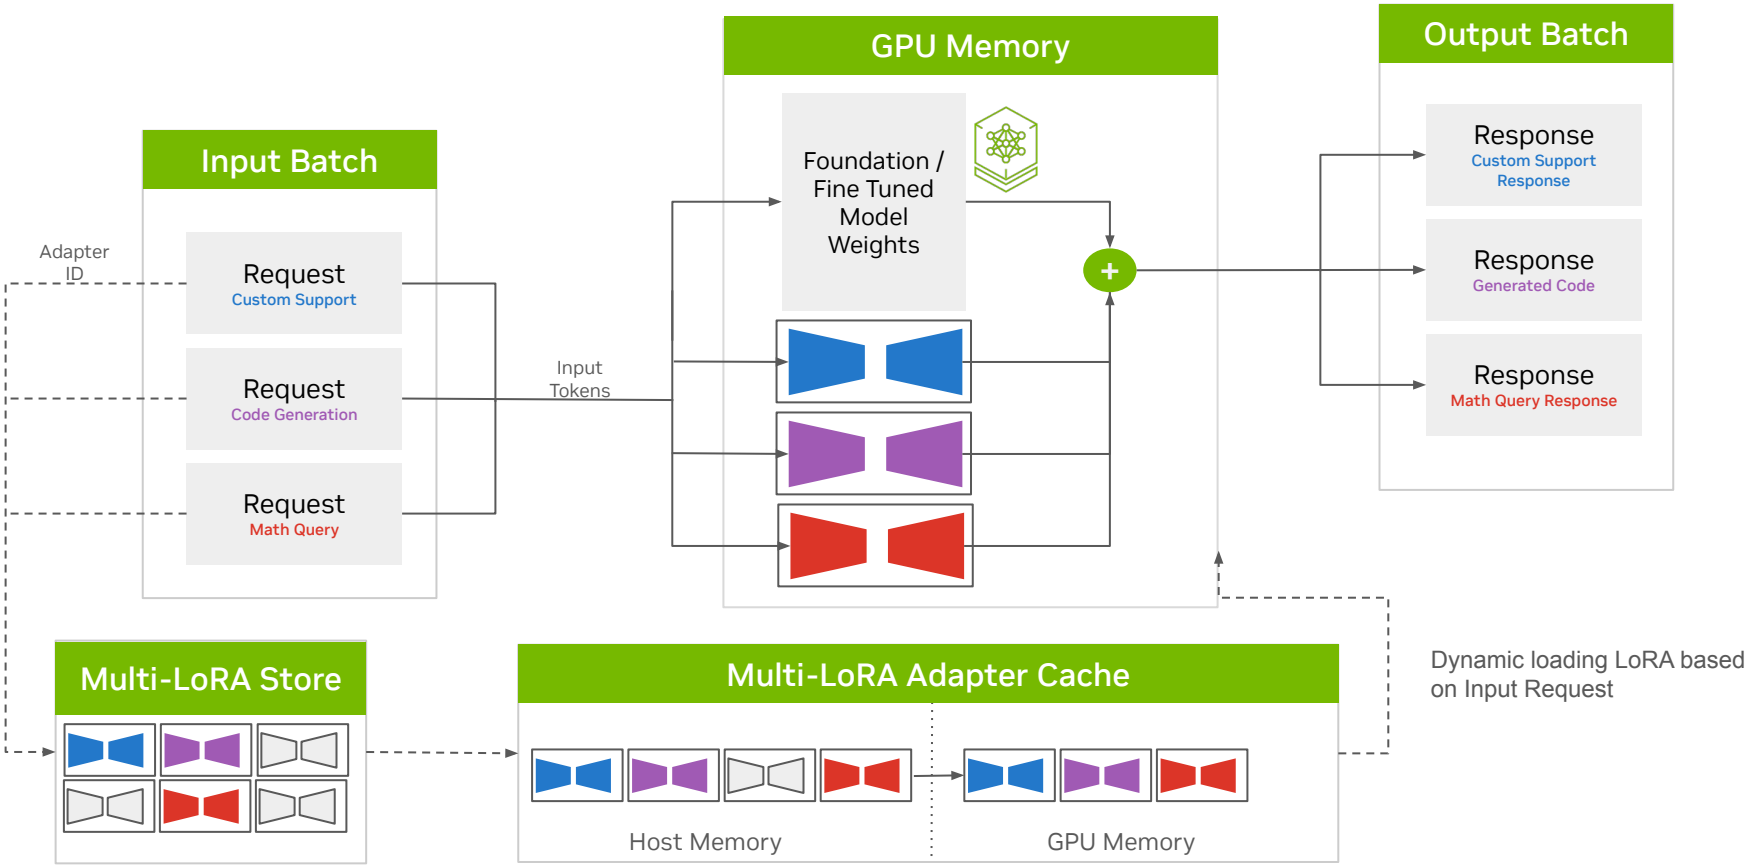

AI Pipelines

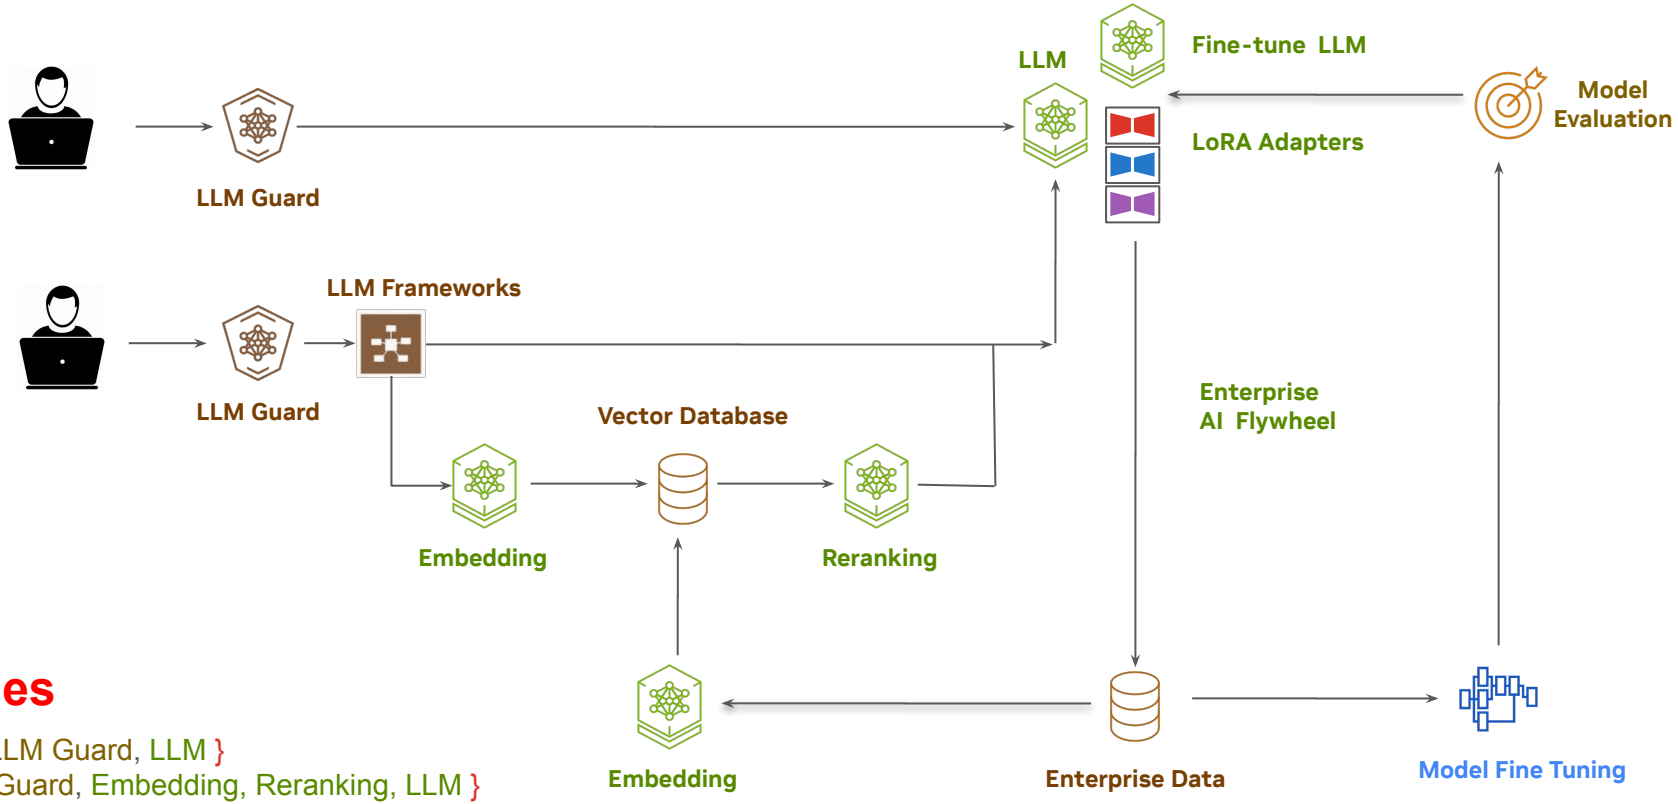

AI Pipelines

- Inference { LLM Guard, LLM }
- RAG { LLM Guard, Embedding, Reranking, LLM }
- Fine Tuning { Model Fine Tuning, Model Evaluation }
- …

Looking Forward…

Multi-node Inference

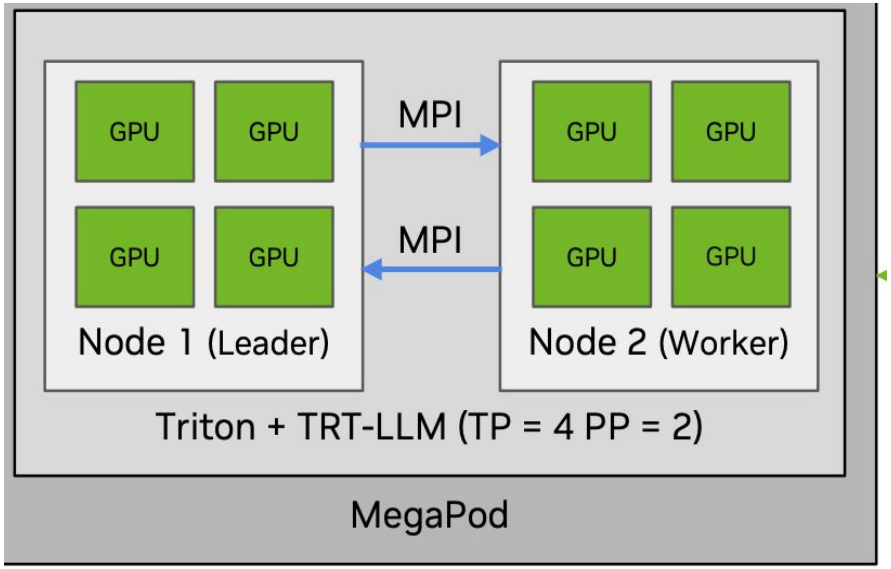

Use Case

- Deploy Massive LLMs
- Automatically Scale and Load Balance

https://github.com/kserve/kserve/pull/3871

Multi-node Inference: Model Sharding

Pipeline Parallelism 2 (PP 2)

Split Into Sets of Layers

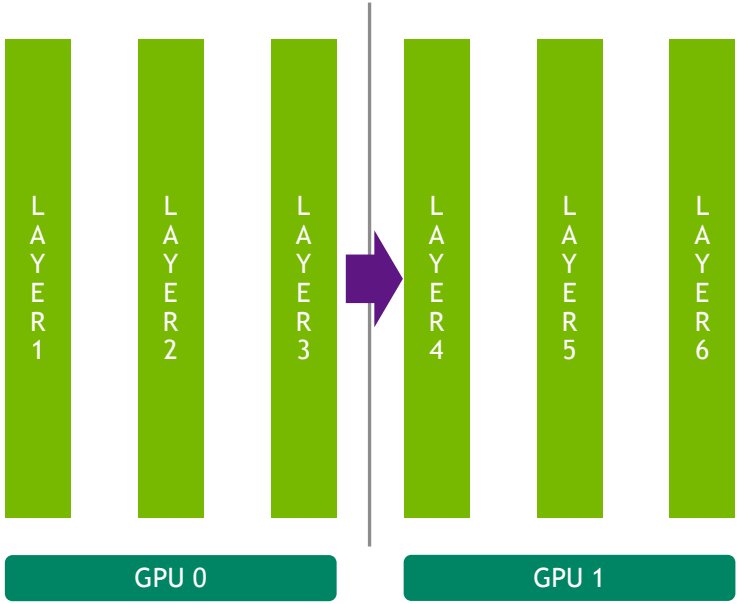

NCCL Send

Tensor Parallelism 2 (TP 2)

Split individual layers across GPUs

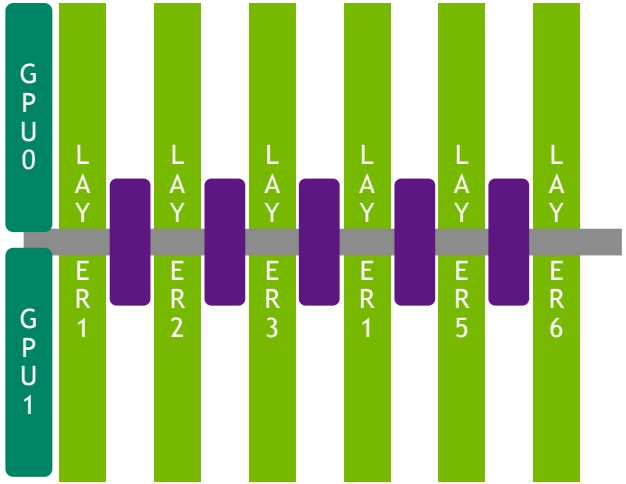

NCCL All Reduce

Multi-node Inference: Key Challenges

- Schedule group of nodes
  - Gang scheduling, Binpacking
- Operations
  - Deploy/Scale group of Nodes (LWS, ..)
  - Multi-node communication (MPI, …)
  - Leader Aware Load Balancing

- Optimizations
  - Accelerate initial loading (caching on shared storage)
  - Accelerate cross-node communication (RoCE / RDMA)

AI Agents

Enable advanced problem solving and automation for improved user experience

AI agents

- **Reasoning and Planning**: Decomposing complex tasks into manageable subgoals through reasoning
- **Memory**
  - Short-term memory in an LLM-powered agent acts as a record of actions and thoughts during a single query
  - Long-term memory logs interactions between the user and agent over extended periods
- **Tools:** Defined executable workflows that agents use to perform tasks

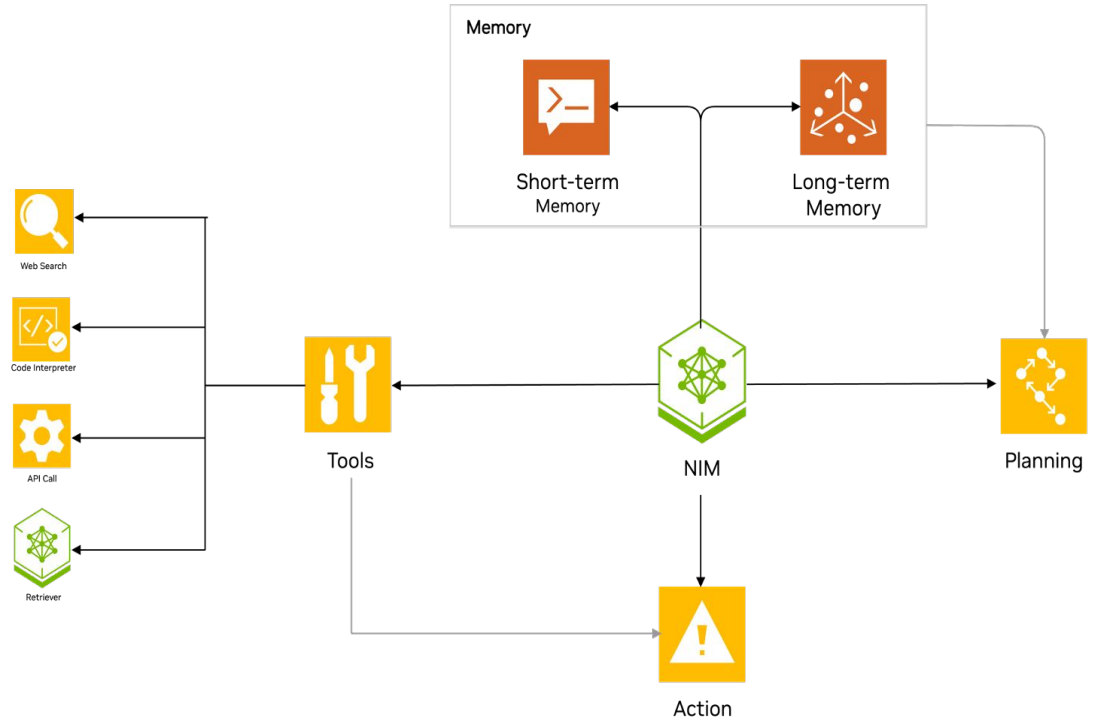

Conclusion

- K8s is a great platform for AI Pipelines
  - Strong GPU and Storage Integrations
  - Advanced Inference and Fine Tuning Platforms
  - Ease of Management and Monitoring

- NVIDIA NIM Operator leverages all these to simplify deployment of AI pipeilnes in Kubernetes

- Community is working on addressing gaps
  - Auto Scaling
  - Model Cache Management
  - LLM Gateway

Thank you and Feedback

KubeCon CloudNativeCon
North America 2024