



**KubeCon**



**CloudNativeCon**

**North America 2024**





KubeCon



CloudNativeCon

North America 2024


# Better Together! GPU, TPU and NIC Topological Alignment with DRA

November, 2024

Patrick Ohly (Intel)

John Belamaric (Google)

# Who are we?



**Patrick Ohly**  
*Principal Engineer*  
*Intel*



**John Belamaric**  
*Sr Staff Software Engineer*  
*Google*

# Dynamic Resource Allocation (DRA) in Kubernetes

- New way of requesting resources
  - Alpha since Kubernetes 1.26
  - Major API changes in 1.31
  - Beta in 1.32
  - GA in ??
- An *alternative* to Device Plugin “count-based” interface
  - `nvidia.com/gpu: 2`
- Provides a much richer API for requesting / configuring resources
- Inspired by the persistent volume API

# Dynamic Resource Allocation (DRA) in Four Parts

Part 1: New Kubernetes API to **describe** devices (ResourceSlice):

*This device is an nvidia.com/gpu, its product ID is A100-SXM4-40GB, it has 40Gi of memory, and 3456 FP64 cores.*

Part 2: New Kubernetes API to **request** devices (ResourceClaim):

*I need an nvidia.com/gpu with at least 30Gi of memory and at least 3000 FP64 cores.*

Part 3: Updated scheduler to **match** requests to devices.

Part 4: New Kubelet API to **actuate** the scheduler's decisions.

# DRA overcomes the limitations of device plugins

Can **subdivide** large devices

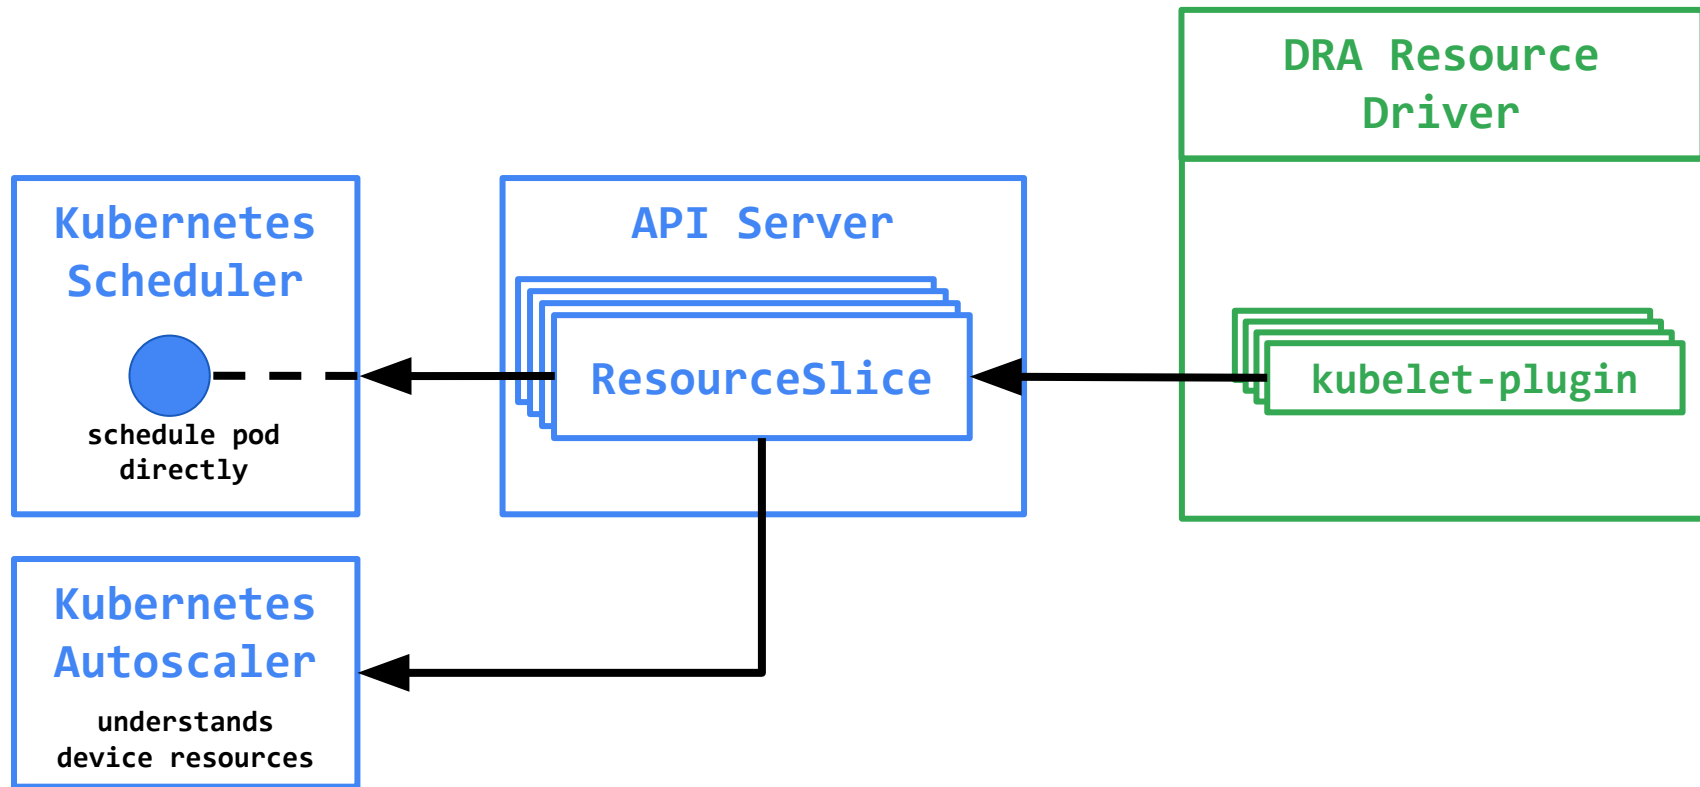
Can **configure** devices individually

Can **share** GPUs in the same node for diverse workloads

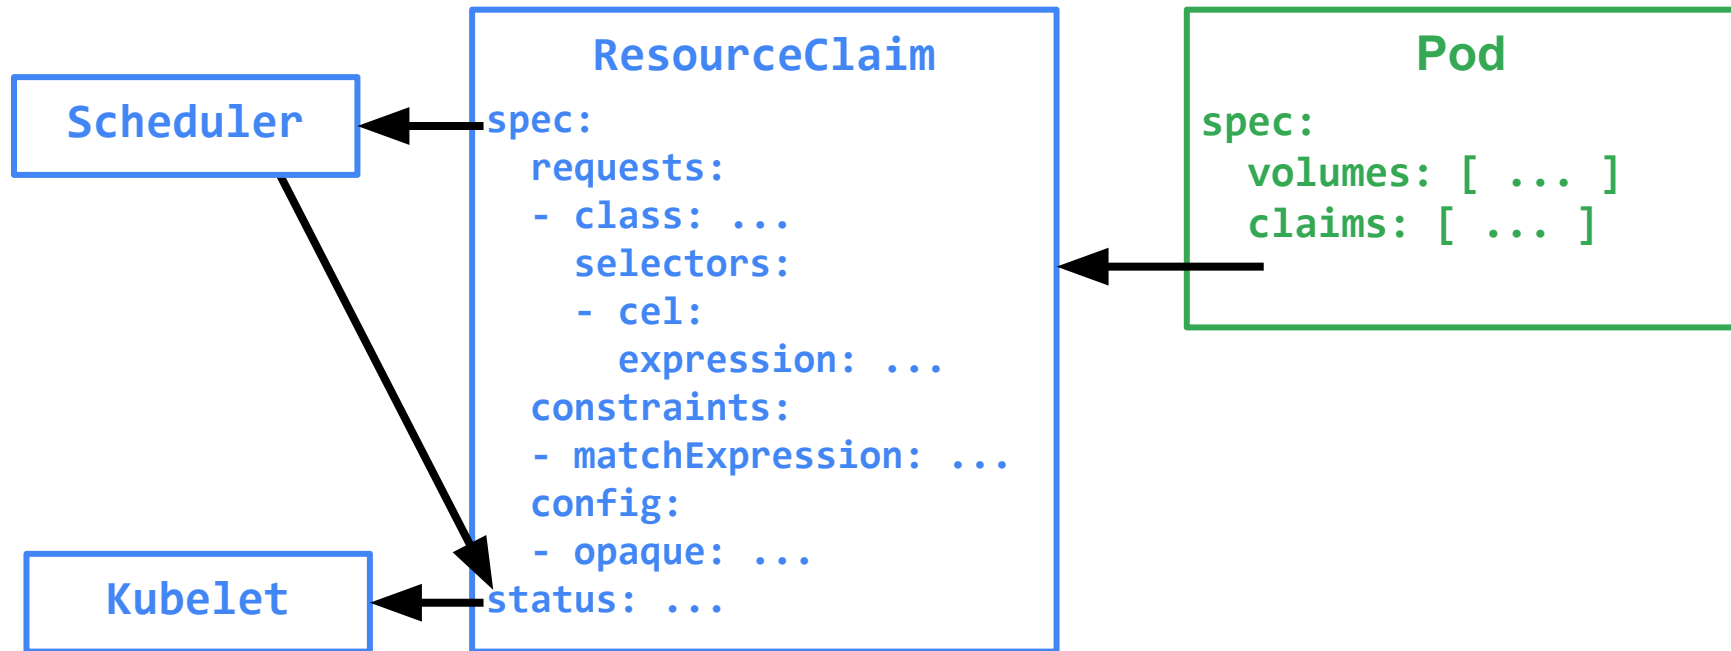
Foundational for new functionality:

- ***Alignment of multiple, independent devices (GPU and NIC alignment on PCIe)***
- Workload-specific accelerator sharing configuration
- Dynamic MIG and TPU
- Consumption of multiple associated devices as a unit

# DRA: Advertising resources



# DRA: Requesting resources







KubeCon

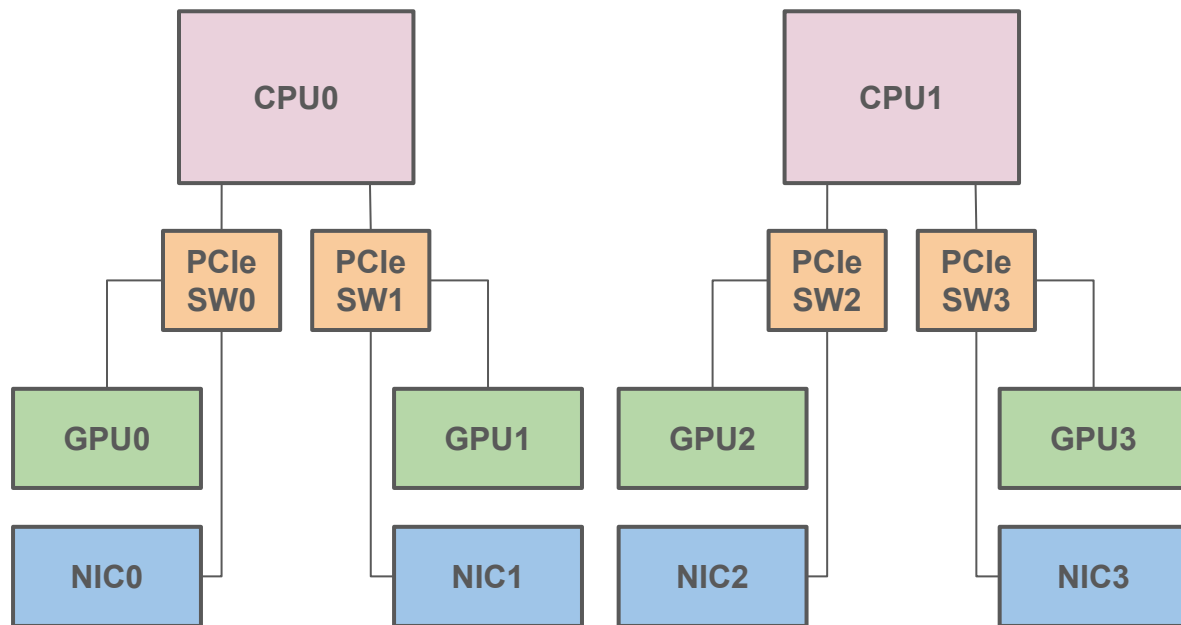


CloudNativeCon

North America 2024

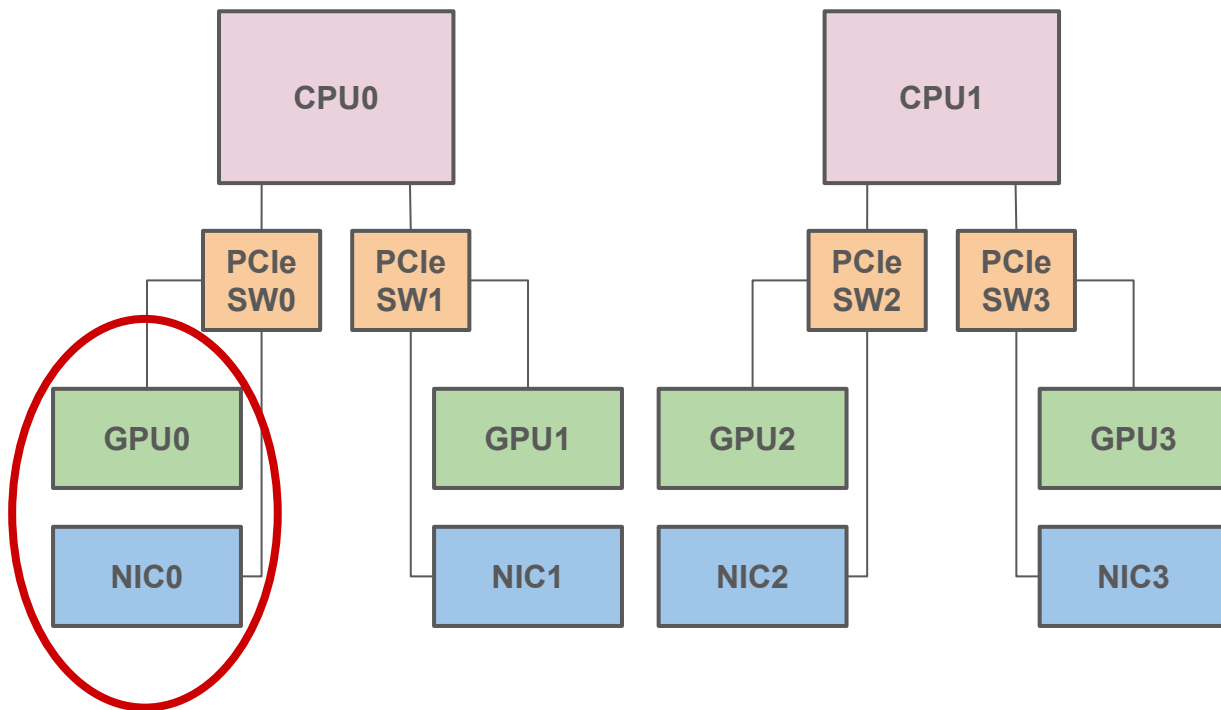
# Device Plugin vs DRA for GPU / NIC Topology

# Simplified Node Topology - GPU / NIC



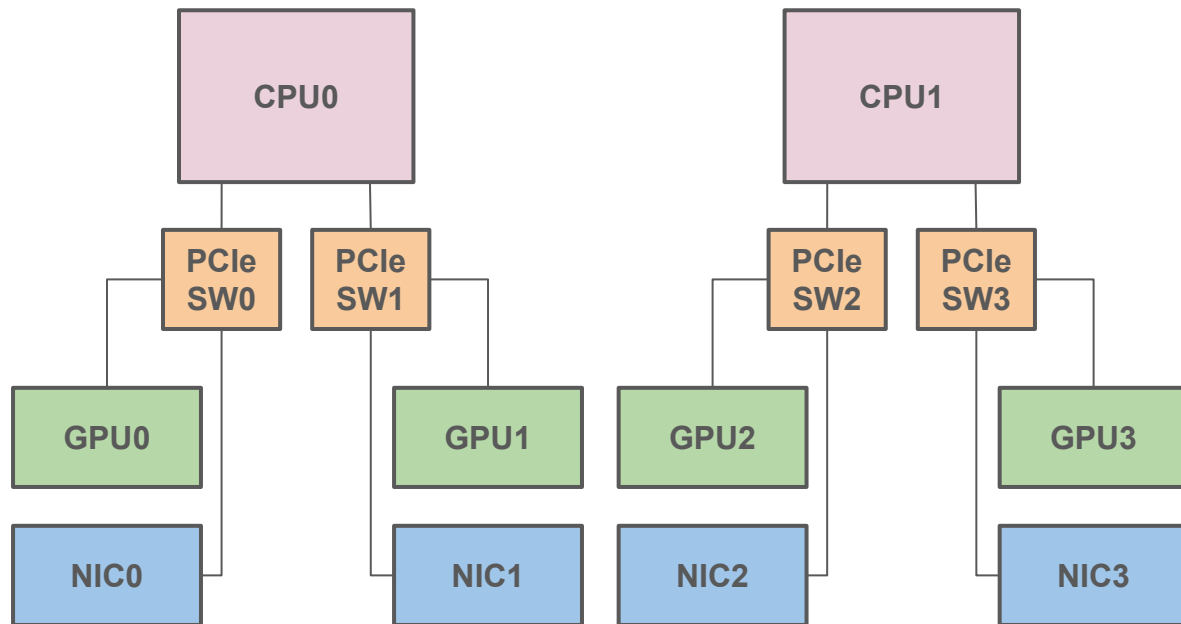
- 2 CPU
- 4 GPU
- 4 NICs
- Internal topology determines performance

# Simplified Node Topology - GPU / NIC

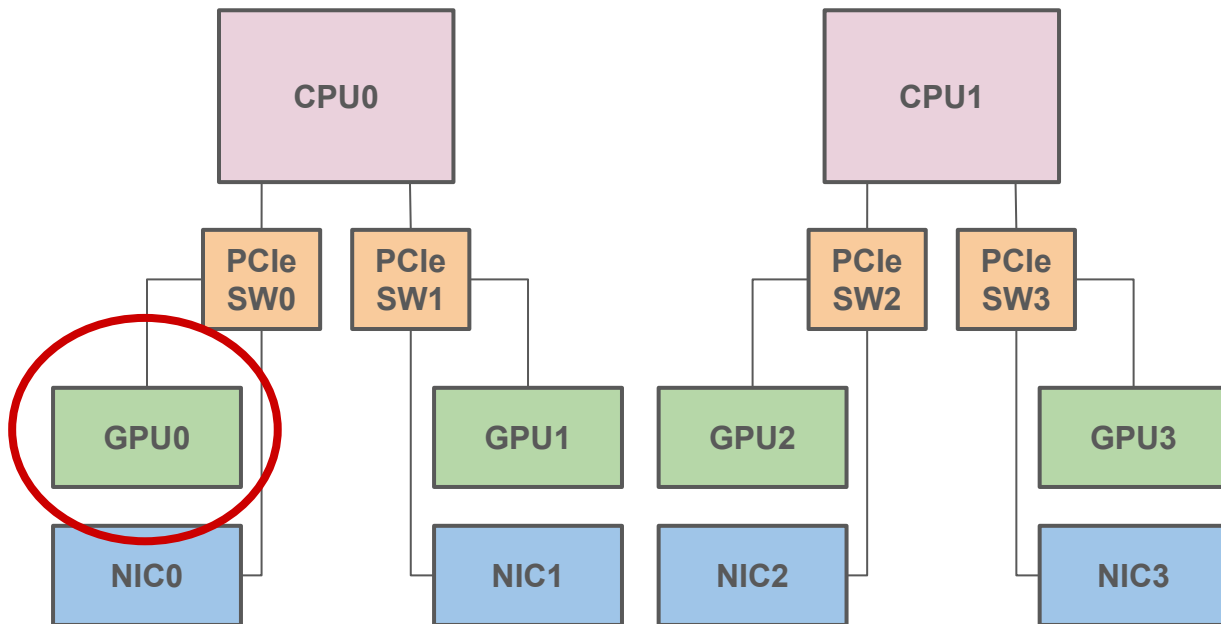


- 2 CPU
- 4 GPU
- 4 NICs
- Internal topology determines performance
- GPUDirect with GPU and NIC are on same root complex

# Allocating Devices with Device Plugin

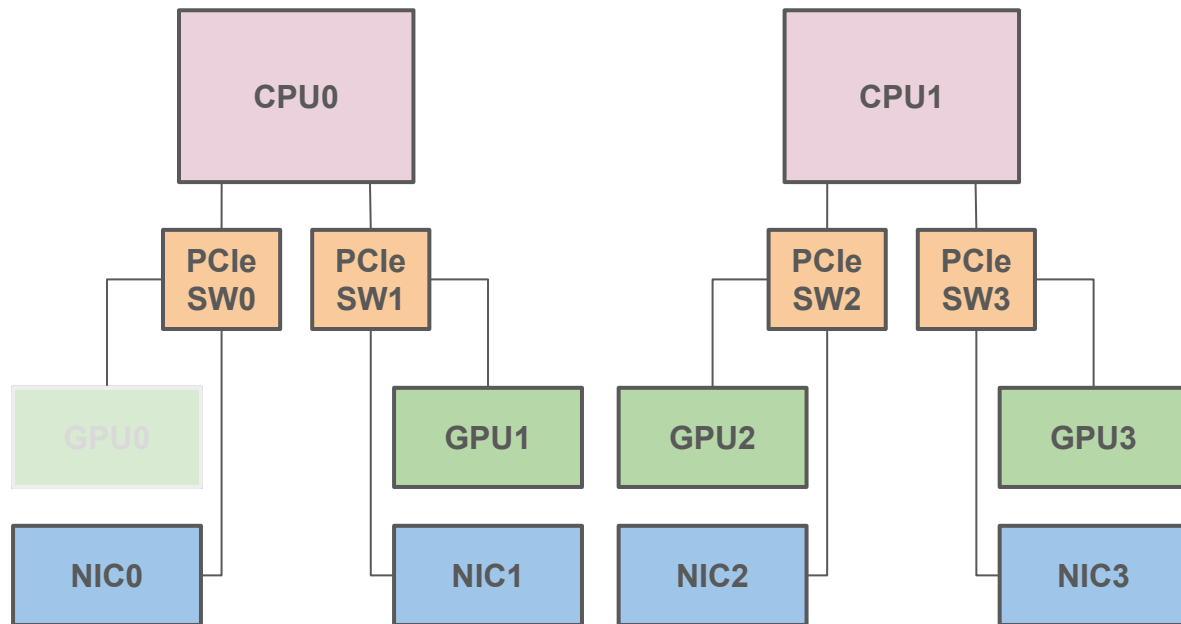


# Allocating Devices with Device Plugin



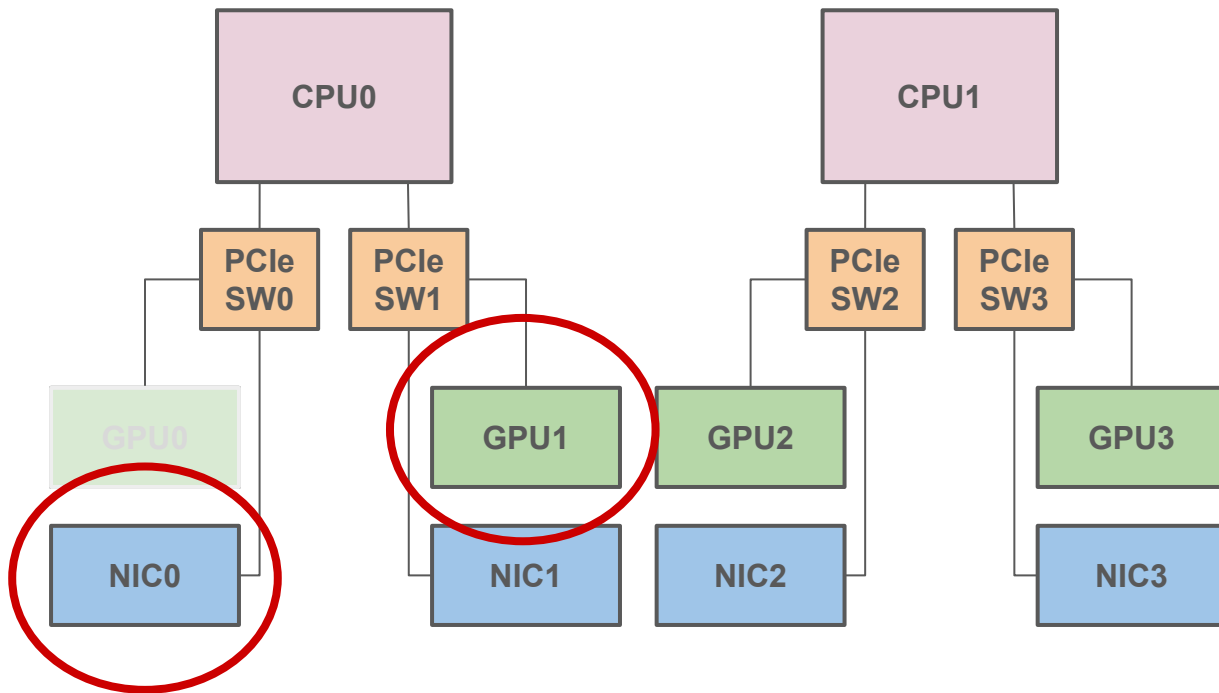
```
resources:  
limits:  
  nvidia.com/gpu: 1
```

# Allocating Devices with Device Plugin



```
resources:  
limits:  
  nvidia.com/gpu: 1
```

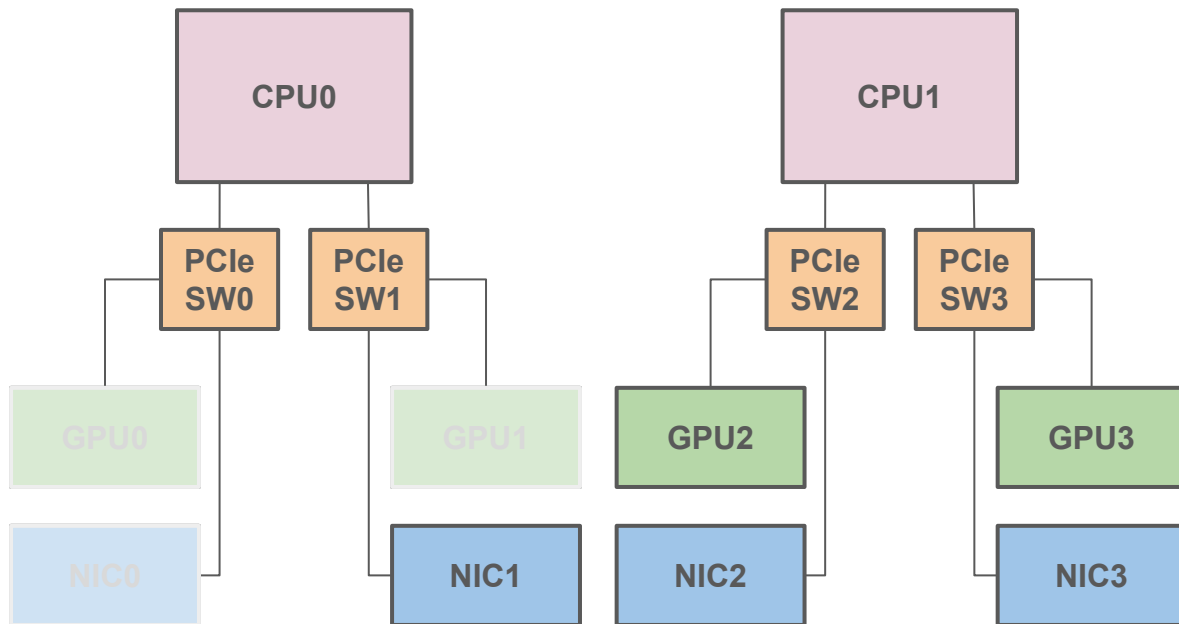
# Allocating Devices with Device Plugin



```
resources:  
limits:  
  nvidia.com/gpu: 1
```

```
resources:  
limits:  
  nvidia.com/gpu: 1  
  rdma/rdma_a: 1
```

# Allocating Devices with Device Plugin

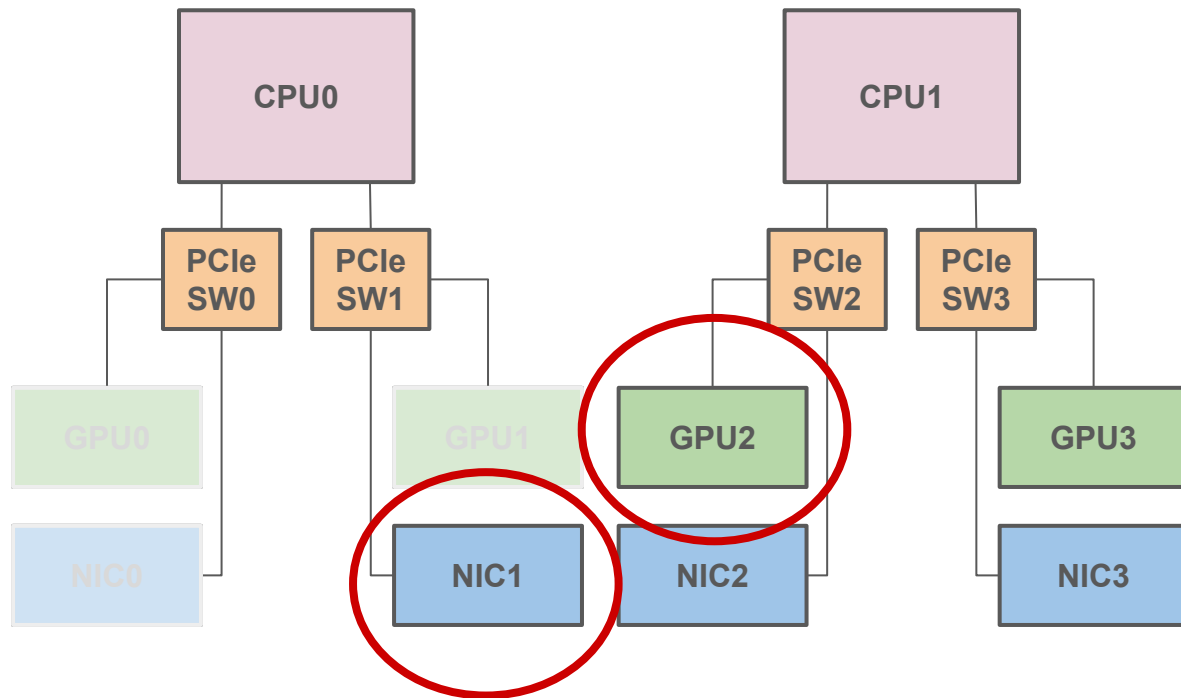


```
resources:  
limits:  
  nvidia.com/gpu: 1
```

```
resources:  
limits:  
  nvidia.com/gpu: 1  
  rdma/rdma_a: 1
```



# Allocating Devices with Device Plugin



```
resources:  
limits:  
  nvidia.com/gpu: 1
```

```
resources:  
limits:  
  nvidia.com/gpu: 1  
  rdma/rdma_a: 1
```

```
resources:  
limits:  
  nvidia.com/gpu: 1  
  rdma/rdma_a: 1
```

# Dynamic Resource Allocation (DRA) in Kubernetes



KubeCon



CloudNativeCon

North America 2024

```
apiVersion: resource.k8s.io/v1alpha2
kind: ResourceClaim
metadata:
  name: big-gpu-with-aligned-nic
spec:
  devices:
    requests:
      - name: gpu
        deviceClassName: gpu.nvidia.com
        selectors:
          - cel:
              expression: "device.capacity['memory'].compareTo(quantity('80Gi')) >= 0"
```

**Give me a GPU with  
at least 80GB of memory**

# Dynamic Resource Allocation (DRA) in Kubernetes



KubeCon



CloudNativeCon

North America 2024

```
apiVersion: resource.k8s.io/v1alpha2
kind: ResourceClaim
metadata:
  name: big-gpu-with-aligned-nic
spec:
  devices:
    requests:
      - name: gpu
        deviceClassName: gpu.nvidia.com
        selectors:
          - cel:
              expression: "device.capacity['memory'].compareTo(quantity('80Gi')) >= 0"

      - name: nic
        deviceClassName: rdma.nvidia.com
        selectors:
          - cel:
              expression: "device.attribute['sriovType'] == 'vf'"
```

**Give me a GPU with  
at least 80GB of memory**

**Together with an  
RDMA virtual function**

# Dynamic Resource Allocation (DRA) in Kubernetes



KubeCon



CloudNativeCon

North America 2024

```
apiVersion: resource.k8s.io/v1alpha2
kind: ResourceClaim
metadata:
  name: big-gpu-with-aligned-nic
spec:
  devices:
    requests:
      - name: gpu
        deviceClassName: gpu.nvidia.com
        selectors:
          - cel:
              expression: "device.capacity['memory'].compareTo(quantity('80Gi')) >= 0"

      - name: nic
        deviceClassName: rdma.nvidia.com
        selectors:
          - cel:
              expression: "device.attribute['sriovType'] == 'vf'"

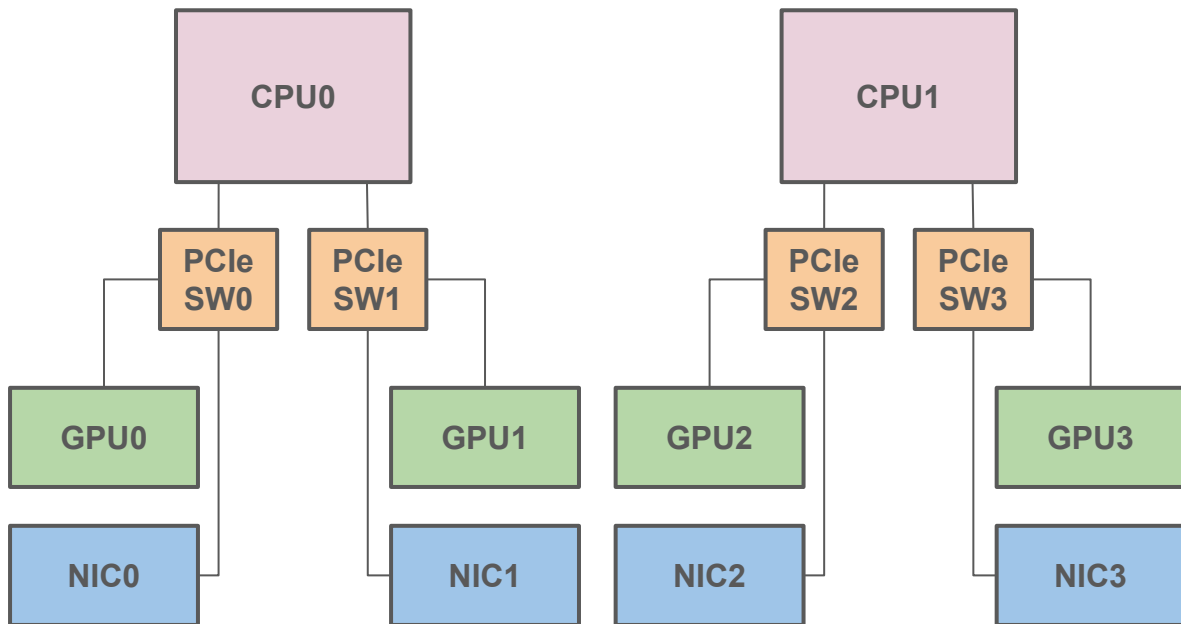
  constraints:
    - requestNames: ["gpu", "nic"]
      matchAttribute: k8s.io/pcieRoot
```

**Give me a GPU with  
at least 80GB of memory**

**Together with an  
RDMA virtual function**

**Make sure the GPU and NIC are aligned  
on the same PCIe root complex**

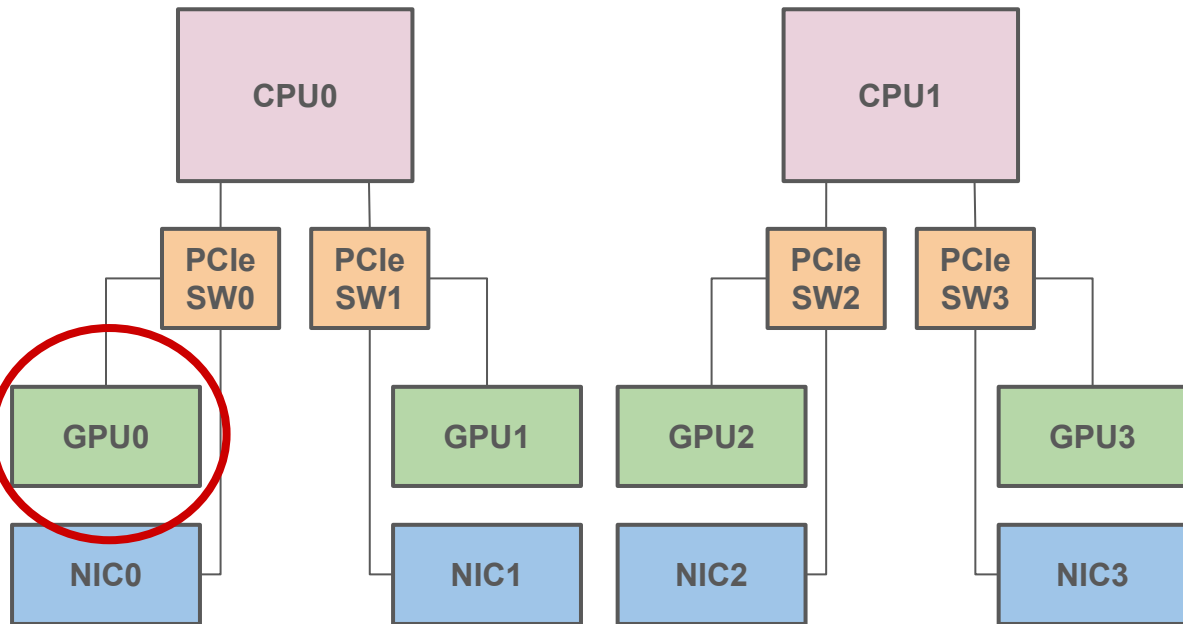
# Allocating Devices with DRA



# Allocating Devices with DRA

requests:

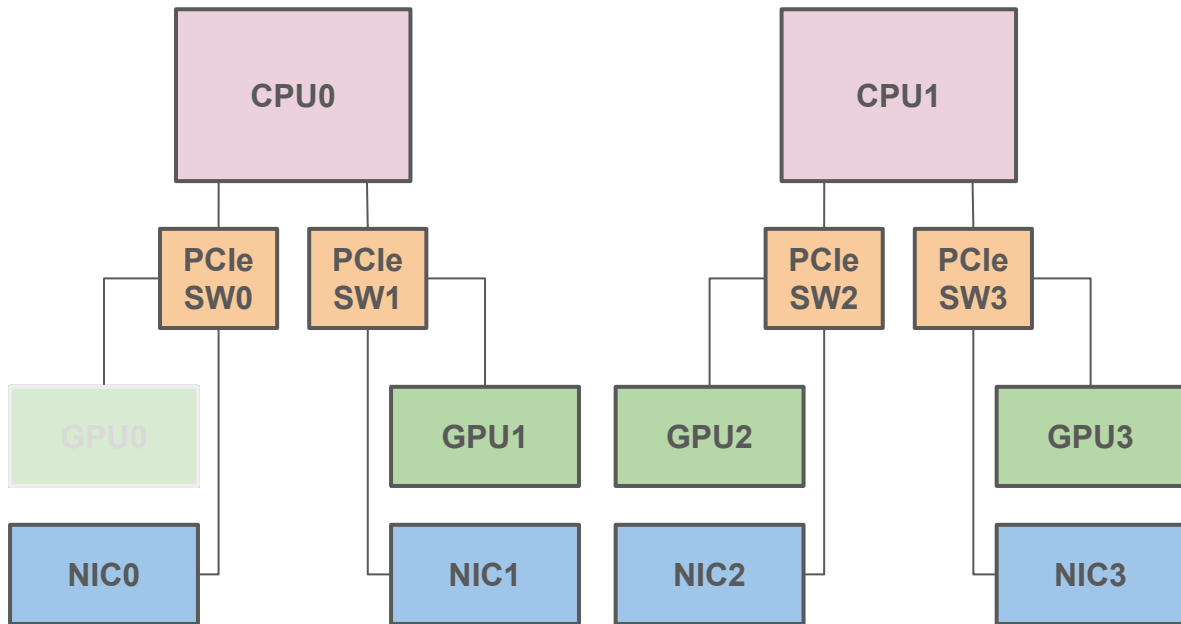
```
- name: gpu  
  deviceClassName: gpu.nvidia.com
```



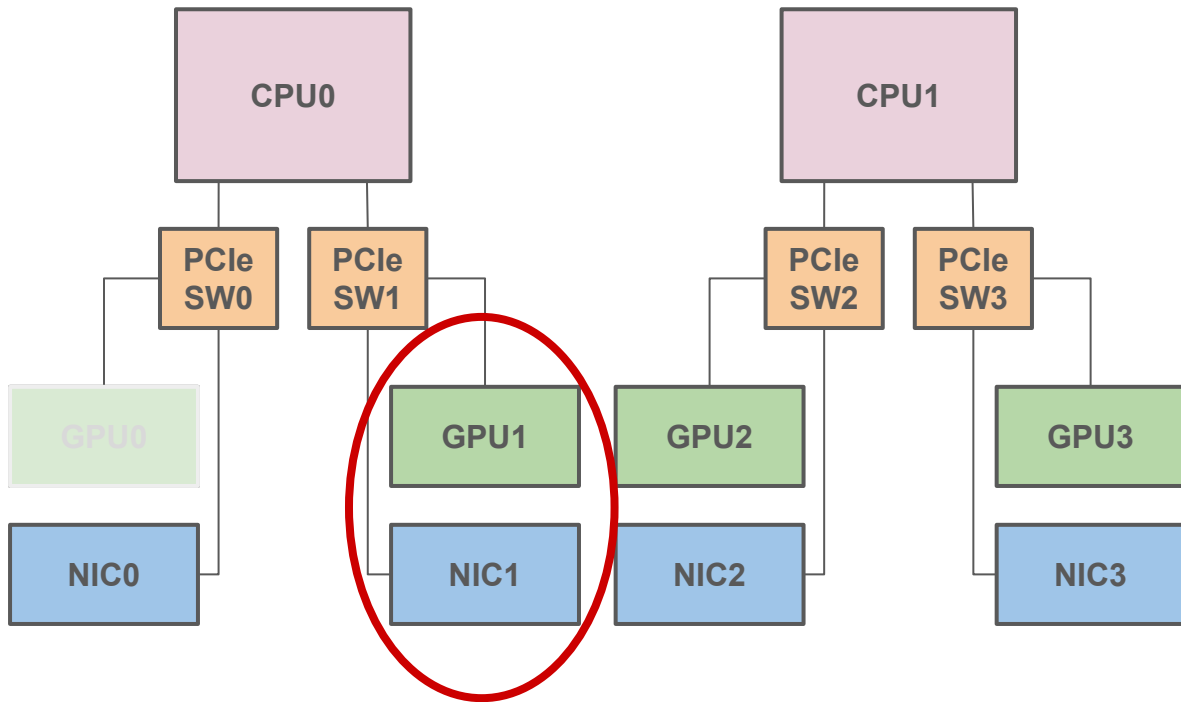
# Allocating Devices with DRA

requests:

```
- name: gpu  
  deviceClassName: gpu.nvidia.com
```



# Allocating Devices with DRA



requests:

- name: gpu  
deviceClassName: gpu.nvidia.com

requests:

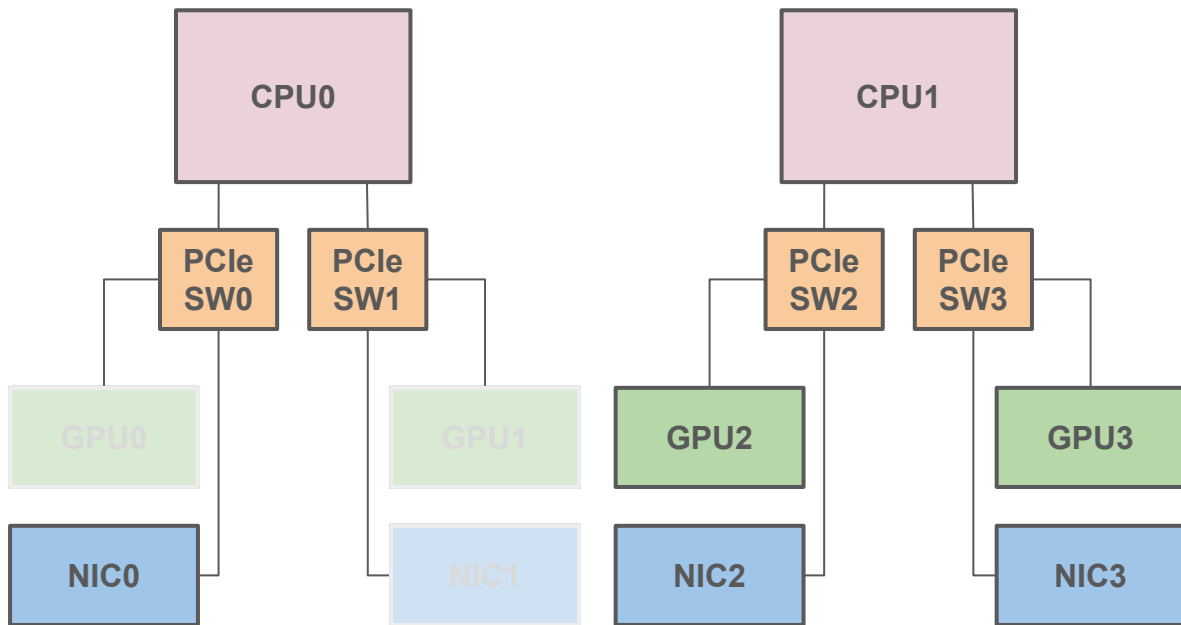
- name: gpu  
deviceClassName: gpu.nvidia.com
- name: nic  
deviceClassName: rdma.nvidia.com

constraints:

- requestNames: ["gpu", "nic"]  
matchAttribute: k8s.io/pcieRoot



# Allocating Devices with Device Plugin



requests:

- name: gpu  
deviceClassName: gpu.nvidia.com

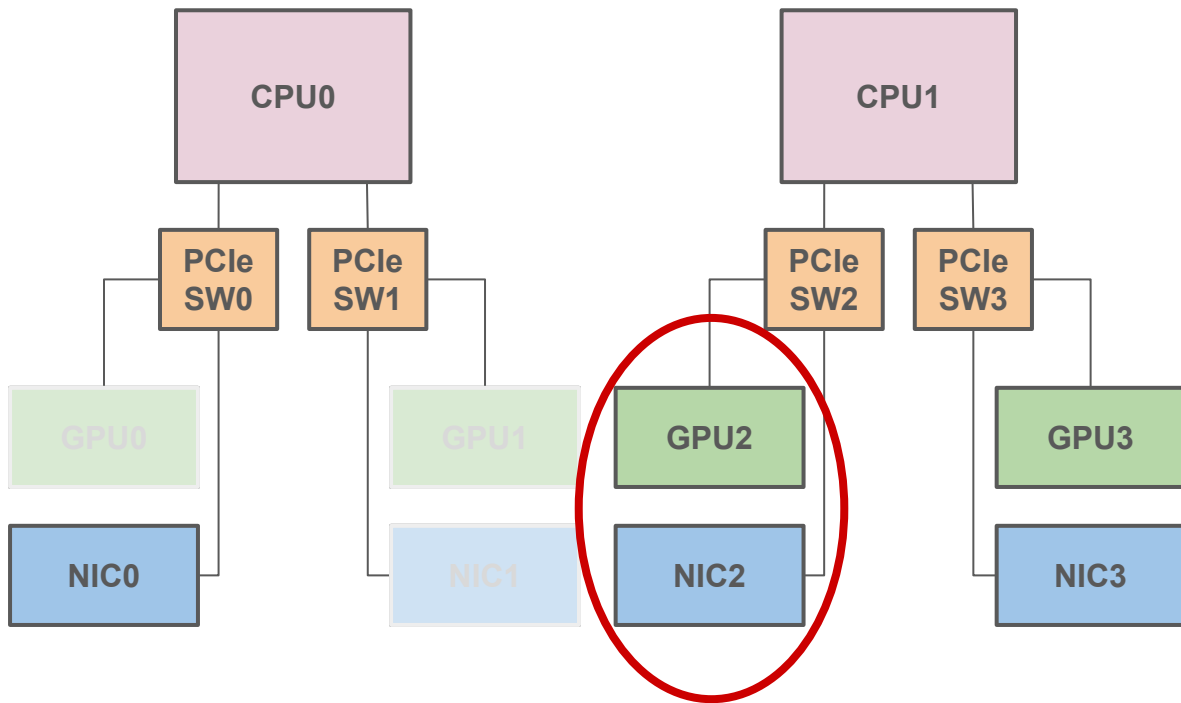
requests:

- name: gpu  
deviceClassName: gpu.nvidia.com
- name: nic  
deviceClassName: rdma.nvidia.com

constraints:

- requestNames: ["gpu", "nic"]  
matchAttribute: k8s.io/pcieRoot

# Allocating Devices with Device Plugin



requests:

- name: gpu  
deviceClassName: gpu.nvidia.com

requests:

- name: gpu  
deviceClassName: gpu.nvidia.com
- name: nic  
deviceClassName: rdma.nvidia.com

constraints:

- requestNames: ["gpu", "nic"]  
matchAttribute: k8s.io/pcieRoot

requests:

- name: gpu  
deviceClassName: gpu.nvidia.com
- name: nic  
deviceClassName: rdma.nvidia.com

constraints:

- requestNames: ["gpu", "nic"]  
matchAttribute: k8s.io/pcieRoot



KubeCon

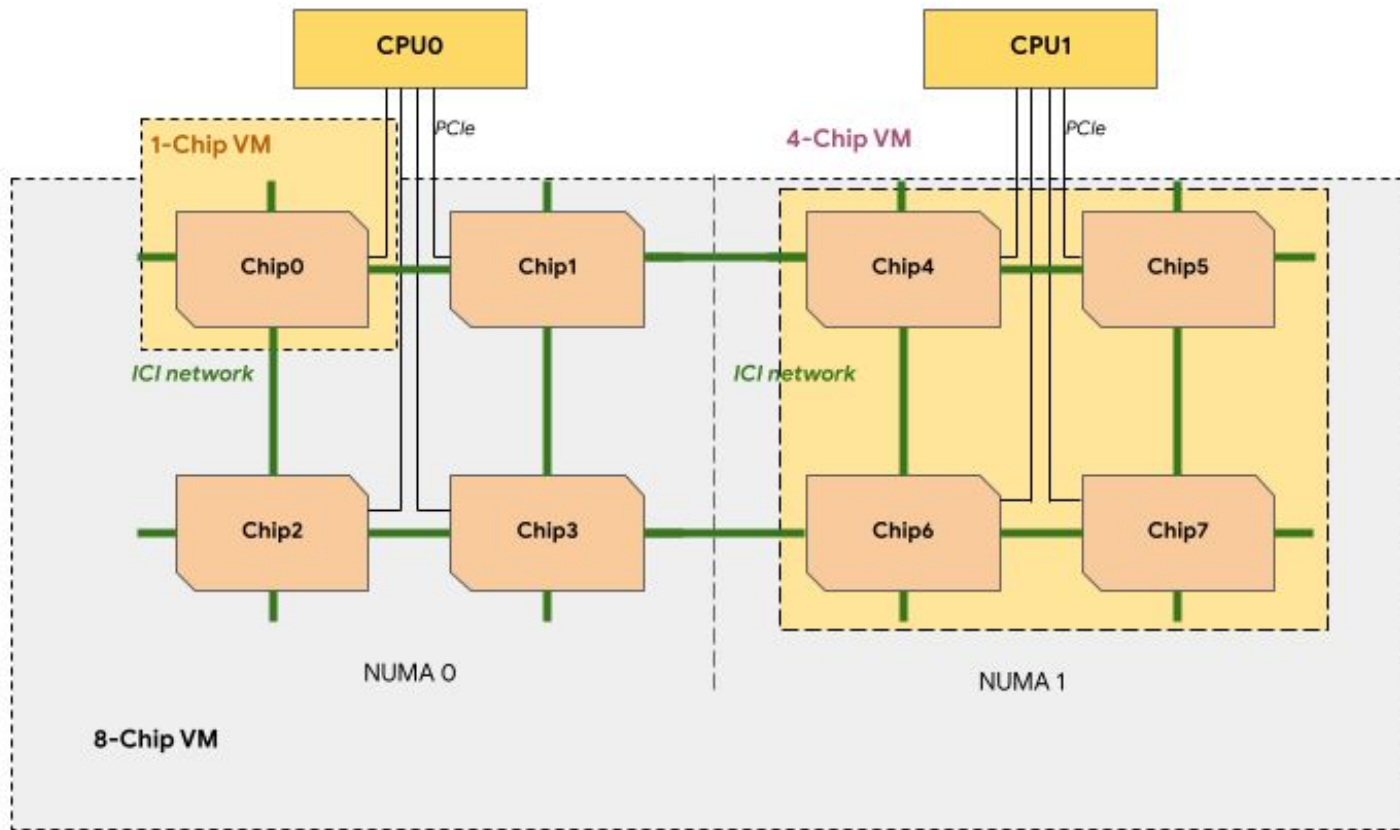


CloudNativeCon

North America 2024

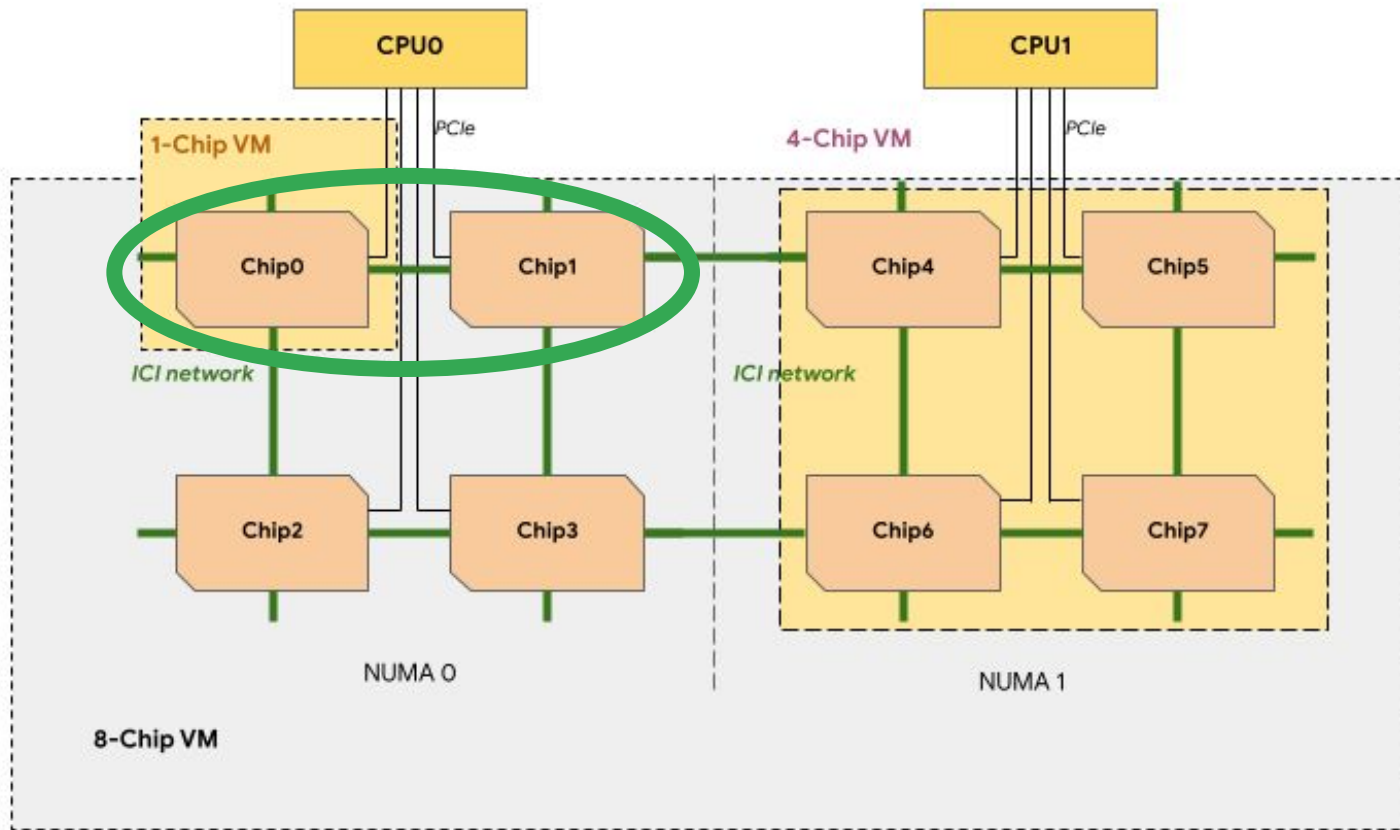
# What about TPU?

# TPU Host v5e Topology



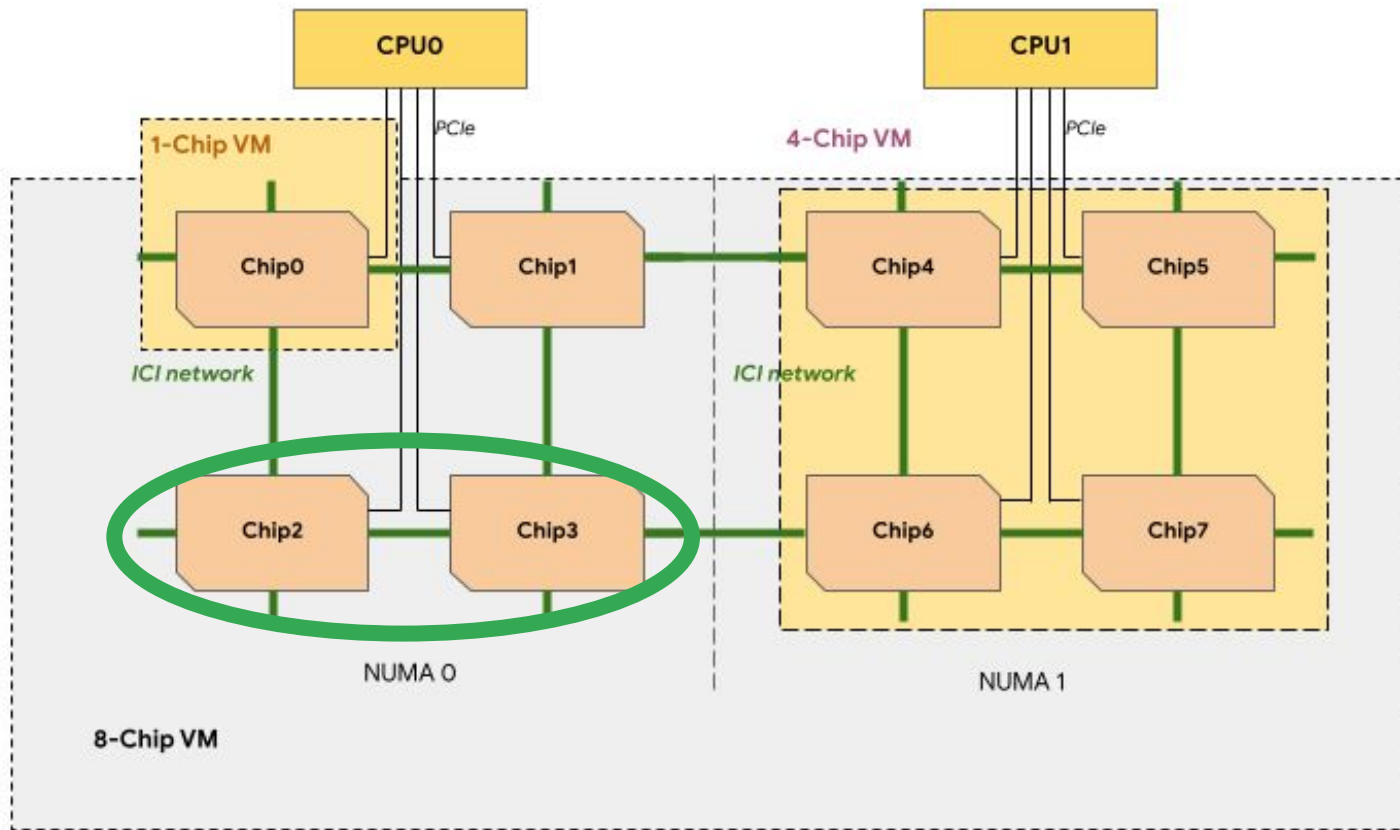
Source: <https://cloud.google.com/tpu/docs/v5e>

# TPU Host v5e Topology



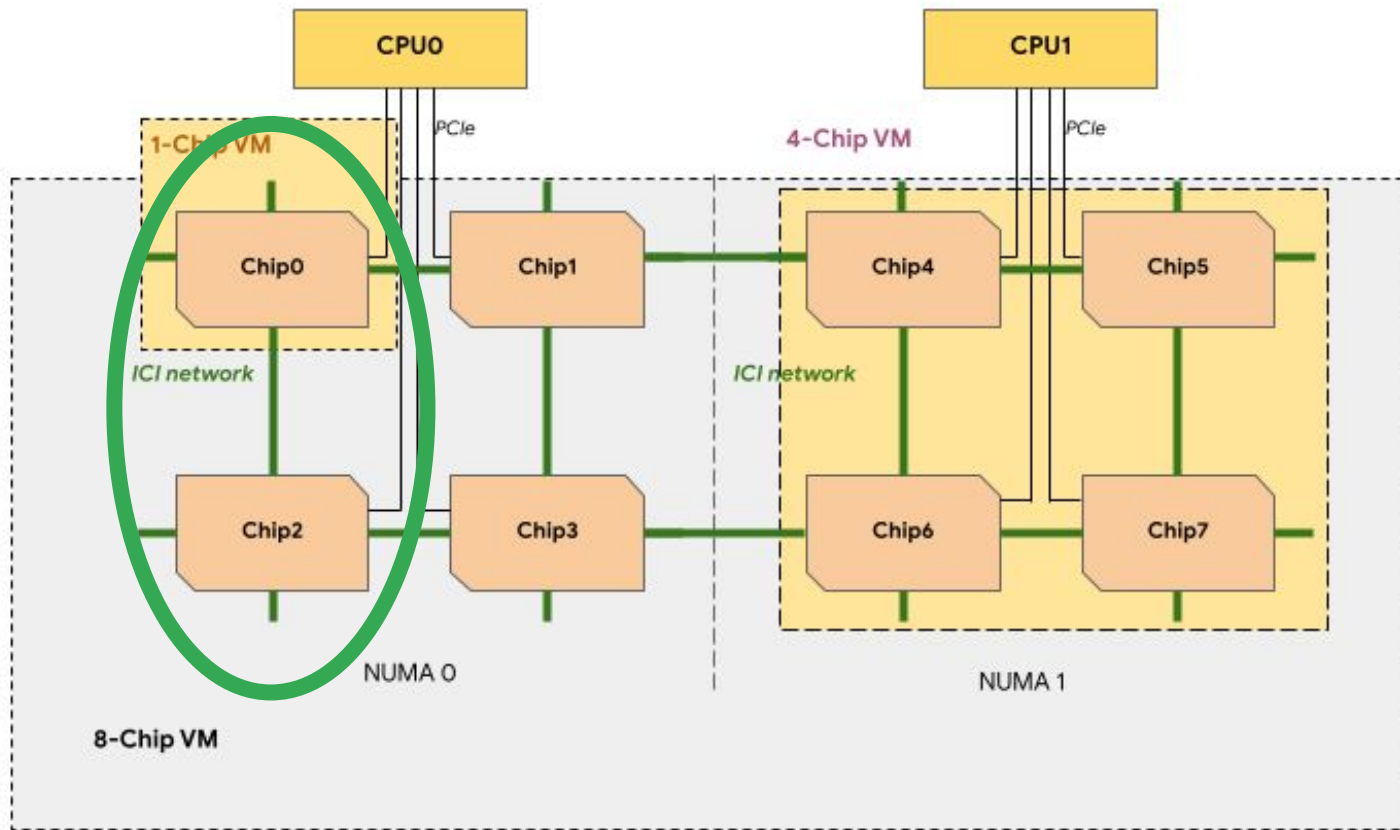
Source: <https://cloud.google.com/tpu/docs/v5e>

# TPU Host v5e Topology



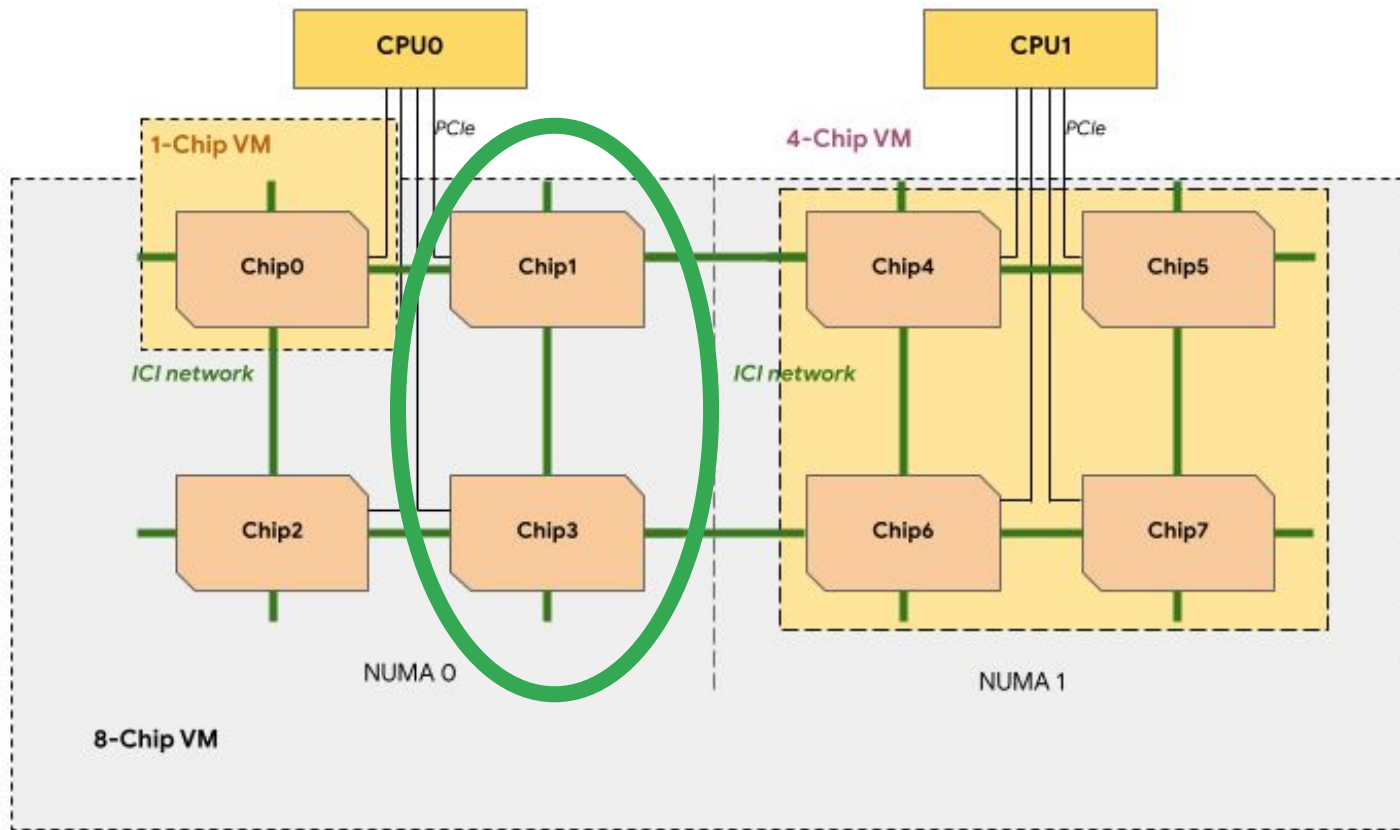
Source: <https://cloud.google.com/tpu/docs/v5e>

# TPU Host v5e Topology



Source: <https://cloud.google.com/tpu/docs/v5e>

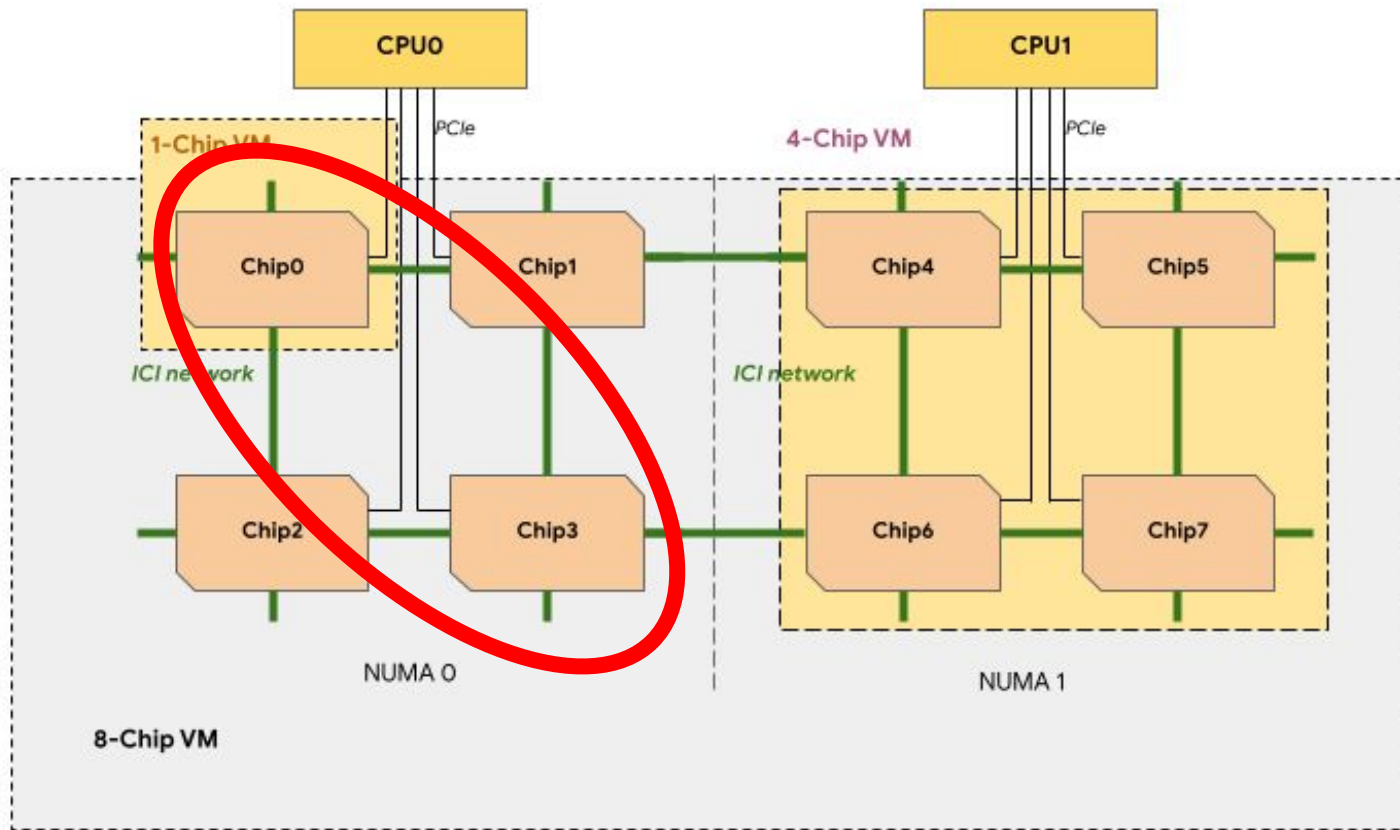
# TPU Host v5e Topology



Source: <https://cloud.google.com/tpu/docs/v5e>



# TPU Host v5e Topology



Source: <https://cloud.google.com/tpu/docs/v5e>

- **matchAttribute** constraints provides exact match semantics
- Can achieve *partial* support for various TPU topologies with **matchAttribute**

```
name: tpu0
attributes:
  2x1:
    string: west
  1x2:
    string: north
```

```
name: tpu1
attributes:
  2x1:
    string: east
  1x2:
    string: north
```

```
name: tpu2
attributes:
  2x1:
    string: west
  1x2:
    string: south
```

```
name: tpu3
attributes:
  2x1:
    string: east
  1x2:
    string: south
```

- User needs to ask for either **2x1** or **1x2**, not just **2**
- Could lead to scheduling failures

## matchExpression

- Proposed feature to use a CEL expression for a match constraint
- Also helps for AWS Neuron cores, which need ***sequential*** chips

All this is pretty hard on the user

- Is there something that can hide the complexity?

## matchExpression

- Proposed feature to use a CEL expression for a match constraint
- Also helps for AWS Neuron cores, which need ***sequential*** chips

All this is pretty hard on the user

- Is there something that can hide the complexity?

YES!

- [DRA: Add support for partitionable devices · Issue #4815](#)
- Alpha feature targeting 1.33
- Driver publishes valid topologies, consumed as a single device
- Much simpler for users - request single “device”
  - By size: tpu1x, tpu2x, tpu4x, tpu8x
  - Only valid topologies can be allocated
  - Underlying chips are counted under the hood
- Will support many vendor partitioning schemes:
  - AWS Neuron Cores
  - Google TPUs
  - Intel® Max Series
  - NVIDIA Multi Instance GPU (MIG)

# Other Talks, Feedback, and Questions

## Kubernetes WG Device Management - Advancing K8s Support for GPUs

*John Belamaric, Google; Patrick Ohly, Intel; Kevin Klues, NVIDIA*

## A Tale of 2 Drivers: GPU Configuration on the Fly Using DRA

*Alay Patel (NVIDIA), Varun Ramachandra Sekar US (NVIDIA)*

## Which GPU Sharing Strategy Is Right for You?

**A Comprehensive Benchmark Study Using DRA**

*Kevin Klues (NVIDIA), Yuan Chen (NVIDIA)*

## **Google Booth Lightning Talk**

**Deploying DRA for AI Infrastructure - Tech Talk & Ask the Experts Panel**

*Laura Lorenz (Google), Kevin Klues (NVIDIA), Tim Hockin (Google), John Belamaric (Google)*

Friday, 12:50pm - 1:05pm MST | Salt Palace | Google Cloud Booth

Feedback

