



**KubeCon**



**CloudNativeCon**

————— **North America 2024** —————





KubeCon



CloudNativeCon

North America 2024

# CoreDNS Plugins: A Deep Dive

*John Belamaric, Google & Yong Tang, Ivanti*

- Flexible DNS server written in Go
- Focus on service discovery
- Plugin based architecture, easily extended
- Default DNS server in Kubernetes
- Supports DNS, DNS over TLS, DNS over gRPC
- AWS Route53, Azure DNS, Google Cloud DNS


- 387 Contributors (Big Thanks!)
- 12,400+ Stars
- 36 Public Adopters
- 28 Maintainers
- Steering Committee

- 1.10.1- 1.11.4
  - **1.11.4** Released November 12, 2024
- New plugins:
  - **timeouts** - allows configuration of server listener timeout
- New features:
  - **acl** - drop queries as an action
  - **dnssec** - new option to load keys from AWS Secrets Manager
  - **template** - create responses with extended DNS errors
  - **cache** - option to serve original record TTLs from cache
  - **forward** - new option next, to try alternate upstreams when receiving specified response codes upstreams on (functions like the external plugin alternate)
  - **loadbalance** - add CNAME target rewrites to plugin
  - **dnstap** - add support for "extra" field in payload
  - **rewrite** - add support response code rewrite, and a new option to revert EDNS0 option rewrites in responses


# More cores != more QPS

- Long standing issue
- Test with cache (no external calls)
- Peaks at ~40k qps with 2 CPUs with this machine
- **What's the bottleneck?**

Cache QPS




# CoreDNS Internal Request Handling



- Internal structure “Server” handles a given port
- Dispatches requests to go routines


# CoreDNS Internal Request Handling



- What happens if we split traffic across ports?
- Aggregate QPS goes up!
- But we can't really force users to use two different ports




# Introducing the multisocket Plugin



- SO\_REUSEPORT
- Multiple sockets on the same port
- kernel distributes packets across the sockets
- Internal changes to create multiple servers for a port (one per socket)
- uber-go/automaxprocs


# Multisocket Plugin

Cache QPS




# Multisocket Plugin

## Cache QPS



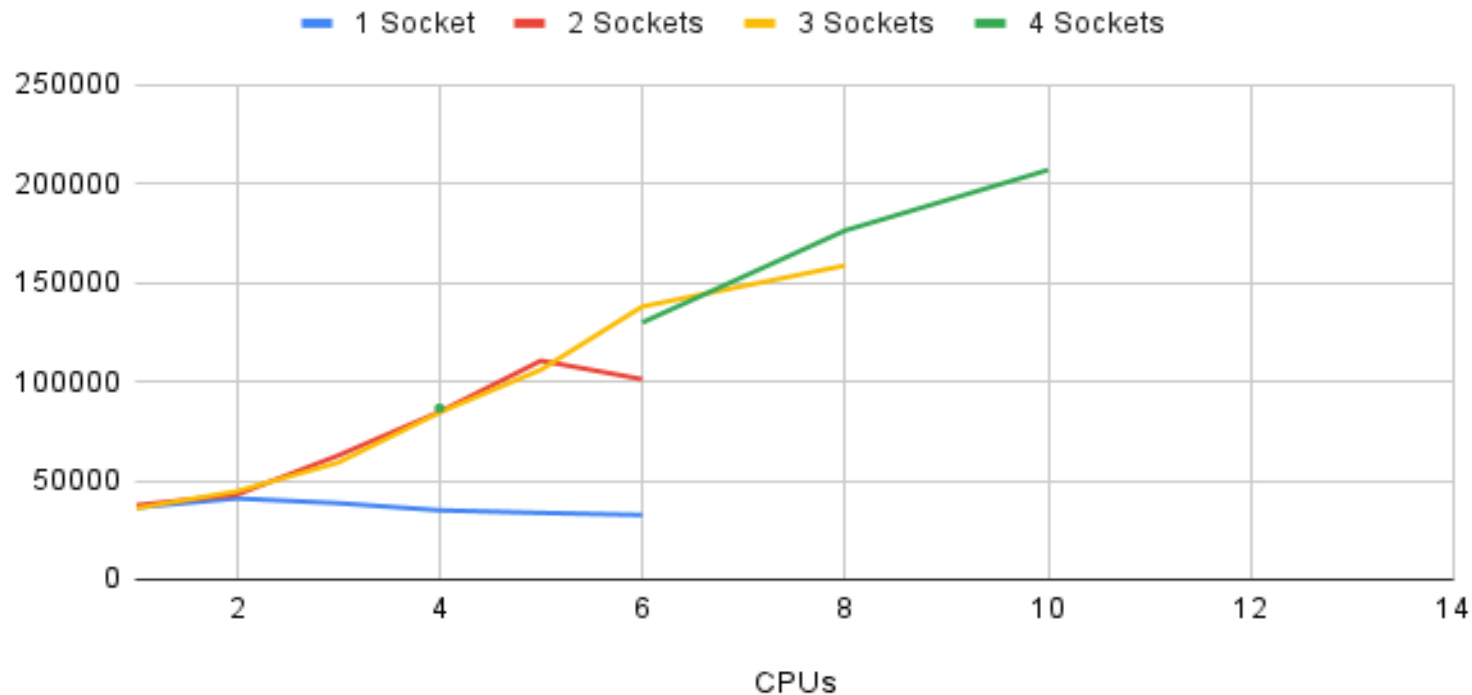
# Multisocket Plugin

## Cache QPS




# Multisocket Plugin

## Cache QPS




# Multisocket Plugin

## Cache QPS




# Multisocket Plugin

## Cache QPS



# Multisocket Plugin


Cache QPS





- Merged in master
- Will be in 1.12
- Targeting release O(weeks)
- Requires SO\_REUSEPORT
- Config
  - multisocket [ NUM ]
- Default NUM == GOMAXPROCS
- May see better performance in specific scenarios with different NUM
  - But GOMAXPROCS is a good default

- Source IP based service discovery
  - Returns "1.1.1.1" for 172.0.0.0/8 or 127.0.0.0/8
  - Returns "8.8.8.8" otherwise



[setup.go]:

`init()`

- Performs one-time initializations, register the setup function with Caddy

`setup(c *caddy.Controller) error`

- Parses the configuration from the file and captures it in a struct
- Adds the handler to the Config object (represents a stanza)
- Called once for each use of the plugin in the Corefile

[demo.go]:

`ServeDNS(context.Context, dns.ResponseWriter,*dns.Msg) (int, error)`

- Processes the DNS request and returns a response, or,
- Passes it down the chain

# Demo Plugin – init()/setup()

```
func init() {  
    caddy.RegisterPlugin("demo", caddy.Plugin{  
        ServerType: "dns",  
        Action:      setup,  
    })  
}  
func setup(c *caddy.Controller) error {  
    c.Next() // 'demo'  
    if c.NextArg() {  
        return plugin.Error("demo", c.ArgErr())  
    }  
    dnsserver.GetConfig(c).AddPlugin(func(next plugin.Handler) plugin.Handler {  
        return Demo{}  
    })  
    return nil  
}
```

# Demo Plugin – ServeDNS()

```
// ServeDNS implements the plugin.Handler.ServeDNS.
func (p Demo) ServeDNS(ctx context.Context, w dns.ResponseWriter, r *dns.Msg) (int, error) {
    state := request.Request{W: w, Req: r}
    qname := state.Name()
    reply := "8.8.8.8"
    if strings.HasPrefix(state.IP(), "172.") || strings.HasPrefix(state.IP(), "127.") {
        reply = "1.1.1.1"
    }
    fmt.Printf("Received query %s from %s, expected to reply %s\n", qname, state.IP(),
reply)
    answers := []dns.RR{}
    ...
}
```

```
.:1053 {  
    #  
    # By default all plugins  
    # are disabled initially,  
    # unless enabled explicitly  
    #  
    demo  
}
```

# Demo Plugin – Build

```
$  
$ # add demo:demo to plugin.cfg  
$  
$ # build with docker (golang:1.22)  
$ docker run --rm -i -t \  
$   -v $PWD:/go/src/github.com/coredns/coredns \  
$   -w /go/src/github.com/coredns/coredns \  
$   golang:1.22 sh -c \  
$       'GOFLAGS="-buildvcs=false" make gen && GOFLAGS="-buildvcs=false" make'  
$  
$ # configure Corefile and run coredns  
$  
$ ./coredns  
$
```

<https://github.com/coredns/demo>



- Star CoreDNS in GitHub:
  - <https://github.com/coredns/coredns>
- Add name to ADOPTERS.MD
- Becomes a contributor:
  - Create a pull request
- Becomes a maintainer:
  - One significant pull request
  - Sponsored by one current maintainer



**KubeCon**



**CloudNativeCon**

North America 2024

# THANK YOU

