



KubeCon

CloudNativeCon

North America 2024





KubeCon



CloudNativeCon

North America 2024

Love thy (Noisy) Neighbor

Strategies for Mitigating Performance Interference in
Cloud-Native Systems

Run 20%-50%
more transactions

Reduce tail latency
by 20-80%





BubbleUp



Heracles



PerfIso



Caladan



Themis



iBench



Dirigent



CPI²



CAT at Scale



PARTIES



Can be made
widely useful

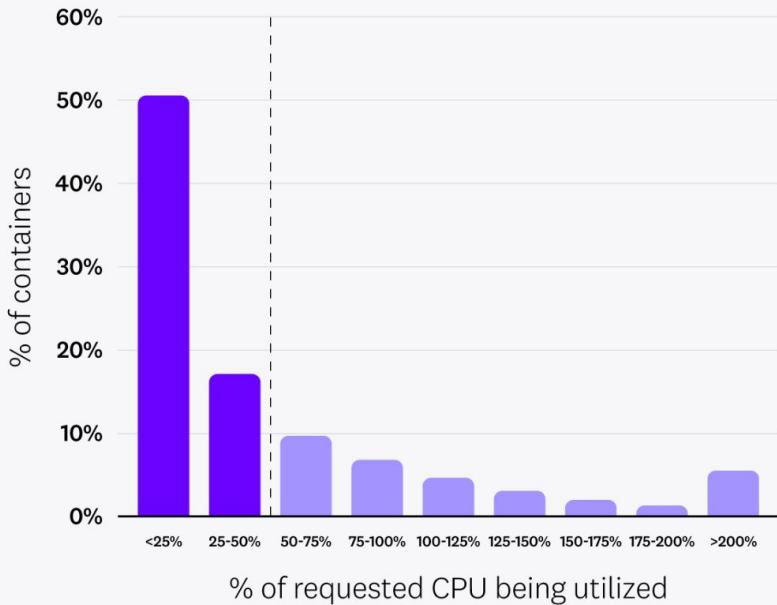
Libra



Alita



Usage of requested CPU



**>65% OF
CONTAINERS
USE LESS
THAN HALF OF
REQUESTED CPU**

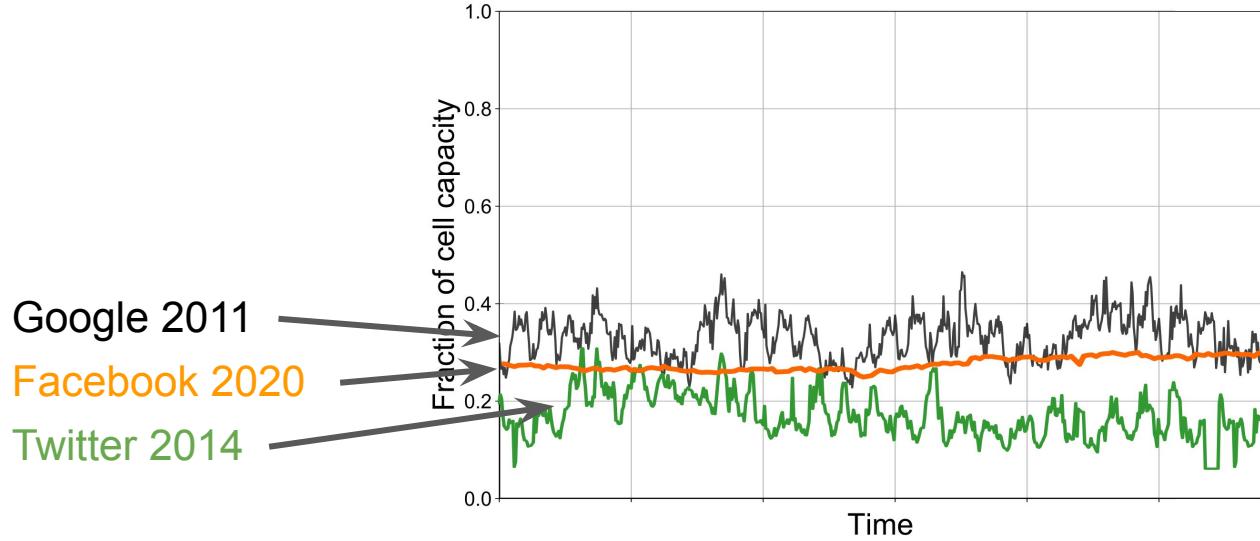
Source: Datadog

[source](#)

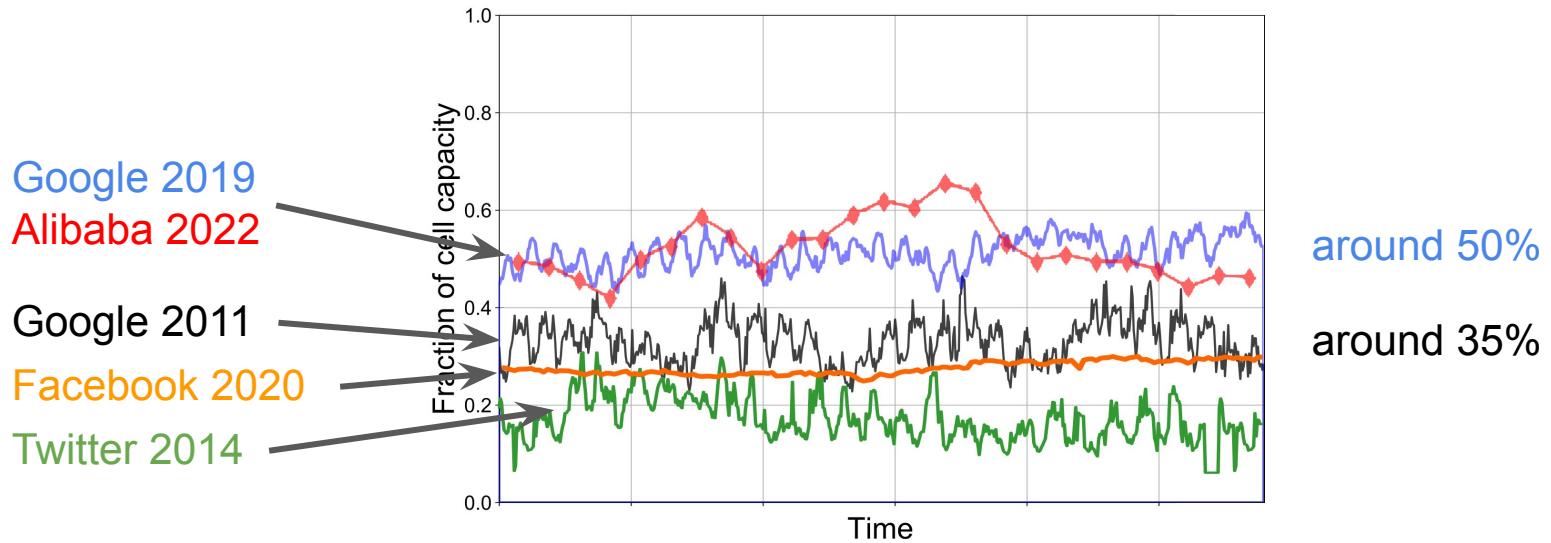
Raise hand if:

- Know prod cluster avg. CPU utilization
- Above 20%?
- Above 30%?
- Above 40%?
- Above 50%?

Hyperscaler CPU was low...



Then.. breakthrough?



Vertical autoscaler

Google's Autopilot (EuroSys'20, not GKE)
 StormForge, PerfectScale, FairWinds, ...

Less wasted resources → Better packing

Handle noisy neighbor

- Reduce cycles per request
- Improve tail latency

→ Less hardware at better SLA

This talk

The
Problem

Current
Support

Mitigation
Systems

Next
Steps

~~deep dive to 1-2 papers~~ broad overview

Problem

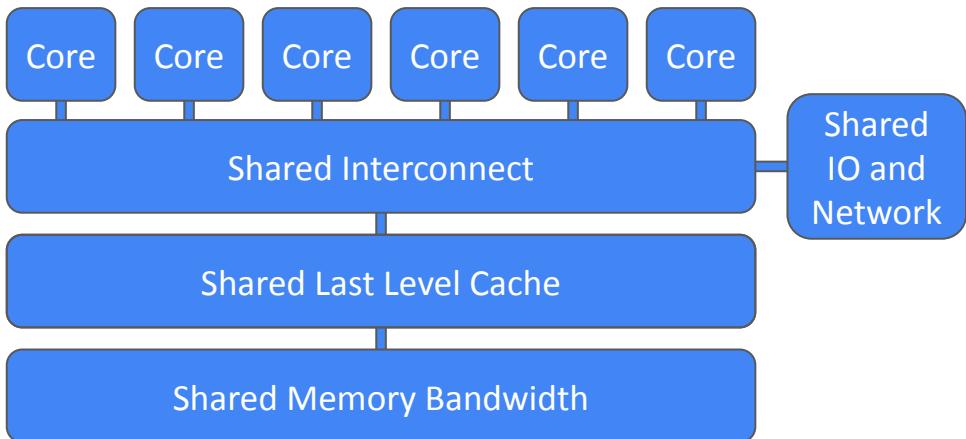
- Memory noisy neighbor
- How does this affect my Pods?
- Do I have it in my cluster?

Noisy neighbor: tragedy of the commons

Apps access physical resources



Shared resources are constrained



Noisy neighbor: tragedy of the commons

Apps access physical resources

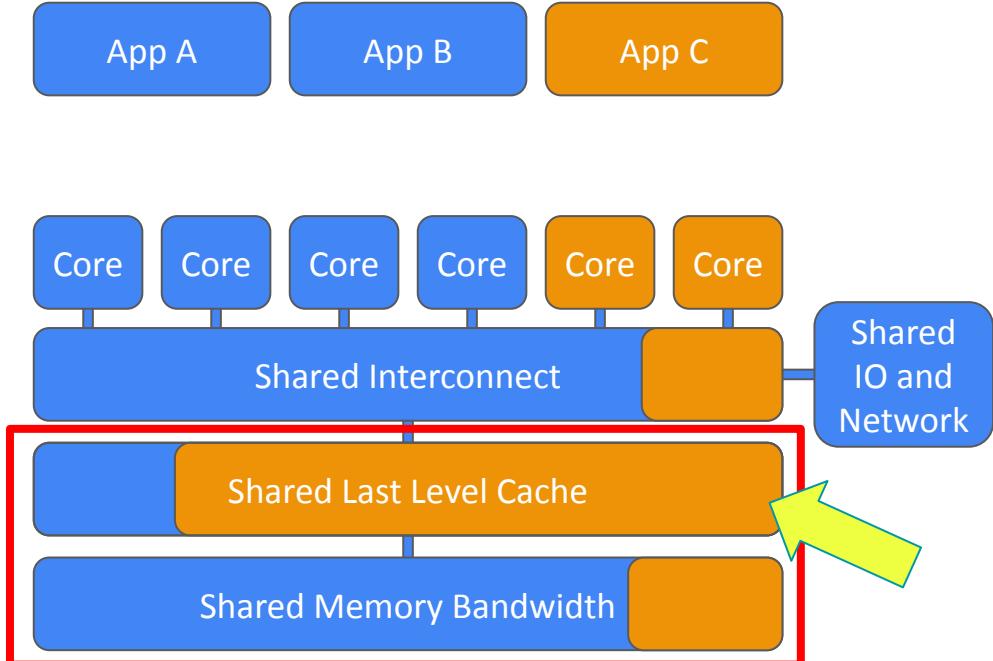


Shared resources are constrained

One App can use more than its fair share, degrading others

This talk:

- Cache
- Memory bandwidth

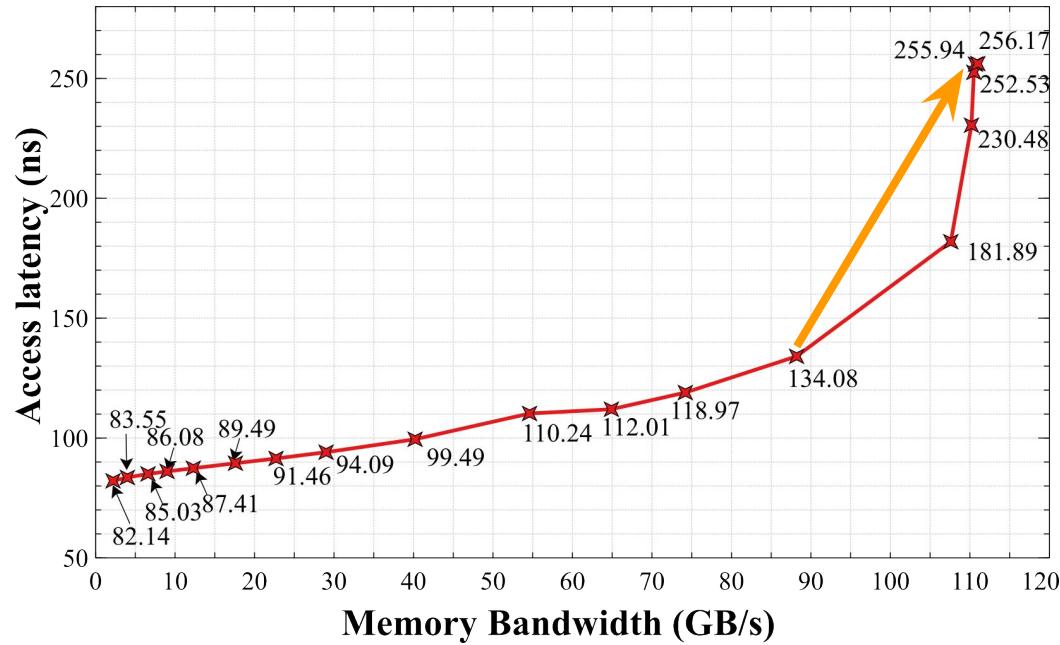


Latency increases with memory bandwidth

Change memory bandwidth,
Measure latency

Knee-point around 80%

80% → 100% bandwidth
latency doubles!



CPUs try to protect us from memory latency:

- Prefetchers
- Reorder buffer (ROB)
- Caches



Are we protected?



...or not?

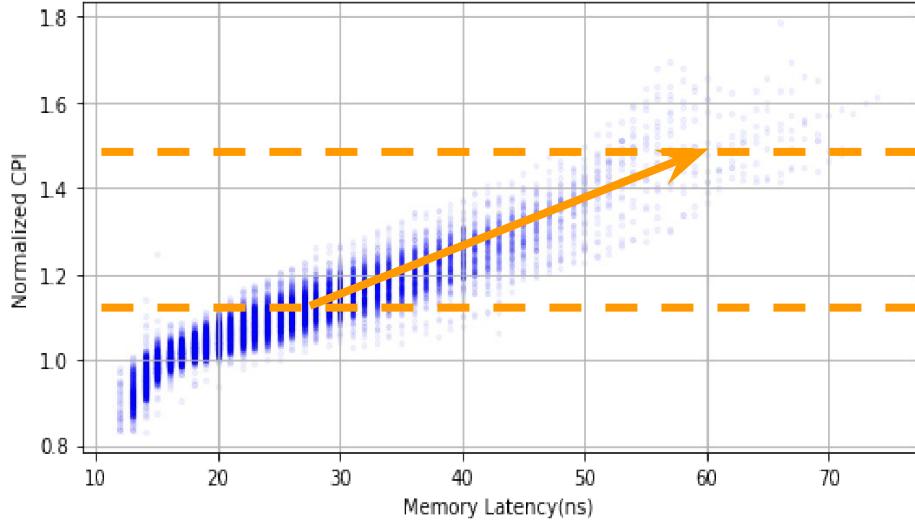
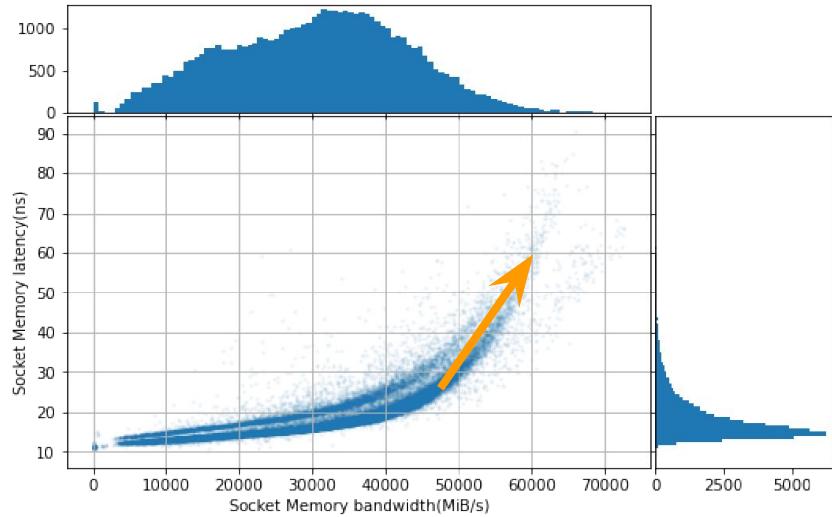
Measuring slowdown: CPI

Does memory latency cause CPU memory stalls?

Cycles Per Instruction – **CPI**

CPU waits for memory → Stall cycles → High CPI

80% bandwidth cap → 25% more compute-efficient



Memory bandwidth only – has cache mitigation

Tail latency can explode with noisy neighbors!



Google, 2015

3 production services

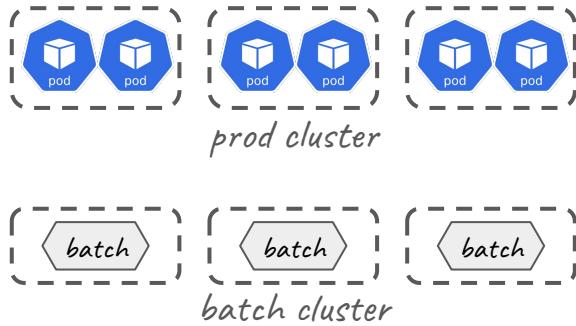
get % of SLO target
(99th / 95th percentile)

w/ synthetic noise
generators

Please raise hand if:

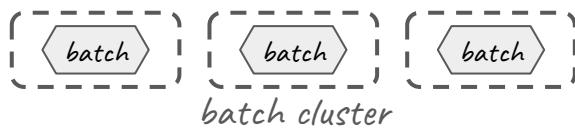
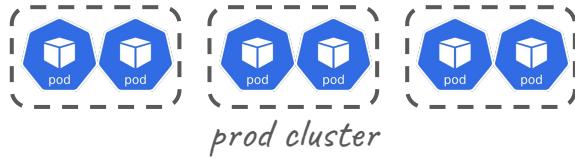
- Know what VMs or bare-metal are used in prod
- Never use fraction of physical CPU

Separating batch clusters



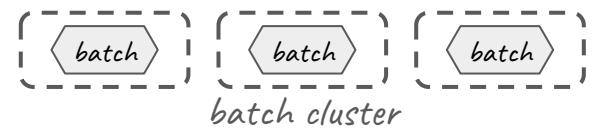
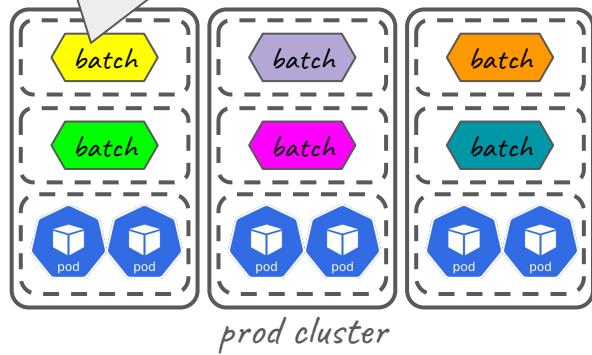
What I think I'm running

Separating batch clusters



What I think I'm running

jobs from some random tenants

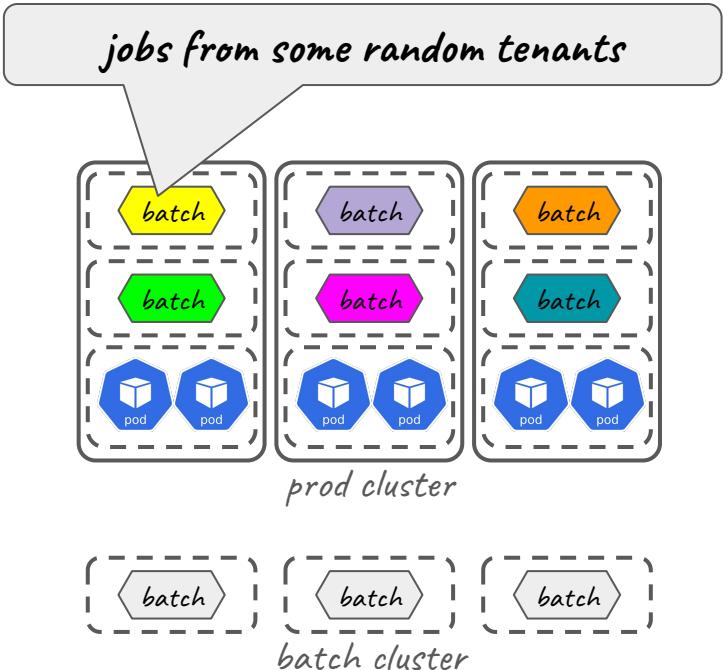


What I'm actually running

Separating batch clusters



Does your provider protect your VMs from other VMs on the same machine?



What I'm actually running



Engineers spend years
optimizing user experience

Does my cluster have noisy neighbors?

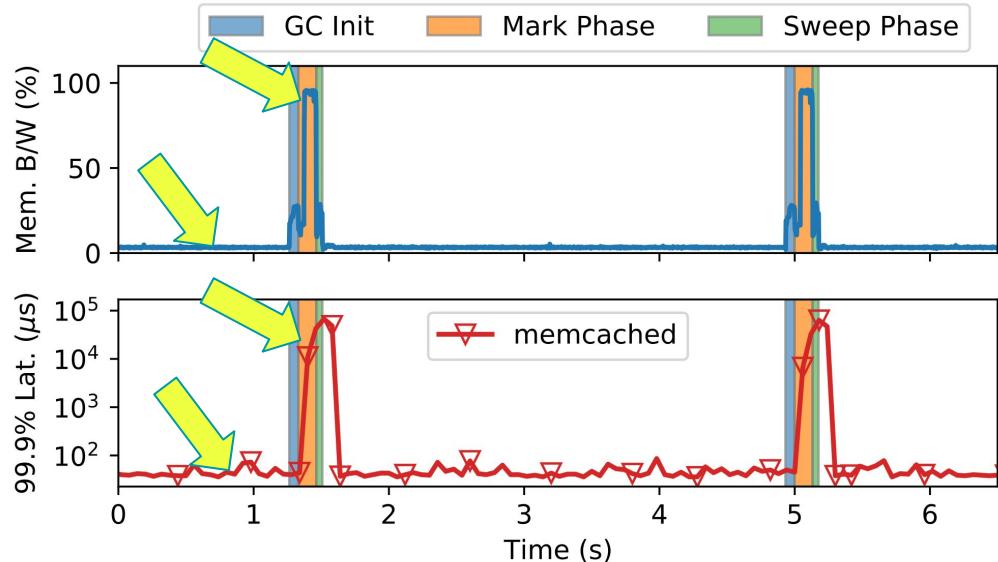
Run:

- Memcached
- garbage-collected workload

Mark Phase is memory-intensive,
causes significant slowdown!

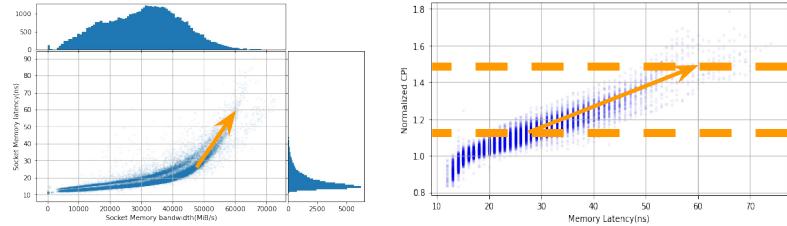
Not only “big-data” is noisy

Also: security scanning, video
streaming, transcoding...



Recap: The Problem

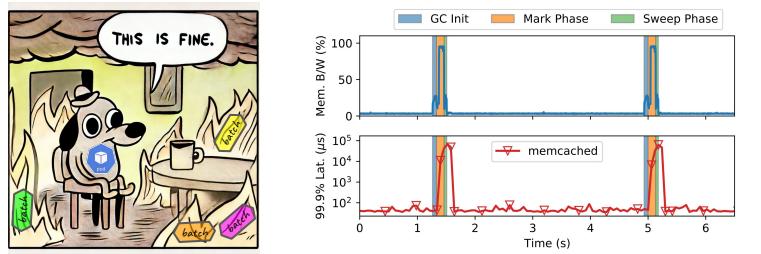
High memory contention
→ High memory latency
→ High CPI (low efficiency)



Tail latency with noisy neighbor increases 4-13x



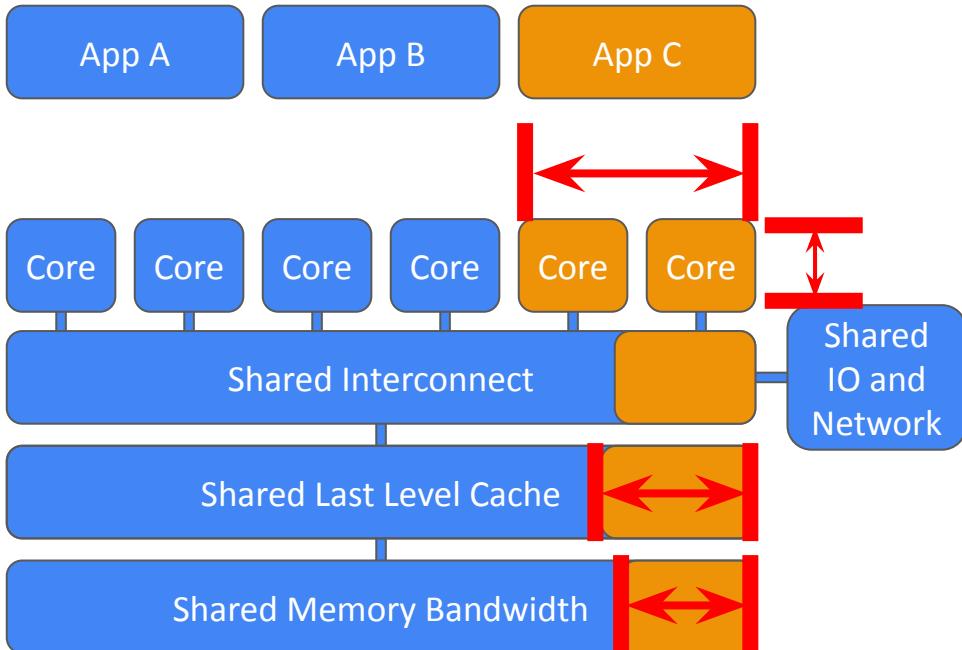
Many workloads can be noisy (e.g., garbage collection)



Current support

- Hardware and software mitigation
- Support in Linux

Memory system mitigation knobs



Reduce demand (#cycles):

- Core pinning
- Frequency scaling

Direct control:

- Cache allocation
- Memory bandwidth limits

Containers (not) to the rescue!

Different focus:

- cycles
- memory capacity (size)

| Cgroup v2 | Responsibility |
|------------|--|
| CPU | Regulates distribution of CPU cycles |
| Memory | Controls distribution and accounting of memory usage |
| IO | Manages distribution of IO resources |
| PID | Limits the number of processes in a cgroup |
| Cpuset | Assigns specific CPU(s) and memory nodes |
| Device | Controls access to device files |
| RDMA | Regulates distribution and accounting of RDMA resources |
| HugeTLB | Limits the HugeTLB usage per control group |
| Perf_event | Allows perf events to be filtered by cgroup path |
| Misc | Provides resource limiting for scalar resources not covered by other controllers |

Linux supports allocation

`egroup resctrl` allocates:

- Memory bandwidth
- Cache space

Evolved:

- Intel (2016)
- AMD (2018)
- ARM (in progress)

Demo:

- Facebook Resource Control Demo by Tejun Heo and collaborators

Linux resctrl CPU support

| | Kernel |
|--|---|
| Intel RDT Intel Resource Director Technology | v4.10 , v4.12 |
| AMD QoS AMD Platform Quality of Service | v5.0 |
| ARM MPAM Memory System Resource Partitioning and Monitoring | WIP , x86→generic Review needed |

Activities Terminal

Oct 2 5:57 PM



fish /home/htejun

Welcome to fish, the friendly interactive shell
Type `help` for instructions on how to use fish
htejun@slm ~> █

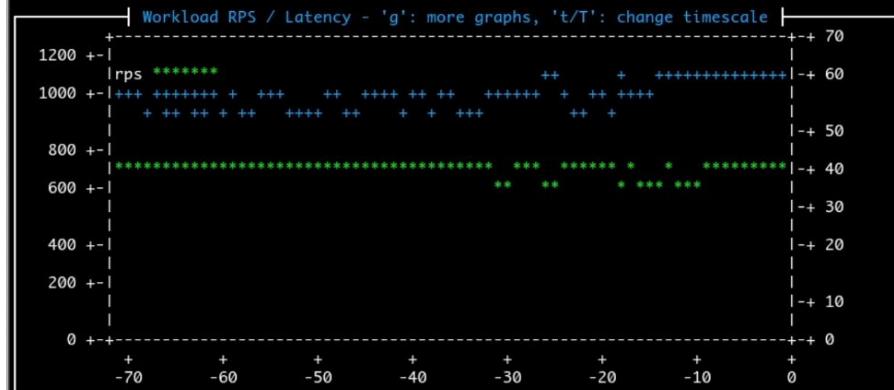
█

Connecting

demo@resctl-demo:~

Facebook Resource Control Demo - 'q': quit

```
[Running] 2020-10-02 10:02:15 PM
[config] satisfied: 17 missed: 0
[oomd] workload: +pressure -senpai system: +pressure -senpai
[sideload] jobs: 1/ 1 failed: 0 cfg_warn: 0 -overload -crit
[sysload] jobs: 0/ 0 failed: 0
[workload] load: 61.6% lat: 60ms cpu: 46.8% mem: 51.1g io: 7.3m
```



Management logs

```
[22:01:13 rd-agent] [INFO] side: "compile-job" started
[22:01:14 rd-agent] [INFO] svc: "rd-sysload-memory-hog.service" started (Running)
[22:01:32 rd-sideloader] OVERLOAD: end, resuming normal operation
[22:01:33 rd-sideloader] JOB: Starting rd-sideload-compile-job.service
[22:01:33 rd-sideloader] Running as unit: rd-sideload-compile-job.service
[22:01:54 rd-agent] [INFO] svc: "rd-sysload-memory-hog.service" transitioned from Running to Other("deactivating:stop-sigterm")
[22:01:54 rd-agent] [INFO] svc: "rd-sysload-memory-hog.service" stopped (NotFound)
```

Other logs

```
[22:02:16 rd-sideload-compile-job] CC security/apparmor/match.o
[22:02:16 rd-sideload-compile-job] CC crypto/scatterwalk.o
[22:02:16 rd-sideload-compile-job] CC arch/x86/events/intel/bts.o
[22:02:16 rd-sideload-compile-job] CC security/selinux/netlink.o
[22:02:16 rd-sideload-compile-job] CC crypto/proc.o
[22:02:16 rd-sideload-compile-job] AR arch/x86/kernel/fpu/built-in.a
[22:02:16 rd-sideload-compile-job] CC arch/x86/kernel/irq_work.o
[22:02:16 rd-sideload-compile-job] HDRTEST usr/include/drm/msm_drm.h
```

| | cpu% | mem | swap | rbps | wbps | cpuP% | memP% | ioP% |
|--------------|------|-------|-------|------|-------|-------|-------|------|
| workload | 46.9 | 51.0g | 43.5m | 7.6m | - | 4.4 | 0.0 | 0.0 |
| sideload | 26.6 | 4.4g | 13.8m | 558K | 1.3m | 88.7 | 0.0 | 0.0 |
| hostcritical | 0.8 | 604m | - | - | 60.0K | 0.4 | - | - |
| system | 0.0 | 140m | 3.7m | - | - | 0.3 | - | - |
| user | - | - | - | - | - | - | - | - |
| - | 78.7 | 56.6g | 130m | 8.2m | 1.4m | 89.7 | 0.0 | 0.0 |

[intro.post-bench] Introduction to resource control demo - 'i': index, 'b': back

memory hog. The former will eat up as many CPU cycles as it can get its hands on along with some memory and IO bandwidth. The latter will keep gobbling up memory causing memory shortage and subsequent IOs once memory is filled up. The combination is a potent antagonist to our interactive rd-hashd.

[Disable resource control and start the competitions]

See the graph for the steep drop in RPS for hashd: That's the competitions taking away its resources: Not good.

Once workload's memory pressure (memP%) in the top right panel starts spiking, you might not have a lot of time before the whole system starts stalling severely. Let's stop them.

[Stop the compile job and memory hog]

Once RPS climbs back up and the memory usage of workload in the top right panel stops growing, start the same competitions but with resource control enabled and the compile job under the supervision of the sideloader:

[Start the competitions under full resource control]

Watch the stable RPS. rd-hashd is now fully protected against the competitions. The compile job and memory hog are throttled. The compile job doesn't seem to be making much progress. This is because sideloads (workloads under the siderloader supervision) are configured to have lower priority than sysloads (workloads under system). Don't worry about the distinction between sideloads and sysloads for now. We'll revisit them later.

Let's stop the memory hog and see what happens.

[Stop the memory hog]

Main workload is doing fine

fine and the compile job is now making reasonable forward loads are now sharing the machine safely and productively, possible before.

Continue reading to learn more about the various components which make this possible.

[Next: Cgroup and Resource Protection]

Hypervisors support allocation

Intel RDT support in hypervisors

- VMware in Telco Cloud Automation
 - Telco + Finance?
- Static

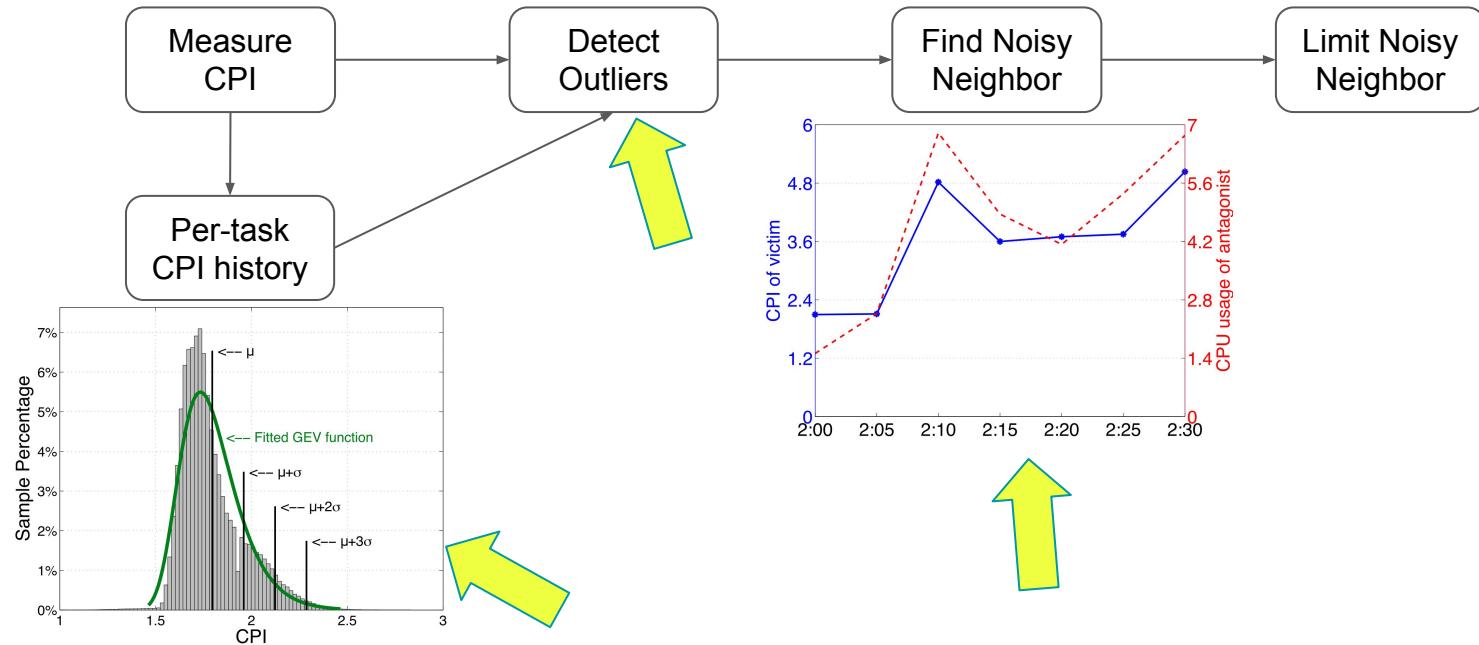
| | Cache partitioning | Memory bandwidth |
|---------|--------------------|-------------------|
| Xen | ✓ | ✓ |
| KVM | ✓ | |
| VMware | ✓ | Monitor (vSphere) |
| Hyper-V | | Monitor PMU |
| ACRN | ✓ | ✓ |

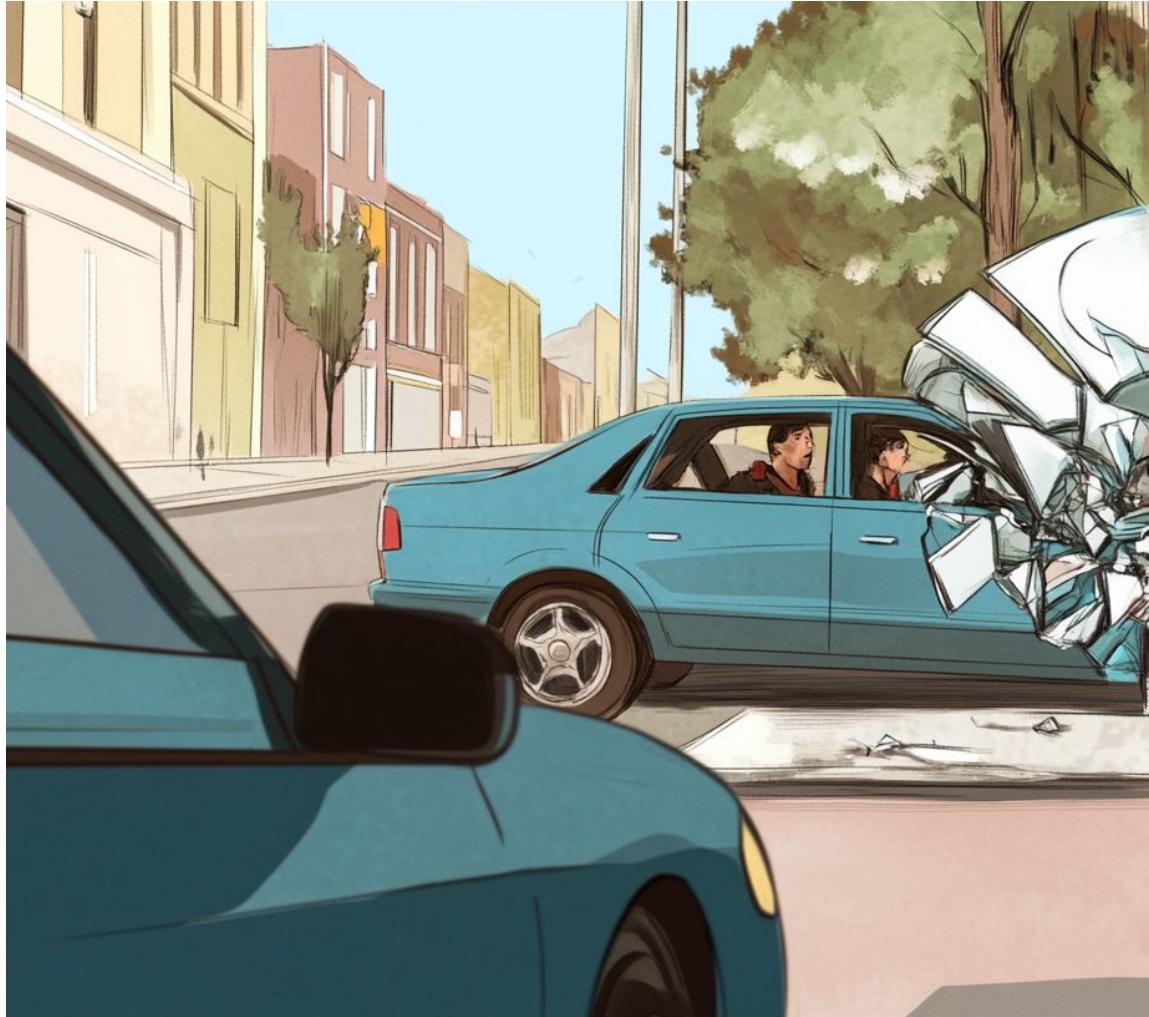
Mitigation Systems

- Type 1: cycles per instruction (CPI)
- Type 2: latency control
- Type 3: usage control

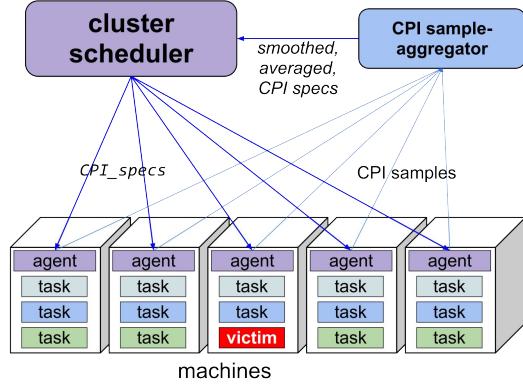
Type 1: Cycles Per Instruction (CPI)

Uses: High interference → high CPI





Type 1: Cycles Per Instruction

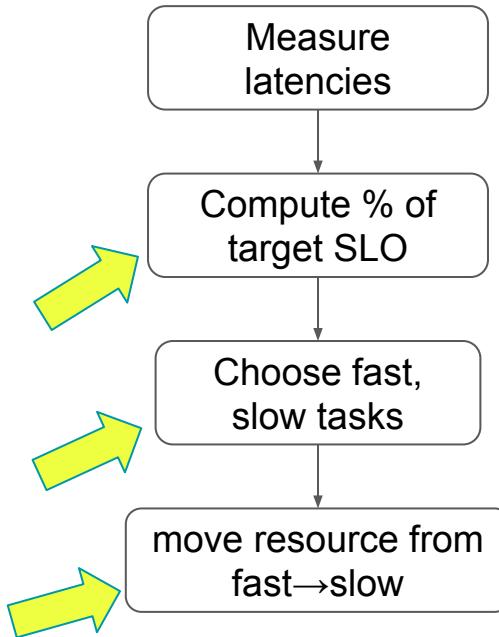


"We have rolled out CPI² to all of Google's shared compute clusters." – paper authors, 2013 @Google

| | Type 1: CPI |
|--------------------|---|
| Measurement | Cycles, Instructions |
| Averaging | High |
| Cluster components | Aggregator |
| Pros | <ul style="list-style-type: none"> Simple to measure |
| Cons | <ul style="list-style-type: none"> Complex deployment Averaging → Slow reaction |

Type 2: Latency Control

Example algorithm:

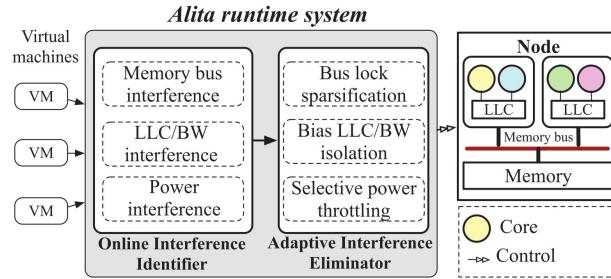


| | Type 2: Latency |
|--------------------|---|
| Measurement | App latency |
| Averaging | Medium |
| Cluster components | Node only |
| Pros | <ul style="list-style-type: none">• No profiling → Node-local• Control what you care about |
| Cons | <ul style="list-style-type: none">• High developer effort• Noisy signal → Averaging |

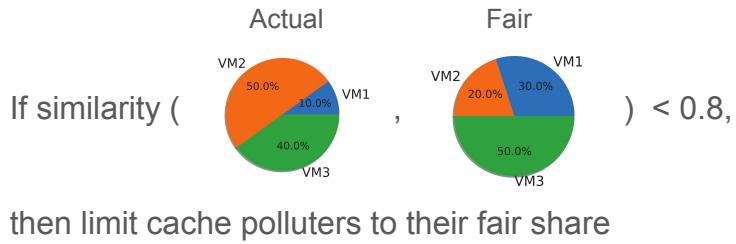
Type 3: Usage Control

How:

- Measure per-app resource usage
- Find unfair allocation
- Limit offender



Cache example:



Do we really want a fair allocation?



Type 3: Usage Control (cont'd)

| | Type 3: Usage |
|--------------------|---|
| Measurement | CPU counters |
| Averaging | Low |
| Cluster components | Node only |
| Pros | <ul style="list-style-type: none"> • Node-local • Measure directly <ul style="list-style-type: none"> → Fast reaction → Easy to reason |
| Cons | <ul style="list-style-type: none"> • Do we really want fair allocation? |

Deployed in production, > 2 years:

- 30k nodes (24 to 48 cores each)
- 250k VMs

- authors, 2020 @Alibaba Cloud

Summary: Mitigation Systems

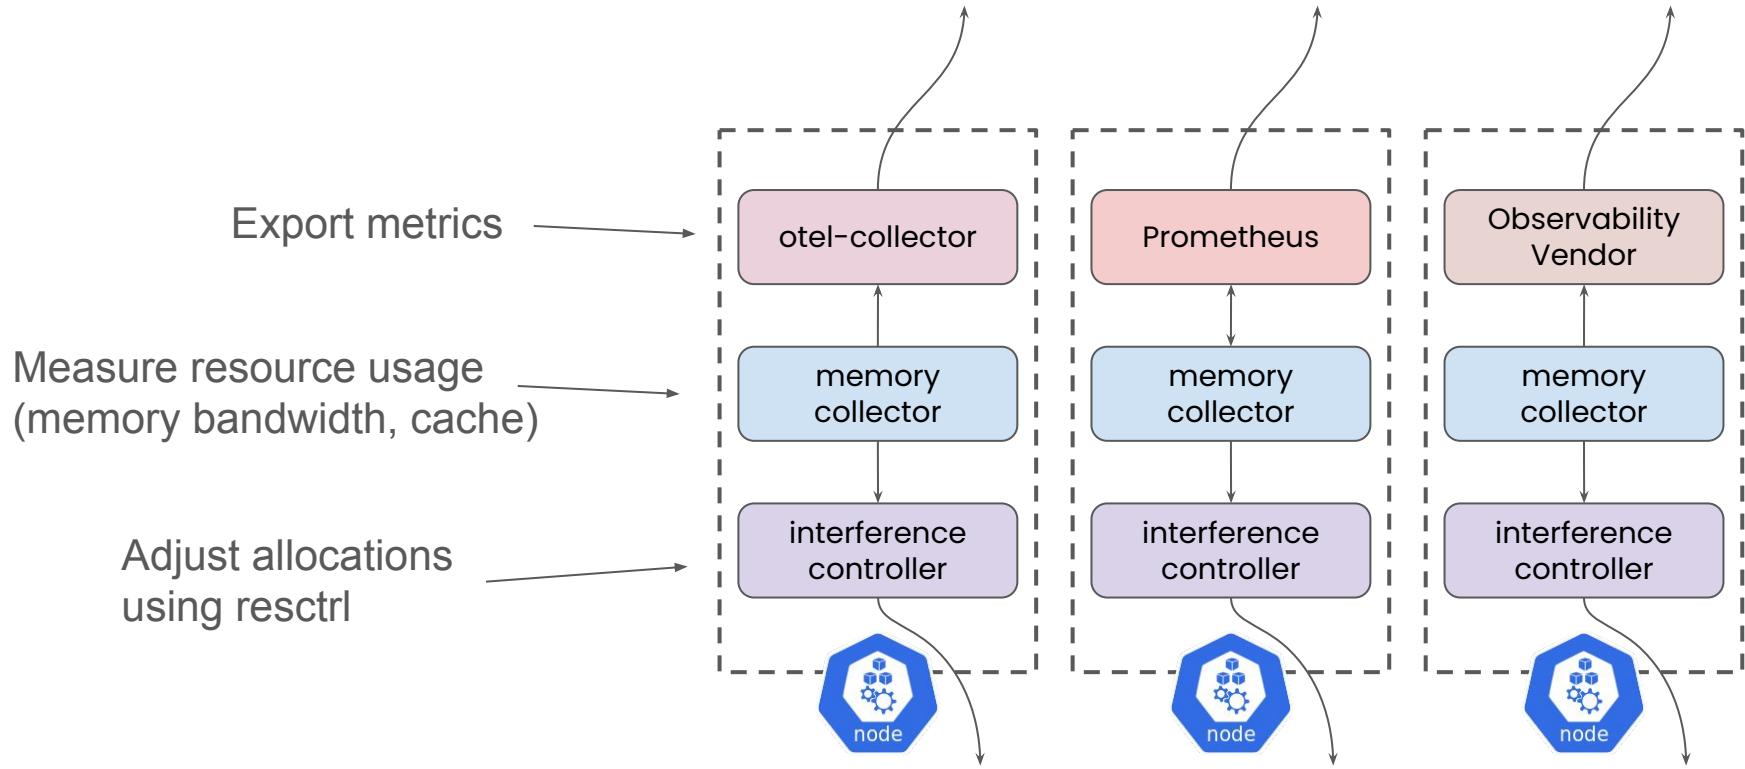
| | Type 1: CPI | Type 2: Latency | Type 3: Usage |
|--------------------|--|---|--|
| Measurement | Cycles, Instructions | App latency | CPU counters |
| Averaging | High | Medium | Low |
| Cluster components | Aggregator | Node only | Node only |
| Pros | <ul style="list-style-type: none">Simple to measure | <ul style="list-style-type: none">No profiling → Node-localControl what you care about | <ul style="list-style-type: none">Node-localMeasure directly → Fast, Easy to reason |
| Cons | <ul style="list-style-type: none">Complex deploymentAveraging → Slow reaction | <ul style="list-style-type: none">High developer effortNoisy signal → Averaging | <ul style="list-style-type: none">Do we really want fair allocation? |

- The three “categories” of published systems
- Many good ideas → general purpose
- Usage control (#3) is promising

Next Steps

- Kubernetes Deployment
- Observability
- Community Next Steps

Kubernetes Deployment

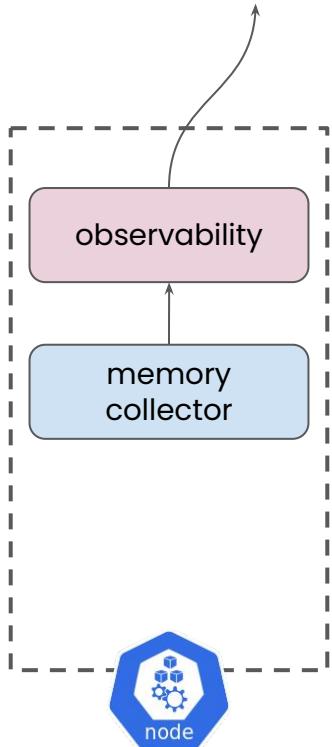


First step

Start with observability, because:

To deploy with confidence, need

- Quantify the benefit
- Metrics to show it is behaving correctly



Community Next Steps



Call for:

- Contributors
 - Collector
 - Kubernetes benchmarks
- Potential users
 - Want to hear more
 - Deploy in test/staging

Contact: yonch@yonch.com

Set up time: yonch.com/collector

