



KubeCon



CloudNativeCon

North America 2024

Powering Automatic Authorization in Envoy Through Live Traffic Inspection

Dom Delnano - CEO/Founder Cosmic and Pixie core maintainer



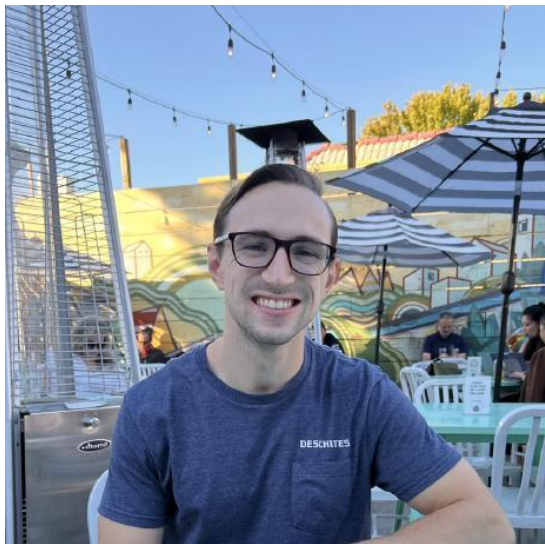
KubeCon



CloudNativeCon

North America 2024

Hi

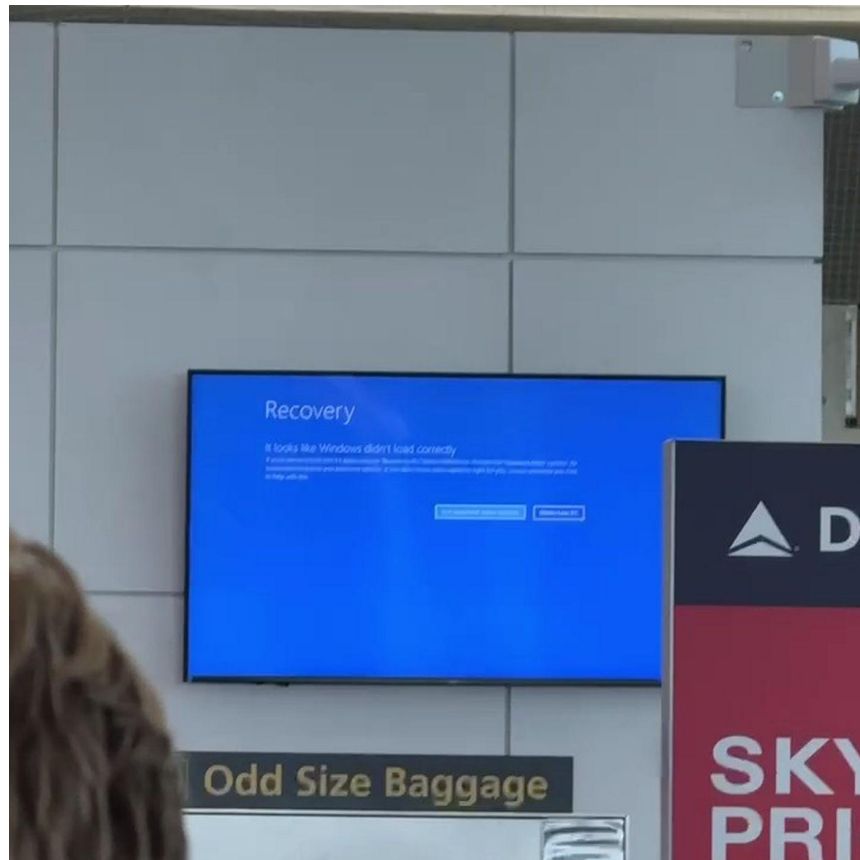


[@ddelnano](https://twitter.com/ddelnano)

Pixie core maintainer
CEO/Founder - Cosmic
Previous: CrowdStrike, New Relic, Twitter



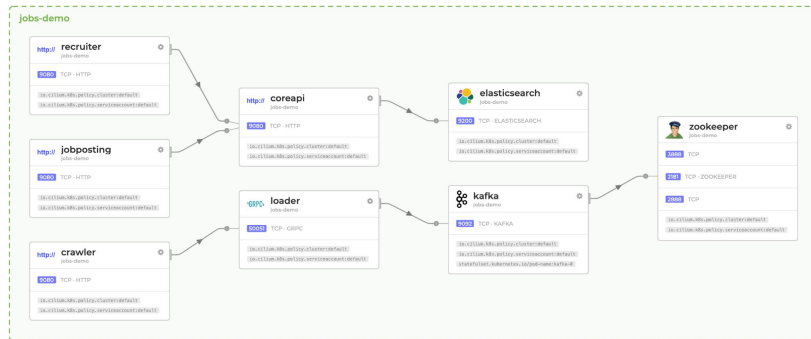
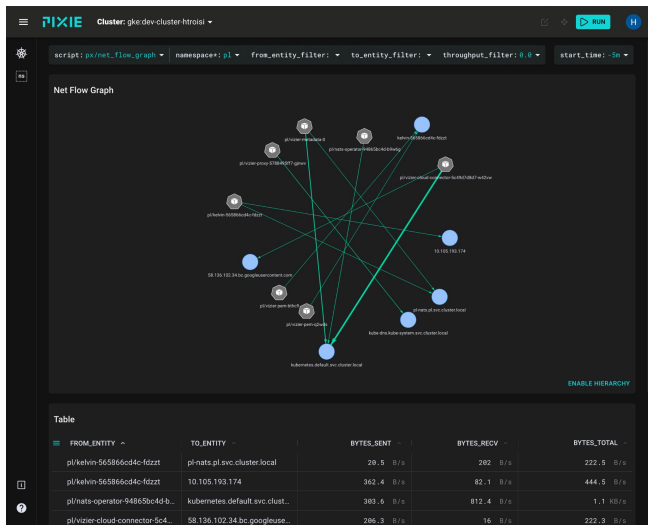
- Security is not a one size fits all solution
- Always consult your security experts



- Security Misconfiguration is #5 in OWSAP's Top 10
- The AuthZ policies of last week, month, year eventually suffer from security/configuration drift
- Imagine a world where AuthZ policies change in response to microservice environment changes

Modern Observability to the Rescue

- Zero-instrumentation Observability tools such as [Pixie](#) and [Hubble](#) provide real time service maps and L7 visibility
- eBPF instruments all network traffic, providing full visibility of an environment
- This unlocks the ability to analyze L7 access patterns and scope access rules accordingly



Goal: Close AuthZ policy gaps as traffic evolves

- Leverage these tools to keep AuthZ policies in sync with existing traffic patterns
 - Provide means for future extension to further restrict L7 access
- This reference architecture will use Pixie and other CNCF components
 - AuthN: mTLS orchestrated by SPIFFE/Spire
 - AuthZ: Envoy with OPA-Envoy
 - Live traffic inspection: Pixie
 - Big data processing: Spark
- This proof of concept lays the foundation for applying Pixie and these principles to your technology stack

Prerequisites

- **Understanding your Authentication**
 - Where in the payload is it? How can the details be accessed?
- **Authorization layer processing:**
 - How should the AuthN be processed/transformed in the AuthZ layer?
 - This part will be auto generated

Implementation

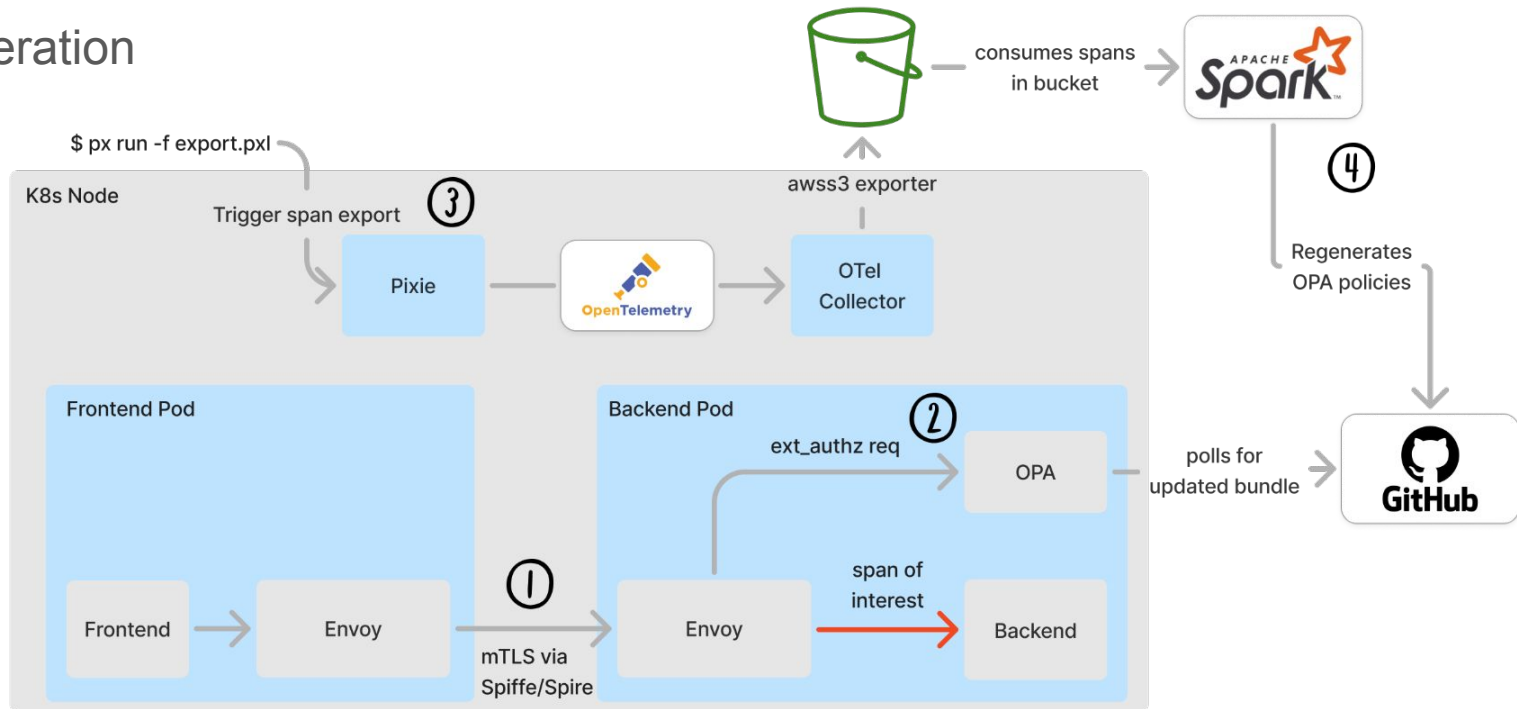
1. **Live traffic inspection and data processing:**
 - Input: Network traffic
 - Output: Protocol spans in [OTLP format](#) annotated w/ AuthN and L7 access
2. **Span data processing and AuthZ rule generation:**
 - Input: Spans in OTLP format stored in an object store
 - Output: OPA policy files that match traffic contained in OTel spans

Automatic Authorization Playbook

Bring Your Own:

1. AuthN
2. AuthZ
4. Policy generation

Required:
3. Pixie



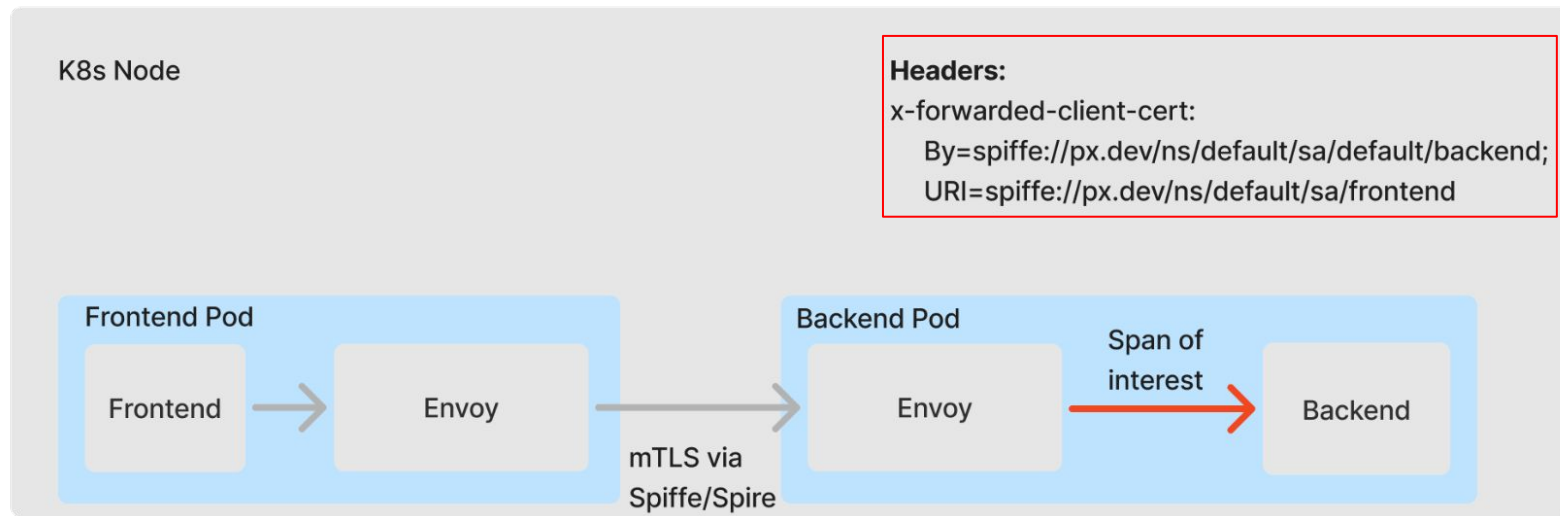
Understanding Your Authentication

- How the AuthN is represented dictates how the AuthZ layer must perform its checks

SPIFFE

- Defines standards for mutual authentication – JWT or X.509 (TLS)
- Easily deployed in cloud native environments w/ Spire – the SPIFFE Runtime Environment
- Spire issues SPIFFE compatible mTLS certs for AuthN and integrates natively with Envoy and other service meshes

Authentication: On the Wire



- Envoy provides a mechanism for performing external AuthZ checks (ext_authz)
 - Authorization service is sent a request to verify if incoming request is authorized
- Open Policy Agent (OPA) is a popular policy engine with a native integration with Envoy (OPA-Envoy)
- OPA policy has access to details from incoming request to make authorization decision
 - AuthN: HTTP headers
 - L7 access: HTTP method, URI

```
package envoy.authz

import rego.v1

import input.attributes.request.http

default allow := false

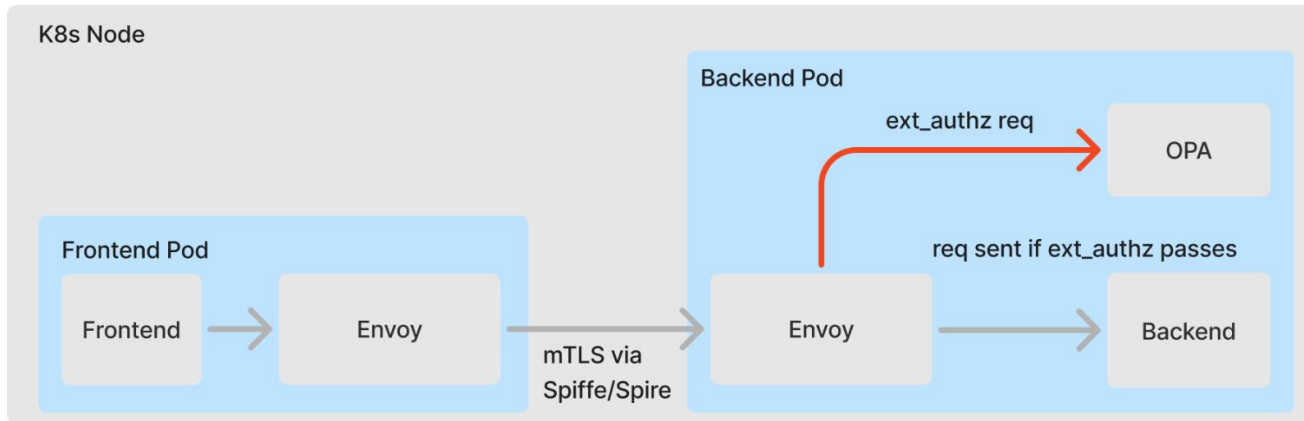
allow if {
    is_token_valid
    action_allowed
}

is_token_valid if {
    token.valid
}

token := {"valid": valid} if {
    [_, _] := split(http.headers.authorization, " ")
    # Use Authorization header to perform the
    # necessary checks
    valid = true
}

action_allowed if {
    http.method == "GET"
    glob.match("/people/*", ["/"], http.path)
}
```

Authorization Architecture



Headers:

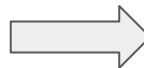
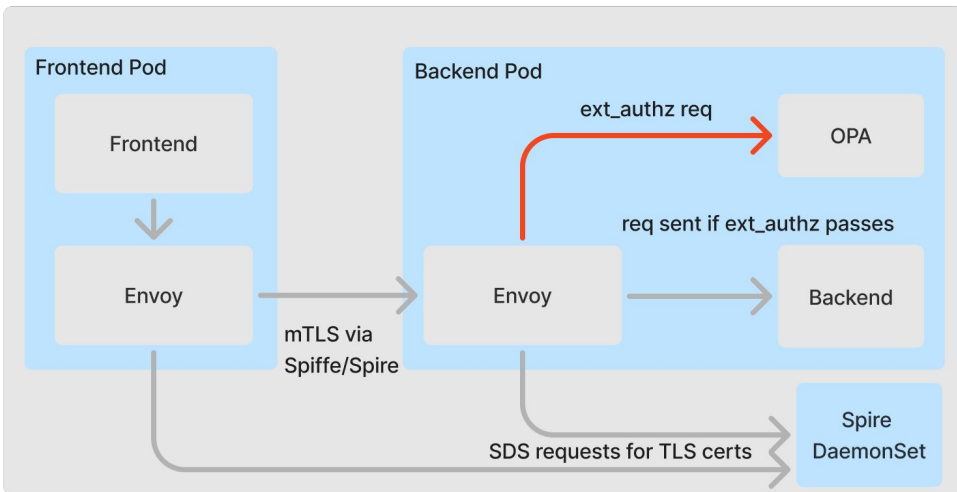
x-forwarded-client-cert:

By=spiffe://px.dev/ns/default/sa/default/backend;

URI=spiffe://px.dev/ns/default/sa/frontend

Authorization: Traffic to OPA Policy

OPA AuthZ policy based on traffic patterns



```
package envoy.authz

import rego.v1

default allow := false

# By=spiffe://.../backend;Hash=<hash>;URI=spiffe://.../frontend
headers := input.attributes.request.http.headers
xfcc := headers["x-forwarded-client-cert"]

# Parse XFCC to pull out src and dest service names
[dest_svc, source_svc] := dest_source_spiffe_svc if {
  # [ ... ]
}

# authz_by_service templated based on client/server details
# in spans and HTTP details (method, URI, etc)
allowed_sources := object.get(authz_by_service, dest_svc, {})
src_svc_http_access := object.get(allowed_sources, src_svc, {})

allow if {
  # Verify src svc is allowed to access dest
  some dest, sources in authz_by_service
  dest == dest_svc

  some src, _ in sources
  src == src_svc
  # Verify L7 access should be permitted
  check_l7_access(src_svc_http_access, http_path, http_method)
}
```



x-forwarded-client-cert:

By=spiffe://px.dev/ns/default/sa/default/backend;
URI=spiffe://px.dev/ns/default/sa/frontend

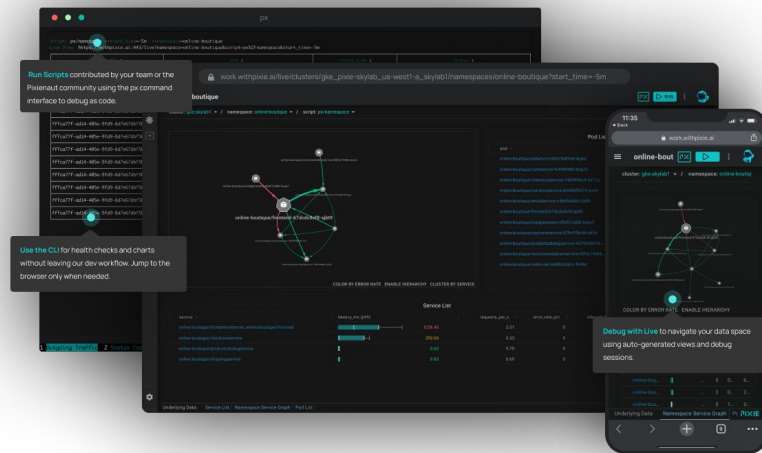
Introduction to Pixie

Goal: Performance debugging without manual instrumentation



- CPU/Memory/Network
- Protocol tracing (spans) w/ latency and payload
- Performance profiles (flamegraphs)

Characteristics

- **Zero-instrumentation:** eBPF captures all data w/o application changes
- **Distributed architecture** with in-memory data store
- **Scriptable Python/Pandas interface** to query and process telemetry



Pixie is Scriptable

- Valid  python™
- Valid  pandas
- Built for data analysis and ML
- UI widgets and data processing completely configurable
- Integrates in the o11y ecosystem
 - OTEL export
 - Grafana data source

```
1 import px
2
3 def http_data():
4     df = px.DataFrame(table='http_events', start_time='-30s')
5
6     df.pod = df.ctx['pod'] # Annotate dataframe with k8s pod
7
8     df.xfcc = px.pluck(df.req_headers, 'x-forwarded-client-cert')
9
10    return df[['pod', 'req_path', 'resp_latency_ns', 'xfcc']]
11
12 px.display(http_data())
```

Live Traffic Inspection: Pixie



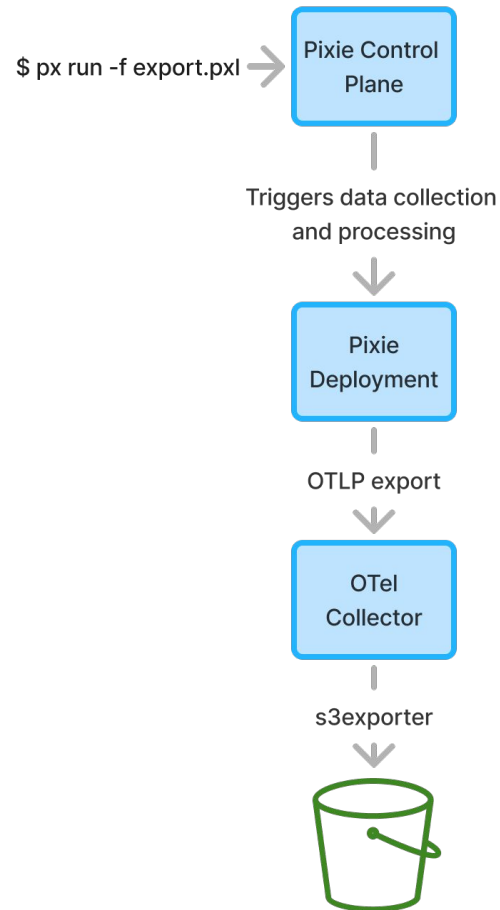
KubeCon



CloudNativeCon

North America 2024

- Pixie's [PxL](#) language provides a means for enriching protocol spans
 - Src and dest k8s service/deployment/pod
 - L7 access (HTTP method, URL)
 - AuthN details – transformation of XFCC header
- Pairing its OTel export with the OTel collector makes it easy to send spans to object storage for big data processing



Span data processing and Policy Generation



OTel collector S3 export path

s3:///.../year=YYYY/month=MM/day=DD/hour=hh/min=mm/*.json

```
1 # Entry for each span. OTLP format highly nested
2 [
3   [
4     {
5       "http.xfcc_by": "By=spiffe://example.org/ns/default/sa/default/backend",
6       "http.xfcc_uri": "URI=spiffe://example.org/ns/default/sa/default/frontend",
7       "http.method": "GET",
8       "http.url": "/balances/balance_1"
9     },
10    {
11      "http.xfcc_by": "By=spiffe://example.org/ns/default/sa/default/backend",
12      "http.xfcc_uri": "URI=spiffe://example.org/ns/default/sa/default/frontend",
13      "http.method": "GET",
14      "http.url": "/transactions/transaction_2"
15    },
16    [ ... ]
17  ]
18 ]
```

Traces of allowed traffic from Pixie



Aggregate and
output OPA
policy files

```
1 $ tree opa
2 opa
3 └── imports
4     ├── backend_a.rego
5     ├── backend_b.rego
6     ├── backend_c.rego
7     └── backend_d.rego
8 └── opa-policy.rego # Main file that uses all imports
9
10 2 directories, 7 files
11
```

OPA Policies that reflect current
live traffic

Span data processing and Policy Generation



KubeCon



CloudNativeCon

North America 2024

```
1 $ tree opa
2 opa
3 |-- imports
4 |   |-- backend_a.rego
5 |   |-- backend_b.rego
6 |   |-- backend_c.rego
7 |   |-- backend_d.rego
8 |-- opa-policy.rego # Main file that uses all imports
9
10 2 directories, 7 files
11
```

```
1 # imports/backend_a.rego
2 package px.service.backend_a
3
4 valid_sources := {
5   "frontend": [
6     {
7       "method": "GET",
8       "prefix": "/profiles/profile_1",
9     },
10    [ ... ]
11  ]
12 }
```

```
1 # opa-policy.rego
2
3 # AUTOGENERATED: imports are added for each destination service
4 import data.px.service.backend_a
5 import data.px.service.backend_b
6 [ ... ]
7
8 authz_by_service := {
9   "backend_a": data.px.service.backend_a.valid_sources,
10  "backend_b": data.px.service.backend_b.valid_sources,
11  [ ... ]
12 }
13 # END AUTOGENERATED
```

OPA Policies that reflect current
live traffic

work.testing.getcosmic.ai/live/clusters/ddelnano-test-cluster?destination_filter=default%2F&max_num_records=1000&script=px%2Fhttp_data&sou...

OSMIC cluster: ddelnano-test-cluster

script: px/http_data source_filter: destination_filter: default/ max_num_records: 1000 start_time: -5m

Table Showing 1 - 15 out of 1000 records

TIME	SOUR...	DESTINATION	LATE...	MAJ...	REQ_PATH	REQ...	REQ...	REQ...
11/14/2...	127.0.0.1	default/frontend-2-c54c84d5c-h6bm5	88.8 ms	1	/	GET	{ Accep...	
11/14/2...	127.0.0.1	default/frontend-2-c54c84d5c-h6bm5	27.8 ms	1	/transactions/transaction_2	GET	{ Accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	180.6 μs	1	/transactions/transaction_2	GET	{ accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	8.5 ms	2	/envoy.service.auth.v3.Authorization/Check	POST	{ :auth...	1: "\nM\...
11/14/2...	127.0.0.1	default/frontend-2-c54c84d5c-h6bm5	16 ms	1	/profiles/profile_2	GET	{ Accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	252.5 μs	1	/profiles/profile_2	GET	{ accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	2.4 ms	2	/envoy.service.auth.v3.Authorization/Check	POST	{ :auth...	1: "\nM\...
11/14/2...	127.0.0.1	default/frontend-2-c54c84d5c-h6bm5	9.9 ms	1	/balances/balance_2	GET	{ Accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	178.5 μs	1	/balances/balance_2	GET	{ accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	2.1 ms	2	/envoy.service.auth.v3.Authorization/Check	POST	{ :auth...	1: "\nM\...
11/14/2...	127.0.0.1	default/frontend-5884c9c8bf-hrxqc	73 ms	1	/	GET	{ Accep...	
11/14/2...	127.0.0.1	default/frontend-5884c9c8bf-hrxqc	15.1 ms	1	/transactions/transaction_1	GET	{ Accep...	
11/14/2...	127.0.0.1	default/backend-57f6f6864c-5552v	168.5 μs	1	/transactions/transaction_1	GET	{ accep...	

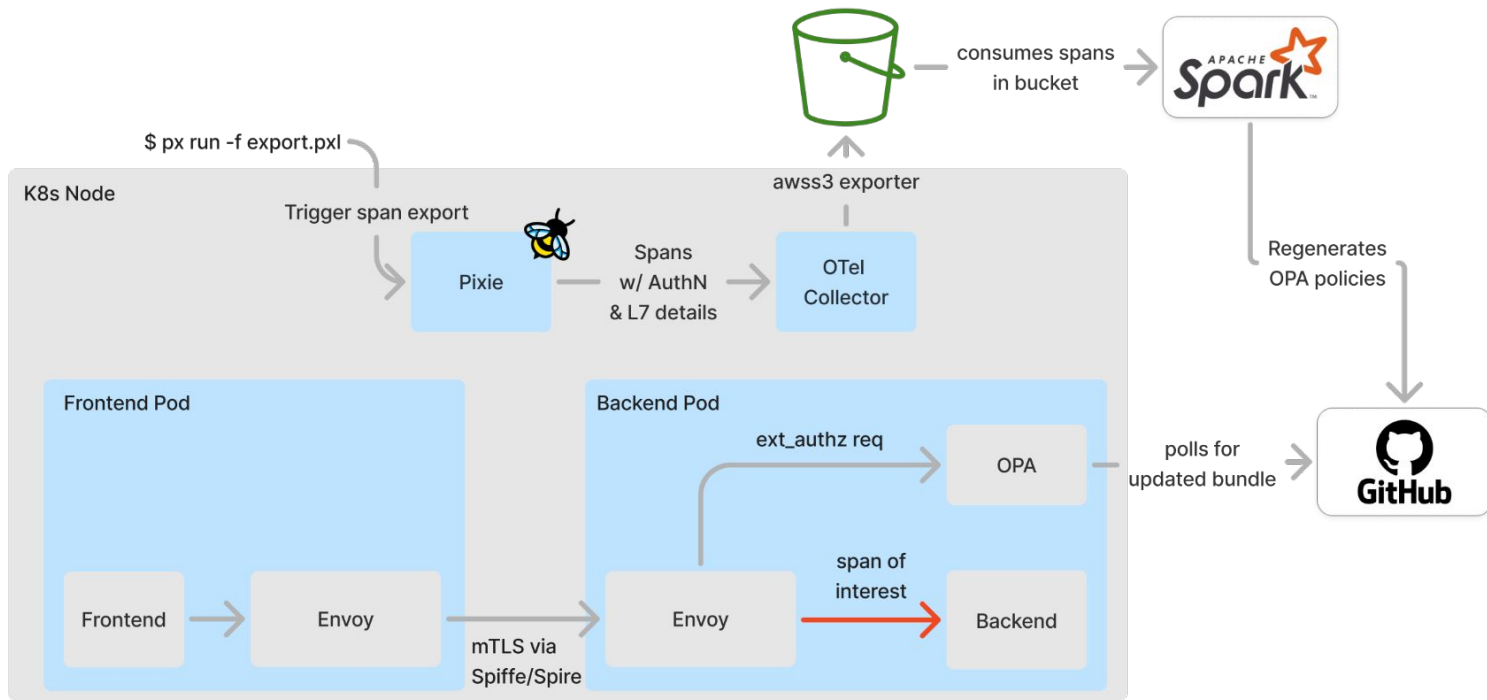


KubeCon



CloudNativeCon

North America 2024



- Integrate rule generation logic with richer service route information
 - Many RPC frameworks, API tooling and generated client libraries provide routing information
 - GRPC
 - OpenAPI (Swagger)
- Using this in the OPA policy file generation can produce least privileged access

- This lays the foundation for continuous enforcement of current traffic patterns
- Eliminate security/configuration drift for AuthZ rules
- Pixie's data processing is key for powering use cases like this
 - Bring Your Own existing AuthN, AuthZ and architecture



KubeCon



CloudNativeCon

North America 2024

Thank you!



- Demo - [pixie-io/pixie-demos#54](https://github.com/pixie-io/pixie-demos#54)