



**KubeCon**



**CloudNativeCon**

**North America 2024**





KubeCon



CloudNativeCon

North America 2024

# Pushing Authorization Further

## CEL, Selectors (and maybe RBAC++)

Anish Ramasekar & Rita Zhang, Microsoft  
Jordan Liggitt, Google



KubeCon



CloudNativeCon

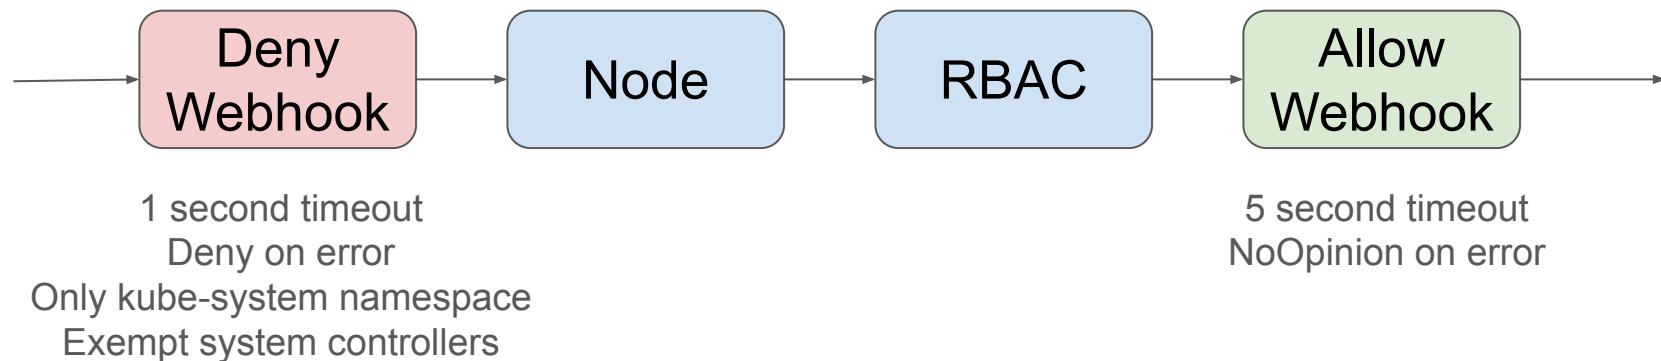
North America 2024

# Graduated Features in 2024

## Structured Authorization Config

[kep.k8s.io/3221](https://kep.k8s.io/3221), beta in v1.30, stable in v1.32

@liggitt, @palnabarun, @ritazh



[2023 Kubecon demo](#)

## Structured Authentication Config

[kep.k8s.io/3331](https://kep.k8s.io/3331), beta in v1.30

@aramase, @enj

### JWT claims:

```
{
```

```
  "username": "jane_doe",
```

```
  "sub": "119abc",
```

```
  "roles": "admin,user",
```

```
  "aud": "kubernetes",
```

```
  ...
```

```
}
```

### Resulting user:

```
username: "jane_doe:external"
```

```
uid: "119abc"
```

```
groups: ["admin", "user"]
```

```
extra:
```

```
  other_attributes: ["example"]
```

```
  client_name: ["kubernetes"]
```

## Label / Field Selector Authorization

[kep.k8s.io/4601](https://kep.k8s.io/4601), beta in v1.32

@deads2k, @liggitt

- Authorizers see label / field selectors for list / watch requests
- Node authorizer limits nodes to listing/watching their own pods
- Webhook authorizers can limit list and watch requests based on the label or field selector used. Examples:
  - “**node1** can only list pods with **.spec.nodeName=node1**”
  - “**myingress** can watch secrets with **controller=myingress**”

## Node Info in Service Account Tokens

[kep.k8s.io/4193](https://kubernetes.io/blog/2023/09/12/service-account-token-claims/), beta in v1.30, stable in v1.32

@enj, @munnerz

- Node name included in service account token claims  
`"node":{"name":"my-node","uid":"..."}`
- Node name surfaces as a user attribute  
`"authentication.kubernetes.io/node-name":["my-node"]`
- Usable by authorization / admission (ValidatingAdmissionPolicy)
- Foundation for future work to keep service account credentials being an escalation path for nodes

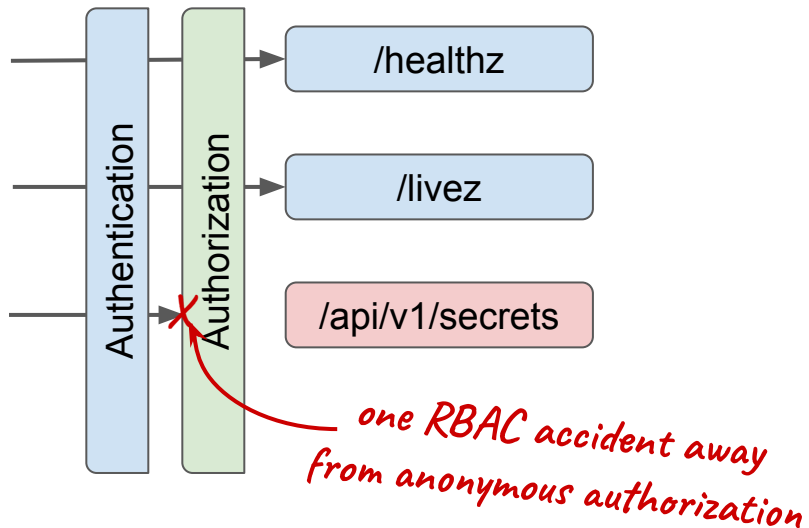
[2024 Kubecon demo](#)

## Restricted Anonymous Authentication

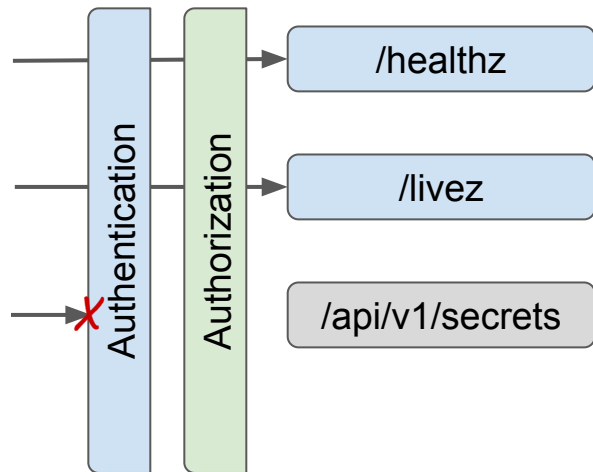
[kep.k8s.io/4633](https://kep.k8s.io/4633), beta in v1.32

@vinayakankugoyal

Before:



After:







KubeCon



CloudNativeCon

North America 2024

# Upcoming Features

## Cluster Trust Bundle

[kep.k8s.io/3257](https://kep.k8s.io/3257), targeting beta in 1.33, @ahmedtd

Publish CA root bundles in API objects, mounted into pods

## Pod Certificates

[kep.k8s.io/4317](https://kep.k8s.io/4317), targeting alpha in 1.33, @ahmedtd

x509 client certificates for service account identities, mounted into pods

## External service account token signing

[kep.k8s.io/740](https://kep.k8s.io/740), alpha in 1.32, @ahmedtd, @HarshalNeelkamal, @liggitt

Token signing: in-process with local private key file → gRPC call

Token verification keys: local public key files → gRPC call

Removes key signing material from the kube-apiserver

Allows dynamic updates of public verifying keys

Can integrate with cloud KMS signers, external audit of token signatures

## Fine-grained Kubelet Authorization

[kep.k8s.io/2862](https://kep.k8s.io/2862), alpha in 1.32, @vinayakankugoyal

Authorize *just* the `/configz`, `/healthz`, or `/pods` kubelet API endpoints

Avoid over-granting permissions to local monitoring processes

## Deleting undecryptable / corrupt resources

[kep.k8s.io/3926](https://kep.k8s.io/3926), alpha in 1.32, @stlaz, @tkashem

The unexpected (both human and cosmic) does happen

Can prevent decoding of an object written to etcd

Failure was *really* bad (unlistable resource type, undeletable object)

Recovery was *really* bad (bypass API and delete directly from etcd)

Allows force delete (with sufficient permissions) of corrupt objects

## Image pull credential verification

[kep.k8s.io/2535](https://kep.k8s.io/2535), targeting alpha in 1.33, @stlaz

Automatically prevent unauthorized use of images already on the node

## Service account image pull credentials

[kep.k8s.io/4412](https://kep.k8s.io/4412), targeting alpha in 1.33, @aramase, @enj

Opt into using service account tokens as image pull credentials

## Node-restricted service accounts

[kep.k8s.io/4935](https://kep.k8s.io/4935), in design, @vinayakankugoyal

Goal: Constrain a service account so it can't be a node escalation path

## Kubelet Serving Certificate Validation

[kep.k8s.io/4872](https://kep.k8s.io/4872), in design, @g-gaston

Additional TLS checks when talking to nodes

Verify node identity in addition to standard IP/DNS validation



## Secrets Store Sync Controller

[sigs.k8s.io/secrets-store-sync-controller](https://sigs.k8s.io/secrets-store-sync-controller), in alpha, @nilekhc

SIG Auth subproject

Controller to sync from external secrets store to Kubernetes secrets



KubeCon



CloudNativeCon

North America 2024

# User Stories for Authorization

- A node can only list pods scheduled to it
  - (already done by node authorizer; this is about being able to express similar policy)
- A node agent can only list a custom resource matching the node it is running on
- A node agent can only get a custom resource with a name matching the node it is running on
- A controller can list all secrets of a particular type
- A controller can list all secrets with a particular label
- Allow access based on extra attributes of user.Info
- Allow access based on namespace or name prefix



KubeCon



CloudNativeCon

North America 2024

these are already achievable  
with webhook authorization



KubeCon



CloudNativeCon

North America 2024

but... declarative?



KubeCon



CloudNativeCon

North America 2024

# RBAC++

# Questions?



<https://git.k8s.io/community/sig-auth>

Bi-weekly meetings (Wednesday at 11am Pacific Time)

Mailing list: [kubernetes-sig-auth@googlegroups.com](mailto:kubernetes-sig-auth@googlegroups.com)

Kubernetes Slack: [#sig-auth](#)