



KubeCon



CloudNativeCon

North America 2024





KubeCon



CloudNativeCon

North America 2024

Rogue No More: Securing Kubernetes with Node-Specific Restrictions

James Munnelly



@munnerz

Anish Ramasekar



@aramase



KubeCon



CloudNativeCon

North America 2024

Background

- Nodes have a set of 'required' permissions
- Kubelet, DaemonSet, possibly others
- What prevents one instance modifying/posting another's information?



KubeCon



CloudNativeCon

North America 2024

kubelet

Limiting kubelet permissions

- kubelets need to manage/mutate:
 - Pods
 - Nodes
 - PersistentVolumeClaims
 - ServiceAccounts
 - Leases
 - CSINodes
 - CertificateSigningRequests

- NodeRestriction admission plugin:
 - Restricts nodes to only be able to mutate objects 'associated' with themselves
 - Inspects authenticated kubelet's `username` and `groups` in the request
 - Matches against existing object or request body

In order to be authorized by the Node authorizer, kubelets must use a credential that identifies them as being in the `system:nodes` group, with a username of `system:node:<nodeName>`. This group and user name format match the identity created for each kubelet as part of [kubelet TLS bootstrapping](#).

<https://kubernetes.io/docs/reference/access-authn-authz/node/#overview>



KubeCon



CloudNativeCon

North America 2024

But what about DaemonSets?

Limiting DaemonSet permissions



- DaemonSet runs across all Nodes in the cluster
- Lack of fine-grained permissions means any instance can modify any object

This attack-vector could allow lateral movement between nodes

CVE-2023-26484

On a compromised node, the virt-handler service account can be used to modify all node specs

Moderate rmohr published GHSA-cp96-jpmq-xrr2 on Mar 15, 2023

Package No package listed	Affected versions all versions	Patched versions None	Severity Moderate
Description Impact <p>If a malicious user has taken over a Kubernetes node where virt-handler (the KubeVirt node-daemon) is running, the virt-handler service account can be used to modify all node specs.</p> <p>This can be misused to lure-in system-level-privileged components (which can for instance read all secrets on the cluster, or can exec into pods on other nodes). This way a compromised node can be used to elevate privileges beyond the node until potentially having full privileged access to the whole cluster.</p> <p>The simplest way to exploit this, once a user could compromise a specific node, is to set with the virt-handler service account all other nodes to unschedulable and simply wait until system-critical components with high privileges appear on its node.</p> <p>Since this requires a node to be compromised first, the severity of this finding is considered Medium.</p>			CVE ID CVE-2023-26484
			Weaknesses No CWEs
			Credits  younaman  XDTG

<https://github.com/kubevirt/kubevirt/security/advisories/GHSA-cp96-jpmq-xrr2>

CVE-2023-30840

On a compromised node, the fluid-csi service account can be used to modify node specs

Moderate cheyang published GHSA-93xx-cvmc-9w3v on May 8, 2023

Package	Affected versions	Patched versions
charts/fluid/templates/role/csi/rbac.yaml (Kubernetes)	< 0.8.6, >= 0.7.0	0.8.6

Description

Impact

If a malicious user gains control of a Kubernetes node running fluid csi pod (controlled by the `csi-nodeplugin-fluid node-daemonset`), he/she can leverage the fluid-csi service account to modify specs of all the nodes in the cluster. However, since this service account lacks "list node" permissions, the attacker may need to use other techniques to identify vulnerable nodes.

Once the attacker identifies and modifies the node specs, he/she can manipulate system-level-privileged components to access all secrets in the cluster or execute pods on other nodes. This allows he/she to elevate privileges beyond the compromised node and potentially gain full privileged access to the whole cluster.

To exploit this vulnerability, the attacker can make all other nodes unschedulable (for example, patch node with taints) and wait for system-critical components with high privilege to appear on the compromised node. However, this attack requires two prerequisites: a compromised node and identifying all vulnerable nodes through other means. Additionally, since the attack is passive and requires patience and luck, the severity of this finding is considered medium.

Severity

Moderate 4.0 / 10

CVSS v3 base metrics

Attack vector	Local
Attack complexity	High
Privileges required	High
User interaction	Required
Scope	Unchanged
Confidentiality	High
Integrity	None
Availability	None

[Learn more about base metrics](#)

CVSS:3.1/AV:L/AC:H/PR:H/UI:R/S:U/C:H/I:N/A:N

CVE ID

CVE-2023-30840



KubeCon



CloudNativeCon

North America 2024

KEP-4193: bound service account token improvements

(we are not very imaginative with naming...)

- **Embedding the name of the Node a Pod is bound to in Pod bound tokens**
- Permit binding directly to Node objects
- Including the JTI field for precise token tracking in audit logs

Changes to service account tokens

```
{
  "aud": [
    "https://kubernetes.default.svc.cluster.local"
  ],
  "exp": 1730388792,
  "iat": 1730385192,
  "iss": "https://kubernetes.default.svc.cluster.local",
  "jti": "0eeb7c44-eba4-4d9c-8b53-51d9c952e28e",
  "kubernetes.io": {
    "namespace": "kube-system",
    "pod": {
      "name": "coredns-55cb58b774-k4ln6",
      "uid": "604fc11d-baf4-4075-9b04-2ac502c15289"
    },
    "serviceaccount": {
      "name": "coredns",
      "uid": "50dd6856-2da2-11e9-9cd9-2afc33b31a7e"
    }
  },
  "nbf": 1730385192,
  "sub": "system:serviceaccount:kube-system:coredns"
}
```



```
{
  "aud": [
    "https://kubernetes.default.svc.cluster.local"
  ],
  "exp": 1730388792,
  "iat": 1730385192,
  "iss": "https://kubernetes.default.svc.cluster.local",
  "jti": "0eeb7c44-eba4-4d9c-8b53-51d9c952e28e",
  "kubernetes.io": {
    "namespace": "kube-system",
    "node": {
      "name": "node1.lab.dev",
      "uid": "646e7c5e-32d6-4d42-9dbd-e504e6cbe6b1"
    },
    "pod": {
      "name": "coredns-55cb58b774-k4ln6",
      "uid": "604fc11d-baf4-4075-9b04-2ac502c15289"
    },
    "serviceaccount": {
      "name": "coredns",
      "uid": "50dd6856-2da2-11e9-9cd9-2afc33b31a7e"
    }
  },
  "nbf": 1730385192,
  "sub": "system:serviceaccount:kube-system:coredns"
}
```

- This allows admission plugins to utilize embedded Node information in Pod-bound tokens when admitting any kind of object!

<http://kep.k8s.io/4193>




KubeCon



CloudNativeCon

North America 2024

Combining with `ValidatingAdmissionPolicy`

 **FEATURE STATE:** Kubernetes v1.30 [stable]

What is Validating Admission Policy?

Validating admission policies offer a declarative, in-process alternative to validating admission webhooks.

Validating admission policies use the Common Expression Language (CEL) to declare the validation rules of a policy. Validation admission policies are highly configurable, enabling policy authors to define policies that can be parameterized and scoped to resources as needed by cluster administrators.

Validating Admission Policy & CEL



KubeCon



CloudNativeCon

North America 2024

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingAdmissionPolicy
metadata:
  name: "only-allow-name-matching-node"
spec:
  failurePolicy: Fail
  matchConstraints:
    resourceRules:
      - apiGroups: [""]
        apiVersions: ["v1"]
        operations: ["UPDATE"]
        resources: ["nodes"]
  matchConditions:
    - name: isRestrictedUser
      expression: >-
        request.userInfo.username == "system:serviceaccount:default:node-patcher-sa"
  variables:
    - name: userNodeName
      expression: >-
        request.userInfo.extra['?authentication.kubernetes.io/node-name'][0].orValue('')
    - name: objectNodeName
      expression: >-
        object.?metadata.name.orValue('')
  validations:
    - expression: variables.userNodeName != ""
      message: >-
        no node association found for user, this user must run in a pod on a node and ServiceAccountTokenPodNodeInfo must be enabled
    - expression: variables.userNodeName == variables.objectNodeName
      messageExpression: >-
        "this user running on node '"+variables.userNodeName+"' may not modify Node '"+variables.objectNodeName +
        "' because the name does not match the node name"
```

```
/ # kubectl auth whoami
ATTRIBUTE                                     VALUE
Username                                     system:serviceaccount:default:node-patcher-sa
UID                                           9b65e954-a358-4894-9bc2-a195937d89b2
Groups                                       [system:serviceaccounts system:serviceaccounts:default system:authenticated]
Extra: authentication.kubernetes.io/credential-id [JTI=1a58c7f2-5b07-449f-bd00-da039856e49f]
Extra: authentication.kubernetes.io/node-name   [kind-worker2]
Extra: authentication.kubernetes.io/node-uid    [78267bd9-a6a7-4052-ab37-cc4bbc27bc78]
Extra: authentication.kubernetes.io/pod-name    [node-patcher-daemonset-4mktr]
Extra: authentication.kubernetes.io/pod-uid     [360ba66a-9578-4deb-a3a1-55697988998e]
```



KubeCon



CloudNativeCon

North America 2024

Demo

- Scoped authentication to image registries
 - [Service account token restriction](#)
 - [Projected Service Account Tokens for Kubelet Credential Providers KEP](#)
- Check out [Squashing Trampoline Pods: The Future of Securely Enabling Hardware Extensions by Joe Betz and David Eads](#)

Give us feedback



KubeCon



CloudNativeCon

North America 2024

