



KubeCon



CloudNativeCon

North America 2024





KubeCon



CloudNativeCon

North America 2024

Share the ride: Robust Multi-Tenancy in Kubernetes at Uber

Apoorva Jindal
Sashank Appireddy

Agenda

- Introduction
- Multi-Tenancy Requirements
- Solution #1: Separate Cluster per Tenant
- Challenges
- Multi Tenant Single Cluster Architecture
- Migration Status
- What works
- Acknowledgements
- Q & A

Platform Overview

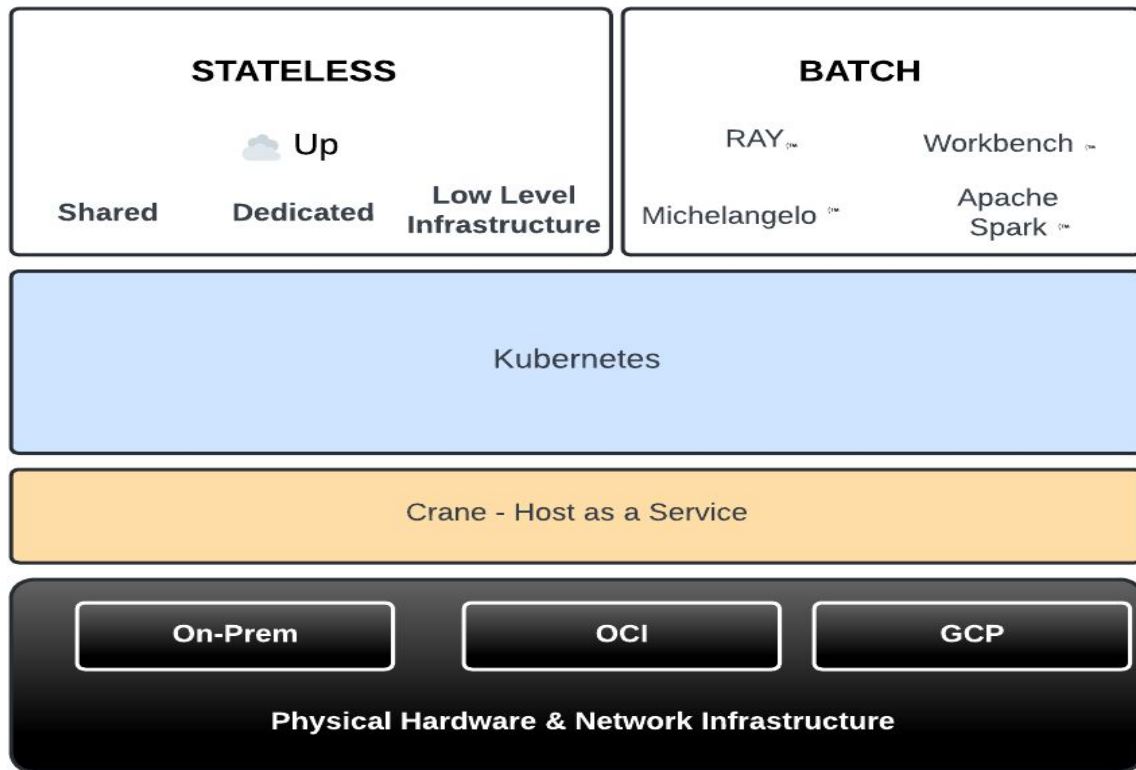


KubeCon




CloudNativeCon

North America 2024



Platform Overview



4000+ microservices

4.5M+ cores

100K+ service deploys per day

1.5M+ containers deployed per day

500K+ containers



KubeCon



CloudNativeCon


North America 2024

Multi-Tenancy Requirements

Requirement #1: Data-Plane Isolation


- Tenants do not share hosts
- Workloads belonging to the same tenant can share the host, but not across tenants

Tenant #1



Node Pool

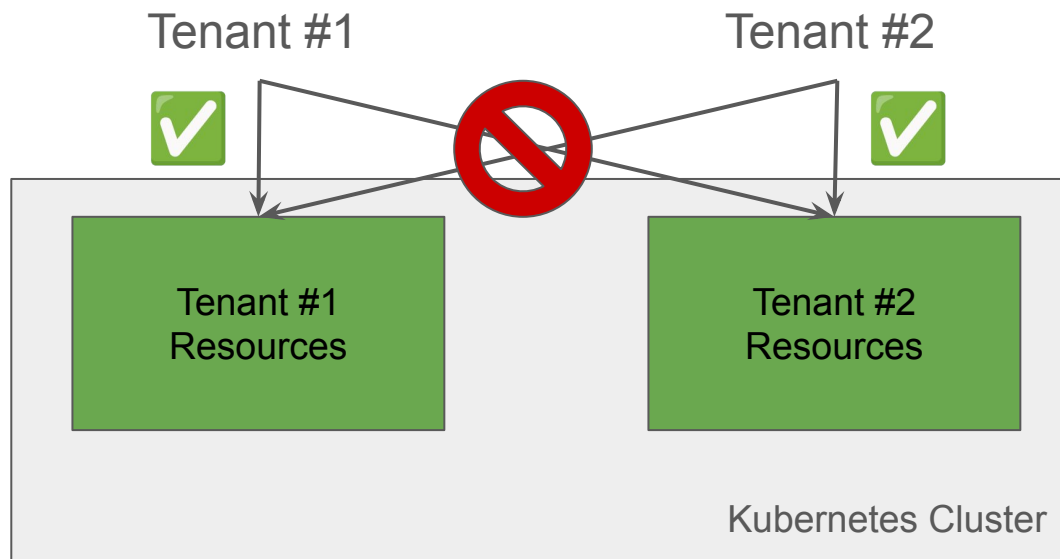
Tenant #2



Node Pool

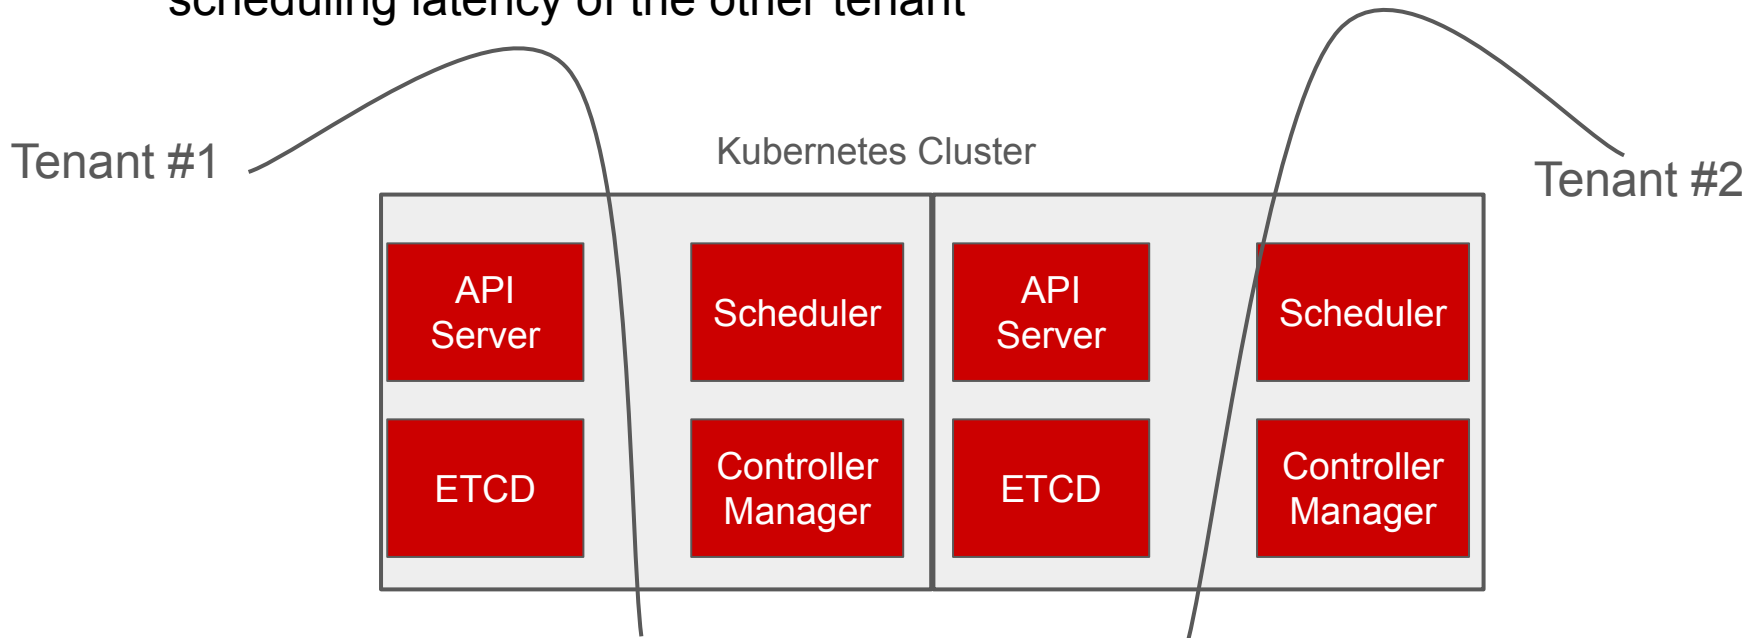
Requirement #2: Access Isolation

- Tenants cannot access resource information about other tenants
 - E.g. pods, nodes, deployments etc



Requirement #3: Control-Plane Isolation

- Tenant is not impacted by other tenants in the control plane
- E.g. a high scheduling throughput of one tenant does not impact low scheduling latency of the other tenant





KubeCon



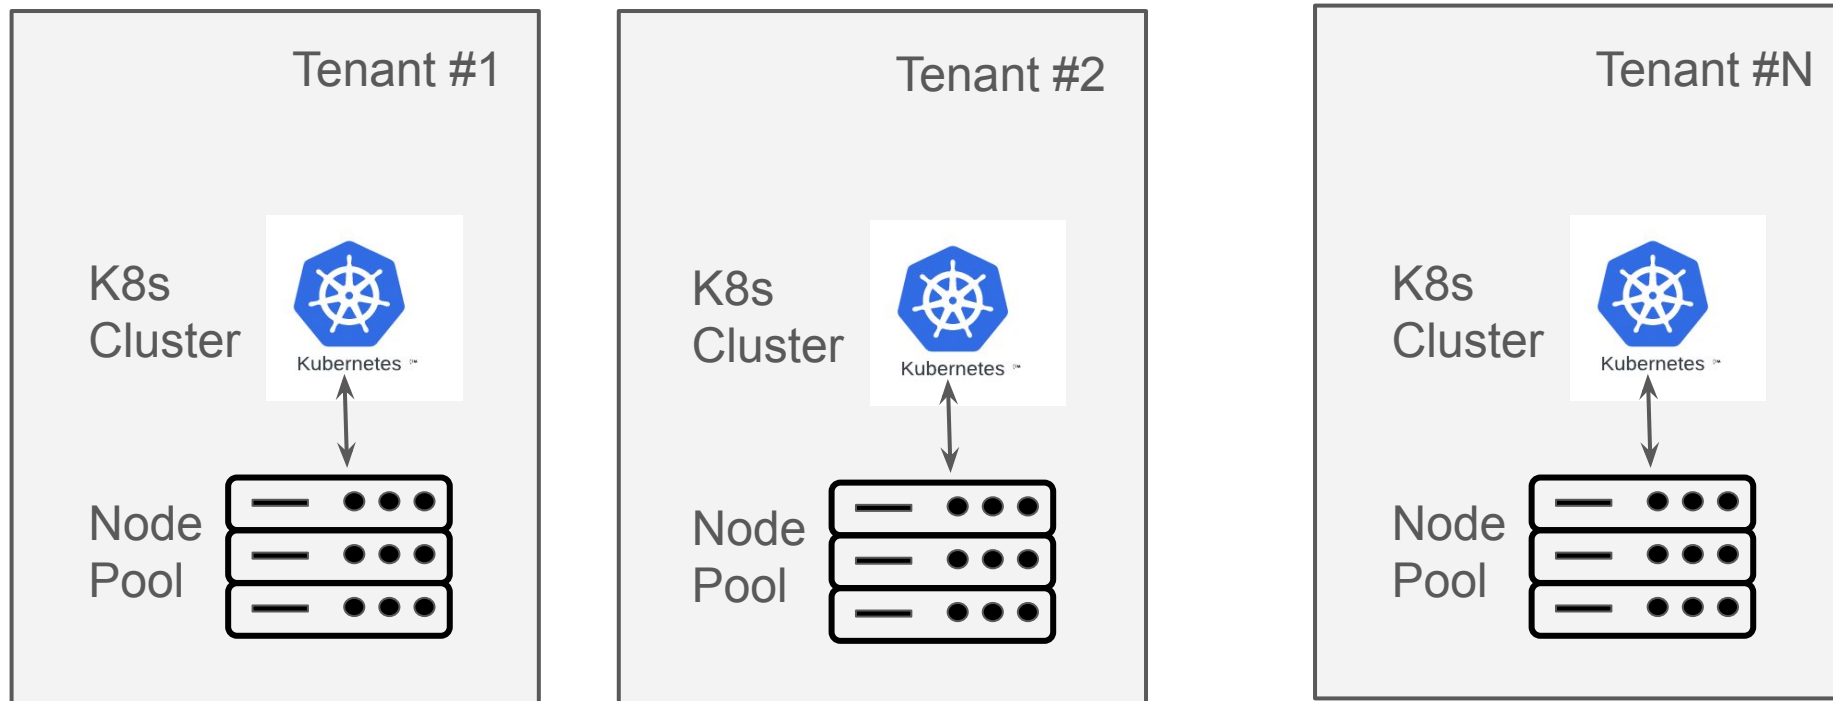
CloudNativeCon

North America 2024

One Cluster per Tenant

Solution #1: Separate Cluster per Tenant

Each Tenant gets its own Kubernetes Cluster and a Dedicated Node Pool



- Certain workloads isolated from others due to security concerns
 - Requires all 3 - data plane isolation, access isolation and control plane isolation


- Certain workloads isolated from others due to security concerns
 - Requires all 3 - data plane isolation, access isolation and control plane isolation
- Isolate noisy neighbors
 - Requires data plane isolation and control plane isolation
- Workload requiring specific hardware
 - Requires data plane isolation
- And many more

- **Manageability**

- Every cluster has a different configuration
- Managed via cluster types
- Error-prone
- Feature mismatch between tenants

- **Operational Concerns**

- Incident mitigation requires understanding cluster types



- **User Experience:**

- Operational cost high for tenants
- Every tenant operation requires a multi-step runbook
- E.g. to grow in a new zone, the tenants first should request their cluster and node pool

- **Efficiency**

- Control plane cost per cluster
- Every cluster maintains its own free pool buffer



KubeCon



CloudNativeCon

North America 2024

Multi-Tenant Single Cluster

Multi Tenant Single Cluster Architecture

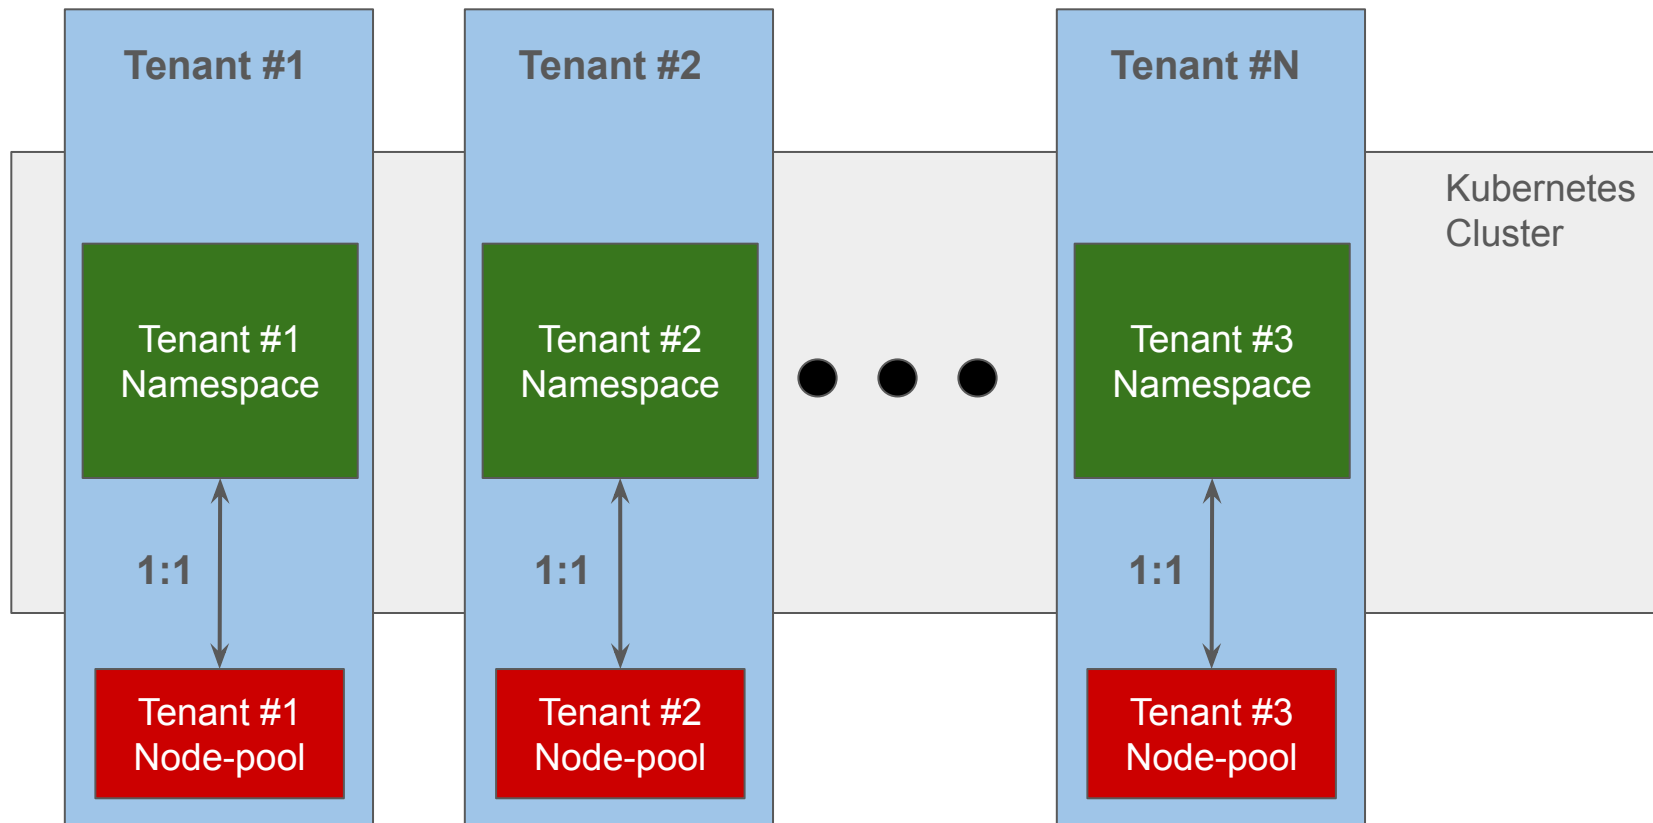


KubeCon




CloudNativeCon


North America 2024



Multi Tenant Single Cluster Architecture



- One cluster supports all tenants
- 1:1 mapping between Kubernetes namespace and Nodepool
 - Node selector added automatically
- All namespaces and nodepools in cluster created upfront
 - Always available
- Auto-scaler manages nodepool capacity
 - User merely add workloads in their namespace in the zone



#1 Access Isolation


■ Default NS Aware

■ Extended Support

Per Tenant:

RBAC's

- Roles:
 - Actions to manage resources
- RoleBindings
 - Associates tenant users to roles



Roles

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: <tenant-name>
  name: <tenant-role-name>
rules:
- apiGroups: ["", "apps"]
  resources: ["pods"]
  verbs: ["get", "list", "watch", "create", "update", "delete"]
```



Role Bindings

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: <tenant-rolebinding-name>
  namespace: <tenant-name>
subjects:
- kind: User
  name: <tenant-user-name>
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: <tenant-role-name>
  apiGroup: rbac.authorization.k8s.io
```

#2 Control Plane Isolation

APF:

- Flow Schemas that control the fair API sharing.
- Priority Settings: Concurrency shares / queues per tenant
- Default rate limits for majority of the tenants.
- Custom overrides for specific tenants.

Network Policies:

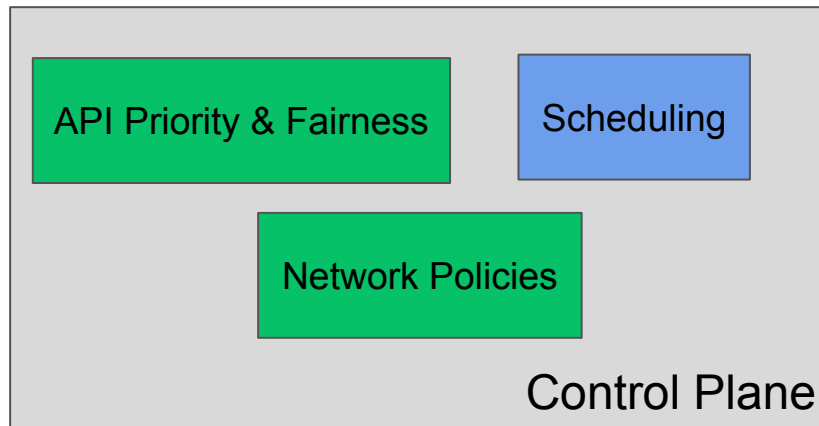
- Prevent inter tenant communication, effectively enforcing network level isolation
- Explicit ingress/egress policies to limit access to/from external endpoints

Scheduling:

- Leveraged default scheduler with extended support for node specific labels
- Actively working on options to achieve scheduler level isolation per tenant

■ Default NS Aware

■ Extended Support



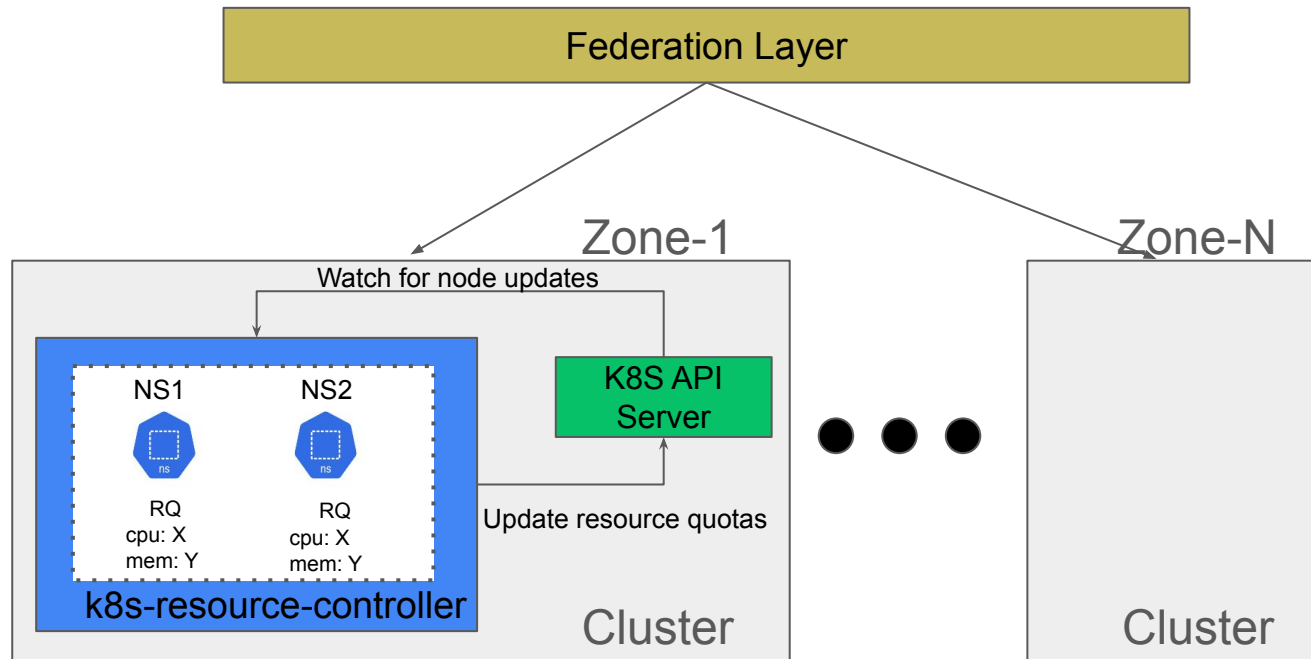
#2a Resource Quota Per Tenant

Federation layer:

- Has global view of the capacity per tenant across zones
- Picks the least loaded zones and schedules the workloads
- Contacts zonal control plane to receive up to date capacity info

Resource Controller:

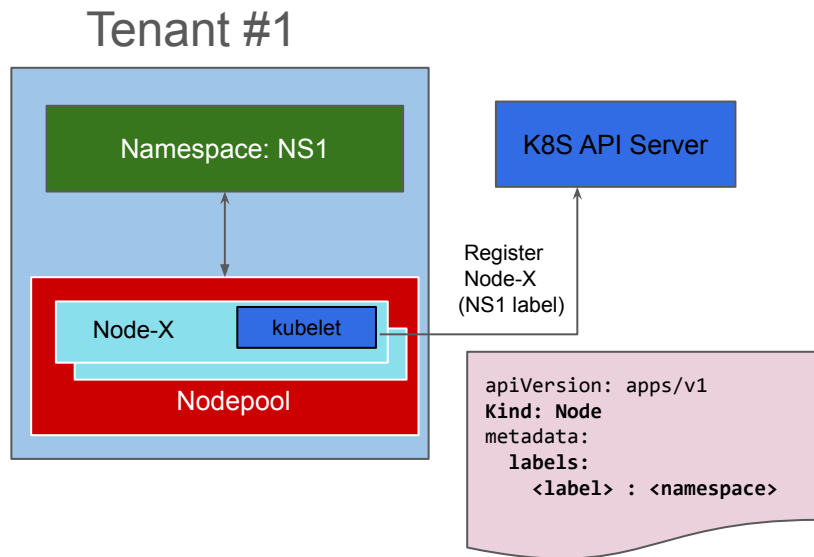
- Has zonal view of the capacity per tenant
- Aggregates resources by filtering nodes matching namespace label



#3 Data Plane Isolation

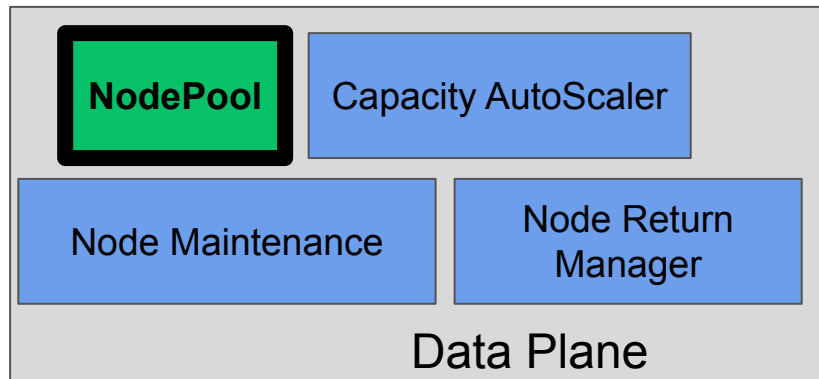
NodePools:

- Per tenant dedicated group of nodes
- 1:1 associated with tenant namespaces
- Rely on node labels to assign nodes to a nodepool



■ Default NS Aware

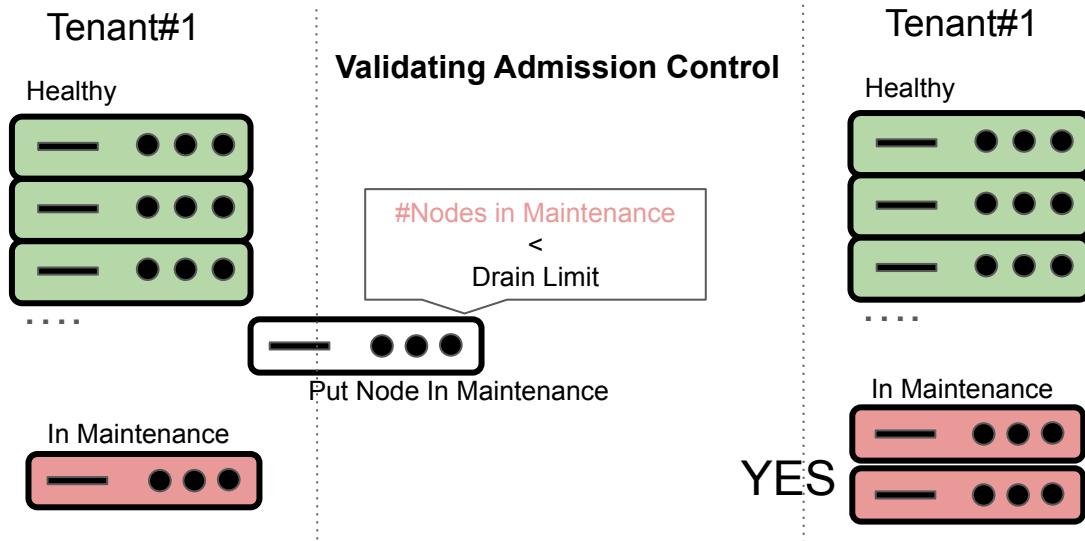

■ Extended Support



#3 Data Plane Isolation (cont'd)

Node Maintenance:

- Perform operational activities (upgrades, etc..) to keep secure, stable, and performant nodes in the cluster.
- Drain Limit: Max Concurrent nodes in maintenance
- Admission controller plugin for validating nodes in maintenance per tenant



#3 Data Plane Isolation (cont'd)

Node Return Manager:

- Evaluator library that calculates number of nodes to be returned safely
- **Safety Threshold:** Max allocation % beyond which workloads become unsafe
- Return criteria policies:
 - Allocation % of the tenant
 - Pod Topology Spread, Failure Domain, etc..


Receives event for processing node returns



Tenant:
Nodes to return < Safety Threshold

■ Default NS Aware

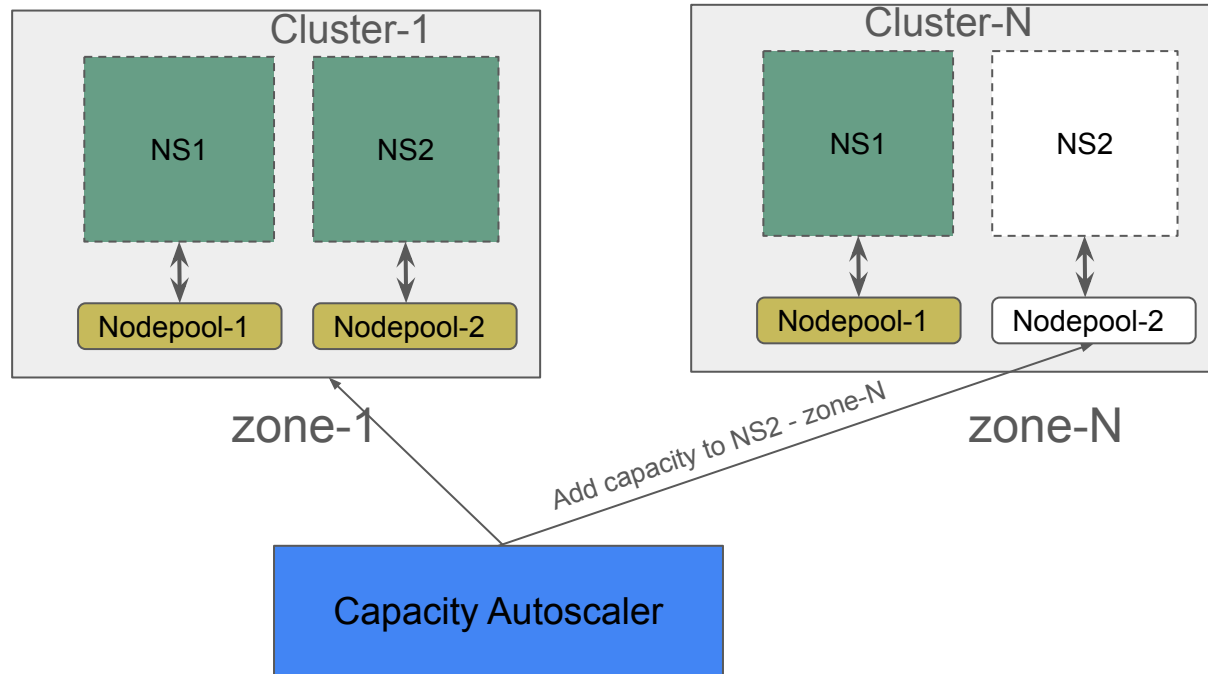
■ Extended Support



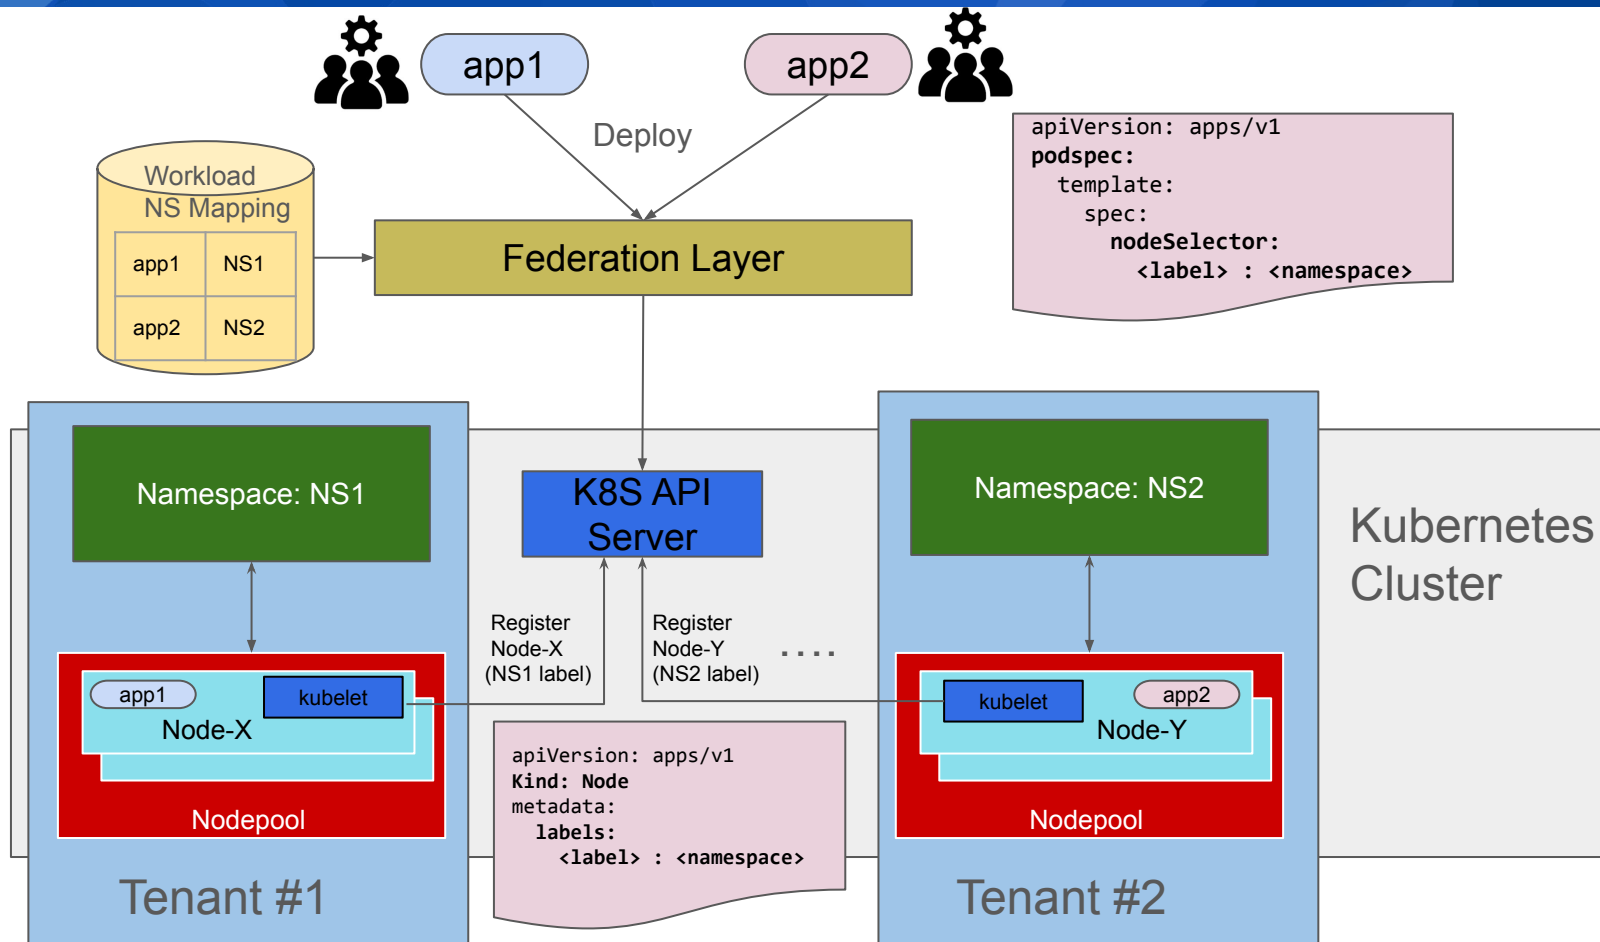
#3 Data Plane Isolation (cont'd)

Capacity AutoScaler:

- Manages the capacity across nodepools / zones.
- **+** capacity when there are unscheduled pods.
- **-** capacity when the allocation drops.
- Improved scheduling latency with hot standby buffer pool
- Inter nodepool swaps are fairly trivial with merely node label swaps



Workload Deploy



Migration Status

- Tens of namespaces 100% migrated to new architecture
- 30% reduction in the number of clusters needed globally
- End Target - 100% by 2025



KubeCon



CloudNativeCon

North America 2024

What Worked & Challenges

What worked vs Challenges



- Operational ease in managing configs
 - Single cluster per zone
- NodeSelector scalability
 - Tens of namespaces live in production.
- Seamless capacity management
- Inter tenant capacity swaps are trivial with label change.
- Reduced control plane cost and shared free pool buffer
- Priority queues and flow control
- No loss in debuggability experience with objects being namespaced



- No native support for nodes to namespaces. Custom controllers to
 - ensure the correctness of the bindings
 - aggregate resource quota
 - drain capacity
- Scheduler is not isolated
- Overhead in managing bootstrap cluster

Problem-1: Scheduler Isolation

Problem:

No native support for scheduler level isolation predominantly

- No guarantees around scheduling latency per tenant
- No per tenant pod queues
- No native support for matching pods to tenant nodes

Options:

- Active - Active Scheduler :
 - Each scheduler partition to operate on a specific tenant nodepool
 - Actively exploring this option
- Per Tenant Queues:
 - Configure separate scheduling queues per tenant
- Open to new ideas

Problem-2: Bootstrap Cluster

Bootstrap Cluster:

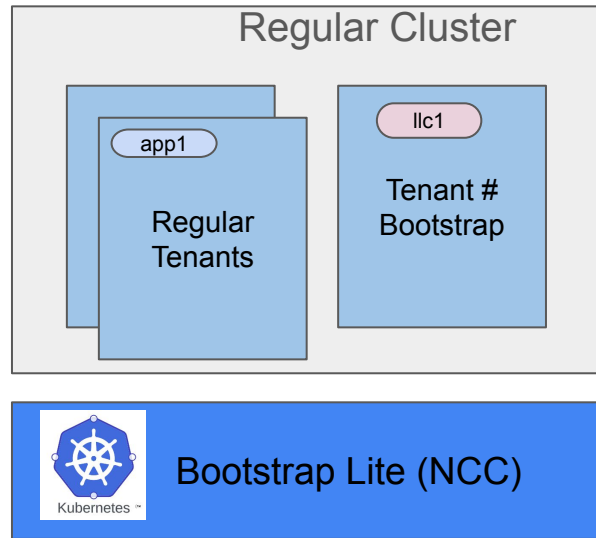
Brings up the low level infrastructure components necessary for the regular cluster to operate.

Components:

- Native Compute Components (NCC)
 - ApiServer
 - Scheduler
 - Controller Manager
- Other low level infra components (LLC):
 - Custom Controllers
 - Metrics infra
 - Logging infra
 -

Problem: Need to manage an additional bootstrap cluster zonally

Proposal



Acknowledgements

- Container Platform
- Service Lifecycle Team
- Host Lifecycle Team
- Security Team
- Observability Team



KubeCon



CloudNativeCon

North America 2024

Q & A