



KubeCon



CloudNativeCon

North America 2024

# Still Don't Do What Charlie Don't Does: Making CRD Changes safer

*Nick Young, @youngnick, Isovalent at Cisco*



@youngnick.net

# Who am I to talk about this?

- Started looking into CRDs in early 2017, when they were called Third-Party Resources or TPRs
- Was involved in building out Contour's **HTTPProxy** resource, that replaced **IngressRoute**
- Have been involved in Gateway API since its inception in 2018 at Kubecon San Diego

# Today's agenda

- Explain some important context about how Kubernetes stores objects and versioning
- Walk through some CRD Change Antipatterns, using “**ChaRlie Don’t**” as our straw-man
- Give you some tips for each on what to do to avoid them



KubeCon




CloudNativeCon

North America 2024

# Why Charlie Don't?

# With thanks to the Simpsons



Images from Finkiac: <https://frinkiac.com/>  
Simpsons Episode 1F06, Season 5, Episode 8, Boy Scoutz 'n the Hood

# Meet Charlie Don't

- Charlie works on a custom controller for Kubernetes at BigCo
- He has the worst luck and always manages to choose the wrong design option
- Poor Charlie!



# Previously on Charlie Don't

- Read the API bibles (API Conventions and API changes docs)
- Think about how your users will use the CRD
  - Use `status` and `status.Conditions`!
- Make as many fields as possible optional, with defaults if it makes sense
- Avoid maps except for labels and annotations. Use `listType=map` instead.
- Avoid `bool` types and bounded enums
- Avoid cross-namespace references, make them need a handshake if you do use them
- Don't make breaking API changes without an API version bump

# Why are API changes so important?

- No matter what you've built your API for, there are a few invariants:
  - You'll *always* need to make some changes to the API
  - There's no such thing as temporary
  - Except for "no". Saying "yes" to a feature affects you approximately forever, saying "no" to a feature is temporary.





KubeCon




CloudNativeCon

North America 2024

Before we hear more about Charlie though...

# The main reference!

The API Changes community documentation is the main reference for this talk.



# Let's talk about Kubernetes object versioning

- Every object has:
  - A **Group**: a domain-like string that identifies the a set of resources - usually an owning group or similar. For example: `gateway.networking.k8s.io` and `cilium.io`
  - A **Kind** and **Resource**: a Kind identifies an object Schema (like `Gateway` or `Node`), and a Resource identifies a resource in a unique way, used via HTTP (like `gateways` or `nodes`)
  - A **Version**: identifies a unique *version* of the resource, looks like `v1`, `v1beta1`, or `v2alpha1`.
- The **Group** and **Version** are often combined into the `apiVersion` field
  - Takes the form `<group>/<version>`, and looks like:
    - `gateway.networking.k8s.io/v1`
    - [`cilium.io/v1alpha2`](#)

# Let's talk about Kubernetes object versioning

- Versioning is Really Important.
- Kubernetes has three classes of *object* stability:
  - **alpha** - objects are experimental and subject to large, breaking changes between versions. No guarantees around conversion either.
  - **beta** - objects are experimental, but more stable than **alpha**. Generally will not have breaking changes between versions, and should include conversion if they do.
  - **stable** - indicated by **v1** or similar, this means that *no* breaking changes will *ever* be made to this API version. *Additive, safe* changes *may* be made. (More on what these are later)
- Making a breaking API change means incrementing the API version!
  - Going from **v1alpha1** to **v1alpha2**, or **v1beta2** to **v1beta3**, etc
  - Making a breaking change to a **v1** object requires a new **v2alpha1** or similar

# Every object also has a **Storage** version

- The **storage** version sets the version of the object's schema that's to be persisted to storage (etcd).
- An object *may* have multiple versions available at any time.
- The apiserver converts between the versions when required.
- When are conversions required?
  - When there are incompatible changes between versions, *and*
  - When a consumer *reads* a version that is not the storage version, *or*
  - When something *writes* to an object that is stored as a version that's *not* the storage version - more on this in a minute
- In-tree resources have all required conversions included in the apiserver code
- CRDs must supply a webhook to handle changes between versions

# Conversion Example

The Widget object on the right is version `v1alpha2`.

It has a `bar` field for storing the number of `bars` that a widget provides.

```
apiVersion: foomake.io/v1alpha2
Kind: Widget
metadata:
  name: testWidget
  namespace: widgethome
spec:
  bar: 20
```

# Conversion Example

Now we decide that we want to support multiple numbers of **bars** with a single **Widget**, so we *pluralize* the **bar** field into **bars**.

Note that we updated the Version as well.

```
apiVersion: foomake.io/v1alpha3
Kind: Widget
metadata:
  name: testWidget
  namespace: widgethome
spec:
  bars:
    - 10
    - 20
```

# Conversion Example

This is not a *compatible* change, so a conversion would be required, which would take the value from the **bar** field, and insert it into the first entry of the **bars** field.

This change is not *compatible* because you can't take a **v1alpha3** version and turn it back into a **v1alpha2** version without possibly losing information.

```
apiVersion: foomake.io/v1alpha3
Kind: Widget
metadata:
  name: testWidget
  namespace: widgethome
spec:
  bars:
    - 10
    - 20
```



# Conversion Example

But, you could write conversion code to safely convert all `v1alpha2` objects to `v1alpha3` objects.

```
apiVersion: foomake.io/v1alpha3
Kind: Widget
metadata:
  name: testWidget
  namespace: widgethome
spec:
  bars:
    - 10
    - 20
```

# You can skip conversions though!

- We can skip having to convert values by making sure changes are *backwards compatible*.
- What makes a backwards compatible change?
- Let's check in with Charlie!

## Charlie Don't makes breaking API changes in the same **version**

- Charlie doesn't understand the rules for API changes.
- Charlie makes a change that's not backwards compatible and doesn't increment his API version
- Users of his CRDs are broken when they upgrade!
- Poor Charlie *and* his users!





KubeCon



CloudNativeCon

North America 2024

So what *can* we do?



KubeCon



CloudNativeCon

North America 2024

# It's complicated!

# The Simplified Backwards Compatibility Rule

If you can take a new object, with only *required* fields set, change the version to the old version, and apply it, and *nothing behaves differently*, then the changes between the versions are **backwards compatible**.

# Let's head back to the real world now though

By that previous definition, almost any change is not backwards compatible!

However, we can make some changes safely, *if and only if*:

- They are *additive*
- We use a feature flag mechanism to control processing of those fields

*Additive* changes are a special class of changes that may break in *limited, understandable* ways, so they *can* be okay to make within the same API version.





KubeCon



CloudNativeCon

North America 2024

# Let's run through some examples

# Charlie Don't adds a new field that is **+required**

- Adding a new field that's *required* is always a breaking change, because it means that the new version can't be serialized to the old version without changing behavior.



# How do we add new fields?

The best way to do this is to make the new field *optional*, and have the default value for the optional field mean “the same behavior as the old version”.

For default values, you can either use the *zero value*, or you can set a default value in the CRD definition.

In this case, if no value is supplied, then the default value will be used instead.

# Adding new fields - scalar fields

- Scalar fields are individual values, like `string` or `int`.
- If you need to tell the difference between “set to a value”, “set to the zero value”, and “unset”, use a pointer!
- Examples
  - Adding a new, optional `newBehavior` string field that defaults to the empty string, which means “don’t do the new behavior”.
  - Adding a new, optional `awesomenessLevel` string field that defaults to `none` (which is how you’re describing the old behavior).
  - Adding a new `timeoutSeconds` field where `0` means “unlimited”, unset means “use the default”, and a value means to use that specific value.

## Charlie Don't uses Enum fields without declaring them open for changes

- An enumerated field is a `string` field that has a set of permitted values - usually supplied in Go types as aliased constants.
- Charlie Don't doesn't make it clear in the field definition that he might add fields later.
- Implementations don't know they need to handle unknown values.



The answer here is straightforward:

- Ensure that your enumerated strings are *always* documented as being open for changes
- You must also document what happens when there's an invalid value!

# Charlie Don't adds values to `bool` fields

- A `bool` field can only ever be `true` or `false`.
- Adding other values to a `bool` field requires changing it to a `string`
- Type changes for fields are breaking changes.
- Don't use `bool` fields!



# What to use instead of `bool` fields?

Use enumerated string fields instead (but make sure you declare them open for changes!)

Don't do:

```
enableAwesomeness: true|false
```

Do:

```
awesomenessLevel: extreme|some|none
```



# Charlie Don't adds **struct** fields that aren't pointers

- For struct fields to be optional, they *must* be pointers.
- If you add a struct field that's not optional, that's a breaking change.
- A good way to have easily expandable struct fields is to use a *union*.



The Widget object on the right is version `v1alpha2`.

It has a `union` setup, where you choose a `type`, and then have a separate config struct per type.

Importantly, we've marked the `type` field as “can be expanded”, since it's an enum field.

But, in this case we're only using a `round` type.

```
apiVersion: foomake.io/v1alpha2
Kind: Widget
metadata:
  name: testWidget
  namespace: widgethome
spec:
  type: round
  roundConfig:
    radius: 10
```

Here, we've added a new value to the `type` field, and an accompanying config struct.

This *can* be an additive change IF:

- the `type` enum allows expansion
- the `squareConfig` struct is added as a pointer, so that it defaults to unset
- The contract for the Widget object includes that invalid values for `type` are processed but recorded as an error in the `status`.

```
apiVersion: foomake.io/v1alpha2
Kind: Widget
metadata:
  name: testWidget
  namespace: widgethome
spec:
  type: square
  roundConfig:
    radius: 10
  squareConfig:
    sideLength: 10
```

## Charlie Don't makes validation rules more strict between versions

- This is backwards incompatible because values that used to be valid are not valid any longer.
- *Loosening* validation can be okay, if you're careful.
- Remember that once you loosen, you can never go back!





KubeCon



CloudNativeCon

North America 2024

# Summing up

# What did we learn from Charlie this time?

- Versioning is important.
- If it can, the apiserver will convert between versions for you. Otherwise, you need a conversion webhook.
- Changes don't need conversion if they are compatible.
- Some ways to make changes compatible:
  - Add new fields as `+optional`
  - When adding an enumerated string field, ensure that you document that values may be added
  - Don't use `bool` fields! Use enumerated strings instead.
  - Make `struct` fields optional properly by making the field be `*struct`
  - Only loosen field validation between versions

**Charlie Don't and I both say**

**Thanks for Listening!**

