

Strimzi:

Data Streaming on Kubernetes with Apache Kafka

Jakub Scholz

Yaodong Yang

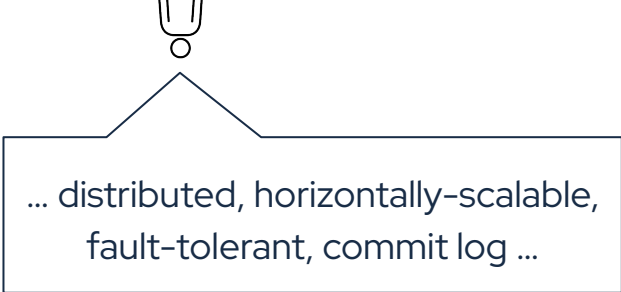


What is Strimzi

- Open Source project / community (Apache License 2.0)
- CNCF Incubating project
- Focuses on Apache Kafka on Kubernetes
 - Based on the operator pattern
 - Provides operators for running and managing Apache Kafka and its components
 - Additional tools to make Apache Kafka easier to use on Kubernetes



Ecosystem



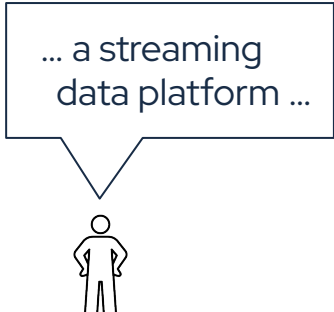
... distributed, horizontally-scalable,
fault-tolerant, commit log ...

Apache Kafka

- Leading distributed message log and data streaming platform
- Open Source project
 - Originally developed by LinkedIn
 - Part of the Apache Software Foundation
 - Licensed under Apache License 2.0



... a publish/subscribe
messaging system ...



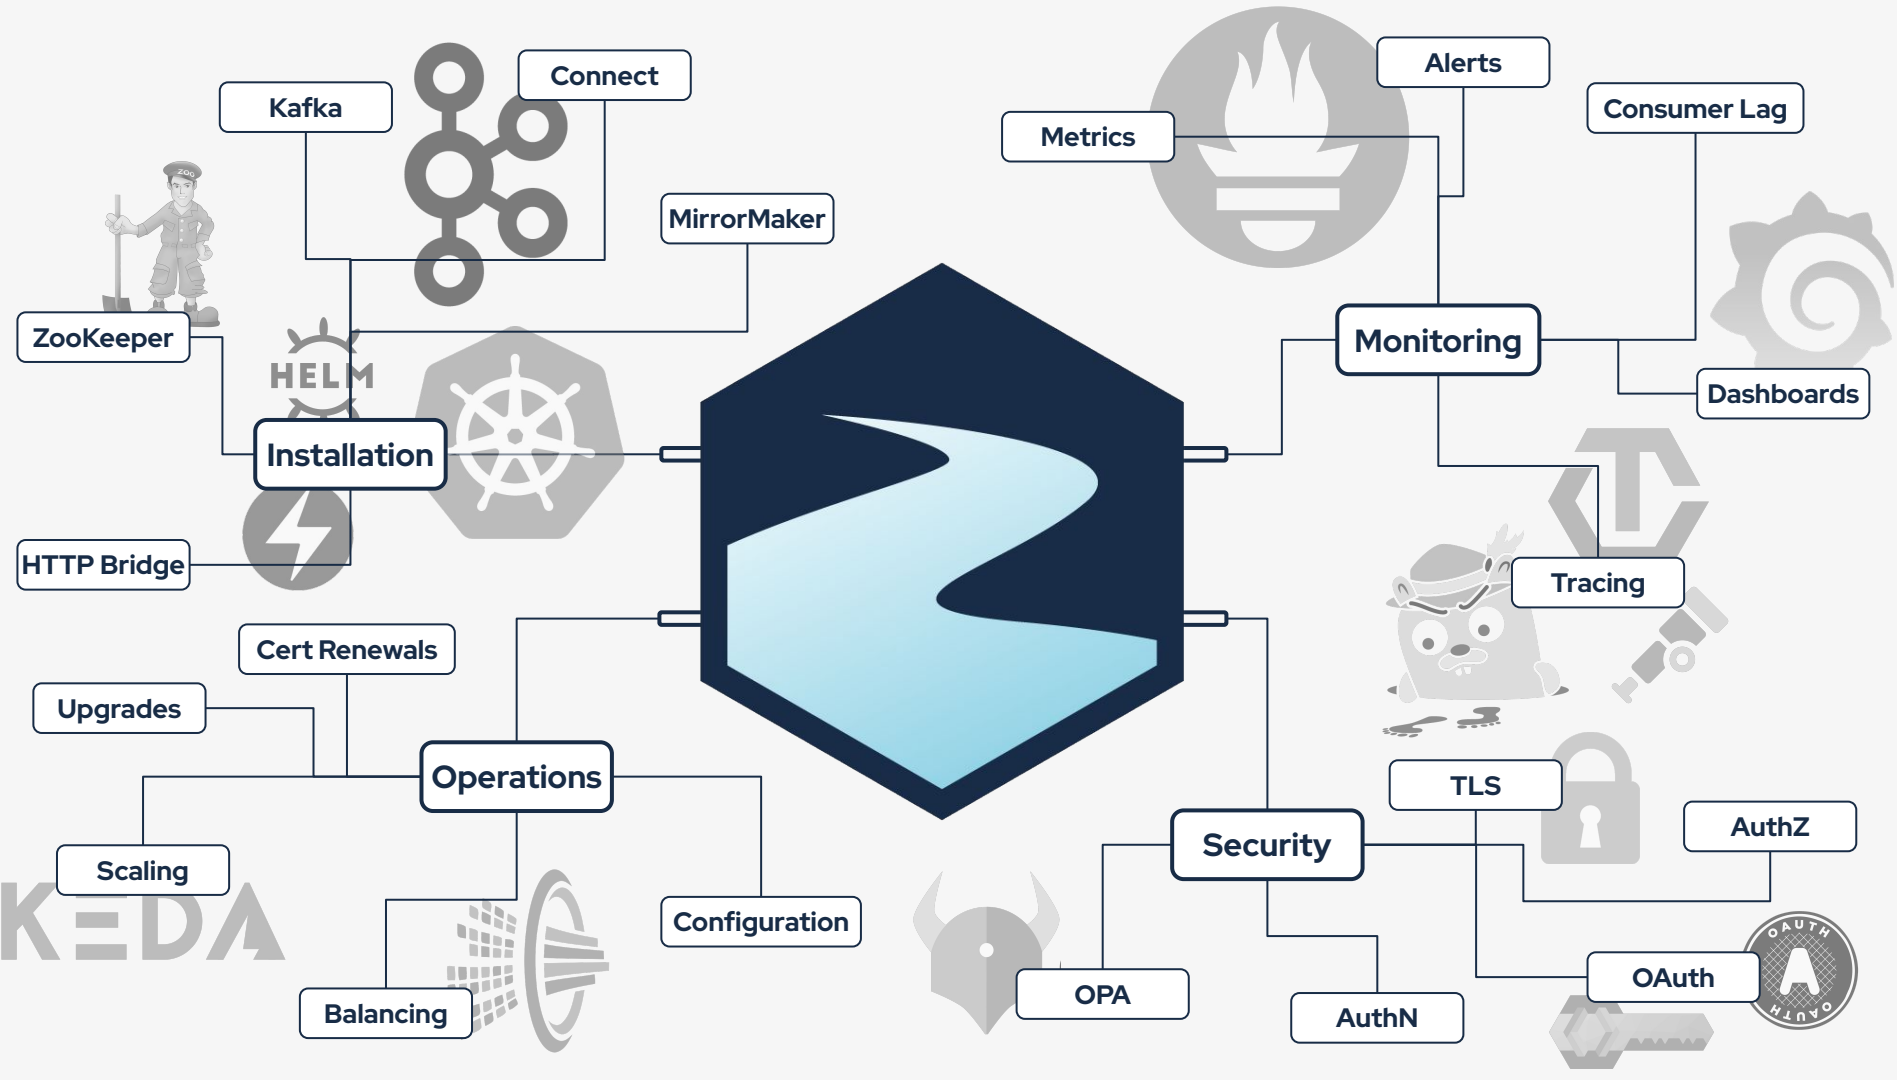
... a streaming
data platform ...

Kubernetes-native Kafka

- Follows the operator pattern
 - Uses CRDs to extend Kubernetes API and define the Kafka resources
 - Encodes the Kafka operations knowledge into the operator code
 - Automates the installation, operations and monitoring tasks
- Examples where Strimzi operators help most
 - Handling the Kafka discovery protocol, Upgrades, Security, Scaling

Kubernetes-native Kafka

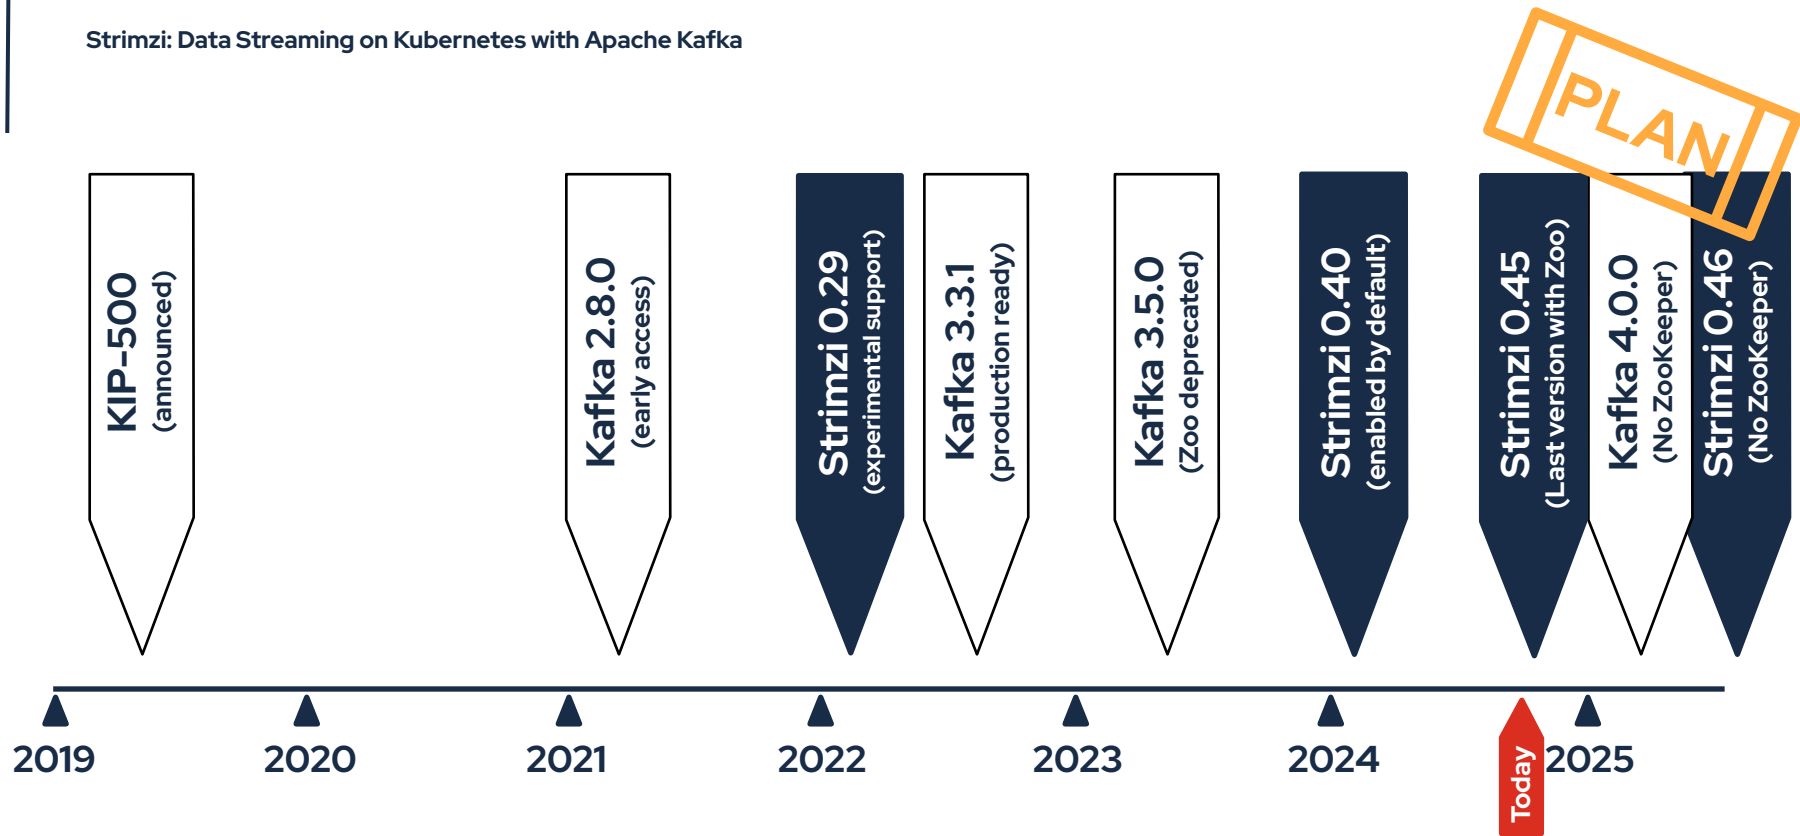
- Support for all Apache Kafka components & more
 - ZooKeeper, Kafka brokers, Kafka Connect, Kafka Mirror Maker (1 and 2)
 - Our own HTTP Bridge
 - Cruise Control for *balancing* the Kafka cluster
 - The components can be used independently or together
- Makes it easier to use Kafka
 - Manages topics, users or connectors through the operator pattern as well







KRaft

KRaft

- Kafka is removing its dependency on ZooKeeper
 - Replaced with Kafka's own implementation based on the Raft protocol
 - Changes the APIs and architecture of the Kafka clusters
 - Lot of effort is dedicated to supporting the new KRaft mode and a smooth transition
 - Coming to an end after 5 years!







Timeline

- Strimzi 0.45.0 (next release, end of 2024)
 - Supports Apache Kafka 3.8.0 and 3.9.0
 - **Last version with ZooKeeper support**
 - We plan to provide “extended support”
 - **ZooKeeper-based Kafka clusters need to be migrated to KRaft**

Timeline

- Strimzi 0.46.0 (early 2025)
 - Supports Apache Kafka 3.9.0 and 4.0.0
 - **Only KRaft-based Kafka clusters supported**
 - Mirror Maker 1 support will be removed as well
 - **ZooKeeper-based Kafka clusters need to be migrated to KRaft before upgrading to Strimzi 0.46.0**

Migration

- Existing ZooKeeper based clusters can be migrated to KRaft
 - Migration to KRaft is driven by users through annotations
 - Cannot be fully automatic because of new configurations, architectures, ...
 - Migration has to happen before upgrading to Strimzi 0.46.0
- KRaft and migration are already supported in Strimzi => you can use it today!

Learn more about KRaft

<https://strimzi.io/kraft/>



Title

- Point one
 - Subpoint
- Point two
 - Subpoint

Tiered Storage

Tiered Storage

- Move data from expensive local disk to remote storage
 - Significantly reduce storage cost
 - Broader use cases adoption
 - Better scalability and reliability
 - Reduced operation overhead

Tiered Storage


- Native support for tiered storage configuration
- Defaulting to use TopicBasedRLMM class for metadata
- Custom Plugin library required for RLM class implementation

Tiered Storage

```
kafka:
  tieredStorage:
    type: custom
    remoteStorageManager:
      className: com.example.kafka.tiered.storage.s3.S3RemoteStorageManager
      classPath: /opt/kafka/plugins/tiered-storage-s3/*
      config:
        # A map with String keys and String values. Key fields will be automatically
        # prefixed with `rsm.config.` and appended to Kafka broker config.
        storage.bucket.name: my-bucket
    config:
      ...
      rlmm.config.remote.log.metadata.topic.replication.factor: 1
```

Tiered Storage

- Limitation
 - Compacted topics not supported
 - Disabling tiered storage supported only in KRaft mode
- Plugins options
 - Aiven's open source plugins: <https://github.com/Aiven-Open/tiered-storage-for-apache-kafka>
 - Amazon S3, Google Cloud Storage, Azure Blob Storage and file storage
 - Custom implementation



Auto-rebalancing

Auto-rebalancing

- Automatically triggers a rebalance when brokers are added or removed
 - Based on Cruise Control integration
 - Makes it easier to scale the Kafka cluster up or down
 - Moves partition replicas to new brokers after they were added
 - Moves partition replicas from brokers before they are removed

Future plans

Future plans

- **v1 CRD APIs and Strimzi 1.0.0**
- Better support for external certificate management tools
- Kafka cluster self-healing
- Built-in Gateway API support for exposing Kafka clusters
- Kafka clusters stretched across multiple Kubernetes clusters
- and more ...

v1 APIs and Strimzi 1.0.0

- Once ZooKeeper is removed
 - We will start working on the new v1 API (CRDs)
 - The main change is expected to be restructuring after ZooKeeper removal
- Strimzi 1.0.0 should follow after the v1 API is rolled out

Better CA abstraction

- Better integration with external providers of TLS certificates
 - Pluggable architecture which can be more easily extended
 - Better support for common tools such as Cert Manager
 - Makes it easier to integrate into user's own certificate management processes and systems
- Possibly disabling TLS completely?

Kafka Cluster self-healing

- Strimzi already integrates with Cruise Control for cluster rebalancing
 - Users interact with it through the `KafkaRebalance` resources
- We plan to integrate Cruise Control self-healing feature as well
 - Uses anomaly detection to find out about partition, disk, or broker failures
 - Where possible uses self-healing to fix the problem

Gateway API support

- Gateway API aims at replacing Ingress
 - Should provide better compatibility between providers
- Strimzi-based Kafka clusters can be already exposed using a Gateway API
 - Requires manual configuration of the Gateway API resources, advertised hosts etc.
- We plan to provide built-in support for Gateway API to make it easier to use!

Stretch clusters

- Different parts of the Kafka cluster running on different Kubernetes clusters
 - Kafka is latency sensitive, so running brokers on different continents might never be possible
 - But should work fine with co-located Kubernetes cluster
 - Makes it easier to migrate between clusters
 - Suitable for *Metropolitan Area Networks / Metro DR*

Join us

<https://strimzi.io/join-us/>



Other talks

Other talks

- Strimzi and the future of Apache Kafka on Kubernetes
 - Project lightning talks, Tue 12th November
- Elastic Data Streaming: Autoscaling Apache Kafka
 - Data Processing + Storage, Thu 14th November



<https://strimzi.io/>



<https://github.com/strimzi>



@strimziio



<https://youtube.com/c/Strimzi>



<https://www.linkedin.com/company/strimzi>

Demo

- External access
- Broker re-configuration
- Cluster rebalancing